

# Video


Lecture 9

<http://www.cs.rutgers.edu/~sn624/352-F24>

Srinivas Narayana

# Digital representation of audio and video


# Digital representation of audio



- Must convert analog signal to digital representation
- Sample
  - How many times (twice the max frequency in the signal)
- Quantize
  - How many levels or bits to represent each sample
  - More levels → more accurate representation of signal
  - More levels → more bits to store & need more bandwidth to transmit
- Compress
  - Compact representation of quantized values


# Audio representation

- analog audio signal sampled at constant rate
  - telephone: 8,000 samples/sec
  - CD music: 44,100 samples/sec
- each sample quantized, i.e., rounded
  - e.g.,  $2^8=256$  possible quantized values
  - each quantized value represented by bits, e.g., 8 bits for 256 values



# Audio representation

- example: 8,000 samples/sec, 256 quantized values
- Bandwidth needed: 64,000 bps
- receiver converts bits back to analog signal:
  - some quality reduction



## Example rates


- CD: 1.411 Mbps
- MP3: 96, 128, 160 Kbps
- Internet telephony: 5.3 Kbps and up

# Video representation

- Video: sequence of images displayed at constant rate
  - e.g., 30 images/sec
  - Appear continuous due to the stroboscopic effect



frame  $i$




frame  $i+1$


# Video representation

- Digital image: array of pixels
  - each pixel represented by bits
  - Encode luminance and color
  - Number of pixels: **resolution**
- Coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)
- Encoding/decoding algorithm often called a **codec**

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$



frame  $i+1$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$  (motion vectors)

# Video codecs: terminology

- **Video bit rate:** effective number of bits per second of the video after encoding
- It depends on many factors
  - Resolution of each image: more pixels = more bits
  - Detail per pixel: more luminance & color detail = more bits
  - Amount of movement in the video. More movement = more bits
  - Quality of overall compression in the codec
- Video bit rate is typically correlated with quality of perception
  - Higher bit rate == better to perceive



# Bit-rates: terminology

- Bit-rate of a video changes over the duration of the video
- **CBR: (constant bit rate):** fixed bit-rate video
- **VBR: (variable bit rate):** different parts of the video have different bit rates, e.g., changes in color, motion, etc.
  - For VBR, we talk about **average bit-rate** over video's duration
- **Examples of average video bit-rates**
  - MPEG 1 (CD-ROM) 1.5 Mbps. MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)
  - In general, one Internet video stream takes up a few Mbit/s (more for HD)

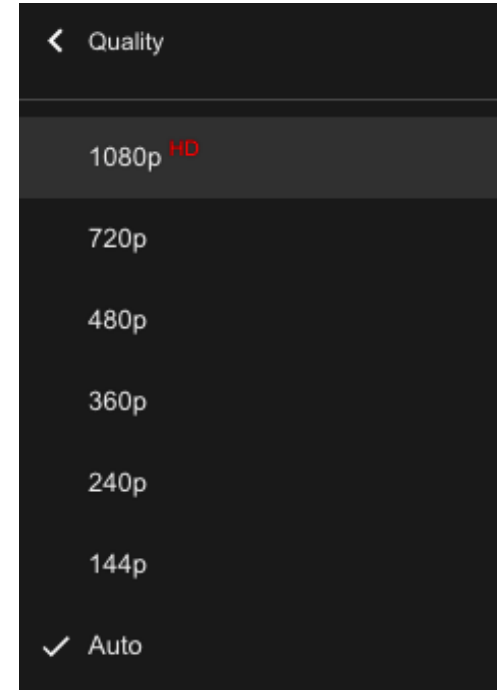
# Networking multimedia: 3 types

- **On-demand streamed video/audio**
  - Can begin playout before downloading the entire file
  - Full video/audio stored at the server: able to transmit faster than audio/video will be rendered (with storing/buffering at client)
  - e.g., Spotify, YouTube, Netflix
- **Conversational** voice or video over IP
  - interactive human-to-human communication limits delay tolerance
  - e.g., Zoom
- **Live streamed** audio, video
  - e.g., sporting event on sky sports
  - Can delay a little, but must be close to the “live edge” of content

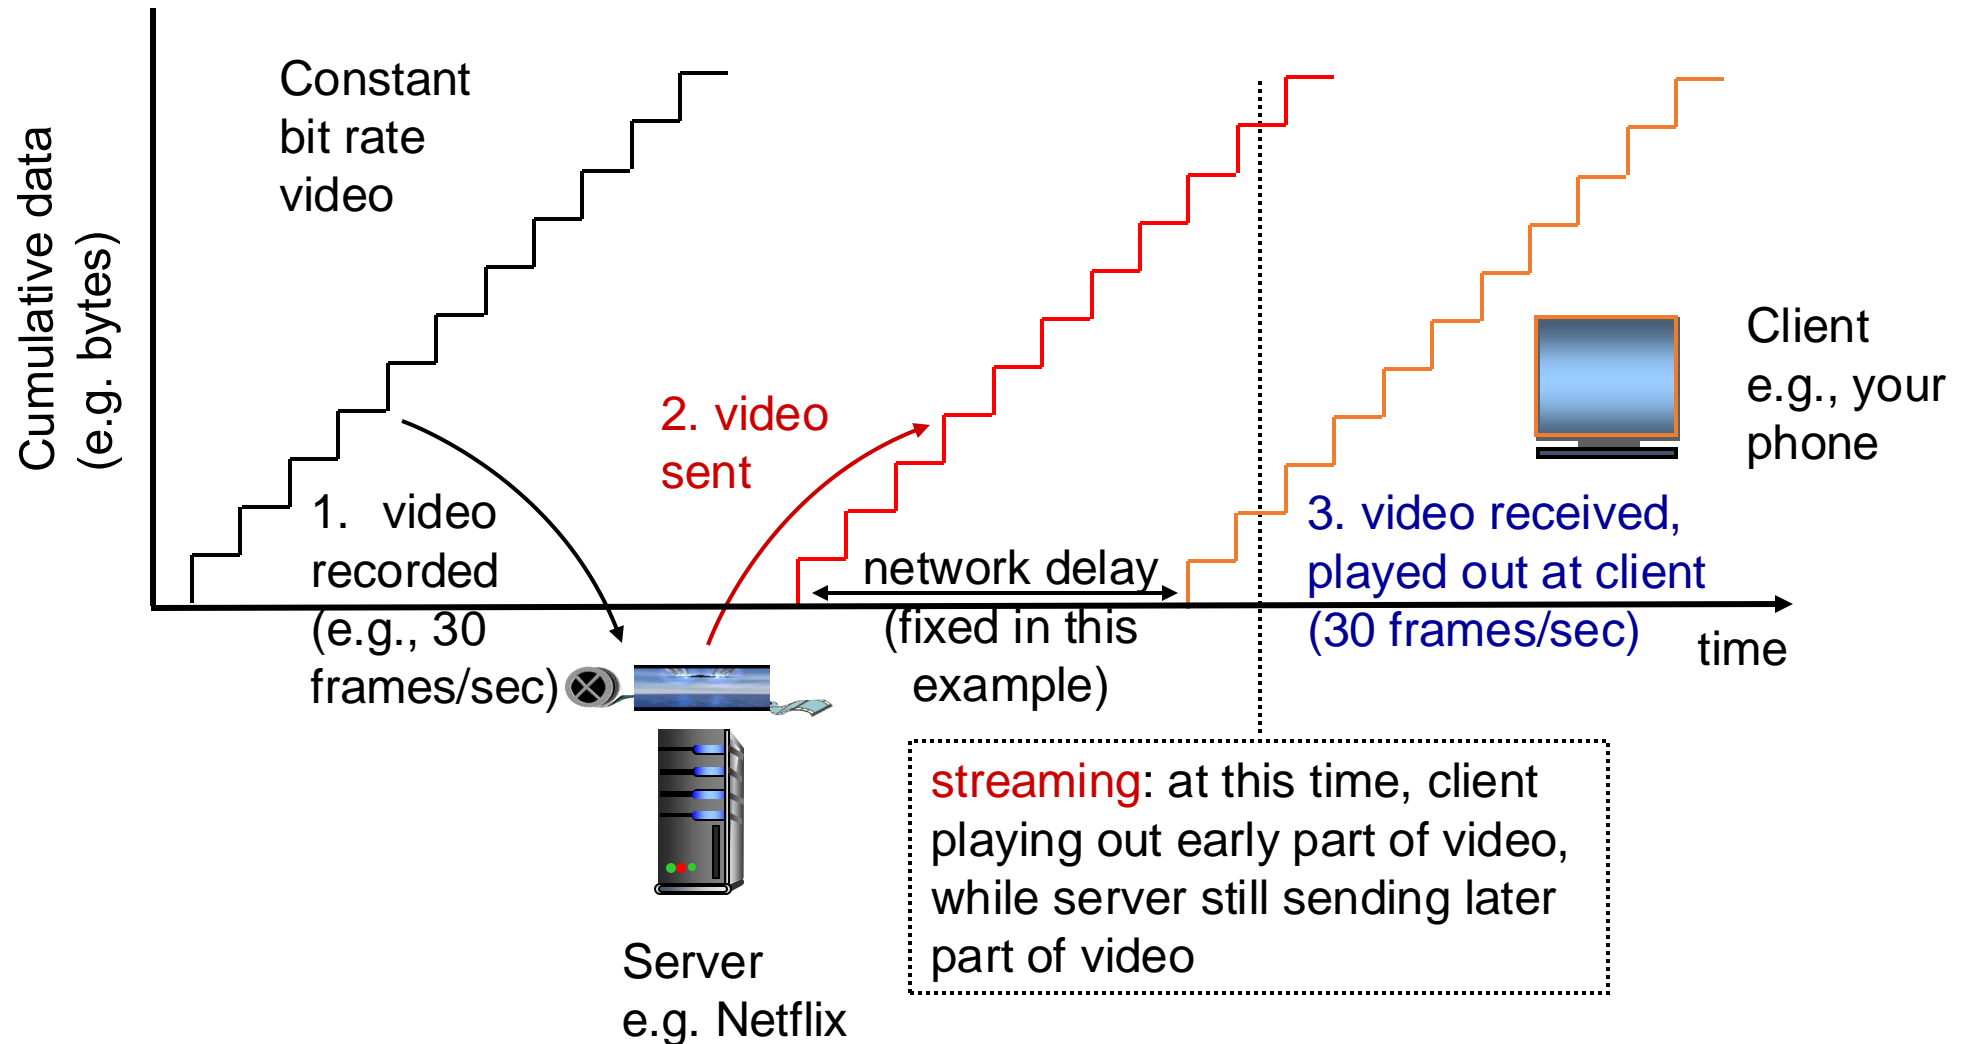
# On-demand Video Streaming

# Streaming (stored) video

- Media is prerecorded at different qualities
  - Available in storage at the server
- Client downloads an initial portion and starts viewing
  - The rest is downloaded as time progresses
  - No need for user to wait for entire content to be downloaded (**streaming**)
- Can change the quality of the content and where it's fetched mid-stream




# Streaming stored video




# Streaming stored video: challenges

- **Continuous playout constraint**: once video playout begins at client, time gap between frames must match the original time gap in the video (why?)
- But **network delays are variable!**
- Clients have a **client-side buffer** of downloaded video to absorb variation in network delay, available bandwidth
- The video buffer also helps with user interactions: pause, fast-forward, rewind, jump through video

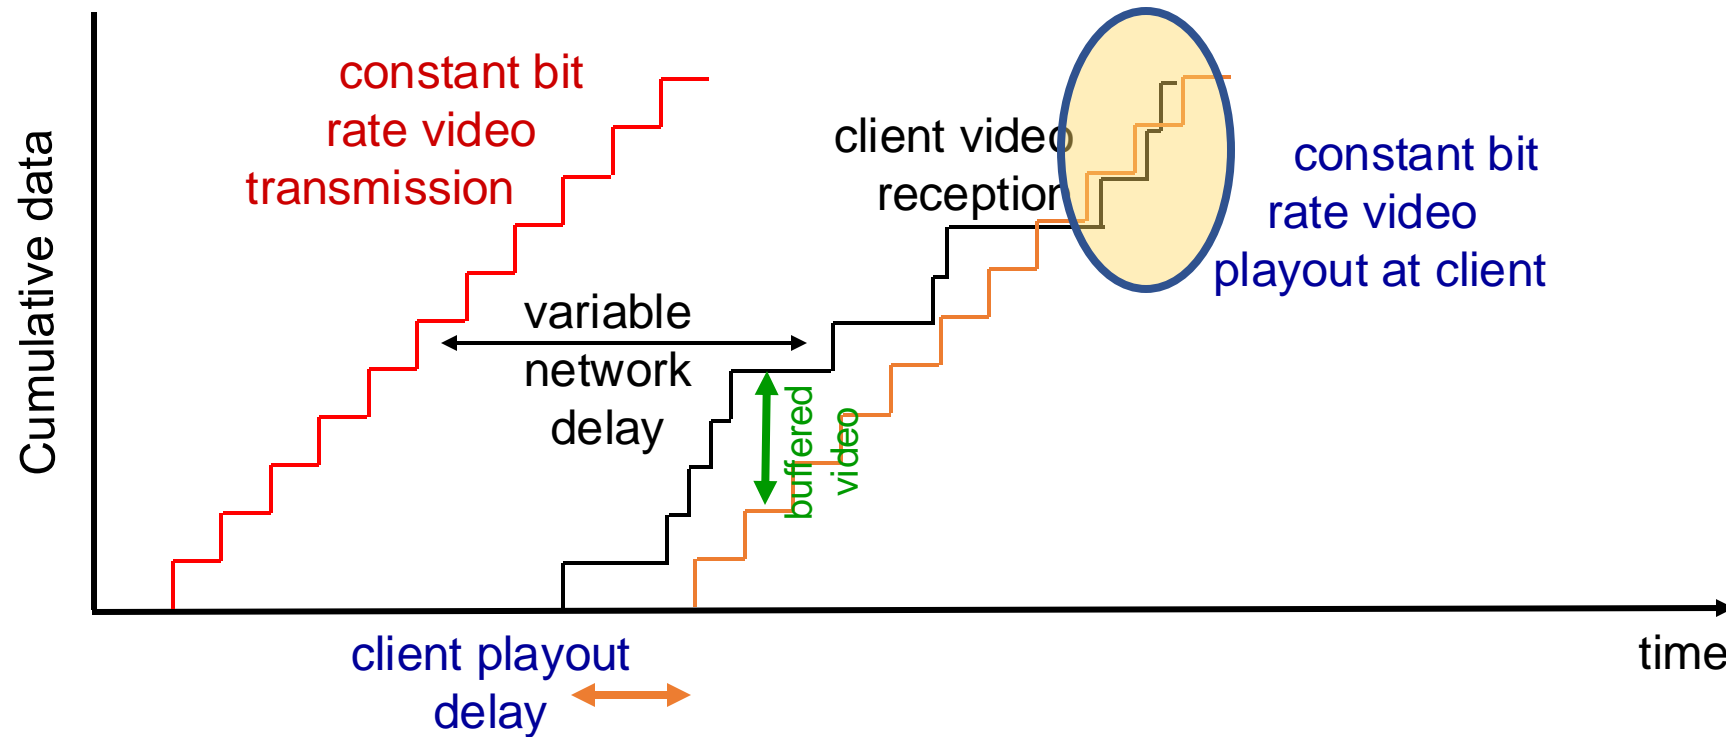


# Introduce a delay for smooth playout



**Client-side buffering with playout delay:**  
compensate for variations in the network delay

# But not too small a delay

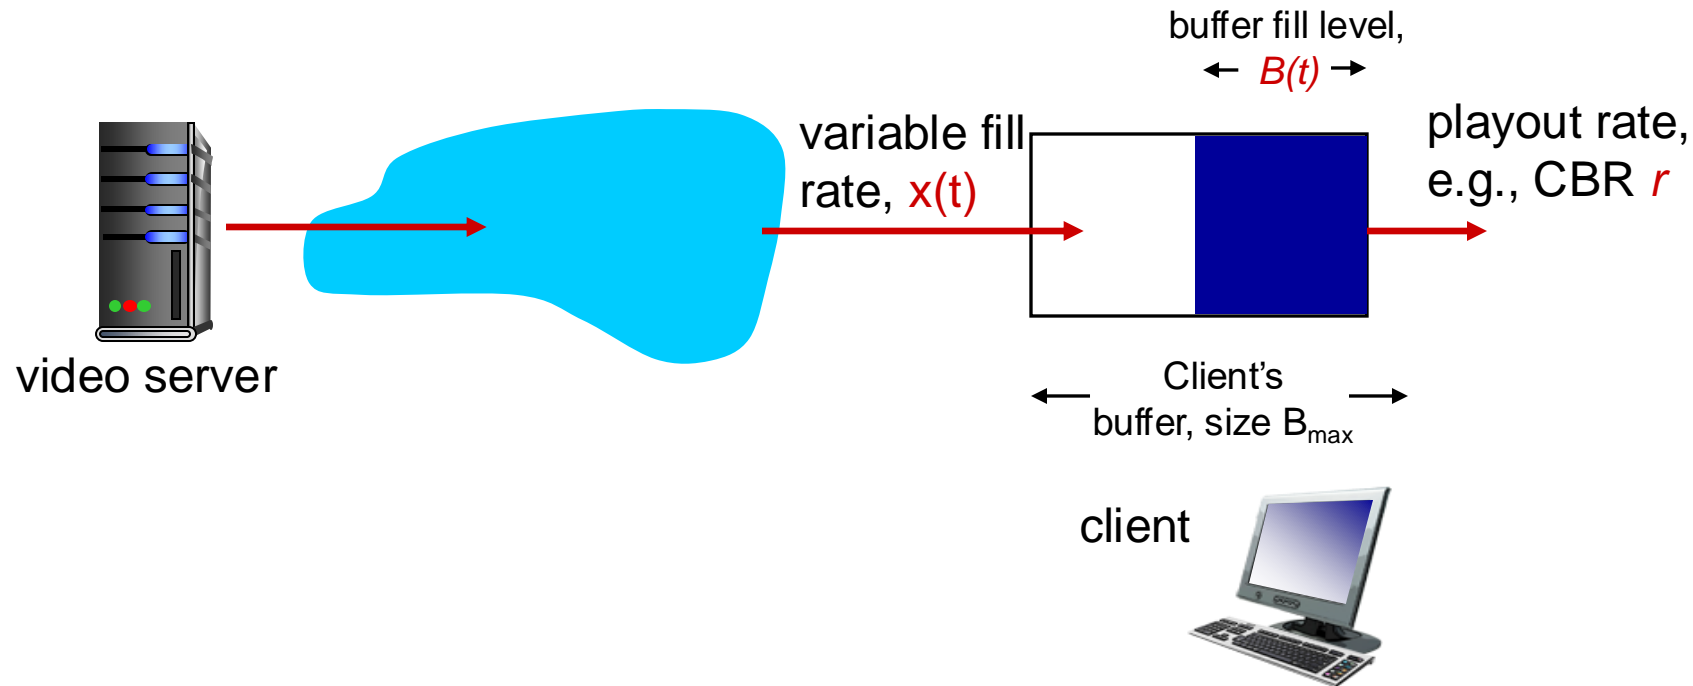


**Playout delay that's too small can cause stalls**

There's nothing in the buffer to show to the user




# Client-side buffering, playout




Most video is broken up in time into multiple **segments**  
Client downloads video segment by segment  
For example: a segment might be 4 seconds worth of video.

# Client-side buffering, playout



1. Initial fill of buffer until playout begins at  $t_p$
2. playout begins at  $t_p$
3. buffer fill level varies over time as fill rate  $x(t)$  varies (assume playout rate  $r$  is constant for now)


# Client-side buffering, playout



*playout buffering: average fill rate ( $\bar{x}$ ), playout rate ( $r$ ):*

- $\bar{x} < r$ : buffer eventually empties for a sufficiently long video. Stall and rebuffering
- $\bar{x} > r$ : buffer will not empty, provided the initial playout delay is large enough to absorb variability in  $x(t)$ 
  - *initial playout delay tradeoff*: buffer starvation less likely with larger delay, but also incur a larger delay until the user begins watching

# Client-side buffering, playout




*playout buffering: average fill rate ( $\bar{x}$ ), playout rate ( $r$ ):*

- is  $\bar{x} < r$  or  $\bar{x} > r$  for a given network connection?
- It can be hard to control  $x$  or even predict it in general
  - Best-effort network inflicts long queues, low bandwidth, loss, etc.
- **How to set the playout rate  $r$ ?**
  - Too low a bit-rate  $r$ : video has poorer quality than needed
  - Too high a bit-rate  $r$ : buffer might empty out. Stall/rebuffering!

# Adaptive bit-rate video


- Motivation: Want to provide high quality video experience, without stalls
- Observations:
  - Videos come in different qualities (average bit rates)
  - Versions of the video for different quality levels readily available
  - Different segments of video can be downloaded separately
- **Adapt bit rate per segment** through collaboration between the video client (e.g., your browser) and the server (e.g., @ Netflix)
- **Adaptive bit-rate (ABR) video**: change the bit-rate (quality) of next video segment based on network and client conditions
- A typical strategy: **Buffer-based rate adaptation**



# Buffer-based bit-rate adaptation

- Key idea: If there is a large stored buffer of video, **optimize for video quality**, i.e., high bit rates
- Else (i.e., buffer has low occupancy), **avoid stalls** by being conservative and ask for a lower quality (bit-rate)
  - Hope: lower bandwidth requirement of a lower quality stream is satisfiable more easily
- Buffer is measured in seconds of playout left before stalling

# Buffer-based bit-rate adaptation



A highly effective method to provide high video quality despite variable and intermittently poor network conditions.

Used by Netflix.

<http://yuba.stanford.edu/~nickm/papers/sigcomm2014-video.pdf>

A Buffer-Based Approach to Rate Adaptation

# Dynamic Adaptive Streaming over HTTP (DASH)




# Streaming multimedia with DASH

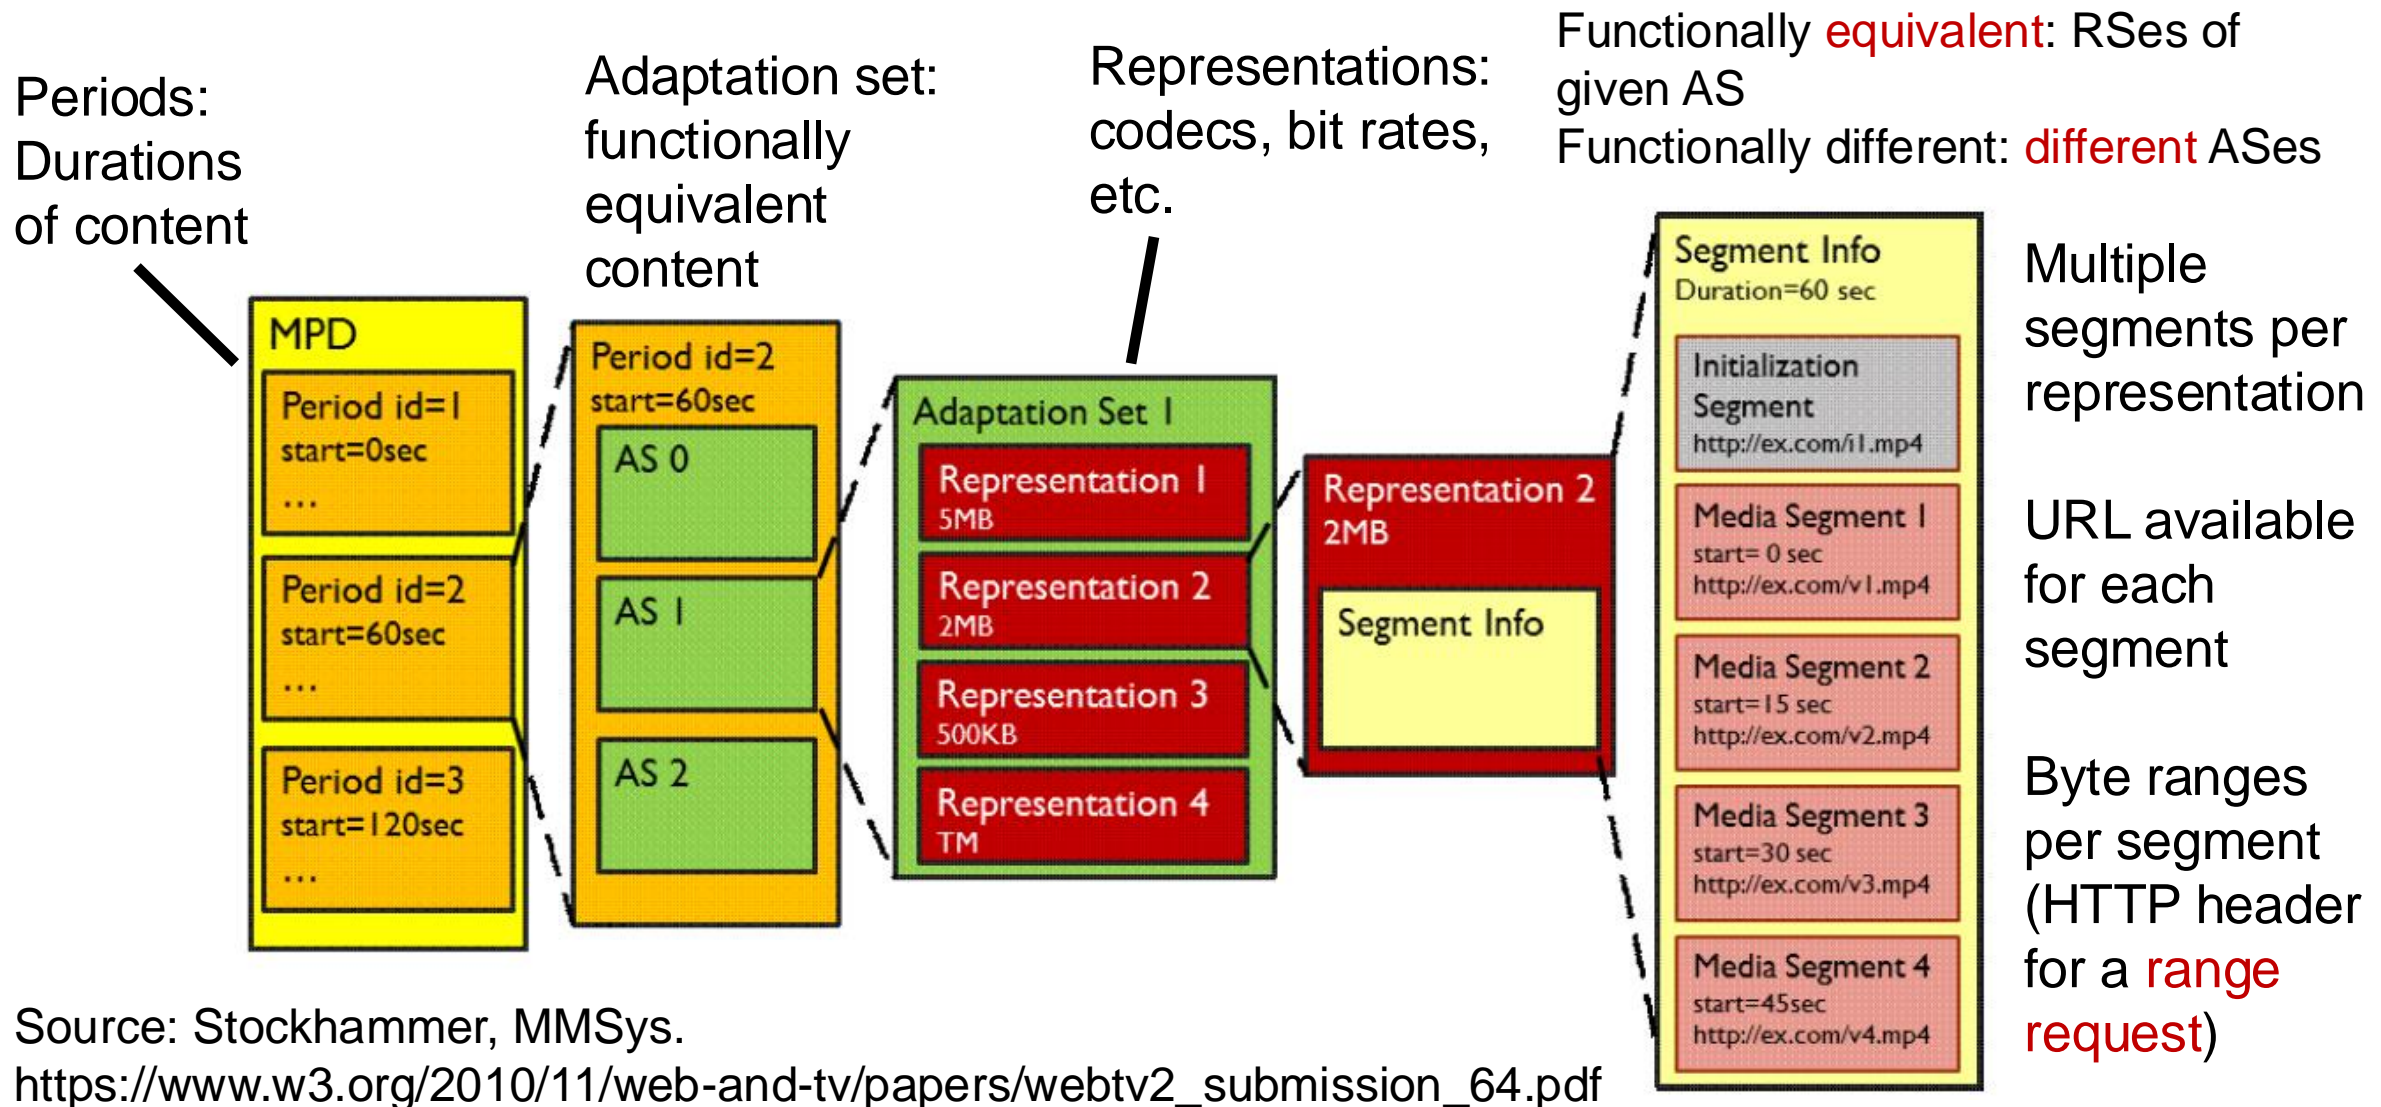
- Dynamic Adaptive Streaming over **HTTP**
  - Used by Netflix and most popular video streaming services
- **Adaptive:** Perform video bit rate adaptation
  - It can be done on the client or the server (with client feedback)
- **Dynamic:** Retrieve a single video from multiple sources
- The DASH video server is just a standard HTTP server
  - Provides video/audio content in multiple formats and encodings
- Leverage existing web infrastructure
  - CDN
  - DNS

# DASH: Key ideas

- Content (video, audio, transcript, etc.) divided into **segments (time)**
- Algorithms to determine and request **varying** attributes (e.g., bitrate, language) for each segment
- Goal: ensure good quality of service, match user prefs, etc.




# What does the manifest contain?



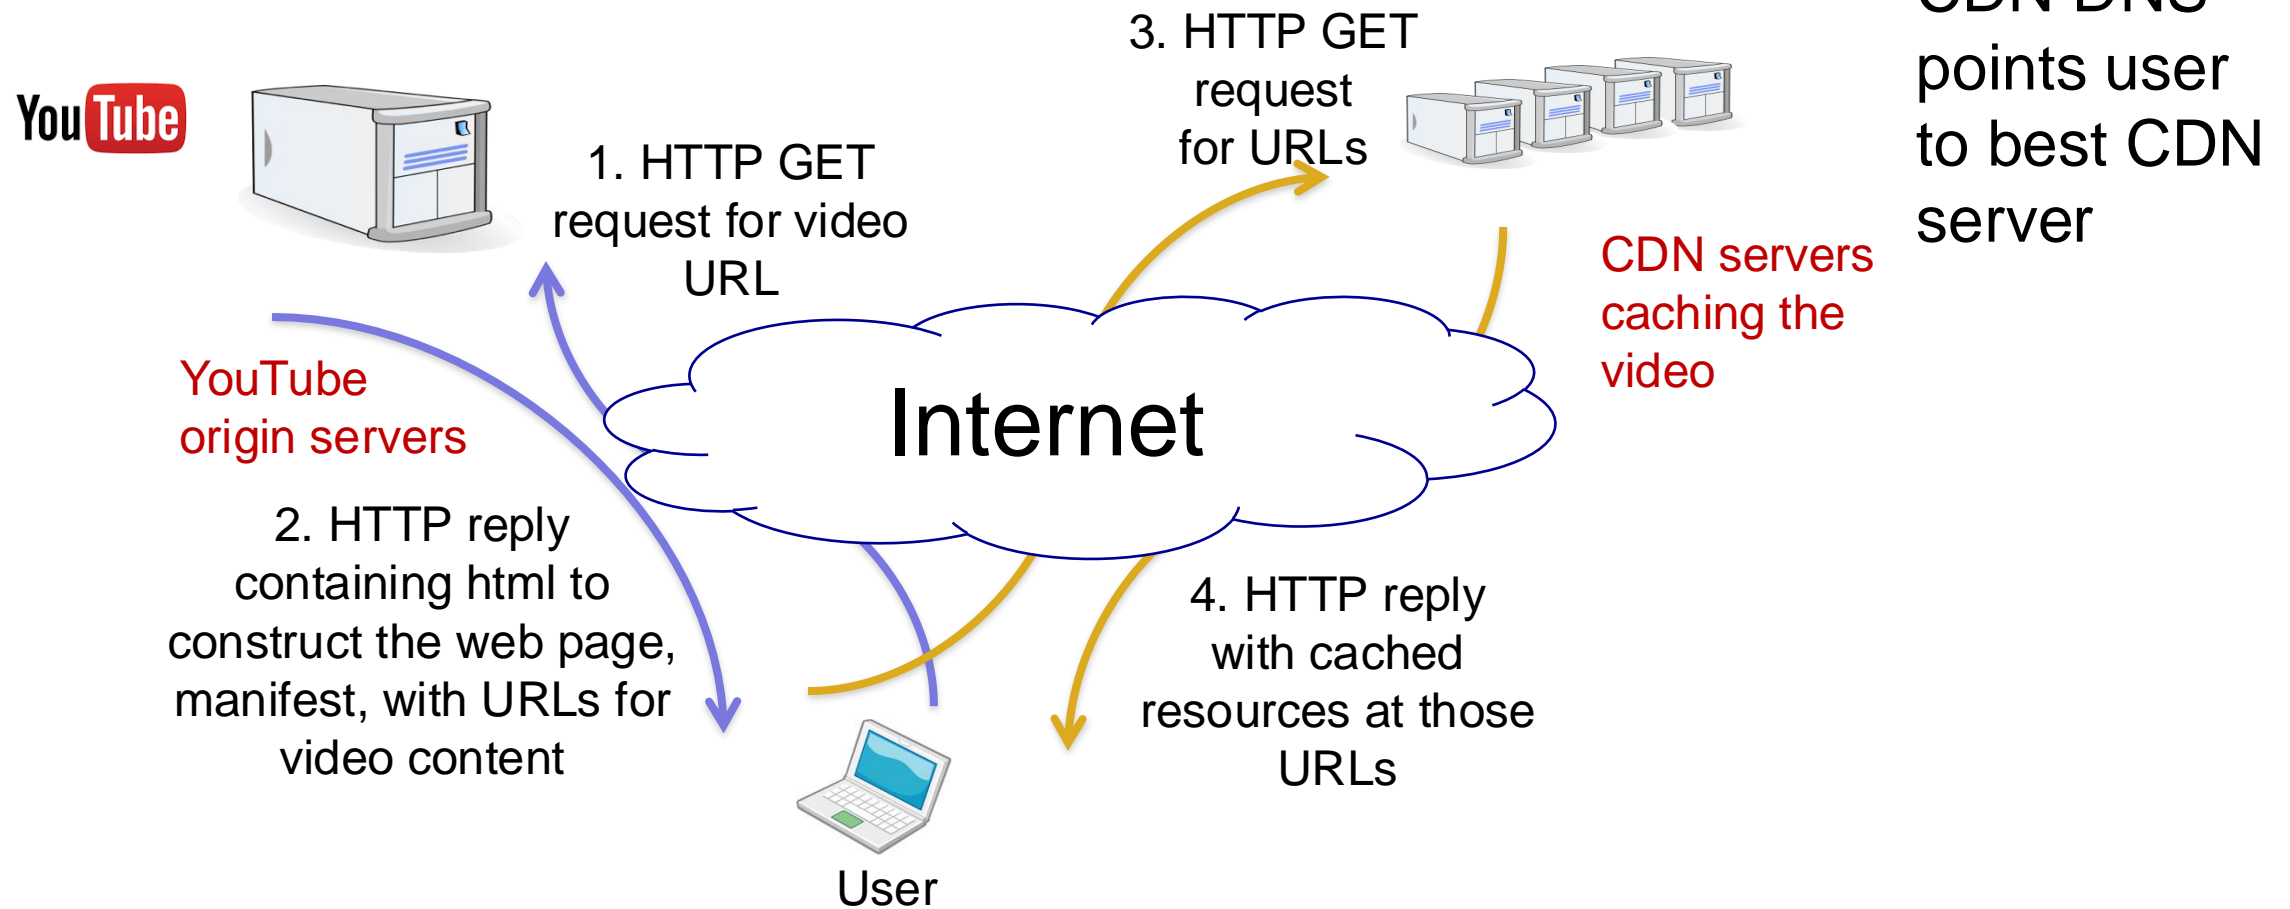


# Dynamic changes in stream quality



# Dynamic changes in stream location

- Just an HTTP request for an HTTP object



# DASH reference player

- <https://reference.dashif.org/dash.js/latest/samples/dash-if-reference-player/index.html>

# DASH Summary

- Piggyback video on HTTP: **widely used**
- Enables independent HTTP requests per segment
  - Choose dynamic quality & preferences over time
  - Independent HTTP byte ranges
- Works well with CDNs
  - Fetch segments from locations other than the origin server
  - Fetch different segments from possibly different locations
- More resources on DASH
  - [https://www.w3.org/2010/11/web-and-tv/papers/webtv2\\_submission\\_64.pdf](https://www.w3.org/2010/11/web-and-tv/papers/webtv2_submission_64.pdf)
  - <https://www.youtube.com/watch?v=xgowGnH5kUE>

# Application Layer: Wrap-up

- Name resolution, the web, video
- Protocols built over the `socket()` abstraction
- Simple designs go a long way
  - Plain text protocols, header-based evolution
- Infrastructure for functionality, performance, ...
  - CDNs, web proxies
- Fit your apps to run on browsers: run almost anywhere (e.g. video)
- Apps are ultimately what users and most engineers care about
- BUT: if you don't understand what's under the hood, you risk bad design and poor performance in Internet applications

