

Video

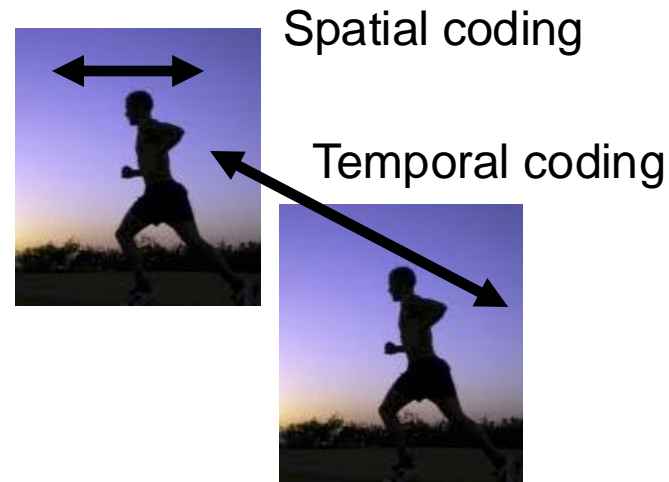
Lecture 10

<http://www.cs.rutgers.edu/~sn624/352-F24>



Srinivas Narayana

Quick recap of concepts


Multimedia



Video **Bitrate**




Client-side buffering, playout




Most video is broken up in time into multiple **segments**
Client downloads video segment by segment
For example: a segment might be 4 seconds worth of video.

Client-side buffering, playout



1. Initial fill of buffer until playout begins at t_p
2. playout begins at t_p
3. buffer fill level varies over time as fill rate $x(t)$ varies (assume playout rate r is constant)


Client-side buffering, playout



playout buffering: average fill rate (\bar{x}), playout rate (r):

- $\bar{x} < r$: buffer eventually empties for a sufficiently long video.
 - Stall and rebuffering
- $\bar{x} > r$: buffer will not empty, provided the initial playout delay is large enough to absorb variability in $x(t)$
 - *initial playout delay tradeoff*: buffer starvation less likely with larger delay, but also incur a larger delay for the user to begin watching

Client-side buffering, playout




playout buffering: average fill rate (\bar{x}), playout rate (r):

- is $\bar{x} < r$ or $\bar{x} > r$ for a given network connection?
- It can be hard to control x or even predict it in general
 - Best-effort network inflicts long queues, low bandwidth, loss, etc.
- **How to set the playout rate r ?**
 - Too low a bit-rate r : video has poorer quality than needed
 - Too high a bit-rate r : buffer might empty out. Stall/rebuffering

Adaptive bit-rate video


- Motivation: Want to provide high quality video experience, without stalls
- Observations:
 - Videos come in different qualities (average bit rates)
 - Versions of the video for different quality levels readily available
 - Different segments of video can be downloaded separately
- **Adapt bit rate per segment** through collaboration between the video client (e.g., your browser) and the server (e.g., @ Netflix)
- **Adaptive bit-rate (ABR) video**: change the bit-rate (quality) of next video segment based on network and client conditions
- A typical strategy: **Buffer-based rate adaptation**



Buffer-based bit-rate adaptation

- Key idea: If there is a large stored buffer of video at the client, **optimize for video quality**, i.e., high bit rates
- Else (i.e., client video buffer has low occupancy), **avoid stalls** by being conservative and asking for a lower quality (bit-rate)
 - The hope: the lower bandwidth requirement of a lower-quality stream is more easily met; stalls averted
- Buffer is measured in seconds of playout left before stalling

Buffer-based bit-rate adaptation



Provide high video quality overall despite variable and intermittently poor network conditions.

<http://yuba.stanford.edu/~nickm/papers/sigcomm2014-video.pdf>

A Buffer-Based Approach to Rate Adaptation
(used by Netflix)

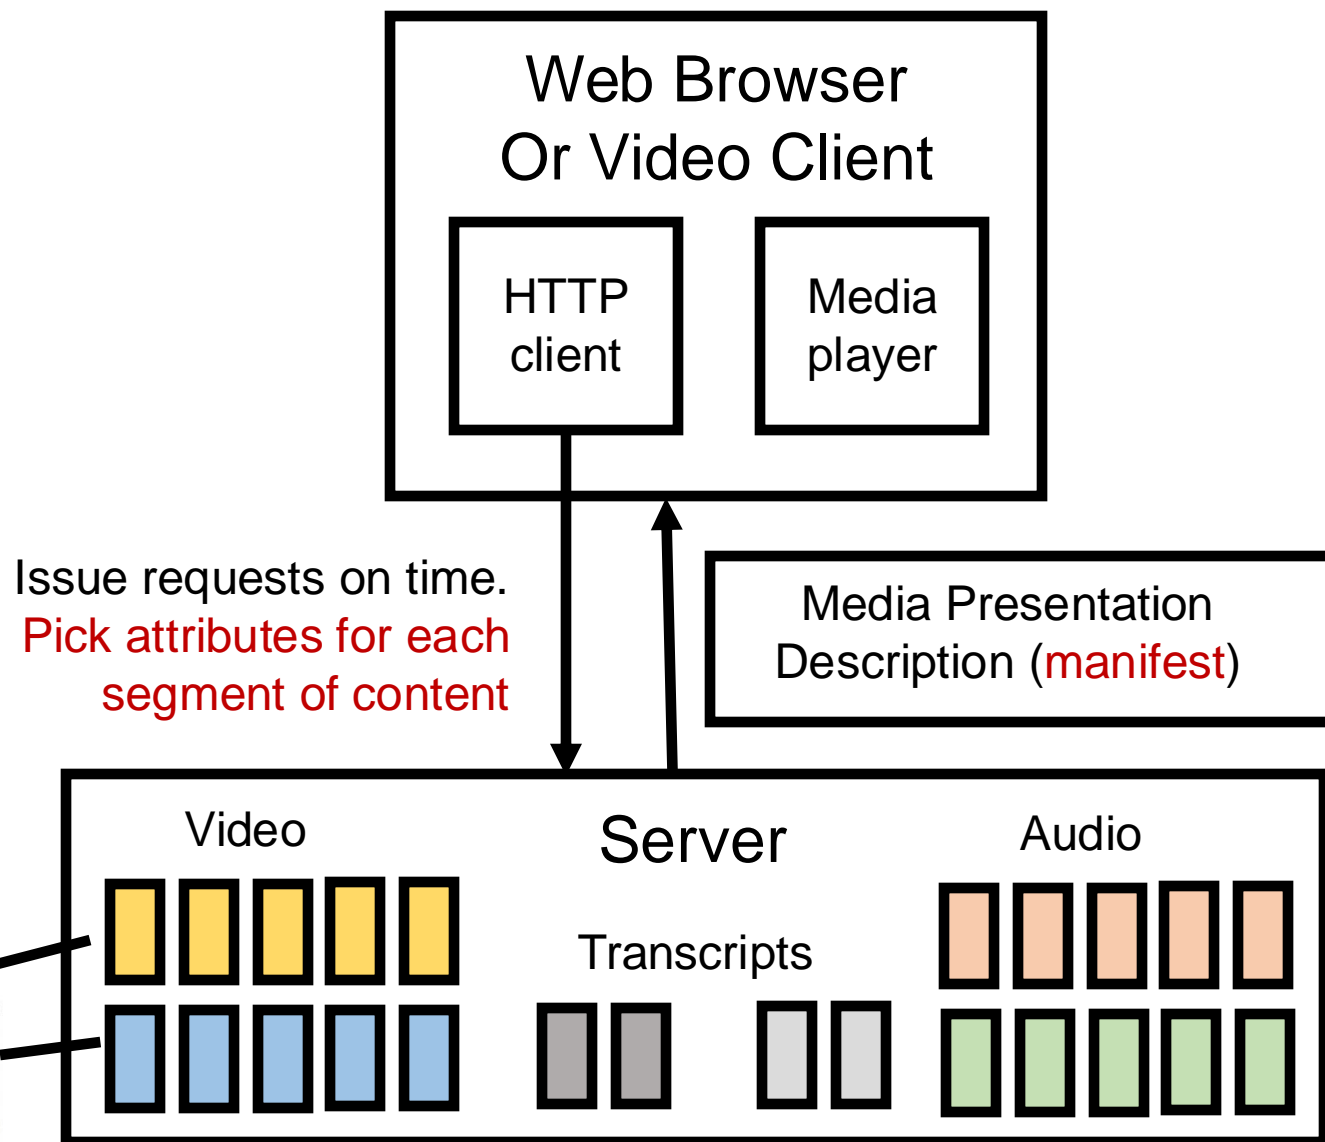
Dynamic Adaptive Streaming over HTTP (DASH)

Streaming multimedia with HTTP

- Early video: basic UDP
 - Problems: reliability, blocking
- Today: repurpose **web** infrastructure and protocols for video
- **DASH: Dynamic Adaptive Streaming over HTTP**
 - Used by Netflix, YouTube, and other video streaming services

DASH: Key ideas

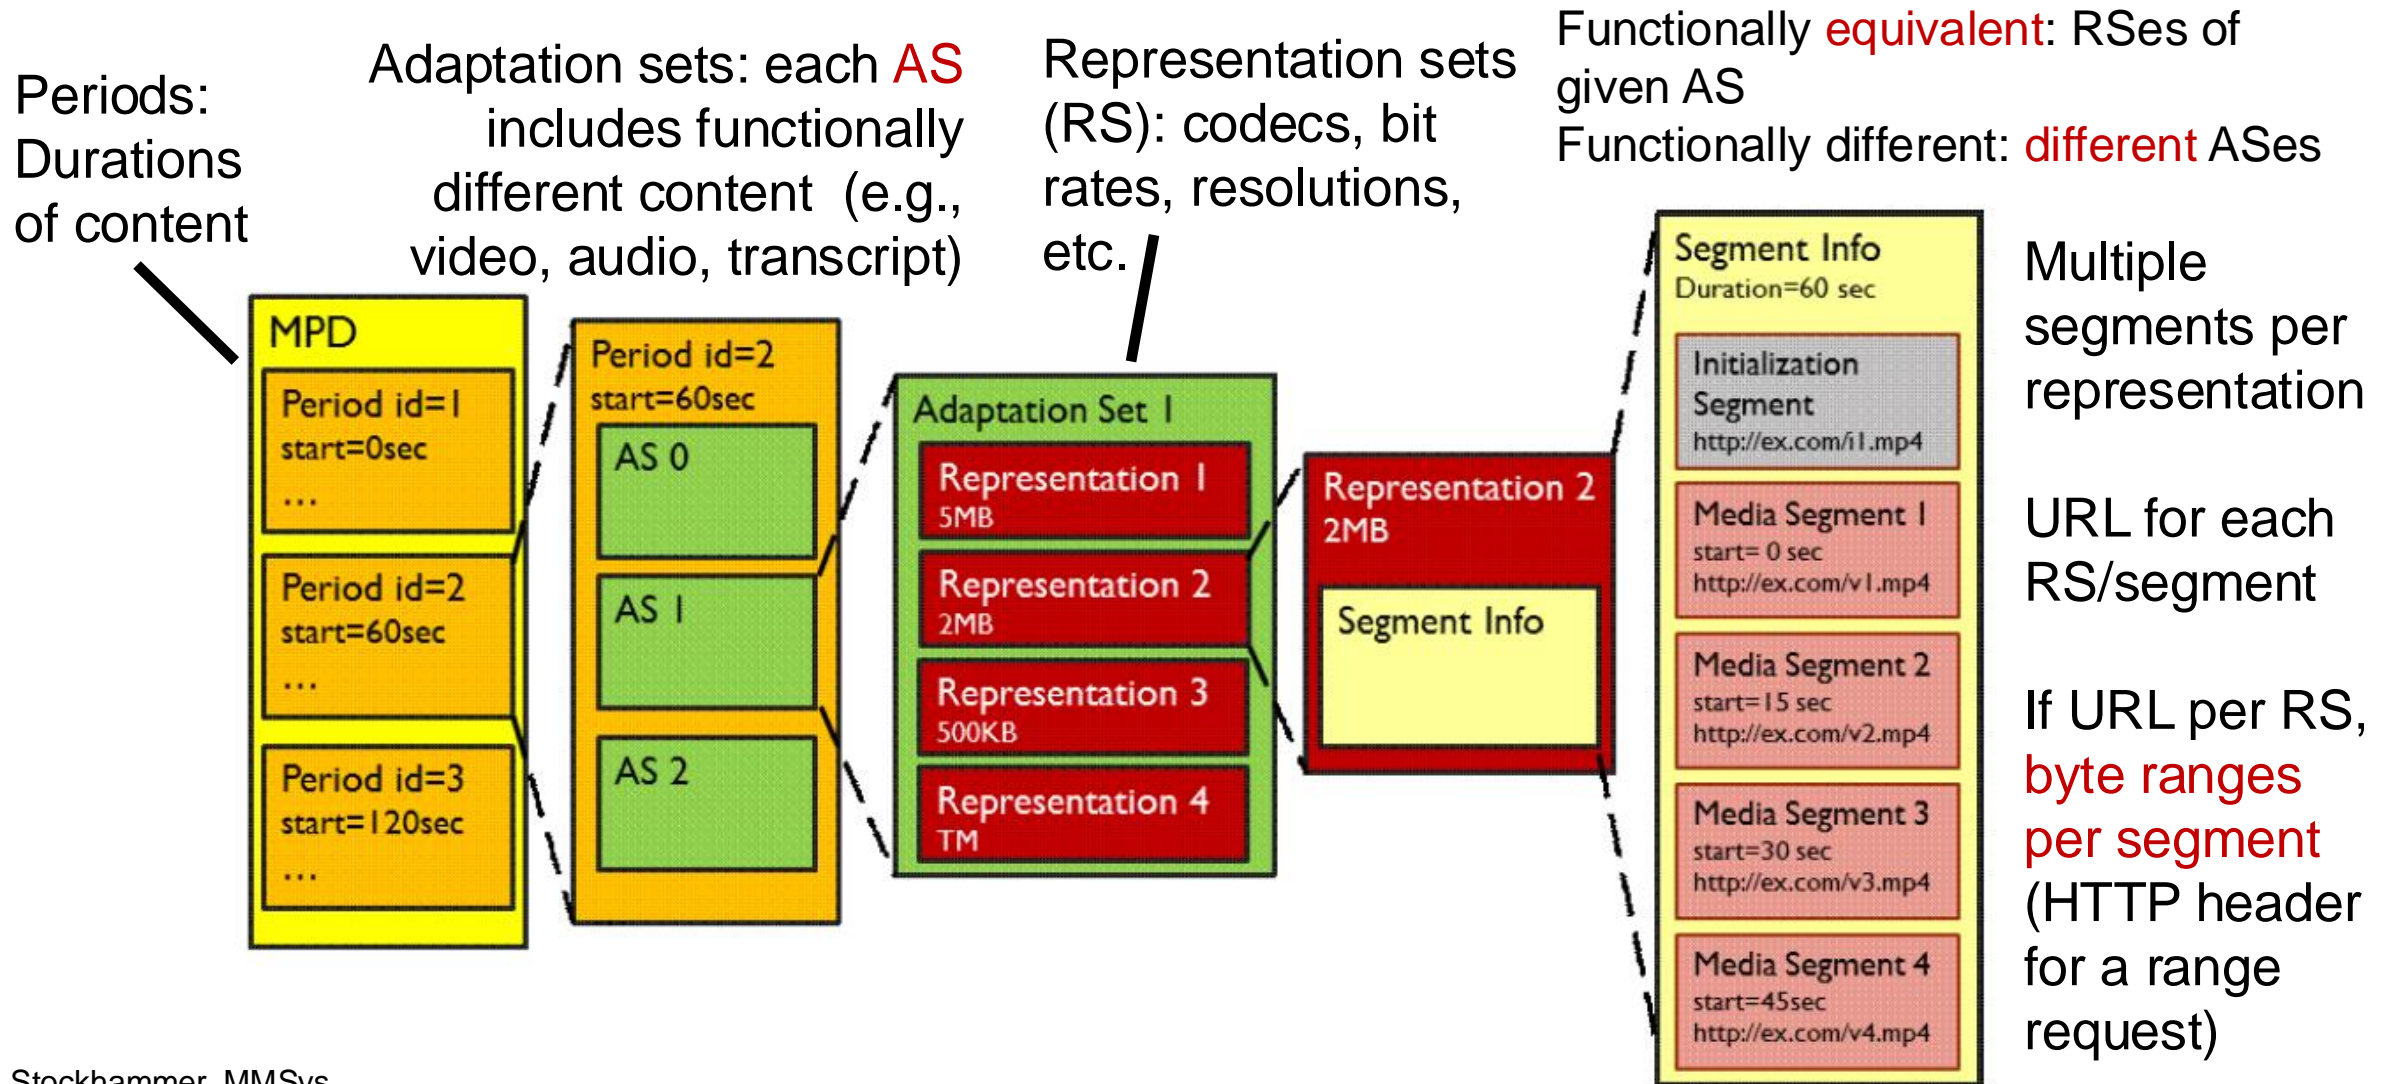
- Content (video, audio, transcript, etc.) divided into **segments (time)**
- Algorithms to determine and request **varying** attributes (e.g., bitrate, language) for each segment
- Goal: ensure good quality of service, match user prefs, etc.




Streaming multimedia with HTTP

- Dividing up a video into multiple segments enables a few things
- Possible to decide the quality **per chunk**
 - Enables **bit rate adaptation**
 - Typically done on the client, but possible on the server (with feedback)
 - Change language, receive sub-titles, etc. mid-stream
- Retrieve segments of a single video from multiple sources
- DASH video server is just a standard HTTP server
 - Client issue HTTP GET requests (typically with the **Range** request header)
 - Leverage existing web infrastructure (CDN, DNS)
- Send different clients to different video sources (CDN)

What does the manifest contain?

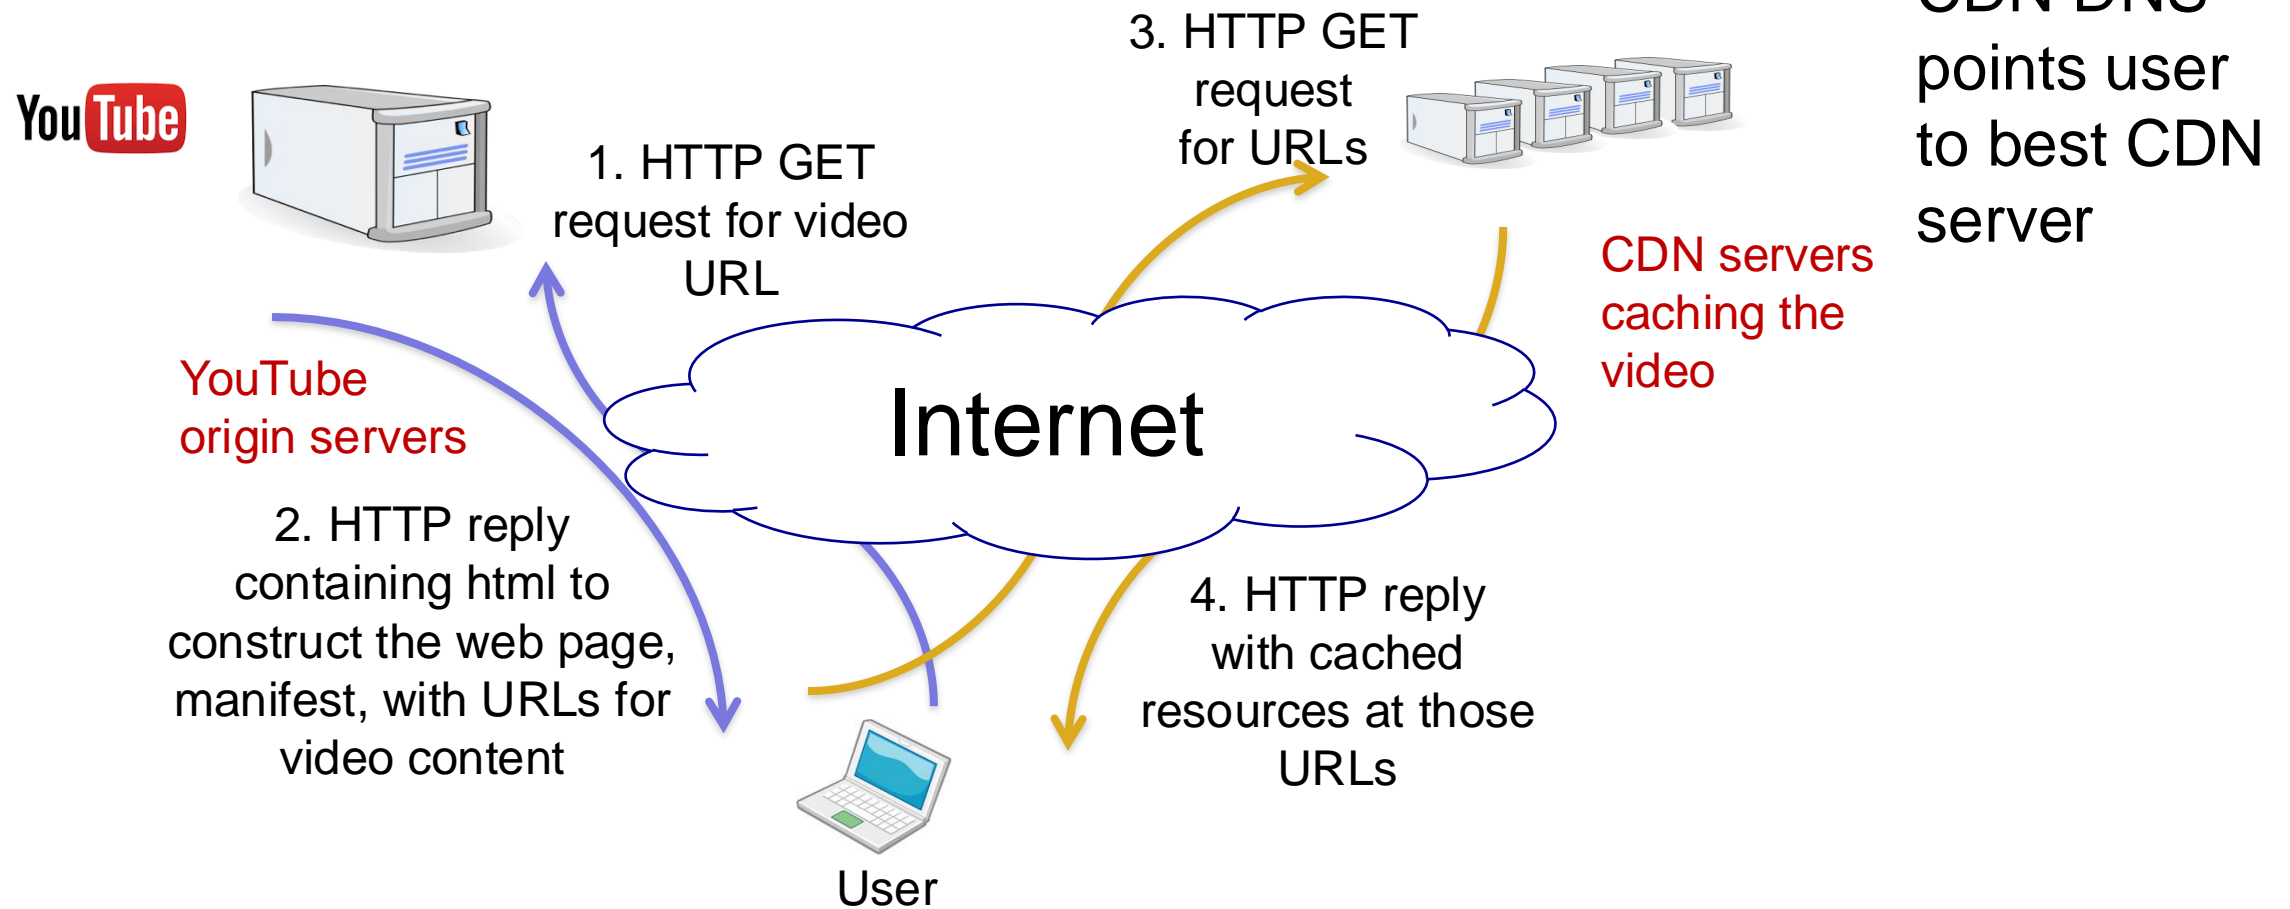


Adaptive changes in quality



Dynamic server selection based on client

- Just an HTTP request for an HTTP object



DASH reference player

- <https://reference.dashif.org/dash.js/latest/samples/dash-if-reference-player/index.html>

DASH Summary

- Piggyback video on HTTP: **widely used**
- Enables independent HTTP requests per segment
 - Works well with CDNs
 - Adapt quality with time, location per client, possibly location over time
 - May use HTTP range requests to ask byte ranges in a segment URL
 - Fetch segments from locations other than the origin server
- More resources on DASH
 - https://www.w3.org/2010/11/web-and-tv/papers/webtv2_submission_64.pdf
 - <https://www.youtube.com/watch?v=xgowGnH5kUE>

Application Layer: Wrap-up

- Name resolution, the web, video
- Protocols built over the `socket()` abstraction
- Simple designs go a long way
 - Plain text protocols, header-based evolution
- Infrastructure for functionality & performance
 - CDNs, web proxies
- Fit your apps to run on browsers: run almost anywhere (e.g. video)
- Apps are ultimately what users and most engineers care about
- But, if you don't understand what's under the hood, you risk bad design and poor performance in Internet applications

