


# Ordering & Flow Control

Lecture 15


<http://www.cs.rutgers.edu/~sn624/352-F24>

Srinivas Narayana

# Review: Reliability





## Stop and Wait



throughput

## Pipelined Reliability



# TCP reliability metadata

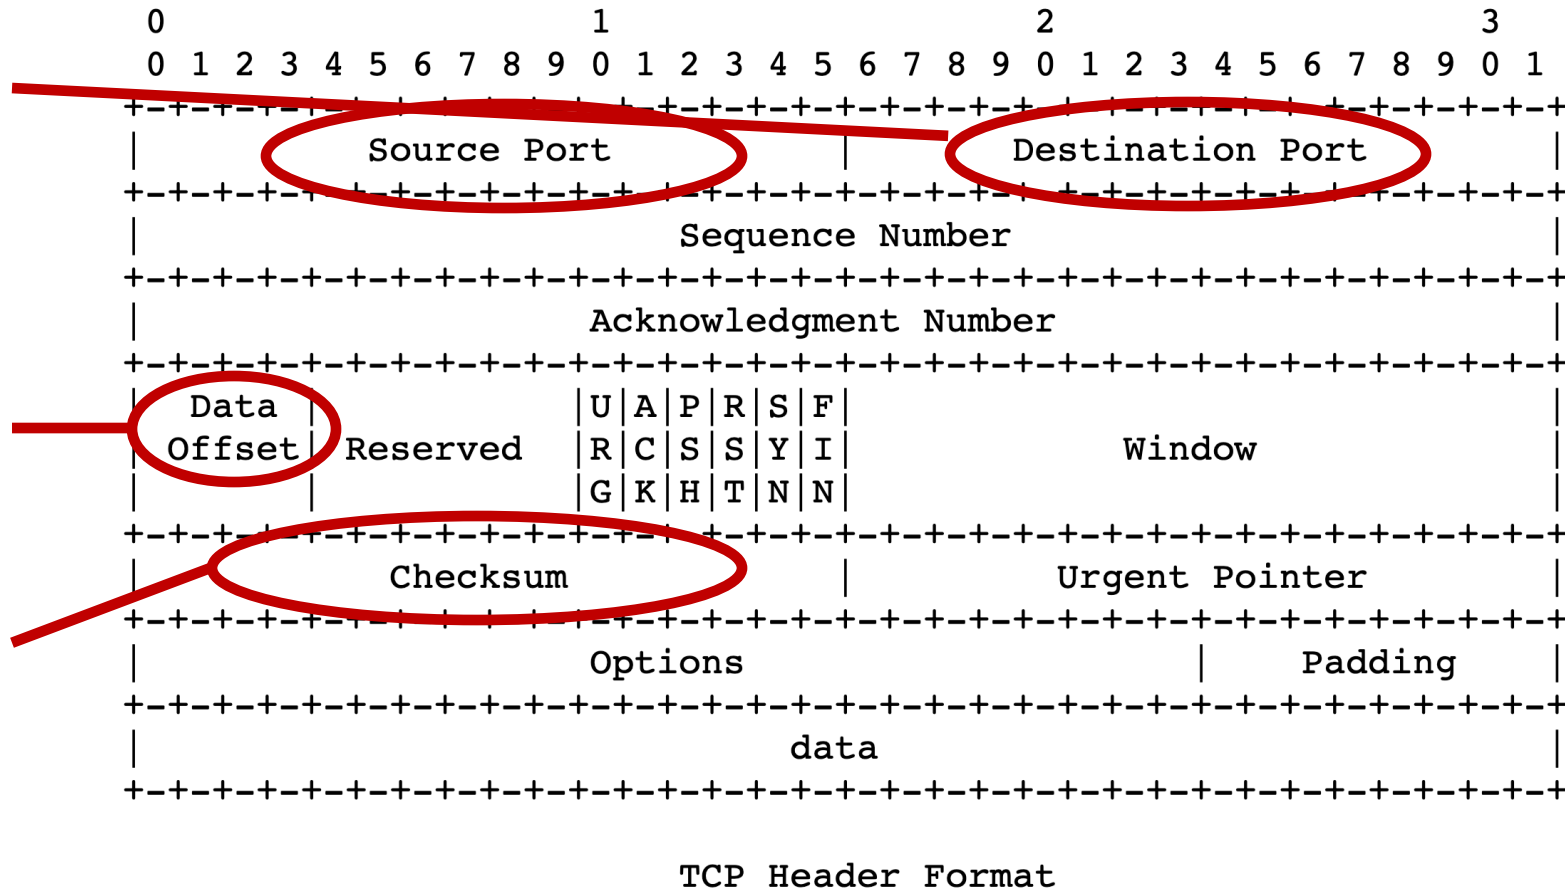
Where are reliability metadata (seq and ack numbers) stored?

# TCP header structure

Source port, destination port (connection demultiplexing)

Size of the TCP header (in 32-bit words)

Basic error detection through checksums (similar to UDP)



Note that one tick mark represents one bit position.


# TCP header structure

Identifies data in the packet  
from the application stream

TCP uses byte seq #s

TCP acks the next seq #  
that the receiver expects.  
(cumulative ACK)

Selective ACKs are written  
into the options field



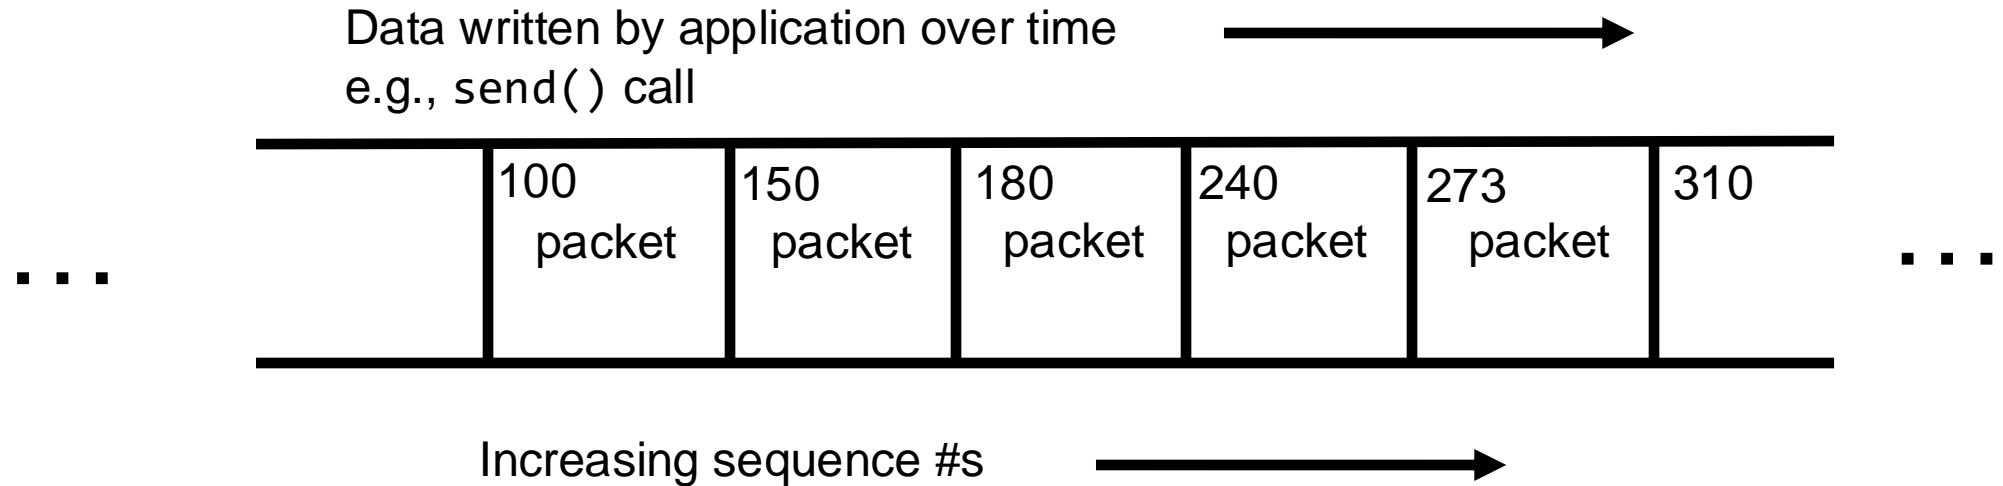
Note that one tick mark represents one bit position.

# Observing a TCP exchange

- `sudo tcpdump -i eno1 tcp portrange 56000-56010`
- `curl --local-port 56000-56010  
https://www.google.com > output.html`
- **Bonus:** Try crafting TCP packets with `scapy`!

# TCP Stream-Oriented Data Transfer


# Sequence numbers in the app's **stream**



TCP uses byte sequence numbers



# Sequence numbers in the app's **stream**

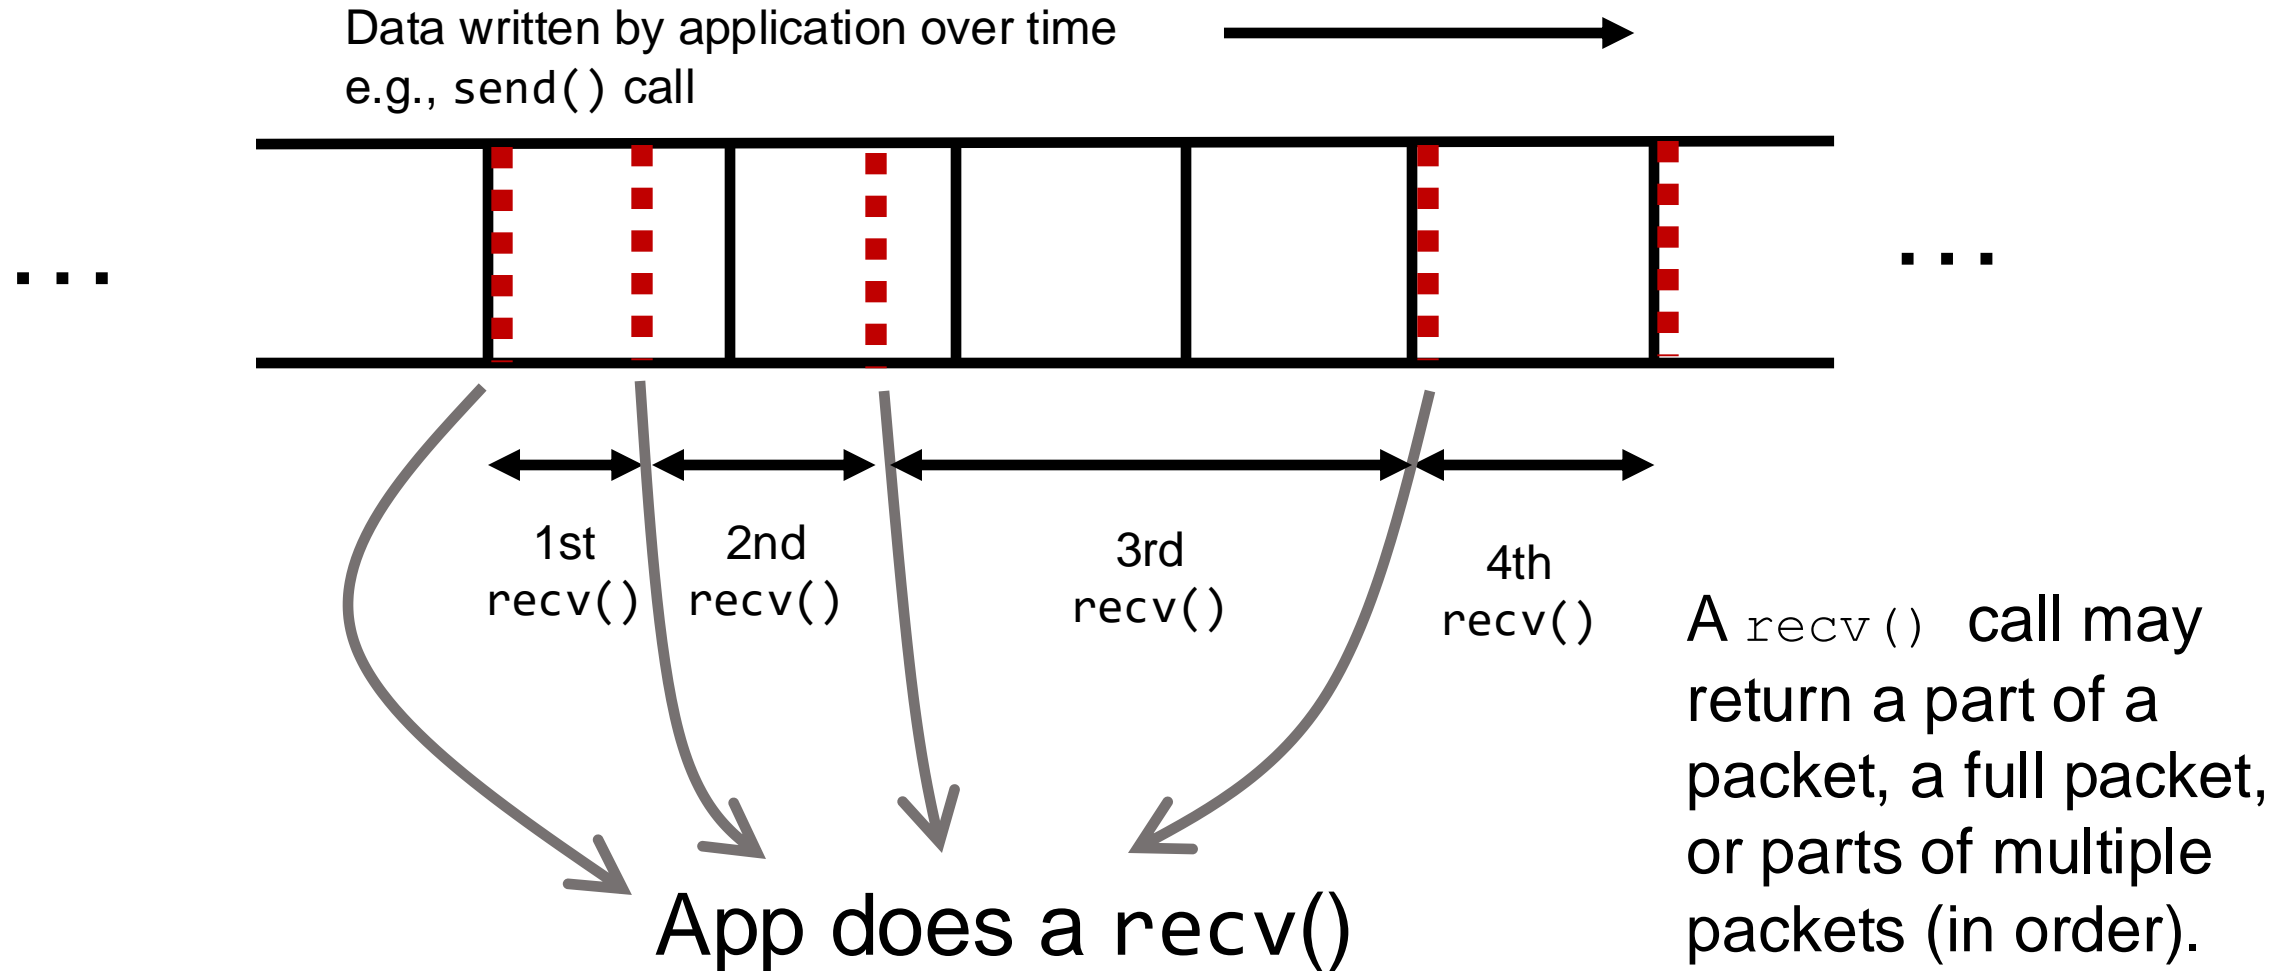


Packet boundaries aren't important for TCP software

TCP is a **stream-oriented** protocol

(We use `SOCK_STREAM` when creating sockets)

# Sequence numbers in the app's **stream**



# Buffering and Ordering in TCP

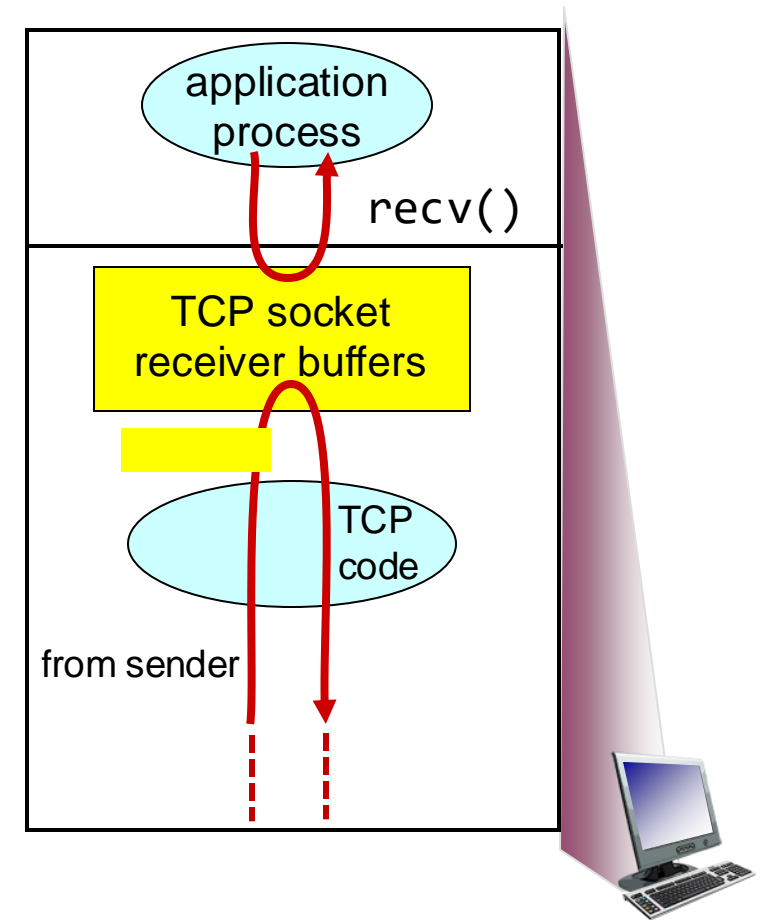
# Memory Buffers at the Transport Layer

# Sockets need receive-side memory buffers

- Since TCP uses selective repeat, the receiver must **buffer** data that is received after loss:
  - e.g., hold packets so that only the “holes” (due to loss) need to be filled in later, without having to retransmit packets that were received successfully
- Apps read from the receive-side socket buffer when you do a `recv()` call.
- Even if data is always reliably received, applications may not always read the data immediately
  - What if you invoked `recv()` in your program infrequently (or never)?
  - For the same reason, UDP sockets also have receive-side buffers

# Receiver app's interaction with TCP

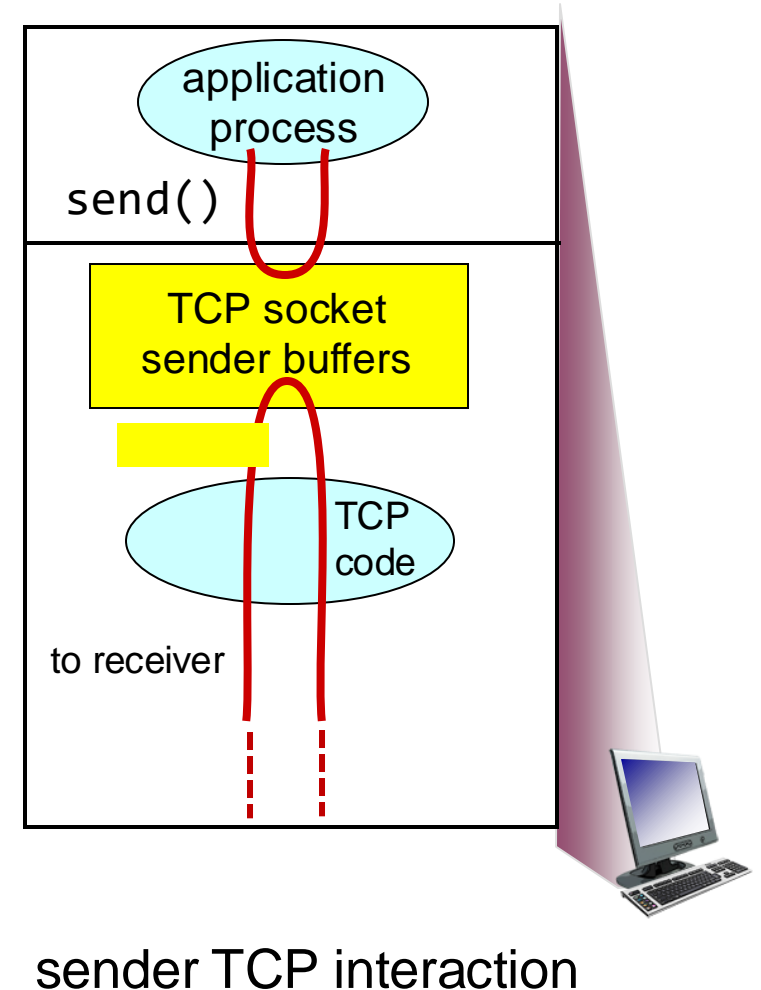
- Upon reception of data, the receiver's TCP stack deposits the data in the receive-side socket buffer
- An app with a TCP socket reads from the TCP receive socket buffer
  - e.g., when you do `data = sock.recv()`



receiver TCP interaction

# Sockets need send-side memory buffers

- The possibility of **packet retransmission** in the future means that data can't be immediately discarded from the sender once transmitted.
- App has issued `send()` and moved on; TCP stack must buffer this data
- Transport layer must wait for ACK of a piece of data before reclaiming (freeing) the memory for that data.




Ordered Delivery




# Reordering packets at the receiver side

- Let's suppose receiver gets packets 1, 2, and 4, but not 3 (dropped)
- Suppose you're trying to download a document containing a report
- What would happen if transport at the receiver directly presents packets 1, 2, and 4 to the application (i.e., receiving 1,2,4 through the `recv()` call)?



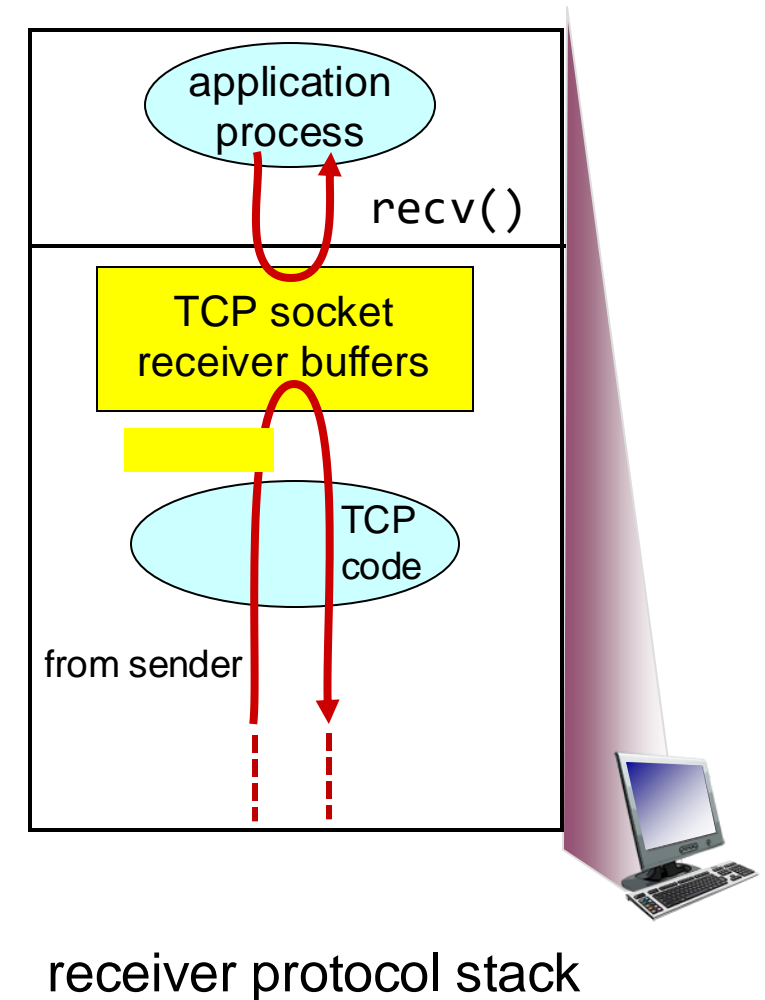
# Reordering packets at the receiver side

- Reordering can happen for a few reasons:
  - Drops
  - Packets taking different paths through a network
- Receiver needs a general strategy to ensure that data is presented to the application **in the same order that the sender pushed it**. Ideas?
- To implement ordered delivery, the receiver uses
  - Sequence numbers
  - Receiver socket buffer
- We've already seen the use of these for reliability; but they can be used to order too!




# Receive-side app and TCP

- TCP receiver software only releases the data from the receive-side socket buffer to the application if:
  - the data is **in order** relative to all other data already read by the application
- This process is called **TCP reassembly**



# TCP Reassembly



# Implications of ordered delivery



- Packets cannot be delivered to the application if there is an **in-order packet missing** from the receiver's buffer
  - The receiver can only buffer so much out-of-order data
  - **Subsequent out-of-order packets dropped**
  - It won't matter that those packets successfully arrive at the receiver from the sender over the network
- **TCP application-level throughput will suffer** if there is too much packet reordering in the network
  - Data may have reached the receiver, but won't be delivered to apps upon a `recv()` (...or may not even be buffered!)

= window size

Proportional to **throughput**


# How much data to keep in flight?

## Stop and Wait



## Pipelined Reliability

# We want to increase throughput, but ...

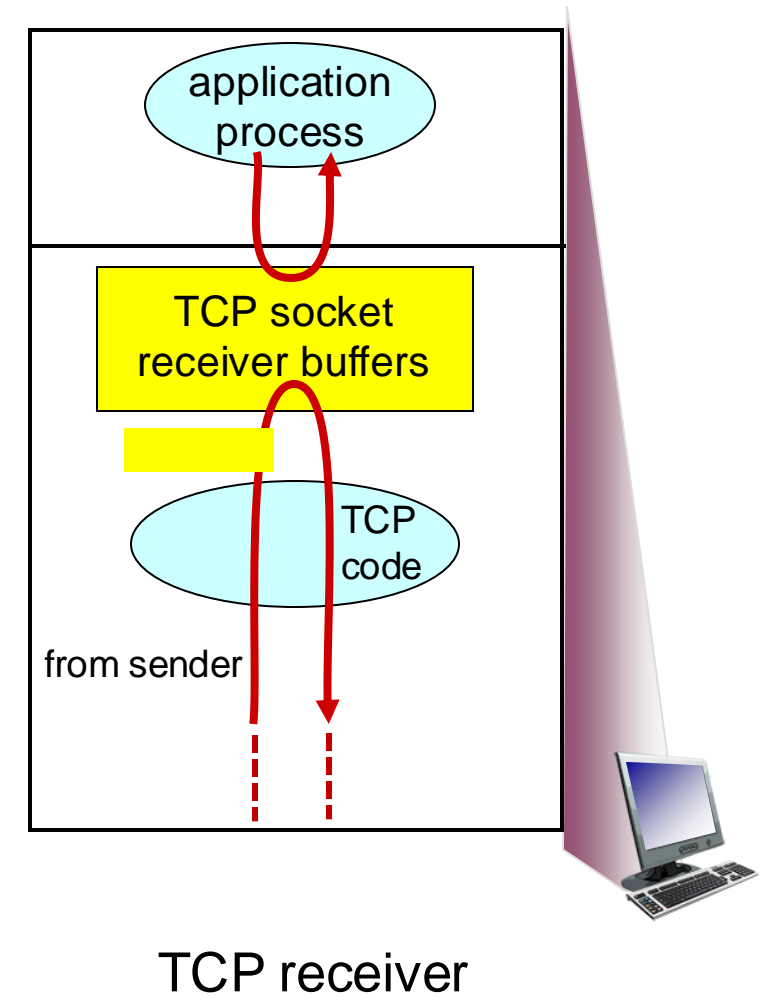


# Flow Control



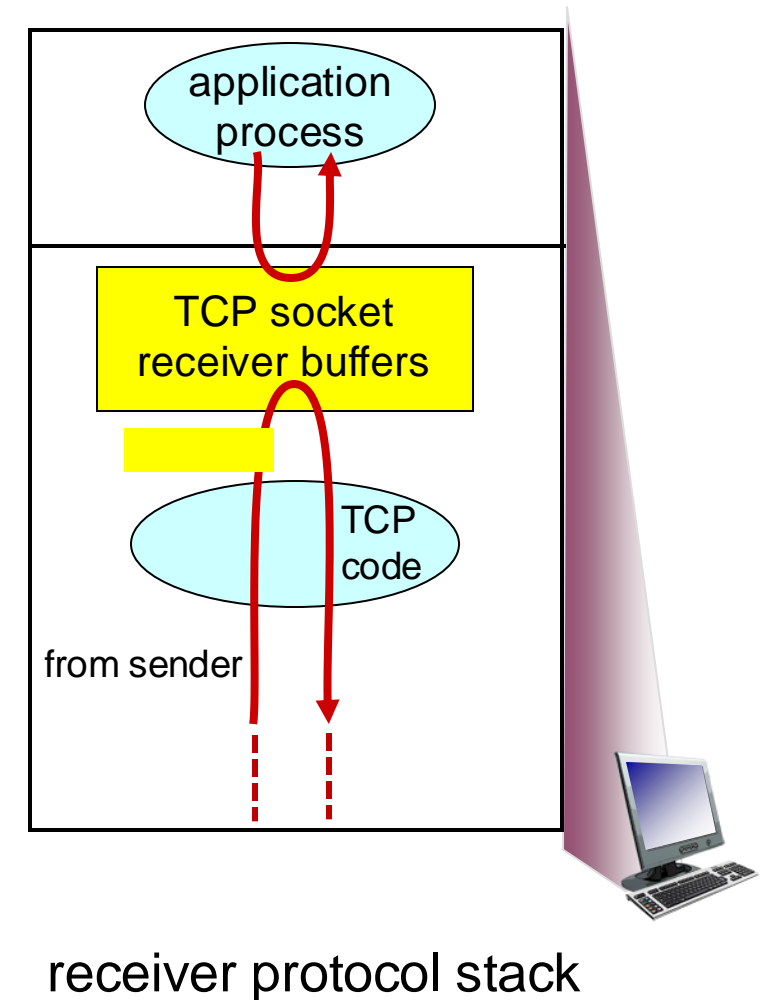
# Socket buffers can become full

- Applications may read data slower than the sender is pushing data in
  - Example: what if an app infrequently or never calls `recv()`?
- There may be too much reordering or packet loss in the network
  - What if the first few bytes of a window are lost or delayed?
- Receivers can only buffer so much before dropping subsequent data




# Goal: avoid drops due to buffer fill

- Have a TCP sender only send as much as the **free buffer space** available at the receiver.
- *Amount of free buffer varies over time!*
- TCP implements **flow control**
- Receiver's ACK contains the amount of data the sender can transmit without running out the receiver's socket buffer
- This number is called the **advertised window size**



# Flow control in TCP headers

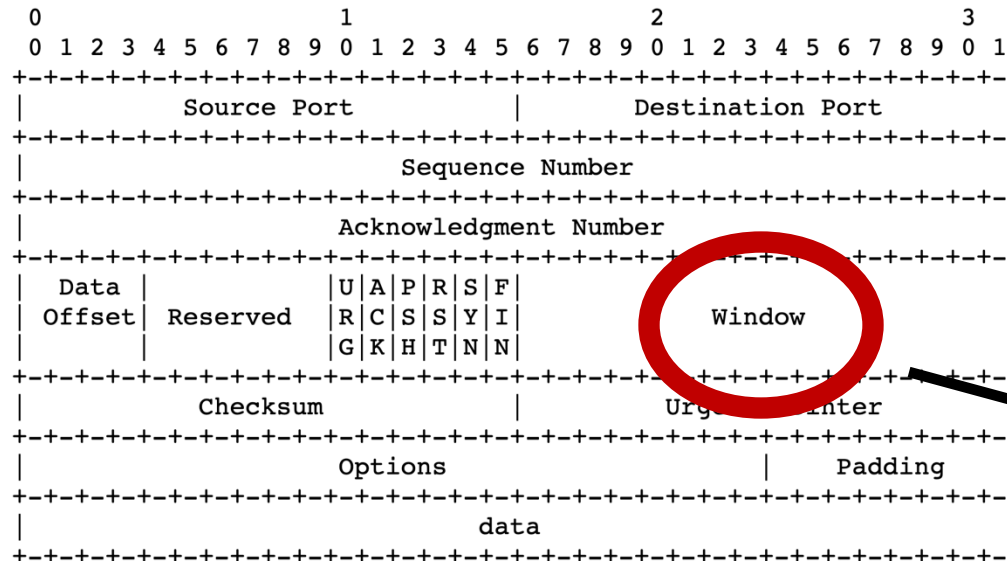


TCP Header Format

Note that one tick mark represents one bit position.


# TCP flow control

- Receiver **advertises** to sender (in the ACK) how much free buffer is available



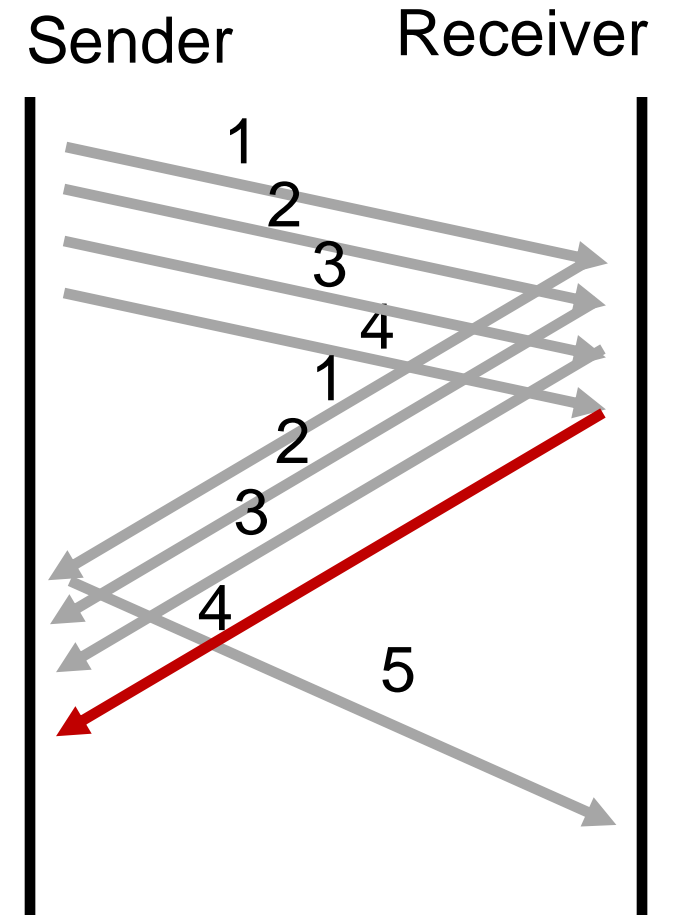
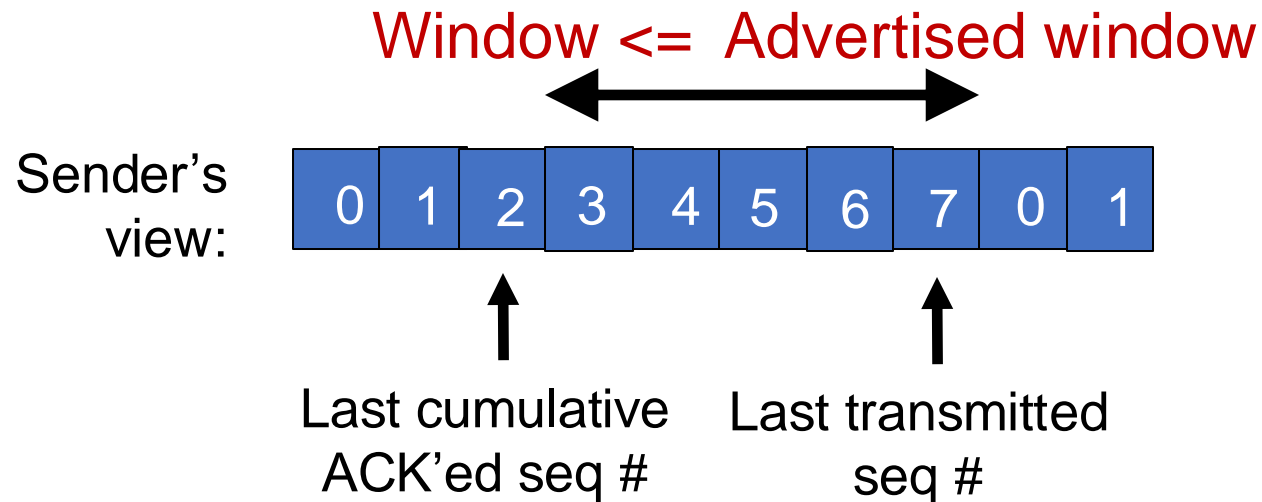
TCP Header Format

Note that one tick mark represents one bit position.





# TCP flow control

- Subsequently, the sender's sliding window cannot be larger than this value
- Restriction on new sequence numbers that can be transmitted
- == restriction on sending rate





# TCP flow control

- If receiver app is too slow reading data:
  - receiver socket buffer fills up
  - => advertised window shrinks
  - => sender's window (sending rate) reduces
  - => sender's socket buffer fills up
  - => sender process put to sleep upon send()



# TCP flow control

Flow control matches the sending process's write speed to the receiving process's read speed.



# Sizing the receiver's socket buffer

- Operating systems have a default receiver socket buffer size
  - Listed among `sysctl -a | grep net.inet.tcp` on MAC
  - Listed among `sysctl -a | grep net.ipv4.tcp` on Linux
- If socket buffer is too small, sender can't keep too many packets in flight → lower throughput
- If socket buffer is too large, too much memory consumed per socket
- How big should the receiver socket buffer be?



# Sizing the receiver's socket buffer

- Case 1: **Suppose the receiving app is reading data too slowly:**
  - no amount of receiver buffer can prevent low sender throughput if the connection is long-lived!

# Sizing the receiver's socket buffer

- Case 2: Suppose the receiving app reads sufficiently fast *on average* to match the sender's writing speed.
  - Assume the sender has a window of size  $W$ .
  - The receiver must use a buffer of size at least  $W$ . Why?
- Captures two cases:
- (1) When the first sequence #s in the window are dropped
  - *Selective repeat*: data in window buffered until the "hole" within the window can be filled by the sender. Advertised window reduces sender's window
- (2) When the sender sends a burst of data of size  $W$ 
  - The receiver may not keep up with the *instantaneous* rate of the sender
- Set receiver socket buffer  $>$  desired window size

# Summary of flow control

- Keep memory buffers available at the receiver whenever the sender transmits data
  - Buffers needed to hold for selective repeat and reassemble data in order
  - Inform the sender on an on-going basis (each ACK)
  - Function: match sender speed to receiver speed
- 
- Correct socket buffer sizing is important for TCP throughput

# Info on (tuning) TCP stack parameters

- [https://www.ibm.com/support/knowledgecenter/linuxonibm/liaag/wkvm/wkvm\\_c\\_tune\\_tcpip.htm](https://www.ibm.com/support/knowledgecenter/linuxonibm/liaag/wkvm/wkvm_c_tune_tcpip.htm)
- <https://cloud.google.com/solutions/tcp-optimization-for-network-performance-in-gcp-and-hybrid>