

# Flow Control; Congestion Control

Lecture 16

<http://www.cs.rutgers.edu/~sn624/352-F24>


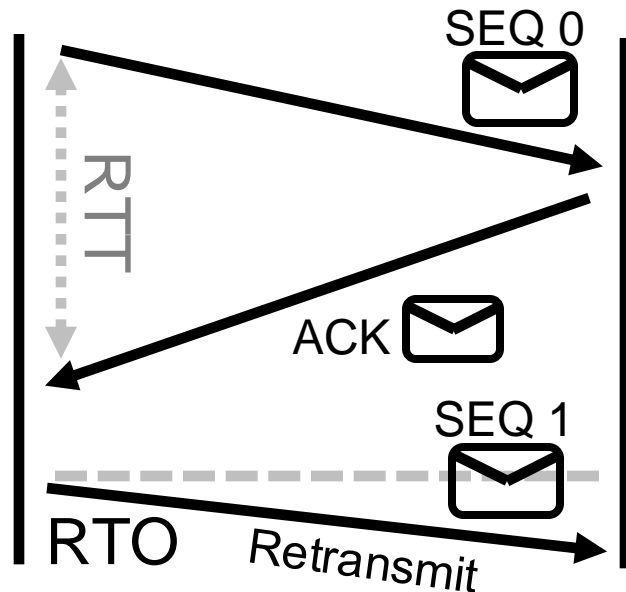
Srinivas Narayana

= window size

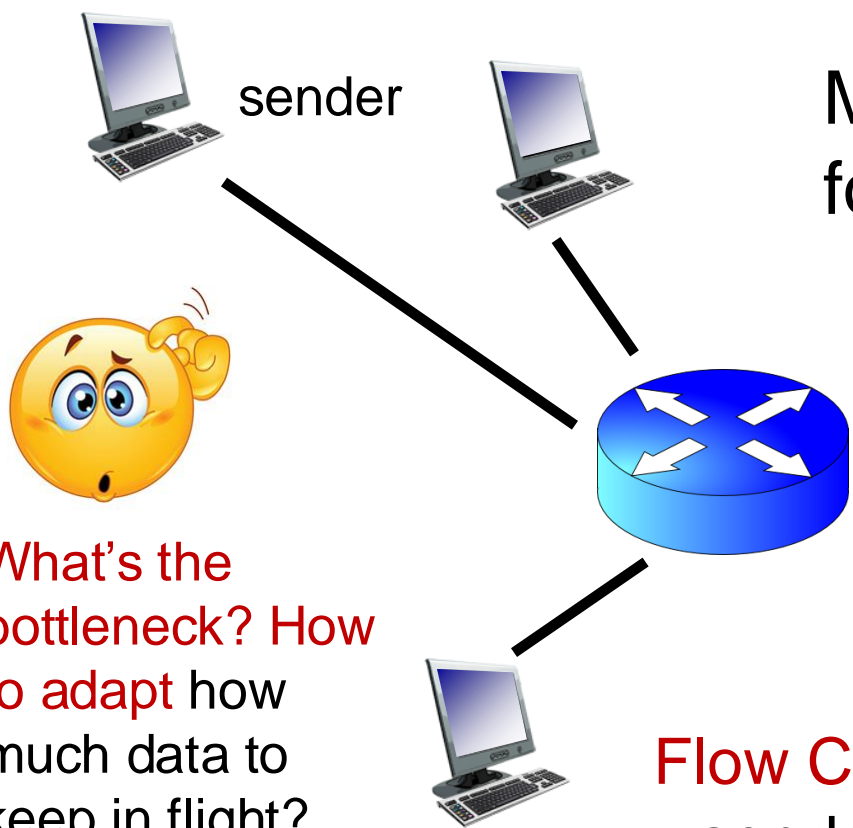
Proportional to **throughput**

# How much data to keep in flight?

## Stop and Wait



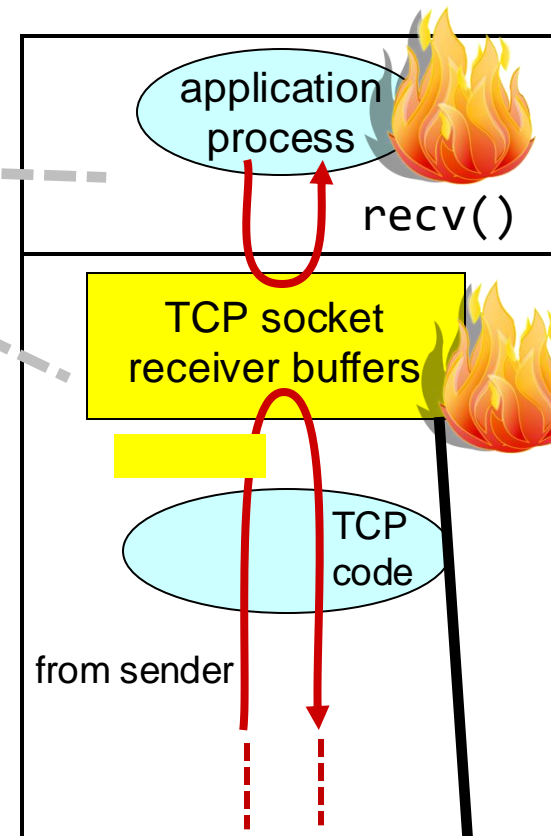
## Pipelined Reliability



Multiple locations for bottlenecks

## Congestion Control

**Flow Control:** Receiver informs sender free buffer over time




Flow Control

0 1 2 3 4 5 6 7 8 9										0 1 2 3 4 5 6 7 8 9										0 1 2 3 4 5 6 7 8 9										0 1 2 3 4 5 6 7 8 9									
Source Port										Destination Port										Sequence Number										Acknowledgment Number									
Data Offset										Reserved										U A P R S F R C S S Y I G K H T N N										Window									
Checksum										Urgent Pointer										Options										Padding									
data																																							

TCP Header Format

Note that one tick mark represents one bit position.

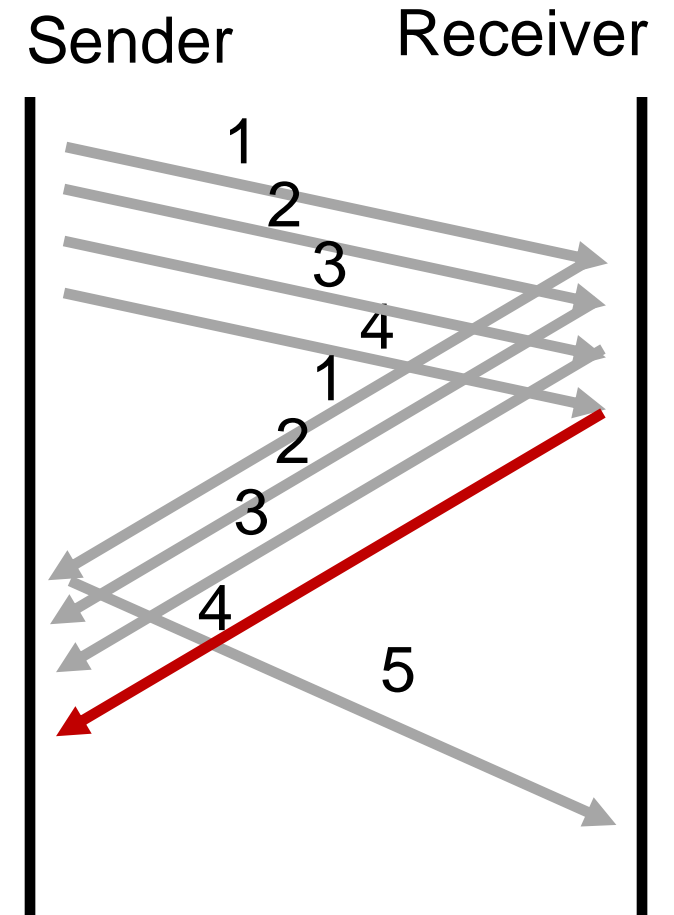
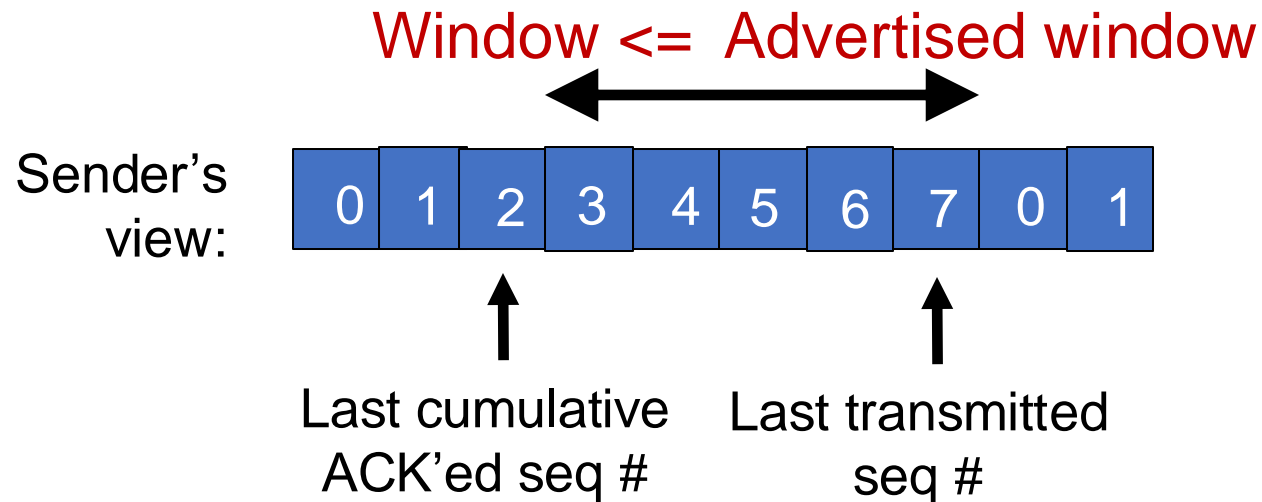


$$\text{Window} \leq \text{Advertised window}$$

How to size this buffer?

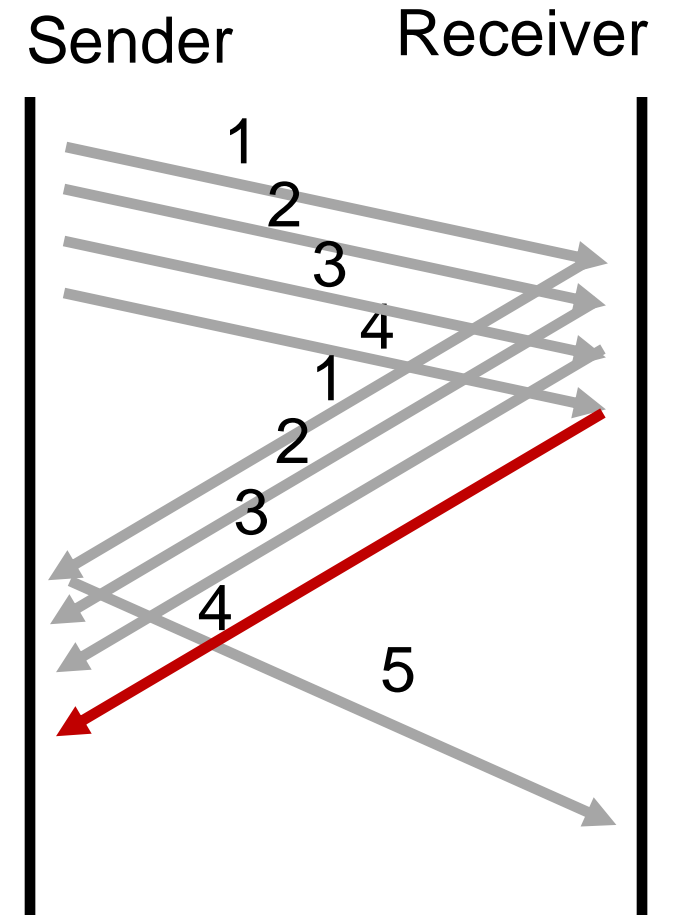
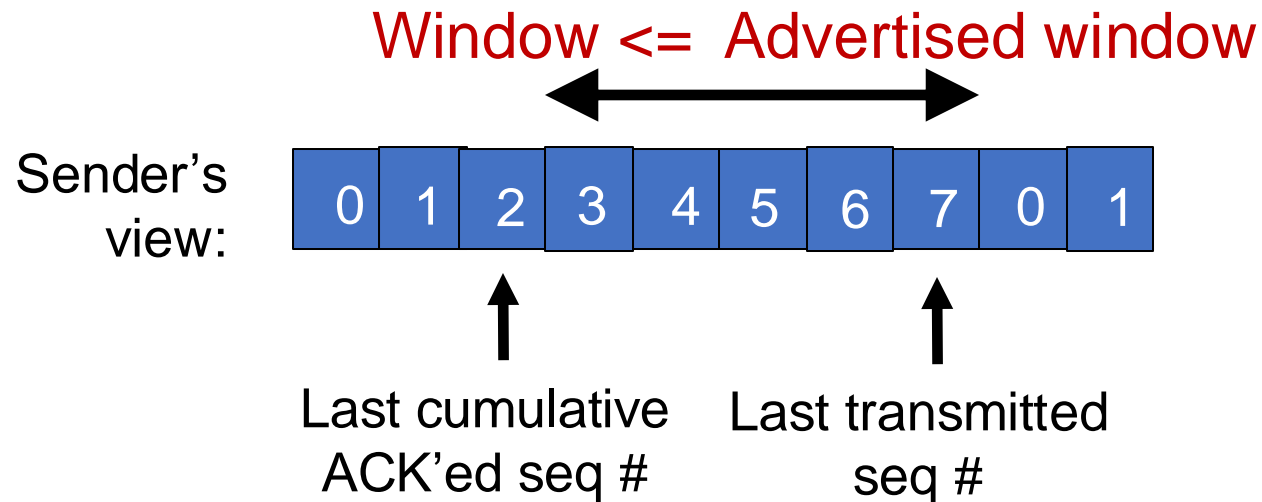
# TCP flow control

- If receiver app is too slow reading data:
  - receiver socket buffer fills up
  - => advertised window shrinks
  - => sender's window (sending rate) reduces
  - => sender's socket buffer fills up
  - => sender process put to sleep upon send()



# TCP flow control

Flow control matches the sending process's write speed to the receiving process's read speed.



# Sizing the receiver's socket buffer

- Operating systems have a default receiver socket buffer size
  - Listed among `sysctl -a | grep net.inet.tcp` on MAC
  - Listed among `sysctl -a | grep net.ipv4.tcp` on Linux
- If socket buffer is too small, sender can't keep too many packets in flight → lower throughput
- If socket buffer is too large, too much memory consumed per socket
- How big should the receiver socket buffer be?

# Sizing the receiver's socket buffer

- Case 1: Suppose the receiving app is reading data too slowly:
- No amount of receiver buffer can prevent low throughput (for a long-lived connection).
- Flow control matches throughput to the receiving app's (low) speed

# Sizing the receiver's socket buffer

- Case 2: Suppose the receiving app reads sufficiently fast *on average* to match the sender's writing speed.
  - Assume the sender desires a window of size  $W$ .
  - The receiver must use a buffer of size at least  $W$ . Why?
- Captures two cases:
- (1) When the first sequence #s in the window are dropped
  - *Selective repeat*: data in window buffered until the "hole" within the window can be filled by the sender. Advertised window reduces sender's window
- (2) When the sender sends a burst of data of size  $W$ 
  - The receiver may not keep up with the *instantaneous* rate of the sender
- Set receiver socket buffer size  $>$  desired window size



# Summary of flow control

- Keep memory buffers available at the receiver whenever the sender transmits data
- Buffers needed to hold for selective repeat, reassembling data in order, and until applications can read data
- Inform available buffer to sender on an ongoing basis, with each ACK
- Function: match sender speed to receiver speed
- **Correct socket buffer sizing is important for TCP throughput**
  - $\text{Throughput} = \text{window size} / \text{RTT} \leq \text{receiver socket buffer} / \text{RTT}$

# Info on (tuning) TCP stack parameters

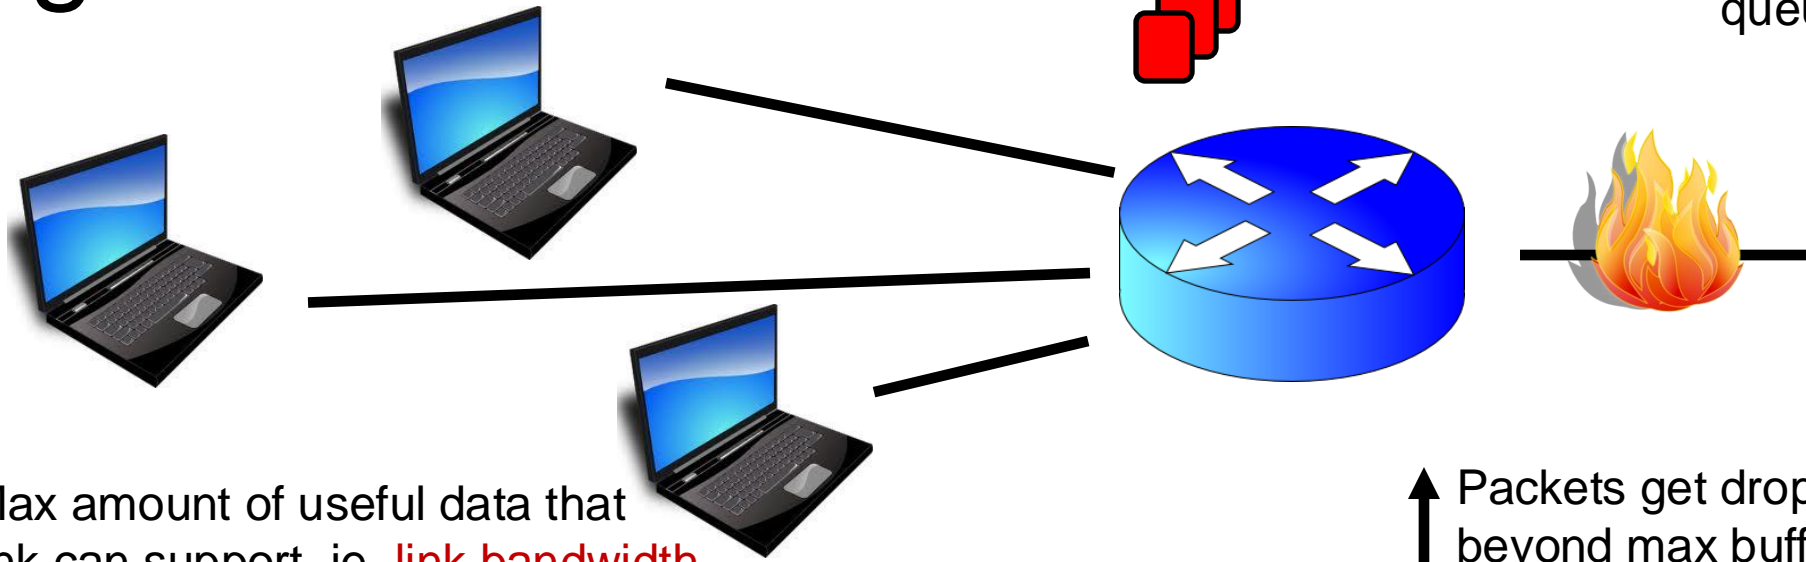
- [https://www.ibm.com/support/knowledgecenter/linuxonibm/liaag/wkvm/wkvm\\_c\\_tune\\_tcpip.htm](https://www.ibm.com/support/knowledgecenter/linuxonibm/liaag/wkvm/wkvm_c_tune_tcpip.htm)
- <https://cloud.google.com/solutions/tcp-optimization-for-network-performance-in-gcp-and-hybrid>

# Playing around with socket buffer sizes

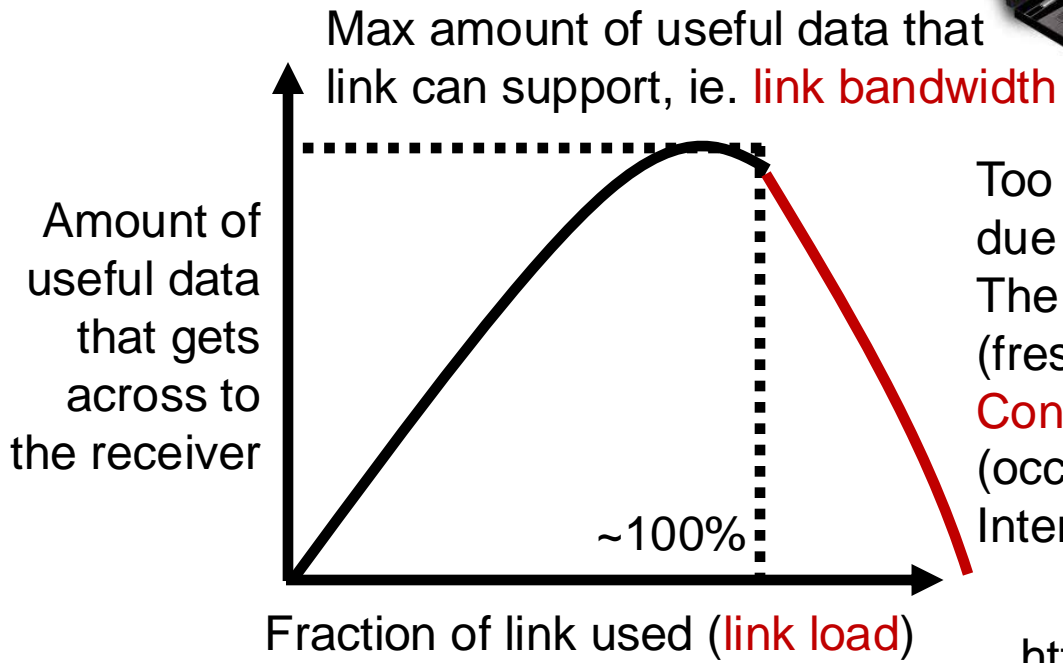
- `iperf -s ; iperf -c localhost -i 1`
- `ping localhost`
- `sudo tc qdisc add dev lo root netem delay 100ms`
- `sudo sysctl net.ipv4.tcp_rmem # min, default, max`
- **Default buffer size 128KB; change e.g., 2.56MB by using**
  - `sudo sysctl net.ipv4.tcp_rmem="4096 2621440 6291456"`
- **Clean up and restore to defaults**
- `sudo tc qdisc del dev lo root netem`
  - **If needed:**
    - `sudo sysctl net.ipv4.tcp_rmem="4096 131072 6291456"`

# Congestion Control

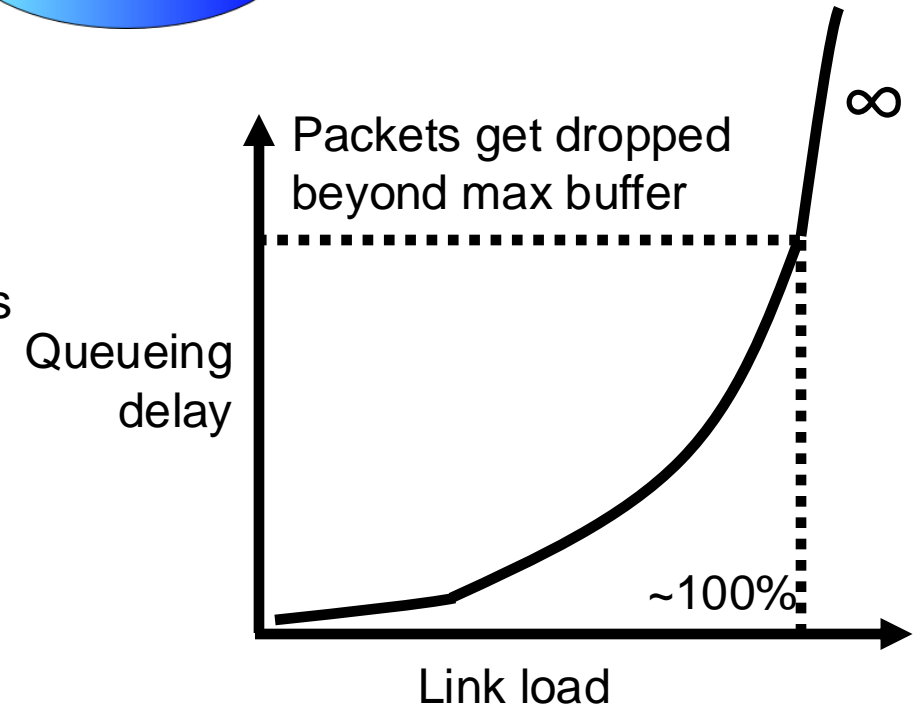
# Congestion



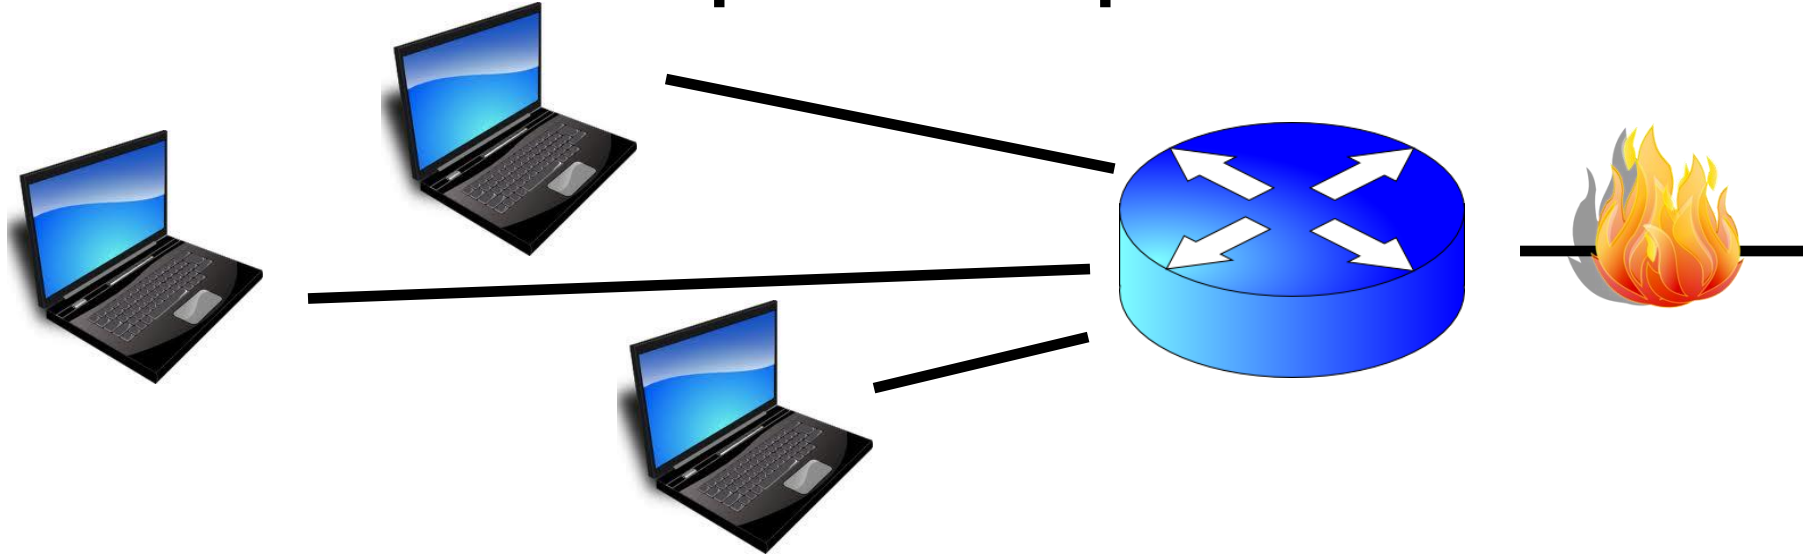
Routers have **buffers** which accommodate queued packets.



Too many retransmissions due to packet drops!  
The amount of useful (fresh) data plummets.  
**Congestion collapse**  
(occurred for real on the Internet in the last 80s)



# How should multiple endpoints share net?



- It is difficult to know where the **bottleneck** link is
- It is difficult to know how many other endpoints are using that link
- Endpoints may join and leave at any time
- Network paths may change over time, leading to different bottleneck links (with different link rates) over time

The approach that the Internet takes is to use a distributed algorithm to converge to an efficient and fair outcome.

The approach that the Internet takes is to use a **distributed algorithm** to converge to an efficient and fair outcome.

No one can centrally view or control all the endpoints and bottlenecks in the Internet.

Every endpoint must try to reach a globally good outcome by itself: i.e., in a distributed fashion.

This also puts a lot of **trust in endpoints**.



The approach that the Internet takes is to use a distributed algorithm to converge to an **efficient** and fair outcome.

If there is spare capacity in the bottleneck link, the endpoints should use it.

The approach that the Internet takes is to use a distributed algorithm to converge to an efficient and **fair** outcome.

If there are  $N$  endpoints sharing a bottleneck link, they should be able to get **equitable** shares of the link's capacity.

For example:  $1/N$ 'th of the link capacity.

# Flow Control vs. Congestion Control

- Avoid overwhelming the receiving application
- Sender is managing the receiver's socket buffer

- Avoid overwhelming the bottleneck network link
- Sender is managing the bottleneck link capacity and bottleneck router buffers

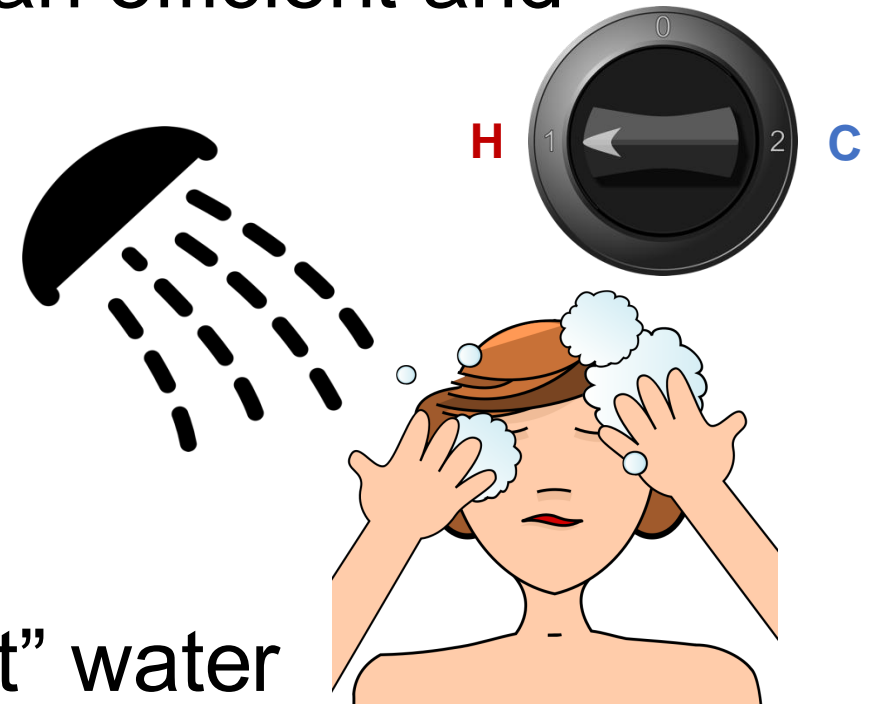
The approach that the Internet takes is to use a distributed algorithm to converge to an efficient and fair outcome.

How to achieve this?

**Approach: sense and react**

Example: showering: Want “just right” water

Use a **feedback loop** with signals and knobs



# Signals and Knobs in Congestion Control

- **Signals**

- Packets being ACK'ed
- Packets being dropped (e.g. RTO fires)
- Packets being delayed (RTT)
- Rate of incoming ACKs

} **Implicit** feedback signals  
measured directly at sender.  
(There are also explicit signals  
that the network might provide.)

- **Knobs**

- What can you change to “probe” the available bottleneck capacity?
  - **Window size**
- Suppose the receiver socket buffer size is unbounded
  - **Congestion window**: window size used for congestion control
- Increase window/sending rate: e.g., add x or multiply by a factor of x
- Decrease window/sending rate: e.g., subtract x or reduce by a factor of x

# Sense and react, sure...but how?

- Where do you want to be?
  - The **steady state**
- How do you get there?
  - Congestion control algorithms
- Sense accurately & react accordingly

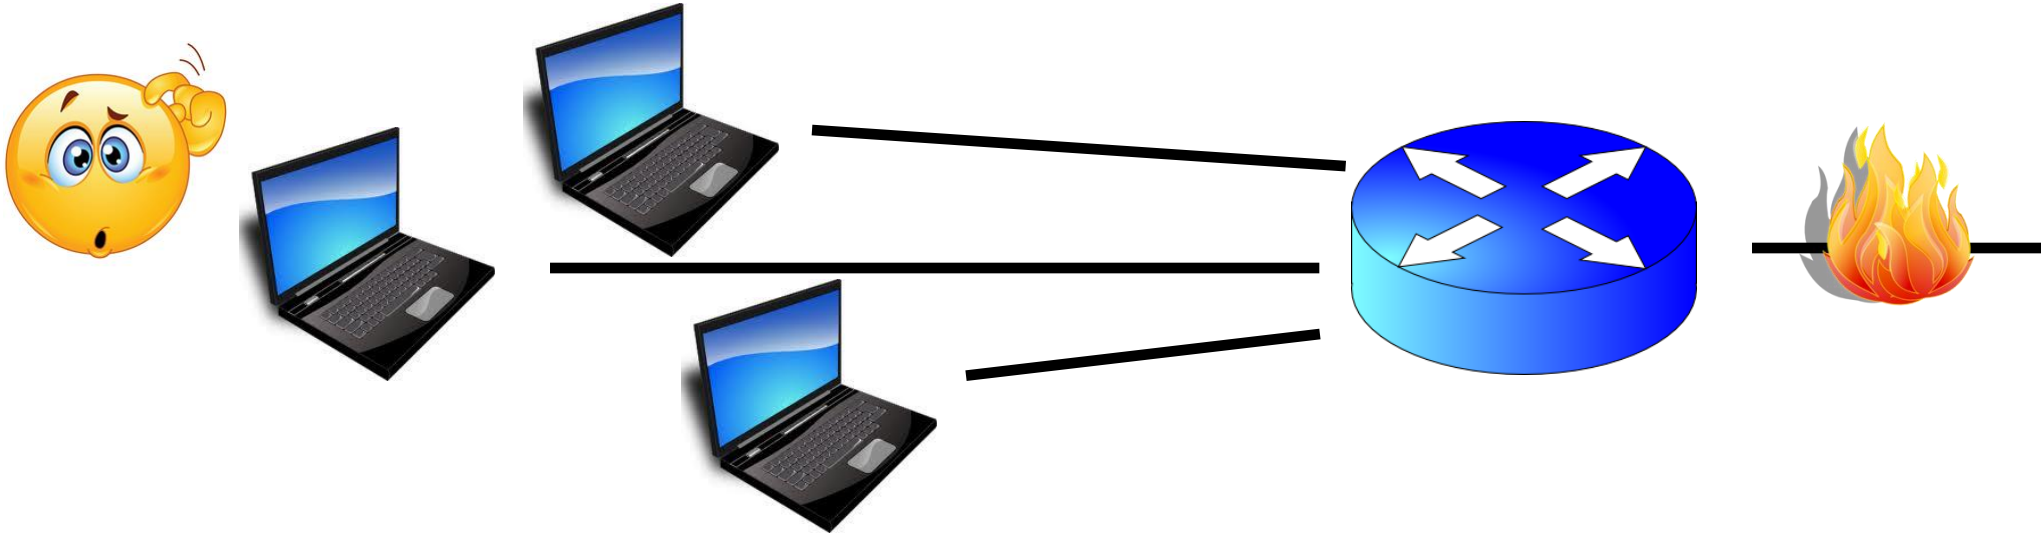


# The Steady State

Efficiency for a single TCP connection

# What does **efficiency** look like?

- Suppose we want to achieve an **efficient** outcome for **one** TCP connection by observing network signals from the endpoint



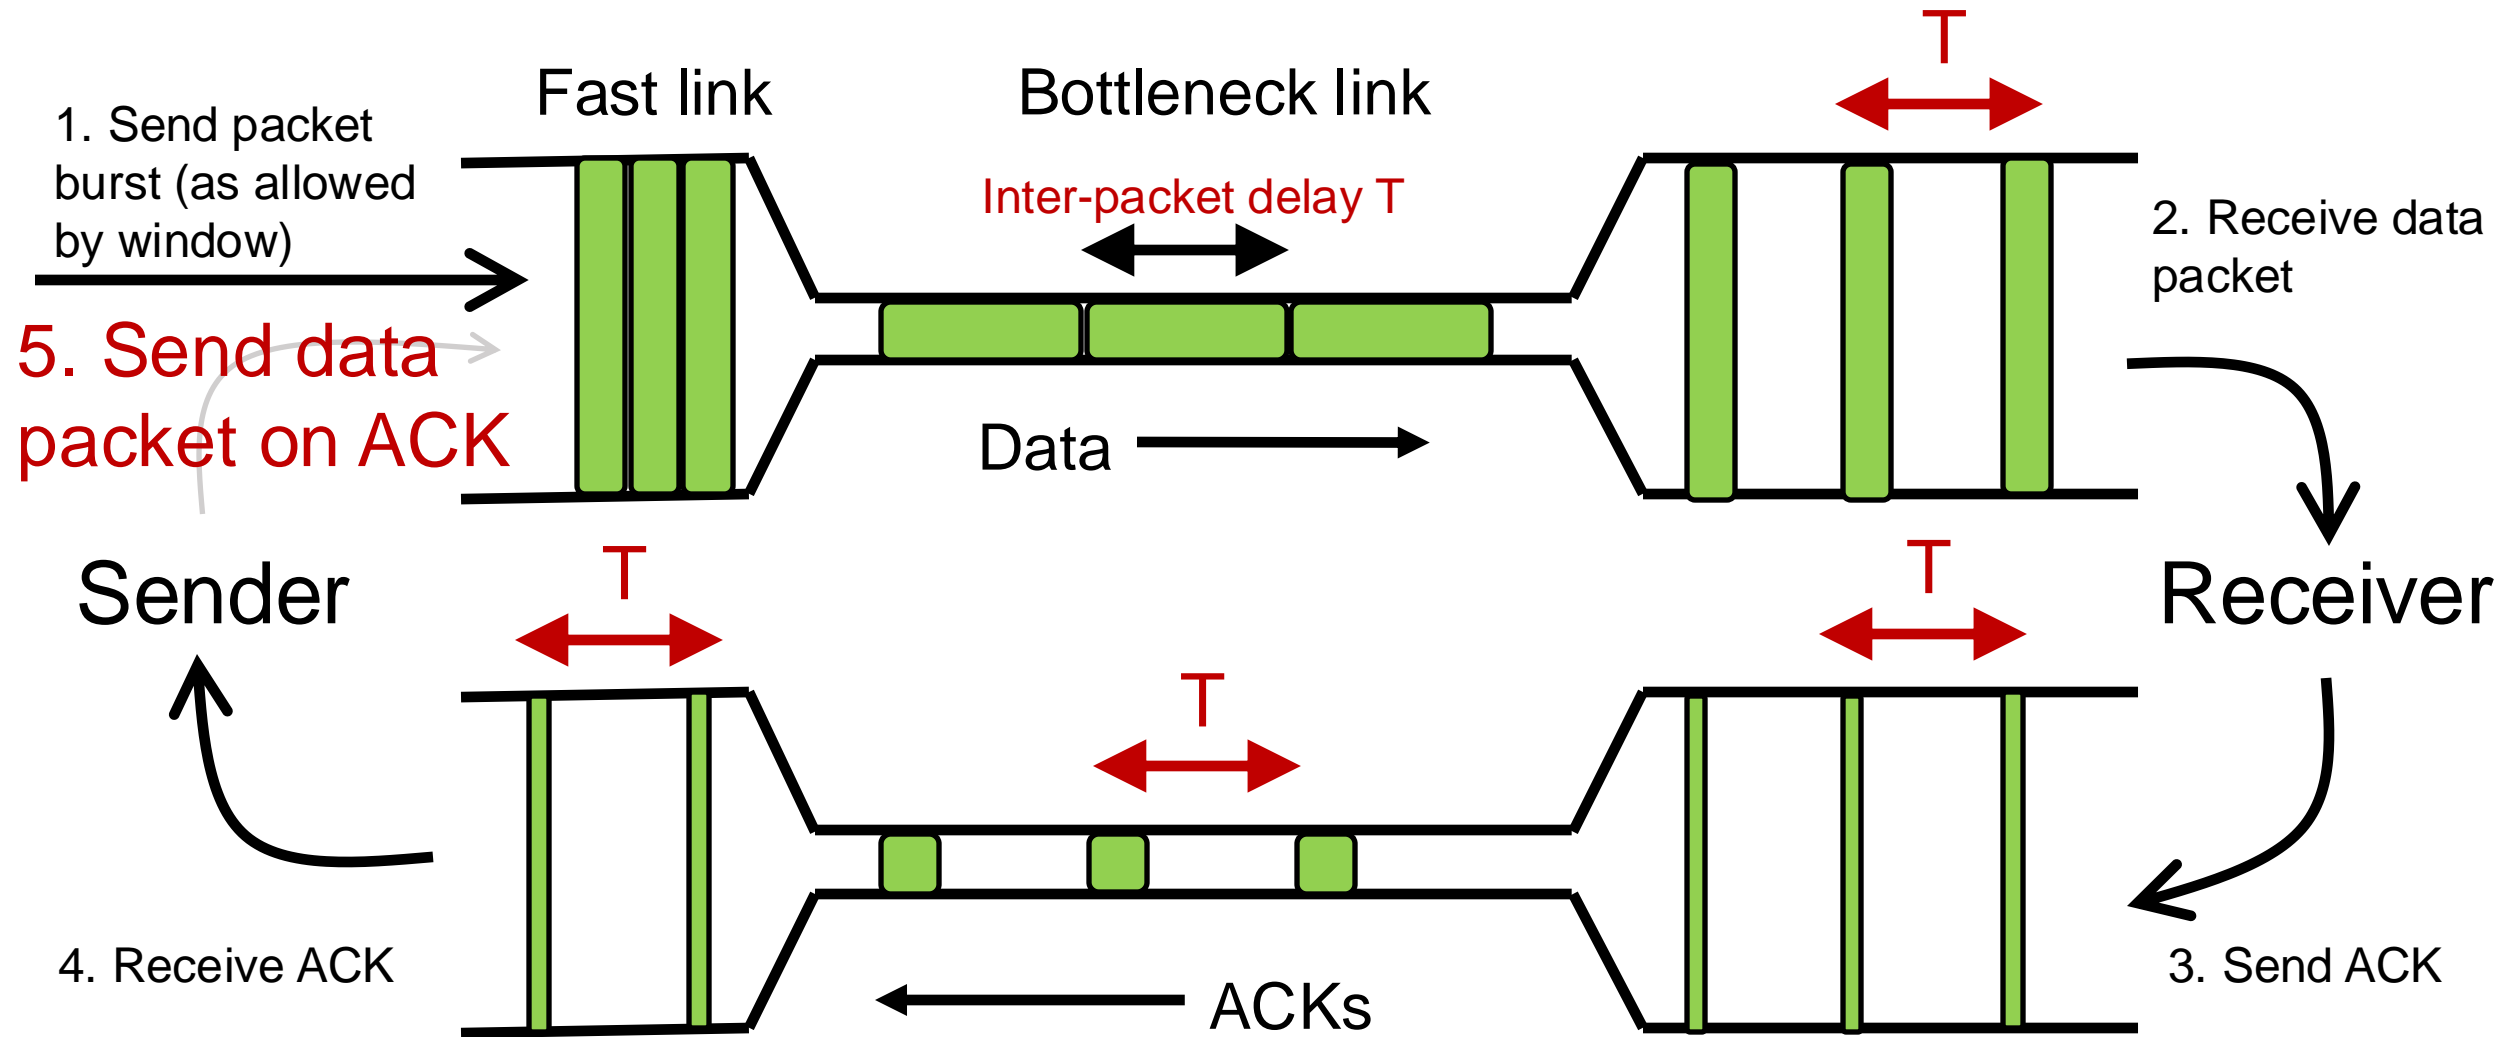
- Q: How should the endpoint behave **at steady state**?
- Challenge: bottleneck link is remotely located



# Steady state: Ideal goal

- **High sending rate:** Use the full capacity of the bottleneck link
- **Low delay:** Minimize the overall delay of packets to get to the receiver
  - Overall delay = propagation + queueing + transmission
  - Assume propagation and transmission components fixed
- “Low delay” reduces to **low queueing delay**
- i.e., don’t push so much data into the network that packets have to wait in queues
- Key question: When to send the next packet?

# When to send the next packet?



# Rationale

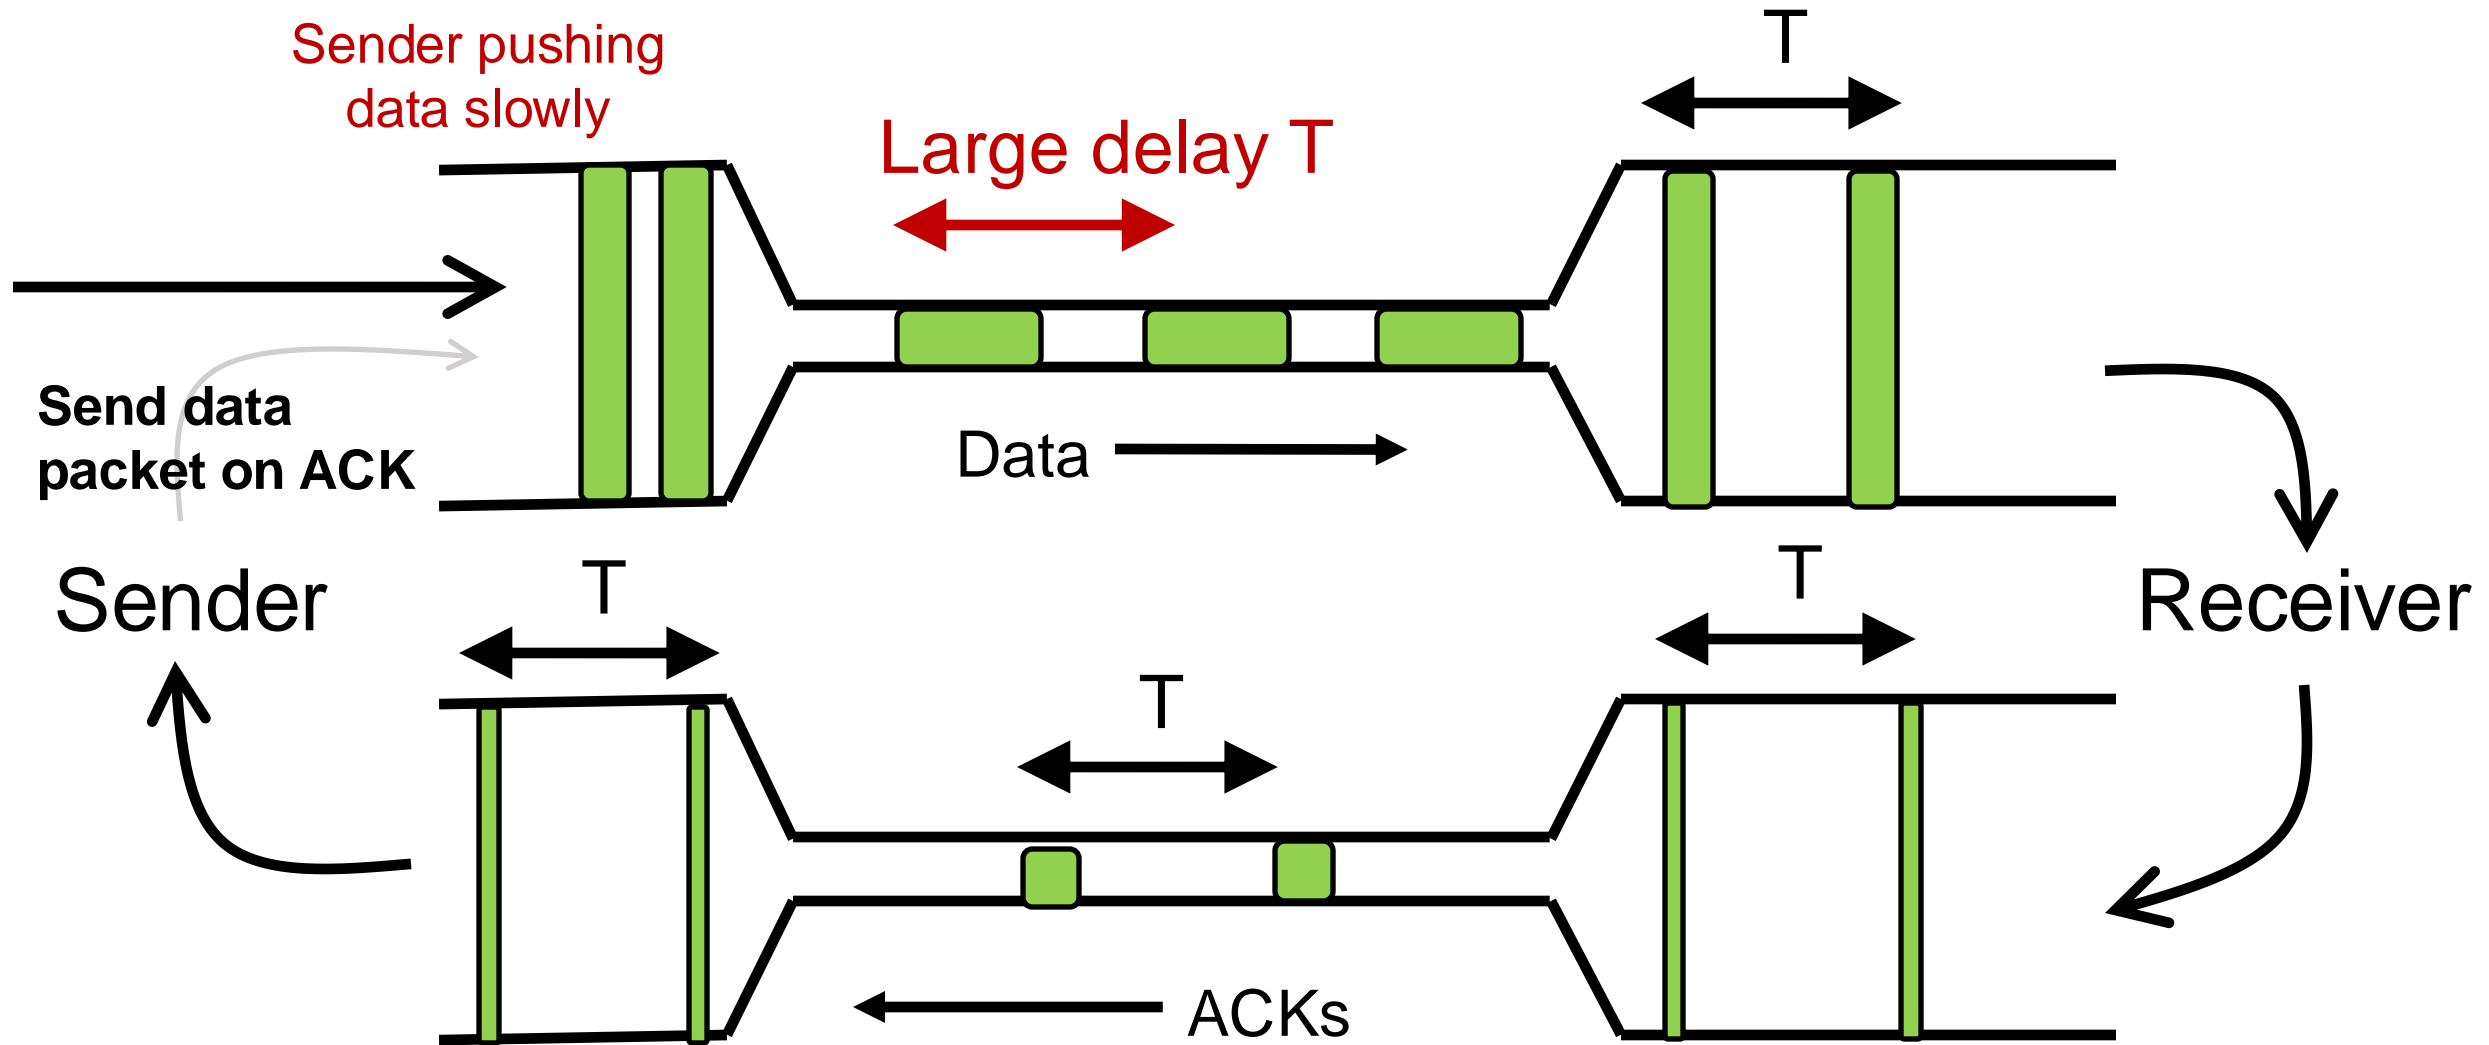
- When the sender receives an ACK, that's a signal that the previous packet has left the bottleneck link (and the rest of the network)
- Hence, it must be safe to send another packet without congesting the bottleneck link
- Such transmissions are said to follow packet conservation
- ACK clocking: “Clock” of ACKs governs packet transmissions

# ACK clocking: analogy

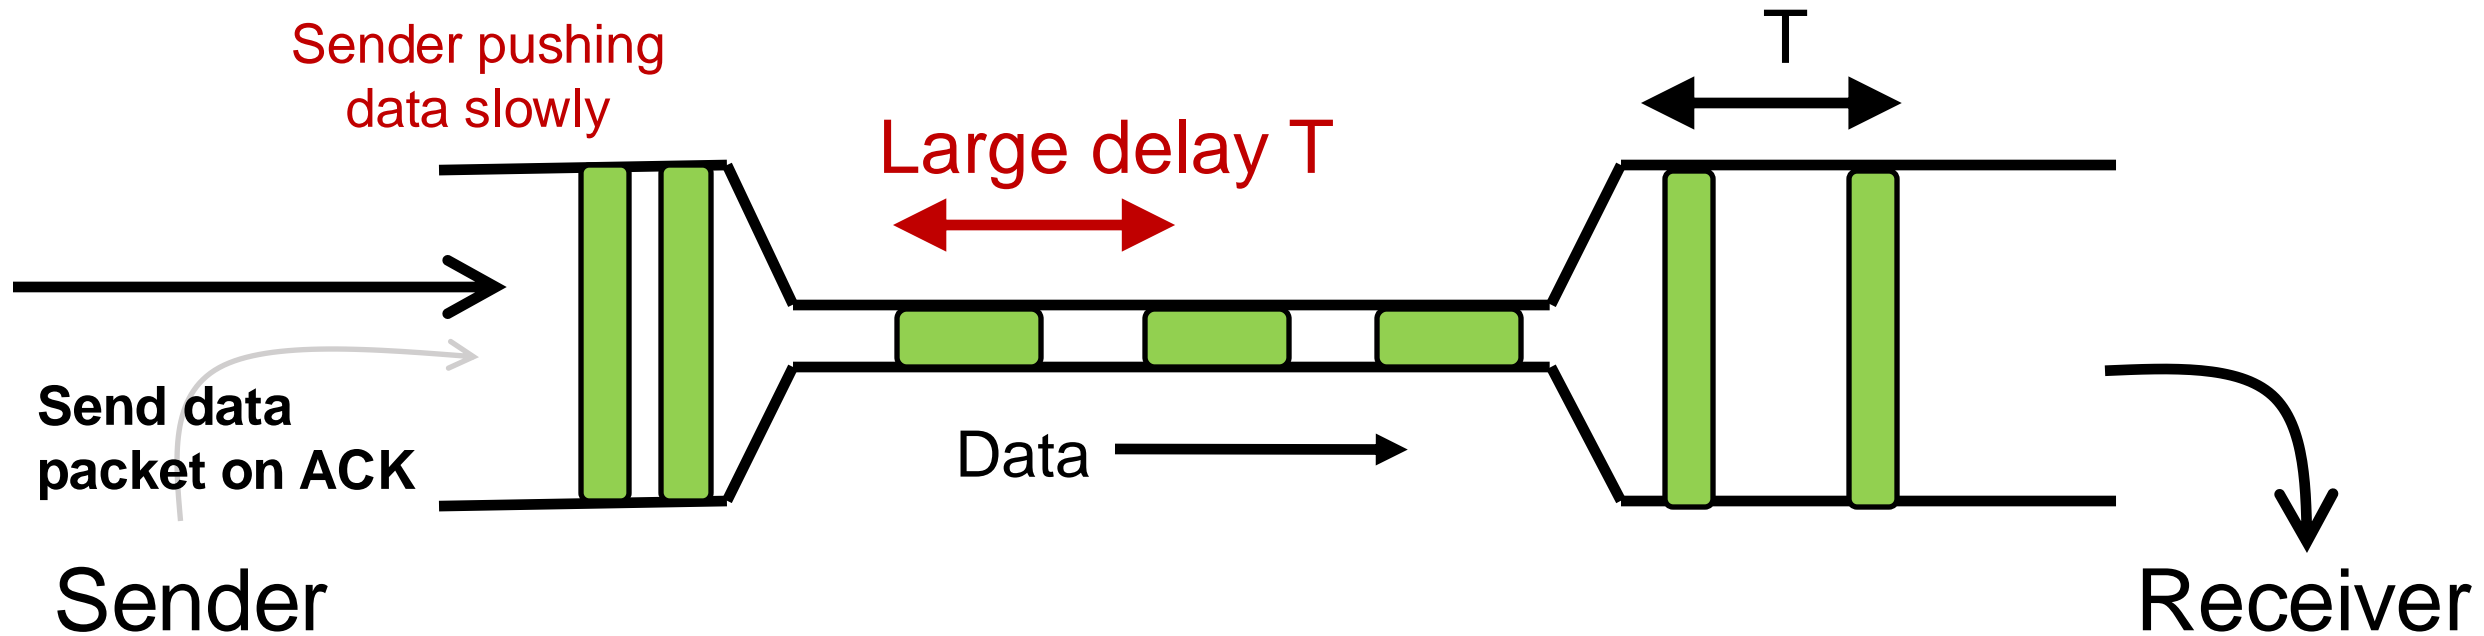
- How to avoid crowding a grocery store?
- Strategy: Send the next waiting customer exactly when a customer exits the store
- However, this strategy alone can lead to inefficient use of resources...



# ACK clocking alone can be inefficient



# ACK clocking alone can be inefficient



The sending rate should be high enough to keep the “pipe” full

Analogy: a grocery store with only 1 customer in entire store

If the store isn't “full”, you're using store space inefficiently

# Steady State of Congestion Control

- Send at the highest rate possible (to keep the pipe full)
- while being ACK-clocked (to avoid congesting the pipe)
- So, how to get to steady state?

# Finding the Right Congestion Window



# Let's play a game

- Suppose I'm thinking of a number (positive integer). You need to guess the number I have in mind.
- Each time you guess, I will tell you whether your number is smaller or larger than (or the same as) the one I'm thinking of
- My number can be very large or small
- How would you go about guessing the number?

# Finding the right congestion window

- TCP congestion control algorithms solve a similar problem!
- There is an **unknown** bottleneck link rate that the sender must match
- If sender sends more than the bottleneck link rate:
  - packet loss, delays, etc.
- If sender sends less than the bottleneck link rate:
  - all packets get through; successful ACKs

# Quickly finding a rate: TCP slow start

- Initially  $cwnd = 1 \text{ MSS}$

- MSS is “maximum segment size”

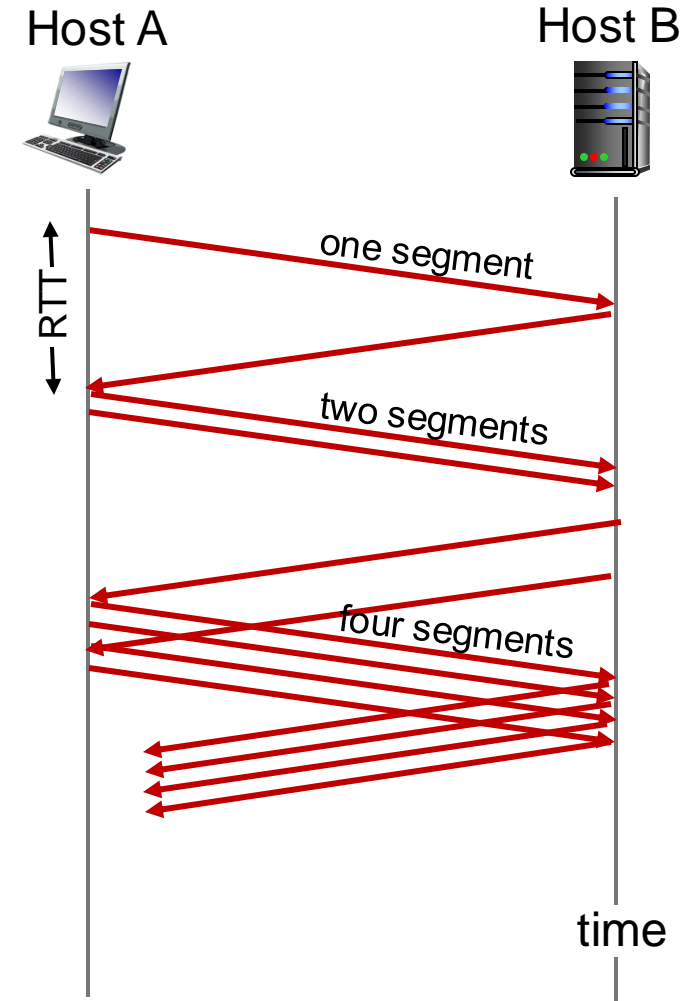


- Upon receiving an ACK of each MSS, increase the  $cwnd$  by 1 MSS

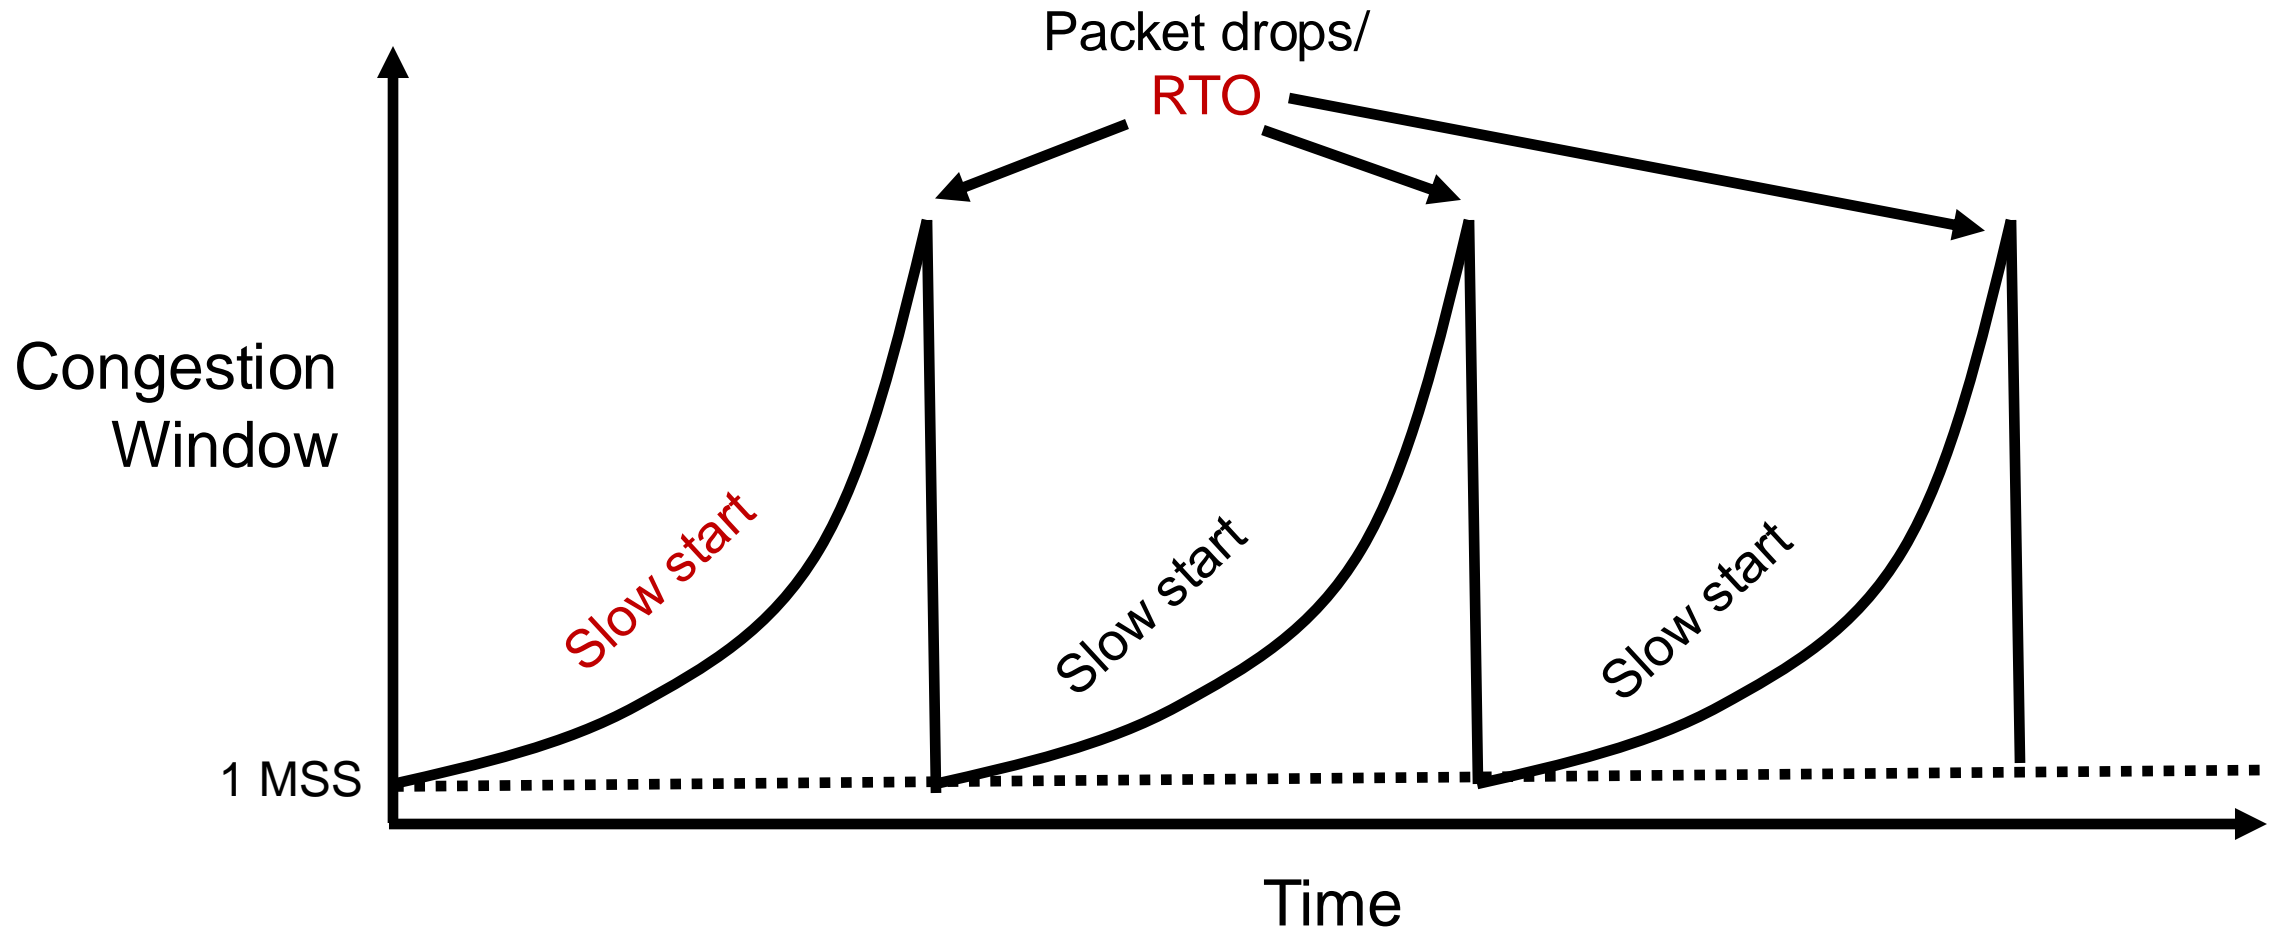
- Effectively, double  $cwnd$  every RTT

- Initial rate is slow but ramps up **exponentially fast**

- On loss (RTO), restart from  $cwnd := 1 \text{ MSS}$



# Behavior of slow start



# Slow start has problems

- Congestion window **increases too rapidly**
  - Example: suppose the “right” window size `cwnd` is 17
  - `cwnd` would go from 16 to 32 and then dropping down to 1
  - Result: massive packet drops
- Congestion window **decreases too rapidly**
  - Suppose the right `cwnd` is 31, and there is a loss when `cwnd` is 32
  - Slow start will resume all the way back from `cwnd` 1
  - Result: unnecessarily low throughput
- Instead, perform **finer adjustments** of `cwnd` based on signals

# Use slow start mainly at the beginning

- You might accelerate your car a lot when you start, but you want to make only small adjustments after.
  - Want a smooth ride, not a jerky one
- Slow start is a good algorithm to get close to the bottleneck link rate when there is little info available about the bottleneck, e.g., the beginning of a connection
- Once close enough to the bottleneck link rate, use a different set of strategies to perform smaller adjustments to the congestion window
  - Called TCP congestion avoidance