

The Transport Layer: De/Multiplexing, Reliability

CS 352, Lecture 6, Spring 2020

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

Course announcements

- Quiz 2 will go online later today
 - Submit on Sakai by Tuesday 10 pm
 - 30-minute time limit
- Project 1 will be released later today
 - Check the Sakai resources section for a tarfile with instructions

Review of concepts

- **Application-layer protocols:** DNS, HTTP, SMTP
- HTTP caching: why?
- Content distribution networks (CDNs):
 - Origin server, CDN DNS server, CDN cache servers, client
 - Use **indirection**
- Simple Mail Transfer Protocol (SMTP): req/resp protocol
 - User agents, sending mail server, receiving mail server
 - SMTP **commands**, mail **headers**, response codes
 - Mail access protocols: POP, IMAP, **HTTP**
 - **Support pull rather than push**

POP vs IMAP


- POP3
- Stateless server
- UA-heavy processing
- UA retrieves email from server, then typically deleted from server
- Latest changes are at the UA
- Simple protocol (list, retr, del within a POP session)

- IMAP4
- Stateful server
- UA and server processing
- Server sees folders, etc. which are visible to UAs
- Changes visible at the server
- Complex protocol

What about web-based email?

- Connect to mail servers via web browser
 - Ex: gmail, outlook, etc.
- Browsers speak HTTP
- Email servers speak SMTP
- Need a bridge to retrieve email using HTTP

Web based email





More themes from app-layer protocols

- **Separation of concerns.** Examples:
 - Content rendering for users (browser, UA) separate from protocol operations (mail server)
 - Reliable mail sending and receiving: mail UA doesn't need to be “always on” to send or receive email reliably
- **In-band vs. out-of-band control:**
 - In-band: headers determine the actions of all the parties of the protocol
 - There are protocols with out-of-band control, e.g., FTP
- **Keep it simple until you really need complexity**
 - ASCII-based design; stateless servers. Then introduce:
 - Cookies for HTTP state
 - IMAP for email organization
 - Security extensions
 - Different methods to set up and use underlying connections, etc.


The Transport Layer

Transport



Transport services and protocols

- Provide **logical communication** between app processes running on different hosts
- Transport protocols run @ hosts
 - send side: breaks app messages into **segments**, passes to network layer
 - recv side: reassembles segments into messages, passes to app layer
- More than one transport protocol available to apps
 - Internet: **TCP** and **UDP**



Transport vs. network layer

- **Network layer:** logical communication between **endpoints**
- **Transport layer:** logical communication between **processes**
 - relies on and enhances network layer services

Household analogy:

12 kids sending letters to 12 kids

- processes = kids
- app messages = letters in envelopes
- endpoints = houses
- transport protocol = Alice and Bob who de/mux to in-house siblings
- network-layer protocol = postal service



What transport guarantees do you want?

- **Reliability**

- Don't drop messages
- Don't corrupt messages

- **Performance**

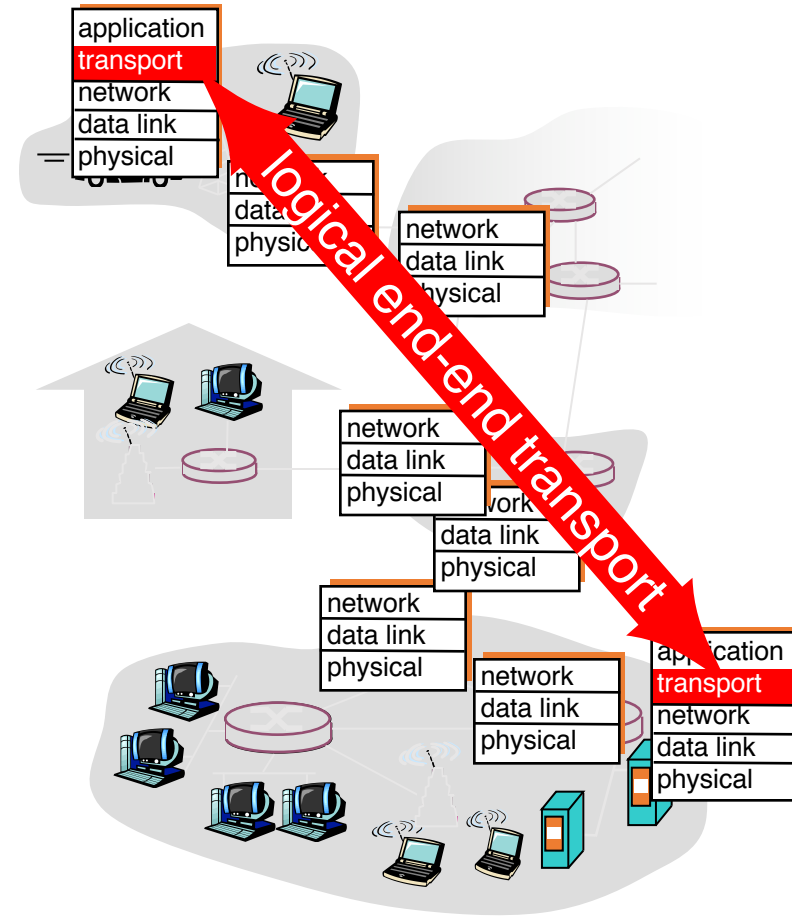
- Get messages to other side as soon as possible
- As cheaply as possible

- **Ordering**


- Don't reorder messages

Internet transport-layer protocols

- Reliable, in-order delivery: Transmission Control Protocol (**TCP**)
 - congestion control, flow control, connection setup
- Unreliable, unordered delivery: User Datagram Protocol (**UDP**)
 - A small extension to underlying network layer protocol, IP
- Services not available:
 - delay guarantees
 - bandwidth guarantees



Layering: in terms of packets



User Datagram Protocol

UDP: User Datagram Protocol [RFC 768]

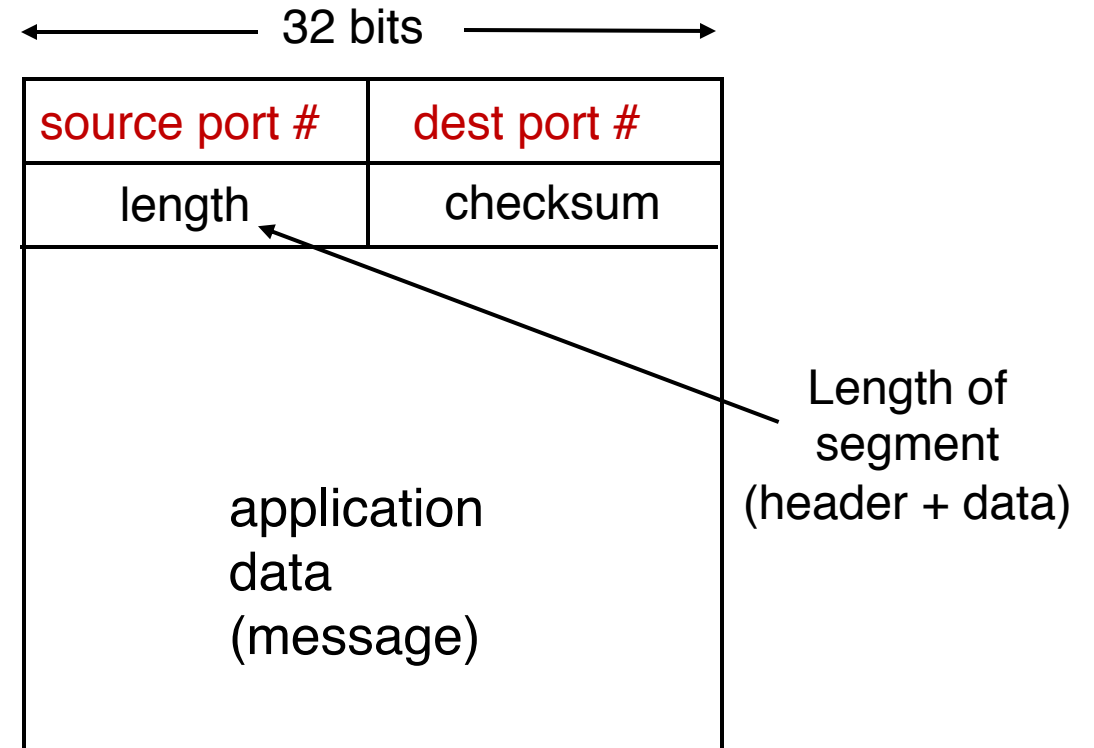
- **Best effort service.** UDP segments may be:
 - Lost
 - Delivered out of order to app
- UDP is **connectionless**
 - Each UDP segment handled **independently** of others (i.e. no “memory” across packets)
- DNS uses UDP
 - UDP suitable for one-off req/resp
 - Also for loss-tolerant delay-sensitive apps, e.g., video calling

Why are UDP's guarantees even okay?

- **Simple & Low overhead**
- No delays due to connection establishment
 - UDP can send data immediately
- No memory for connection state at sender & receiver
- Small segment header
- UDP can blast away data as fast as desired
 - UDP has no congestion control

How demultiplexing works

- Host receives IP datagrams
 - Datagram contains a transport-level **segment**
 - each segment has source IP address, destination IP address
 - each segment has source, destination port number
- Host uses IP addresses & port numbers to direct segment to appropriate app-level **socket**



UDP segment format

Time for an activity

Connectionless demultiplexing

- Create sockets with endpoint-local port numbers to receive data

```
// Example: Python UDP socket
sock = socket(AF_INET, SOCK_DGRAM);
// can bind sock to specific local port
```

- When creating data to send into a UDP socket, you must specify the remote IP address and port.

```
sock.sendto(msg, ("128.1.2.3", 4500));
```

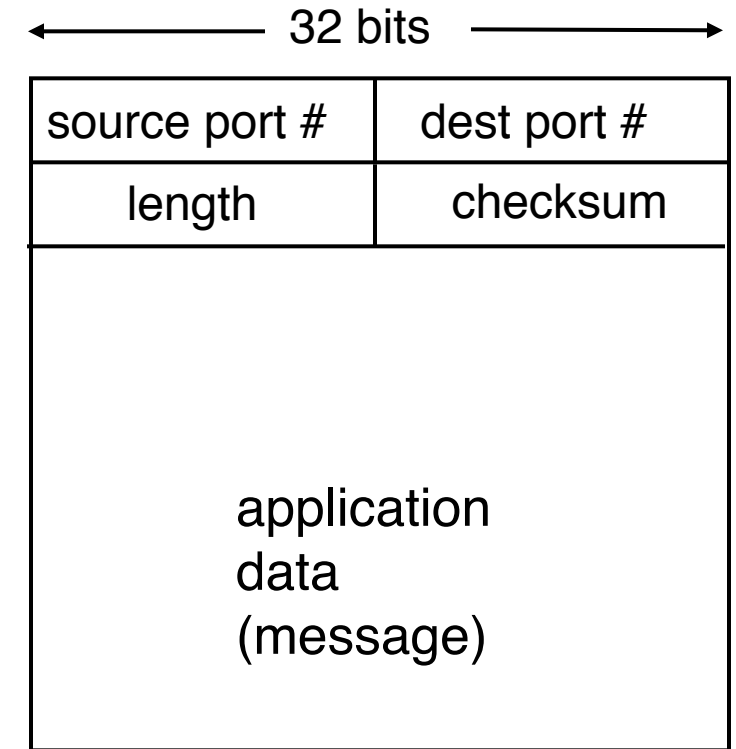
- When endpoint receives UDP segment:
 - Lookup a table with the key (destination IP, destination port) to check if a corresponding socket exists
 - Direct UDP segment to socket
- Datagrams with different source IP addresses and/or source port numbers are directed to the same socket
 - Subtle difference from TCP demultiplexing.

Error Detection

Necessary, but insufficient, for reliability

Data may get corrupted along the way...

- Bits flipped from $0 \rightarrow 1$ or $1 \rightarrow 0$
- Packet bits lost
- How to detect errors?
- Idea: compute a function over the data
 - Store the result along with the data
- Function must be **easy to compute**
- Function & stored data efficient to **verify**
- Ideas for functions?



UDP segment format

From the UDP specification (RFC 768)

- Checksum is the 16-bit one's complement of the one's complement sum of a pseudo header of information from the IP header, the UDP header, and the data, padded with zero octets at the end (if necessary) to make a multiple of two octets.
- The pseudo header conceptually prefixed to the UDP header contains the source address, the destination address, the protocol, and the UDP length.

UDP Checksum

Sender:

- treat segment contents as sequence of 16-bit integers
- checksum: addition (1's complement sum) of segment contents
- sender puts checksum value into UDP checksum field

Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
- NO - error detected
- YES - no error detected.

UDP checksum example

- Note: when adding numbers, a carryout from the most significant bit needs to be added to the result
- Example: add two 16-bit integers

	1 1 1 0 0 1 1 0 0 1 1 0 0 1 1 0
	1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1
	<hr/>
wraparound	1 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1
	<hr/>
sum	1 0 1 1 1 0 1 1 1 0 1 1 1 1 0 0
checksum	0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 1

User Datagram Protocol

- A thin shim around best-effort network delivery
 - Lightweight to send one-off request/response messages
 - Lightweight transport for loss-tolerant delay-sensitive applications
- Provides basic **multiplexing/demultiplexing** for applications
- No reliability, performance, or ordering guarantees
- Can do basic error detection (bit flips) using checksums
 - Error detection is a necessary condition for reliability
 - But need to do much more for reliability. Subject of the next lecture.

