

# The Transport Layer: Reliability

CS 352, Lecture 7, Spring 2020

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Course announcements

- Project 1 released. Need partners?
- Quiz 2 completed, no quiz 3
- **Mid-term 1** next Wednesday
  - In class, 1 h 20 m
  - Covers lectures 1 through 8 (coming Friday)
  - Closed book
  - Only calculators allowed. **No cell phones**
- Pattern: multiple choice, reasoning, problem solving
  - A review will be released later this week

# Review of concepts

- Mail access protocols: POP, IMAP, HTTP
- **Transport-layer protocols:** UDP, TCP
- Support communication between **processes**
  - vs. network layer: **endpoints**
- Transport guarantees: **reliability, ordering, performance**
  - You get two of these at most
- **Demultiplexing:** map packet to app-level socket
- User Datagram Protocol (**UDP**): thin wrapper around net layer
  - Connectionless, best-effort service, simple
  - Demultiplex using **dst IP address, dst port**
  - Who uses UDP?

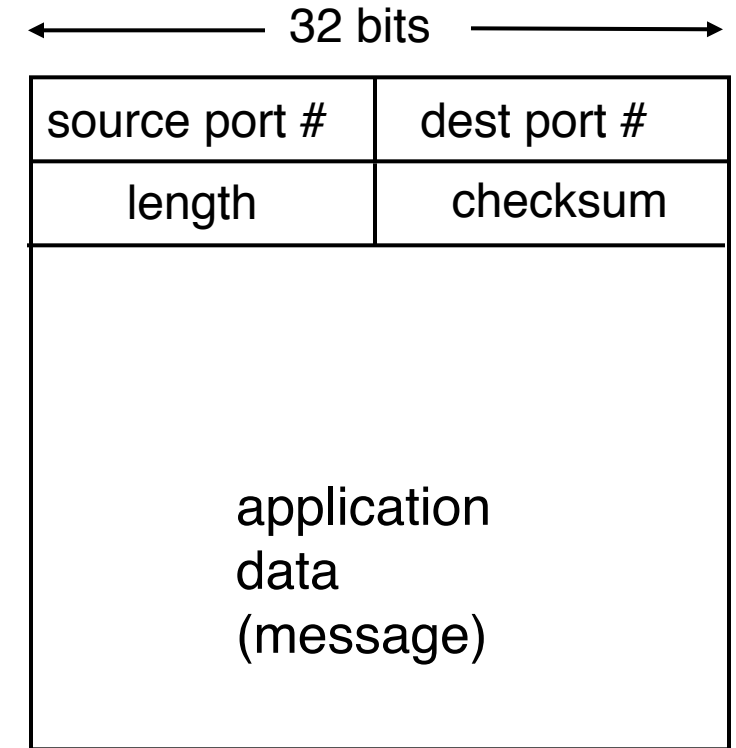


# Error Detection

Necessary, but insufficient, for reliability

# Data may get corrupted along the way...

- Bits flipped from  $0 \rightarrow 1$  or  $1 \rightarrow 0$
- Packet bits lost
- How to detect errors?
- Idea: compute a function over the data
  - Store the result along with the data
- Function must be **easy to compute**
- Function & stored data efficient to **verify**
- Ideas for functions?



UDP segment format

# From the UDP specification (RFC 768)

- Checksum is the 16-bit one's complement of the one's complement sum of a pseudo header of information from the IP header, the UDP header, and the data, padded with zero octets at the end (if necessary) to make a multiple of two octets.
- The pseudo header conceptually prefixed to the UDP header contains the source address, the destination address, the protocol, and the UDP length.

# UDP Checksum

## Sender:

- treat segment contents as sequence of 16-bit integers
- checksum: addition (1's complement sum) of segment contents
- sender puts checksum value into UDP checksum field

## Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
- NO - error detected
- YES - no error detected.

# UDP checksum example

- Note: when adding numbers, a carryout from the most significant bit needs to be added to the result
- Example: add two 16-bit integers

	1 1 1 0 0 1 1 0 0 1 1 0 0 1 1 0
	1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1
	<hr/>
wraparound	1 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1
	<hr/>
sum	1 0 1 1 1 0 1 1 1 0 1 1 1 1 0 0
checksum	0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 1




# User Datagram Protocol

- A thin shim around best-effort network delivery
  - Lightweight to send one-off request/response messages
  - Lightweight transport for loss-tolerant delay-sensitive applications
- Provides basic **multiplexing/demultiplexing** for applications
- No reliability, performance, or ordering guarantees
  - **Need Transmission Control Protocol (TCP)**
- Can do basic error detection (bit flips) using checksums
  - Error detection is a necessary condition for reliability
- But need to do a lot more to achieve reliable data delivery
  - Subject of the rest of this lecture
  - Mechanisms in general; some of them implemented in TCP

# Reliable data delivery

Stop and Wait


# Packet loss



- How might a sender and receiver ensure that data is delivered reliably (despite some packets being lost)?
- TCP uses three mechanisms


# Coping with packet loss: (1) ACK

- Key idea: Receiver returns an **acknowledgment** (ACK) per packet sent
- If sender receives an ACK, it knows that the receiver got the packet.




# Coping with packet **corruption**: (1) ACK

- ACKs also work to detect packet corruption on the way to the receiver
  - A receiver could send a negative acknowledgment, or a **NAK**, if it receives a corrupted packet
- TCP only uses positive acknowledgments (ACKs).
- What if a packet was lost and ACK never arrives?




# Coping with packet loss: (2) RTO

- Key idea: Wait for a duration of time (called **retransmission timeout** or RTO) before **re-sending** the packet
- In TCP, **the onus is on the sender** to retransmit lost data when ACKs are not received
- Note that retransmission works also if **ACKs are lost** or delayed



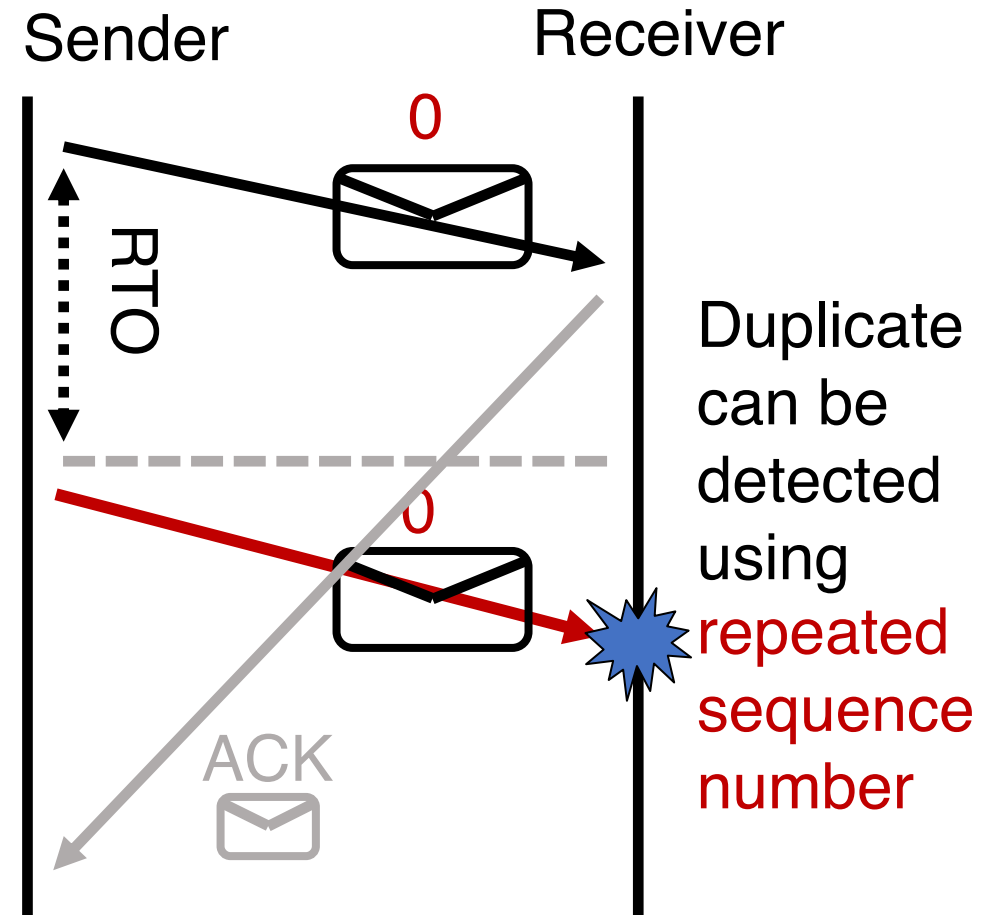
# Coping with packet loss

- What if ACKs are delayed?
  - Sender may retransmit the same data
  - Receiver wouldn't know that it just received duplicate data from this retransmitted packet
- Add some identification to each packet to help distinguish between adjacent transmissions
  - This is known as the **sequence number**



# Coping with packet loss: (3) Seq nums

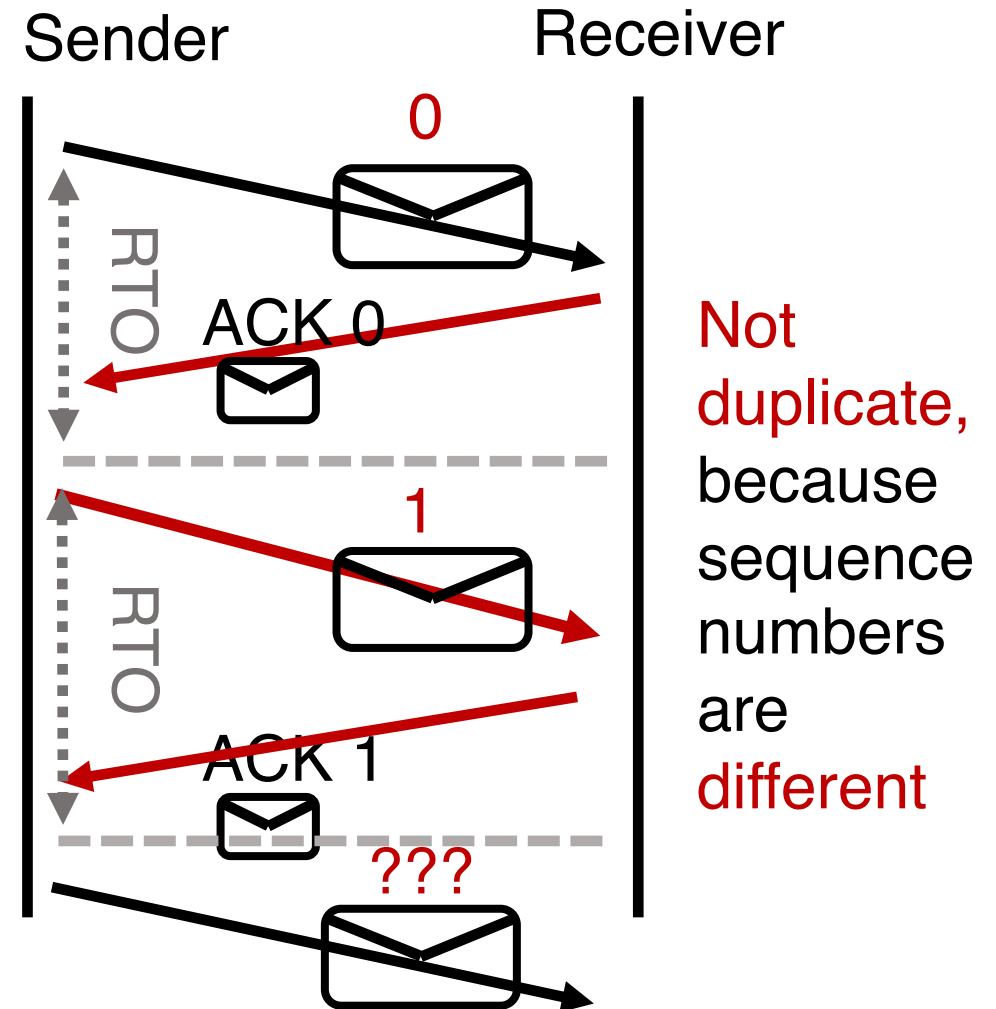
- Bad cases: packet dropped or ACK dropped or ACK delayed beyond RTO
- Sequence numbers of adjacent transmitted packets are exactly the same!
  - Receiver can disambiguate a fresh packet from a retransmission






# Coping with packet loss: (3) Seq nums

- Good case: packet received and ACK received within the RTO
- Sequence numbers of adjacent transmitted packets are different
- Q1: where are the sequence numbers written to, exactly?
- Q2: what is the seq# of third packet?



# Stop-and-Wait Reliability

- Scheme so far: sender waits for an ACK/RTO before sending another packet
  - Also called “stop and wait” reliability
- Suppose no packets are dropped
  - Round-trip-time: 100 milliseconds
  - Packet size: 12,000 bits
  - Link rate: 12 Mega bits/s
- At what rate is the sender getting data across to the receiver?



# Reliable data delivery


Pipelined Transmission, a.k.a. Sliding Window Protocols

# Amount of “in-flight” data

- We term the amount of unACKed data as data “in flight”
- With just one packet in flight, the data rate is limited by the packet delay (RTT) rather than available bandwidth (link rate)
- Idea: **Keep many packets in flight**
  - Also referred to as **pipelined transmission**
- More packets in flight improves **throughput**
  - Throughput is the amount of data delivered per unit time.

# Keeping many packets in flight


- With link rate 12 Mega bits/s and packet size 12,000 bits, what is the **transmission delay** of a single packet?
- Remember the earlier throughput of 120 Kbit/s. If there are, say 4 packets in flight, throughput is 480 Kbits/s!
- We just improved the throughput 4 times by keeping 4 packets in flight




# TCP is a **pipelined transmission protocol**

Sender allows multiple, “in-flight”, yet-to-be-acknowledged packets

A few packets are on the way while, concurrently, new packets are transmitted




(a) a stop-and-wait protocol in operation




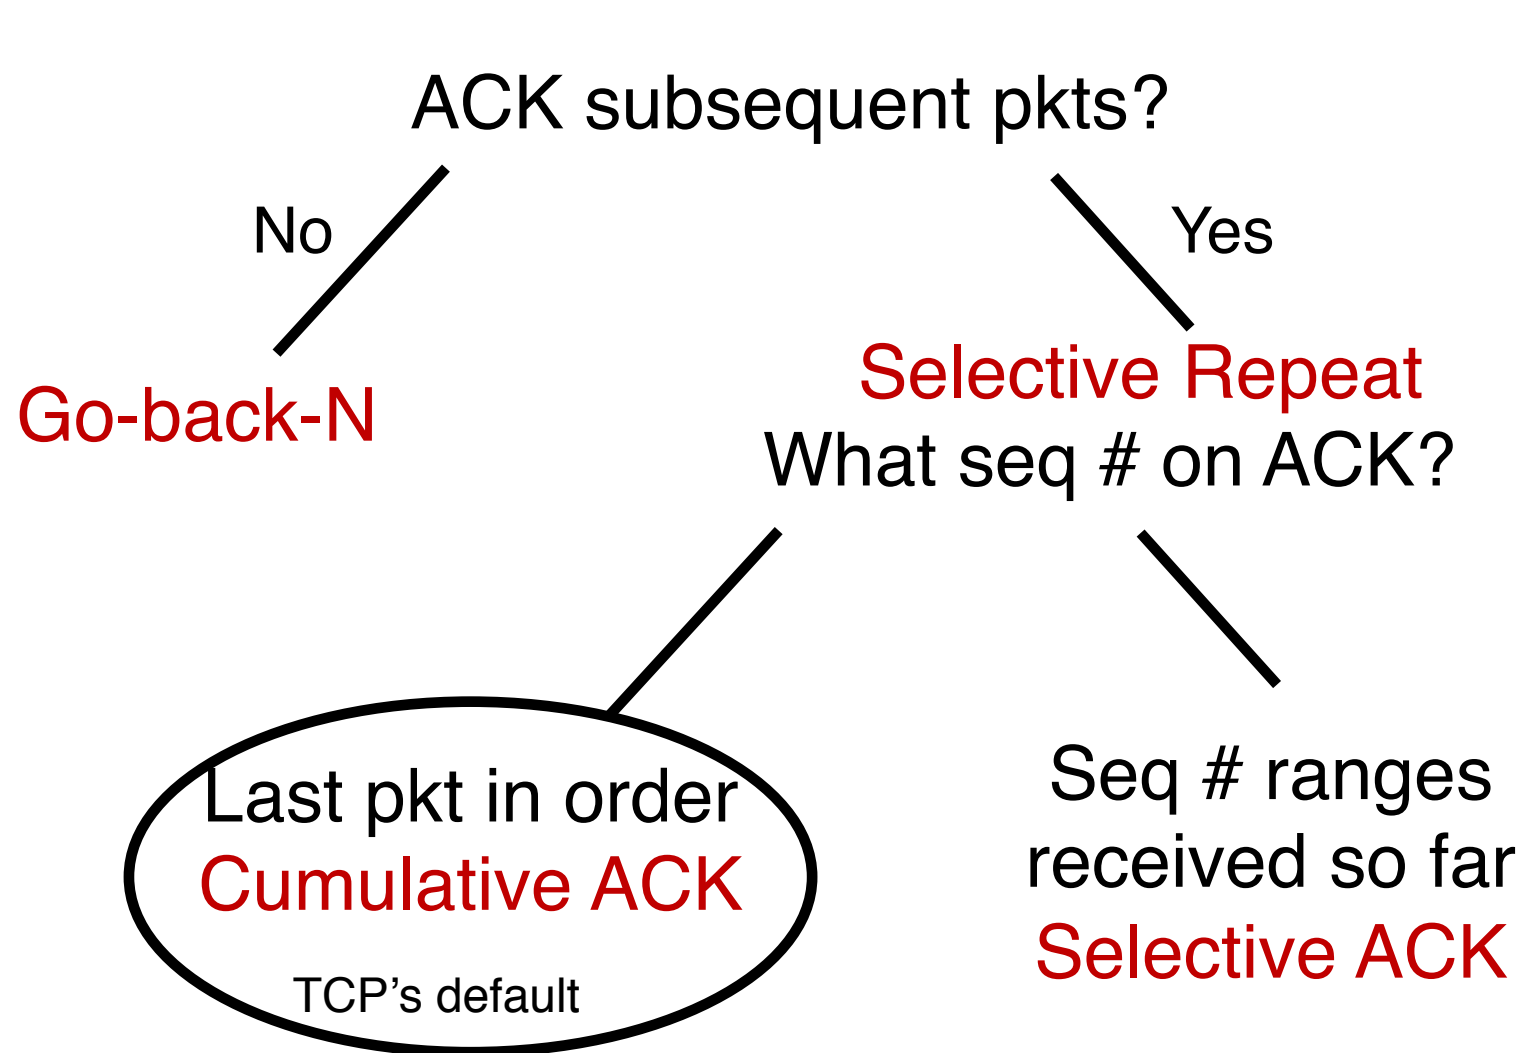
(b) a pipelined protocol in operation

# What if some packets/ACKs dropped?

- **Sequence numbers** help associate an ACK with its packet
  - Note: In TCP, every byte has a sequence #
  - We will often simplify our examples by assuming each packet has a sequence #
- In TCP, the ACK contains the sequence number of the next byte expected
  - Note: example uses packet seq #s
- Q1: If a packet is dropped, should the receiver ACK subsequent packets?
  - Q2: If so, with what sequence number?



# Receiver strategies upon packet loss






# Sliding Window with Go Back N

- When the receiver notices a missing or erroneous frame:
- It simply discards all frames with greater sequence numbers
  - The receiver will send no ACK
- The sender will eventually time out and retransmit all the frames in its sending window

# Go back N



# Go Back N *(cont'd)*

Go Back N can recover from erroneous or missing frames

But...

It is wasteful. If there are errors, the sender will spend time retransmitting frames the receiver has already seen

# Selective repeat with cumulative ACK

Idea: sender should only retransmit dropped/corrupted segments.

- The receiver **stores** all the correct frames that arrive following the bad one. (Note that the receiver requires a **memory buffer** for each sequence number in its receiver window.)
- When the receiver notices a skipped sequence number, it keeps acknowledging the **last good sequence number, . i.e., cumulative ACK**
- When the sender times out waiting for an acknowledgement, it **just retransmits the one unacknowledged frame**, not all its successors.

# Selective repeat

