

Network Layer: DHCP, ICMP, NAT, IPv6

CS 352, Lecture 11

<http://www.cs.rutgers.edu/~sn624/352-S19>

Srinivas Narayana

(heavily adapted from slides by Prof. Badri Nath and the textbook authors)



IP addresses: how to get one?

Q: How does a *host* get IP address?

- Hard-coded by system admin in a file
 - UNIX: /etc/network/interfaces
 - Windows: controlpanel -> network -> configuration -> tcp/ip -> properties
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from server
 - “plug-and-play”

Similar bootstrapping problems

- How does a host get its IP address?
- How does a host know its local DNS server?
- How does a host know its subnet mask?
- How does a host know which router is its “gateway” to other networks?

DHCP: Dynamic Host Configuration Protocol

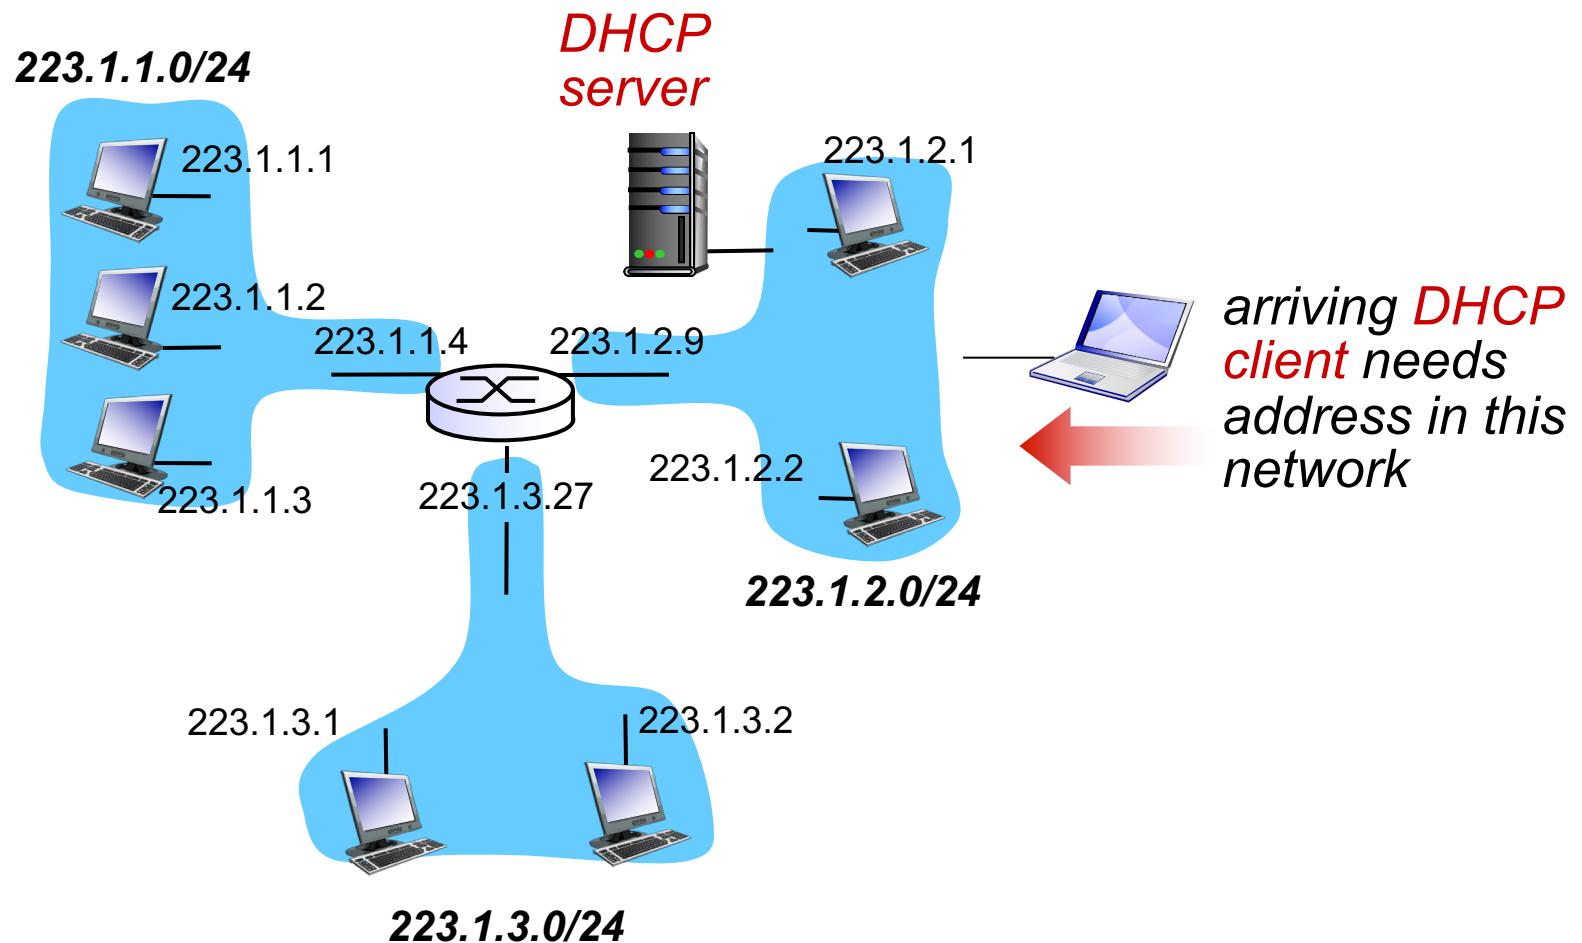
goal: allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network

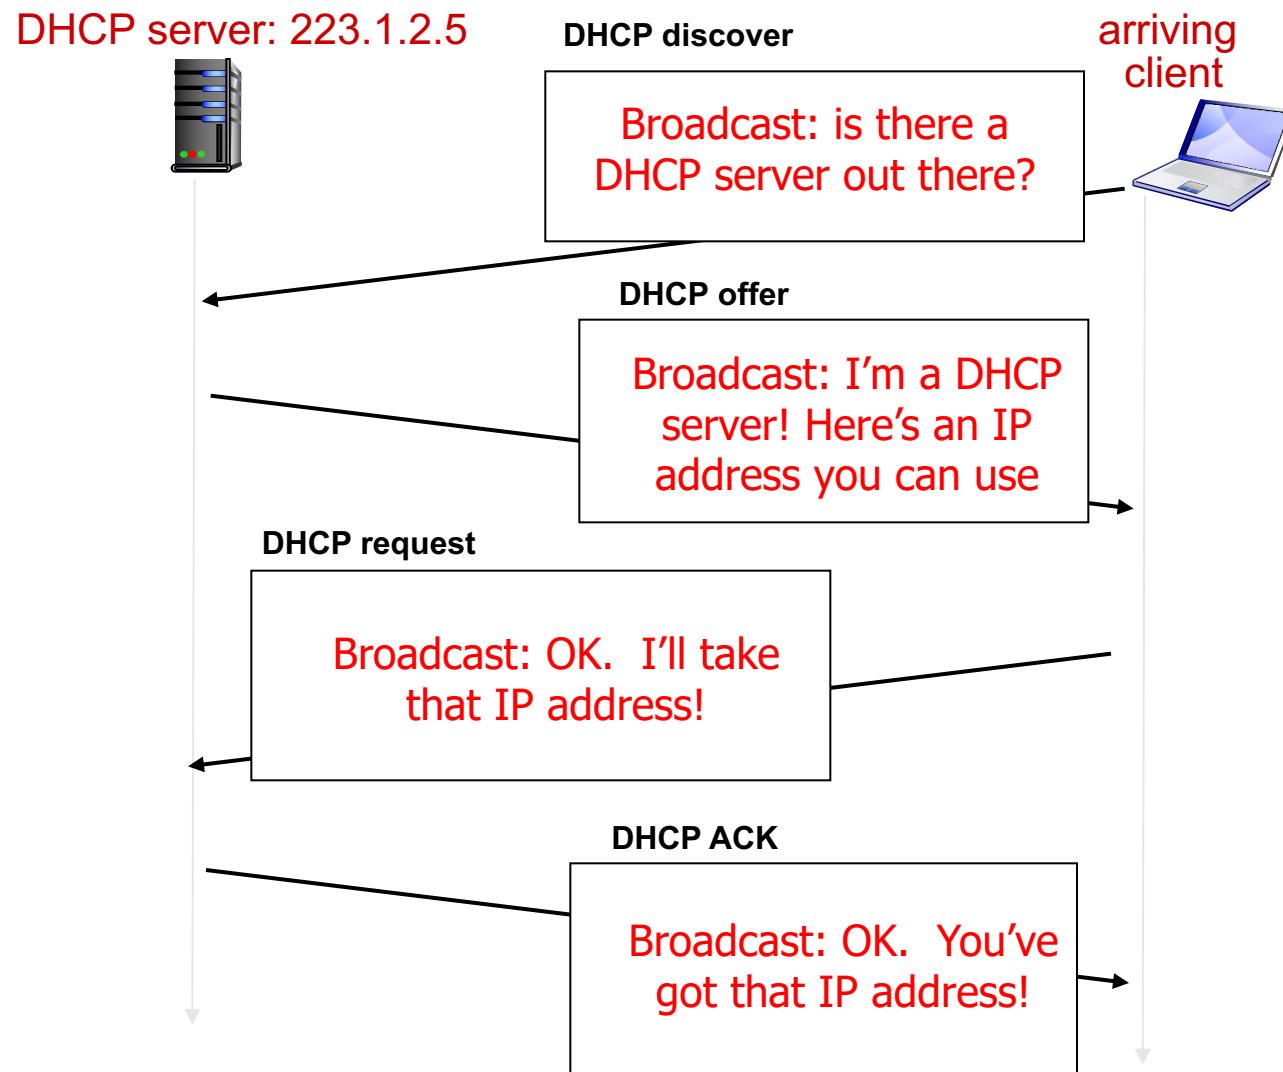
DHCP overview:

- host broadcasts “**DHCP discover**” msg [optional]
- DHCP server responds with “**DHCP offer**” msg [optional]
- host requests IP address: “**DHCP request**” msg
- DHCP server sends address: “**DHCP ack**” msg

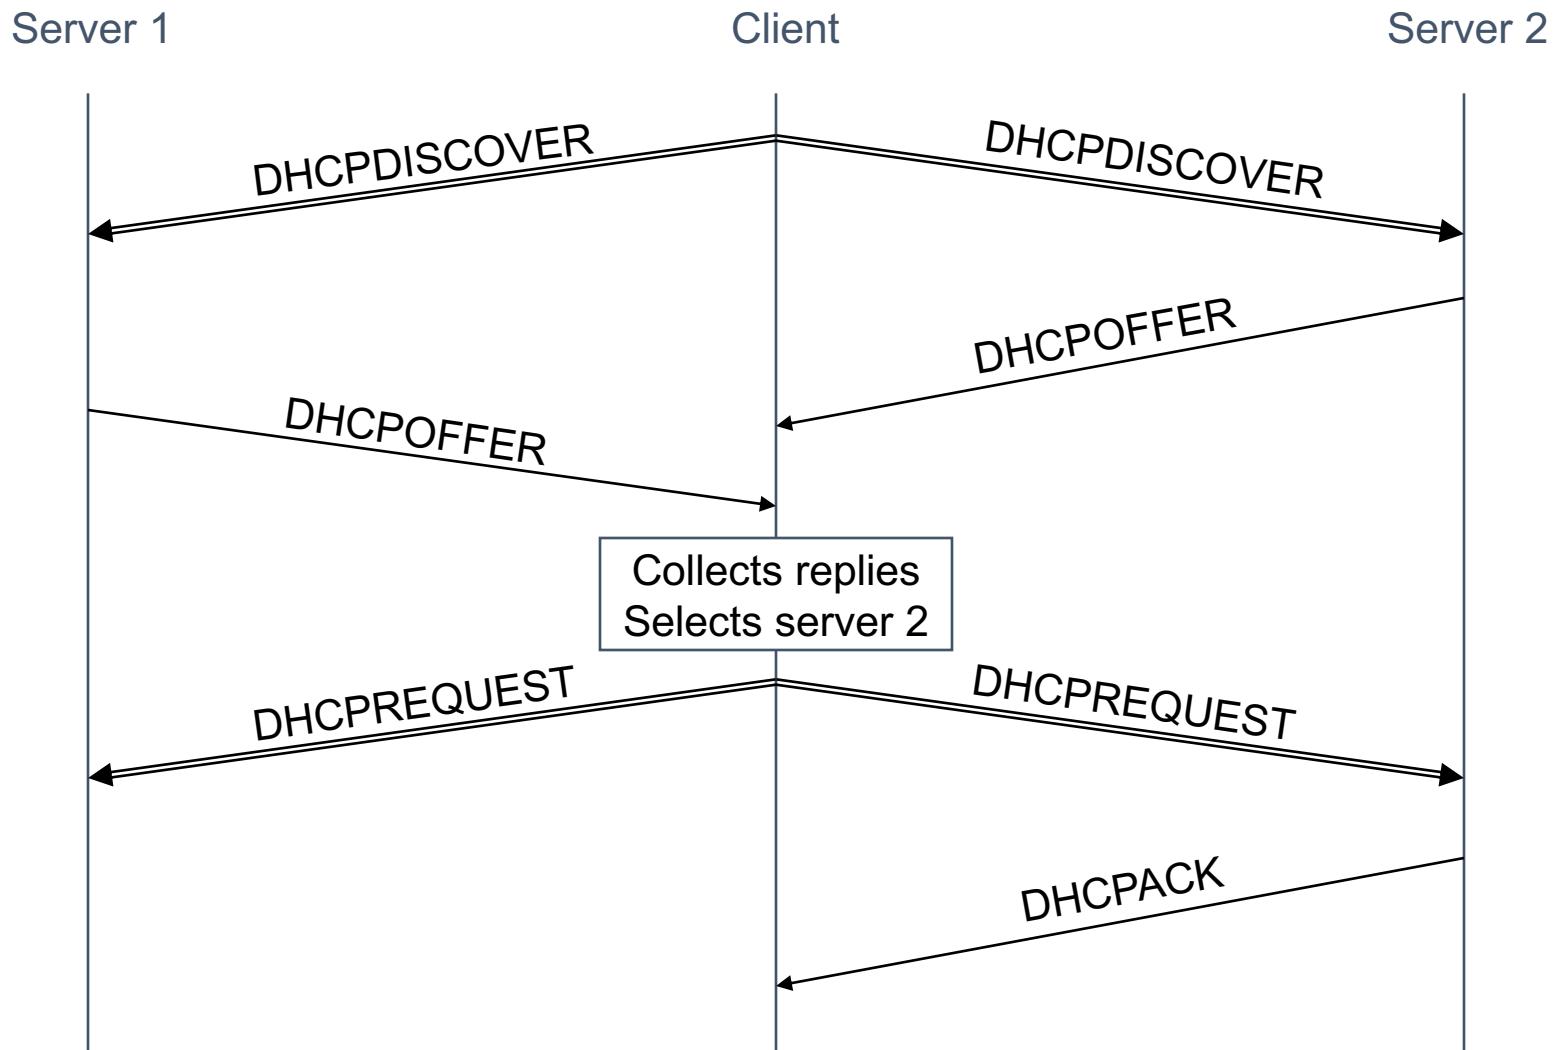
DHCP client-server scenario



DHCP client-server scenario



DHCP Protocol



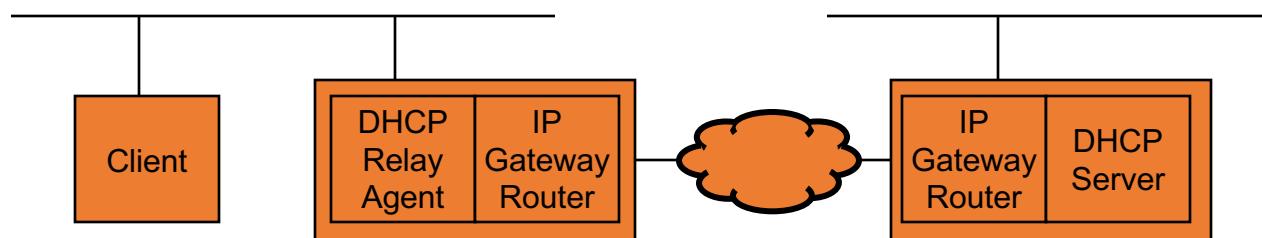
DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:

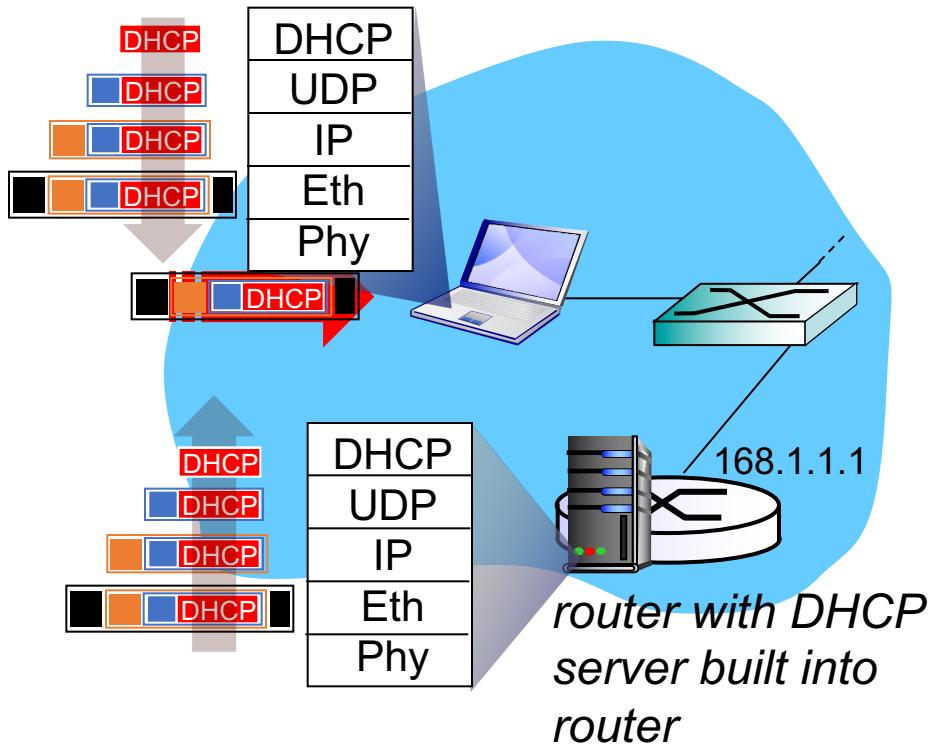
- address of first-hop router for client to reach other subnets
- name and IP address of the local DNS sever
- subnet mask

DHCP Relay Agents

- DHCP relay agents allow DHCP servers to handle requests from other subnets

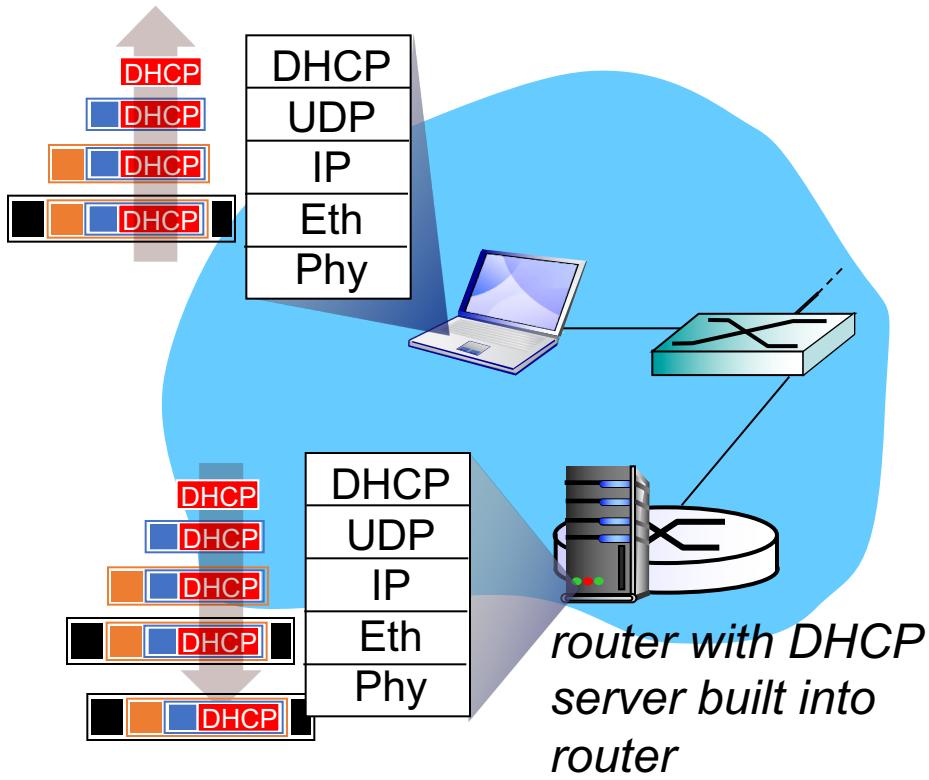


DHCP: example



- connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- Ethernet frame broadcast (dest: FFFFFFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demuxed to IP demuxed, UDP demuxed to DHCP

DHCP: example



- DCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router

Summary

- IP addresses don't have to be manually configured into hosts
- DHCP allows "ignorant" hosts to receive IP addresses (and more) at start-up time
- DHCP solves important bootstrapping problems in attaching new hosts to a network

Internet Control Message Protocol (ICMP)

ICMP

- Protocol for error detection and reporting
 - tightly coupled with IP, unreliable
- ICMP messages delivered in IP packets
- ICMP functions:
 - Announce reachability and network errors
 - Announce “time exceeded” errors for IP packets
 - Announce network congestion
- ICMP assists network troubleshooting in general

ICMP message

IP header

Source, Destination Address, TTL, ...

ICMP MSG

Message type, Code, Checksum,
Data

ICMP: Internet Control Message Protocol

Type	Code	description
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

Time for an activity!

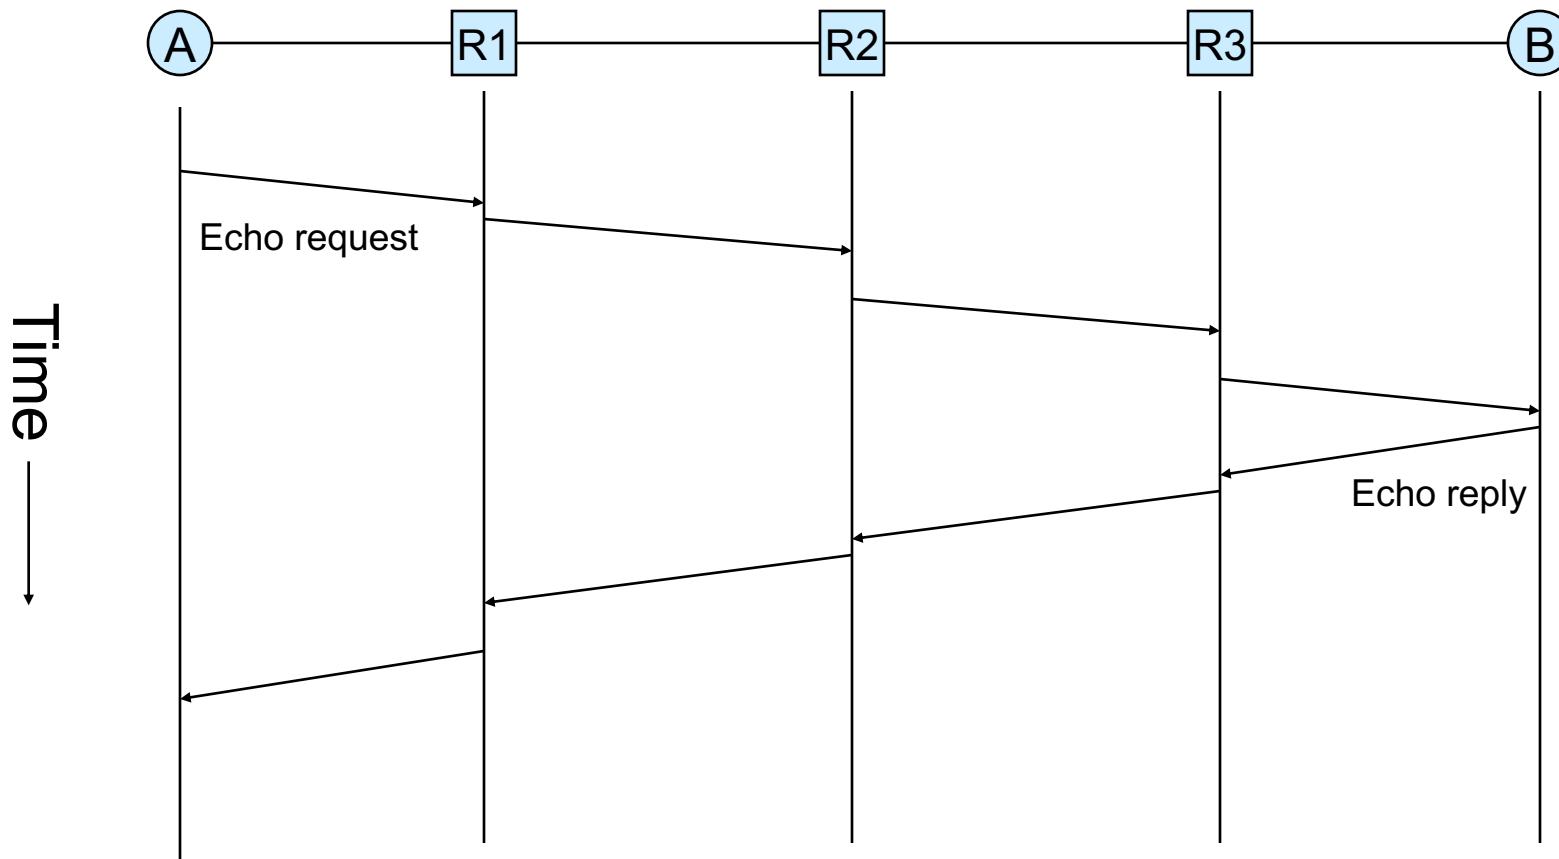
Specific uses of ICMP

- Echo request reply
 - Can be used to check if a host is alive
- Destination unreachable
 - Invalid address and/or port
- TTL expired
 - Routing loops, or too far away

Ping

- Uses ICMP echo request/reply
- Source sends ICMP echo request message to the destination address
- Destination replies with an ICMP echo reply message containing the data in the original echo request message
- Source can calculate round trip time (RTT) of packets
- If no echo reply comes back then the destination is unreachable

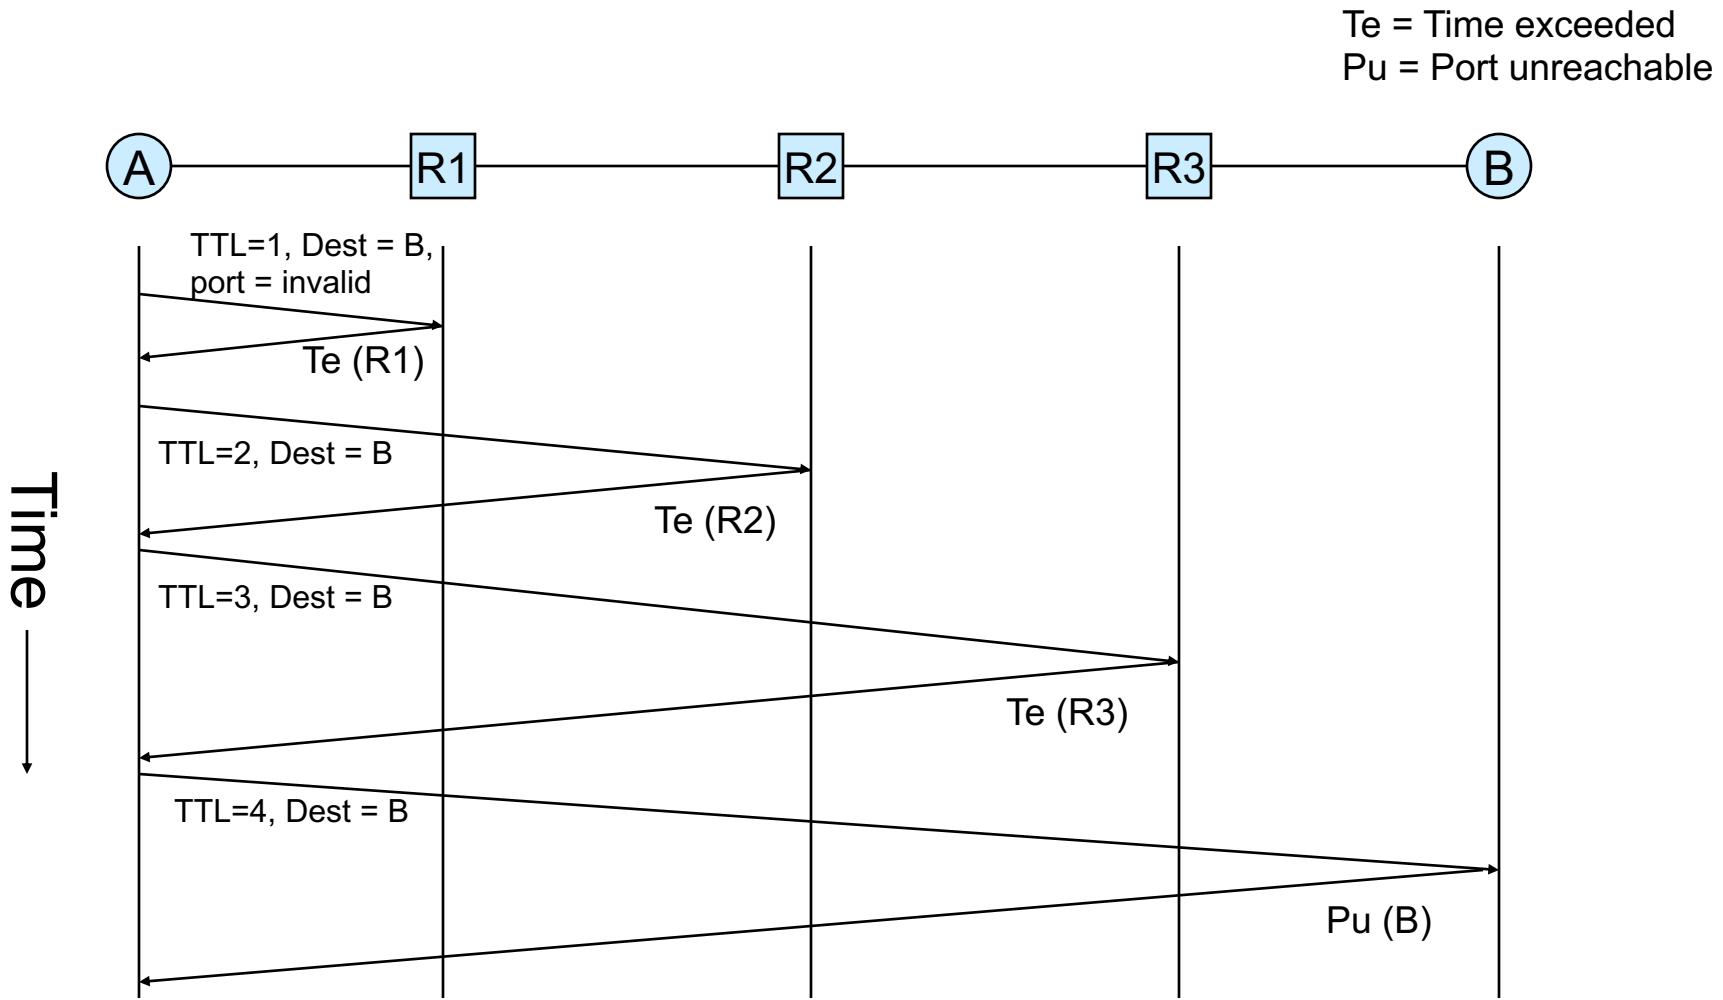
Ping (*cont'd*)



Traceroute

- Traceroute records the route that packets take
- A clever use of the TTL field
- When a router receives a packet, it decrements TTL
- If TTL=0, it sends an ICMP time exceeded message back to the sender
- To determine the route, progressively increase TTL
 - Every time an ICMP time exceeded message is received, record the sender's (router's) address
 - Repeat until the destination host is reached or an error message occurs
- If packet reaches the destination, the dest host usually sends an ICMP port unreachable

Traceroute (*cont'd*)

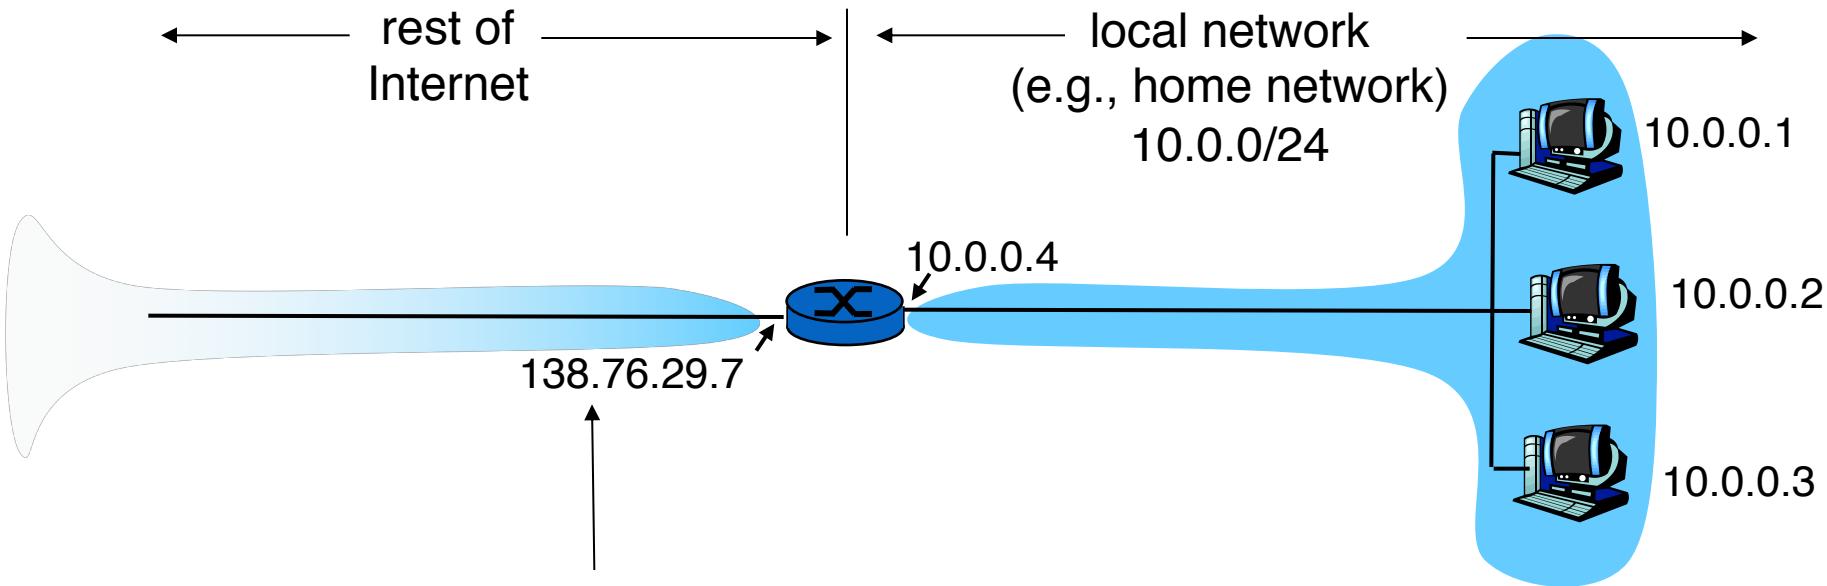


Traceroute example

```
[flow:352]$ traceroute google.com
traceroute to google.com (172.217.10.238), 64 hops max, 52 byte packets
 1  vlan451-sr03-hill-nbp.runet.rutgers.net (172.25.112.1)  9.278 ms  3.210 ms  3.124 ms
 2  xe-1-3-0-0-cr02-hill-nbp.runet.rutgers.net (172.29.6.65)  37.125 ms  2.912 ms  2.899 ms
 3  ae1-2000-cr10-hill-nbp.runet.rutgers.net (172.29.6.42)  3.078 ms  3.086 ms  3.016 ms
 4  ae5-2000-cr02-halsey-nwk.runet.rutgers.net (172.29.6.55)  3.693 ms  3.707 ms  3.793 ms
 5  ae2-0-er10-halsey-ext.runet.rutgers.net (172.29.8.6)  3.699 ms  3.693 ms  3.766 ms
 6  ae1-0-fw01-halsey-nwk.runet.rutgers.net (172.29.8.41)  4.019 ms  3.909 ms  3.750 ms
 7  et-2-2-0-0-er10-halsey-ext.runet.rutgers.net (172.29.8.46)  4.310 ms  4.181 ms  3.948 ms
 8  gateway-pni.google.com (128.6.1.114)  4.426 ms  3.901 ms  4.161 ms
 9  108.170.248.65 (108.170.248.65)  5.024 ms
     108.170.248.1 (108.170.248.1)  6.147 ms  6.165 ms
10  72.14.233.201 (72.14.233.201)  5.316 ms  5.426 ms  5.359 ms
11  lga25s59-in-f14.1e100.net (172.217.10.238)  5.410 ms  5.156 ms  5.135 ms
[flow:352]$
```

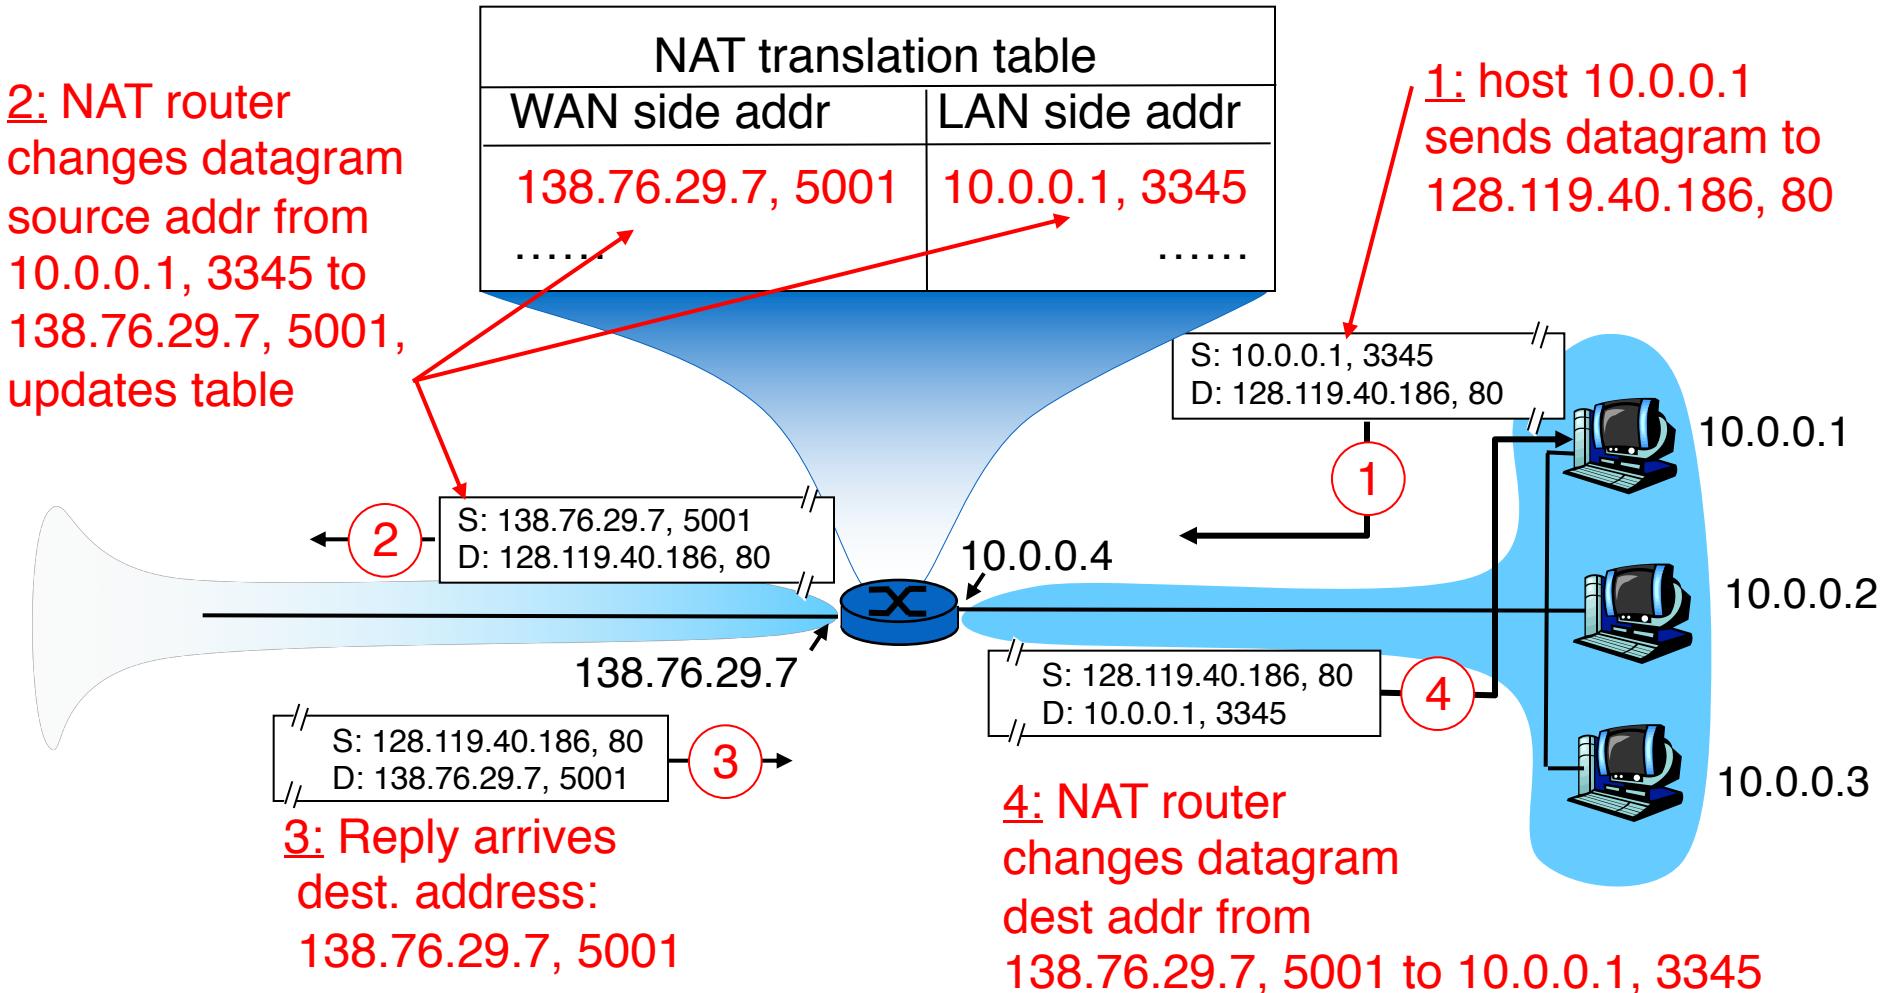
Network Address Translation (NAT)

NAT: Network Address Translation



All datagrams *leaving* local
network have **same** single source
NAT IP address: 138.76.29.7,
different source port numbers

NAT: Network Address Translation



NAT: Network Address Translation

- **Features:** local network uses just one IP address as far as outside world is concerned:
 - range of addresses not needed from ISP: just one IP address for all devices
 - can change addresses of devices in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - devices inside local net not explicitly addressable, visible by outside world (a security plus).

Think about...

- How do the hosts inside the home network get their IP addresses?
- How does your home router get its externally visible IP address?

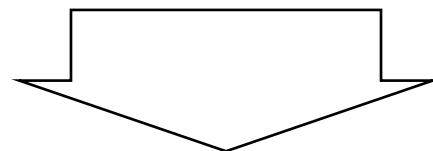
NAT: Network Address Translation

- 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
 - Routers should only work upto the network layer, not transport ports!
 - violates “end-to-end argument”
 - NAT must be taken into account by app designers
 - e.g., P2P applications like skype
 - Purists: address shortage should instead be solved by IPv6

Internet Protocol v6 (IPv6)

Recent Developments: IPv6

- IPv4 has limited address space (32 bits) and is running out of addresses. 32 bits are not enough!
- More devices: phones, watches, your refrigerator(!), ...
- Real-time traffic and mobile users are also becoming more common



IP version 6

IPv6: Main changes from IPv4

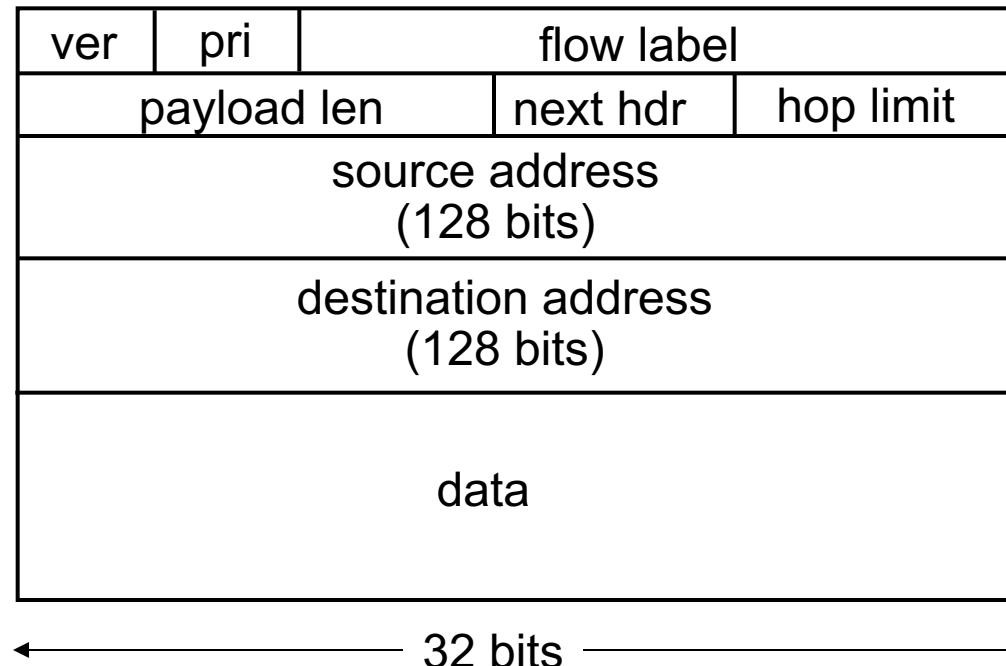
- Large address space:
 - 128-bit addresses (16 bytes)
 - Allows up to 340,282,366,920,938,463,463,374,607,431,768,211,456 unique addresses (3.4×10^{38})
- Fixed length headers (40 bytes)
 - Improves the speed of packet processing in routers
- IPv6 “options” processing happens through a separate mechanism

IPv6 datagram format

priority: identify priority among datagrams in flow

flow Label: identify datagrams in same “flow”
(concept of “flow” left undefined)

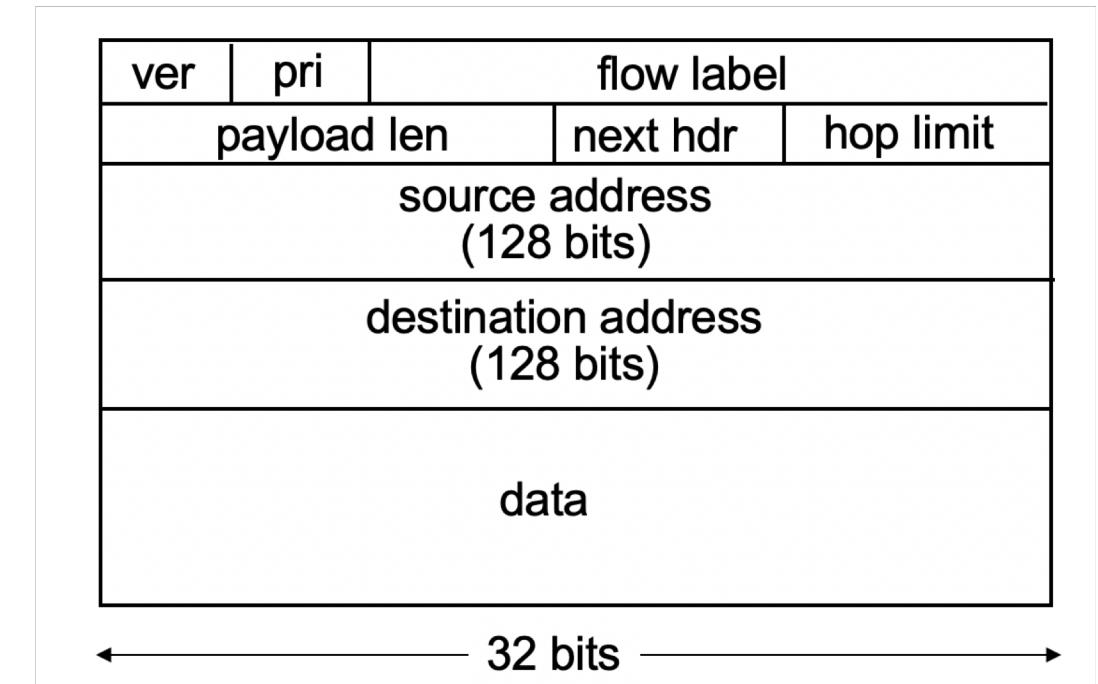
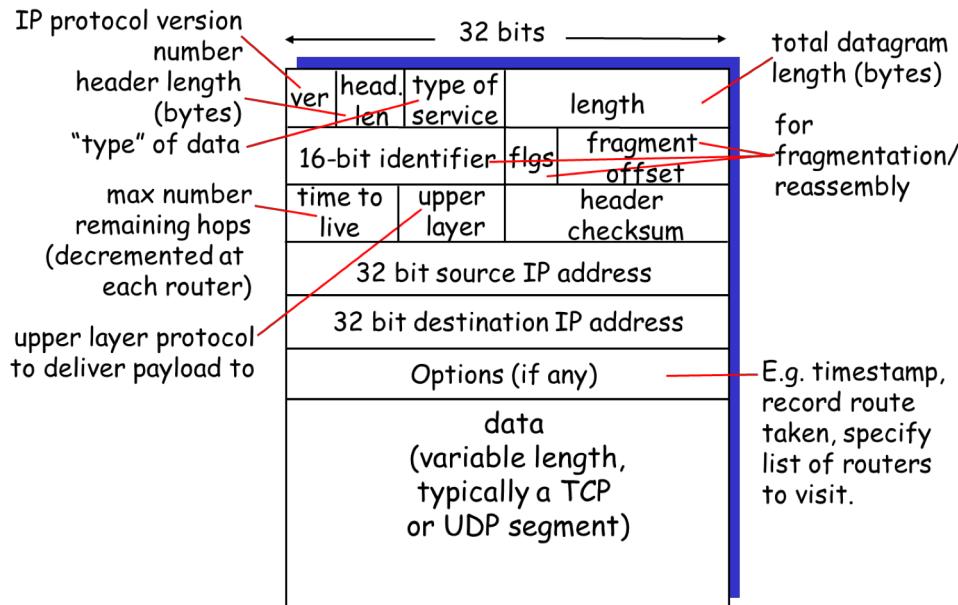
next header: identify upper layer protocol for data



Other changes from IPv4

- *checksum*: removed entirely to reduce processing time at each hop
- *options*: allowed, but outside of header, indicated by “Next Header” field
- *ICMPv6*: new version of ICMP
 - additional message types, e.g. “Packet Too Big”
 - multicast group management functions

IPv4 vs IPv6: Can you tell the differences?



IPv6 Flows

- Support for “flows”
 - Flows help support real-time service in the Internet
 - A “flow” is a number in the IPv6 header that can be used by routers to see which packets belong to the same stream
 - Guarantees can then be assigned to certain flows
 - Example:
 - Packets from flow 10 should receive rapid delivery
 - Packets from flow 12 should receive reliable delivery

IPv6 Addresses

- Classless addressing/routing (similar to CIDR)
- Notation: **xx:xx:xx:xx:xx:xx:xx:xx**
 - x = 4-bit hex number
 - contiguous 0s are compressed: 47CD::A456:0124
 - IPv6 compatible IPv4 address: ::128.64.18.87
 - First 96 bits are 0
 - Global unicast addresses start with 001....
 - 2000::/3 prefix

IPv6: Adoption

- Google: 8% of clients access services via IPv6
- NIST: 1/3 of all US government domains are IPv6 capable
- *Long (long!) time for deployment, use*
 - 20 years and counting!
- Think of application-level changes in last 20 years: WWW, Facebook, streaming media, Skype, ...
- *Why?*

Traceroute Examle

```
1 lcsr-gw (128.6.13.21) 1.206 ms 0.973 ms 0.782 ms
2 rucs-gw (165.230.212.129) 0.697 ms 0.569 ms 0.571 ms
3 transition2-gw (165.230.12.145) 2.786 ms 0.994 ms 0.769 ms
4 rutgers-gw.Rutgers.EDU (198.151.130.209) 1.726 ms 2.048 ms Vl1000-sr02-hil
1012-svcs.Rutgers.EDU (198.151.130.14) 1.278 ms
5 rutgers-gw.Rutgers.EDU (198.151.130.209) 1.755 ms 1.241 ms 1.828 ms
6 198.151.130.226 (198.151.130.226) 2.748 ms 3.070 ms 2.640 ms
7 clev-nycm.abilene.ucaid.edu (198.32.8.29) 15.162 ms 14.619 ms 14.663 ms
8 ipls-clev.abilene.ucaid.edu (198.32.8.25) 21.220 ms 22.497 ms 21.450 ms
9 kscy-ipls.abilene.ucaid.edu (198.32.8.5) 30.257 ms 30.604 ms 30.969 ms
10 dnvr-kscy.abilene.ucaid.edu (198.32.8.13) 40.823 ms 41.181 ms 41.076 ms
11 snva-dnvr.abilene.ucaid.edu (198.32.8.1) 65.436 ms 66.068 ms 65.569 ms
12 198.32.249.161 (198.32.249.161) 65.673 ms 65.771 ms 66.006 ms
13 BERK--SUNV.POS.calren2.net (198.32.249.13) 67.183 ms 67.131 ms 66.858 ms
14 pos1-0.inr-000-eva.Berkeley.EDU (128.32.0.89) 67.192 ms 66.749 ms 67.720
ms
15 vlan198.inr-201-eva.Berkeley.EDU (128.32.0.194) 67.373 ms 67.067 ms 67.82
1 ms
16 fast8-0-0.inr-210-cory.Berkeley.EDU (128.32.255.122) 67.634 ms 68.735 ms
68.413 ms
17 GE.cory-gw.EECS.Berkeley.EDU (169.229.1.46) 67.575 ms 68.222 ms 67.772 ms
18 gig8-1.snr1.CS.Berkeley.EDU (169.229.3.66) 67.454 ms 67.988 ms 67.177 ms
19 now.CS.Berkeley.EDU (128.32.44.96) 67.892 ms * 67.818 ms
```