

# The Transport Layer: TCP Timeouts and Connection management

CS 352, Lecture 11, Spring 2020

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Course announcements

- Project 2 will go online today
- Quiz 4 will go online today
  - Due Tue 03/10 at 10 PM
- Mid-term grades will be released this weekend
  - Papers in class on Wednesday

# Review of concepts

- TCP congestion control: need distributed, efficient, fair
  - signals (loss) and knobs (congestion window)
- ACK clocking
- Slow start
- Additive increase and the slow start threshold
- Multiplicative decrease
  - Triple duplicate ACKs and fast retransmit


TCP performs additive increase and multiplicative decrease of its congestion window.

This is often termed **AIMD**.

AIMD results in the so-called **TCP sawtooth**.

# Why AIMD?

- Converges to fairness
- Can also show it converges to efficiency
  - Intuition: Increments to rate get smaller as fairness increases



# Calculating the TCP timeout


# TCP timeout (RTO)

- Useful for reliable delivery and congestion control
- **How to pick the RTO value?**
  - Too long: slow reaction to loss
  - Too short: premature unnecessary retransmissions
- Intuition: somehow use the observed RTT (sampleRTT)
  - Can we just directly set the latest RTT as the RTO?
- RTT can vary significantly!
  - Intermittent congestion, path changes, signal quality changes on wireless channel, etc.

# Estimated RTT

- Exponential weighted moving average (typical  $\alpha = 1/8$ )


$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$





# Timeout == estimated RTT + **safety**

- Estimated RTT can have a large **variance**
  - Use a larger safety margin if larger variance



Timeout == estimated RTT + **safety**

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”

# Managing a single timer

## data rcvd from app:

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: `TimeoutInterval`


## timeout:

- retransmit segment that caused timeout
- restart timer

## ack rcvd:

- if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - restart timer if there are still unacked segments

# Problem with sampleRTT calculation



# Retransmission ambiguity

- If you retransmitted, how do you measure `sampleRTT` for it?
  - Measure RTT from original data segment?
  - Measure RTT from most recent (retransmitted) segment?
- There could be an error in RTT estimate, since we can't be sure
- One solution
  - **Never update RTT measurements** based on acknowledgements from retransmitted packets
- Problem: **Sudden change in RTT**, coupled with many retransmissions, can cause system to update RTT very late
  - Ex: Primary path failure leads to a slower secondary path

# Karn's algorithm


- Use back-off as part of `sampleRTT` computation
- Whenever packet loss, RTO is increased by a factor
- Use this increased RTO as RTO estimate for the next segment (**not from EstimatedRTT**)
- Only after an acknowledgment received for a successful transmission is the timer set to new RTT obtained from EstimatedRTT

# Connection Management


# Connection Management

before exchanging data, sender/receiver “handshake”:

- agree to establish connection
- agree on connection parameters



```
Socket clientSocket =  
    newSocket("hostname", "port  
number");
```





```
Socket connectionSocket =  
    welcomeSocket.accept();
```



# Agreeing to establish a connection

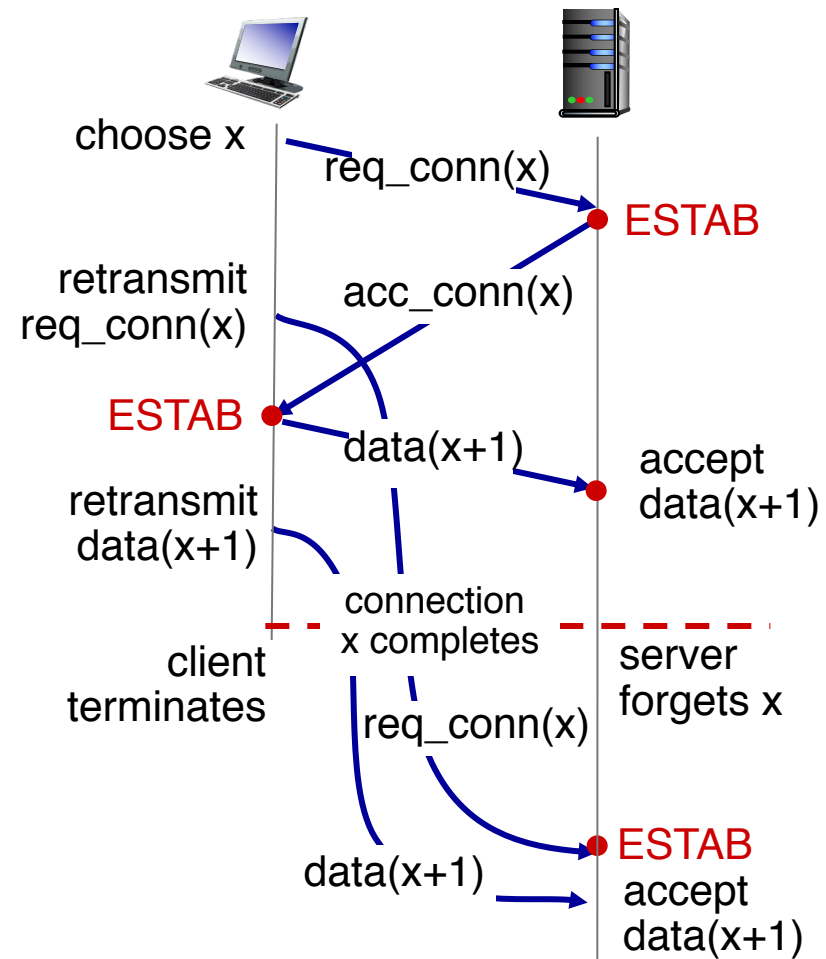
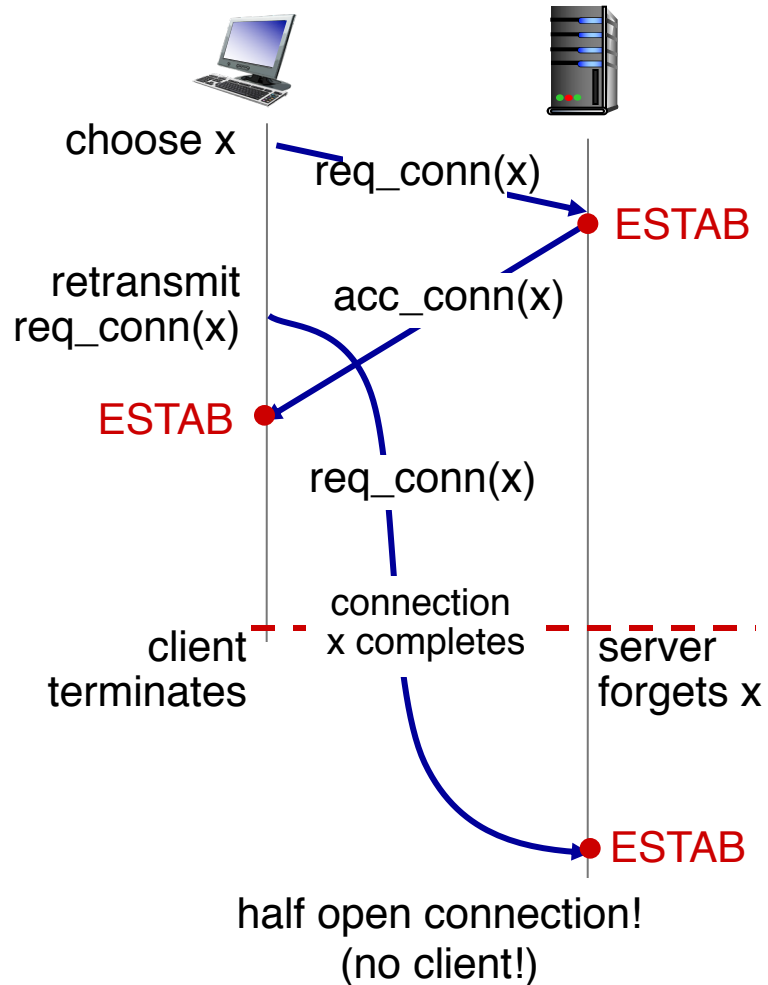
2-way handshake:



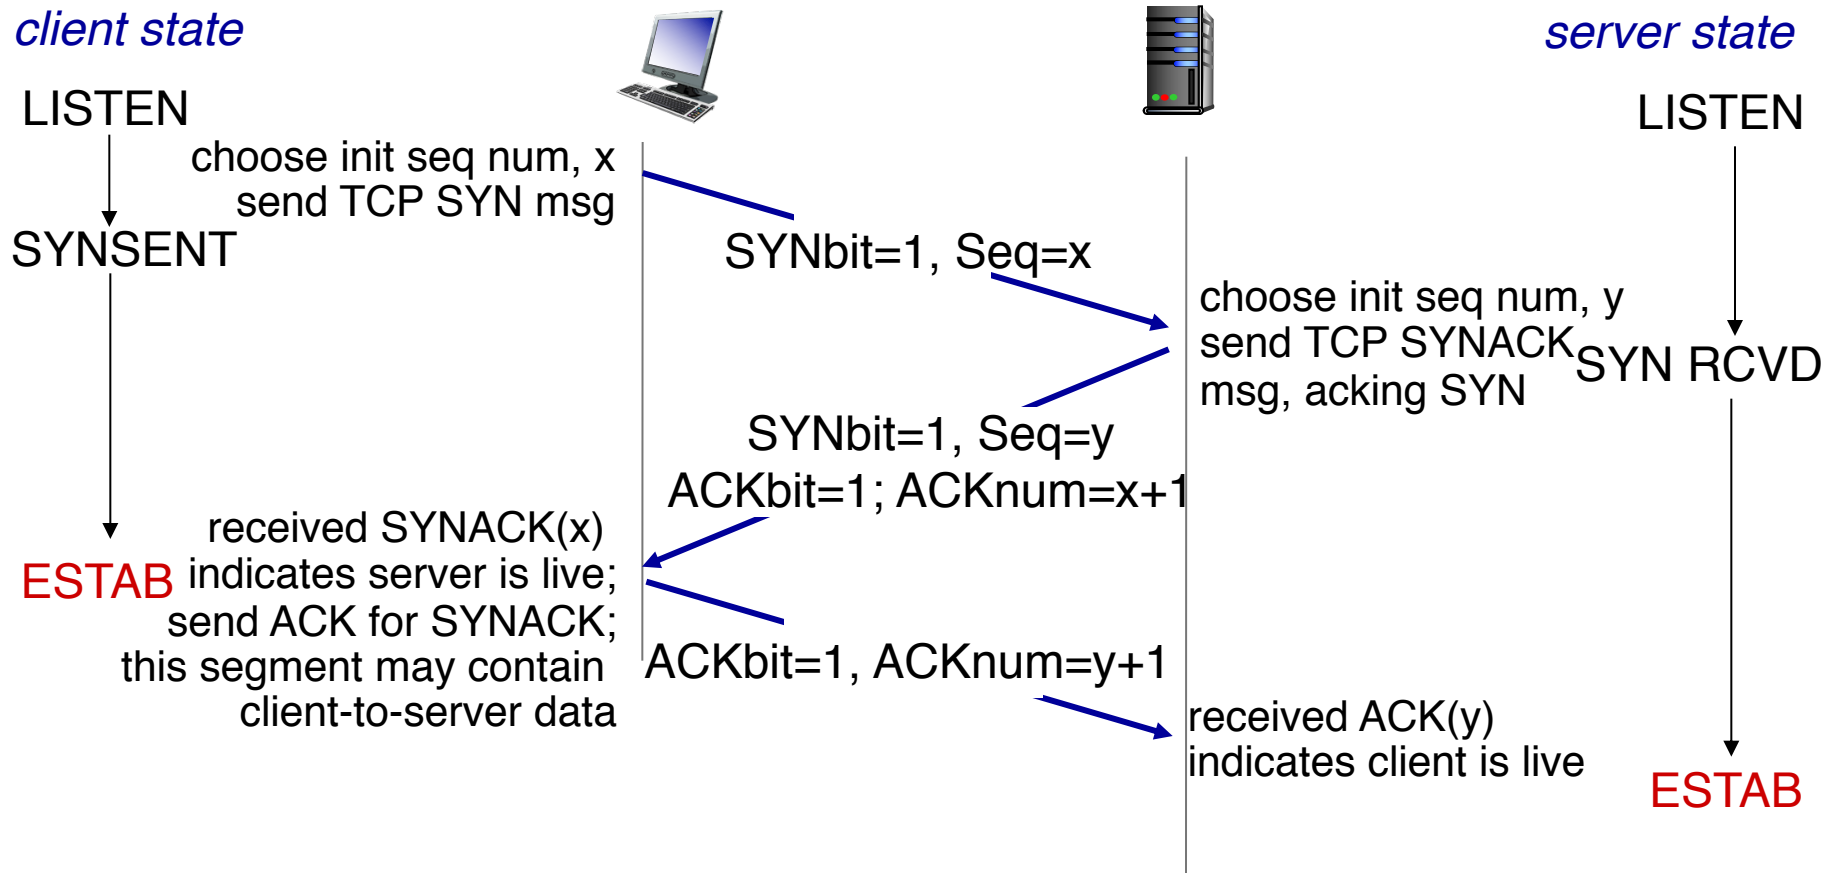
Q: will 2-way handshake always work in network?

- variable delays
- retransmitted messages (e.g. req\_conn(x)) due to message loss
- message reordering
- can't "see" other side

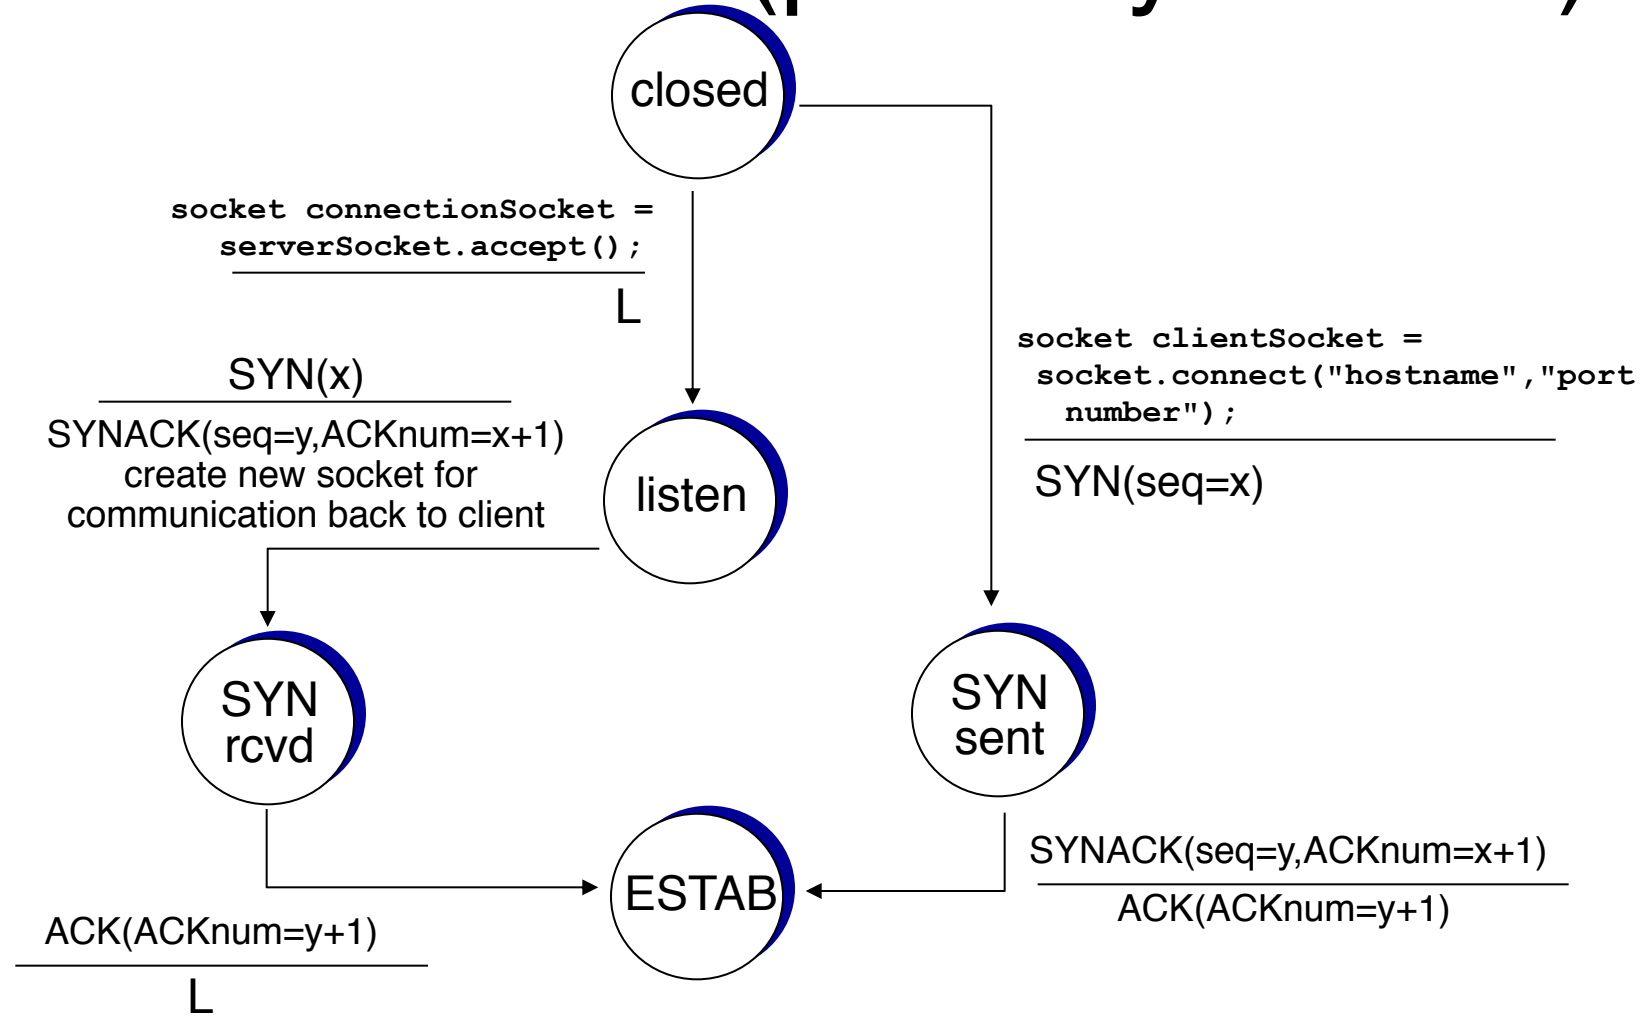
# 2-way handshake failure scenarios



# TCP 3-way handshake



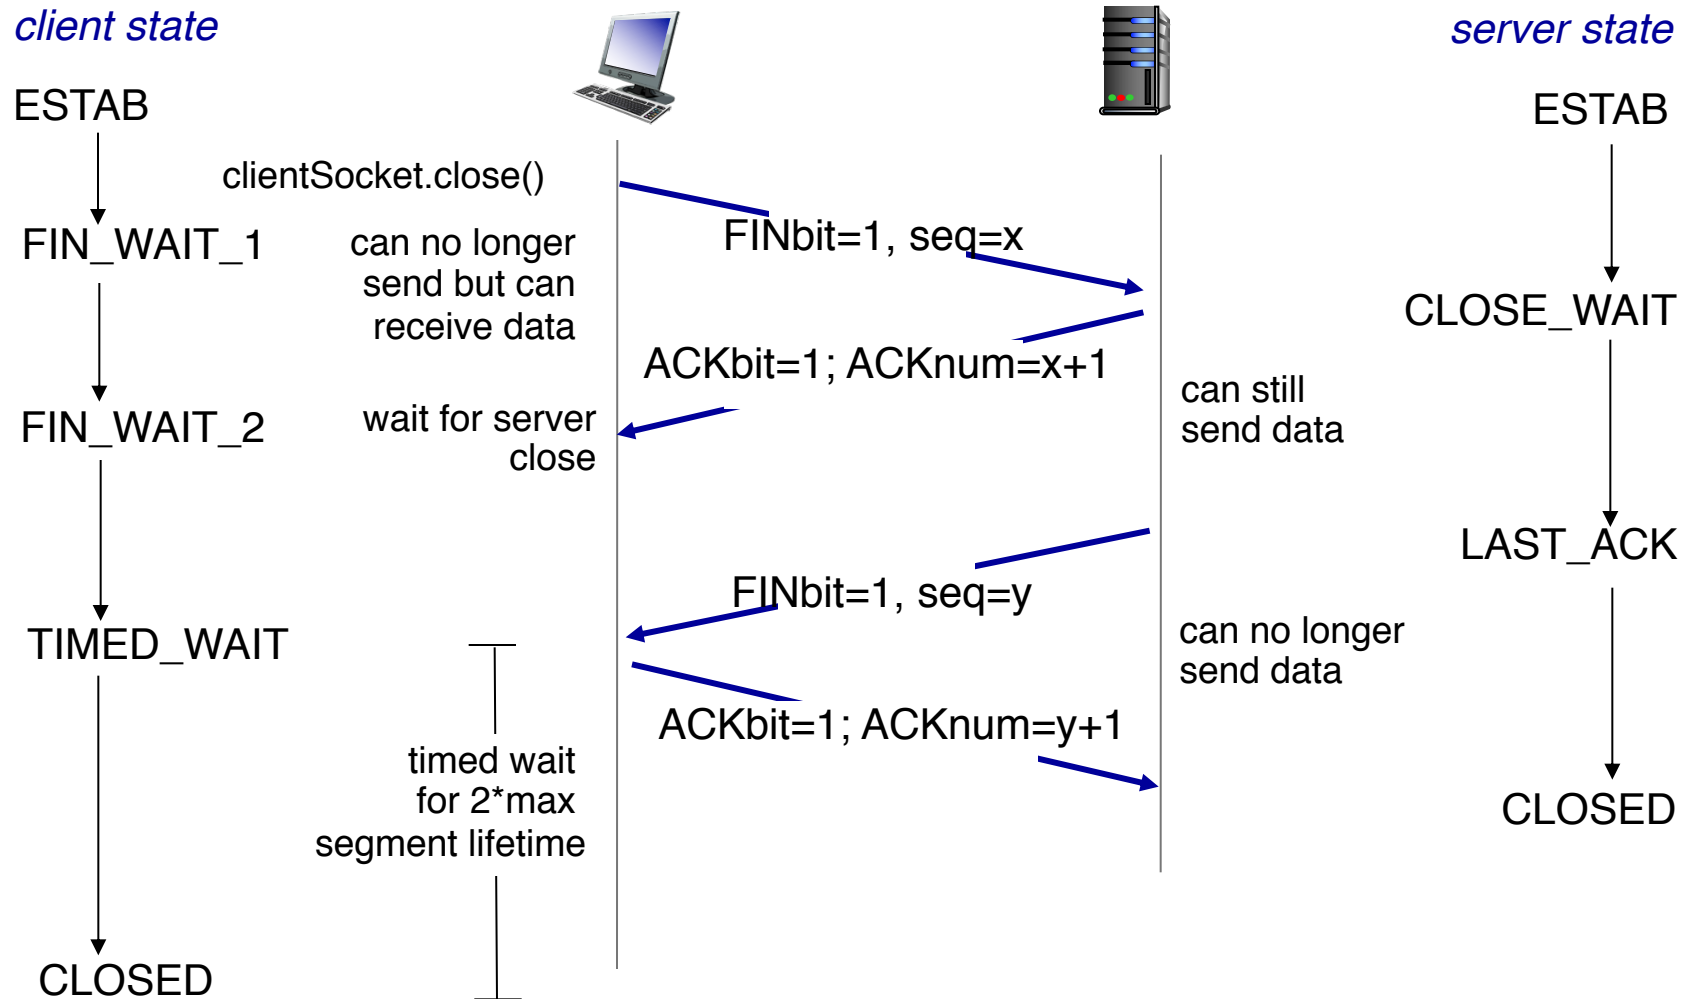
# TCP state machine (partially shown)



# TCP: closing a connection

- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- In general, **TCP is full-duplex**: both sides can send
- But **FIN is unidirectional**: stop one side of the communication
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

# TCP: closing a connection



# TCP summary

- Reliability
- Ordering
- Flow control
- Congestion control
- Timeout computation
- Connection management, state machine