

# The Network Layer: Protocols DHCP, ICMP, NAT, IPv6

CS 352, Lecture 14, Spring 2020


<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Course announcements

- Quiz 5 due Tuesday
  - Will be released later today
- Recording of Wednesday's lecture available
  - See Piazza for details
  - We are working on converting video into a format that's more accessible than WebEx's ARF... stay tuned

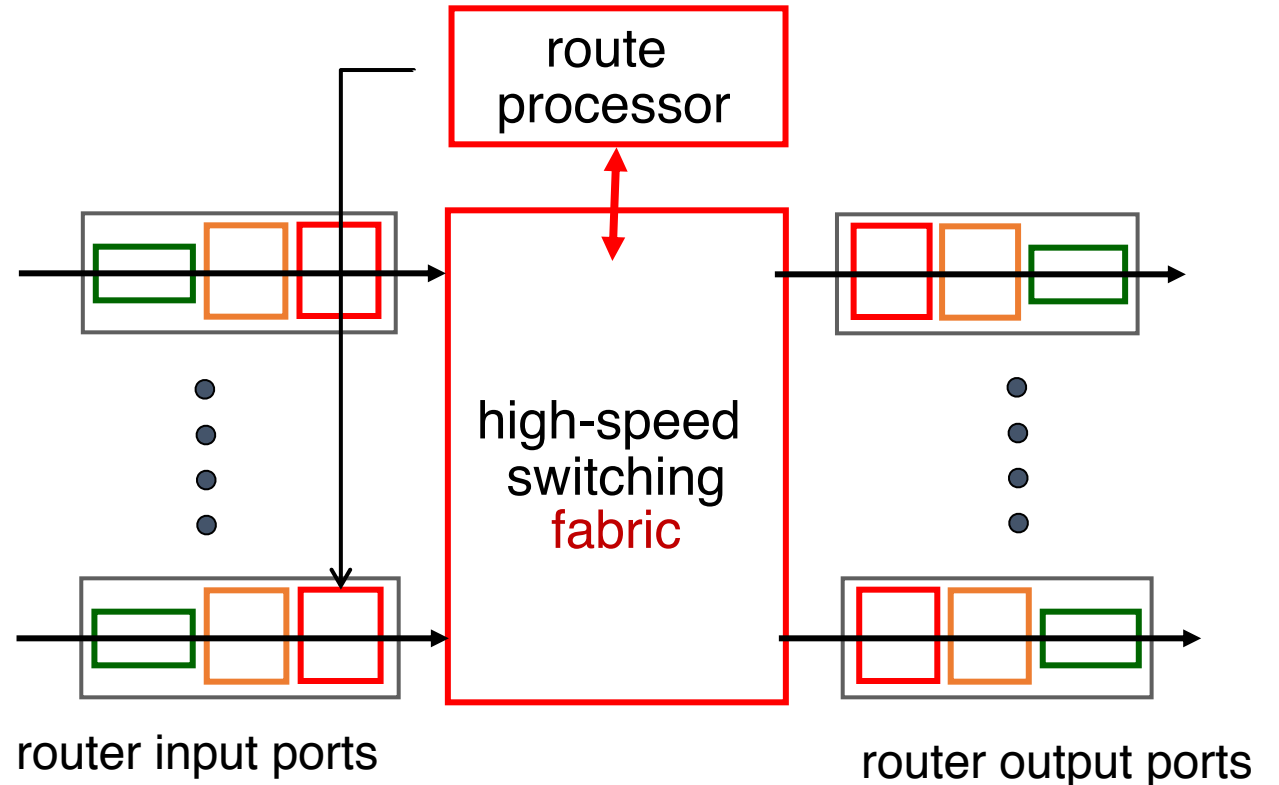
# Where we are: The network layer



The network layer exists on every endpoint and router.

# Review of concepts

- Router components
- Input port: line termination, forwarding function
  - Forwarding: based on IP destination.
  - Longest-prefix matching
- Switching fabric: memory, bus, crossbar
- Output port: buffer management and scheduling



# Poll #1

- If an ISP X owns two prefixes 128.0.0.0/16 and 128.1.0.0/16, which IP prefix can the rest of the Internet use in its forwarding rules to route data towards X?
  - (a) 128.0.0.0/8
  - (b) 128.0.0.0/24
  - (c) 128.0.0.0/15
  - (d) None of the above

# Poll #2

- Suppose a router has two forwarding rules. Which port should pkt with destination IP 192.168.0.56 go out of?

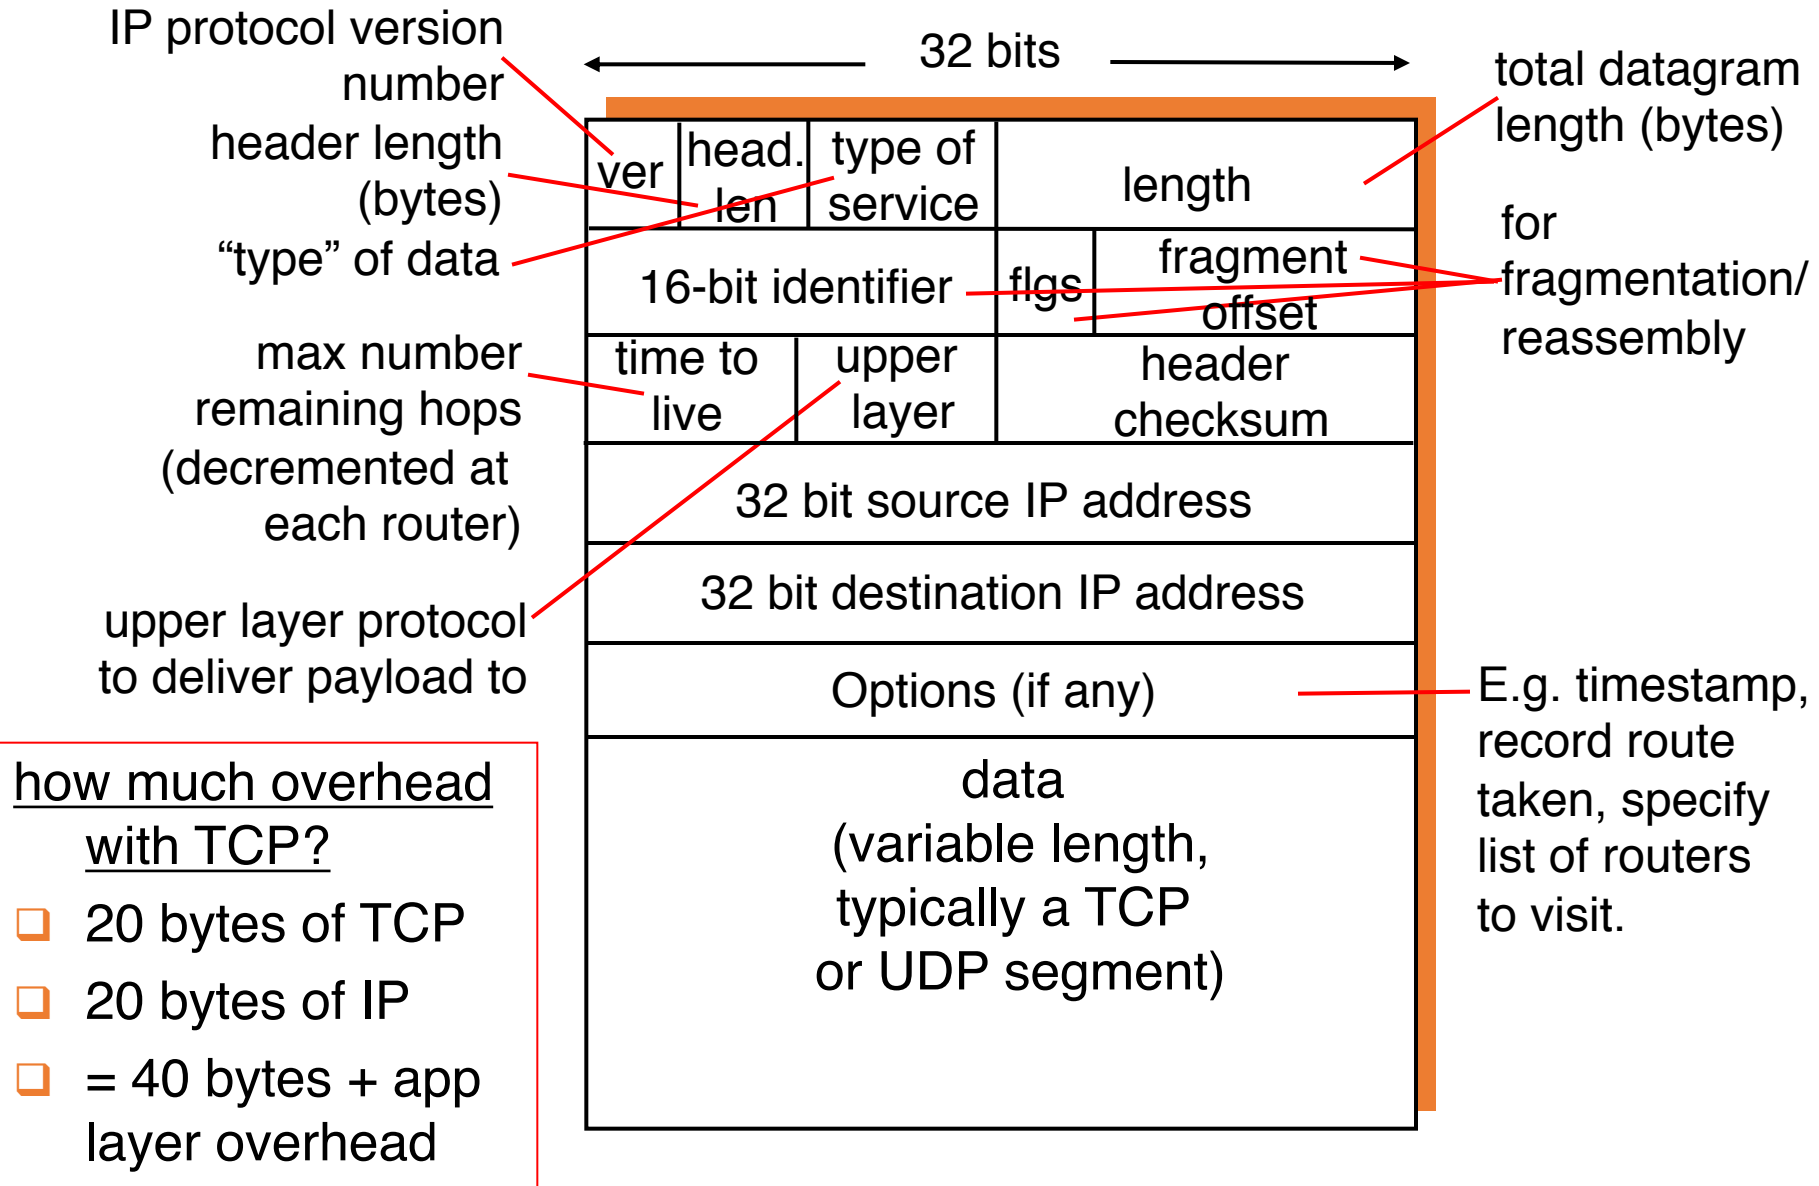
192.168.0.0/16 -> port 1

192.168.0.0/24 -> port 2

- (a) port 1
- (b) port 2
- (c) neither port
- (d) both ports

# The Internet Protocol (IP)


# IP datagram format





# IP Fragmentation & Reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame.
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - reassembled only at the destination, at the IP layer
  - IP header bits used to identify, order related fragments



# IP Fragmentation and Reassembly

## Example

- ❑ 4000 byte datagram
- ❑ MTU = 1500 bytes
- ❑ (includes IP headers)

1480 bytes in  
data field

offset =  
 $1480/8$

	length =4000	ID =x	fragflag =0	offset =0	
--	-----------------	----------	----------------	--------------	--

One large datagram becomes  
several smaller datagrams

	length =1500	ID =x	fragflag =1	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =185	
--	-----------------	----------	----------------	----------------	--

	length =1040	ID =x	fragflag =0	offset =370	
--	-----------------	----------	----------------	----------------	--

# Poll #3

- Which header does the IP source and destination live in?
  - (a) TCP header
  - (b) IP header
  - (c) Application-layer header
  - (d) None of the above

# Dynamic Host Configuration

How does an endpoint get its IP address?

# IP addresses: how to get one?

**Q:** How does a *host* get IP address?

- Hard-coded by system admin in a file
  - UNIX: /etc/network/interfaces
  - Windows: controlpanel -> network -> configuration -> tcp/ip -> properties
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from as server
  - “plug-and-play”

# Similar bootstrapping problems

- How does a host get its IP address?
- How does a host know its local DNS server?
- How does a host know its subnet mask?
- How does a host know which router is its “gateway” to other networks?

# DHCP: Dynamic Host Configuration Protocol


*goal:* allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network

## *DHCP overview:*

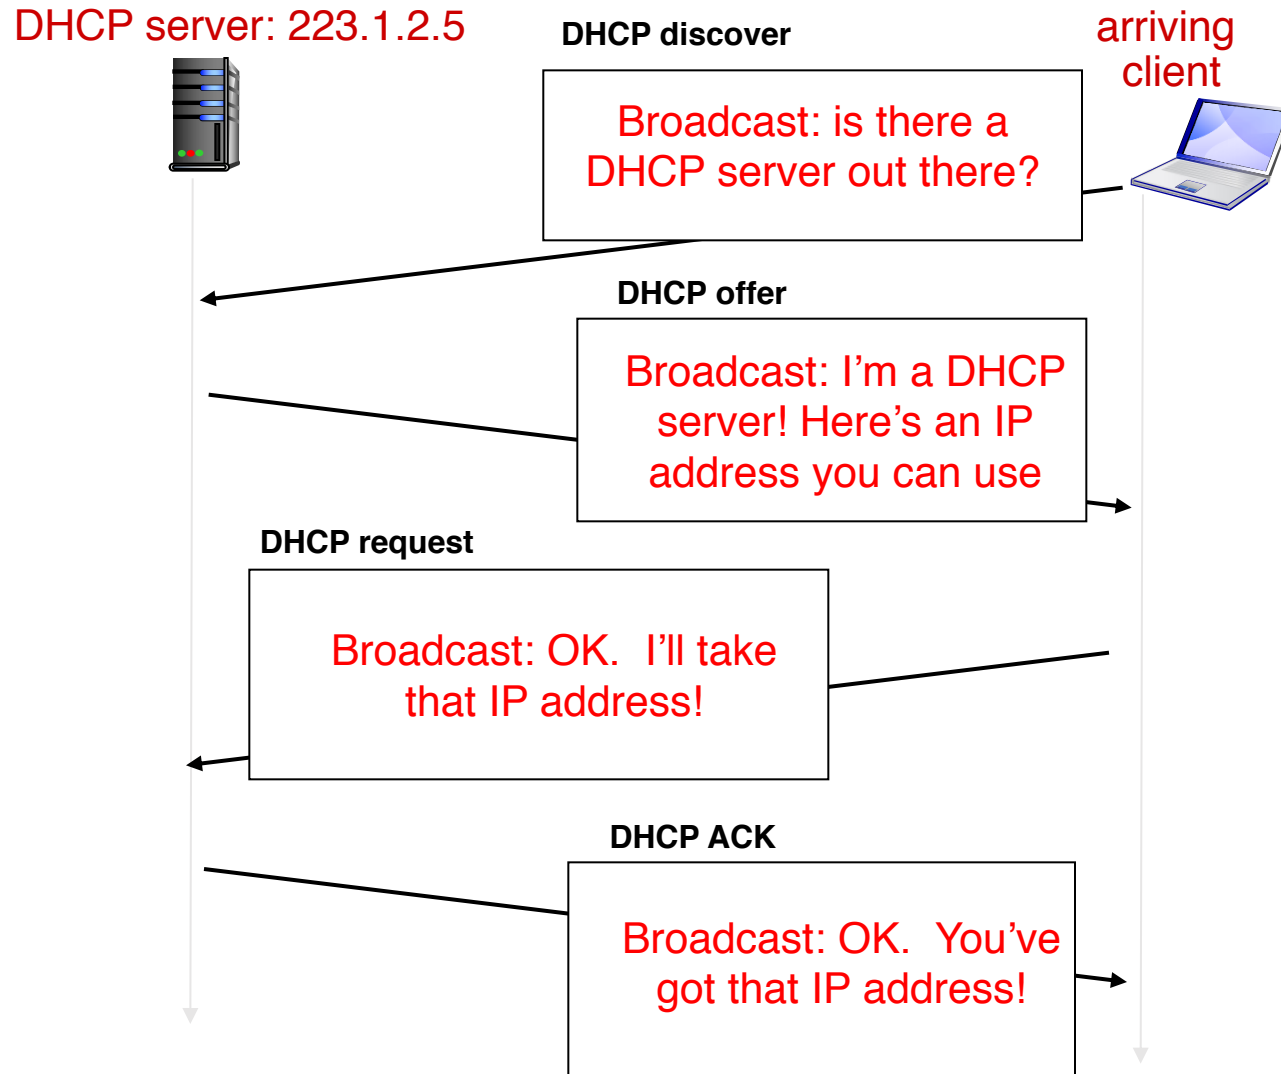
- host broadcasts “DHCP discover” msg [optional]
- DHCP server responds with “DHCP offer” msg [optional]
- host requests IP address: “DHCP request” msg
- DHCP server sends address: “DHCP ack” msg

# DHCP client-server scenario

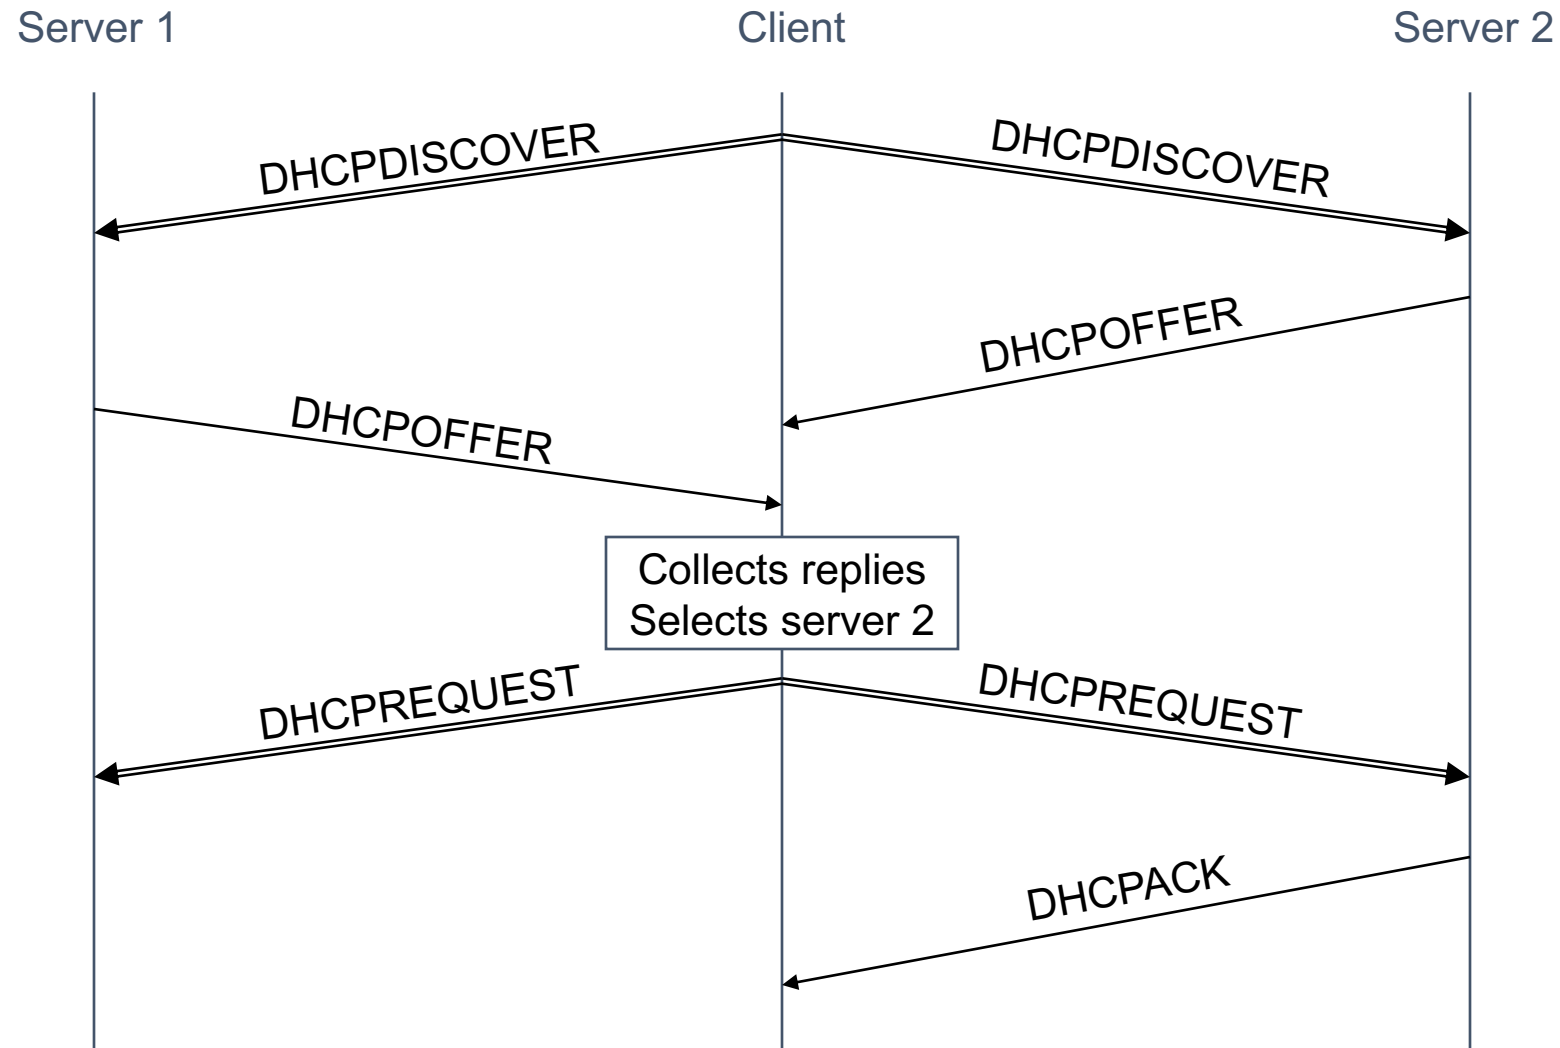




# DHCP client-server scenario



# DHCP Protocol




# DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:


- address of **first-hop router** for client to reach other subnets
  - Also called the **gateway router**
- name and IP address of the **local DNS server**
- **subnet mask** of the IP network the host is on
  - Useful to know whether other endpoint is inside or outside the current IP network

# DHCP Relay Agents

- DHCP relay agents allow DHCP servers to handle requests from other subnets




# DHCP: An example



- connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- Ethernet frame broadcast (dest: FFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demuxed to IP demuxed, UDP demuxed to DHCP

# DHCP: An example



- DHCP server formulates DHCP ACK containing client's IP address, subnet mask, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router

# Poll #4

- What functions does the DHCP protocol serve?
  - (a) discovering a usable IP address when a host joins
  - (b) discovering the local DNS server address
  - (c) discovering the gateway router of the local IP network
  - (d) all of the above

# Poll #5

- When a host joins a new network, how does it know the IP address of the DHCP server?
  - (1) this is a static configuration on all endpoints
  - (2) it broadcasts a message to discover DHCP servers
  - (3) neither of the above



# Summary

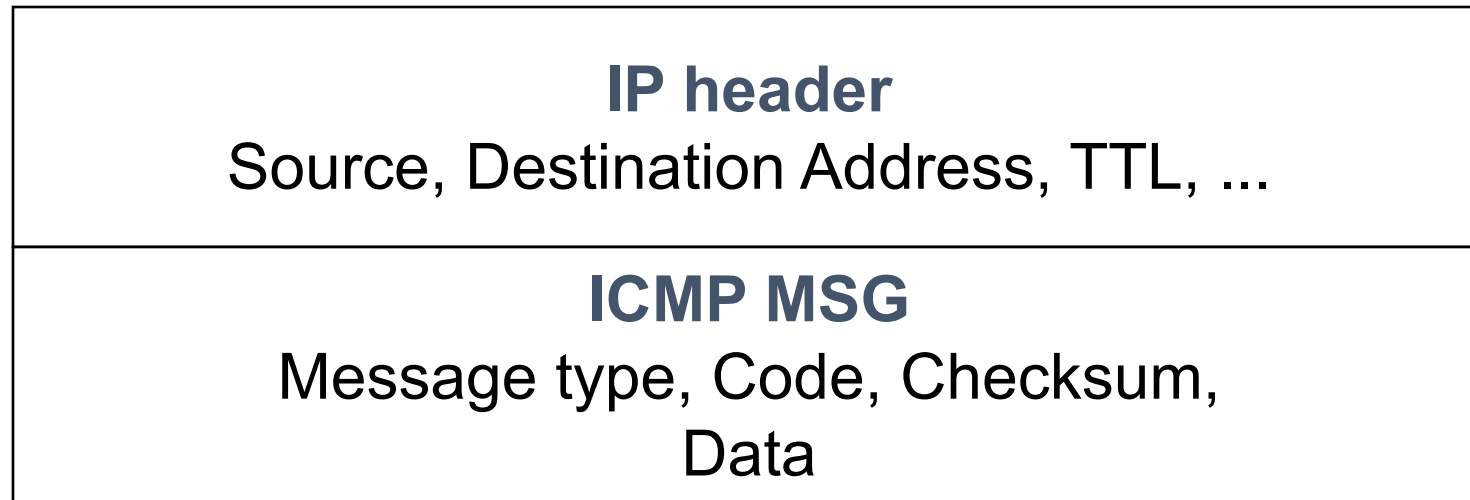
- IP addresses don't have to be manually configured into hosts
- DHCP allows “ignorant” hosts to receive IP addresses (and more) at start-up time
- DHCP solves important bootstrapping problems in attaching new hosts to a network

# Internet Control Message Protocol (ICMP)

# ICMP

- Protocol for error detection and reporting
  - tightly coupled with IP, unreliable
- ICMP messages delivered in IP packets
- ICMP functions:
  - Announce reachability and network errors
  - Announce “time exceeded” errors for IP packets
  - Announce network congestion
- ICMP assists **network troubleshooting** in general

# ICMP message



# ICMP: Internet Control Message Protocol

<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

Time for an  
activity!


# Specific uses of ICMP

- Echo request reply
  - Can be used to check if a host is alive
- Destination unreachable
  - Invalid address and/or port
- TTL expired
  - Routing loops, or too far away

# Ping

- Uses ICMP echo request/reply
- Source sends ICMP **echo request** message to the destination address
- Destination replies with an ICMP **echo reply** message containing the data in the original echo request message
- Source can calculate round trip time (RTT) of packets
- If no echo reply comes back, then the destination is **unreachable**

# Ping *(cont'd)*





# Ping example

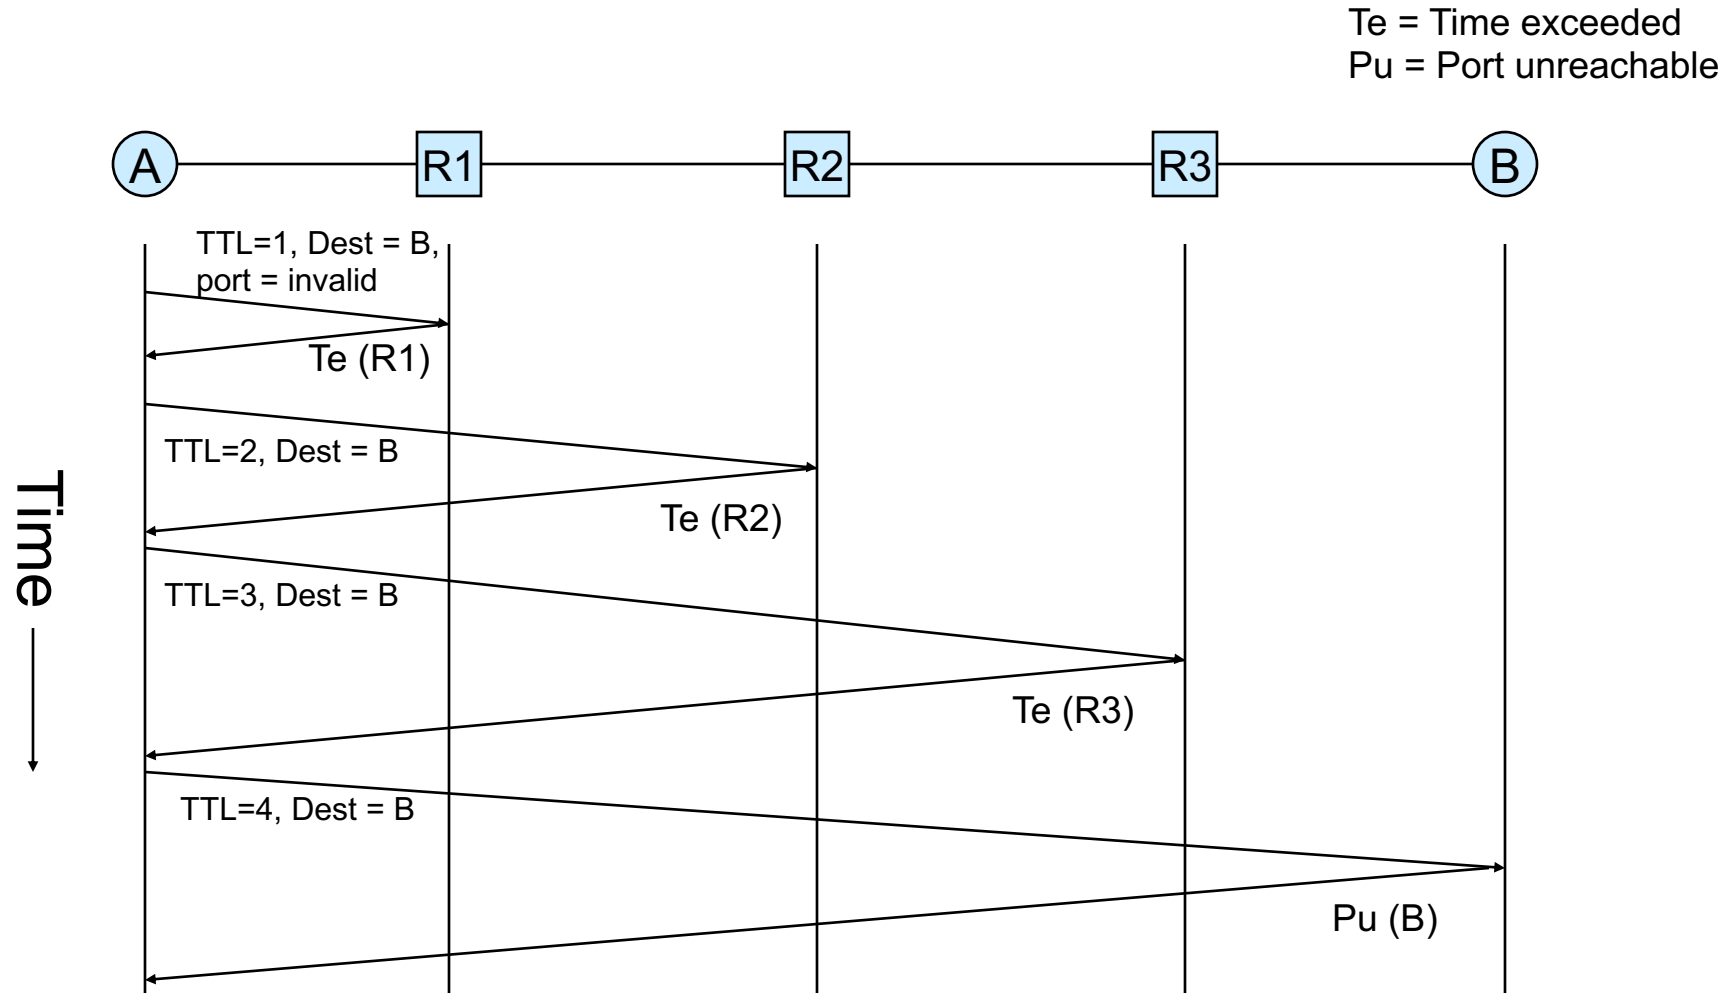
```
[flow:352-S20]$ ping google.com
PING google.com (172.217.10.110): 56 data bytes
64 bytes from 172.217.10.110: icmp_seq=0 ttl=56 time=5.727 ms
64 bytes from 172.217.10.110: icmp_seq=1 ttl=56 time=4.701 ms
64 bytes from 172.217.10.110: icmp_seq=2 ttl=56 time=5.954 ms
64 bytes from 172.217.10.110: icmp_seq=3 ttl=56 time=11.981 ms
64 bytes from 172.217.10.110: icmp_seq=4 ttl=56 time=6.084 ms
64 bytes from 172.217.10.110: icmp_seq=5 ttl=56 time=5.829 ms
64 bytes from 172.217.10.110: icmp_seq=6 ttl=56 time=8.667 ms
64 bytes from 172.217.10.110: icmp_seq=7 ttl=56 time=8.916 ms
64 bytes from 172.217.10.110: icmp_seq=8 ttl=56 time=4.537 ms
64 bytes from 172.217.10.110: icmp_seq=9 ttl=56 time=4.980 ms
```



# Traceroute

- Traceroute records the route that packets take
- A clever use of the IP **TTL** (time to live) field
- In general, when a router receives a packet, **it decrements the TTL** on that packet
- If TTL=0, router sends an ICMP **time exceeded** message back to sender
- traceroute **progressively increases** TTL of the packets it sends out
  - Every time an ICMP **time exceeded** message is received, record the sender's (router's) address
  - Repeat until the destination host is reached or an error message occurs
- If packet reaches the destination, the dest host usually sends an ICMP port unreachable

# Traceroute *(cont'd)*



# Traceroute example

```
[flow:352]$ traceroute google.com
traceroute to google.com (172.217.10.238), 64 hops max, 52 byte packets
 1  vlan451-sr03-hill-nbp.runet.rutgers.net (172.25.112.1)  9.278 ms  3.210 ms  3.124 ms
 2  xe-1-3-0-0-cr02-hill-nbp.runet.rutgers.net (172.29.6.65)  37.125 ms  2.912 ms  2.899 ms
 3  ae1-2000-cr10-hill-nbp.runet.rutgers.net (172.29.6.42)  3.078 ms  3.086 ms  3.016 ms
 4  ae5-2000-cr02-halsey-nwk.runet.rutgers.net (172.29.6.55)  3.693 ms  3.707 ms  3.793 ms
 5  ae2-0-er10-halsey-ext.runet.rutgers.net (172.29.8.6)  3.699 ms  3.693 ms  3.766 ms
 6  ae1-0-fw01-halsey-nwk.runet.rutgers.net (172.29.8.41)  4.019 ms  3.909 ms  3.750 ms
 7  et-2-2-0-0-er10-halsey-ext.runet.rutgers.net (172.29.8.46)  4.310 ms  4.181 ms  3.948 ms
 8  gateway-pni.google.com (128.6.1.114)  4.426 ms  3.901 ms  4.161 ms
 9  108.170.248.65 (108.170.248.65)  5.024 ms
    108.170.248.1 (108.170.248.1)  6.147 ms  6.165 ms
10  72.14.233.201 (72.14.233.201)  5.316 ms  5.426 ms  5.359 ms
11  lga25s59-in-f14.1e100.net (172.217.10.238)  5.410 ms  5.156 ms  5.135 ms
[flow:352]$ █
```


# Poll #6

- What's the IP header field that traceroute changes for each packet that it sends out to a given destination?
  - (1) IP source address
  - (2) IP source port
  - (3) IP TTL
  - (4) IP length

# Network Address Translation (NAT)


How do you survive in a society where names are scarce?

# NAT: Network Address Translation



**All** datagrams **leaving** local network have **same** source IP address: 138.76.29.7, with **different source IP port** numbers

# NAT: Network Address Translation





# NAT: Network Address Translation

- **Features:** local network uses just one IP address as far as outside world is concerned:
  - range of addresses not needed from ISP: **just one IP address for all devices**
  - can change addresses of devices in local network without notifying outside world
  - can change ISP without changing addresses of devices in local network
  - Devices inside local network not explicitly addressable
  - Devices inside local network invisible to the outside world (a security plus) unless the device inside connects first.

**Your home WiFi router uses NATs.**

**If you're home, you're behind a NAT right now.**

# The impact of NATs

```
[flow:352-S20]$ ifconfig en0
en0: flags=8863<UP,BROADCAST,SMART,RUNNING,SIMPLEX,MULTICAST> mtu 1500
    ether f0:18:98:1c:fc:36
    inet6 fe80::1036:7dea:82ee:e868%en0 prefixlen 64 secured scopeid 0xa
    inet 192.168.1.151 netmask 0xffffffff broadcast 192.168.1.255
    nd6 options=201<PERFORMNUD,DAD>
    media: autoselect
    status: active
[flow:352-S20]$ █
```



what's my ip address



All Images Videos News Maps | Answer

Settings ▼

Your IP address is 74.102.79.209 in [New Brunswick, New Jersey, United States \(08901\)](#)

# NAT: Network Address Translation

- 16-bit port-number field:
  - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - Routers should only work upto the network layer, not transport ports!
  - violates “end-to-end argument”
    - NAT must be taken into account by app designers
    - e.g., P2P applications like skype
  - Purists: address shortage should instead be solved by IPv6

# Think about...

- How do the hosts inside the home network get their IP addresses?
- How does your home router get its externally visible IP address?

# Poll #7

- The translation table at the NAT router contains
  - (1) Local IP address
  - (2) Local IP port
  - (3) Remote IP address
  - (4) Remote IP port
  - (5) All of the above


# Poll #8

- **Network Address Translation**
  - (1) reduces the demand for fresh IPv4 addresses
  - (2) keeps internal IP addresses hidden from the outside world
  - (3) requires support from the NAT gateway router
  - (4) all of the above

# Internet Protocol v6 (IPv6)

# Recent Developments: IPv6

- IPv4 has limited address space (32 bits) and is running out of addresses. 32 bits are not enough!
- More devices: phones, watches, your refrigerator(!), ...
- Real-time traffic and mobile users are also becoming more common



IP version 6



# IPv6: Main changes from IPv4


- Large address space:
  - 128-bit addresses (16 bytes)
  - Allows up to 340,282,366,920,938,463,463,374,607,431,768,211,456 unique addresses ( $3.4 \times 10^{38}$ )
- Fixed length headers (40 bytes)
  - Improves the speed of packet processing in routers
- IPv6 “options” processing happens through a separate mechanism

# IPv6 datagram format

*priority:* identify priority among datagrams in flow

*flow Label:* identify datagrams in same “flow”  
(concept of “flow” left undefined)

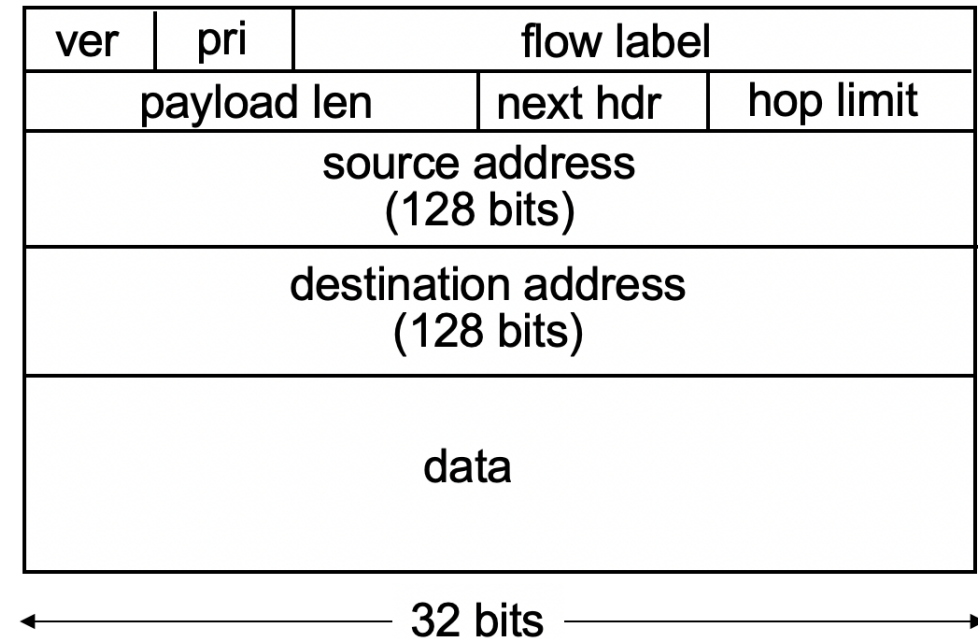
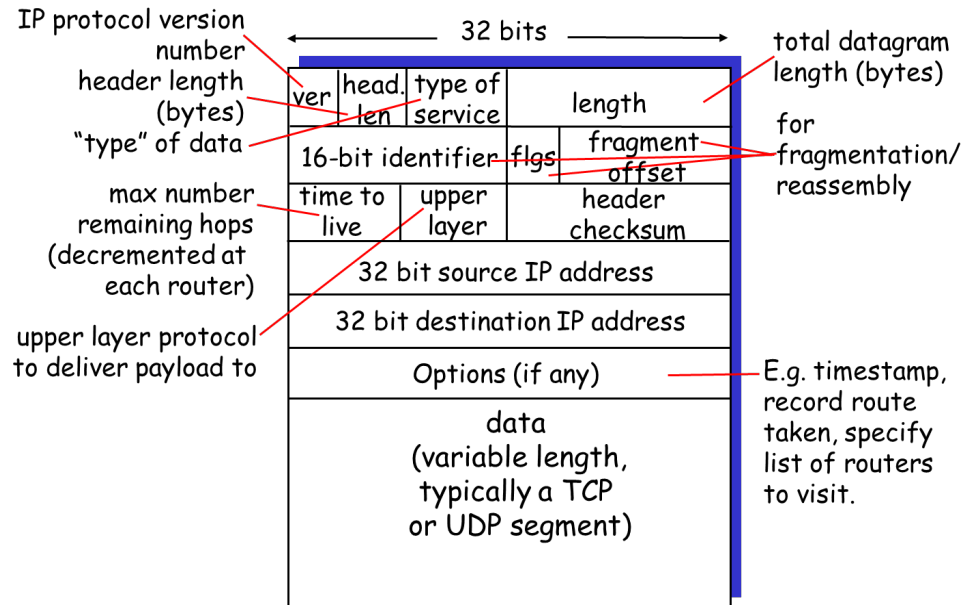
*next header:* identify upper layer protocol for data



# Other changes from IPv4

- *checksum*: removed entirely to reduce processing time at each hop
- *options*: allowed, but outside of header, indicated by “Next Header” field
- *ICMPv6*: new version of ICMP
  - additional message types, e.g. “Packet Too Big”
  - multicast group management functions

# IPv4 vs IPv6: Can you tell the differences?



# IPv6 Flows

- Support for “flows”
  - Flows help support real-time service in the Internet
  - A “flow” is a number in the IPv6 header that can be used by routers to see which packets belong to the same stream
  - Guarantees can then be assigned to certain flows
  - Example:
    - Packets from flow 10 should receive rapid delivery
    - Packets from flow 12 should receive reliable delivery

# IPv6 Addresses

- Classless addressing/routing (similar to CIDR)
- Notation: xx:xx:xx:xx:xx:xx:xx:xx
  - x = 4-bit hex number
  - contiguous 0s are compressed: 47CD::A456:0124
  - IPv6 compatible IPv4 address: ::128.64.18.87
    - First 96 bits are 0
  - Global unicast addresses start with 001....
  - 2000::/3 prefix

# IPv6: Adoption

- Google: 8% of clients access services via IPv6
- NIST: 1/3 of all US government domains are IPv6 capable
- *Long (long!) time for deployment, use*
  - 20 years and counting!
- Think of application-level changes in last 20 years: WWW, Facebook, streaming media, Skype, ...
- *Why?*