

# The Network Layer: Quality of Service

CS 352, Lecture 18, Spring 2020

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Course announcements

- Quiz 6 due on Tuesday at 10 PM
  - Covers lectures 17 and 18
- Project 2 due later today
- Project 3 will go out this weekend
  - Start early


# Review of concepts

- Border Gateway Protocol: a path-vector protocol
  - Announce **paths** to specific destinations through prefix + attributes
    - AS Path, next hop
- Business relationships influence route export and selection
- BGP paths get propagated inside ASes using iBGP
  - Used together with intra-domain routing to compute forwarding table
  - Two tables used together (BGP next hop, intra-domain lookup)
- Intra- and inter-AS routing protocols address different concerns
  - Policy vs. performance
- **Quality of service** within the network
  - ... when best-effort isn't enough for apps (e.g., robotic surgery over the net)
  - Traffic identified using **packet marking; now how to arbitrate traffic?**

# Quality of Service

How can the network make application performance better?


# Multiple classes of service: scenario



- example: 1Mbps VoIP, HTTP share 1.5 Mbps link.
  - HTTP bursts can congest router, cause audio loss
  - want to give **priority** to audio over HTTP

# Principles for QOS guarantees

- what if applications misbehave (VoIP sends higher than declared rate)
  - policing: force source adherence to bandwidth allocations
- *marking*, *policing* at network edge




## Principle 1: Isolation

Provide protection for one class from others

# Principles for QoS guarantees

- allocating *fixed* (non-sharable) bandwidth to flow: *inefficient* use of bandwidth if flows doesn't use its allocation



**Principle 2: Work conservation**  
While providing isolation, use resources as efficiently as possible

# Poll #1

- Where does contention between different traffic classes (e.g., resulting in long queues) typically occur within routers?
  - (a) switch fabric
  - (b) input line termination
  - (c) output port buffers
  - (d) forwarding table



# Poll #2

- What router mechanisms might you use to implement quality of service mechanisms?
  - (a) forwarding
  - (b) scheduling
  - (c) buffer management
  - (d) switching

# Packet scheduling for QoS

Shaping and policing

# Why care about packet scheduling?

- Significantly influences how packets are treated regardless of the endpoint transport
  - Implementations of **Quality of Service (QoS)** within large networks
  - Implications for **net neutrality** debates
- Intellectually interesting, foundational problem that a network solves
  - Classic Demers et al paper on scheduling (WFQ) has **~ 1500** citations
  - Important connections to other literature (e.g., job scheduling)
- Scheduling algorithms influence many daily life decisions 😊



# Scheduling vs. Buffer Management

- How packets **enter** vs. how packets **leave** the switch buffer
  - Common architecture: **shared-buffer switches**, where buffer memory is shared between different output ports
  - Typical buffer management strategy: tail-drop
- How should buffer memory be partitioned across ports?
- Static partitioning?
  - **Inefficient**: even if port 1 has nothing to send, might drop port 2
- Also want **fair sharing** of buffer
  - If output port 1 is congested, why should port 2 traffic suffer?
- State of the art: **demand-aware buffer sharing algos**



# Packet scheduling for QoS

- Choose next queued packet to send on outgoing link
- Common scheduling disciplines
  - first come first served (FIFO)
  - simple multi-class priority
  - round robin
  - weighted fair queueing (WFQ)
  - Rate limiting
- All except rate-limiting are work-conserving



# Poll #3

- When cars from two lanes on a congested freeway take turns to merge into one lane, what scheduling discipline is being implemented?
  - (a) priority-queueing
  - (b) weighted fair queueing
  - (c) rate limiting
  - (d) none of the above

# Isolation by Rate Limits

# Providing Isolation through Rate Limiting

Three commonly used terms:

- *(long term) average rate*: how many pkts can be sent per unit time (in the long run)
  - crucial question: **what is the interval length?** 100 packets per sec or 6000 packets per min have same average! Instantaneous behaviors of the two may look very different.
- *peak rate*: e.g., 6000 pkts per min (ppm) avg.; 1500 ppm peak rate
- *(max.) burst size*: max number of pkts sent consecutively (with no intervening idle)




# Shaping and Policing

---

Policing	<p>Enforces rate by <i>dropping</i> excess packets immediately</p> <ul style="list-style-type: none"><li>– Can result in high loss rates</li><li>+ Does not require memory buffer</li><li>+ No RTT inflation</li></ul>
Shaping	<p>Enforces rate by <i>queueing</i> excess packets</p> <ul style="list-style-type: none"><li>+ Only drops packets when buffer is full</li><li>– Requires memory to buffer packets</li><li>– Can inflate RTTs due to high queueing delay</li></ul>

---

# Mechanism (1): Leaky Bucket Shaper



# Mechanism (1): Leaky Bucket Shaper

- Packets may be generated in a **bursty** manner, but after they pass through the leaky bucket, they enter the network **evenly spaced**
  - The “bucket” buffers packets up to a certain point
  - If the bucket is full, packets are dropped.
- May be used in conjunction with resource reservation to police the host’s resource use
  - E.g., at the host-network interface, allow packets into the network at a constant rate
- May be used in the core of a network to limit bandwidth use

# Shaping traffic with leaky buckets

- The leaky bucket is a **traffic shaper**: It changes the downstream arrival times & characteristics of packets passing through it
- Traffic shaping makes traffic more manageable and more predictable downstream (e.g., always under a certain rate)
- Usually, a system/network administrator would set the rate at which packets may be sent through the leaky bucket
- Administrator also sets up policies to map any connection that started up to a leaky bucket (and rate) of its own

# Issues with a leaky bucket


- Many short transfers just have a few packets
  - E.g., web requests and responses
  - Enforcing rate limit for those can significantly delay completion
- For a leaky bucket shaper, **average rate == peak rate**
- Sometimes, we wish to have peak rate higher than avg rate
- For this purpose we use a **token bucket**, which is a burst-tolerant version of a leaky bucket

# Token buckets

# Mechanism (2): Token Bucket Shaper

*token bucket*: limit input to specified *burst size* and *average rate*

- Tokens generated at rate  $r$  tokens/sec, put into bucket
- Bucket can hold  $b$  tokens. Tokens are dropped if bucket is full
- A packet can leave as long as a token is available
- Over interval of time  $t$ : number of packets that leave the token bucket is less than or equal to  $(r * t + b)$



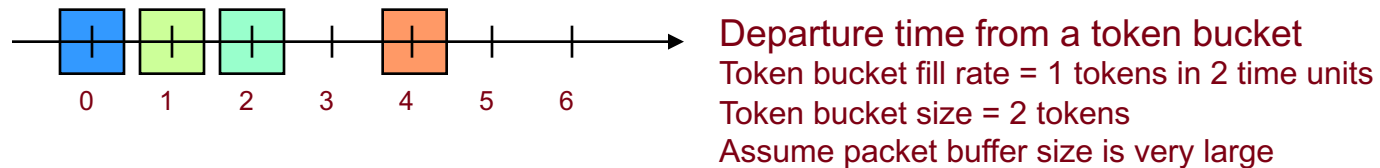


# Token Bucket Shaper

- There is a **bucket for tokens** and a **packet buffer for packets**
  - Packets will be dropped if the associated packet buffer is full
  - A token bucket **policer** doesn't contain a packet buffer: any packet arriving without a token is immediately dropped
- Bucket of tokens enables small bursts to go through unscathed
  - Short flows can exceed rate limit  $r$ , since they use up the reserve in the bucket
  - Long flows will fit into the rate limit  $r$  over longer periods of time, since once they use up the reserve, they are limited by the token-fill rate  $r$




# Token Bucket vs. Leaky Bucket

## Case 1: Small burst of packet arrivals




# Token Bucket vs. Leaky Bucket

## Case 2: Large burst of packet arrivals




Arrival time at bucket



Departure time from a leaky bucket shaper

Leaky bucket rate = 1 packet in 2 time units

Assume packet buffer size is very large



Departure time from a token bucket shaper

Token bucket fill rate = 1 token in 2 time units

Token bucket size = 2 tokens

Assume packet buffer size is very large


??

Departure time from a token bucket policer

Token bucket fill rate = 1 token in 2 time units

Token bucket size = 2 tokens



# Poll #4



- Suppose the packets above arrived at a token bucket policer with a fill rate of 1 token per 2 time units, and a token bucket size of 2 tokens. What happens to packet 3 above as it arrives?
  - (a) Buffered into the packet buffer
  - (b) Transmitted out of the token bucket
  - (c) Packet is dropped
  - (d) None of the above

# Token buckets can be used to provide delay QoS guarantees

- token bucket + WFQ combine to provide a guaranteed upper bound on delay



# Impact of Token Bucket Policers


The Internet has tons of them, and they affect application performance

# Google study from 2016

Region	Policed segments		Loss rate	
	(among lossy)	(overall)	(policed)	(non-pol.)
India	6.8%	1.4%	28.2%	3.9%
Africa	6.2%	1.3%	27.5%	4.1%
Asia (w/o India)	6.5%	1.2%	22.8%	2.3%
South America	4.1%	0.7%	22.8%	2.3%
Europe	5.0%	0.7%	20.4%	1.3%
Australia	2.0%	0.4%	21.0%	1.8%
North America	2.6%	0.2%	22.5%	1.0%

**Table 2: % segments policed among lossy segments ( $\geq 15$  losses, the threshold to trigger the policing detector), and overall. Avg. loss rates for policed and unpoliced segments.**


# Impact on TCP



**Figure 1: TCP sequence graph for a policed flow: (1 and 4) high throughput until token bucket empties, (2 and 5) multiple rounds of retransmissions to adjust to the policing rate, (3) idle period between chunks pushed by the application.**

# Policing losses impact applications

- Video rebuffer rate: rebuffer time / overall watch time



**Figure 9: Rebuffer to watch time ratios for video playbacks. Each had at least one chunk with a goodput of 300 kbps, 1.5 Mbps, or 3 Mbps ( $\pm 15\%$ ).**