

# Multimedia: Streaming Video & Audio

CS 352, Lecture 22, Spring 2020

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Course announcements

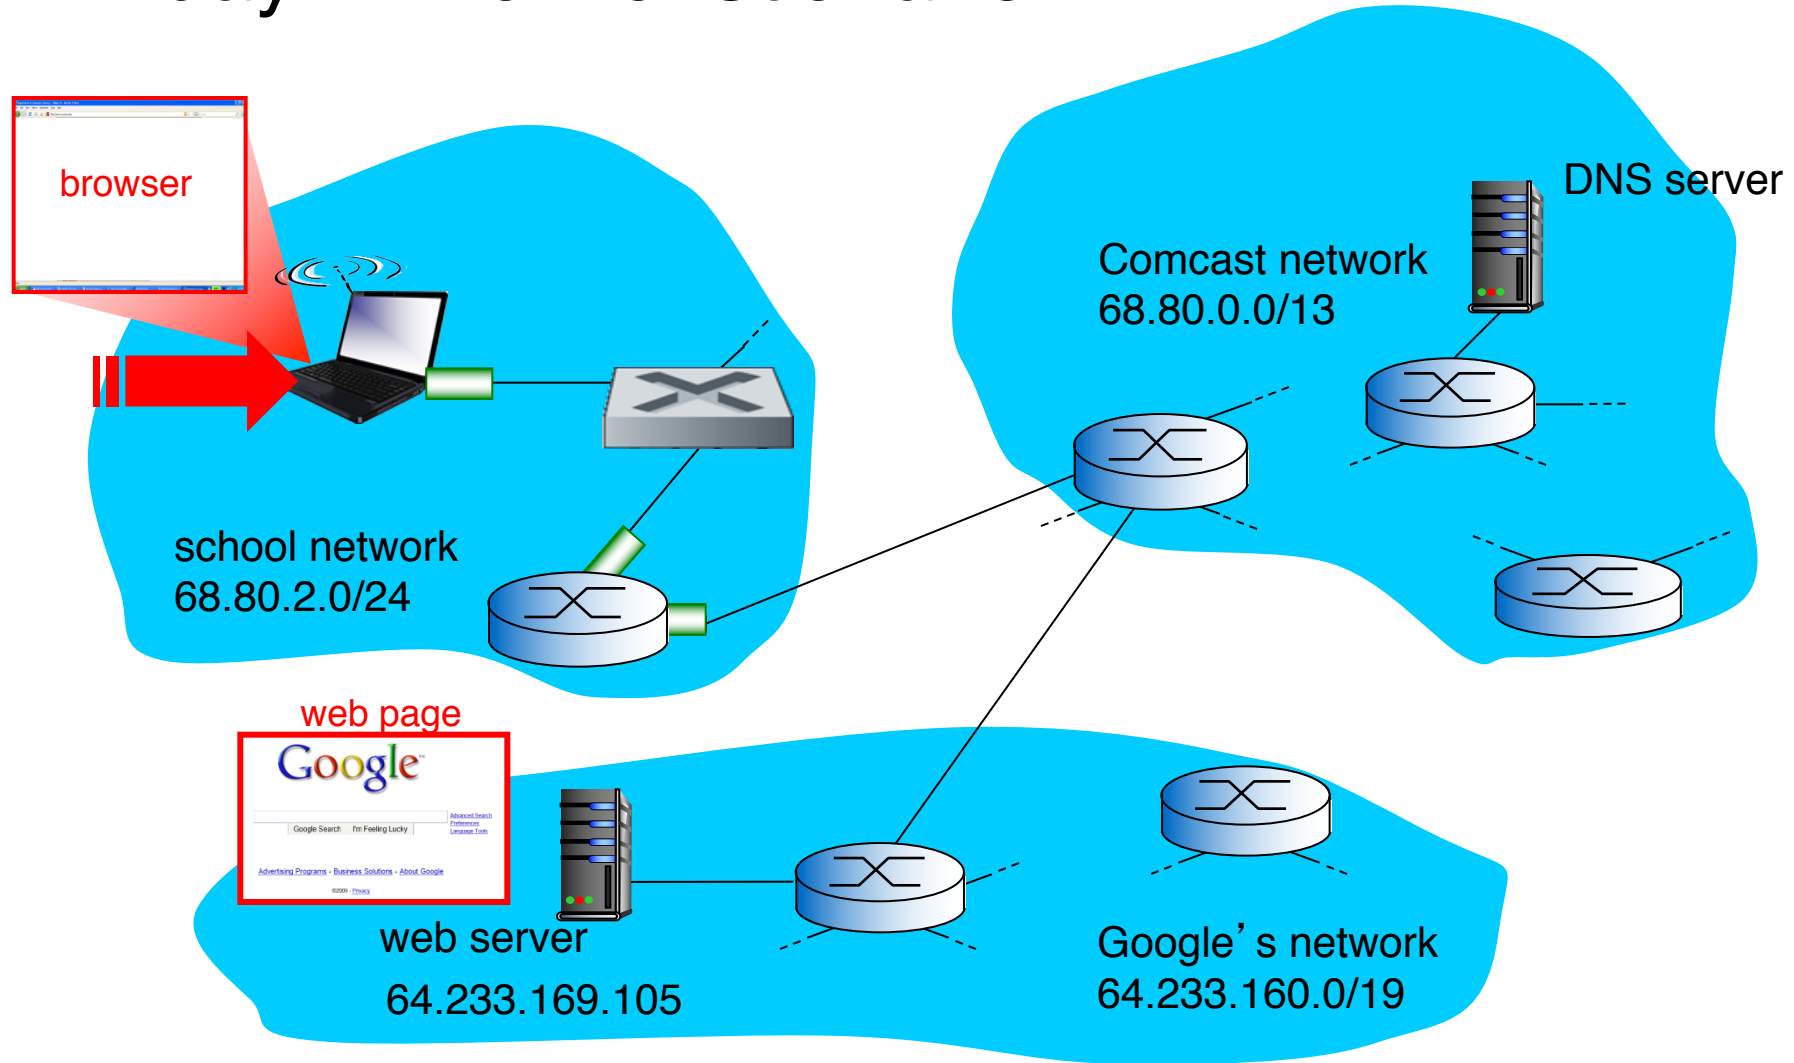
- Quiz 8 (the last one!) will go online later today
  - Due Tuesday at 10 PM on Sakai
- Project 3 due next Friday

# Synthesis of protocols


## *Synthesis:* a day in the life of a web request

- Our journey down protocol stack complete!
  - application, transport, network, link
- putting-it-all-together: synthesis!
  - *goal:* identify, review, understand protocols (at all layers) involved in seemingly simple scenario: requesting www page
  - *scenario:* student attaches laptop to campus network, requests/receives [www.google.com](http://www.google.com)

# A day in the life: scenario




# A day in the life... connecting to the Internet



- connecting laptop needs to get its own IP address, addr of first-hop router, addr of DNS server: use *DHCP*
- DHCP request *encapsulated* in *UDP*, encapsulated in *IP*, encapsulated in *802.3* Ethernet
- Ethernet frame *broadcast* (dest: FFFFFFFFFFFFFFFF) on LAN, received at router running *DHCP* server
- Ethernet *demuxed* to IP demuxed, UDP demuxed to DHCP


# A day in the life... connecting to the Internet



- DHCP server formulates *DHCP ACK* containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation at DHCP server, frame forwarded through LAN, demultiplexing at client
- DHCP client receives DHCP ACK reply

*Client now has IP address, knows name & addr of DNS server, IP address of its first-hop router*

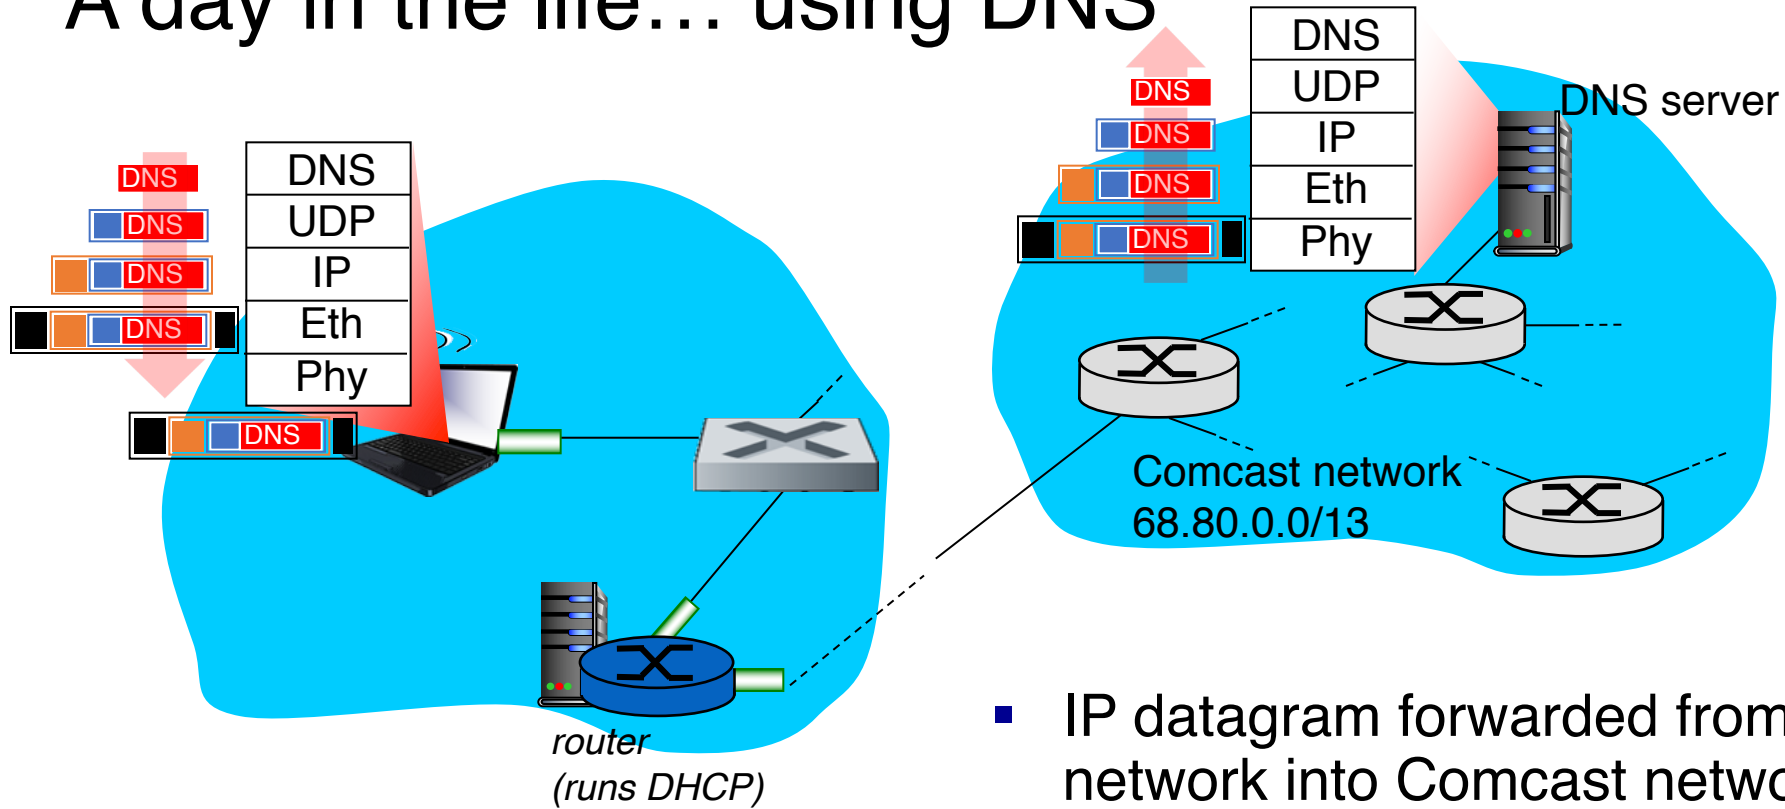
# A day in the life... ARP (before DNS, before HTTP)



- before sending *HTTP* request, need IP address of `www.google.com`: *DNS*
- DNS query created, encapsulated in UDP, encapsulated in IP, encapsulated in Eth. To send frame to router, need MAC address of router interface: *ARP*
- *ARP query* broadcast, received by router, which replies with *ARP reply* giving MAC address of router interface
- client now knows MAC address of first hop router, so can now send frame containing DNS query



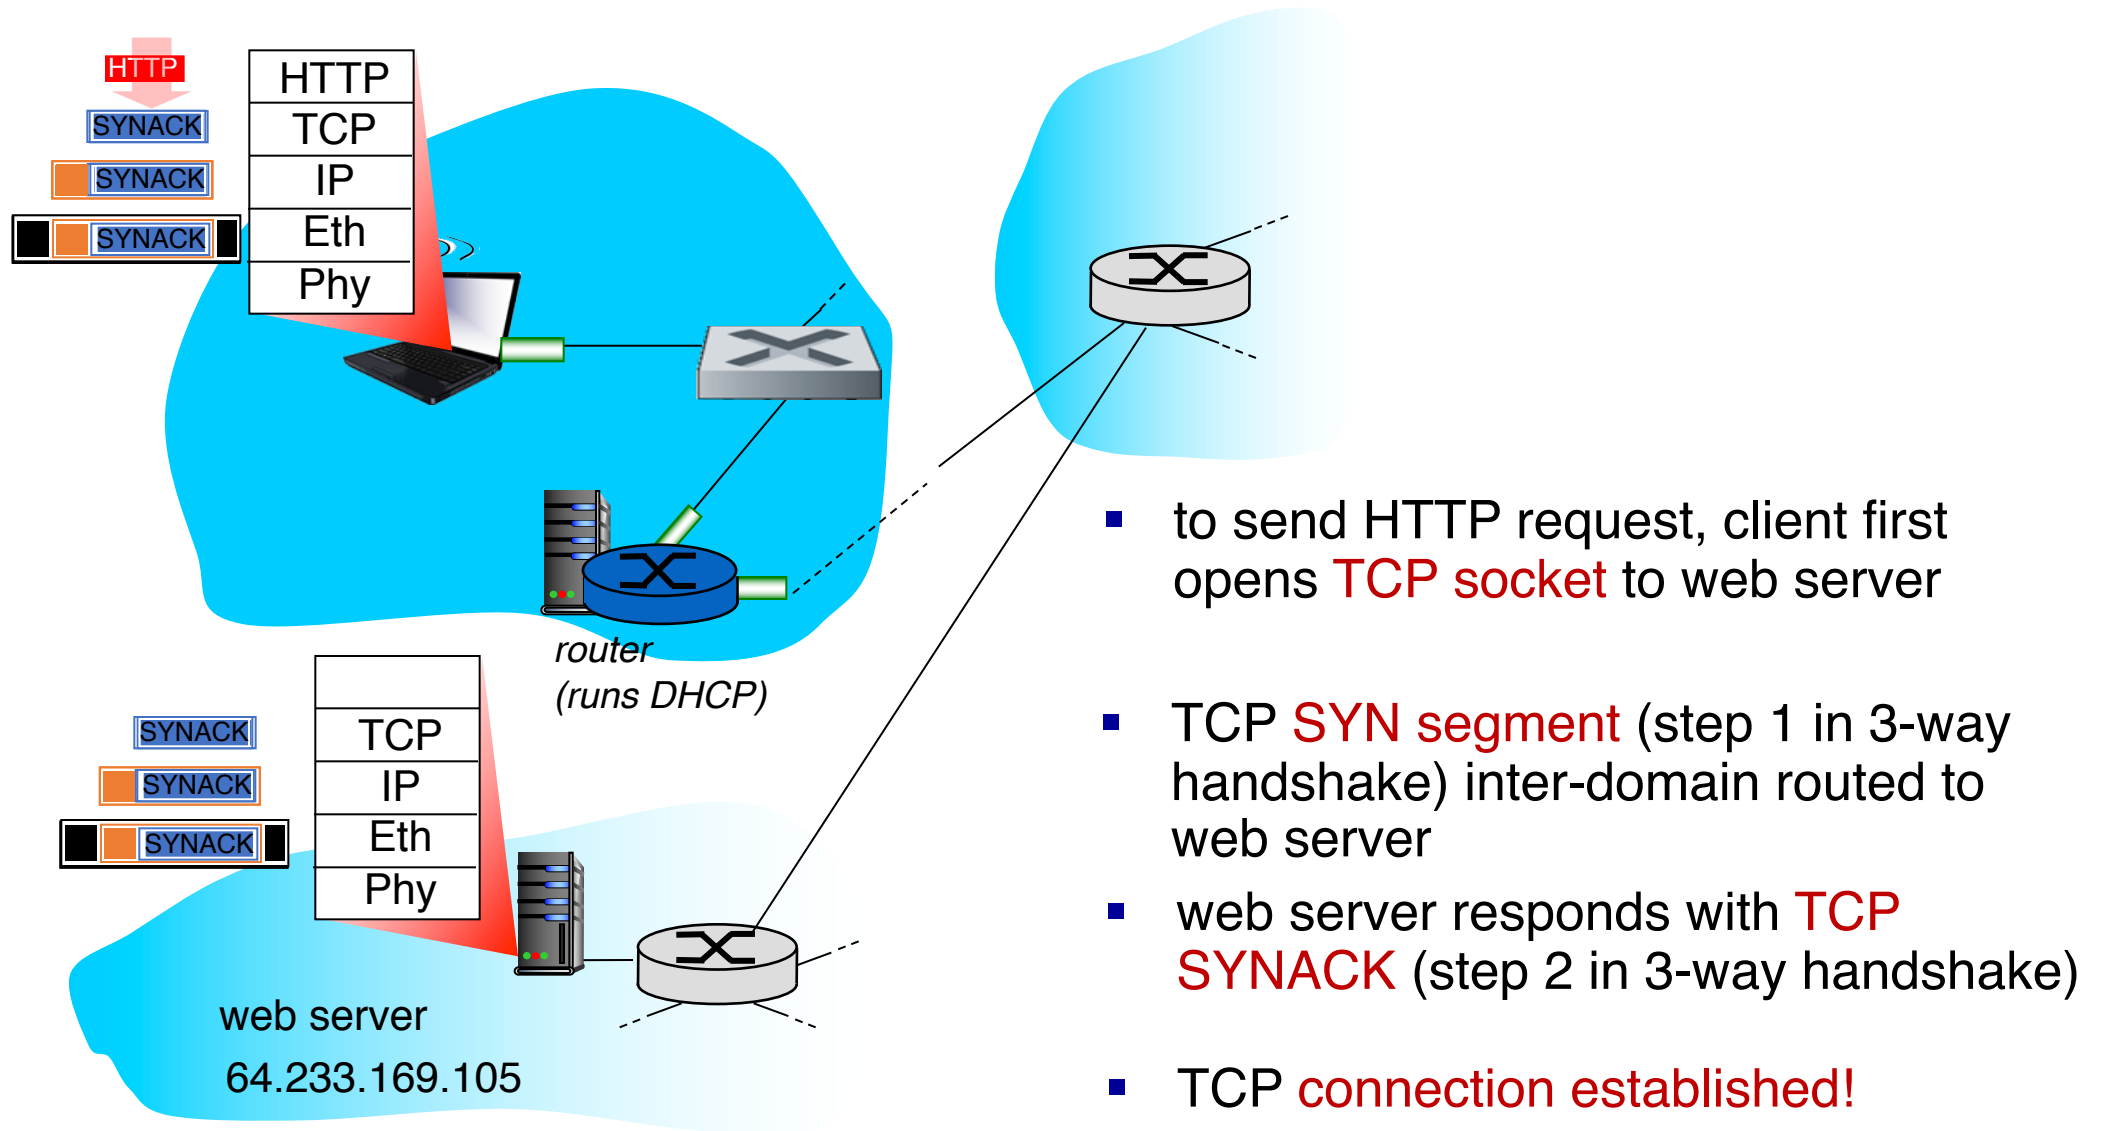
# A day in the life... using DNS



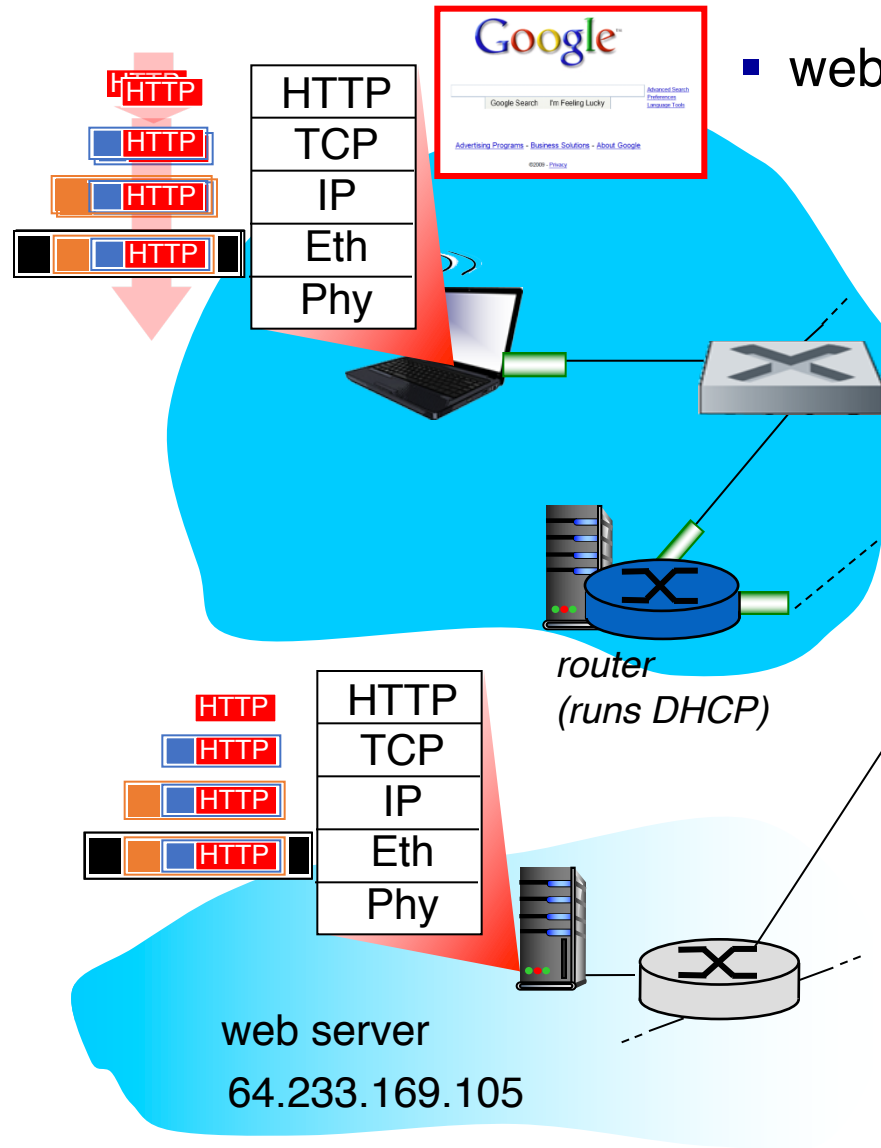
- IP datagram containing DNS query forwarded via LAN switch from client to 1<sup>st</sup> hop router

- IP datagram forwarded from campus network into Comcast network, routed (tables created by **RIP**, **OSPF**, **IS-IS** and/or **BGP** routing protocols) to DNS server
- demuxed to DNS server
- DNS server replies to client with IP address of [www.google.com](http://www.google.com)

# A day in the life...TCP connection carrying HTTP



# A day in the life... HTTP request/reply



- web page **finally (!!!)** displayed

- **HTTP request** sent into TCP socket
- IP datagram containing HTTP request routed to `www.google.com`
- web server responds with **HTTP reply** (containing web page)
- IP datagram containing HTTP reply routed back to client

# Multimedia Networking

# Multimedia networking

- Many applications on the Internet use audio or video
- IP video traffic will be 82 percent of all IP traffic [...] by 2022, up from 75 percent in 2017
- Internet video surveillance traffic will increase sevenfold between 2017 to 2022
- Internet video to TV will increase threefold between 2017 to 2022.
- Consumer Video-on-Demand (VoD) traffic will nearly double by 2022

Source: Cisco visual networking index 2017--22




# What's different about these applications?

- Traditional applications (HTTP(S), SMTP)
  - Delay tolerant but not loss tolerant
  - Data used *after* transfer complete
- But multimedia applications are often “real time”
  - Data delivery time *during transfer* has implications
- Video/audio streaming
  - Delay-sensitive
- Real-time audio and video
  - Delays  $> 400$  ms for audio is a bad user experience
  - Somewhat loss tolerant

# Digital representation of audio and video

# Digital representation of audio




- Must convert analog signal to digital representation
- Sample
  - How many times (twice the max frequency in the signal)
- Quantize
  - How many levels or bits to represent each sample
  - More levels → more accuracy
  - More levels → more bits to store & more bandwidth to transmit
- Compress
  - Compact representation of quantized values



# Audio representation

- analog audio signal sampled at constant rate
  - telephone: 8,000 samples/sec
  - CD music: 44,100 samples/sec
- each sample quantized, i.e., rounded
  - e.g.,  $2^8=256$  possible quantized values
  - each quantized value represented by bits, e.g., 8 bits for 256 values




# Audio representation

- example: 8,000 samples/sec, 256 quantized values
- Bandwidth needed: 64,000 bps
- receiver converts bits back to analog signal:
  - some quality reduction

## Example rates


- CD: 1.411 Mbps
- MP3: 96, 128, 160 Kbps
- Internet telephony: 5.3 Kbps and up



# Video representation


- Video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- Digital image: array of pixels
  - each pixel represented by bits
- Coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)
- Coding/decoding algorithm often called a **codec**

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$




frame  $i+1$

# Video representation


- **Video bit rate:** effective bits per second of the video after encoding
- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
  - MPEG 1 (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$

# Multimedia networking: 3 application types

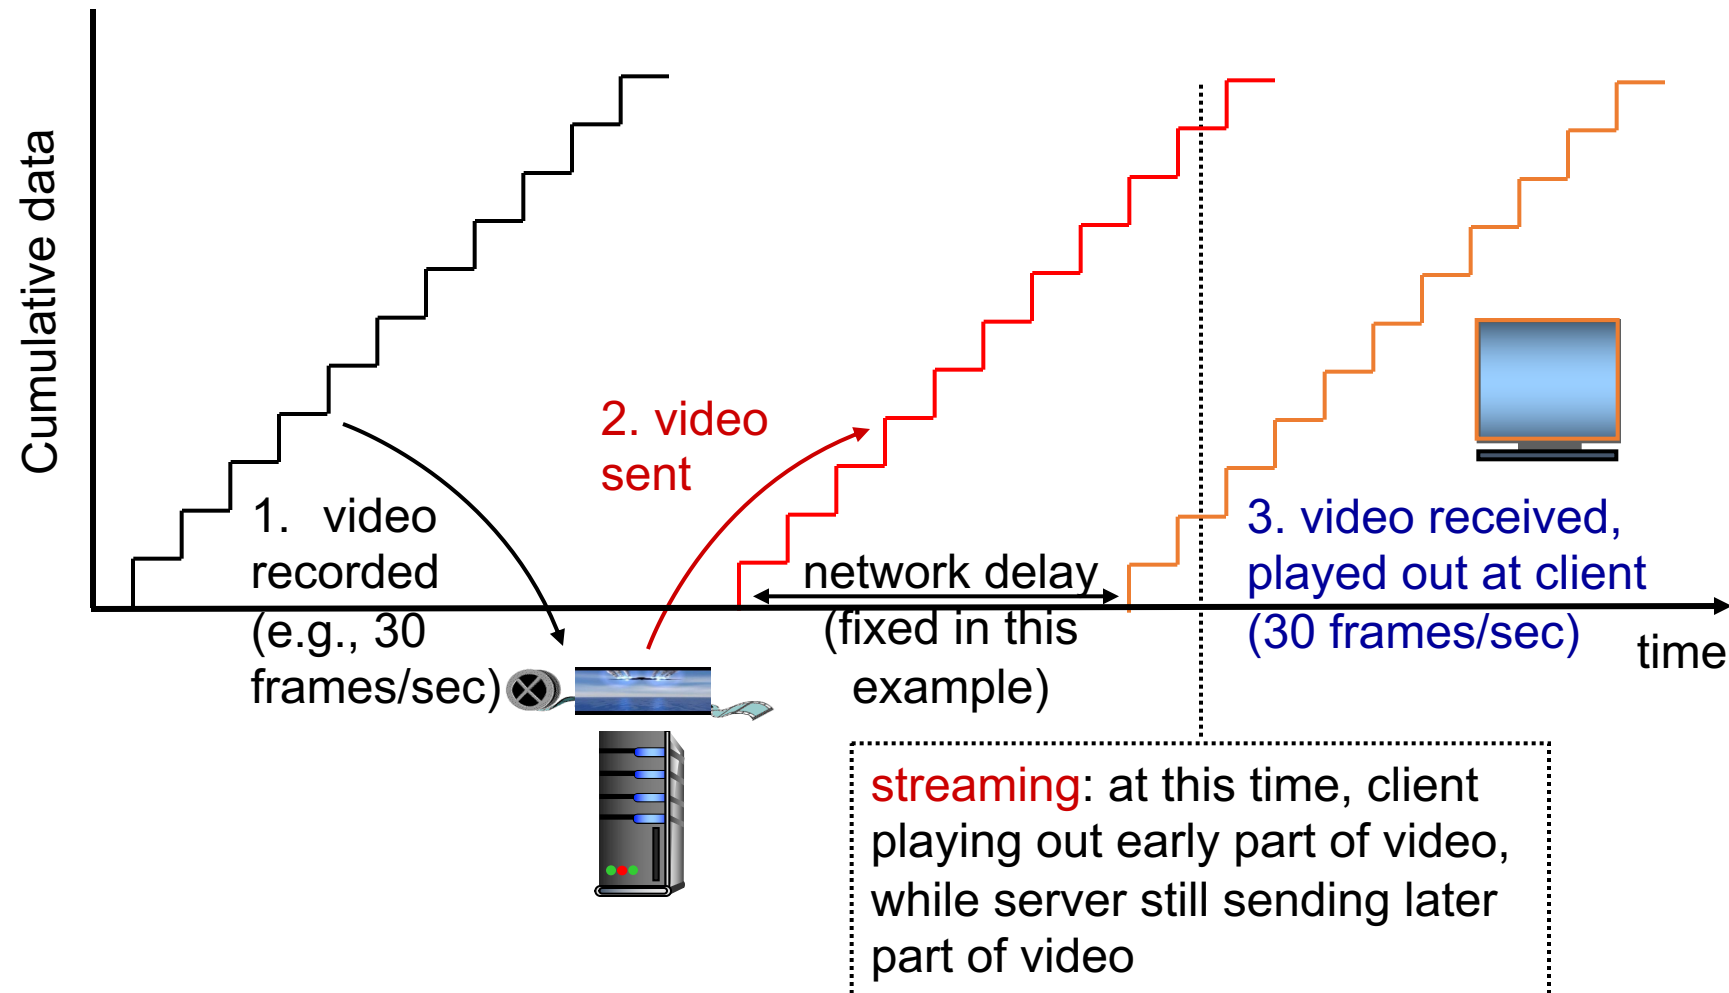
- *streaming, stored* audio, video
  - *streaming*: can begin playout before downloading entire file
  - *stored (at server)*: can transmit faster than audio/video will be rendered (implies storing/buffering at client)
  - e.g., YouTube, AmazonPrime, Disney, Netflix, Hulu
- *conversational* voice/video over IP
  - interactive nature of human-to-human conversation limits delay tolerance
  - e.g., Skype
- *streaming live* audio, video
  - e.g., live sporting event

# Streaming video

# Streaming stored content

- Media is prerecorded
- Client downloads an initial portion and starts viewing
- Rest downloaded as time progresses
- No need to wait for entire content to be downloaded
- Can change content sites mid-stream based on network conditions

# Streaming stored video:






# Streaming stored video: challenges


- **continuous playout constraint**: once client playout begins, playback must match original timing
  - ... but **network delays are variable** (jitter), so will need **client-side buffer** to match playout requirements
- other challenges:
  - client interactivity: pause, fast-forward, rewind, jump through video
  - video packets may be lost, retransmitted

# Streaming stored video: revisited

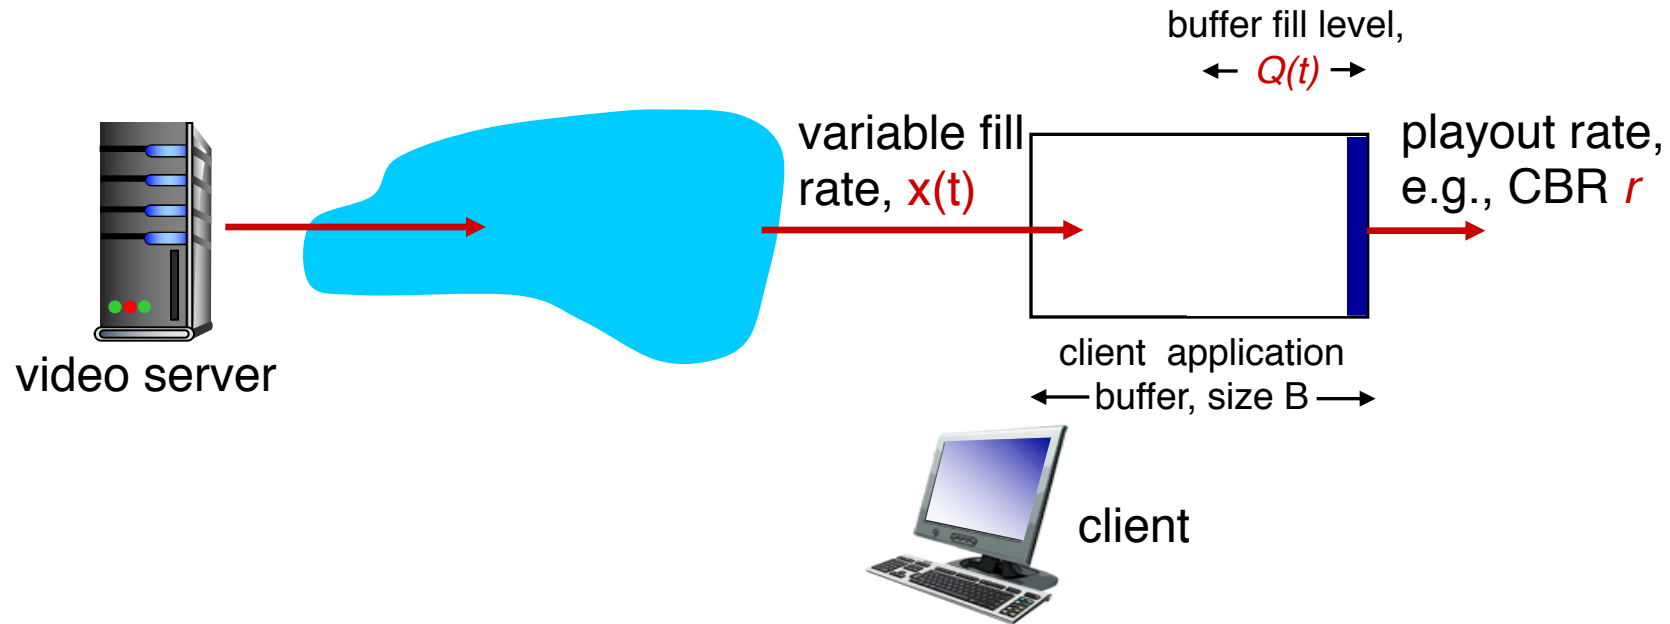


- *client-side buffering and playout delay:*  
compensate for network-added delay, delay jitter

# Client-side buffering, playout




# Client-side buffering, playout



1. Initial fill of buffer until playout begins at  $t_p$
2. playout begins at  $t_p$ ,
3. buffer fill level varies over time as fill rate  $x(t)$  varies and playout rate  $r$  is constant

# Client-side buffering, playout



*playout buffering: average fill rate ( $\bar{x}$ ), playout rate ( $r$ ):*


- $\bar{x} < r$ : buffer eventually empties (causing freezing of video playout until buffer again fills)
- $\bar{x} > r$ : buffer will not empty, provided initial playout delay is large enough to absorb variability in  $x(t)$ 
  - *initial playout delay tradeoff*: buffer starvation less likely with larger delay, but larger delay until user begins watching

# Streaming multimedia: UDP

- server sends at rate appropriate for client
  - often: send rate = encoding rate = constant rate
  - transmission rate can be oblivious to congestion levels
- short playout delay (2-5 seconds) to remove network jitter
- error recovery: application-level, time permitting
- RTP [RFC 2326]: multimedia payload types
- UDP traffic may *not* get through firewalls

# Streaming multimedia: HTTP/TCP

- multimedia file retrieved via HTTP GET
- send at maximum possible rate under TCP



- fill rate fluctuates due to TCP congestion control, retransmissions (in-order delivery)
- larger playout delay: smooth TCP delivery rate
- HTTP/TCP passes more easily through firewalls

# Streaming multimedia with DASH

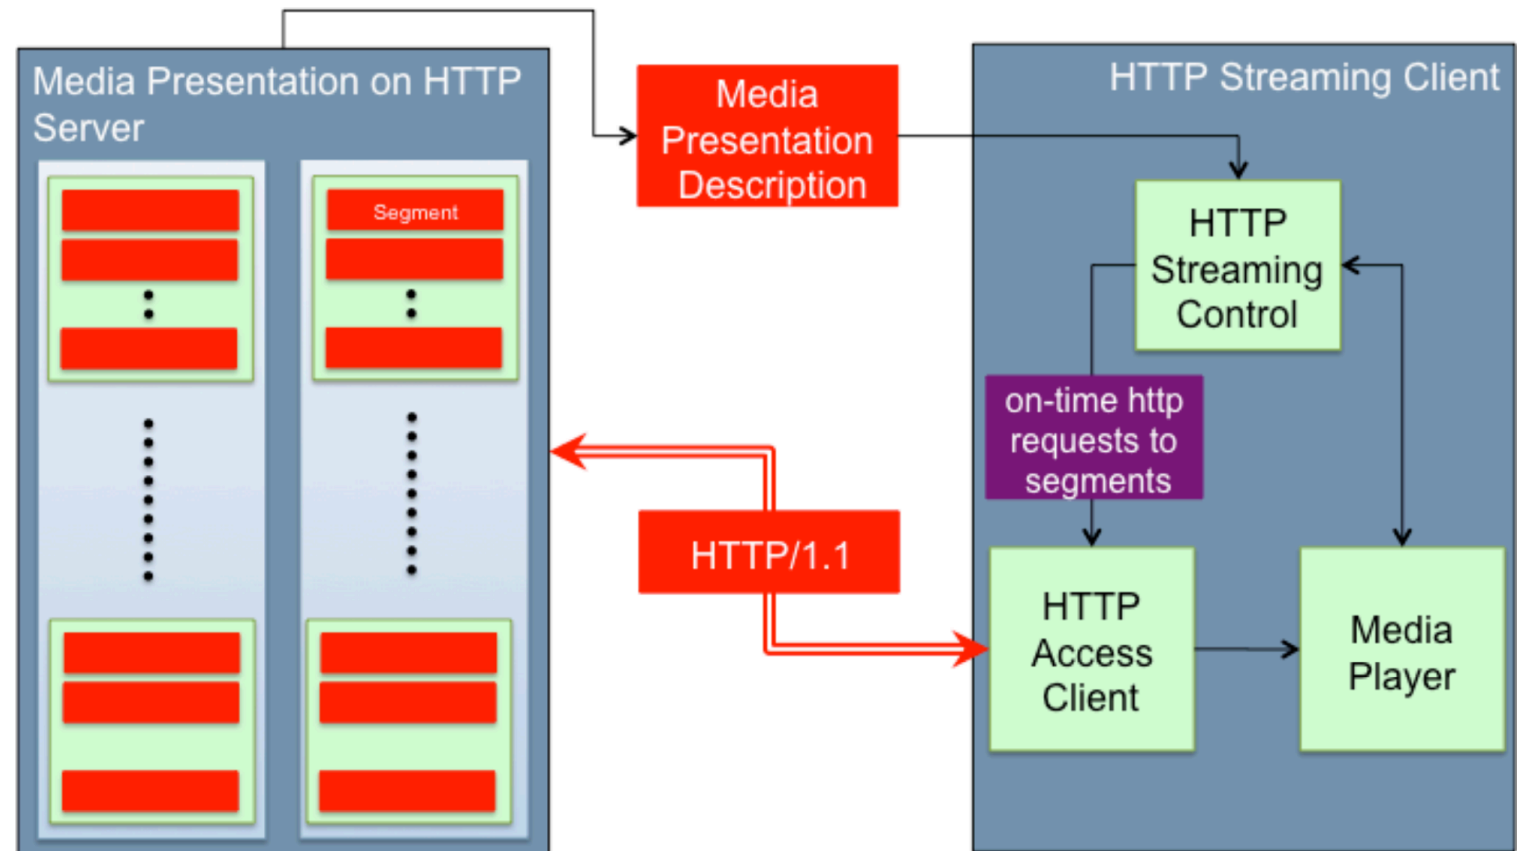
- Dynamic Adaptive Streaming over HTTP
- Used by Netflix and other video streaming services
- Client-centric approach to video delivery
  - **Adaptive:** Client performs video bit rate adaptation
  - **Dynamic:** Can retrieve a single video from multiple sources
- Retain benefits of existing Internet and end host systems
- Server is standard HTTP server
  - Provides video/audio content in multiple formats and encodings
  - DASH allows the use of CDNs for data delivery




# Dynamic Adaptive Streaming over HTTP (DASH)

# DASH: Key ideas

- Content **chunks**
- Each chunk can be **independently** retrieved
- Client-side algorithms to determine and request a **varying** bit rate for each chunk
  - Goal: ensure good quality of service



# DASH Data model




Media has several periods

Each period has several **Adaptation Sets**: Audio, video, close caption

Several **Representations** (ex: codecs, bit rates) per Adaptation set


Several **Chunks/Segments** per Representation

# Dynamic bit rate changing of streams

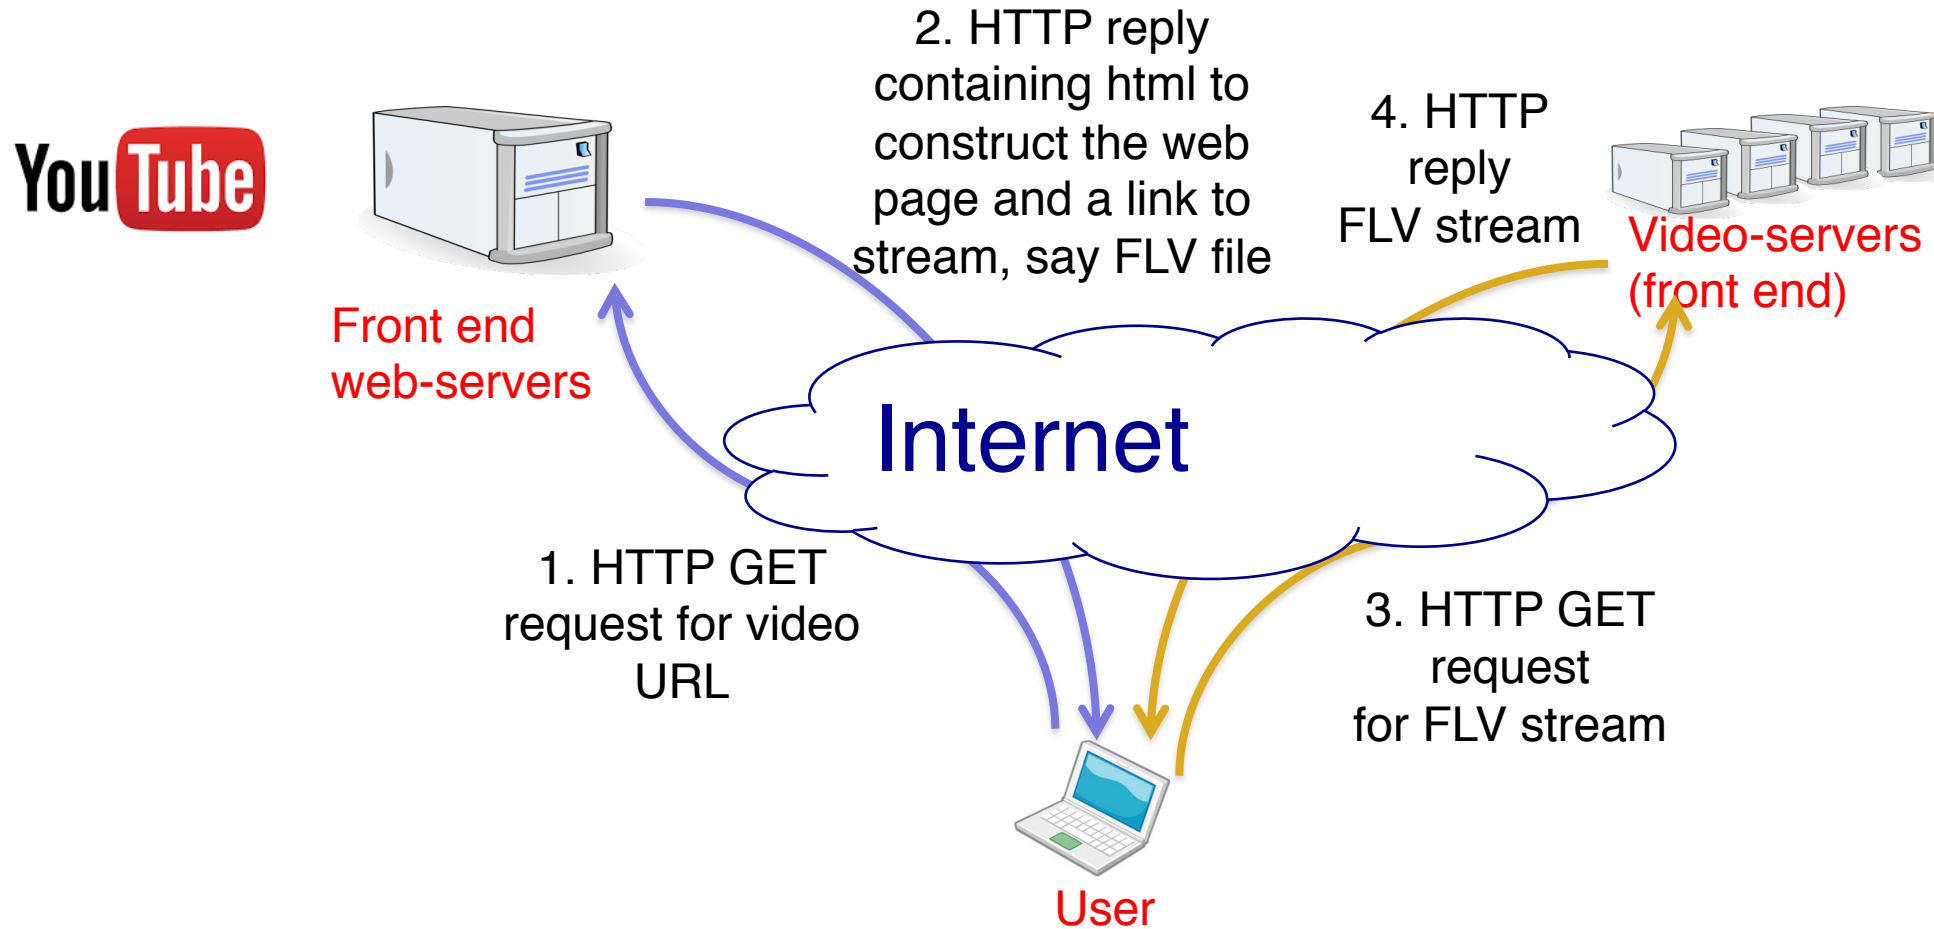


# Media Presentation Descriptor (MPD)

- MPD requested over http
  - Also called “manifest”
- Describes all segments
- Timing information and byte ranges of chunks
- Client uses HTTP GET RANGE from a given AS + representation to ask a given bit rate
- Client could use a different URL for each AS + representation



# Video Delivery using CDN





# Server Selection

- File → server mapping done in at least three ways
- Dynamic DNS resolution
  - DNS returns different IP addresses for a given DNS name
- HTTP redirect
  - Use HTTP status code 3xx [with new URL]
  - Web browser does a GET from the new site
- IP anycast
  - Use BGP to announce the same IP address from different locations
  - Client reaches “nearest” location according to inter-domain routing

# DASH Summary

- Widely used in video streaming services
- Allows independent requests per segment
  - Hence, independent segment quality and data sizes
  - Encoded through separate HTTP objects and corresponding HTTP byte ranges
  - Combined or separate audio & video streams
- Works well with CDNs
  - Independent representations or chunks can be queried from different locations if needed