

CS 352


Link Layer: Introduction

CS 352, Lecture 22.1

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

Link layer




The main function of the link layer is **link-local delivery**: getting packets from one side of the link to the other.

Link layer: introduction

Terminology:

- Endpoints and routers: **nodes**
- communication channels that connect adjacent nodes along communication path: **links**
 - wired links, wireless links, local area networks (LANs)
- layer-2 packet: **frame**, encapsulates (IP) datagram

Link layer has the responsibility of transferring datagram from one node to **physically adjacent** node over a link.



Slides in this presentation were heavily adapted from those of Kurose and Ross, who own the copyright on the material.

Link layer: context

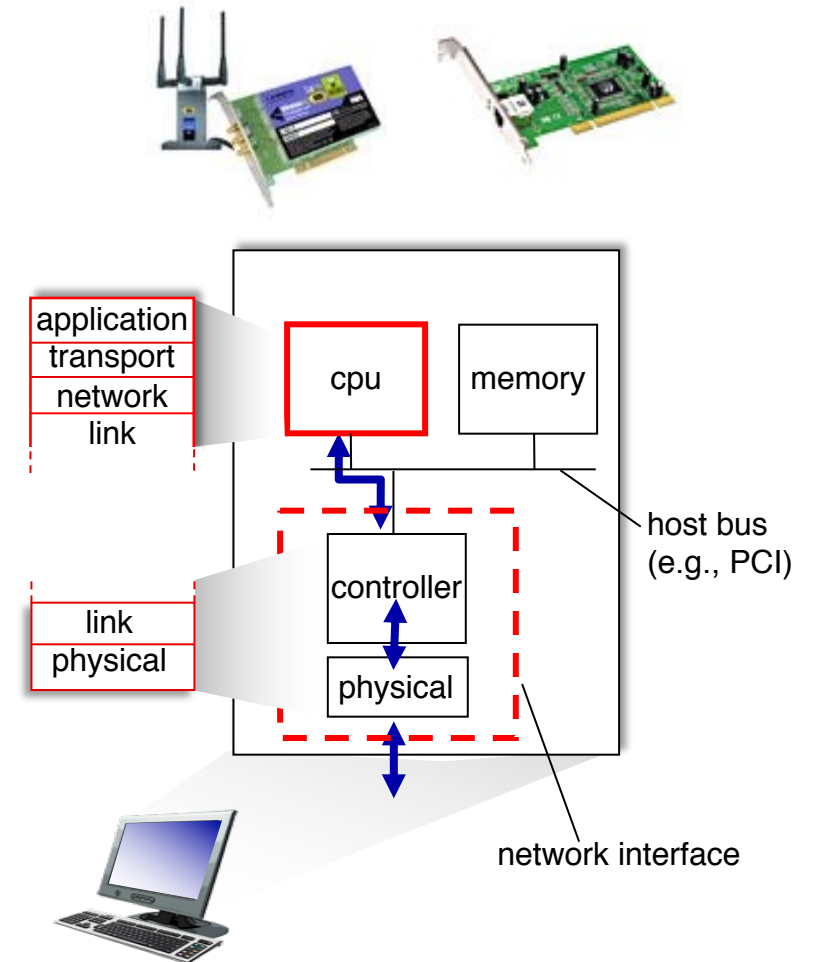
- Datagram transferred by different link layer protocols over different links
 - e.g., Ethernet on first link, frame relay on intermediate links, 802.11 on last link
- Each link layer protocol may provide different services
 - e.g., some links may provide additional reliability mechanisms

Analogy:

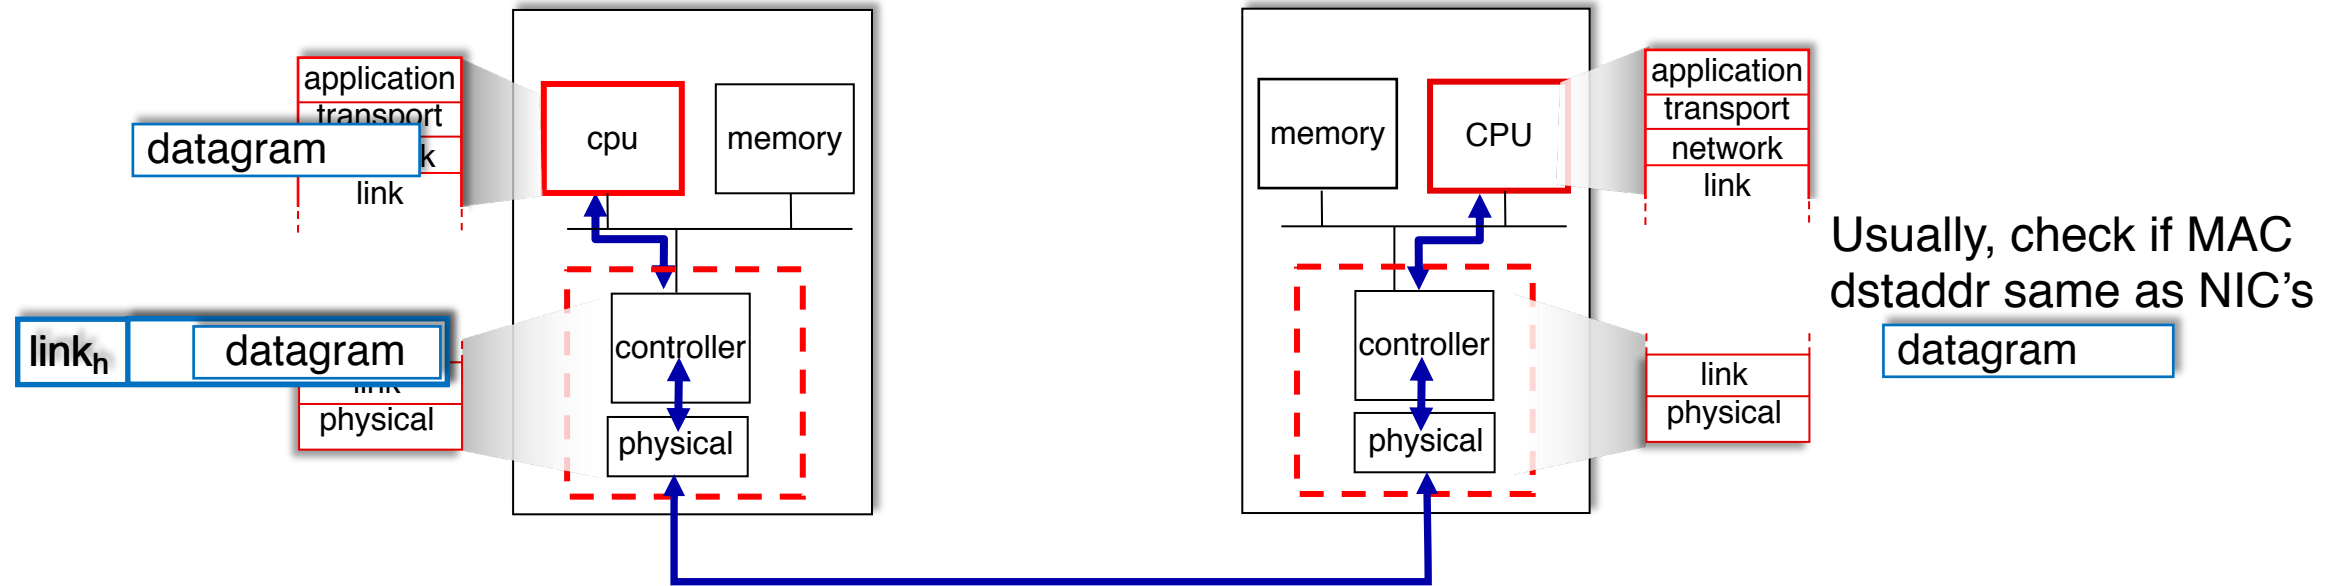
- trip from Piscataway, NJ to Palo Alto, CA
 - limo: Piscataway to JFK
 - plane: JFK to San Francisco
 - train: San Francisco to Palo Alto
- tourist = **datagram**
- transport segment (road/flight/rail) = **communication link**
- transportation mode (car/plane/train) = **link layer protocol**

Where is the link layer implemented?

- in every endpoint & router
- link layer implemented in **network interface card (NIC)**
 - Ethernet, WiFi card or chip
 - Router input and output ports
- At endpoint, attaches into its system buses (e.g., PCIe)
- combination of hardware, software, firmware



Interfaces communicating



sending side:

- encapsulates datagram in frame
- adds error checking bits, reliable data transfer, flow control, etc.

receiving side:

- looks for errors, reliable data transfer, flow control, etc.
- extracts datagram, passes to upper layer at receiving side

Link layer (MAC) addresses


Link layer addresses

- Review: we've looked at 32-bit IPv4 addresses
 - network-layer address for interface
 - used for layer 3 (network layer) forwarding
 - e.g.: 128.119.40.136
- **MAC or physical or link-layer address**
 - Used “locally” to get frame from one interface to another physically-connected interface (same subnet, in IP-addressing sense)
 - This course: **Ethernet family** of link layer protocols
 - **48-bit** Ethernet MAC address burned in NIC ROM
 - Sometimes, the address can be set in software
 - e.g.: 1A-2F-BB-76-09-AD ————— *hexadecimal (base 16) notation
(each “numeral” represents 4 bits)*

MAC addresses

each interface on LAN

- has unique 48-bit MAC address
- has a locally unique 32-bit IP address (as we've seen)



MAC addresses

- MAC address allocation administered by IEEE
- manufacturer buys portion of MAC address space (to assure uniqueness)
- Analogy:
 - MAC address: like Social Security Number
 - IP address: like postal address
- MAC is a **flat address**: portability
 - e.g., can move interface from one Ethernet LAN to another
 - Recall: IP address is not portable: depends on IP network to which node is attached

Link layer services

Link layer: services

- **Line termination (serialize/deserialize)**
 - Physical signaling (into the wire) & extracting digital signal (out of the wire)
 - Encapsulate datagram into frame (**framing**). Add header, trailer
- **Error detection:**
 - errors caused by signal attenuation, noise.
 - receiver detects errors, signals retransmission, or drops frame
- **Error correction:**
 - receiver identifies **and corrects** bit error(s) without retransmission
- **Medium access control**
 - channel **medium access** in a shared medium:
 - “Who should talk?”

Link layer: services

- Line termination (serialize/deserialize)
 - Physical signaling (into the wire) & extracting digital signal (out of the wire)
 - Encapsulate datagram into frame (framing). Add header, trailer
- Error detection:
 - errors caused by signal attenuation, noise
 - receiver detects errors, signals retransmission, or drops frame
- Error correction:
 - receiver identifies and corrects bit error(s) without retransmission

This lecture

- Medium access control
 - channel medium access in a shared medium
 - “Who should talk?”

Next 2 lectures

Link layer: services: there's more!

- **Reliable delivery between adjacent nodes**
 - Seldom used on **low bit-error** links
 - Not strictly needed for functionality
 - **An optimization that significantly improves performance** over pure end-to-end reliable delivery over high-error-rate links (e.g., wireless)
- **Flow control:**
 - pacing between adjacent sending and receiving nodes
 - Used in **lossless** link layer technologies (e.g., Infiniband, lossless Ethernet)

Not covered in this course

Link layer: services

Next module:

Encoding, error detection, and error correction

- Line termination (serialize/deserialize)
 - Physical signaling (into the wire) & extracting digital signal (out of the wire)
 - Encapsulate datagram into frame (framing), Add header, trailer
- Error detection:
 - errors caused by signal attenuation, noise
 - receiver detects errors, signals retransmission, or drops frame
- Error correction
 - receiver identifies and corrects bit error(s) without retransmission
- Medium access control
 - channel medium access in a shared medium:
 - “Who should talk?”

CS 352

Encoding, Error Detection, and Correction

CS 352, Lecture 22.2

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

Recall: Link layer services

- **Line termination (serialize/deserialize)**
 - Physical signaling (into the wire) & extracting digital signal (out of the wire)
 - Encapsulate datagram into frame (**framing**). Add header, trailer
- **Error detection:**
 - errors caused by signal attenuation, noise.
 - receiver detects errors, signals retransmission, or drops frame
- **Error correction:**
 - receiver identifies **and corrects** bit error(s) without retransmission
- **Medium access control**
 - channel **medium access** in a shared medium:
 - “Who should talk?”


Link layer: services

- Line termination (serialize/deserialize)
 - Physical signaling (into the wire) & extracting digital signal (out of the wire)
 - Encapsulate datagram into frame (adding) header, trailer
- **Encoding, error detection, and error correction**
 - **Error detection:**
 - errors caused by signal attenuation, noise
 - receiver detects errors, signals retransmission, or drops frame
 - **Error correction:**
 - receiver identifies and corrects bit error(s) without retransmission
- Medium access control
 - channel medium access in a shared medium:
 - “Who should talk?”

Handling digital and physical Information



Encoding and Decoding

- Signals propagate over a physical medium
 - modulate electromagnetic waves
 - e.g., vary voltage
- Encode binary data onto signals
 - e.g., 0 as low signal and 1 as high signal
 - known as **Non-Return to zero** (NRZ)




Clock sync and recovery

- Receiver needs to know when a symbol begins and ends
- One approach: send a clock signal together with data signal
 - Lowers data rate by 2x!
- Another approach: look for **transitions** in the data signal to re-synchronize the clock
- **Long strings of 0s and 1s** make synchronization challenging
- It's like trying to dance in sync without a beat



How many 1s?

Self-clocking encoding



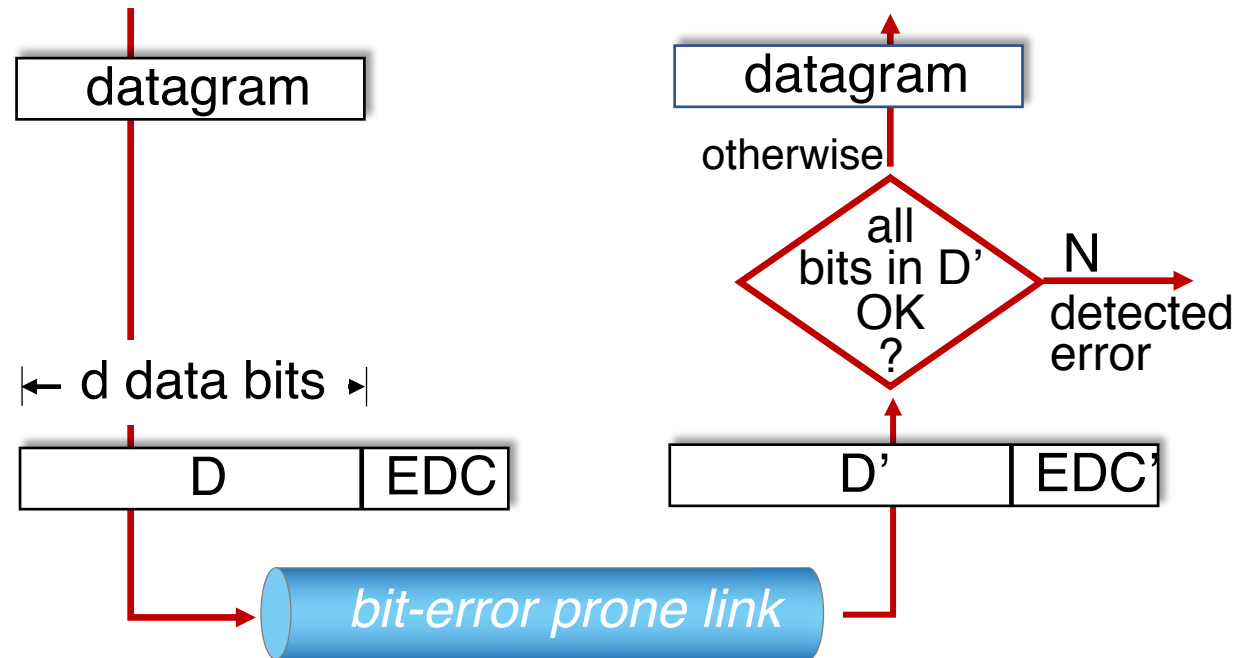
- **Manchester encoding**
- 0 encoded by a positive transition, 1 by a negative transition
- Construct using $\text{XOR}(\text{bit}, \text{clock})$
- Used in early Ethernet standards (up to 10 Mbit/s)

Error Detection & Correction

Error detection

EDC: error detection and correction bits (e.g., redundancy)

D: data protected by error checking, may include header fields




Error detection not 100% reliable!

- protocol may miss some errors, but rarely
- larger EDC field yields better detection and correction

Parity checking

single bit parity:


- detect single bit errors



Even parity: set parity bit so there is an even number of 1's

two-dimensional bit parity:

- detect **and correct** single bit errors



no errors:

1	0	1	0	1	1
1	1	1	1	0	0
0	1	1	1	0	1
1	0	1	0	1	0

detected and correctable single-bit error:

1	0	1	0	1	1
1	0	1	1	0	0
0	1	1	1	0	1
1	0	1	0	1	0

parity error

Review: Transport (UDP/TCP) checksum

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

sender:

- treat contents of transport segment as sequence of 16-bit integers
- **checksum:** addition (roughly) of segment content
- checksum value put into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - Yes: *assume* no error
 - No: declare error

Cyclic Redundancy Check (CRC)

- more powerful error-detection coding than checksums
- D: data bits (given, think of these as a binary number)
- G: bit pattern (**generator**), of $r+1$ bits (given)



Goal: choose r CRC bits, R , such that $\langle D, R \rangle$ exactly divisible by G

Do all arithmetic mod 2: all additions, subtractions replaced by XOR

- receiver knows G , divides $\langle D, R \rangle$ by G . If non-zero remainder: error detected!
- can detect all **burst errors (continuous bit errors)** less than $r+1$ bits
- widely used in practice (Ethernet, 802.11 WiFi)

Cyclic Redundancy Check (CRC): example

We want:

$$D \cdot 2^r \text{ XOR } R = nG$$

or equivalently:

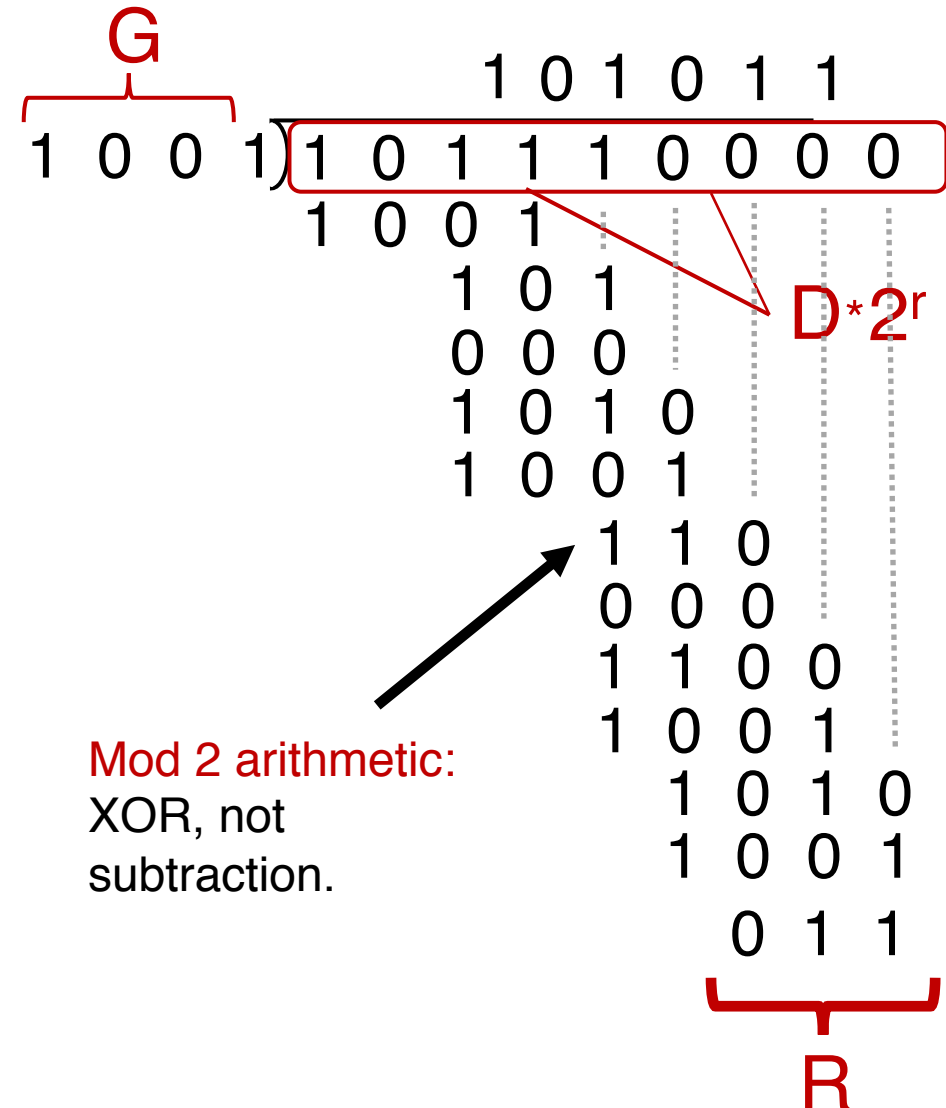
if we divide (**mod 2**) $D \cdot 2^r$ by G , want remainder R to satisfy:

$$R = \text{remainder} \left[\frac{D \cdot 2^r}{G} \right]$$

Perform long division to compute the remainder.
E.g., $D = 101110$, $G = 1001$

How CRC is computed in software:

<https://www.kernel.org/doc/html/latest/staging/crc32.html>



Summary

- Self-clocking encoding useful to synchronize sender & receiver
- Error detection and correction mechanisms:
 - Parity bits: single or a few bits of error
 - CRCs: bursty errors up to a certain size
- Error detection and correction codes widely used across many computer systems

CS 352

Connecting Multiple Endpoints into a Single Network

CS 352, Lecture 22.3

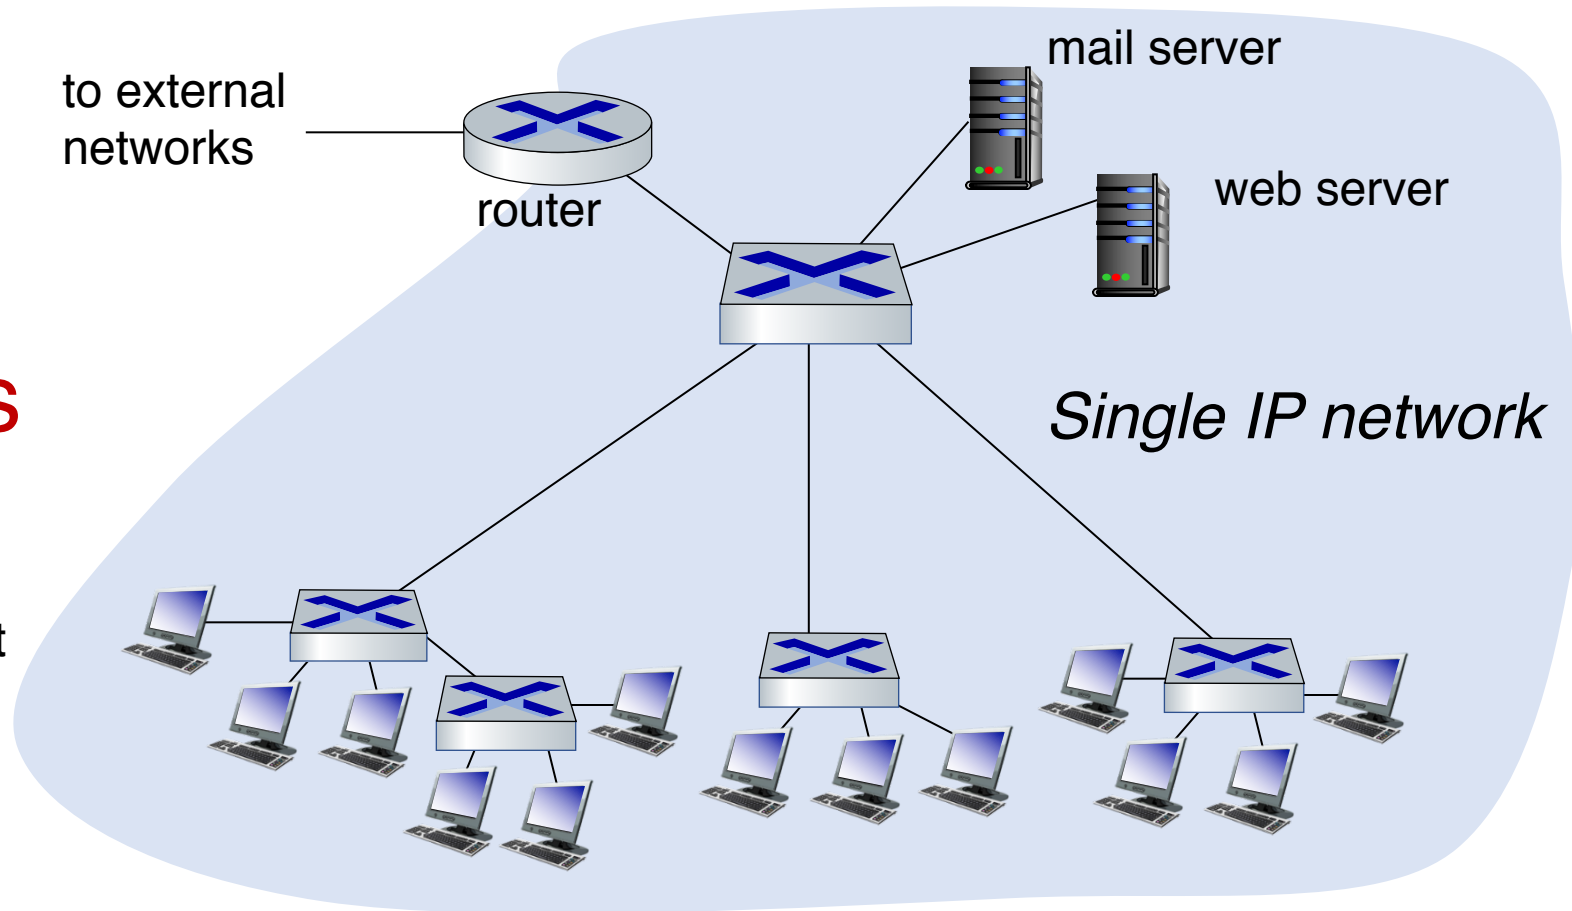
<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

A small organizational network today

Switches

Located in switching closets, connect directly to endpoints.




Switching

- A **switch** is a link-layer device
 - Examine incoming frame's destination MAC address
 - Selectively forward frame to one-or-more outgoing links when frame is to be forwarded
 - Can store link layer frames in switch buffers
- **Transparent**: hosts unaware of presence of switches
- Plug-and-play, **self-learning**
 - Switches do not need to be configured


Switch: multiple simultaneous transmissions

- hosts have dedicated, direct connection to switch
- switches **buffer** packets
- Same link layer protocol used on each incoming link
 - full duplex links
 - No medium access control needed (more next lecture)
- switching: A-to-A' and B-to-B' can transmit **simultaneously**



Switch: multiple simultaneous transmissions

- hosts have dedicated, direct connection to switch
- switches **buffer** packets
- Same link layer protocol used on each incoming link
 - full duplex links
 - No medium access control needed (more next lecture)
- switching: A-to-A' and B-to-B' can transmit **simultaneously**
- However, A → A' and C → A' can't happen simultaneously



Switched LAN

- If switches don't need to be configured, how can switches route to the correct endpoints?
- Process known as **MAC learning** or **layer-2 bridging**
 - a configuration-free, **learning-based** routing protocol


Switch forwarding table

Ex: how does switch know A' reachable via interface 4, B' reachable via interface 5?

Each switch has a MAC table.

Each entry:


- (MAC address of host, interface to reach host, timestamp)
- **looks like a forwarding table!**



How are entries created and maintained in the MAC table?

MAC learning

- switch **learns** which hosts can be reached through which interfaces
 - when frame received, switch “learns” location of sender: incoming LAN segment
 - records sender/location pair in switch table



MAC addr	interface	TTL
A	1	60

*Switch table
(initially empty)*

MAC learning: frame forwarding

when a frame received at switch:

1. record incoming link, **source MAC** address
2. index switch table using **destination MAC** address

3. if **entry found** for destination
then {

- if destination on link from which frame arrived
then **drop frame**

- else **forward frame** on interface indicated by entry

- }

- else **flood** /* forward on all ports except arriving interface */

MAC	port	TTL
A	1	60

Flooding is only acceptable because all endpoints are in the same org/IP network.

Forwarding: Example


- frame destination A'.
Interface unknown
(not in table)

flood

- destination A location
known:

Selectively send on one link

- Subsequent $A \leftrightarrow A'$
packets not flooded




MAC addr	interface	TTL
A	1	60
A'	4	60

*switch table
(initially empty)*

Interconnecting switches

- MAC learning switches can be connected together
 - The algorithm works the same way!
- Complication: what if there are **loops** in the switch topology?
 - Flooding may result in the same packet arriving from multiple ports
- Ethernet: **spanning tree protocol**
 - Switches discover the switch-level graph
 - Process akin to link state advertisements flooding
 - Then, switches use a **loop-free** subset of links
 - A spanning tree of the network graph



Switches vs. routers

Both can store, buffer, and forward.

- **routers**: network-layer devices (examine network-layer headers)
- **switches**: link-layer devices (examine link-layer headers)

Both have forwarding tables.

- **routers**: compute forwarding tables using routing algorithms, link configurations, and announced IP addresses
- **switches**: learn forwarding table using flooding and learning MAC addresses

Summary

- Enterprises often use switches for their ease of configuration and plug-and-play nature
 - Switched Ethernet: popular in dorms and office buildings
- Switches can discover endpoints
- **Flooding** facilitates reachability across endpoints. Only possible as all endpoints part of the same IP network
- **MAC learning** records where endpoints send from, enabling the discovery of endpoint-port associations without prior knowledge