

CS 352



Internet Protocol (IP)

CS 352, Lecture 16.1

<http://www.cs.rutgers.edu/~sn624/352>

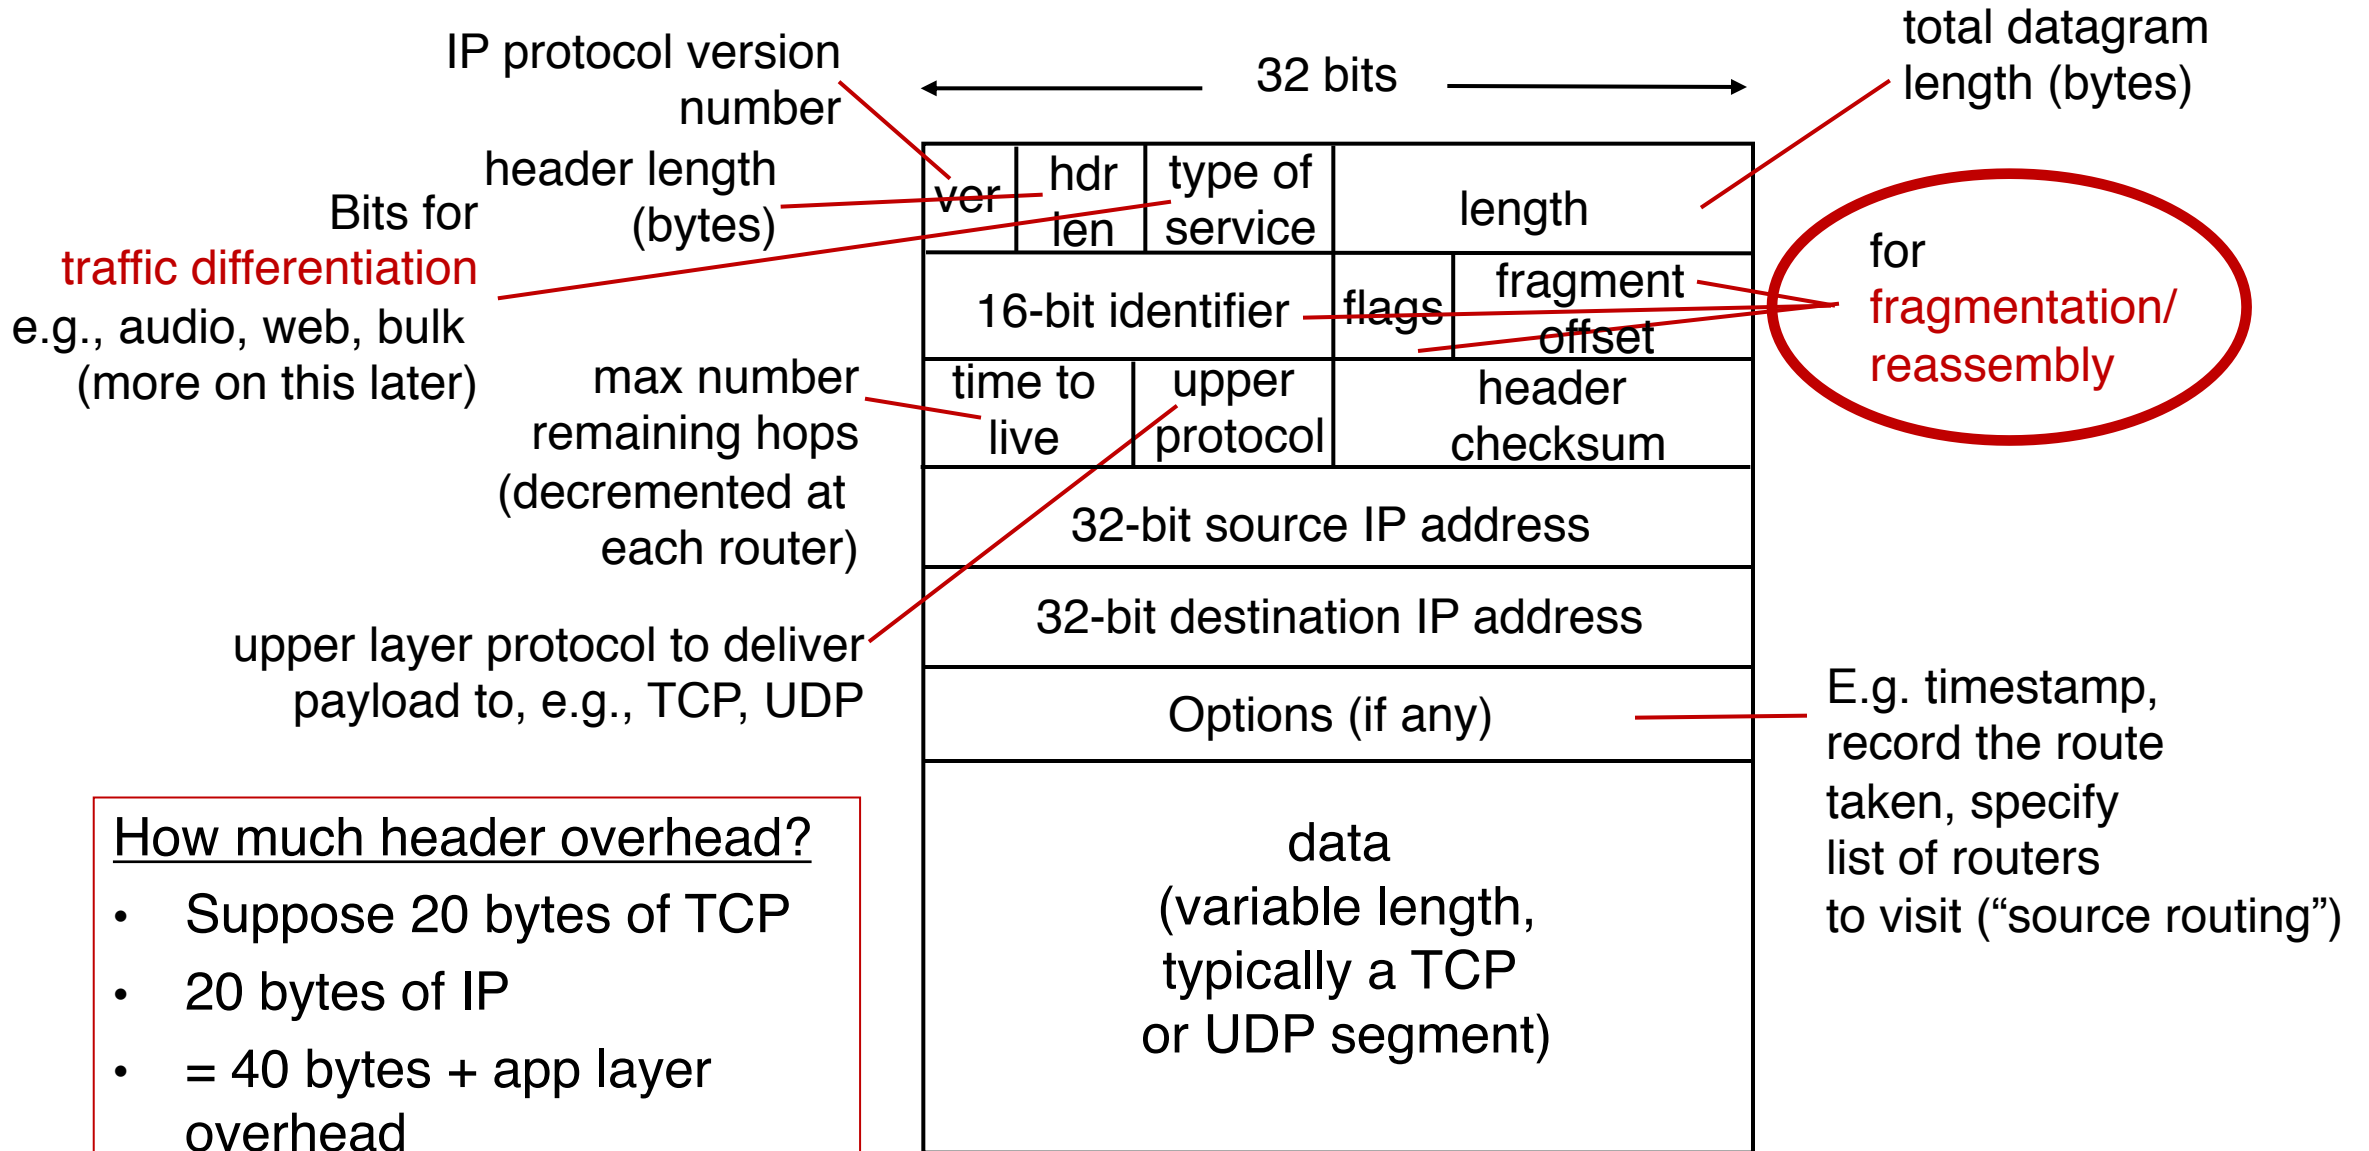
Srinivas Narayana

Network



The main function of the network layer is to
move packets from one endpoint to another.

IPv4 Datagram Format




How much header overhead?

- Suppose 20 bytes of TCP
- 20 bytes of IP
- = 40 bytes + app layer overhead

IP fragmentation and reassembly

- Links and transmission media have **MTUs (maximum transmission unit)**:
 - Largest possible link-level frame
- On a network path, a packet might traverse links with different MTUs
- This may result in a large IP datagram to be divided (**fragmented**) by a router
 - Fragments reassembled only at the destination endpoint, at the IP layer
 - IP header bits used to identify and **reassemble** related fragments



IP fragmentation and reassembly

- Suppose a large 4000-byte datagram reaches a router. The next link has MTU 1500 bytes.
 - Note: MTU includes IP headers, so does the length field of the IP header. IP payload = 3980 bytes.
- Result: 3 datagrams of length 1500, 1500, 1040 bytes resp.
 - IP payload = 1480, 1480, 1020 bytes resp. (adds to 3980)
- Offset field = index of payload byte / 8

	length =4000	ID =x	fragflag =0	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =185	
--	-----------------	----------	----------------	----------------	--

1480/8

	length =1040	ID =x	fragflag =0	offset =370	
--	-----------------	----------	----------------	----------------	--

2960/8

IP fragmentation and reassembly

- At the destination endpoints, the fragments are reassembled using the IP **identifier** field
 - Fragments of the same original datagram share the same IP ID
- The fragmentation flag is set to 0 for the terminal fragment, and 1, if other fragments follow
- The offset field allows the IP stack to reassemble the fragments in order into a single IP datagram

	length =4000	ID =x	fragflag =0	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =185	
--	-----------------	----------	----------------	----------------	--

	length =1040	ID =x	fragflag =0	offset =370	
--	-----------------	----------	----------------	----------------	--

1480/8

2960/8

The rest of this lecture and the next

- We'll talk about some **support protocols** and mechanisms for the network layer
 - Protocols: DHCP, ICMP, ARP
 - Mechanisms: NAT
 - We'll also talk about IP version 6 (IPv6)
- Some of these protocols use an IP header underneath their own header (ICMP) or replace the IP header with their own (ARP)
 - But these shouldn't be construed as transport/network protocols
 - They are fundamental to supporting IP/network layer functionality
 - More appropriately discussed as support protocols for the network layer

CS 352

Dynamic Host Configuration

CS 352, Lecture 16.2

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

How does an endpoint get its IP addr?

- One possibility: hard-code the IP address on the endpoint
 - e.g., a system admin writing addresses in a file
 - UNIX: /etc/network/interfaces
 - Windows: control panel → network → configuration → TCP/IP → properties
- Another possibility: dynamically receive an address “from the network”
 - **DHCP**: **D**ynamic **H**ost **C**onfiguration **P**rotocol
 - Provide plug-and-play functionality for endpoints (e.g., phones, laptops)

Many similar bootstrapping problems

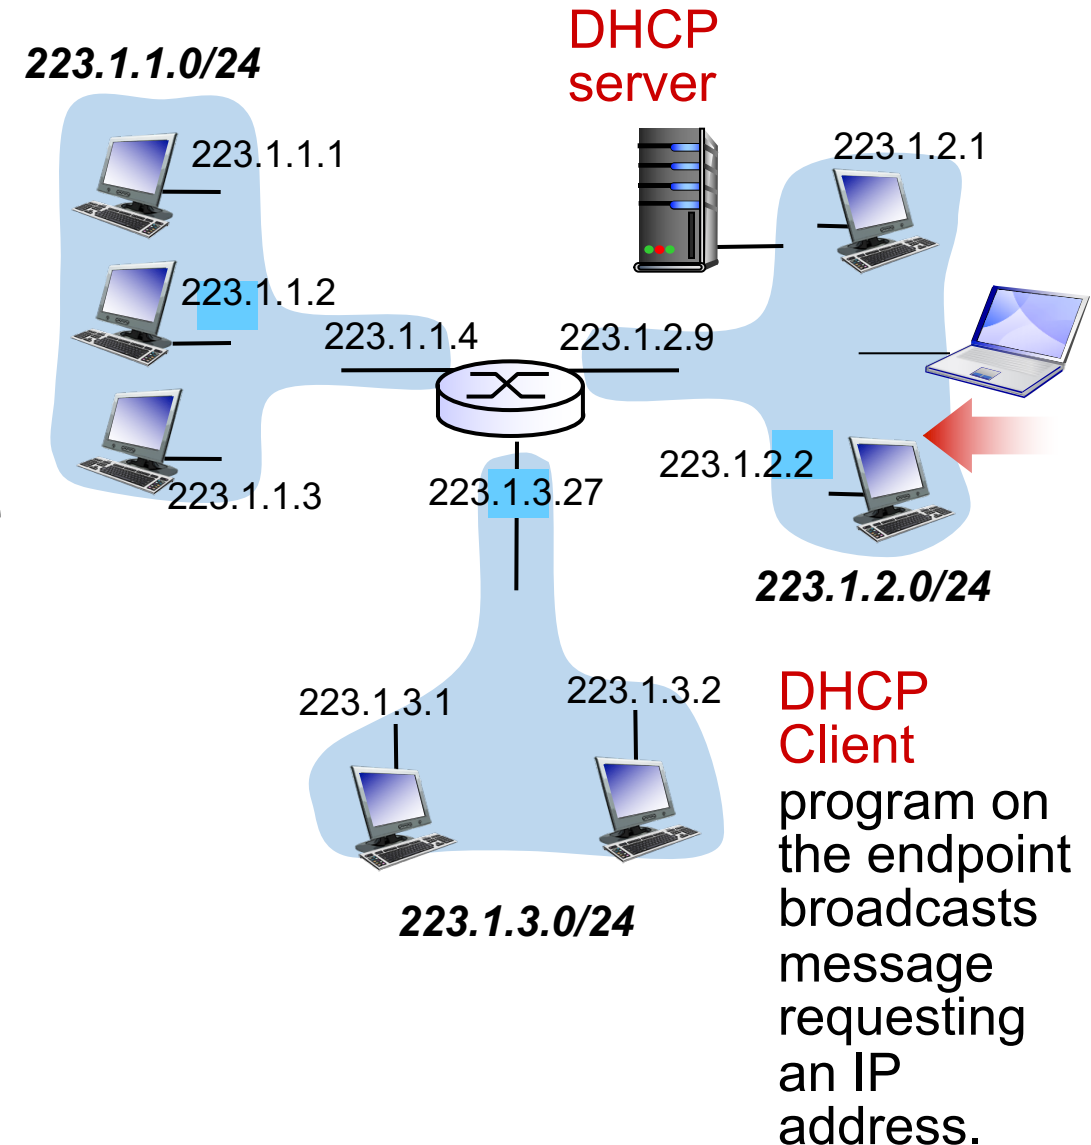
- How does a host get its IP address?
- How does a host know its local DNS server?
- How does a host know its netmask?
 - i.e., so that it can know which other hosts are in the same network
- How does a host know how to reach other networks?
 - i.e., which router is at the “border” of the current network?
 - This router is also called the **gateway router**: crucial for an endpoint to communicate with another endpoint external to the network

How DHCP works

- A new endpoint that just joins a network knows nothing about the network
 - It doesn't even have a network address for its point of attachment
- It makes no sense to have it contact a “known” server to receive this information.
- The only known mechanism that might work is **broadcast**:
 - Ask everyone!

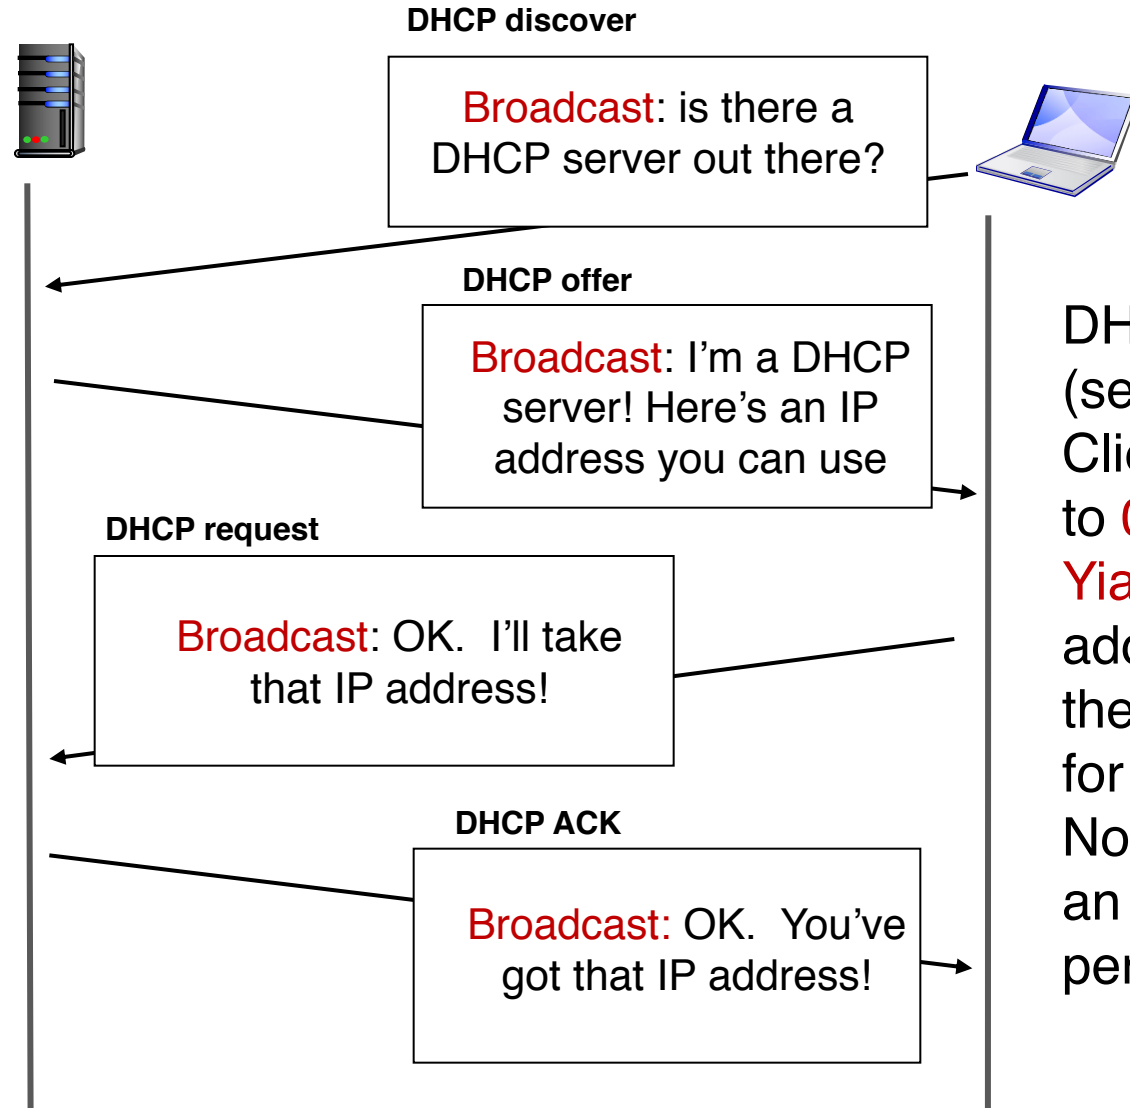
How DHCP works

- DHCP allows a host to dynamically obtain its IP address from a **server** on a network when it joins the network
- DHCP can allow a host to be mobile across different networks, obtaining IP addresses as needed
- DHCP uses **leases** on addresses
 - Host must renew lease periodically
 - Allows network to reuse an IP with an expired lease, reclaiming addresses from inactive hosts



DHCP client-server scenario

DHCP server:
223.1.2.5



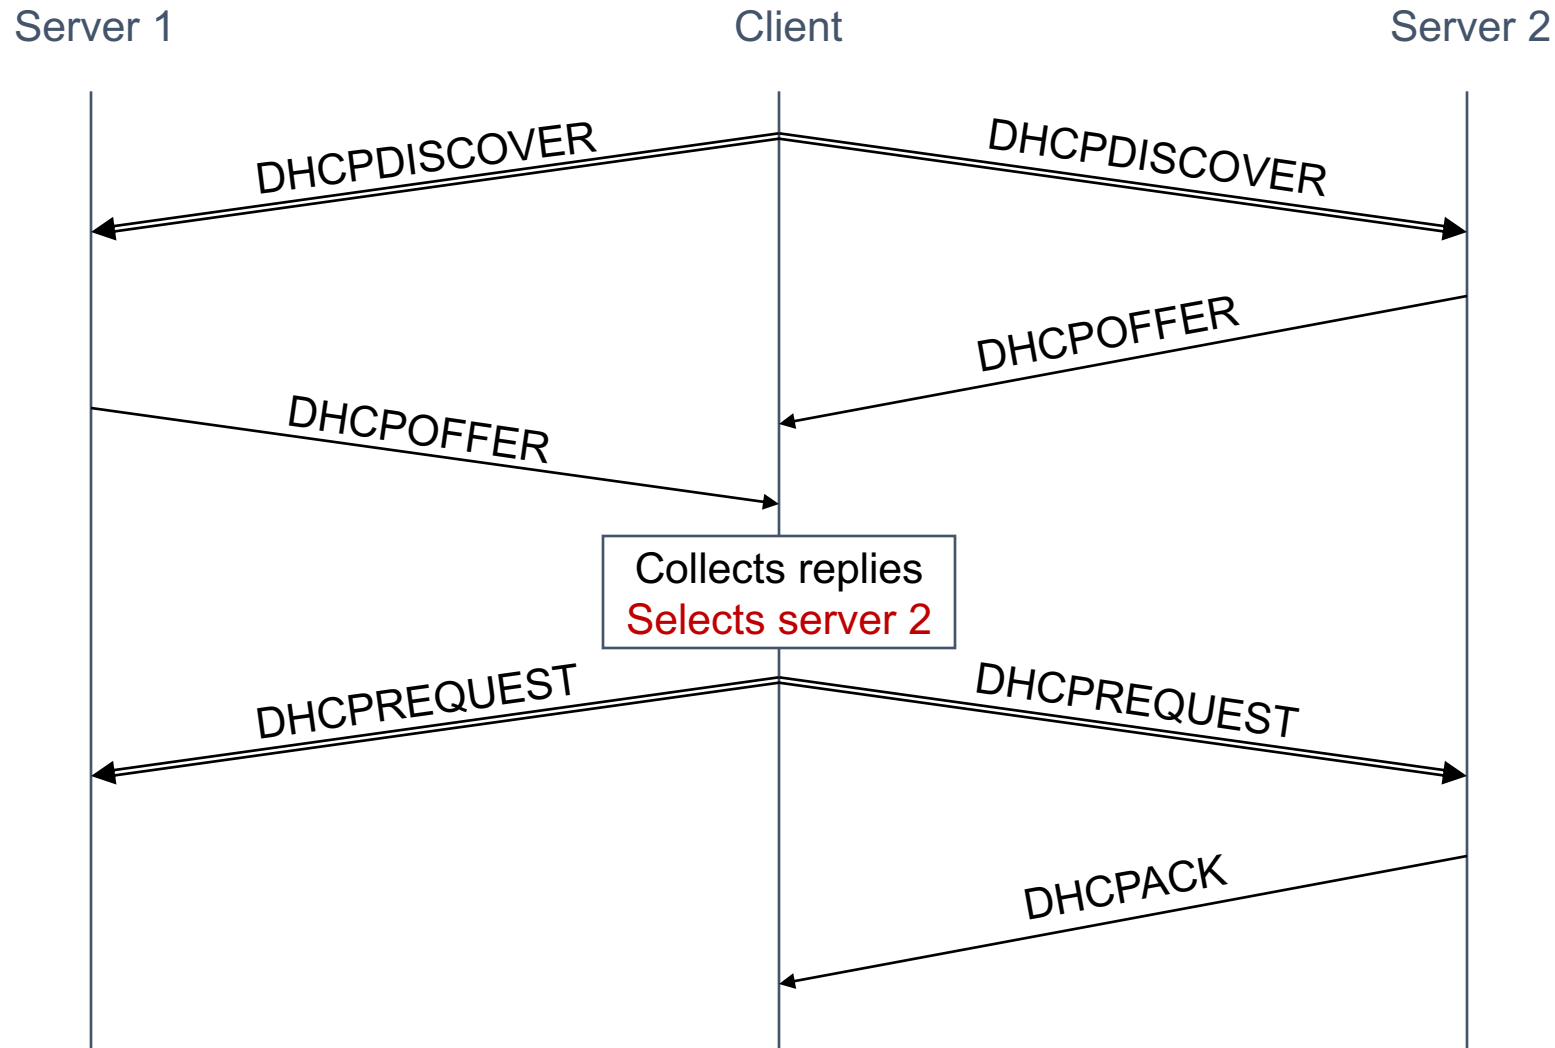
223.1.2.4
Arriving
client

DHCP runs on UDP ports 67 (server) and 68 (client)
Client's initial IP address is set to 0.0.0.0

Yiaddr stands for "your IP address" – an address value the server sends to the client for consideration

Note that the IP allocation has an associated **lifetime** (lease period)

Multiple DHCP servers can coexist



DHCP returns more than an IP address

- Name and IP address of the **local DNS server**
- **Netmask** of the IP network the host is on
 - Useful to know whether another endpoint is inside or outside the current IP network
- Address of the **gateway router** to enable the endpoint to reach other IP networks

Your home router runs DHCP

- Likely, your home devices (laptops, tablets, phones) are all using DHCP-assigned IP addresses
- The DHCP server is running on the control processor of your home's access router (e.g., WiFi router)
- You can access the DHCP client program on Linux using the command `dhclient` and on Linux using `sudo ipconfig <interface> DHCP`

Summary of DHCP

- Want endpoints to have plug and play functionality
 - Avoid tedious manual configuration of IP addresses and other information
- DHCP: a general bootstrapping mechanism for critical information required for network layer functionality
- Hosts can be simple: receive information from DHCP servers by **broadcasting** over the network

CS 352

Internet Control Message Protocol

CS 352, Lecture 16.3

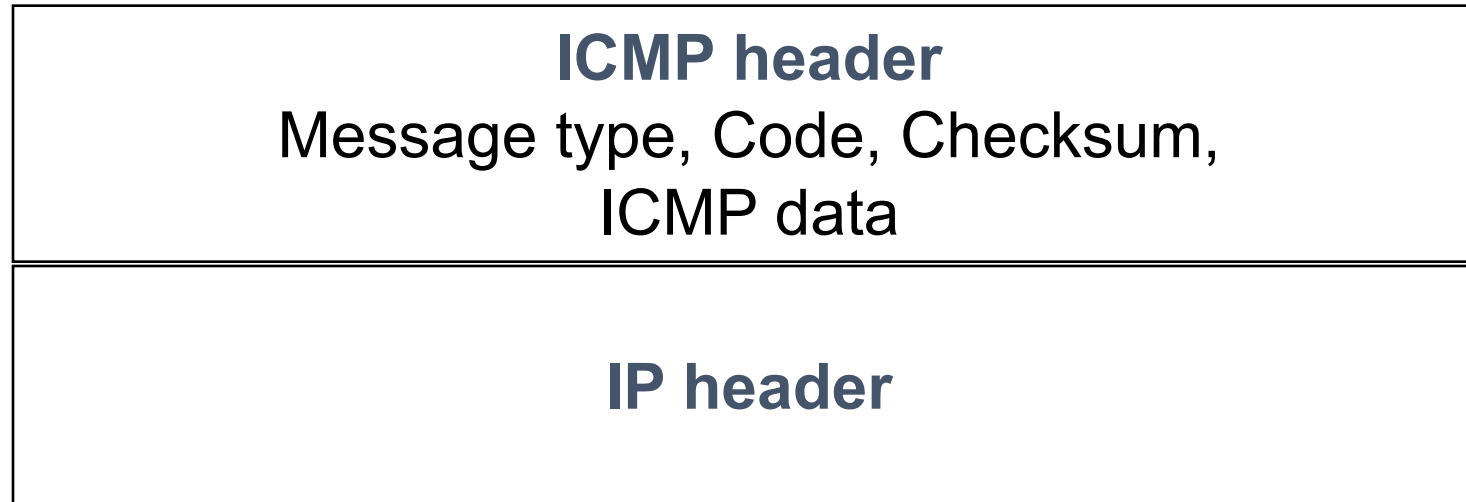
<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

Internet Control Message Protocol

- A protocol for **troubleshooting** and diagnostics
- Works over IP: **unreliable delivery** of packets
- Some functions of ICMP:
 - Determine reachability and network errors
 - Specify that packets have been in the network for too long

ICMP message



https://en.wikipedia.org/wiki/Internet_Control_Message_Protocol#Control_messages


Specific uses of ICMP

- Echo request reply
 - Check remotely if an endpoint is alive and connected
- An unreachable destination
 - Invalid address and/or port
- Knowing if packet's IP time-to-live expired
 - Example, due to routing loops
- Look at two tools built using ICMP: **ping** and **traceroute**

Ping

- Uses ICMP echo request (type=8, code=0) and reply (type=0, code=0)
- Source sends ICMP **echo request** message to dst address
- Destination replies with an ICMP **echo reply** message containing the data in the original echo request message
- Source can calculate round trip time (RTT) of packets
- If no echo reply comes back, then the destination is **unreachable**
- Don't need to have a server program running on the other side
 - In general, the remote endpoint can be completely outside your control

Ping



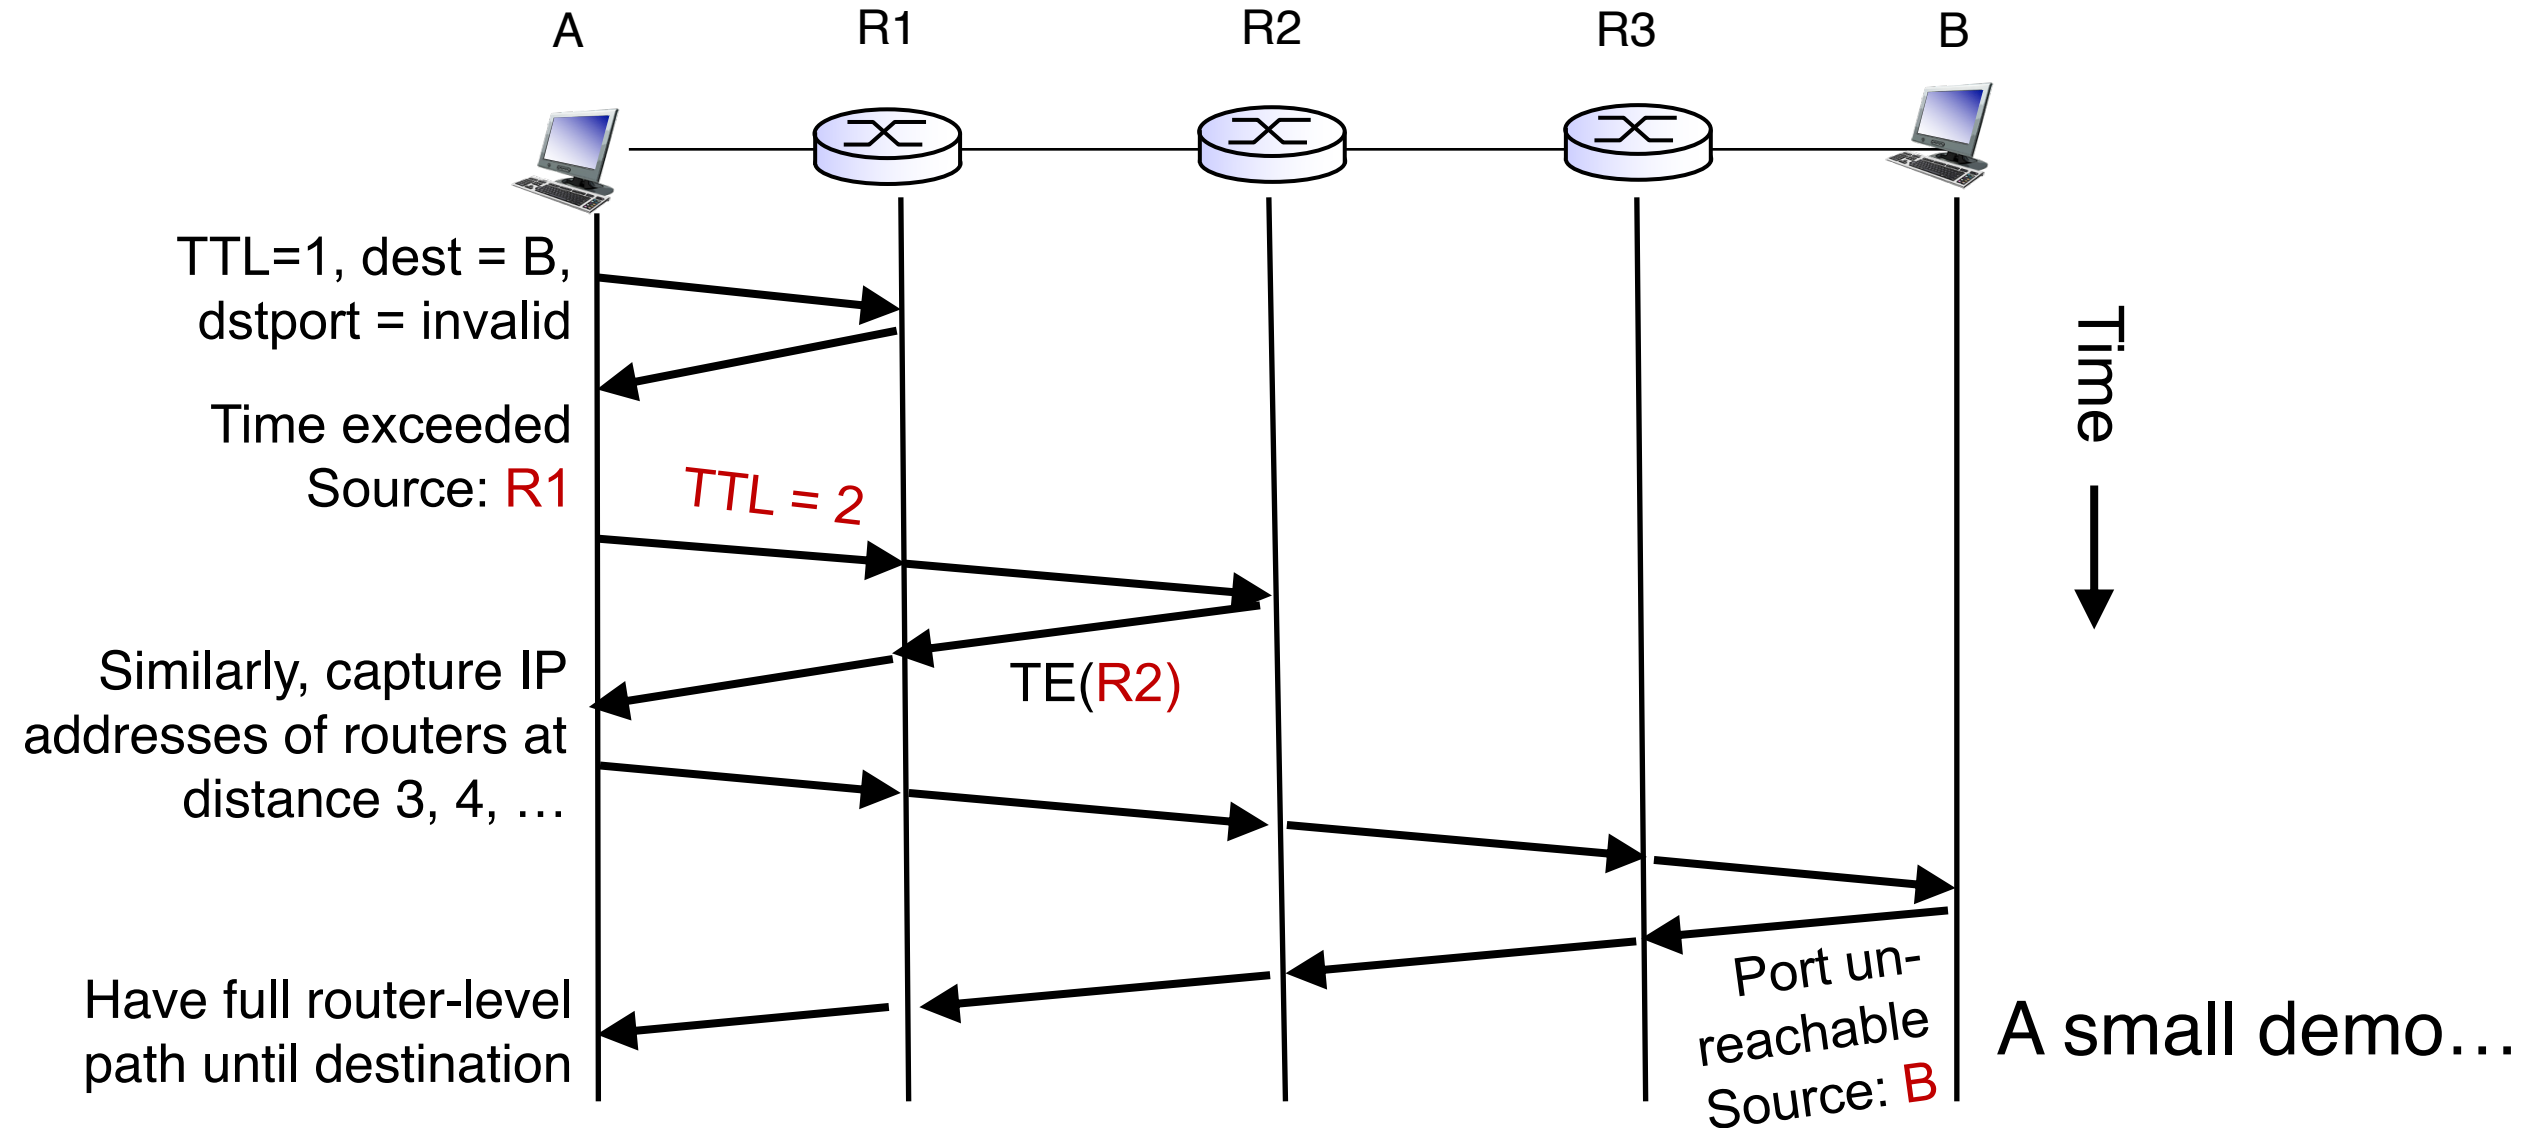
Traceroute

- A tool that can record the router-level path taken by packets
- A clever use of the IP **time-to-live** (TTL) field
- In general, when a router receives an IP packet, it decrements the TTL field on the packet
 - A failsafe mechanism to ensure packets don't keep taking up network resources for too long
- If a router receives a packet with TTL=0, it sends an **ICMP time exceeded** message to the source endpoint

Traceroute

- Traceroute sends multiple packets to a destination endpoint
- But it progressively increases the TTL on those packets: 1, 2, ...
- Every time a time exceeded message is received, record the router's IP address
- Process repeated until the destination endpoint is reached
- If the packet reaches the destination endpoint (i.e.: TTL is high enough), then the endpoint sends a port unreachable message

Traceroute



Summary of ICMP

- A protocol for network diagnostics and troubleshooting
- Two useful tools: **ping** and **traceroute**
- Ping: test connectivity to a machine totally outside your control
 - Use ICMP echo request and reply
- Traceroute: determine router-level path to a remote endpoint
 - A smart use of the TTL field in the IP header

