

# CS 352



# Network Service Guarantees

CS 352, Lecture 20.1

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Network




The network layer moves packets from one place to another.  
**The network will put in its best effort but makes no guarantees.**

# Network support for applications

- A **best effort** Internet architecture does not offer any guarantees on delay, bandwidth, and loss
  - Network may drop, reorder, corrupt packets
  - Network may treat traffic randomly regardless of their “importance”
- However, many apps require special treatment & guarantees
  - E.g., voice over IP (phone calls) require strict delay guarantees
  - E.g., HD video requires a reasonable minimum bandwidth
  - E.g., remote surgery with 3D-vision requires strict sync & latency
- **Q: How to provide quality of service (QoS) for apps?**

# Why best effort isn't enough: Contention



- Resource contention occurs in the **core** of the network
- Congestion control will react, but may be too little & too late:
  - Congestion control can't prevent packet drops “now”
  - Congestion control won't prevent high-sending-rate flows from inflicting large delays or recurring drops

Can networks help improve the quality of service for applications?

Yes, but networks must become better than best-effort.

# Approach 1: Provision more capacity

- If you're an ISP (e.g., AT&T), you might **deploy enough capacity** so that contention doesn't occur any more
  - Low complexity: can use current "best effort" network
- However, this approach incurs **high costs (e.g., bandwidth)**
- A key challenge: estimating how much bandwidth is enough
  - Need to estimate demand over time
  - Network operators can do this quite well usually
  - But there are exceptional circumstances: pandemics, Superbowl, etc.

# Approach 2: Classes of service


- Have the network treat different traffic differently
  - Also called **traffic differentiation**
- Analogy: lines at an airport (e.g., first class vs. economy)
- Partition traffic into classes and offer service guarantees **per class** and **across classes**
  - Classes may be indicated using the IP type of service header bits
  - Classes may be inferred from IP & transport headers (e.g., src/dst/ports)
- **Packet classification:** assigning packets to classes
  - (Not in scope: we won't discuss packet classification)

# Kinds of Service Guarantees




# (1) Strict prioritization

Transmitted immediately regardless of HTTP packets




- Suppose a 1Mbps interactive flow and an HTTP connection share a 1.5 Mbps link.
- A network operator (e.g., Rutgers admin) might choose to **prioritize** the interactive app strictly over the HTTP flow.

## (2) Rate limiting



- What if a flow doesn't respect its allocation?
  - Example: Say, conf call flow goes beyond 1 Mbit/s
  - Don't want to starve HTTP flow!
- An operator might want to limit a flow to a certain max rate
- **Isolation:** HTTP should not be impacted by the conf call


### (3) Weighted fair sharing



- An operator might want to partition the link's rate  $C$  into separate allocations for each class
  - Partitions may have **weights**  $w$  (example: 2, 1)
- Usually, class  $i$  gets the illusion of traversing a logical link of rate
$$w_i * C / \sum_j w_j$$


### (3) Weighted fair sharing

- Customary to think of different classes as belonging to different **queues**
- For this reason, weighted fair sharing is also called **weighted fair queueing (WFQ)**
- Each queue is first-in-first-out (FIFO)
- The link multiplexes among these queues
- Intuitively, packets of one queue should not influence the behavior of other queues
- Hence, fair queueing is also a form of **isolation** across traffic classes



### (3) Weighted fair sharing

- But what if one class doesn't use its share?
  - Can other classes use the spare capacity?
- Yes! WFQ is **work-conserving**: a router implementing WFQ will allow other classes to use the unused capacity
- Work conservation makes WFQ different from rate limits applied separately to each class
  - Class  $i$ 's usage can exceed  $w_i * C / \sum_j w_j$
  - (only if spare capacity is available, of course.)



# Q: Where are guarantees enforced?

- We've seen three kinds of service guarantees: prioritization, rate limiting, and fair sharing
- Common goal: allocate the bottleneck link capacity across packets from traffic classes
- This allocation occurs in the **packet scheduler** in the bottleneck router
  - Recall: scheduling is the task of choosing the packet (among buffered packets) which is transmitted over the output link
- A router is said to implement packet **scheduling policies**

# Why care about service guarantees?

- Influences how packets are treated at contentious resources in the core of the network
  - Regardless of the endpoint transport
- Service guarantees: prioritization, **rate limiting**, fair sharing
- Implementations of **scheduling (QoS)** within large networks have implications for debates on **network neutrality**
- Scheduling is a fundamental problem in computer networks

Next lecture: a deeper look at mechanisms for one kind of service guarantee







# CS 352

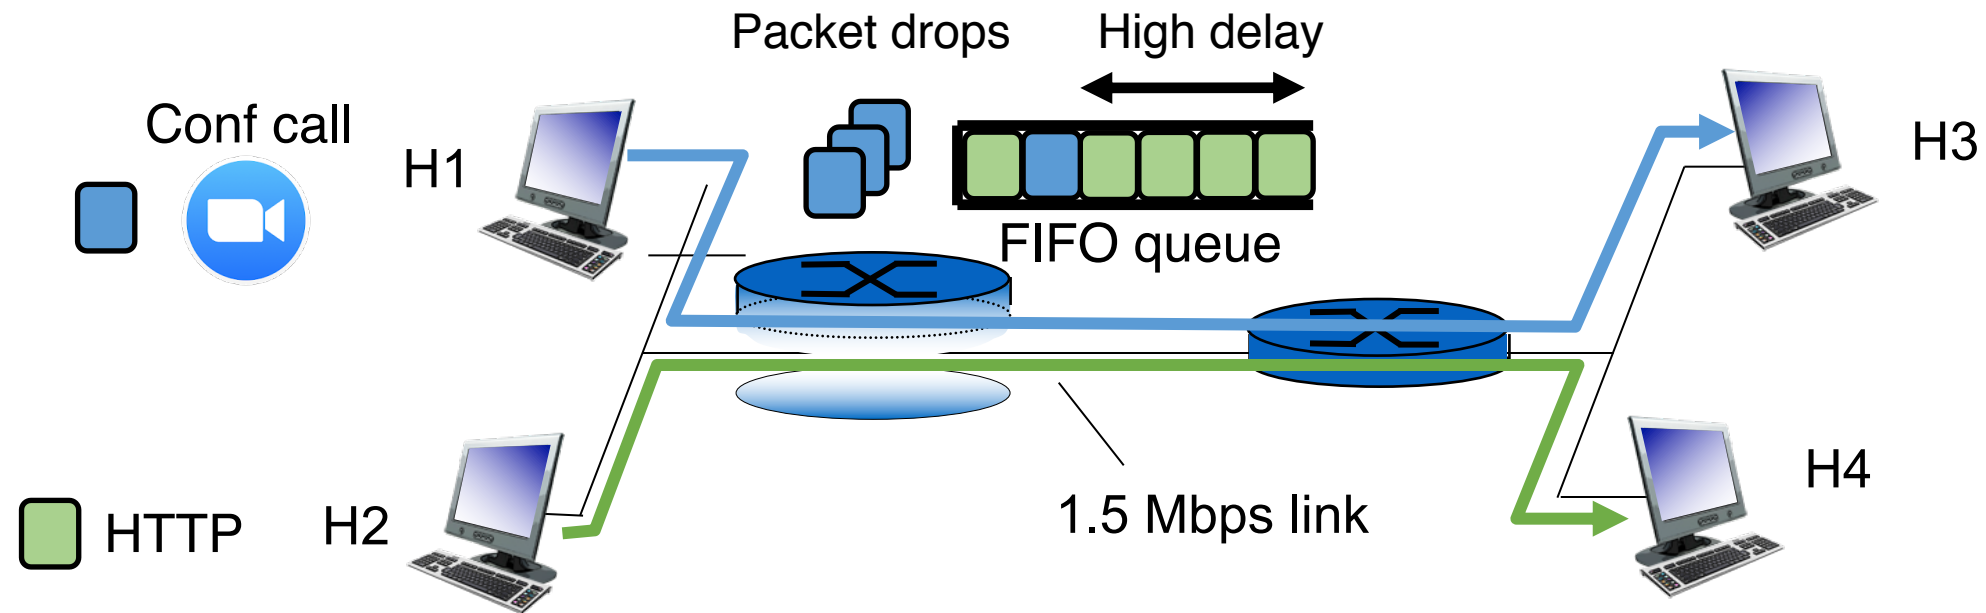
# Rate Limiting

CS 352, Lecture 20.2

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana


# Review



- Best-effort network isn't enough
  - Applications might require more guarantees
- 3 mechanisms: prioritization, rate limiting, and fair sharing
- This module: **implementation of rate limiting**

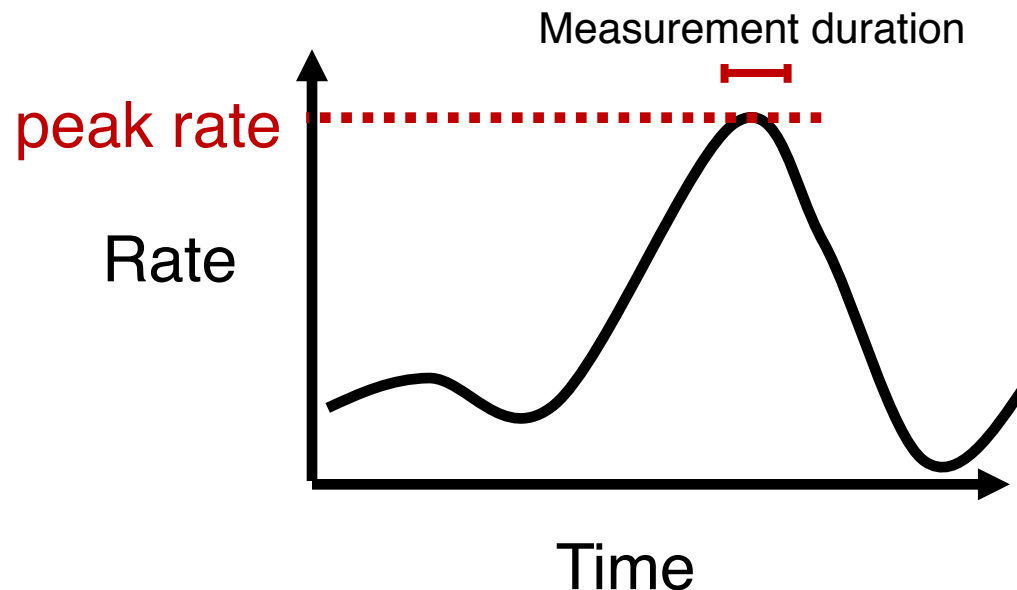
# Measures of transmission rate

- **Long-term/average rate:** data rate transmitted per unit time, over a long period
  - Crucial question: **what is the time interval** over which rate is measured?
- Average and instantaneous behaviors can be very different



# Measures of transmission rate

- **Peak rate:** largest instantaneous rate that is transmitted
  - Measurement duration is typically very small
- **Burst size:** maximum amount of data sent consecutively without any intervening idle periods



# Rate enforcement

- There are two kinds of rate enforcement policies:
  - shaping and policing
- Two specific mechanisms to implement those:
  - leaky buckets and token buckets

# Shaping

vs.

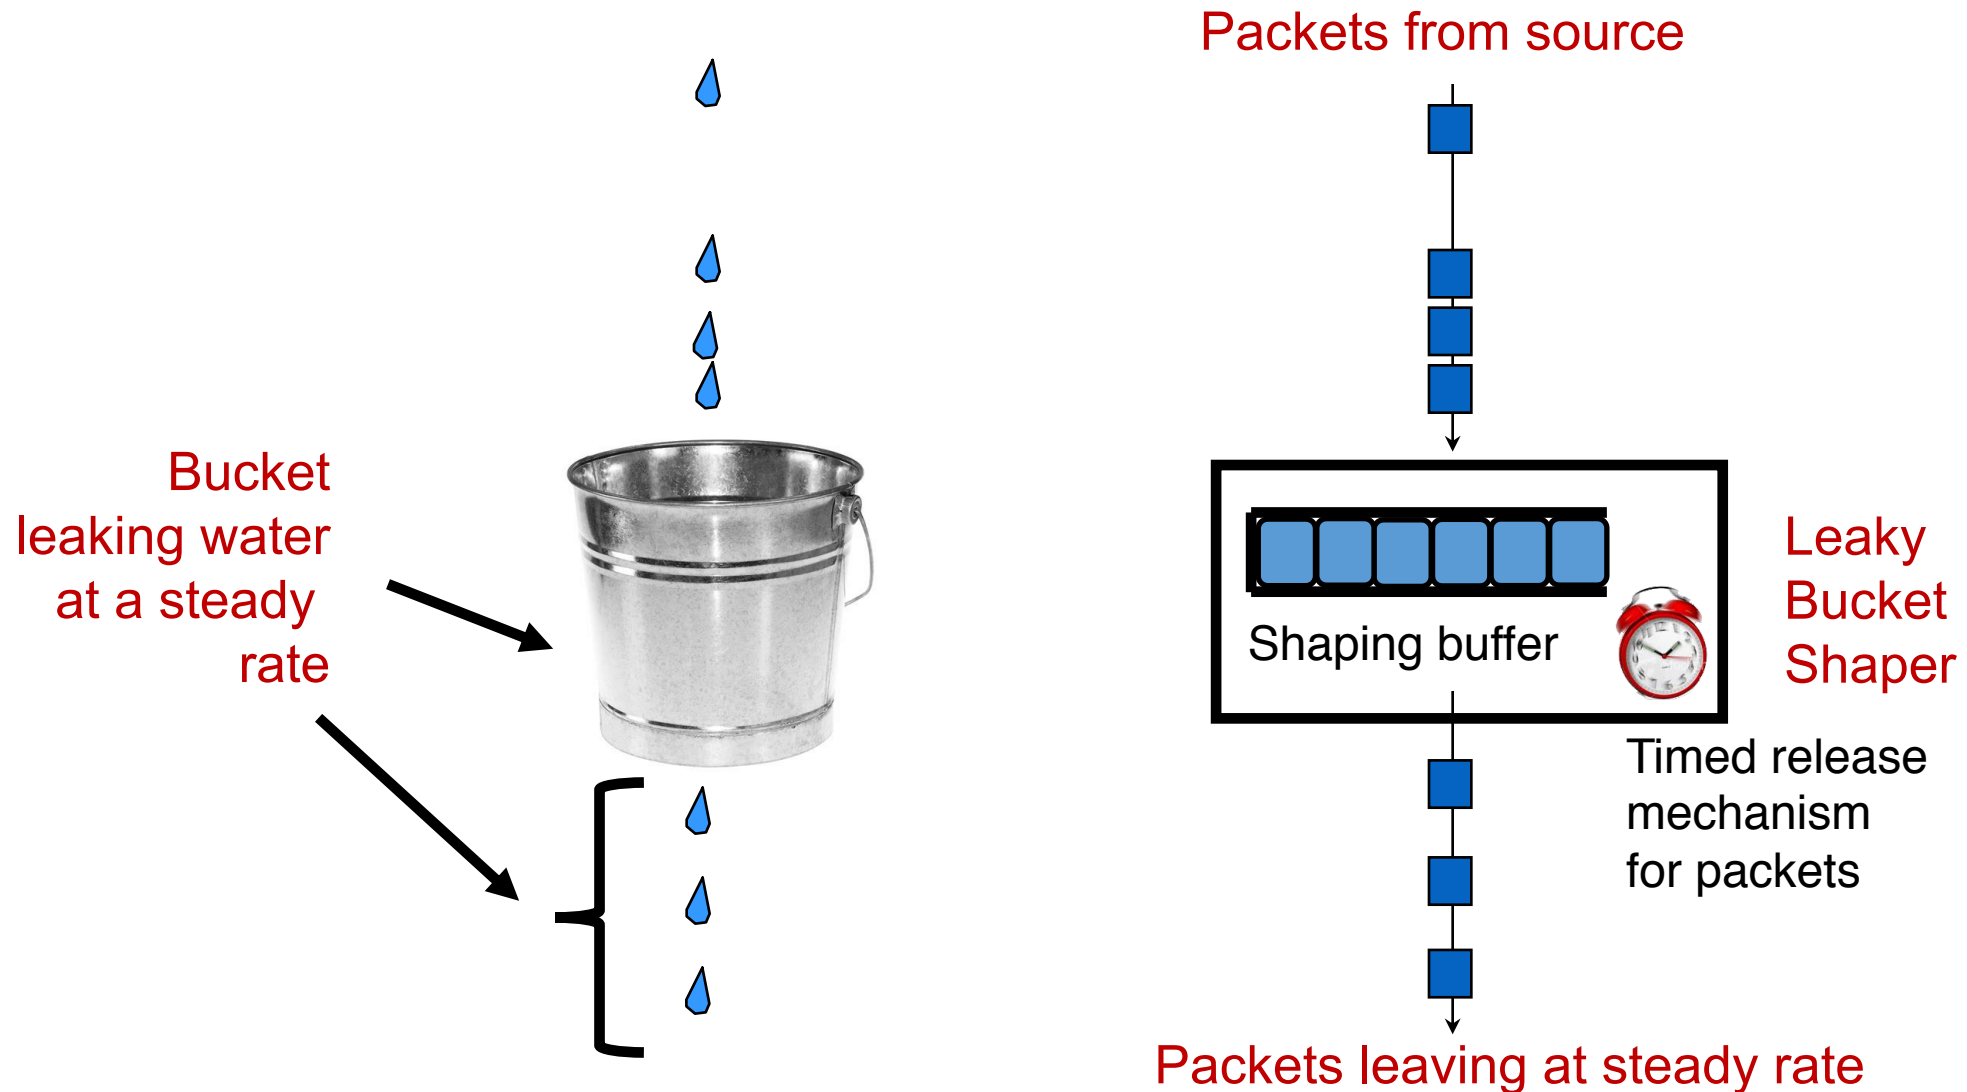
# Policing

- Enforces rate by **queueing** excess packets in a buffer
  - Drop only if buffer is full
- Requires memory to buffer packets
- Can inflate round-trip time (queueing in shaping buffer)

- Enforces rate by **dropping** excess packets immediately
  - Can result in high loss rates
- Does not require a memory buffer
- No additional inflation in round-trip times

Leaky bucket shaper

# Intuition: release packets at steady rate





# Leaky Bucket Shaper

- Packets may enter in a **bursty** manner
- However, once they pass through the leaky bucket, they are **evenly spaced**
- The **shaping buffer** holds packets up to a certain point
  - If the buffer is full, packets are dropped (true of any shaper)
- Setting the rate is a policy concern
  - Assume an admin provides us the rate
- Shapers may be used in the core of a network to limit bandwidth use, or at the edge to pace packets entering the network in the first place


# Leaky Bucket Shaper

- For a leaky bucket shaper, assume **average rate == peak rate**
- However, many Internet transfers just have a few packets
  - For example, web requests and responses
  - Enforcing rate limit for those can significantly delay completion
- We often wish to have peak rate higher than avg rate
  - If so, use a **token bucket**: burst-tolerant version of a leaky bucket

Token bucket shaper


# Token bucket shaper

- Limits traffic class to a specified average rate  $r$  and burst size  $B$
- Tokens are filled in at rate  $r$
- The token bucket can hold a maximum of  $B$  tokens. Further tokens dropped
  - Note: distinct from shaping buffer size
- Suppose a packet is at the head of the shaping buffer
- If a token exists in the bucket, **remove** token, and **transmit** the packet
  - If not, **wait**.



# Token bucket shaper


- In time  $t$ , the maximum number of packets that depart the shaper is  $(r * t) + B$
- A full bucket of tokens would allow small flows to go through unaffected
  - A maximum burst of  $B$  packets
- Longer flows have average rate  $r$ 
  - Bucket emptied initially, the rest of the flow must respect the token fill rate
  - As  $t \rightarrow \infty$ , the average rate approaches  $r$
  - That is,  $(1/t) * (r*t + B) \rightarrow r$



Token bucket policers

# Token bucket policer

- A token bucket policer is just a token bucket shaper without the shaping buffer
- No place for packets to wait if there are no tokens
- If token exists, packet transmitted.
- If not, packet **dropped**
- Simple and efficient to implement.
- The internet has tons of token bucket policers



# Google study from 2016

Region	Policed segments		Loss rate	
	(among lossy)	(overall)	(policed)	(non pol.)
India	6.8%	1.4%	28.2%	3.9%
Africa	6.2%	1.3%	27.5%	4.1%
Asia (w/o India)	6.5%	1.2%	22.8%	2.3%
South America	4.1%	0.7%	22.8%	2.3%
Europe	5.0%	0.7%	20.4%	1.3%
Australia	2.0%	0.4%	21.0%	1.8%
North America	2.6%	0.2%	22.5%	1.0%


Small but  
non-trivial  
fraction of  
policed links

Significant  
impact on  
packet loss  
rate

**Table 2:** % segments policed among lossy segments ( $\geq 15$  losses, the threshold to trigger the policing detector), and overall. Avg. loss rates for policed and unpoliced segments.




# Impact on TCP



**Figure 1: TCP sequence graph for a policed flow: (1 and 4) high throughput until token bucket empties, (2 and 5) multiple rounds of retransmissions to adjust to the policing rate, (3) idle period between chunks pushed by the application.**

# Effect on actual apps: YouTube

- Video rebuffer rate: rebuffer time / overall watch time



**Figure 9: Rebuffer to watch time ratios for video playbacks. Each had at least one chunk with a goodput of 300 kbps, 1.5 Mbps, or 3 Mbps ( $\pm 15\%$ ).**

# Summary of rate limiting

- Rate limiting is a useful mechanism to isolate traffic classes from each other
- Two strategies: policing and shaping
- Leaky bucket and token bucket
- The Internet has a lot of token bucket policers, causing real impact on TCP connections and app performance

