

# CS 352



# Bandwidth-Delay Product

CS 352, Lecture 13.1

<http://www.cs.rutgers.edu/~sn624/352>


Srinivas Narayana

# Transport



# How do apps get perf guarantees?

- The network core provides no guarantees on packet delivery



- Transport software on the endpoint oversees implementing guarantees on top of a best-effort network

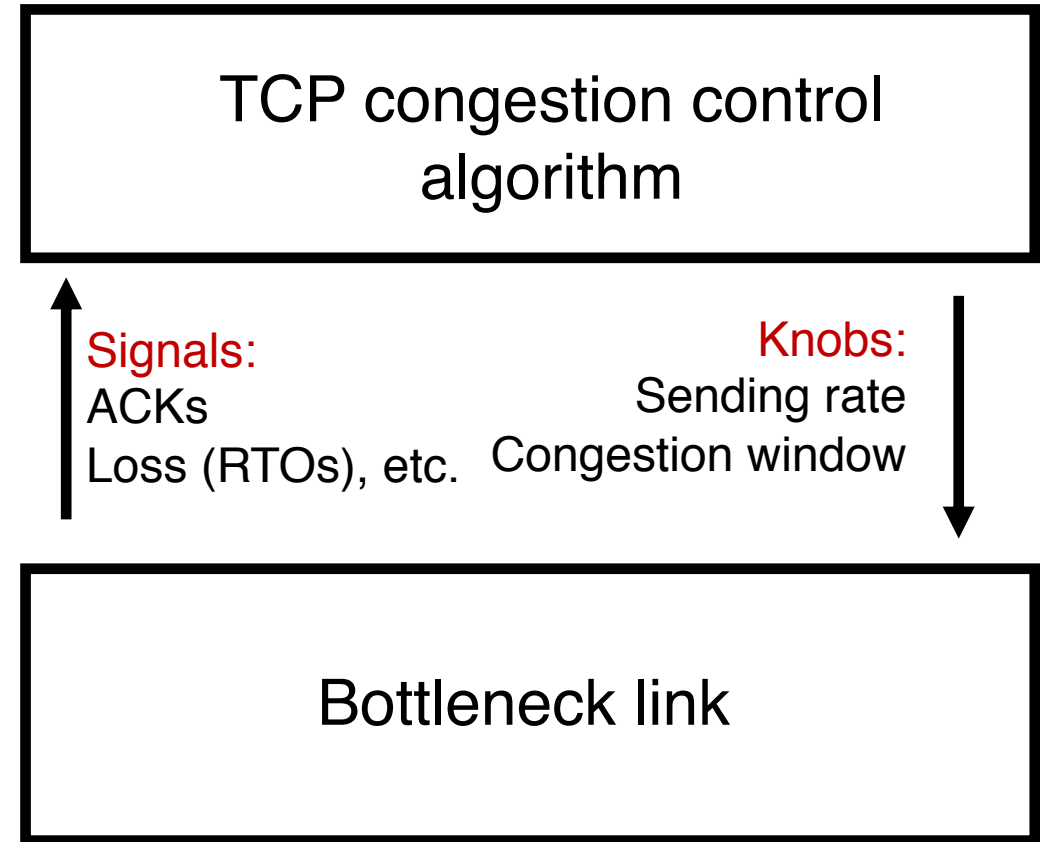
- Three important kinds of guarantees

- Reliability
- Ordered delivery
- Resource sharing in the network core

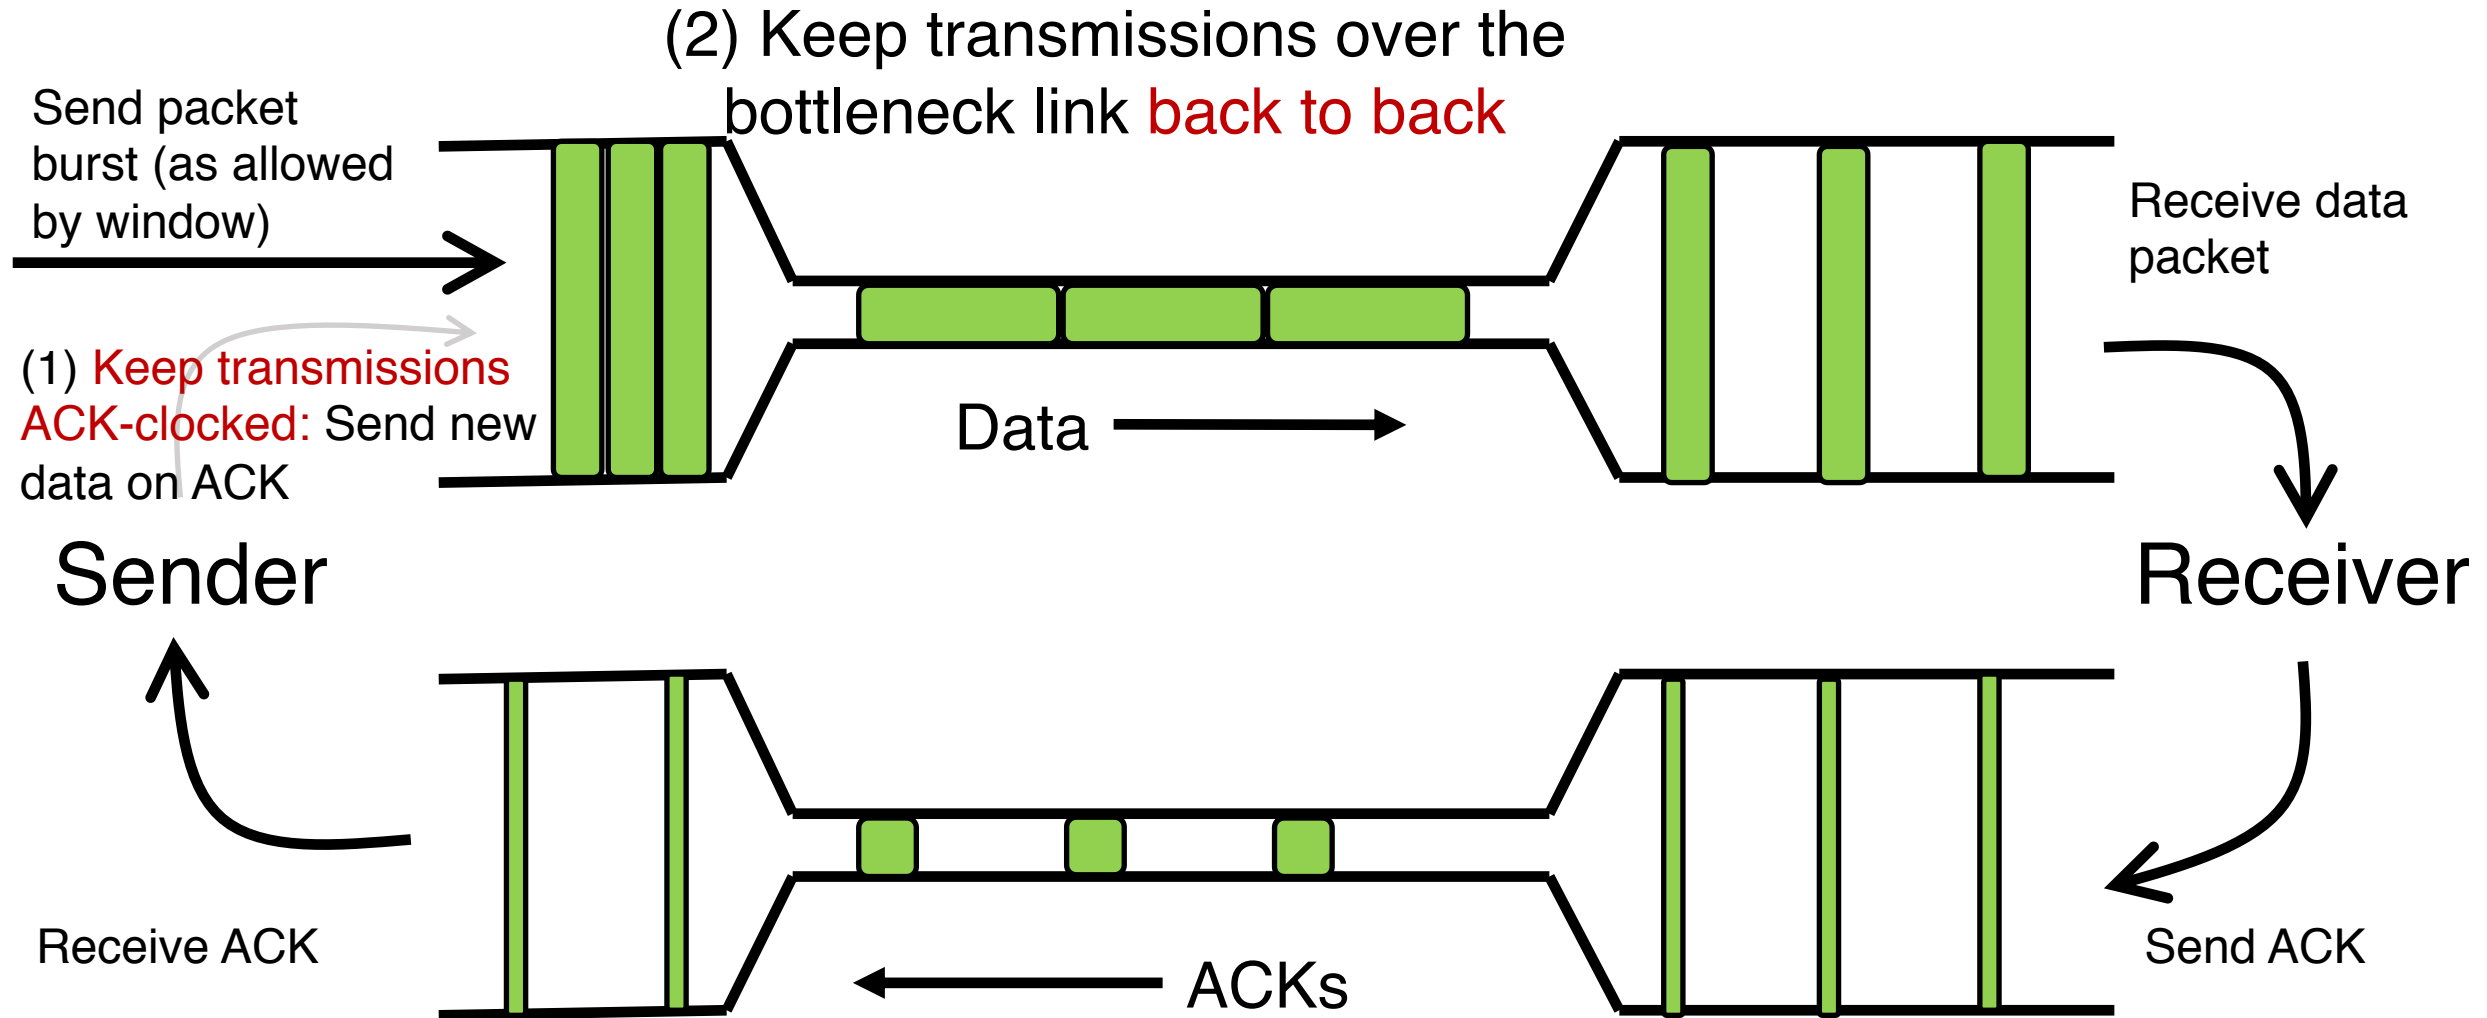
} Transmission  
Control Protocol  
(TCP)

# Review: Congestion control so far

- Algorithm by which multiple endpoints **efficiently** and **fairly** share bottleneck link
- So far, we've looked at just efficiency.
- Steady state: **ACK clocking** (keep the pipe full, but don't congest it)
- Getting to steady state:
  - Slow start: exponential increase
  - TCP New Reno: Additive increase
  - TCP BBR: gain cycling & filters



# Goal of steady state operation




# Steady state cwnd for a single flow

- Suppose the bottleneck link has rate  $C$
- Suppose the propagation round-trip delay (propRTT) between sender and receiver is  $T$
- Ignore transmission delays for this example;
- Assume steady state i.e., high sending rate, no bottleneck congestion
- Q: how much data is in flight over a single RTT?
- $C * T$  data i.e., amount of data unACKed at any point in time
- ACKs take time  $T$  to arrive (without any queueing). In the meantime, sender is transmitting at rate  $C$


# The Bandwidth-Delay Product

- $C * T$  = **bandwidth-delay product**:
  - The amount of data in flight for a sender transmitting at the ideal rate during the ideal round-trip delay of a packet
- Note: this is just the amount of data “on the pipe”



# The Bandwidth-Delay Product


- Q: What happens if  $cwnd > C * T$ ?
  - i.e., where are the rest of the in-flight packets?
- A: Waiting at the bottleneck router queues





# Router buffers and the max cwnd

- Router buffer memory is finite: queues can only be so long
  - If the router buffer size is  $B$ , there is at most  $B$  data waiting in the queue
- If cwnd increases beyond  $C * T + B$ , data is dropped!



# Summary

- Bandwidth-Delay Product (BDP) governs the window size of a single flow at steady state
- The bottleneck router buffer size governs how much the `cwnd` can exceed the BDP before packet drops occur



# CS 352

# Detecting & Reacting to Losses

CS 352, Lecture 13.2

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Detecting packet loss

- So far, all the algorithms we've studied have a coarse loss detection mechanism: RTO timer expiration
  - Let the RTO expire, drop `cwnd` all the way to 1 MSS
- Analogy: you're driving a car
  - You're waiting until the next car in front is super close to you (RTO) and then hitting the brakes really hard (set `cwnd := 1`)
  - Q: Can you see obstacles from afar and slow down proportionately?
- That is, can the sender see packet loss coming in advance?
  - And reduce `cwnd` more gently?

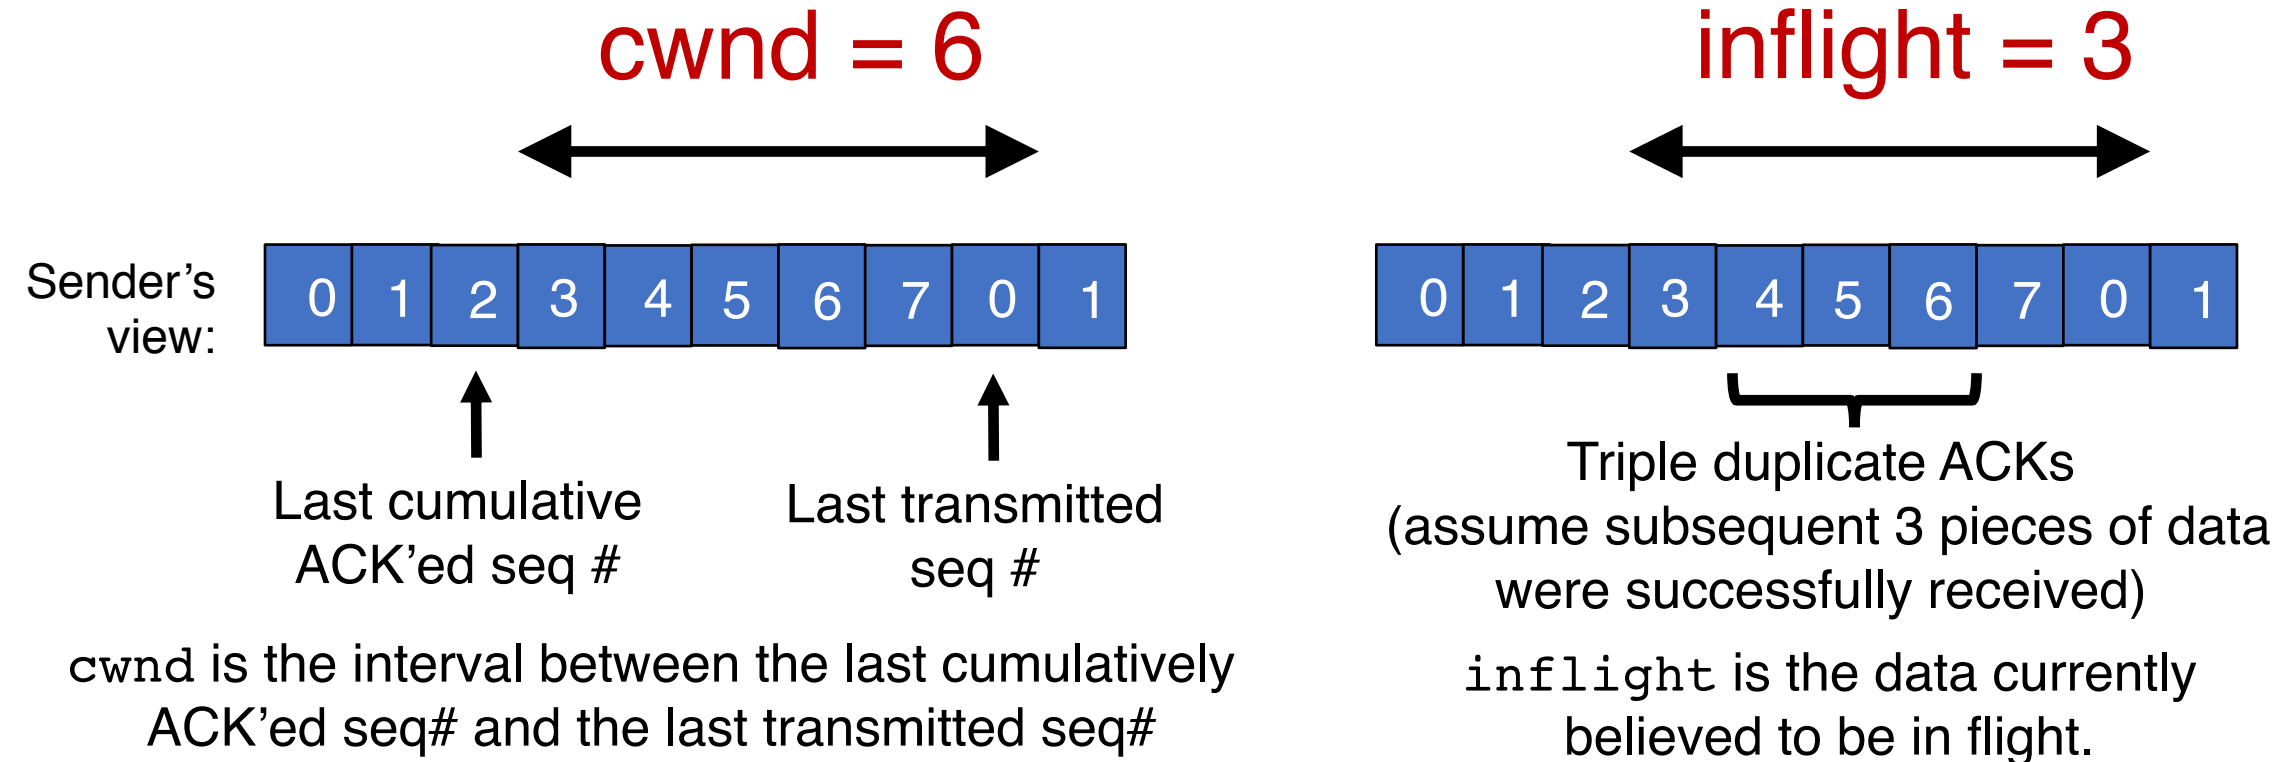
# Can we detect loss earlier than RTO?

- Key idea: use the information in the ACKs. **How?**
- Suppose successive (cumulative) ACKs contain the same ACK#
  - Also called **duplicate ACKs**
  - Occur when network is reordering packets, or one (but not most) packets in the window were lost
- Reduce `cwnd` when you see many duplicate ACKs
  - Consider many dup ACKs a strong indication that packet was lost
  - Default threshold: 3 dup ACKs, i.e., **triple duplicate ACK**
  - **Make `cwnd` reduction gentler than setting `cwnd = 1`; recover faster**

# Fast Retransmit & Fast Recovery

# Distinction: In-flight versus window

- So far, window and in-flight referred to the same data
- Fast retransmit & fast recovery differentiate the two notions





# TCP **fast retransmit** (RFC 2581)

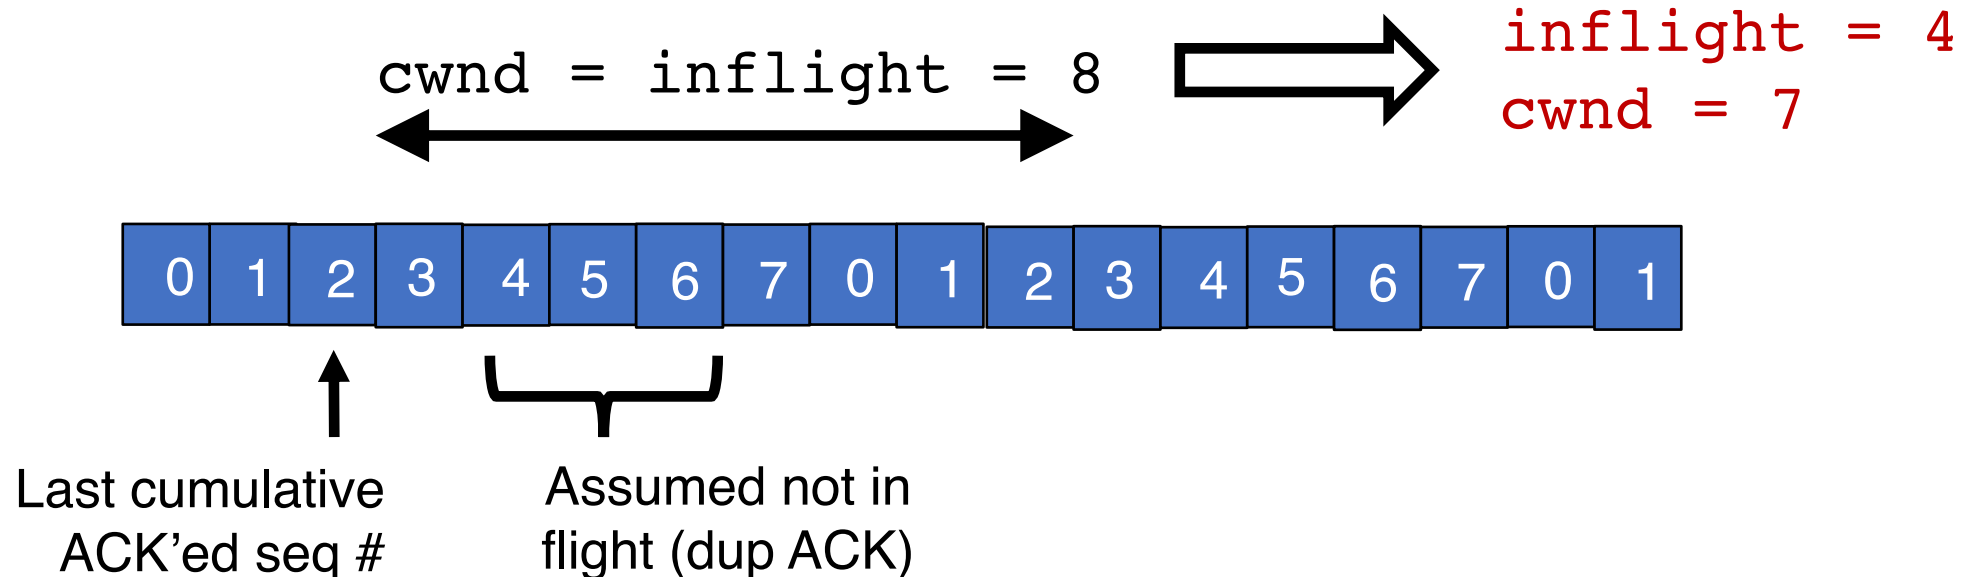
- The fact that ACKs are coming means that data is getting delivered to the receiver, albeit with some loss.
- Note: Before the dup ACKs arrive, we assume `inflight = cwnd`
- TCP sender does two actions with fast retransmit

# TCP fast retransmit (RFC 2581)

- (1) Reduce the `cwnd` and `in-flight` gently
  - Don't drop `cwnd` all the way down to 1 MSS
- Reduce the amount of in-flight data **multiplicatively**
  - Set `inflight`  $\rightarrow$  `inflight / 2`
  - That is, set `cwnd`  $=$  `(inflight / 2) + 3MSS`
  - This step is called **multiplicative decrease**
  - Algorithm also sets `ssthresh` to `inflight / 2`

# TCP **fast retransmit** (RFC 2581)

- Example: Suppose `cwnd` and `inflight` (before triple dup ACK) were both 8 MSS.
- After triple dup ACK, reduce `inflight` to 4 MSS
- *Assume* 3 of those 8 MSS no longer in flight; set `cwnd` = 7 MSS



# TCP fast retransmit (RFC 2581)


- (2) The seq# from dup ACKs is immediately retransmitted
- That is, don't wait for an RTO if there is sufficiently strong evidence that a packet was lost

# TCP **fast recovery** (RFC 2581)

- Sender keeps the reduced `inflight` until a **new ACK** arrives
  - New ACK: an ACK for the `seq#` that was just retransmitted
  - May also include the (three or more) pieces of data that were subsequently delivered to generate the duplicate ACKs
- **Conserve packets in flight**: transmit *some* data over lossy periods (rather than no data, which would happen if `cwnd := 1`)

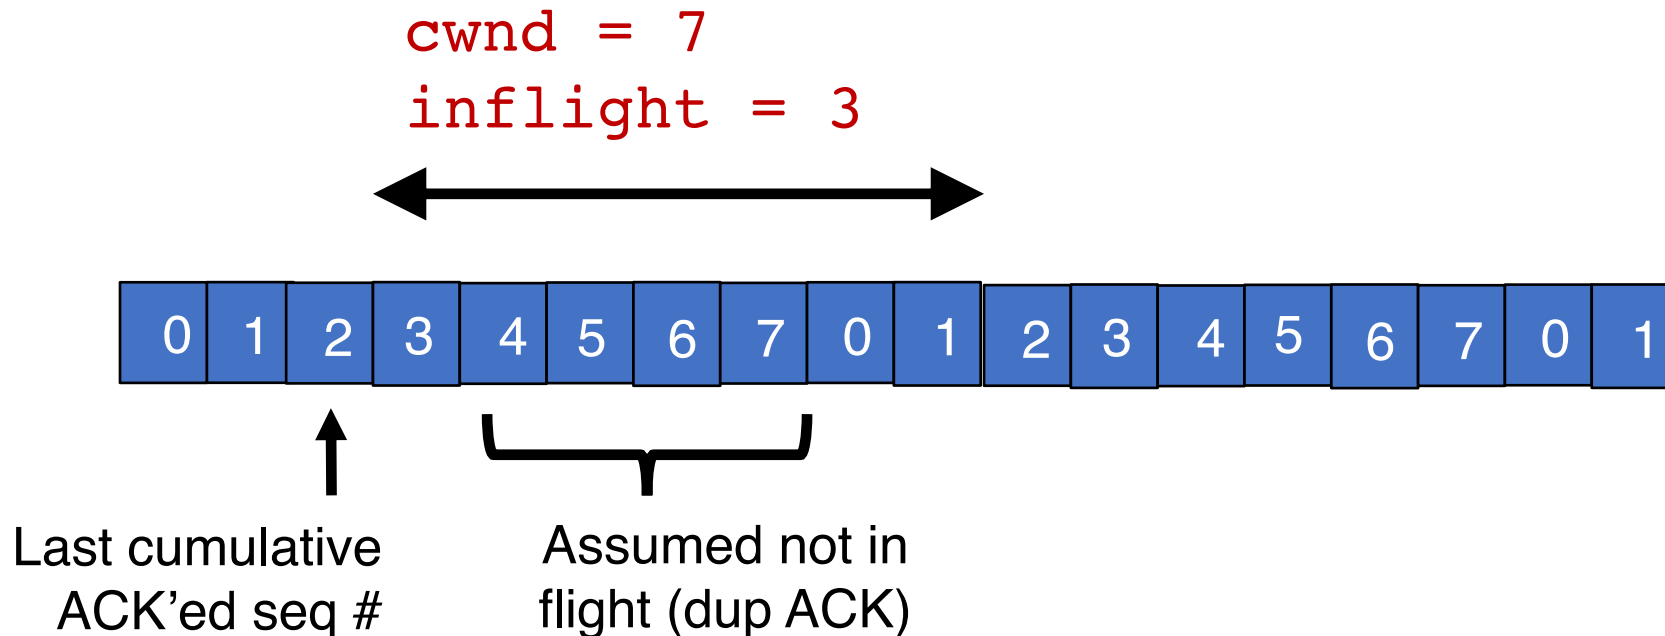
# TCP **fast recovery** (RFC 2581)

- Keep incrementing `cwnd` by 1 MSS for each dup ACK



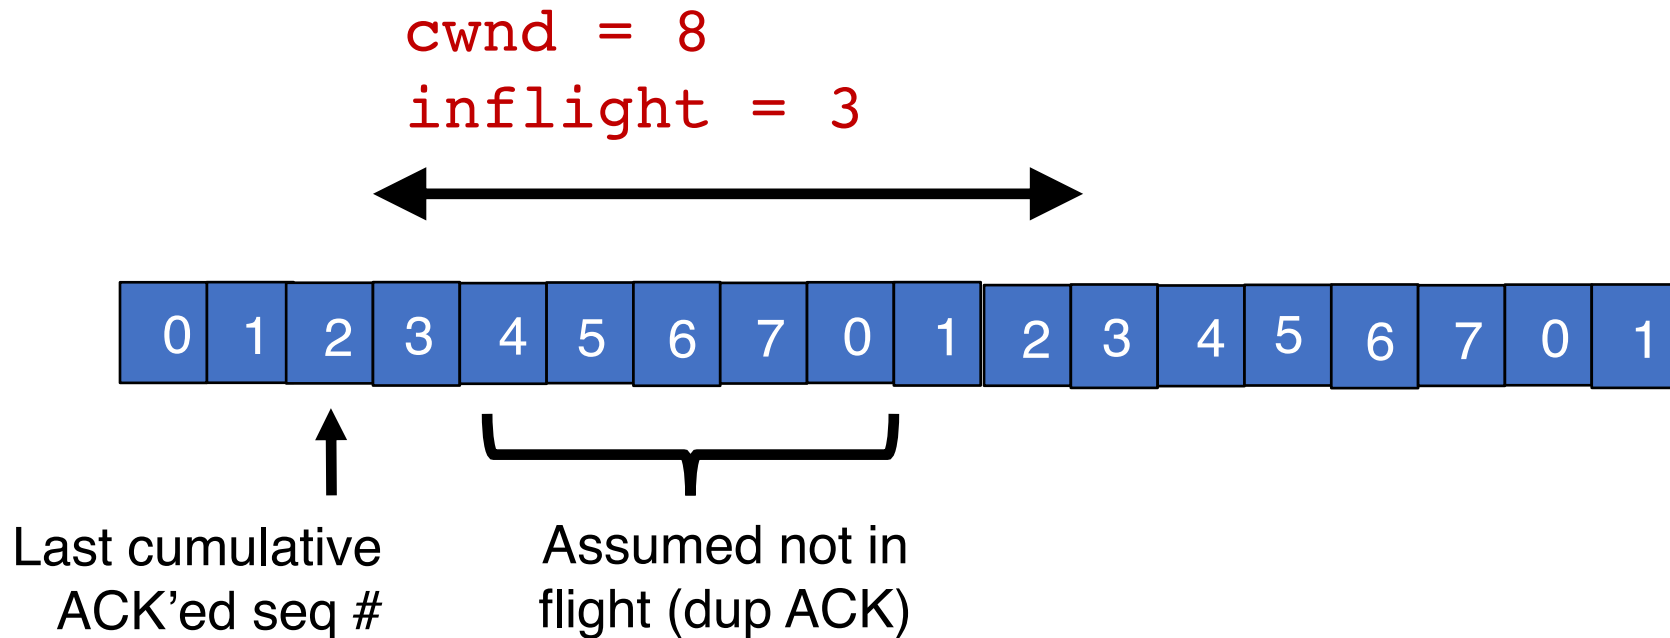
# TCP **fast recovery** (RFC 2581)

- Keep incrementing cwnd by 1 MSS for each dup ACK



# TCP **fast recovery** (RFC 2581)


- Keep incrementing cwnd by 1 MSS for each dup ACK





# TCP **fast recovery** (RFC 2581)

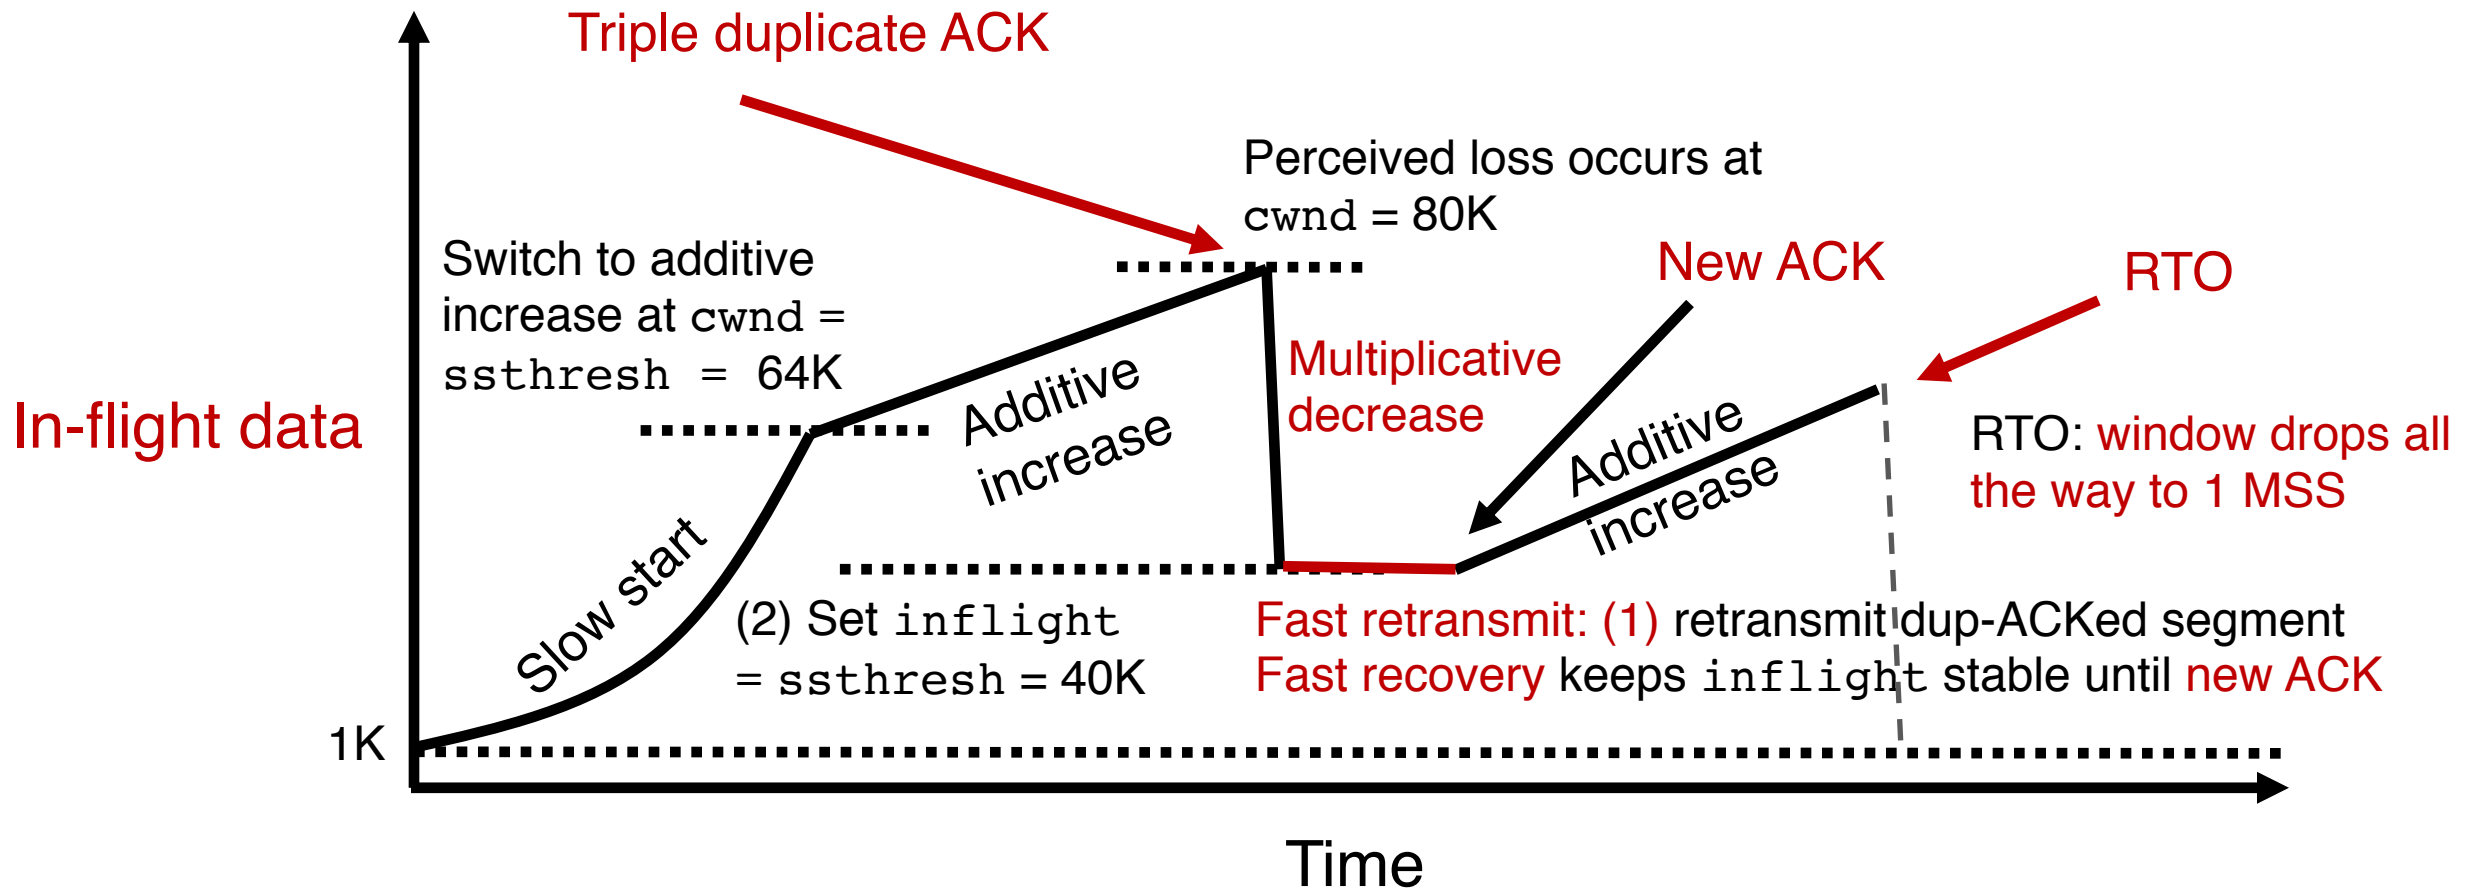
- Eventually a **new ACK** arrives, acknowledging the retransmitted data and all data in between
- Deflate cwnd to half of cwnd before fast retransmit.
  - cwnd and inflight are aligned and equal once again
- Perform **additive increase** from this point!



# Additive Increase/Multiplicative Decrease

Say MSS = 1 KByte

Default ssthresh = 64KB = 64 MSS



TCP New Reno performs additive increase and multiplicative decrease of its congestion window.

In short, we often refer to this as **AIMD**.

**Multiplicative decrease** is a part of all TCP algorithms, including BBR.

[It is necessary for **fairness** across TCP flows.]

# Summary of TCP loss detection

- Don't wait for an RTO and then set the `cwnd` to 1 MSS
  - Tantamount to waiting to get super close to the car in front and then jamming the brakes really hard
- Instead, react proportionately by sensing pkt loss in advance

## Fast Retransmit

- **Triple dup ACK:** sufficiently strong signal that network has dropped data, before RTO
- Immediately retransmit data
- Multiplicatively decrease in-flight data to **half** of its value

## Fast Recovery

- Maintain this reduced amount of in-flight data as long as dup ACKs arrive
  - Data is successfully getting delivered
- When **new ACK** arrives, do **additive increase** from there on



# CS 352

## Computing the Retransmit Timeout

CS 352, Lecture 13.3

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana


# TCP timeout (RTO)

- Useful for reliable delivery and congestion control
- How to pick the RTO value?
  - Too long: slow reaction to loss
  - Too short: premature retransmissions which are wasteful
- Want: RTO must predict the upper bound of RTTs resulting from a successful packet + ACK
- Intuition: somehow use the observed RTT (sampleRTT)
  - Can we just directly set the latest RTT as the RTO?
- No. RTT can vary significantly!
  - Intermittent congestion, path changes, signal quality changes on wireless channel, etc.

# Estimate an “average” RTT

- Exponential weighted moving average (typical  $\alpha = 1/8$ )


$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$





# Accounting for RTT variance

- RTT samples can have a large **variance**
- Use a **safety margin** in the RTO estimate to account for variance



# TCP timeout computation

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )



$$\text{RTO} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



average RTT



safety margin


Conceptually, there is an RTO timer for each seq #.

# Too many timers?

- **Timers are expensive** – we don't want one per sequence #
  - Interrupts, OS data structures, and book-keeping
- The TCP stack maintains just one “real” timer per connection
- When a packet is transmitted, its transmission time is recorded
- The only real timer in the system is the RTO for the first unACK'ed segment
  - Expiration interval: RTO
- **If ACK before RTO fires:** set timer for next unACK'ed segment, based on recorded transmission time of that segment
- **If RTO fires:** retransmit the segment, restart RTO timer

Retransmission ambiguity

# Real RTT of a retransmitted segment?



# How to estimate RTT/RTO despite retransmit?

- One solution: **Never update RTT measurements** based on ACKs from retransmitted packets
- Problem: **Sudden change in RTT**, coupled with many retransmissions, can cause system to update RTT very late
  - Ex: Primary path failure leads to a high-RTT secondary path
- If RTT estimates are not updated, the RTO estimate isn't, and that leads to a host of other problems.
  - Ex: Unnecessary retransmissions since RTOs needlessly expire

# Karn's algorithm

- Use **back-off** as part of the `sampleRTT` computation
- Whenever packet loss (RTO), RTO is increased by a factor
  - Conservatively assume that RTT may have increased since the last unambiguous RTT samples were obtained
- Use this increased RTO as RTO estimate for the next segment
  - Don't use the `estimatedRTT` from stale `sampleRTT`
- Only after an ACK is received for a successful transmission is the RTO timer set to a value obtained from `EstimatedRTT`

# Summary

- RTO computation is an important part of TCP's behavior under loss
- TCP uses both an average RTT as well as the variance to obtain a safe prediction of an upper bound of a successful RTT
- Resolve retransmission ambiguity under path changes by avoiding `sampleRTT` measurements and multiplicatively increasing the RTO each time





# CS 352

# TCP Connection Management

CS 352, Lecture 13.4

<http://www.cs.rutgers.edu/~sn624/352>


Srinivas Narayana

# TCP connections need lots of bookkeeping

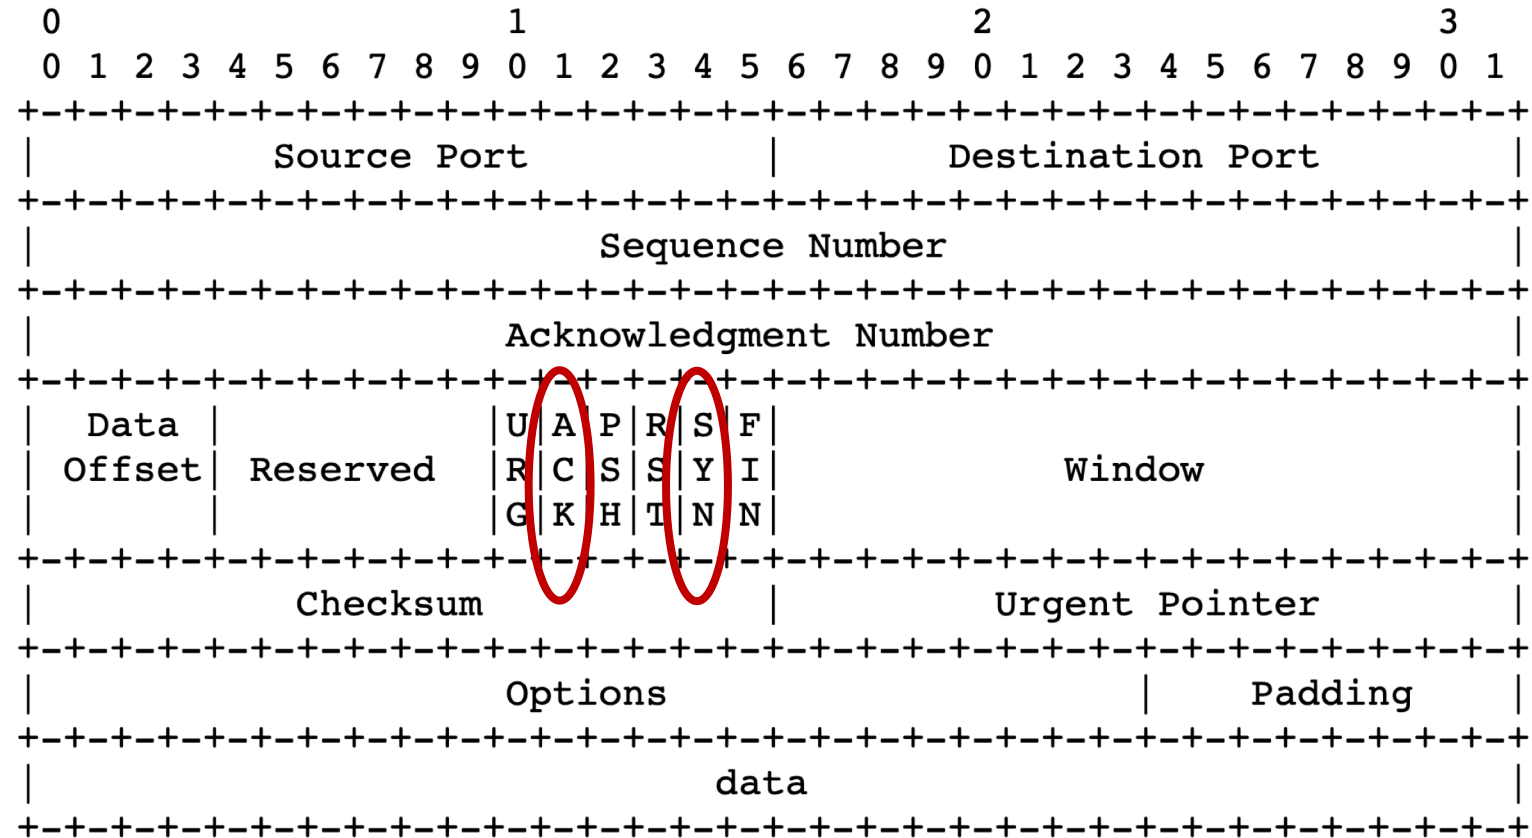
- Socket buffer memory
- Entries in connection lookup tables
- Data structures and parameters (e.g., sequence numbers) in the operating system kernel
- These resources can get expensive on machines running many connections, e.g., web servers

# Handshake

- Before starting data transmission, TCP client and server perform a **handshake** and agree on parameters
- TCP is **bidirectional**: independent set of sequence numbers for each direction
- Sequence numbers start from a random initial value
- Specific TCP **flags** indicate connection initiation and acceptance



# TCP flags in the header




TCP Header Format

Note that one tick mark represents one bit position.


# 2-way handshake not enough

- Suppose the server receives the first SYN packet and decides to allocate all the resources needed for the connection.
- What happens if a malicious client sends a ton of SYN packets?
- **Asymmetric work:** client doesn't need to allocate any resources of its own
  - Just have to send a well-crafted packet
- However, server's resources exhausted!
- **SYN flood attack:** a form of **denial of service**



# Consequences

- The server should not allocate resources upon receiving the first client message (SYN)
- The server cannot carry any application data in SYN/ACK
  - Server hasn't yet allocated all necessary resources
- Client cannot send any data in the SYN packet
- Recall: HTTP requires an RTT for the handshake before sending HTTP request

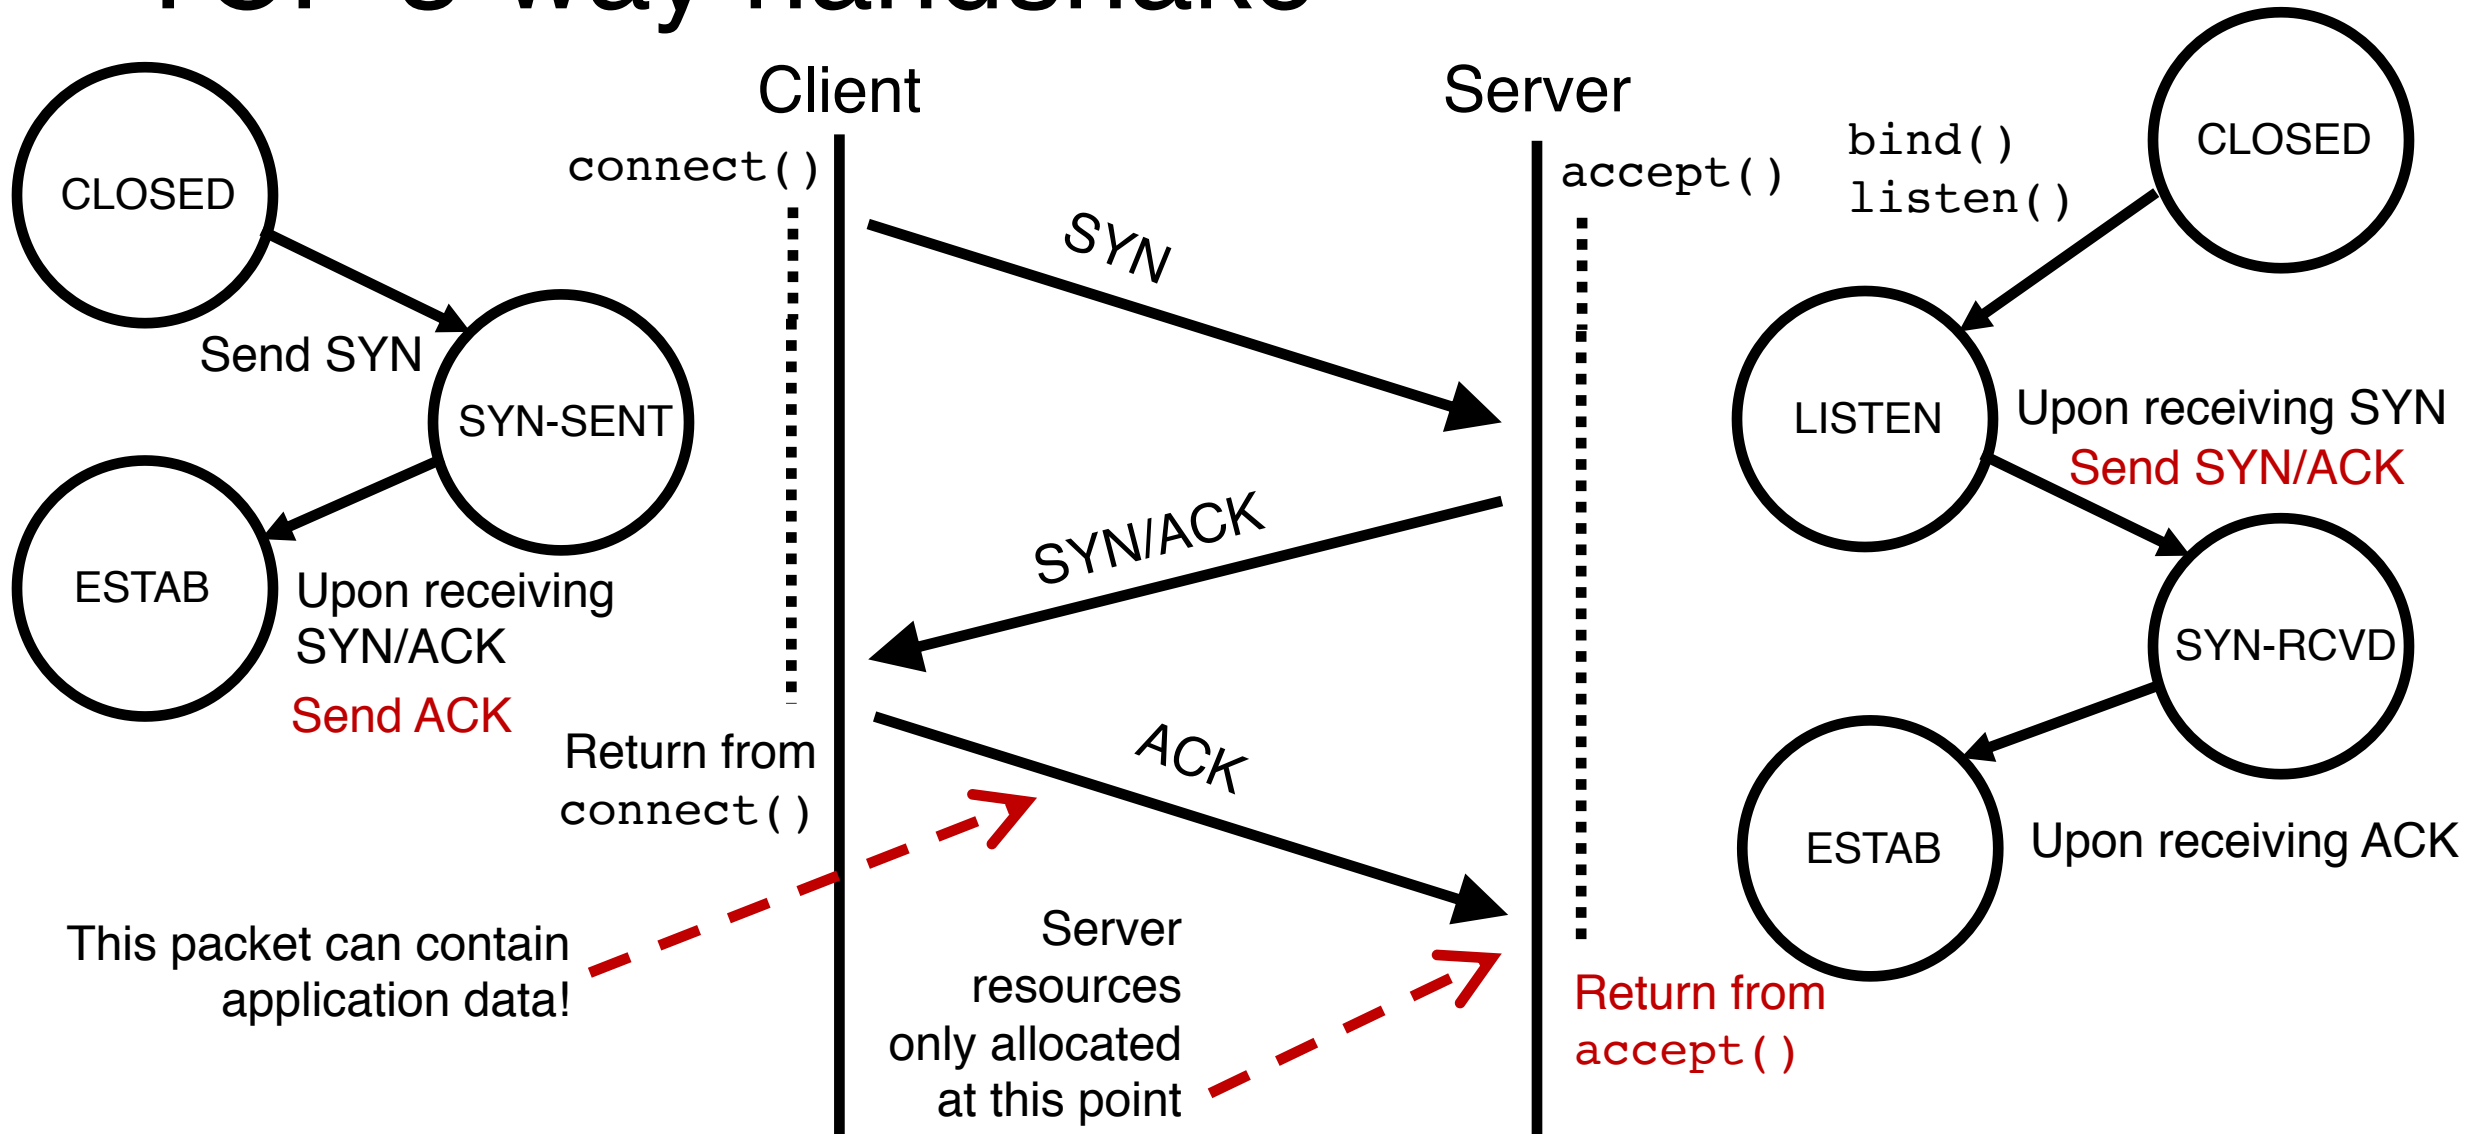


# Mitigating the denial of service problem

- Key idea: **Make the client do more work** before allocating server resources
- The client should send at least one more packet, responding to the data in the server's SYN/ACK, before the server decides to call the connection **established**
  - That is, before all required server resources like buffers are allocated
- Result: 3-way handshake
- Per-connection **finite state machine** tracks this process



# TCP 3-way handshake



# TCP: Closing a connection

- Client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- In general, **TCP is full-duplex**: both sides can send
- **However, FIN is unidirectional**: stop one side of the communication
- Respond to received FIN with ACK
  - On receiving FIN, ACK can be combined with own FIN
- Simultaneous FIN exchanges can be handled

# Summary of TCP connection management

- TCP connections have associated resources: managing them requires book-keeping the establishment of a connection carefully
- Simple 2-way handshakes suffer from denial of service vulnerability
  - Moral: don't allocate resources on the first client message
- 3-way handshake mitigates this issue by making client work harder
  - Client must send ACK to server's SYN/ACK before server can handle data
  - The cost: increased time before sending application data from client

