

# CS 352



# Detecting & Reacting to Losses

CS 352, Lecture 13.1

<http://www.cs.rutgers.edu/~sn624/352>


Srinivas Narayana

# Transport



# How do apps get perf guarantees?

- The network core provides no guarantees on packet delivery



- Transport software on the endpoint oversees implementing guarantees on top of a best-effort network

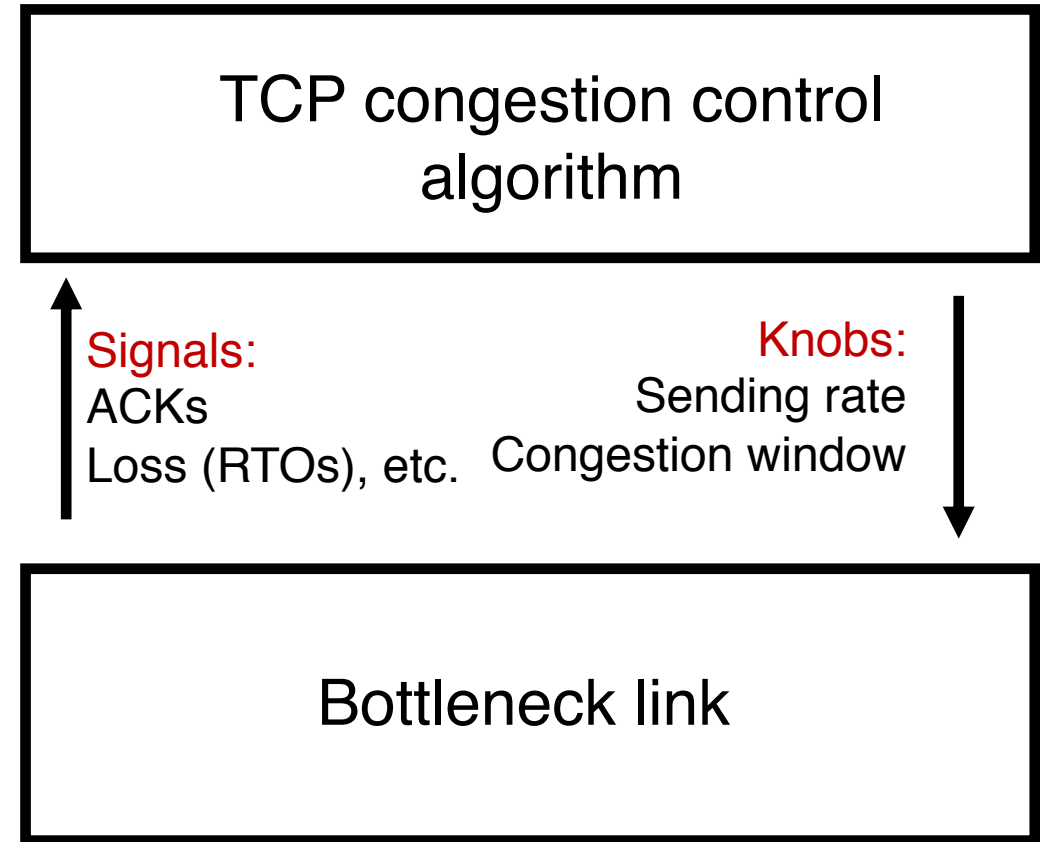
- Three important kinds of guarantees

- Reliability
- Ordered delivery
- Resource sharing in the network core

} Transmission  
Control Protocol  
(TCP)

# Review: Congestion control so far

- Algorithm by which multiple endpoints **efficiently** and **fairly** share bottleneck link
- So far, we've looked at just efficiency.
- Steady state: **ACK clocking** (keep the pipe full, but don't congest it)
- Getting to steady state:
  - Slow start: exponential increase
  - TCP New Reno: Additive increase
  - TCP BBR: gain cycling & filters



# Detecting packet loss

- So far, all the algorithms we've studied have a coarse loss detection mechanism: RTO timer expiration
  - Let the RTO expire, drop `cwnd` all the way to 1 MSS
- Analogy: you're driving a car
  - You're waiting until the next car in front is super close to you (RTO) and then hitting the brakes really hard (set `cwnd := 1`)
  - Q: Can you see obstacles from afar and slow down proportionately?
- That is, can the sender see packet loss coming in advance?
  - And reduce `cwnd` more gently?

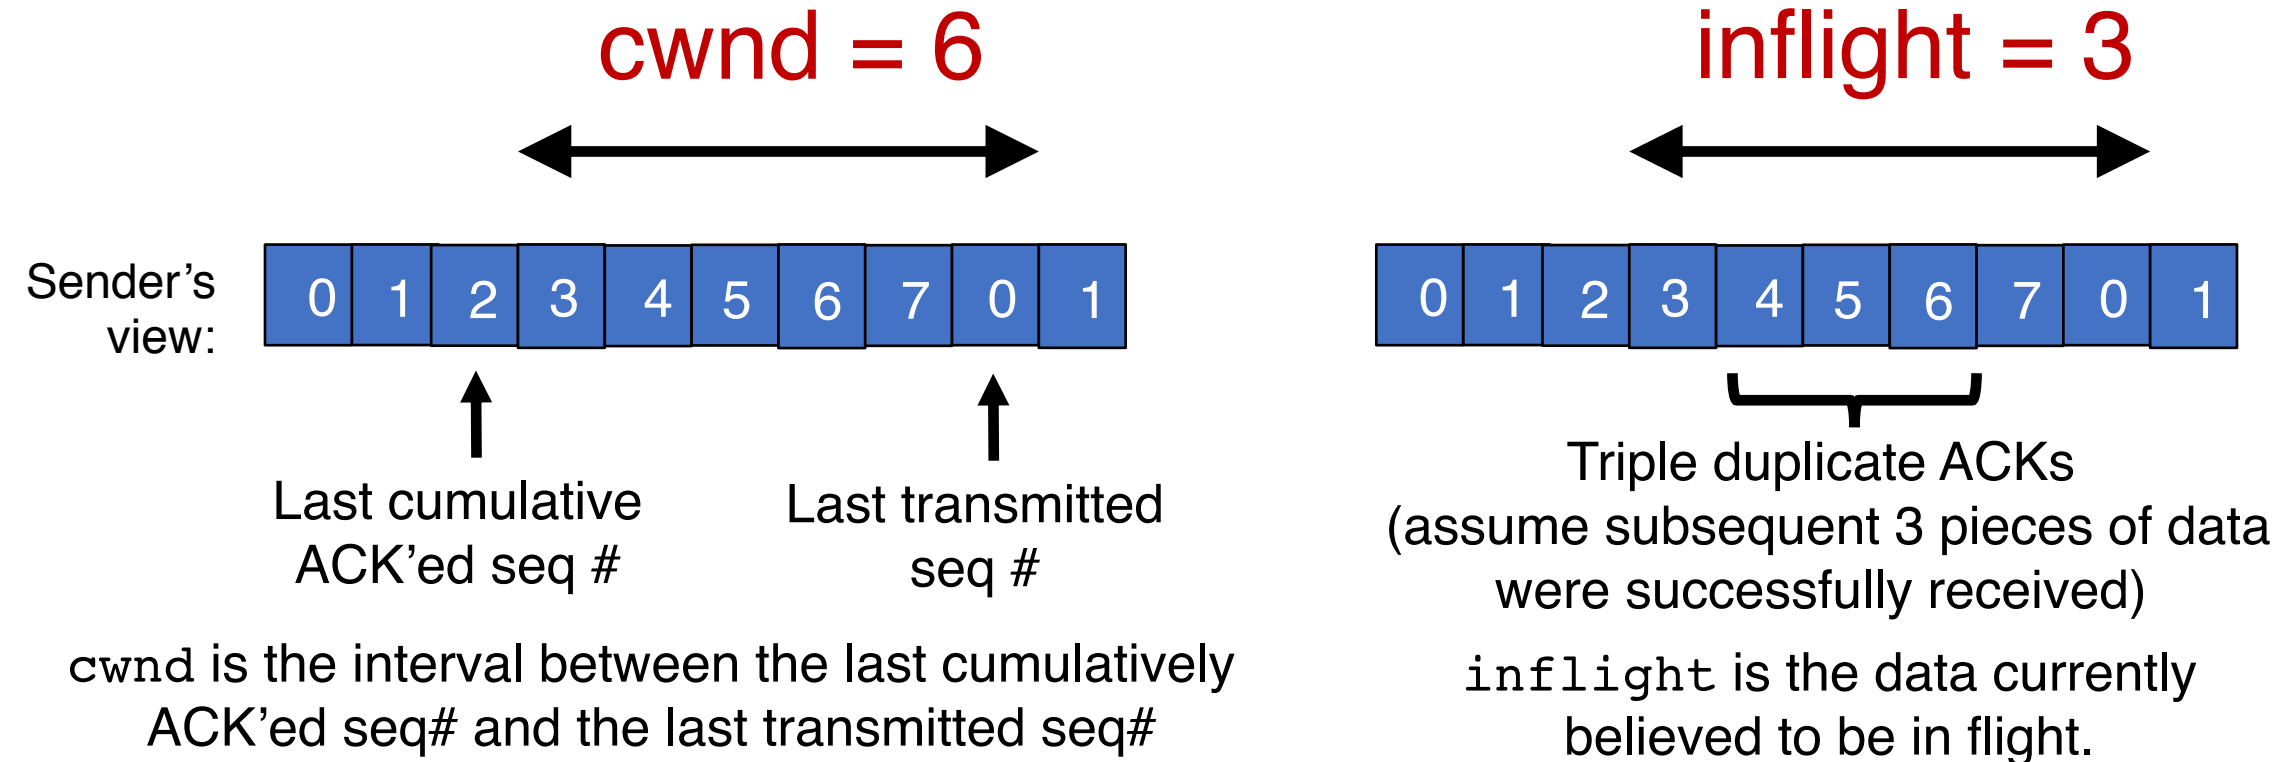
# Can we detect loss earlier than RTO?

- Key idea: use the information in the ACKs. **How?**
- Suppose successive (cumulative) ACKs contain the same ACK#
  - Also called **duplicate ACKs**
  - Occur when network is reordering packets, or one (but not most) packets in the window were lost
- Reduce `cwnd` when you see many duplicate ACKs
  - Consider many dup ACKs a strong indication that packet was lost
  - Default threshold: 3 dup ACKs, i.e., **triple duplicate ACK**
  - **Make `cwnd` reduction gentler than setting `cwnd = 1`; recover faster**

# Fast Retransmit & Fast Recovery

# Distinction: In-flight versus window

- So far, window and in-flight referred to the same data
- Fast retransmit & fast recovery differentiate the two notions





# TCP **fast retransmit** (RFC 2581)

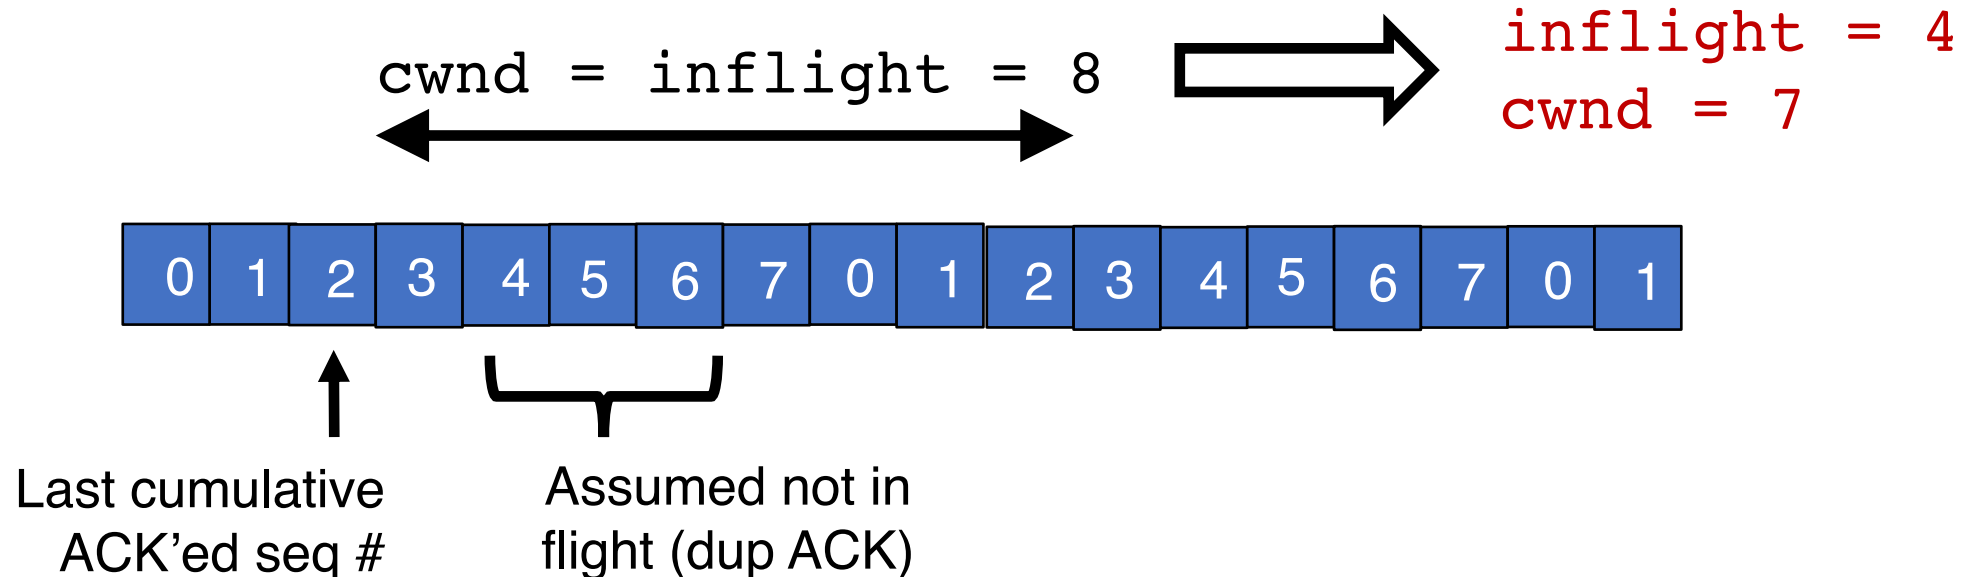
- The fact that ACKs are coming means that data is getting delivered to the receiver, albeit with some loss.
- Note: Before the dup ACKs arrive, we assume `inflight = cwnd`
- TCP sender does two actions with fast retransmit

# TCP fast retransmit (RFC 2581)

- (1) Reduce the `cwnd` and `in-flight` gently
  - Don't drop `cwnd` all the way down to 1 MSS
- Reduce the amount of in-flight data **multiplicatively**
  - Set `inflight`  $\rightarrow$  `inflight / 2`
  - That is, set `cwnd`  $= (\text{inflight} / 2) + 3\text{MSS}$
  - This step is called **multiplicative decrease**
  - Algorithm also sets `ssthresh` to `inflight / 2`

# TCP **fast retransmit** (RFC 2581)

- Example: Suppose `cwnd` and `inflight` (before triple dup ACK) were both 8 MSS.
- After triple dup ACK, reduce `inflight` to 4 MSS
- *Assume* 3 of those 8 MSS no longer in flight; set `cwnd` = 7 MSS



# TCP fast retransmit (RFC 2581)


- (2) The seq# from dup ACKs is immediately retransmitted
- That is, don't wait for an RTO if there is sufficiently strong evidence that a packet was lost

# TCP **fast recovery** (RFC 2581)

- Sender keeps the reduced `inflight` until a **new ACK** arrives
  - New ACK: an ACK for the `seq#` that was just retransmitted
  - May also include the (three or more) pieces of data that were subsequently delivered to generate the duplicate ACKs
- **Conserve packets in flight**: transmit *some* data over lossy periods (rather than no data, which would happen if `cwnd := 1`)

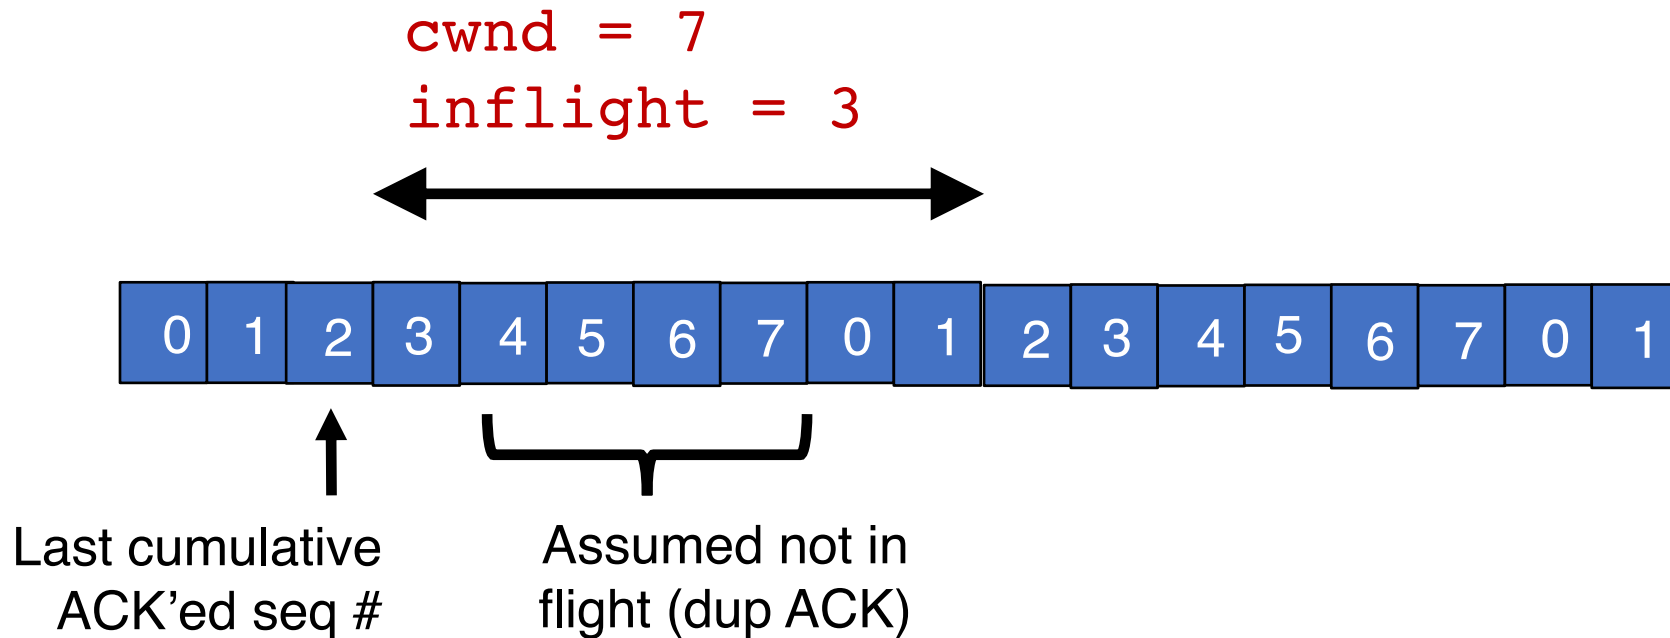
# TCP **fast recovery** (RFC 2581)

- Keep incrementing cwnd by 1 MSS for each dup ACK



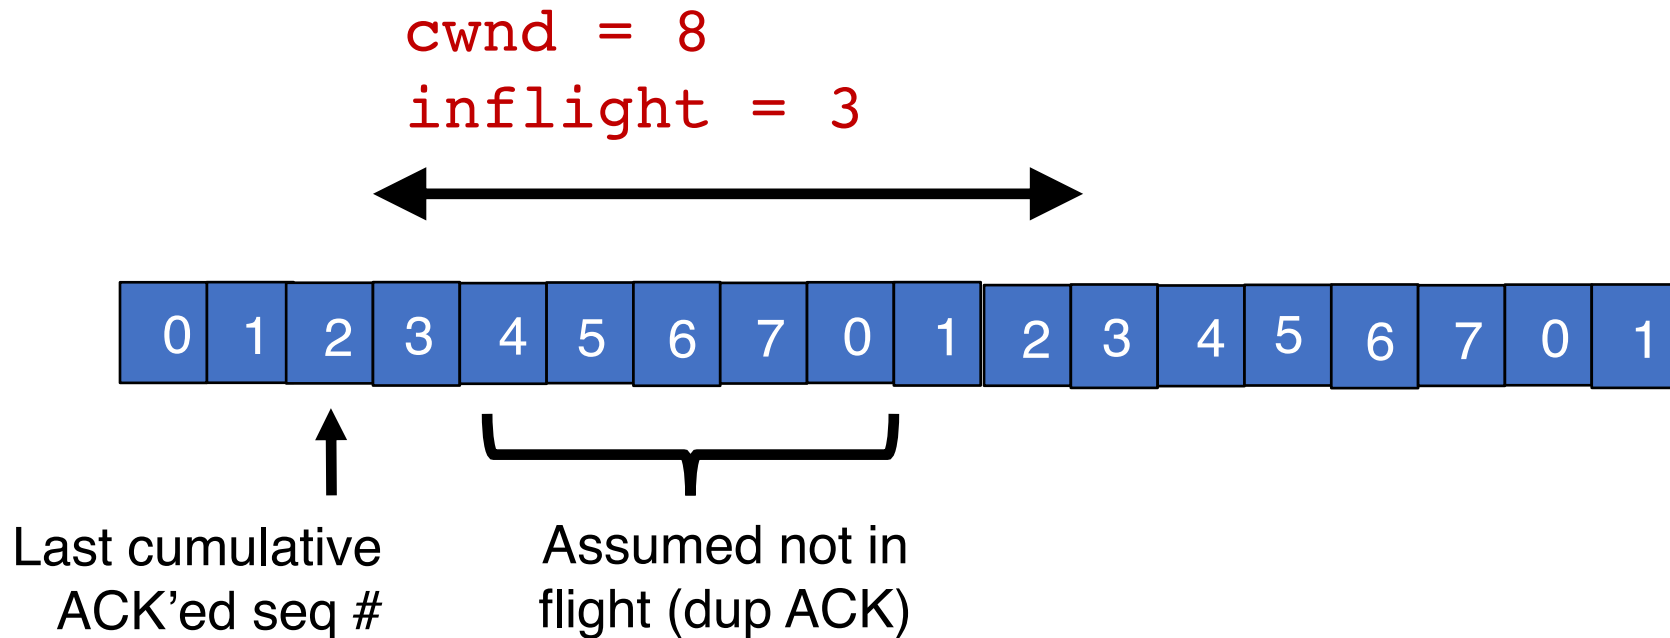
# TCP **fast recovery** (RFC 2581)

- Keep incrementing cwnd by 1 MSS for each dup ACK



# TCP **fast recovery** (RFC 2581)


- Keep incrementing `cwnd` by 1 MSS for each dup ACK





# TCP **fast recovery** (RFC 2581)

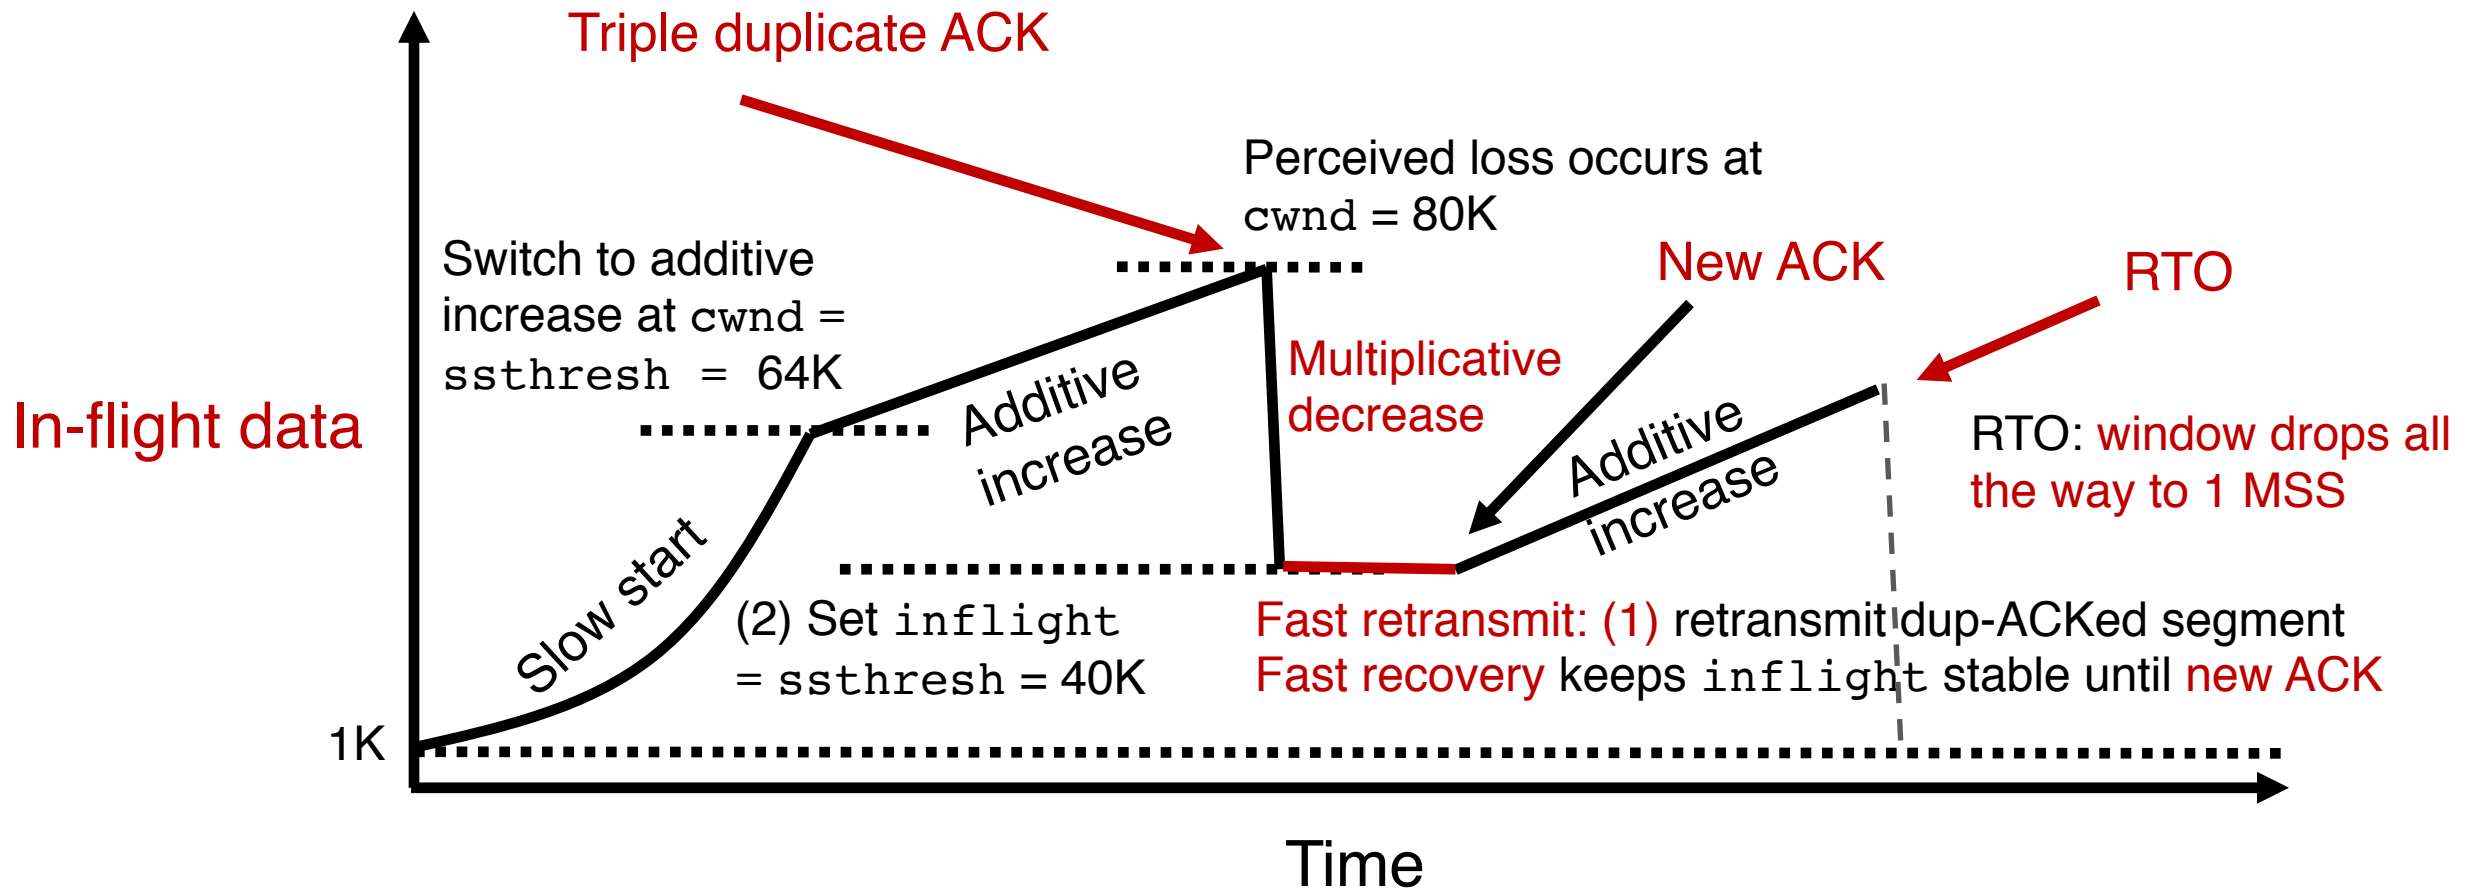
- Eventually a **new ACK** arrives, acknowledging the retransmitted data and all data in between
- Deflate cwnd to half of cwnd before fast retransmit.
  - cwnd and inflight are aligned and equal once again
- Perform **additive increase** from this point!



# Additive Increase/Multiplicative Decrease

Say MSS = 1 KByte

Default ssthresh = 64KB = 64 MSS



TCP New Reno performs additive increase and multiplicative decrease of its congestion window.

In short, we often refer to this as **AIMD**.

**Multiplicative decrease** is a part of all TCP algorithms, including BBR.

[It is necessary for **fairness** across TCP flows.]

# Summary of TCP loss detection

- Don't wait for an RTO and then set the `cwnd` to 1 MSS
  - Tantamount to waiting to get super close to the car in front and then jamming the brakes really hard
- Instead, react proportionately by sensing pkt loss in advance

## Fast Retransmit

- **Triple dup ACK:** sufficiently strong signal that network has dropped data, before RTO
- Immediately retransmit data
- Multiplicatively decrease in-flight data to **half** of its value

## Fast Recovery

- Maintain this reduced amount of in-flight data as long as dup ACKs arrive
  - Data is successfully getting delivered
- When **new ACK** arrives, do **additive increase** from there on



# CS 352

## Computing the Retransmit Timeout

CS 352, Lecture 13.2

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana


# TCP timeout (RTO)

- Useful for reliable delivery and congestion control
- How to pick the RTO value?
  - Too long: slow reaction to loss
  - Too short: premature retransmissions which are wasteful
- Intuition: somehow use the observed RTT (sampleRTT)
  - Can we just directly set the latest RTT as the RTO?
- RTT can vary significantly!
  - Intermittent congestion, path changes, signal quality changes on wireless channel, etc.

# Estimated RTT

- Exponential weighted moving average (typical  $\alpha = 1/8$ )


$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$





# Timeout == estimated RTT + **safety**

- Estimated RTT can have a large **variance**
  - Use a larger safety margin if larger variance



Timeout == estimated RTT + **safety**

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



estimated RTT



safety margin

# Managing a single timer

## data rcvd from app:

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: `TimeoutInterval`


## timeout:

- retransmit segment that caused timeout
- restart timer

## ack rcvd:

- if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - restart timer if there are still unacked segments

# Problem with sampleRTT calculation



# Retransmission ambiguity

- If you retransmitted, how do you measure sampleRTT for it?
  - Measure RTT from original data segment?
  - Measure RTT from most recent (retransmitted) segment?
- There could be an error in RTT estimate, since we can't be sure
- One solution
  - **Never update RTT measurements** based on acknowledgements from retransmitted packets
- Problem: **Sudden change in RTT**, coupled with many retransmissions, can cause system to update RTT very late
  - Ex: Primary path failure leads to a slower secondary path

# Karn's algorithm

- Use back-off as part of `sampleRTT` computation
- Whenever packet loss, RTO is increased by a factor
- Use this increased RTO as RTO estimate for the next segment (**not from EstimatedRTT**)
- Only after an acknowledgment received for a successful transmission is the timer set to new RTT obtained from EstimatedRTT



# CS 352

# TCP Connection Establishment

CS 352, Lecture 13.3

<http://www.cs.rutgers.edu/~sn624/352>


Srinivas Narayana




# Connection Management

before exchanging data, sender/receiver “handshake”:

- agree to establish connection
- agree on connection parameters





```
Socket clientSocket =  
    newSocket("hostname", "port  
number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

# Agreeing to establish a connection

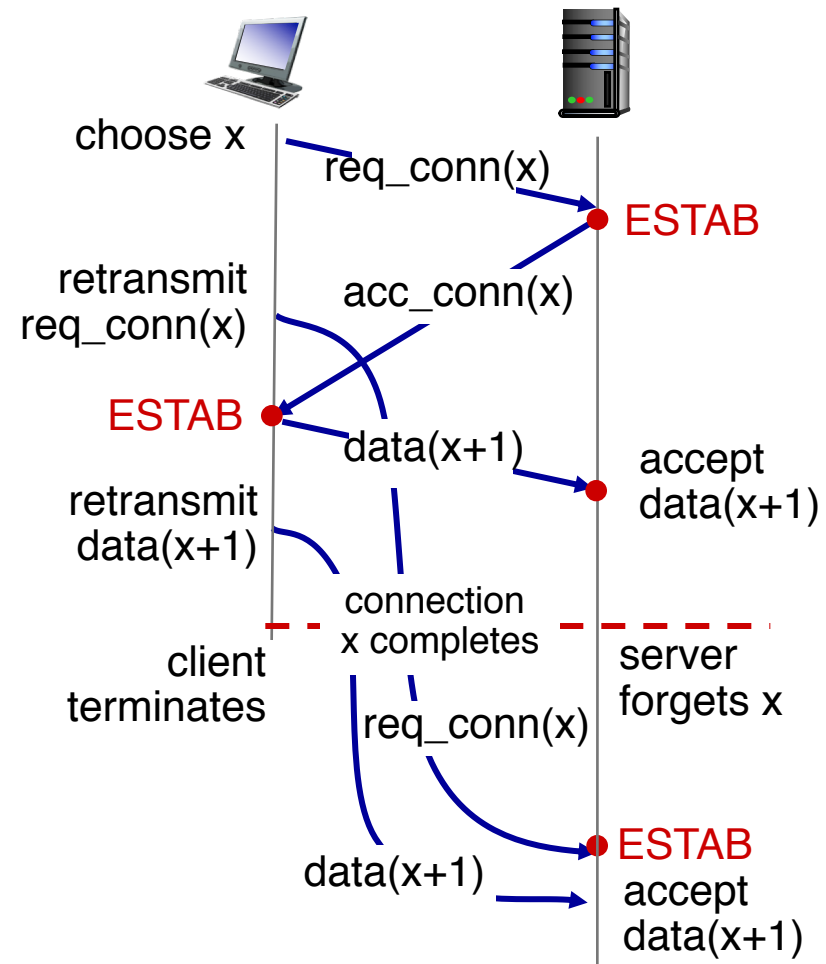
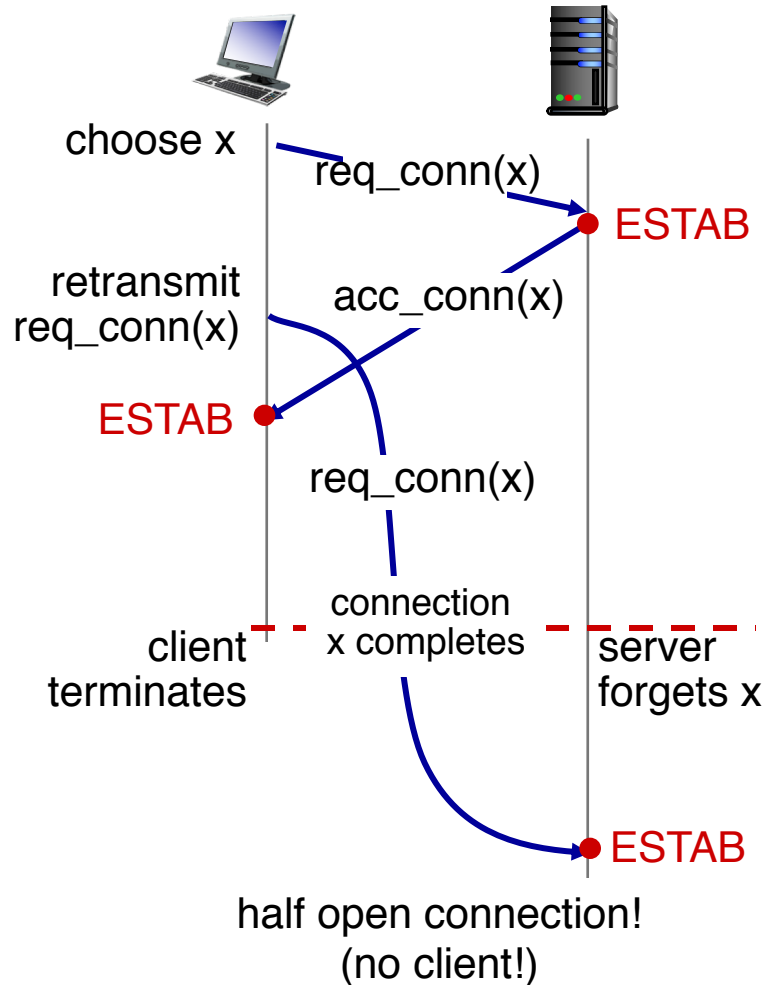
2-way handshake:



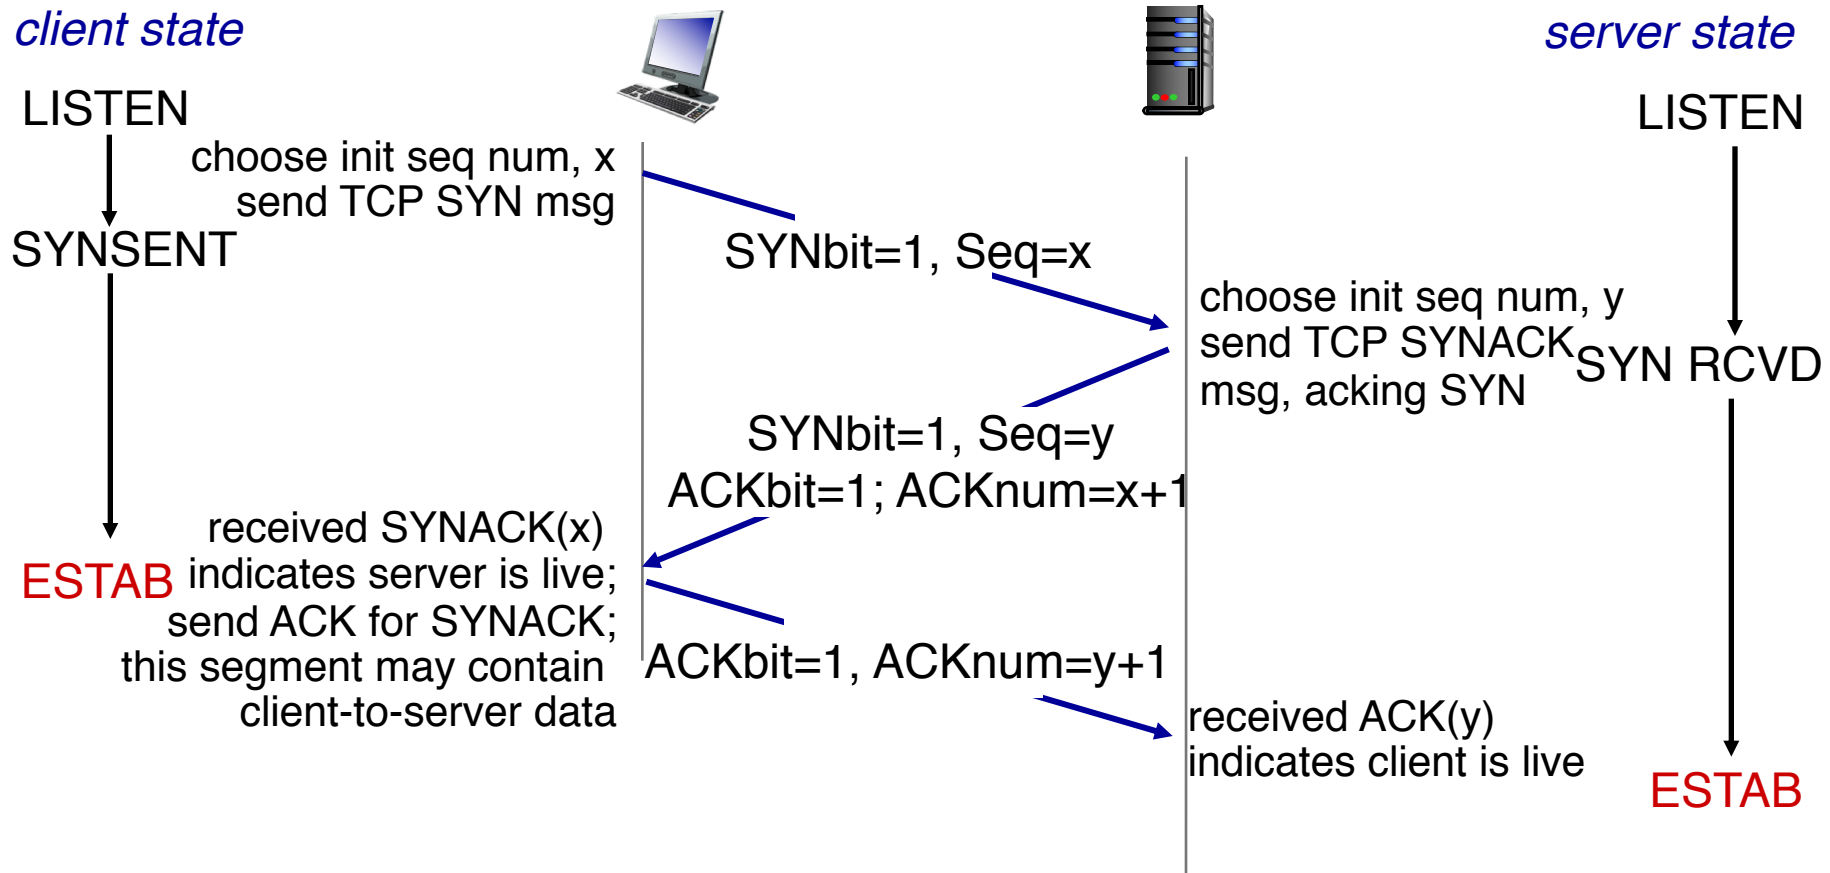
Q: will 2-way handshake always work in network?

- variable delays
- retransmitted messages (e.g. req\_conn(x)) due to message loss
- message reordering
- can't "see" other side

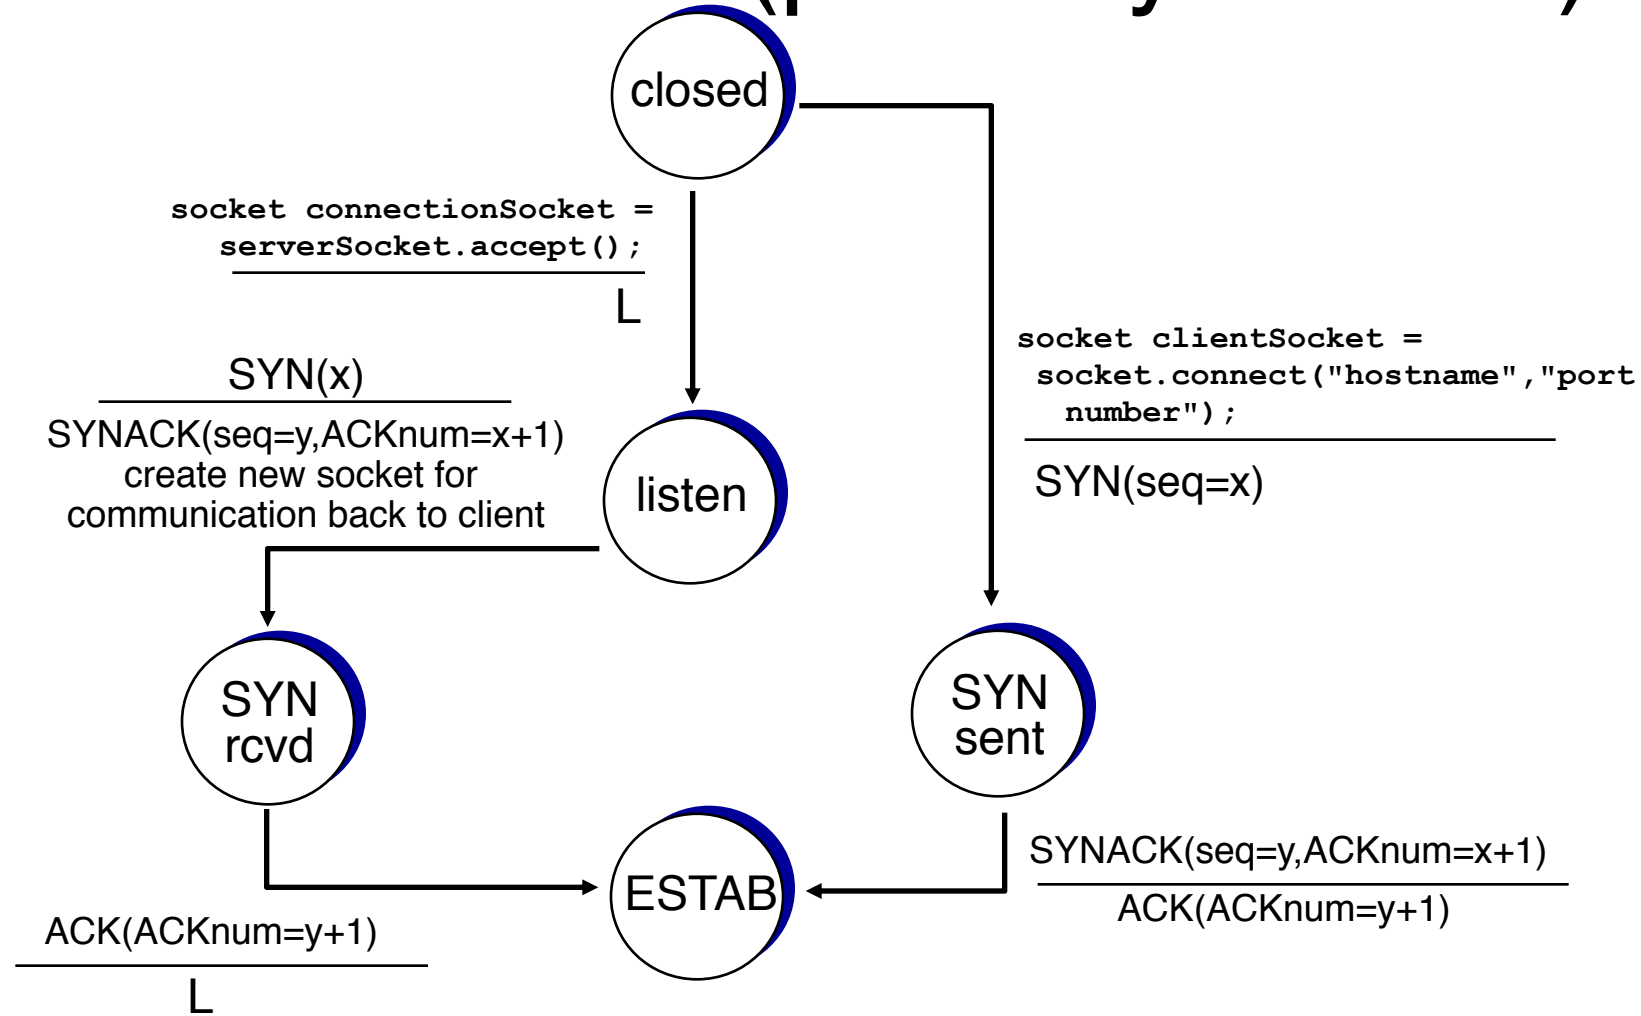
# 2-way handshake failure scenarios



# TCP 3-way handshake



# TCP state machine (partially shown)



# TCP: closing a connection

- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- In general, **TCP is full-duplex**: both sides can send
- But **FIN is unidirectional**: stop one side of the communication
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

# TCP: closing a connection

