

# Reliability (wrap-up); Ordering

Lecture 13


<http://www.cs.rutgers.edu/~sn624/352-S22>

Srinivas Narayana

# Quick recap of concepts




## Stop and Wait



TCP: Connection-oriented

Q1. Which packets are currently in flight?




ACK pkts after a drop?

No / Yes

Go-back-N / Selective repeat

Cumulative ACK / Selective ACK


## Pipelined Reliability



Q2. Which packets were successfully delivered?


Q3. Which packets should the sender retransmit?

# Review: Cumulative ACK



Subtle: Even if there were multiple drops, retransmission after an RTO only includes the first dropped sequence number. Recovering each drop will require one RTO after corresponding packet was transmitted.


# Selective repeat with selective ACK



This slide assumes retransmissions are only triggered by an RTO.  
If other signals were to be used to retransmit earlier (e.g., triple dup ACK -- more on this soon),  
SACK significantly reduces the number of duplicate transmissions compared to cumulative-only ACKs.

# TCP: Cumulative & Selective ACKs

- Sender retransmits the seq #s it thinks aren't received successfully yet
- Pros & cons: selective vs. cumulative ACKs
  - Precision of info available to sender
  - Redundancy of retransmissions
  - Packet header space
  - Complexity (and bugs) in transport software
- On modern Linux, TCP uses selective ACKs by default



TCP reliability metadata

# Metadata on TCP packets for Reliability


- TCP uses metadata in the form of sequence #s and ACK #s
- Where are these stored? Naturally, in the packet header!

# TCP header structure

Source port, destination port (connection demultiplexing)

Size of the TCP header (in 32-bit words)

Basic error detection through checksums (similar to UDP)



Note that one tick mark represents one bit position.




# TCP header structure

Identifies data in the packet from sender's perspective

TCP uses byte seq #s

Identifies the data being ACKed from the receiver's perspective.

TCP uses next seq # that the receiver is expecting.



TCP Header Format

Note that one tick mark represents one bit position.

# Observing a TCP exchange

- `sudo tcpdump -i enol tcp portrange 56000-56010`
- `curl --local-port 56000-56010  
https://www.google.com > output.html`
- Bonus: Try crafting TCP packets with `scapy`!

# Buffering and Ordering in TCP

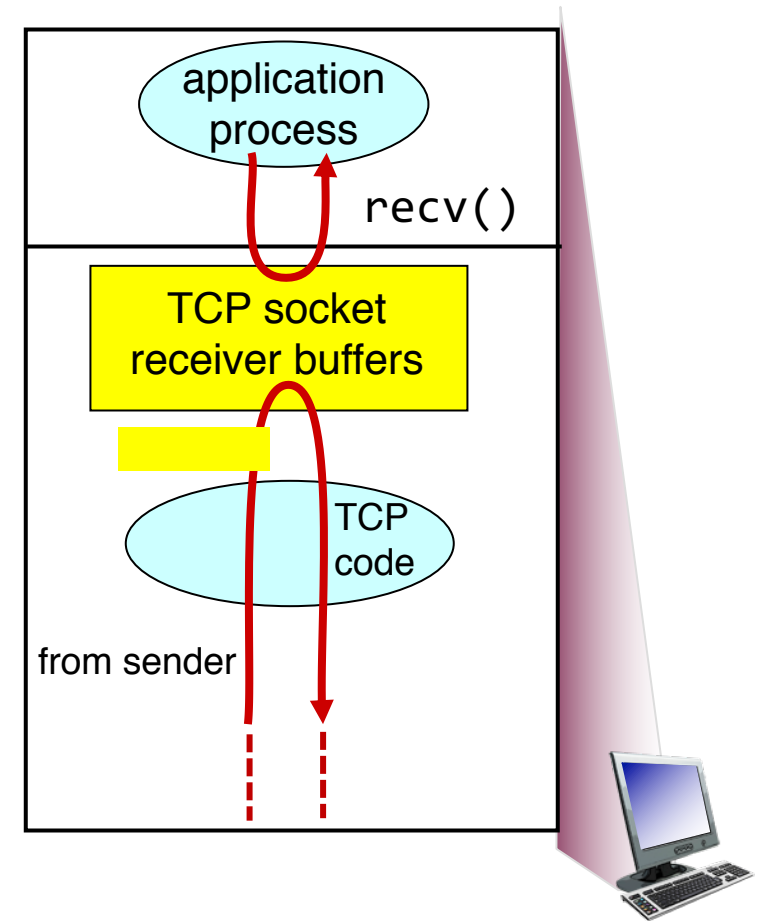
# Memory Buffers at the Transport Layer

# Sockets need receive-side memory buffers

- Since TCP uses selective repeat, the receiver must **buffer** data that is received after loss:
  - e.g., hold packets so that only the “holes” (due to loss) need to be filled in later, without having to retransmit packets that were received successfully
- Apps read from the receive-side socket buffer when you do a `recv()` call.
- Even if data is always reliably received, applications may not always read the data immediately
  - What if you invoked `recv()` in your program infrequently (or never)?
  - For the same reason, UDP sockets also have receive-side buffers

# Receiver app's interaction with TCP

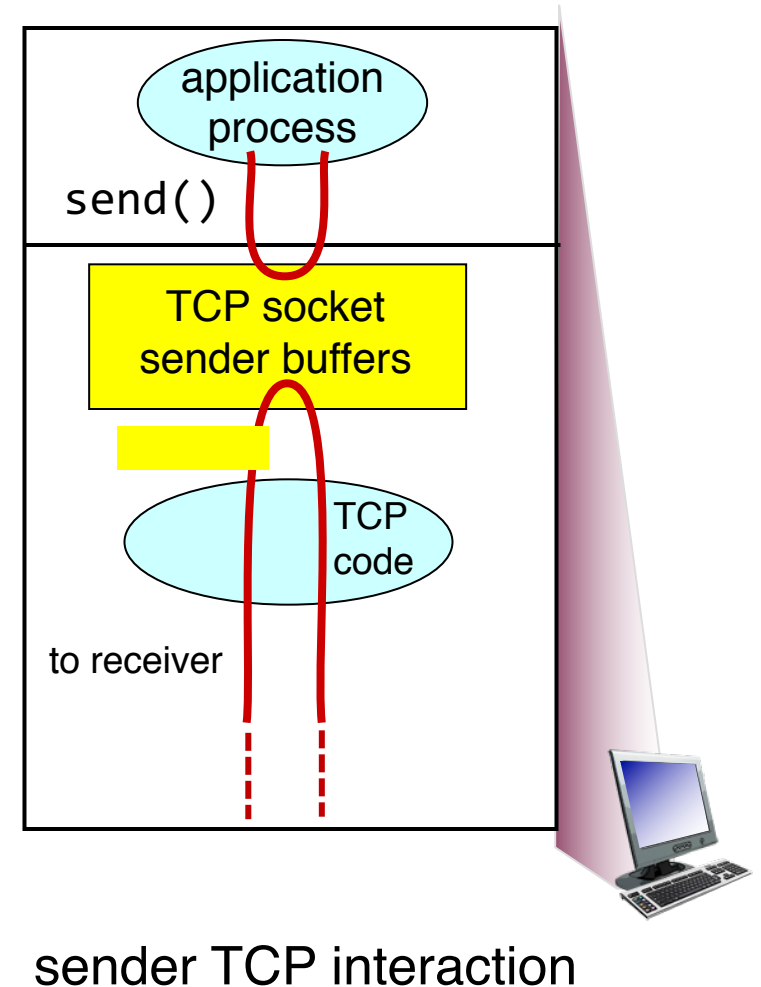
- Upon reception of data, the receiver's TCP stack deposits the data in the receive-side socket buffer
- An app with a TCP socket reads from the TCP receive socket buffer
  - e.g., when you do `data = sock.recv()`



receiver TCP interaction

# Sockets need send-side memory buffers

- The possibility of **packet retransmission** in the future means that data can't be immediately discarded from the sender once transmitted.
- App has issued `send( )` and moved on; TCP stack must buffer this data
- Transport layer must wait for ACK of a piece of data before reclaiming (freeing) the memory for that data.




Ordered Delivery




# Reordering packets at the receiver side

- Let's suppose receiver gets packets 1, 2, and 4, but not 3 (dropped)
- Suppose you're trying to download a document containing a report
- What would happen if transport at the receiver directly presents packets 1, 2, and 4 to the application (i.e., receiving 1,2,4 through the `recv( )` call)?



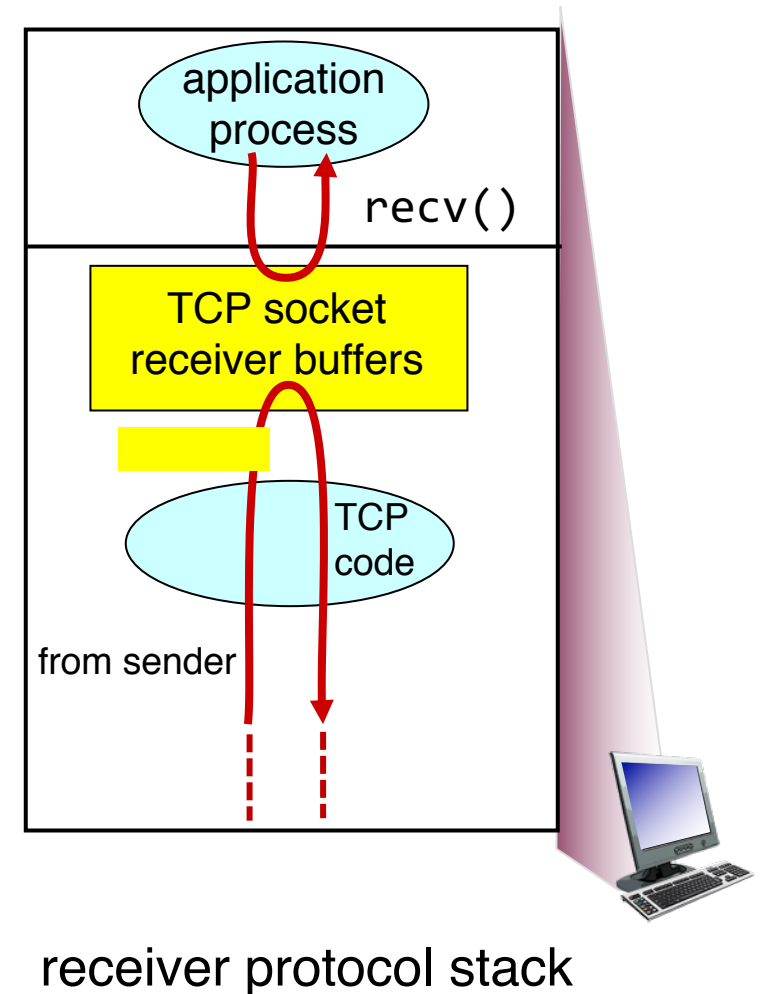
# Reordering packets at the receiver side

- Reordering can happen for a few reasons:
  - Drops
  - Packets taking different paths through a network
- Receiver needs a general strategy to ensure that data is presented to the application **in the same order that the sender pushed it**
- To implement ordered delivery, the receiver uses
  - Sequence numbers
  - Receiver socket buffer
- We've already seen the use of these for reliability; but they can be used to order too!




# Receive-side app and TCP

- TCP receiver software only releases the data from the receive-side socket buffer to the application if:
  - the data is **in order** relative to all other data already read by the application
- This process is called **TCP reassembly**



# TCP Reassembly

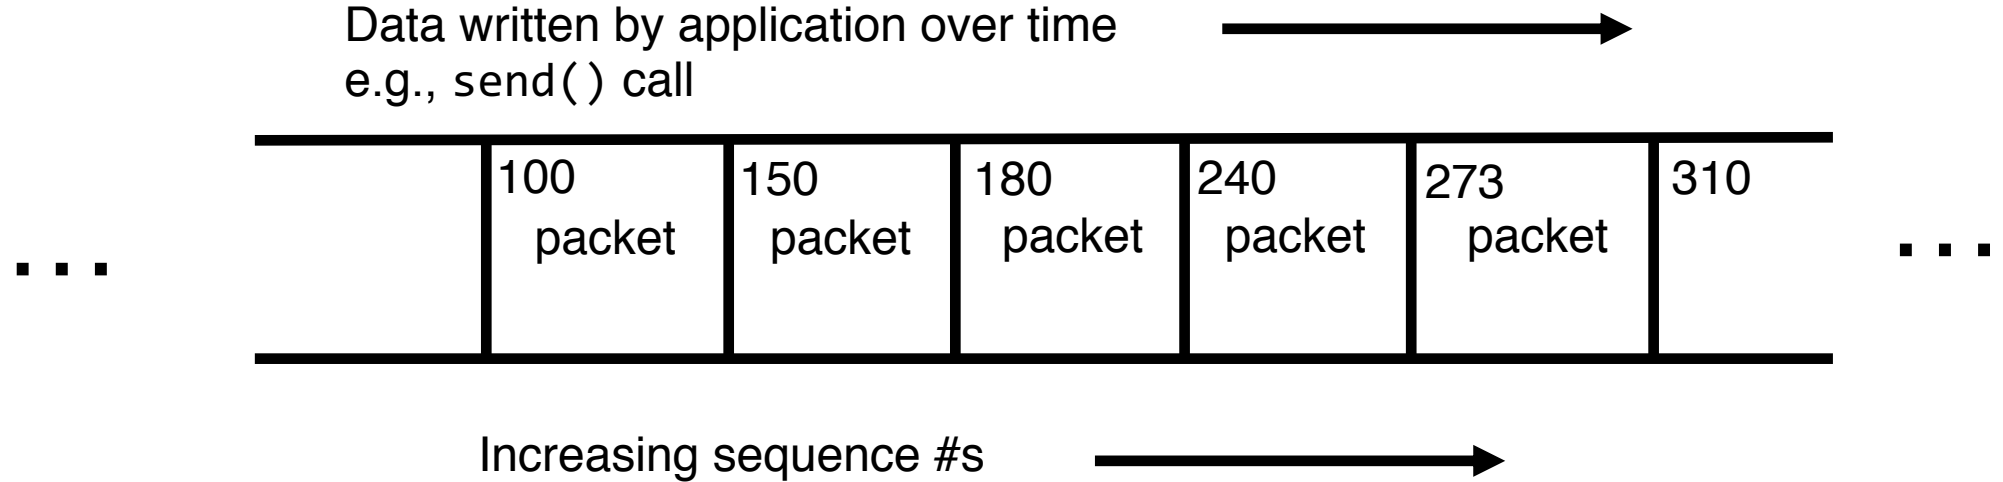


# Implications of ordered delivery

- Packets cannot be delivered to the application if there is an **in-order packet missing** from the receiver's buffer
  - The receiver can only buffer so much out-of-order data
  - **Subsequent out-of-order packets dropped**
  - It won't matter that those packets successfully arrive at the receiver from the sender over the network
- **TCP application-level throughput will suffer** if there is too much packet reordering in the network
  - Data may have reached the receiver, but won't be delivered to apps upon a `recv()` (...or may not even be buffered!)

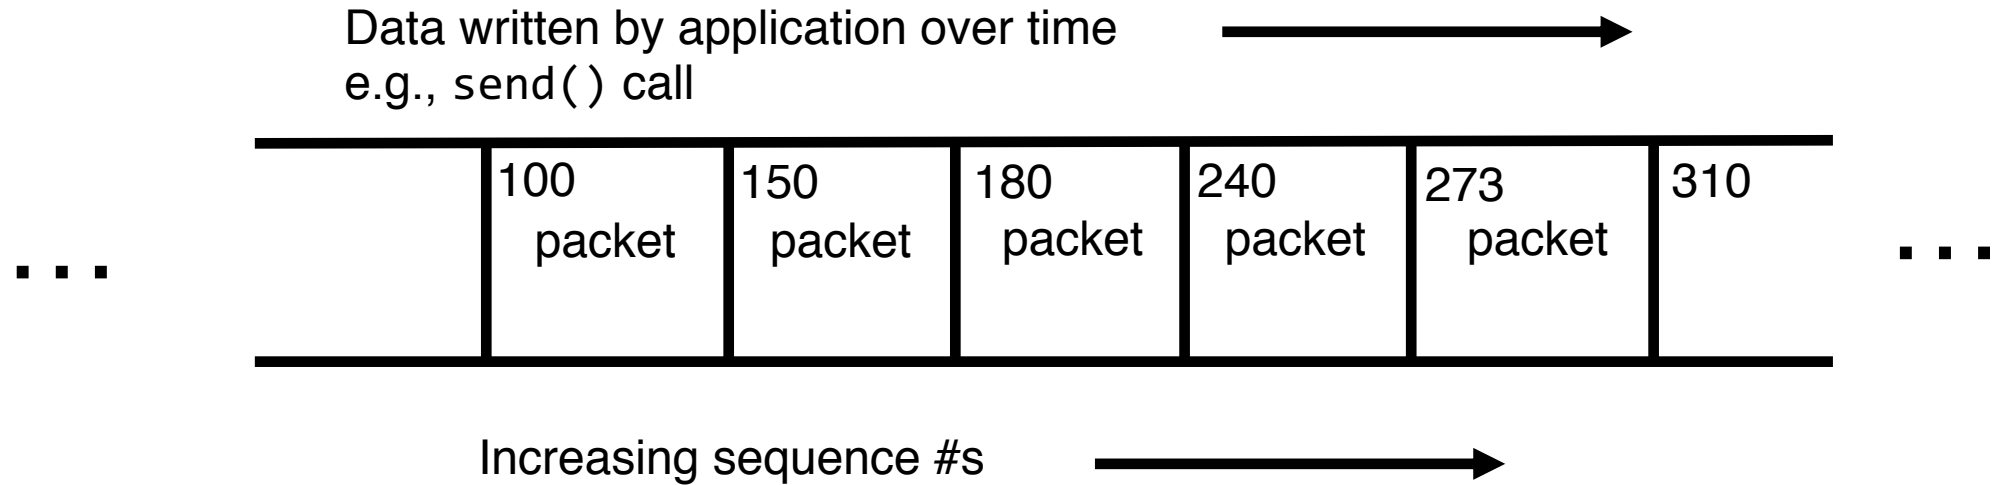
# Stream-Oriented Data Transfer

# Sequence numbers in the app's **stream**



TCP uses byte sequence numbers

# Sequence numbers in the app's **stream**



Packet boundaries aren't important for TCP software

TCP is a **stream-oriented** protocol

(We use `SOCK_STREAM` when creating sockets)



# Sequence numbers in the app's **stream**

