

Flow Control

Lecture 14

<http://www.cs.rutgers.edu/~sn624/352-S22>

Srinivas Narayana

Quick recap of concepts

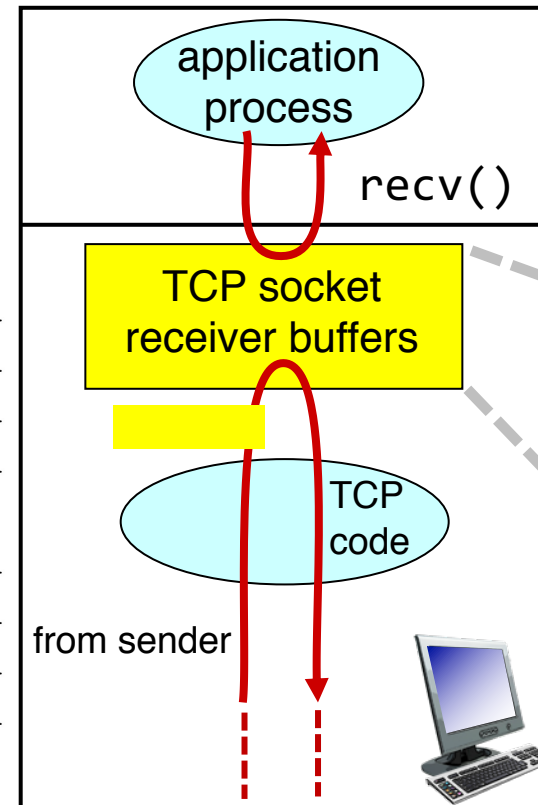


TCP:
Connection-oriented

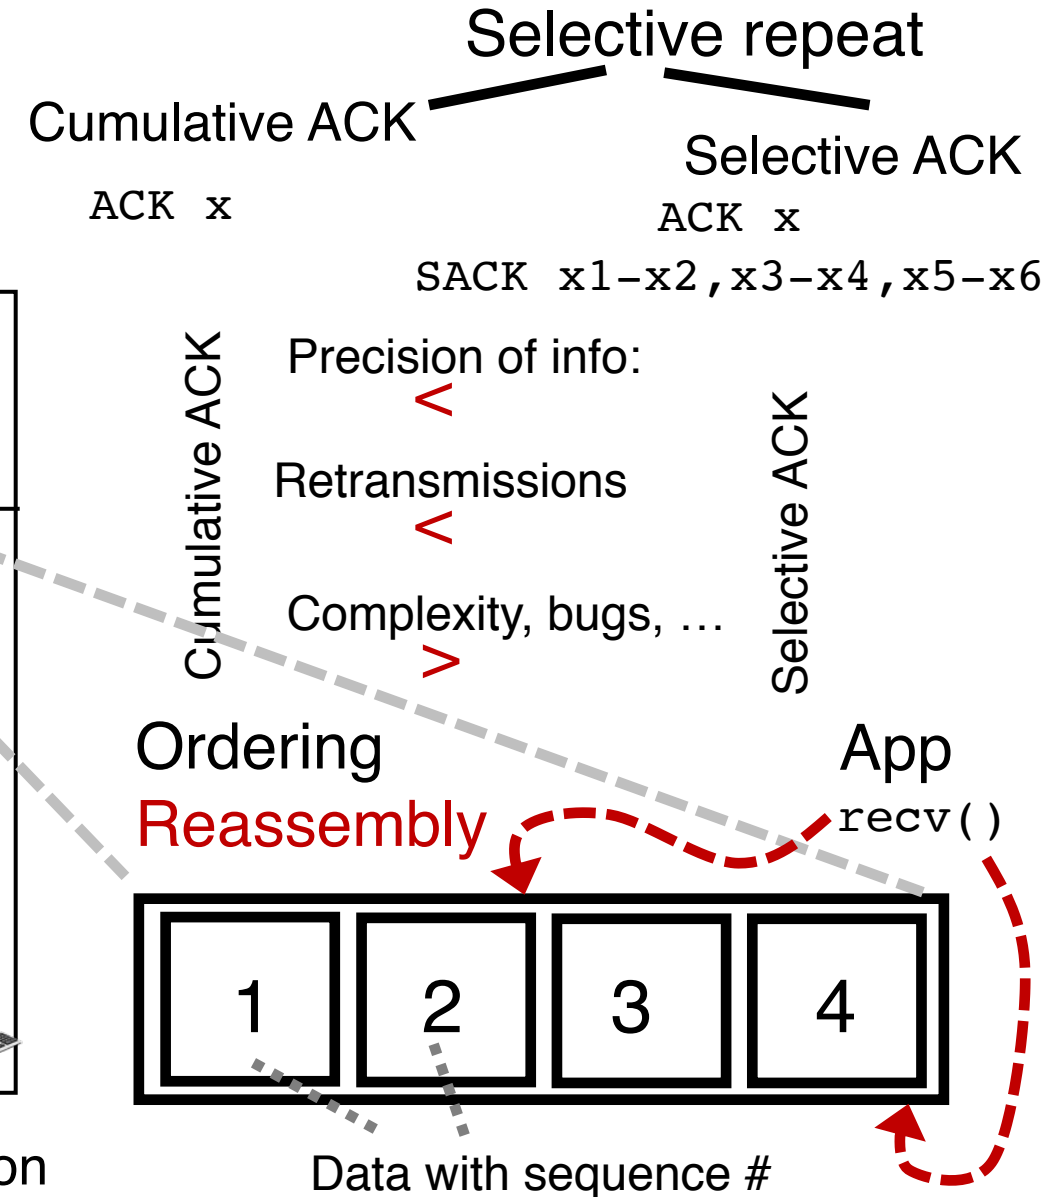
0										1										2										3											
0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1
Source Port										Destination Port																															
Sequence Number																																									
Acknowledgment Number																																									
Data Offset					Reserved					R	C	S	S	Y	I	Window																									
										G	K	H	T	N	N																										
Checksum										Urgent Pointer																															
Options										Padding																															
data																																									

TCP Header Format

Note that one tick mark represents one bit position.



receiver TCP interaction




Implications of ordered delivery

- Packets cannot be delivered to the application if there is an **in-order packet missing** from the receiver's buffer
 - The receiver can only buffer so much out-of-order data
 - **Subsequent out-of-order packets dropped**
 - It won't matter that those packets successfully arrive at the receiver from the sender over the network
- **TCP application-level throughput will suffer** if there is too much packet reordering in the network
 - Data may have reached the receiver but won't be delivered to apps upon a `recv()` (. . . or may not even be buffered!)


Stream-Oriented Data Transfer

Sequence numbers in the app's **stream**



TCP uses byte sequence numbers

Sequence numbers in the app's **stream**

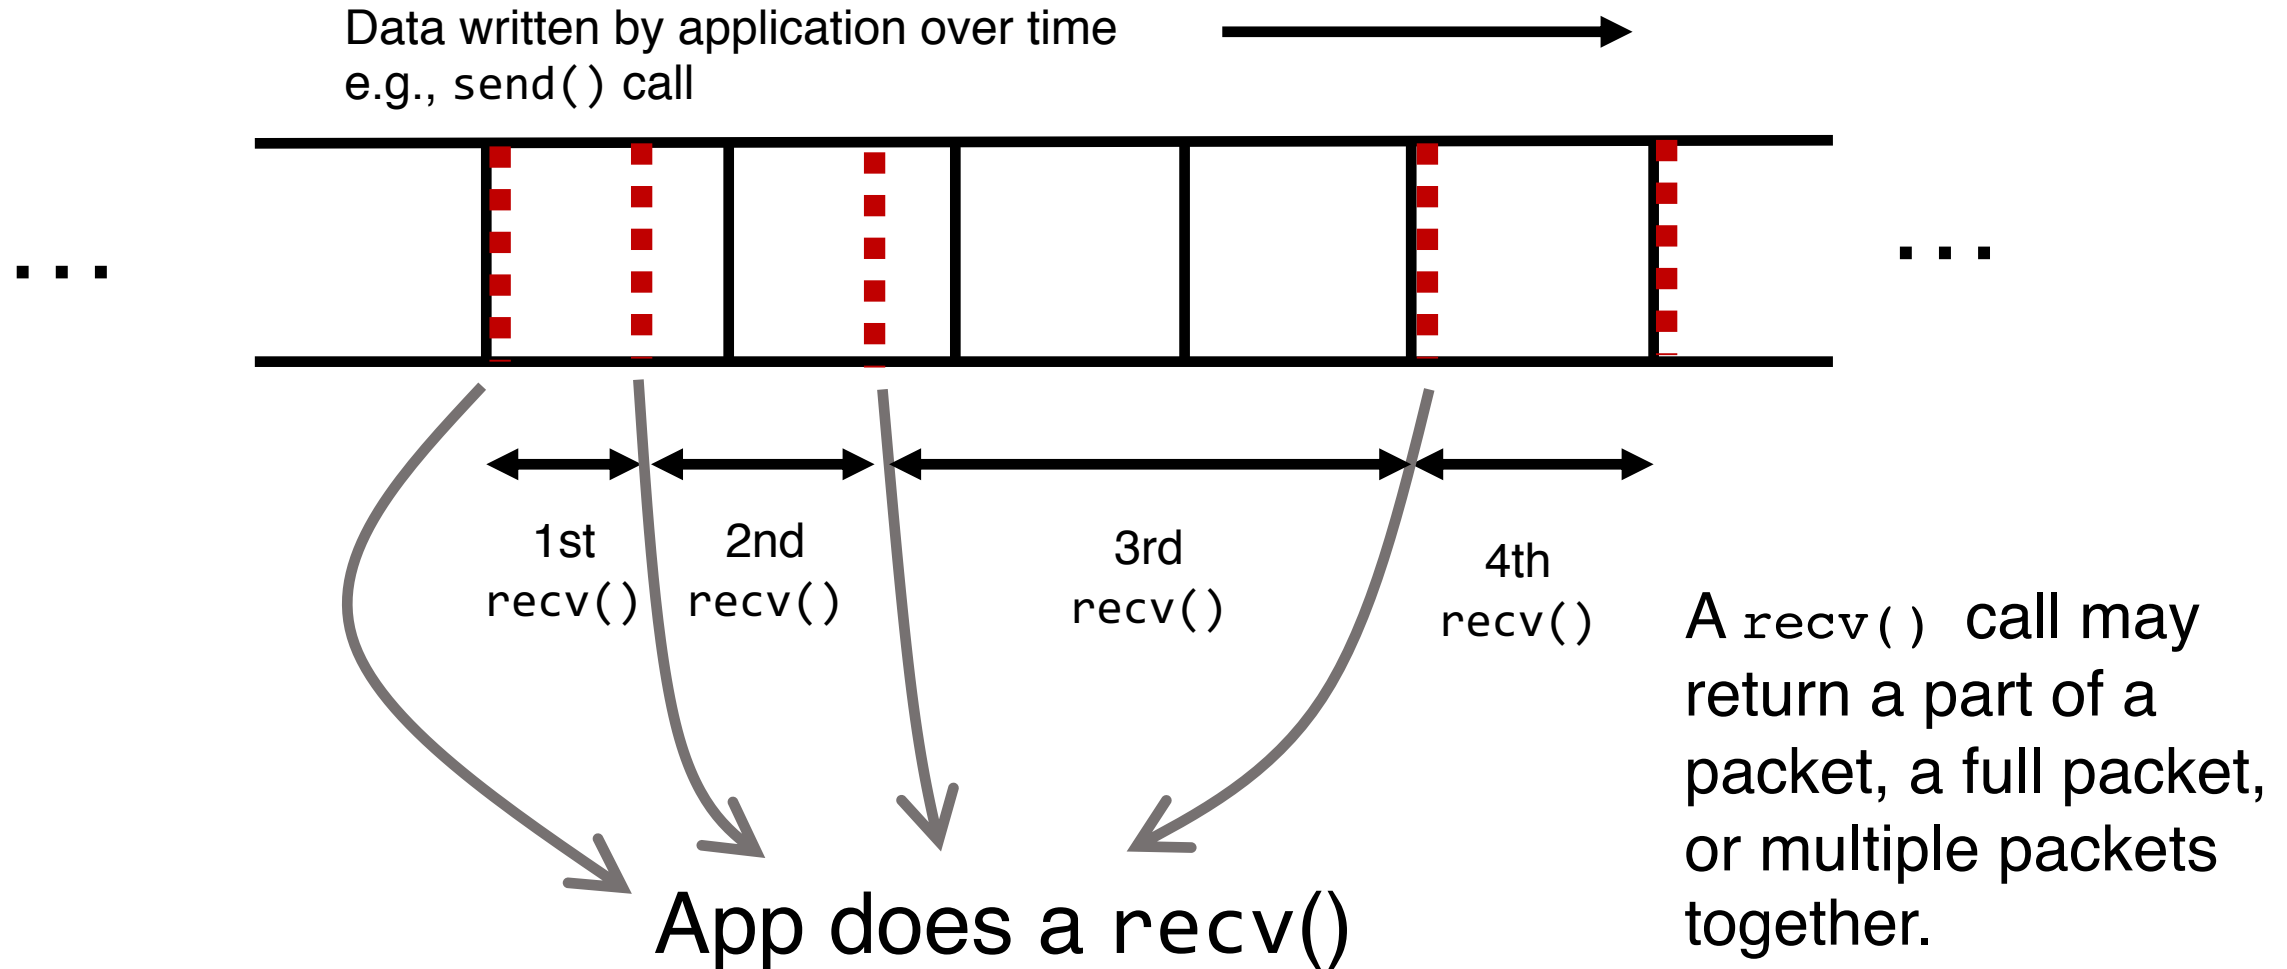


Packet boundaries aren't important for TCP software

TCP is a **stream-oriented** protocol

(We use `SOCK_STREAM` when creating sockets)

Sequence numbers in the app's **stream**





= window size

Proportional to **throughput**

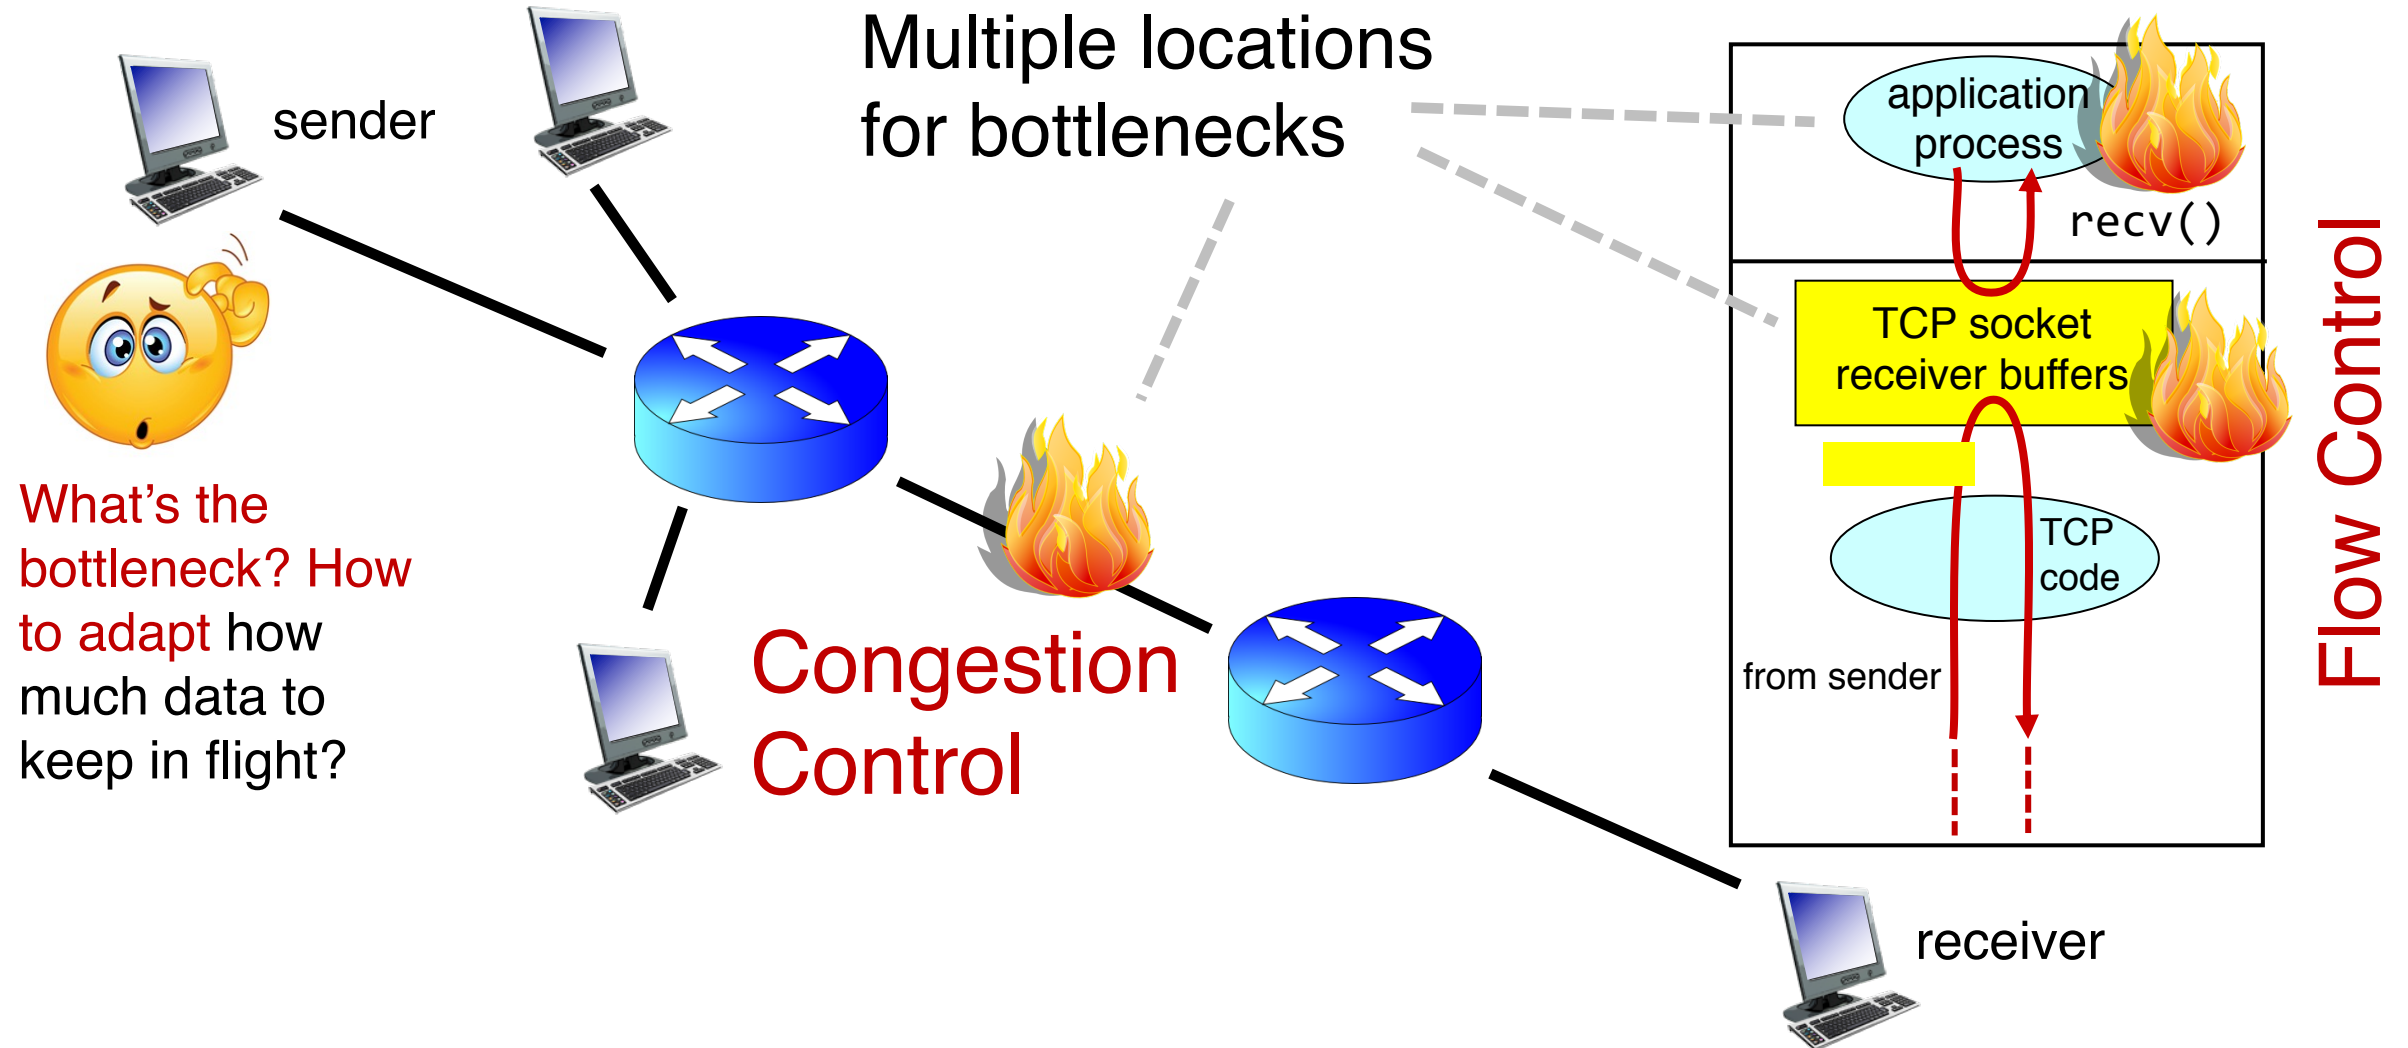
How much data to keep in flight?

Stop and Wait



Pipelined Reliability

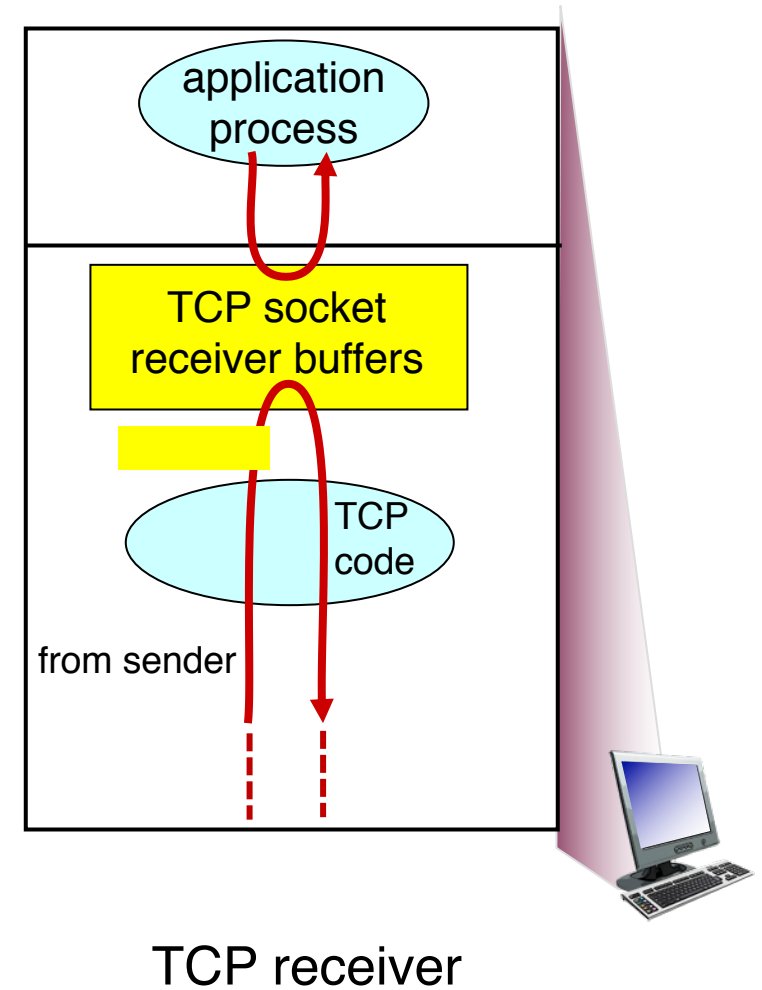
We want to increase throughput, but ...



Flow Control

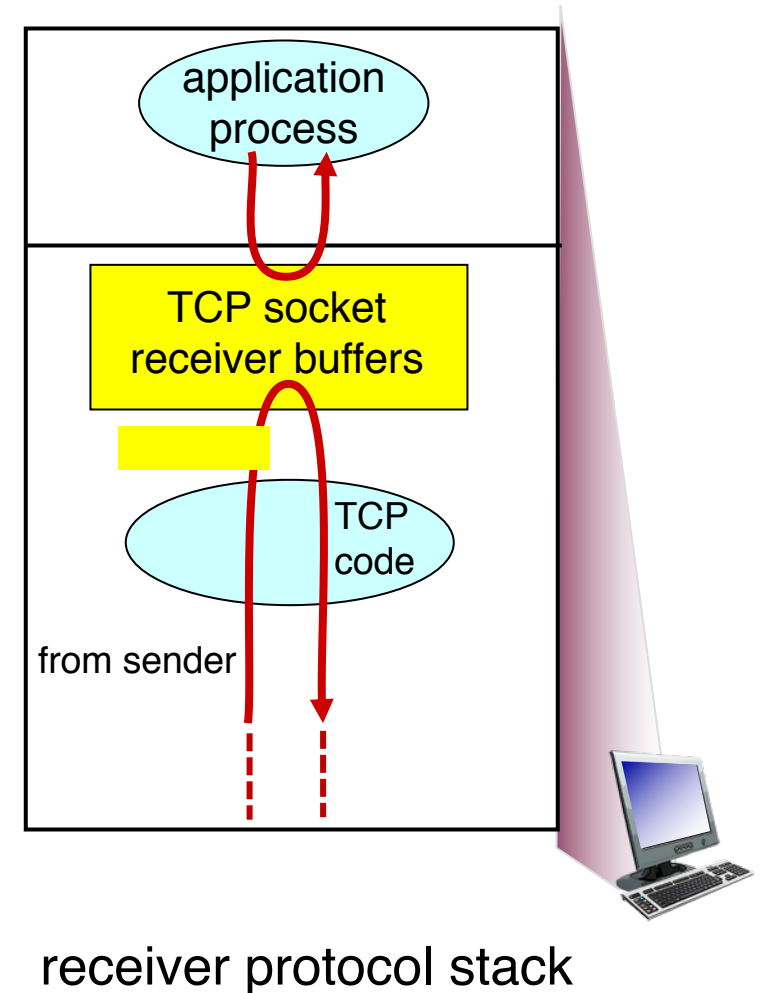
Socket buffers can become full

- Applications may read data slower than the sender is pushing data in
 - Example: what if an app infrequently or never calls `recv()`?
- There may be too much reordering or packet loss in the network
 - What if the first few bytes of a window are lost or delayed?
- Receivers can only buffer so much before dropping subsequent data

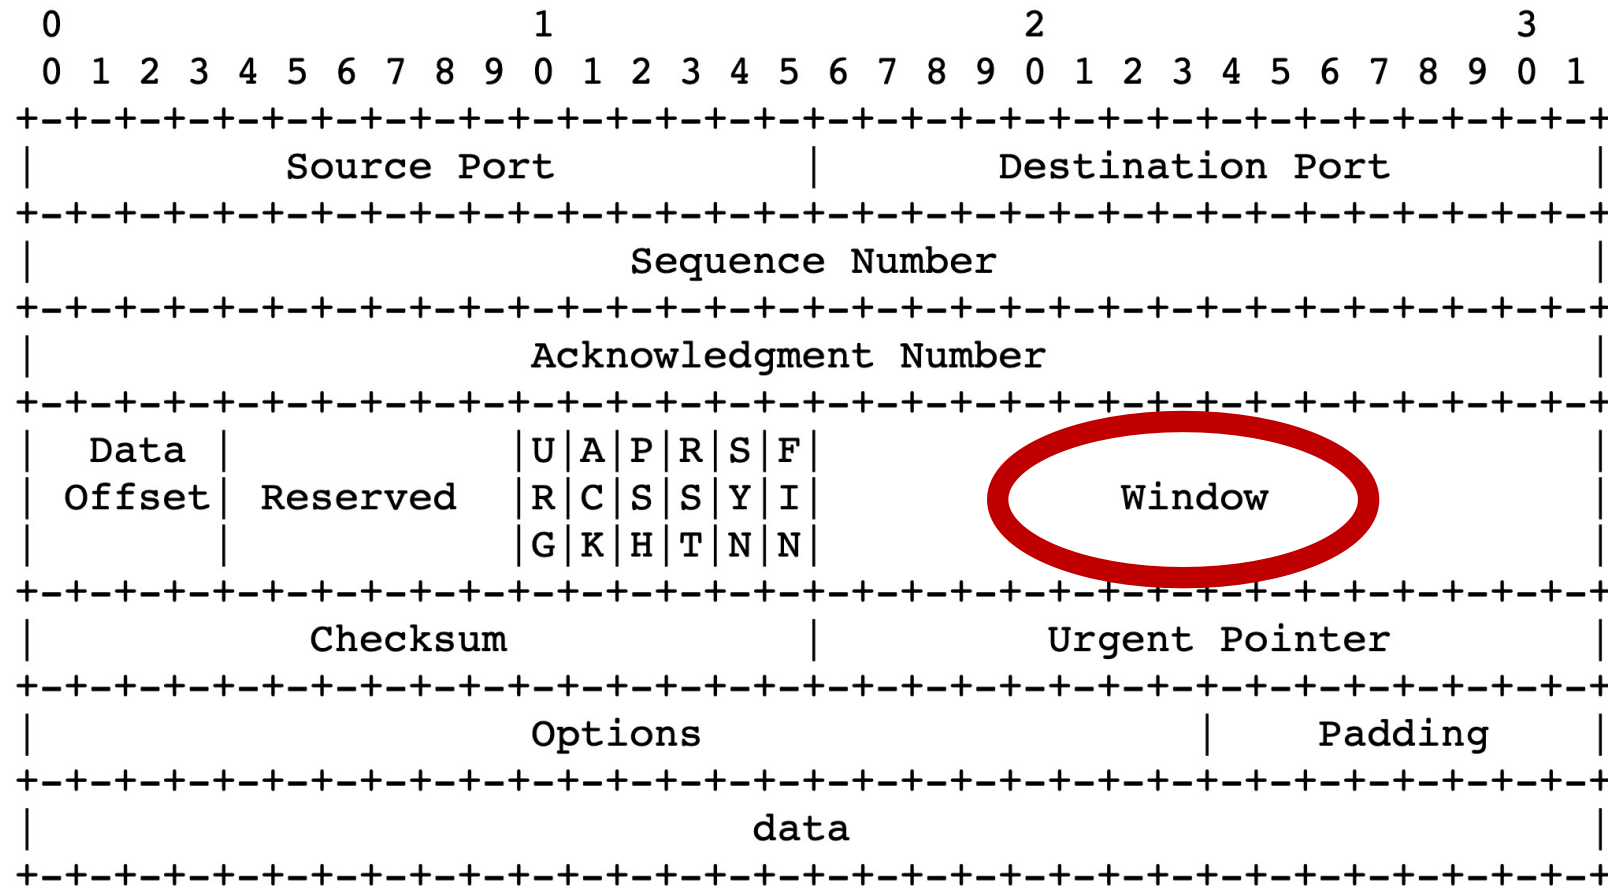


Goal: avoid drops due to buffer fill

- Have a TCP sender only send as much as the **free buffer space** available at the receiver.
- *Amount of free buffer varies over time!*
- TCP implements **flow control**
- Receiver's ACK contains the amount of data the sender can transmit without running out the receiver's socket buffer
- This number is called the **advertised window size**



Flow control in TCP headers




TCP Header Format

Note that one tick mark represents one bit position.


TCP flow control

- Receiver **advertises** to sender (in the ACK) how much free buffer is available





TCP Header Format

Note that one tick mark represents one bit position.





TCP flow control

- Subsequently, the sender's sliding window cannot be larger than this value
- Restriction on new sequence numbers that can be transmitted
- == restriction on sending rate!





TCP flow control

- If receiver app is too slow reading data:
 - receiver socket buffer fills up
 - So, advertised window shrinks
 - So, sender's window shrinks
 - So, sender's sending rate reduces



TCP flow control

Flow control matches the sender's write speed to the receiver's read speed.



Sizing the receiver's socket buffer

- Operating systems have a default receiver socket buffer size
 - Listed among `sysctl -a | grep net.inet.tcp` on MAC
 - Listed among `sysctl -a | grep net.ipv4.tcp` on Linux
- If socket buffer is too small, sender can't keep too many packets in flight → lower throughput
- If socket buffer is too large, too much memory consumed per socket
- How big should the receiver socket buffer be?

Sizing the receiver's socket buffer

- Case 1: **Suppose the receiving app is reading data too slowly:**
 - no amount of receiver buffer can prevent low sender throughput if the connection is long-lived!

Sizing the receiver's socket buffer

- Case 2: Suppose the receiving app reads sufficiently fast *on average* to match the sender's writing speed.
 - Assume the sender has a window of size W .
 - The receiver must use a buffer of size at least W . Why?
- Captures two cases:
- (1) When the first sequence #s in the window are dropped
 - *Selective repeat*: data in window buffered until the ACKs of delivered data (within window) reach sender. Adv. win reduces sender's window
- (2) When the sender sends a burst of data of size W
 - Receiver may not match the *instantaneous* rate of the sender

Summary of flow control

- Keep memory buffers available at the receiver whenever the sender transmits data
 - Inform the sender on an on-going basis (each ACK)
 - Function #1: match sender speed to receiver speed
 - Function #2: reassemble data in order and hold for selective repeat
-
- Correct socket buffer sizing is important for TCP throughput

Info on (tuning) TCP stack parameters

- https://www.ibm.com/support/knowledgecenter/linuxonibm/liaag/wkvm/wkvm_c_tune_tcpip.htm
- <https://cloud.google.com/solutions/tcp-optimization-for-network-performance-in-gcp-and-hybrid>