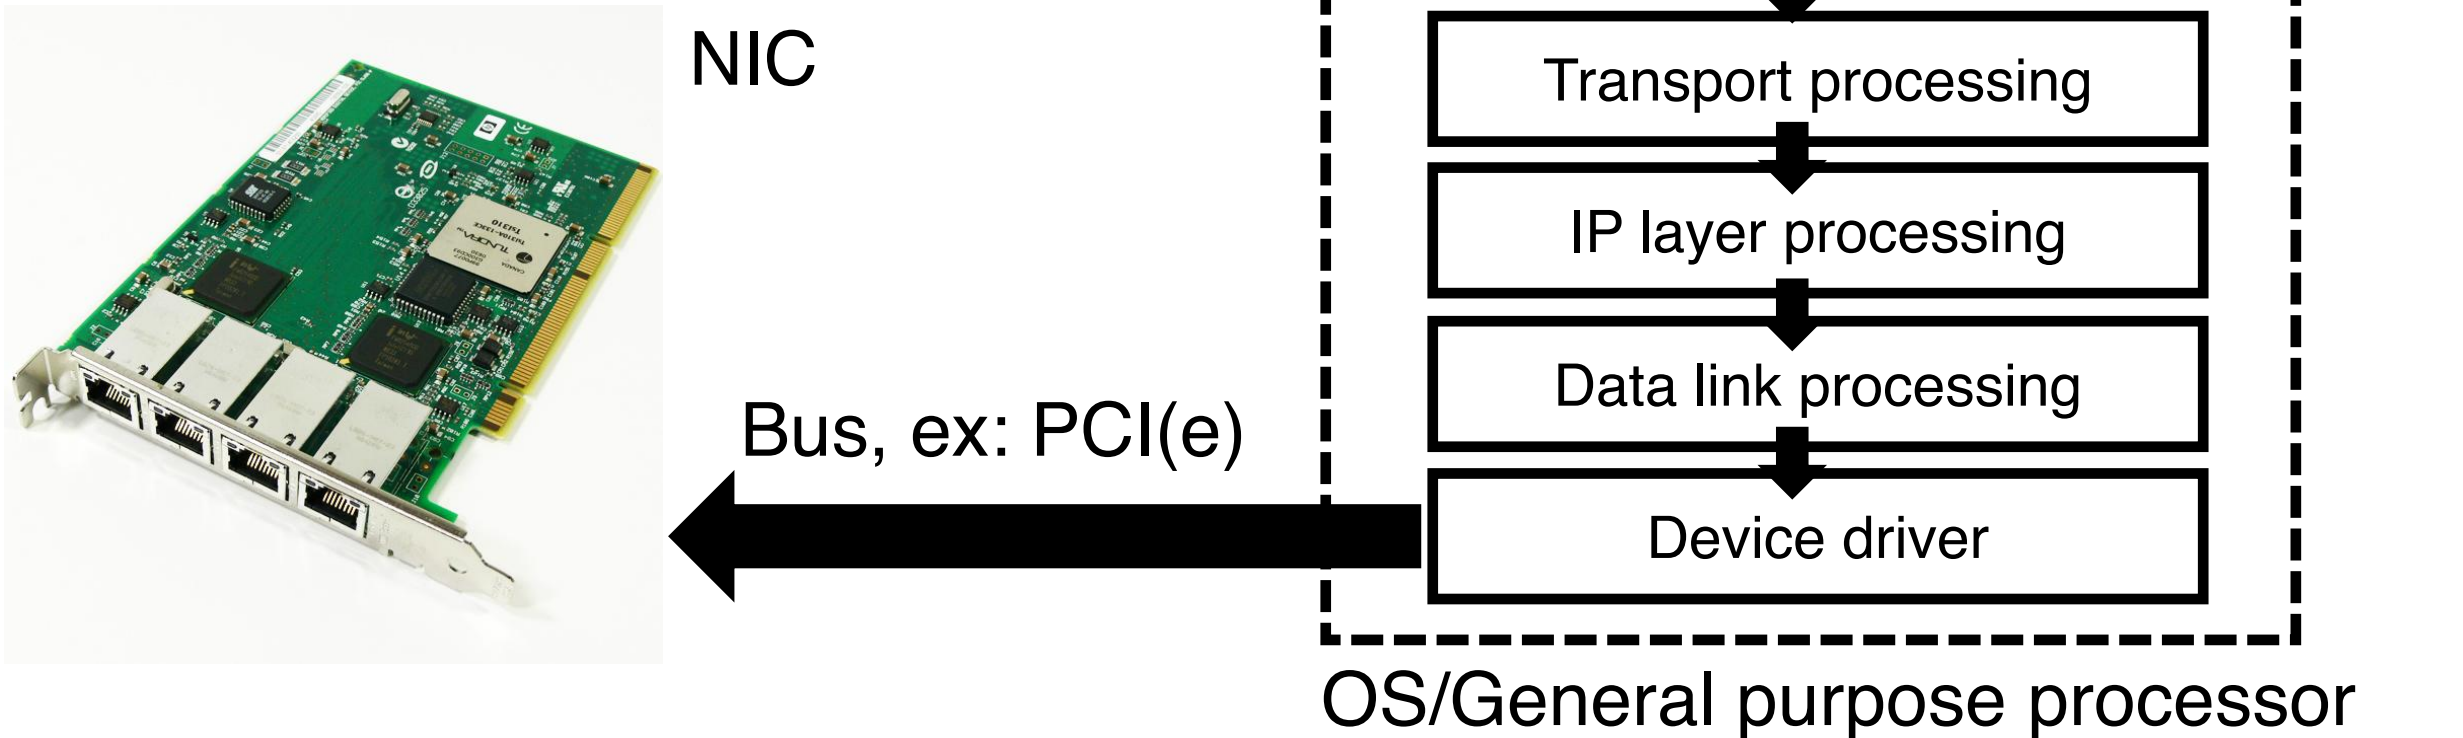


Programmable NICs

Lecture 14, Computer Networks (198:552)

Network Interface Cards (NICs)

- The physical interface between a machine and the wire
- Life of a transmitted packet



Data transfers between the host & NIC

- Receive (*rx*) path
 - Polling by the CPU
 - Interrupt the CPU whenever a packet arrives
- Transmit (*tx*) path
 - Programmed input/output: CPU explicitly moves data
 - Direct Memory Access (DMA)
- Multiple queues to transmit and receive data
 - Allow scheduling/prioritization, isolation, and load balancing

NIC speeds have been going up...

- 2009: 1 Gbit/s
- 2012: 10 Gbit/s
- 2015: 40 Gbit/s
- 2018: 100 Gbit/s
- Soon: 200 and 400 Gbit/s standards

But OS stack has significant overheads!

- Measurements from [netmap'12]
 - FreeBSD, i7 (2.93 GHz), ixgbe driver
- Sources of overhead:
 - System call
 - Memory allocation
 - IP routing/header lookup
 - (surprisingly) MAC lookups
 - Moving data from device memory to NIC memory
- Affects throughput, delay, CPU util

File	Function/description	time ns	delta ns
user program	sendto system call	8	96
uipc_syscalls.c	sys_sendto	104	137
uipc_syscalls.c	sendit	111	
uipc_syscalls.c	kern_sendit	118	
uipc_socket.c	sosend	—	
uipc_socket.c	sosend_dgram sockbuf locking, mbuf allocation, copyin	146	
udp_usrreq.c	udp_send	273	57
udp_usrreq.c	udp_output	273	
ip_output.c	ip_output route lookup, ip header setup	330	198
if_ethersubr.c	ether_output MAC header lookup and copy, loopback	528	162
if_ethersubr.c	ether_output_frame	690	
ixgbe.c	ixgbe_mq_start	698	
ixgbe.c	ixgbe_mq_start_locked	720	
ixgbe.c	ixgbe_xmit mbuf mangling, device programming	730	220
—	on wire	950	

CPU processing won't get much better

- The end of “Dennard scaling” and “Moore’s law”
- Transistor power dissipation isn’t going down with generations
 - Can’t clock x86 cores with higher frequencies
 - Implication: performance for serialized computation has stagnated
- Transistor feature sizes are approaching quantum limits
 - Number of transistors may increase in the near future
 - But it may stagnate pretty soon
 - Implication: there is only so much parallelism to be extracted

Approaches to keep up with high speed

- 1) Optimize the common case (in the existing network stack)
 - 2) Harden common stateless operations in the NIC ASIC
 - 3) Use parallelism by processing packets with multiple cores
 - 4) Bypass the OS kernel entirely
-

- 5) Bypass the hypervisor entirely (in virtualized environments)
- 6) Offload packet processing to some high-perf substrate
 - ASIC, SoC, FPGAs

(1) Optimize the common case

- Fast and slow path separation
 - Ex: TCP ACK processing (e.g., 30 instructions by Van Jacobson)
 - Ex: IP route caching
- Batching:
 - NIC: Interrupt coalescing
 - NIC: Receive segment coalescing
 - Stack: Transmit packets in a bunch until no more to transmit

(2) Harden stateless operations in NIC

- Move some of the fixed per-packet *independent* processing to NIC
 - Historically the main reason why software kept up with link speeds
- Examples:
 - TCP and IP checksum computations
 - “Scatter gather” I/O
 - Segmentation offload (tx and rx)
 - ... and many more
- Often built first in software and moved to NIC hardware

(3) Process packets with multiple cores


- Receive-Side Scaling (RSS): on packet receive,
 - Hash the packet into one of N buckets
 - Each bucket mapped to a core on the host
 - Interrupt the corresponding core
- May use separate queue for each core
- May separate the packets by connection
 - Single transport connection is the granularity of maintaining state
- Analogous: Transmit-side steering

(4) Bypass the OS kernel altogether

- Use a custom device driver that polls the network device
 - Give up a core or two simply to poll the NIC!
- Share memory directly between driver and user-space
 - Flexible way to implement many different features in user-space
- Intel's Data Plane Development Kit (DPDK)
 - Many high performance “network functions”
- Cons: breaks OS abstractions and isolation between processes


Virtualized cluster environments

- Clusters with multiple *tenants*
- Hypervisor mediates net access
- Additional sources of overhead
 - Copy packet into VM memory
 - Simulate an interrupt to the VM
- VM continues processing packet with its own OS/network stack



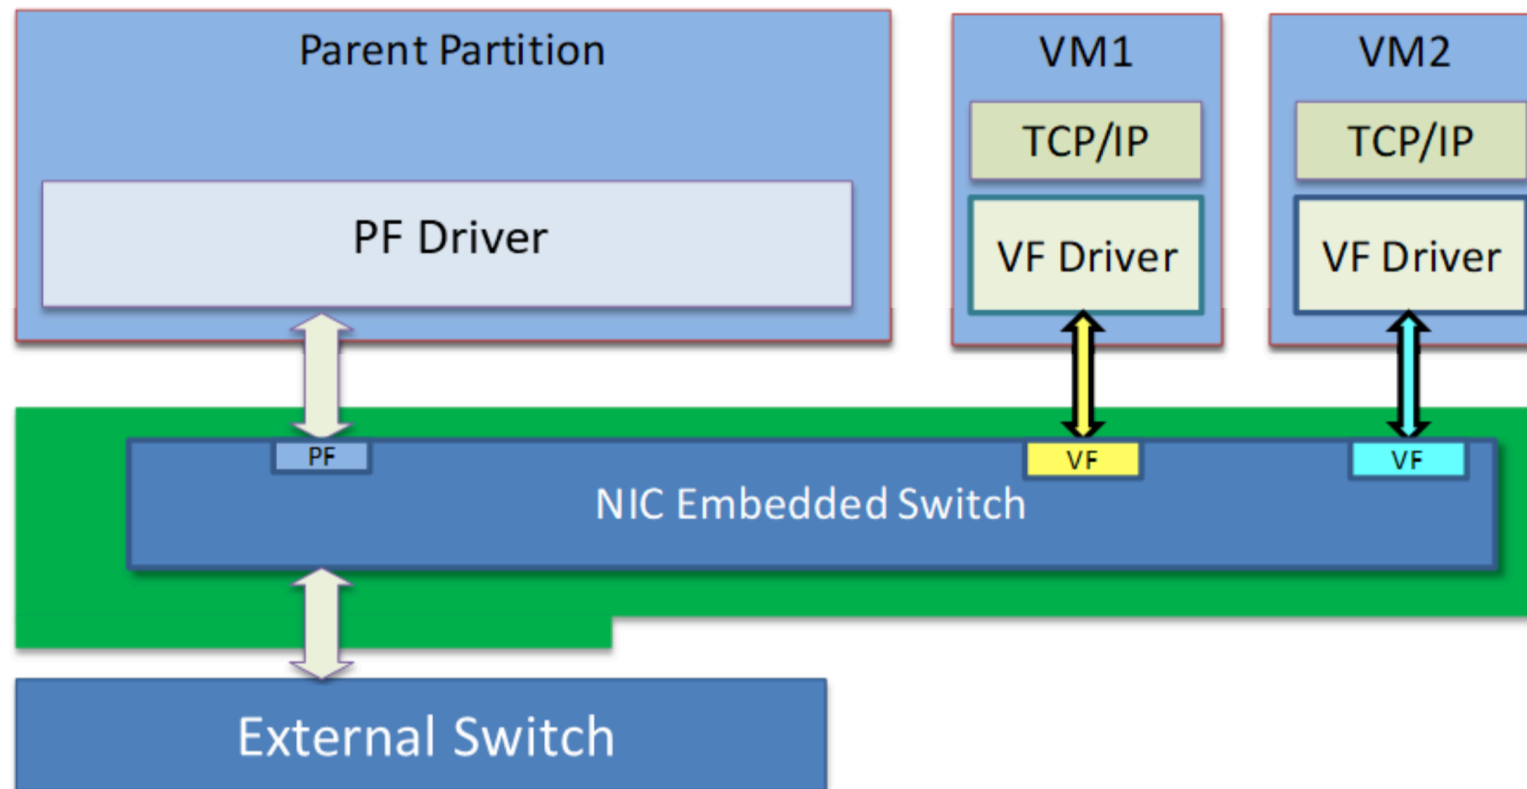
Edge-core partition with virtualization

- Racks contain a bunch of servers
- Top-of-rack (ToR) switch connects
 - Servers to themselves
 - This rack to other racks
- Some new constraints...
- Complex tenant-level policies
 - Moved from core to the edge
- Use cores for apps as much as possible
 - Incentives to reduce network-related CPU



(5) Bypass the hypervisor

- Single-Root I/O Virtualization (SR-IOV)
 - Extension to PCIe bus protocol: support virtual & physical endpoints



(6) Offload to high-perf substrates

- Offloads are all about partitioning of functionality
 - “What do you run where?”
- Key idea: move fast-path processing to a faster compute substrate
 - That is, faster than x86 for the specific purpose of pkt processing 😊
- A regular NIC is an ASIC for network processing
 - High performance, low power, but not programmable
- Other ASICs: TCP offload engines (TOE), RDMA

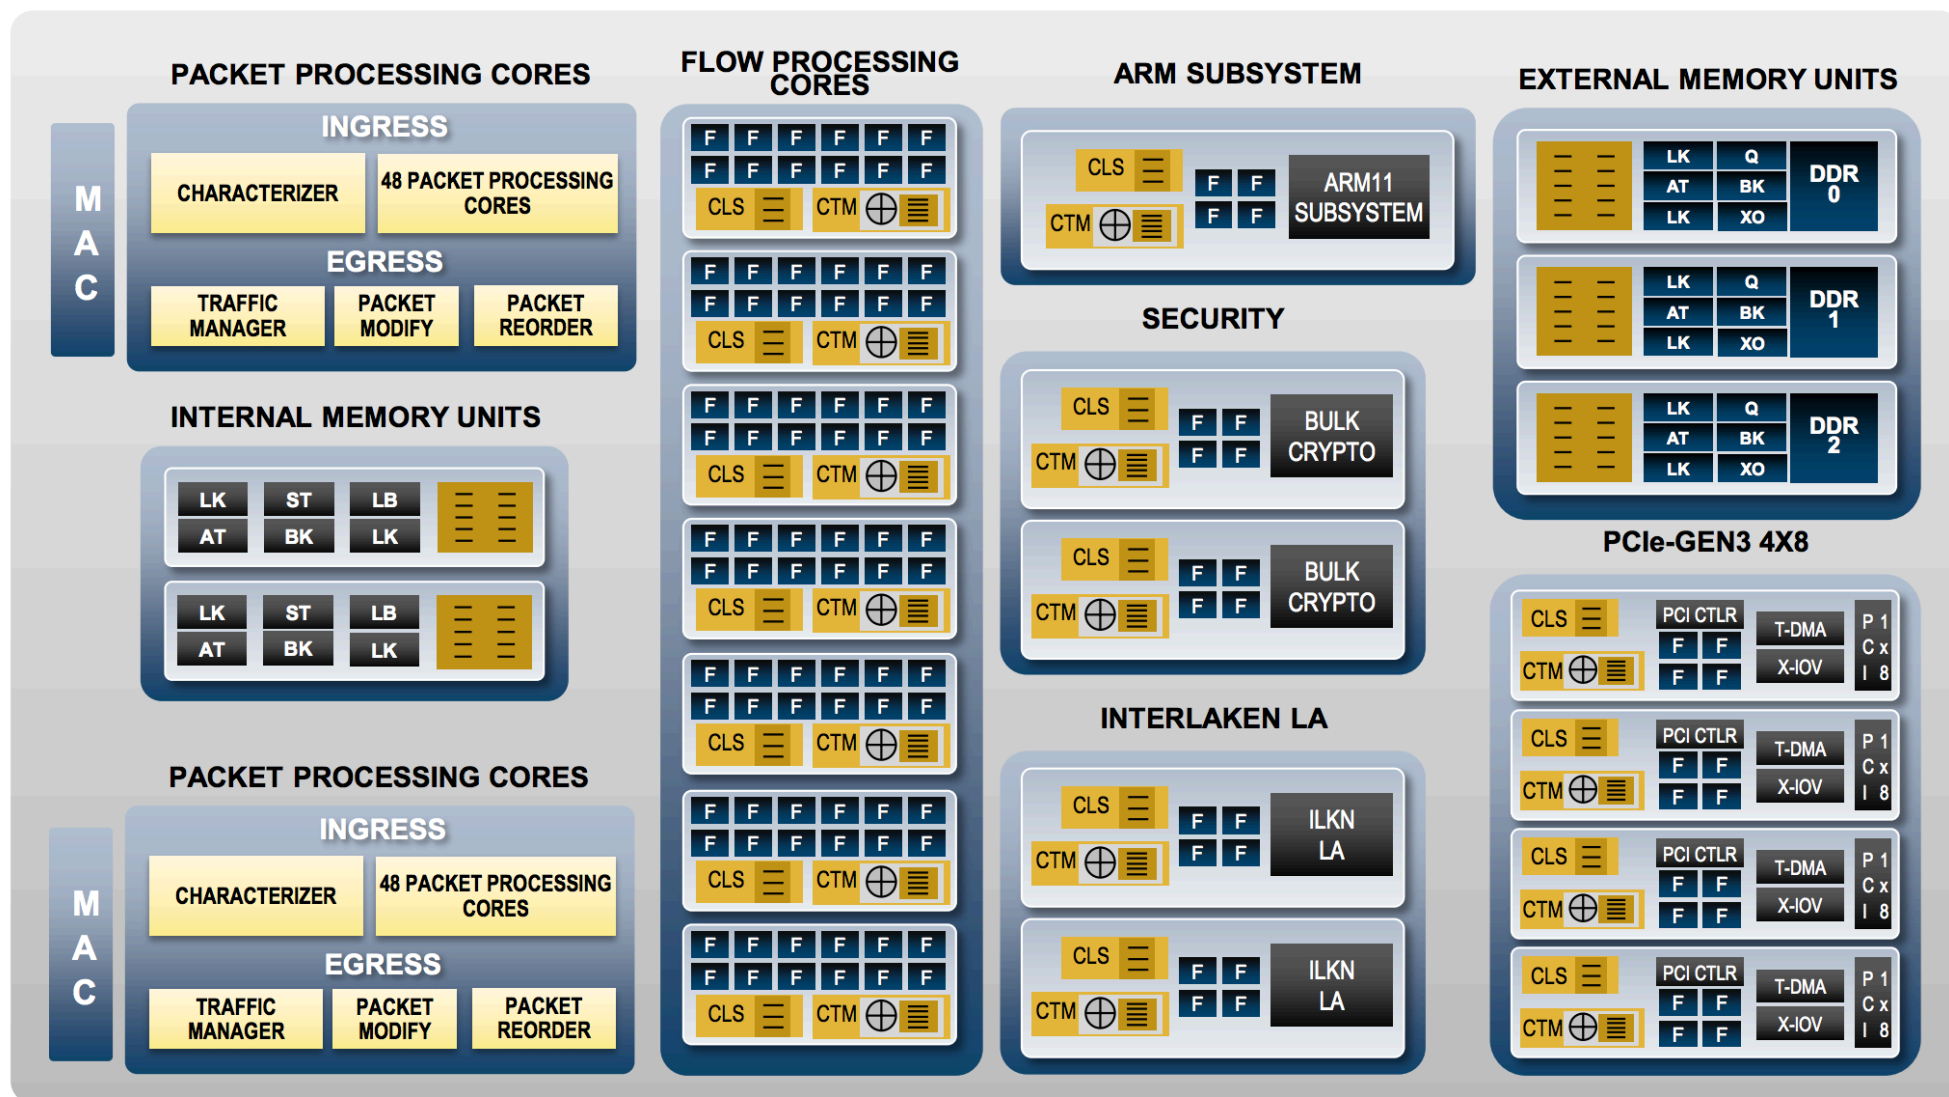
(6) Multi-core System On Chip (SoC)

- SoC: many “small” CPU cores each of which is weaker than x86
- Many fixed-function components for specific functions
 - Cryptography, compression, hashing, etc.
 - ... but general-purpose cores as flexible as x86!

(6) Multi-core System On Chip (SoC)


- Achieve perf by computing on packets *spatially* and *temporally*
- Spray packets across the available cores to achieve high rates
- Stateful processing: all flow's packets must go to same core
 - Troublesome for single high-rate (ex: 100 Gbit/s) connections

(6) Example: Netronome NFP-6xxx




(6) FPGAs

- “Field Programmable Gate Arrays”
- Small programmable components called Lookup Tables (LUTs)
 - Implement any circuit of your choice
- Interconnect moves data from LUT to LUT
 - Interconnect is also configurable



(6) FPGAs

- Achieve perf by programming functionality directly into the LUTs
 - Spatial computation
 - Thread-level parallelism
 - Pipelined parallelism
- *Usually* need circuit-level programming (ex: System Verilog) to achieve high performance
- Challenging for software developers ☹️



(6) Example: Altera Arria 10



Today's papers

- Azure NIC: offloading complex policies to FPGAs
 - Process packets without spinning CPU cores
- HotCocoa: offload congestion control to a NIC

Azure Accelerated Networking: SmartNICs in the Public Cloud

Usenix NSDI '18

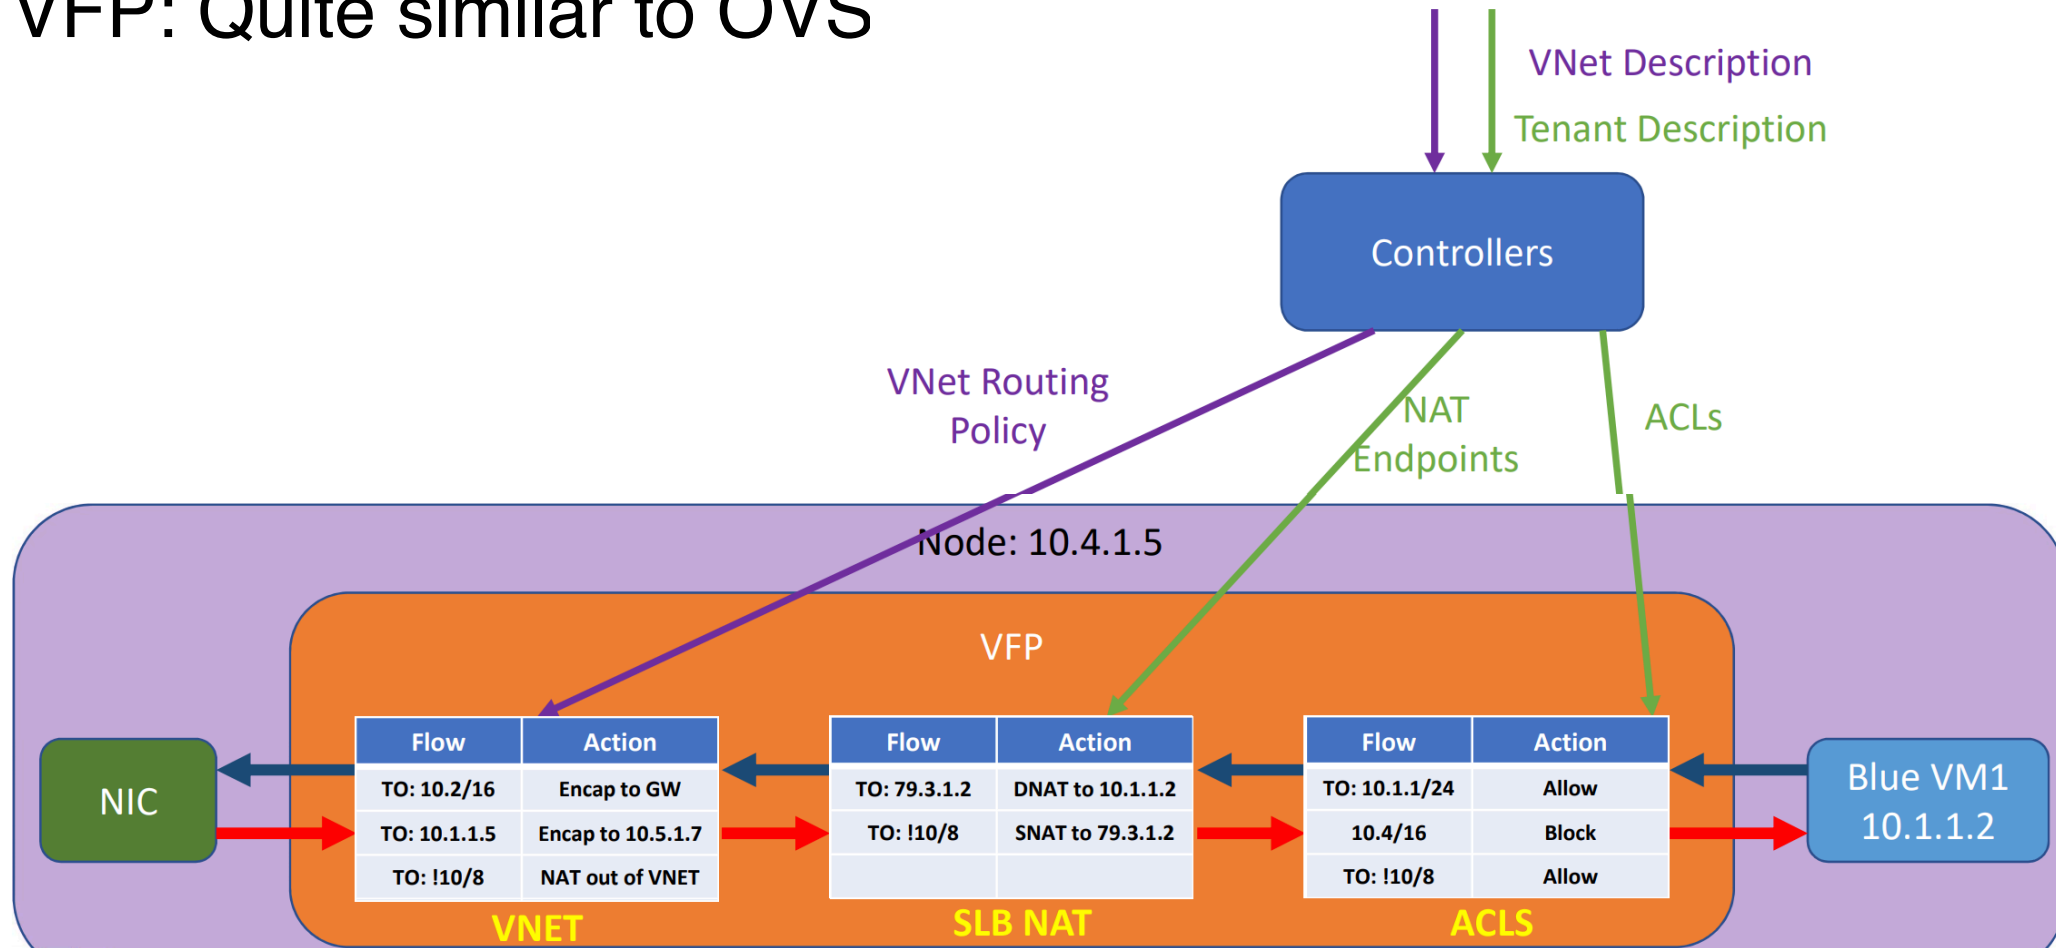
Daniel Firestone et al.

Azure host SDN

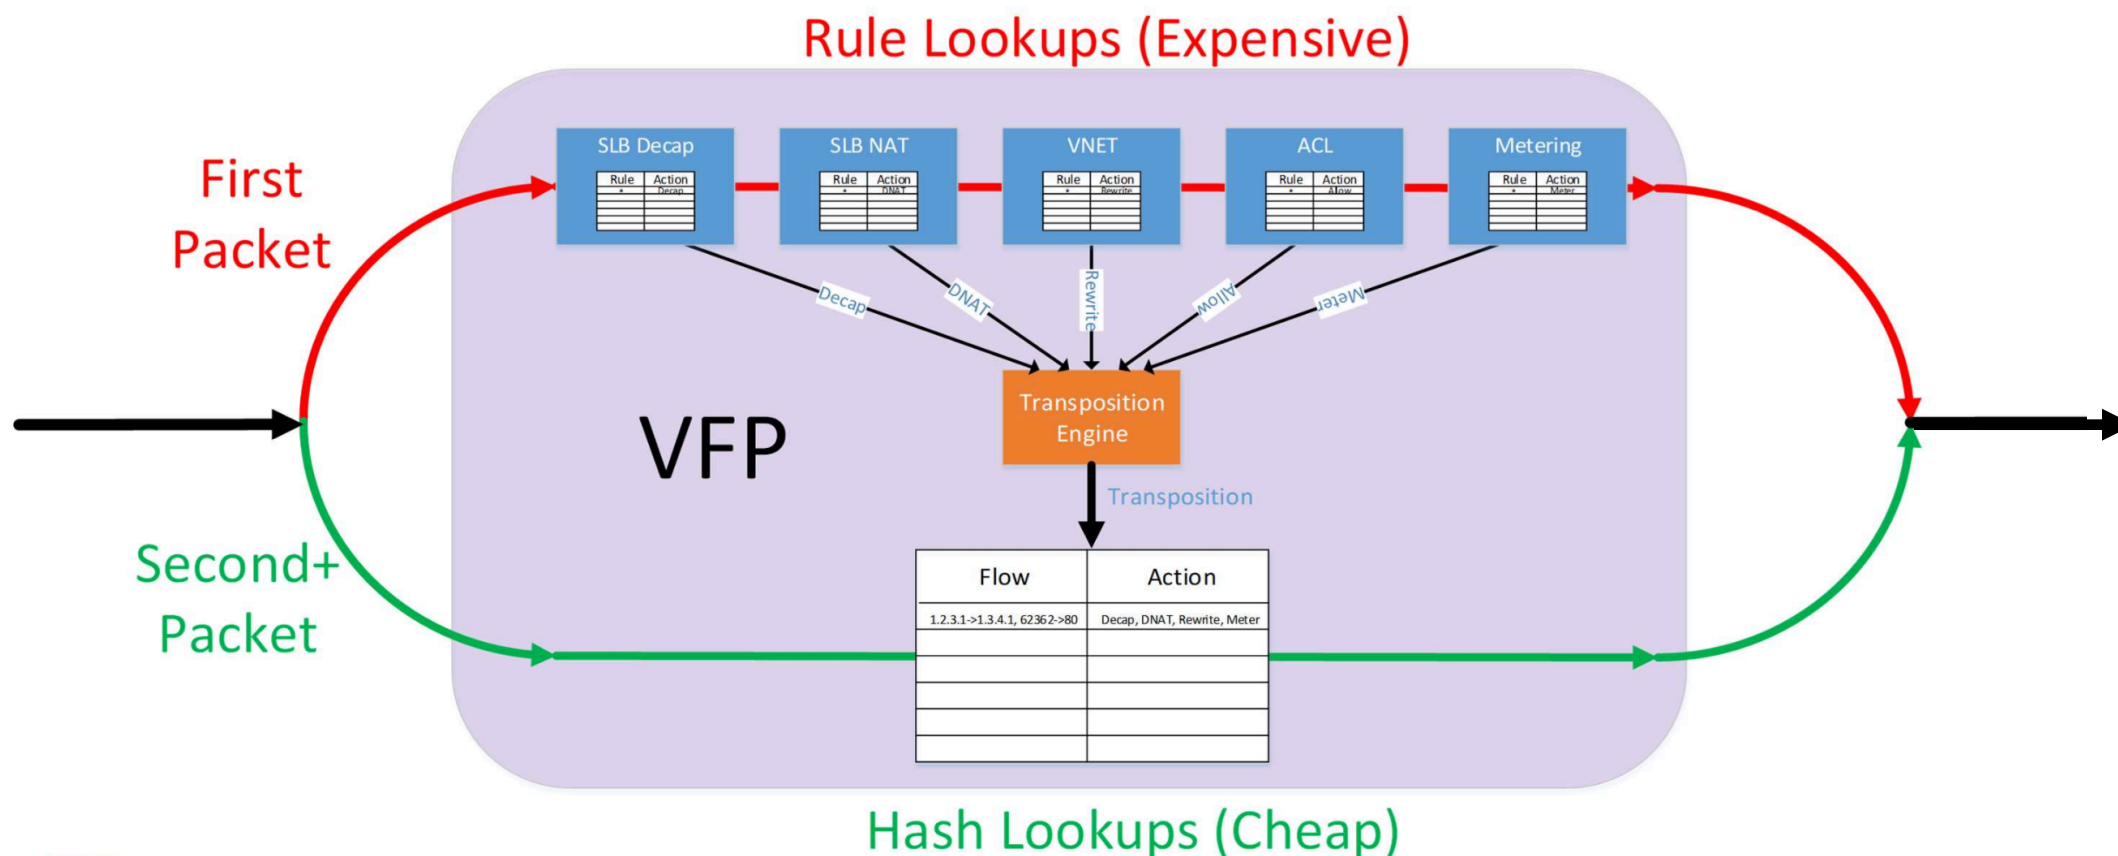
- Endpoint policies implemented on a virtual (software) switch
 - Address virtualization
 - Server load balancing
 - ACLs
 - Metering
 - Security guards
- Programmable flow tables with multiple controllers
 - Cloud provider and multiple tenants

Azure host SDN: Match-Action Tables

- VFP: Quite similar to OVS



Unified flow table: A fast cache for actions

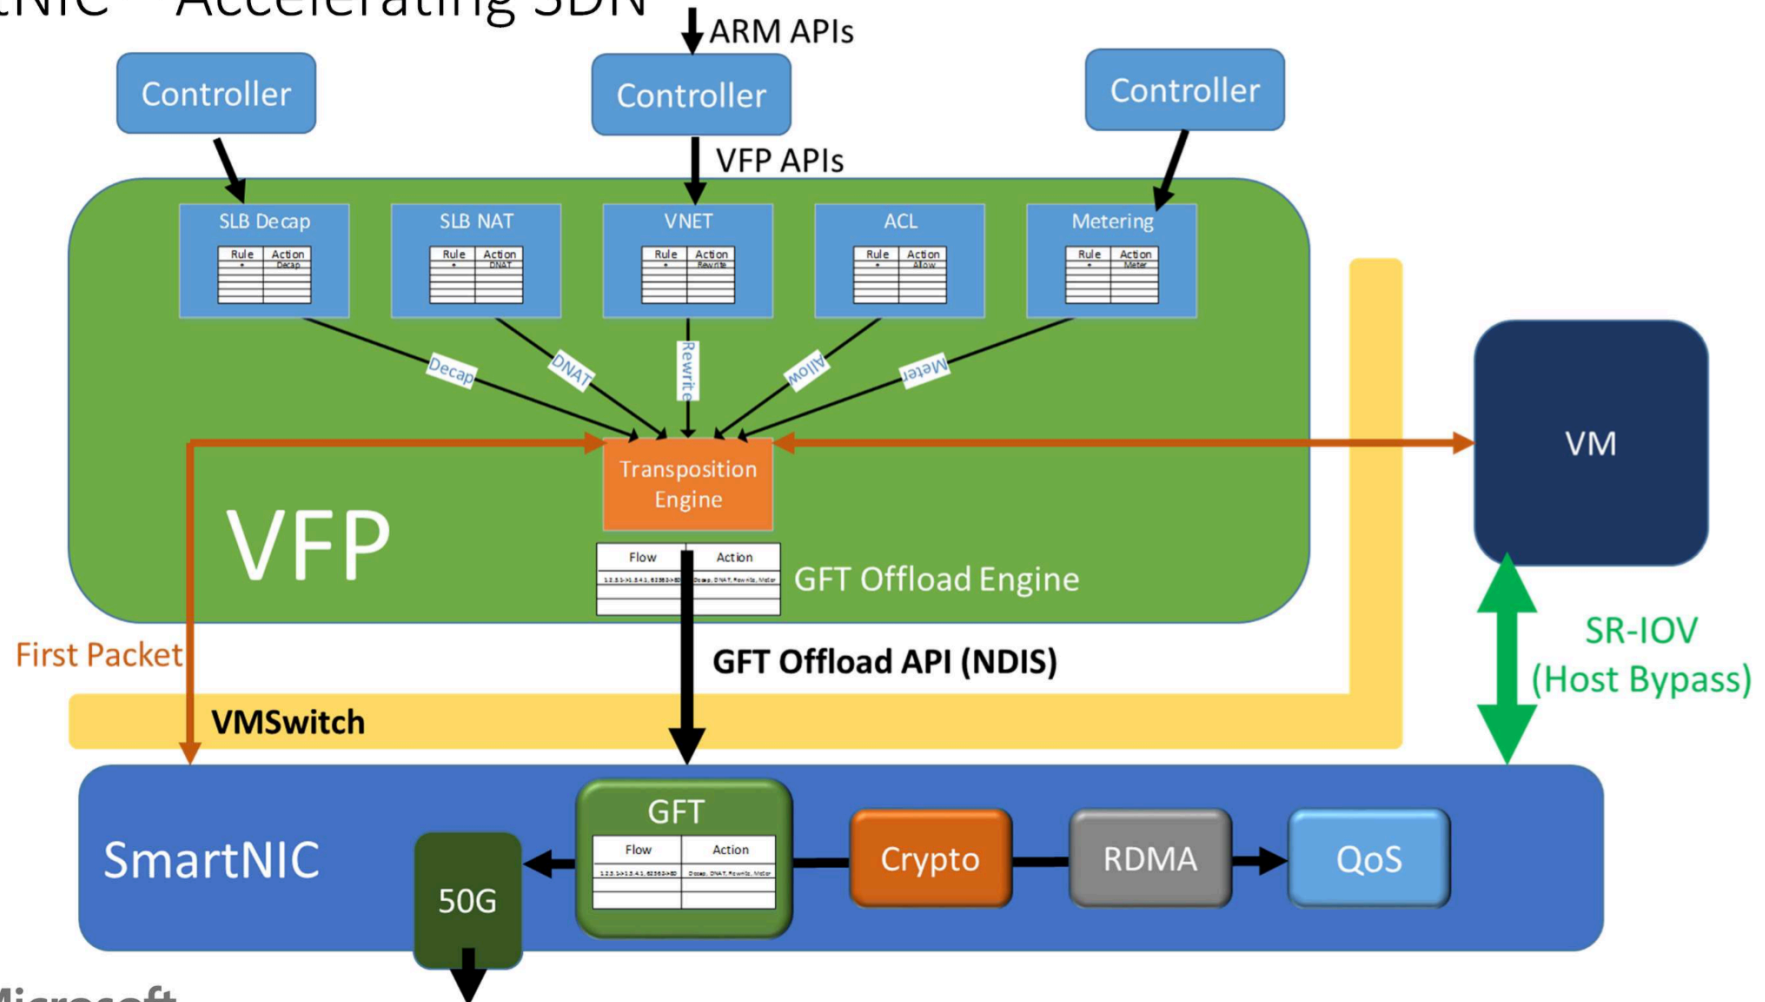


Problems with existing solutions

- Can't scale software to 100 Gbit/s and beyond...
- SR-IOV: an “all or nothing” hardware offload
- Even if one action (e.g., TCP connection state tracking) isn't supported, all processing must fall back to software
- *Key Q:* is it possible to achieve
 - the programmability of the host SDN software solution
 - ... with the performance of hardware?

Azure NIC: an FPGA “smart NIC”

SmartNIC – Accelerating SDN



Hardest challenge: “Serviceability”

- Update parts of the system without breaking connectivity or rebooting VMs
 - ... even if performance suffers (transiently)
- Includes servicing the FPGA circuitry, NIC driver, vSwitch, and the SR-IOV virtual functions
 - And, of course, updating the flow table 😊

HotCocoa: Hardware Congestion Control Abstractions

ACM HotNets '17

Mina Tahmasbi, Monia Ghobadi, Jennifer Rexford, and David Walker

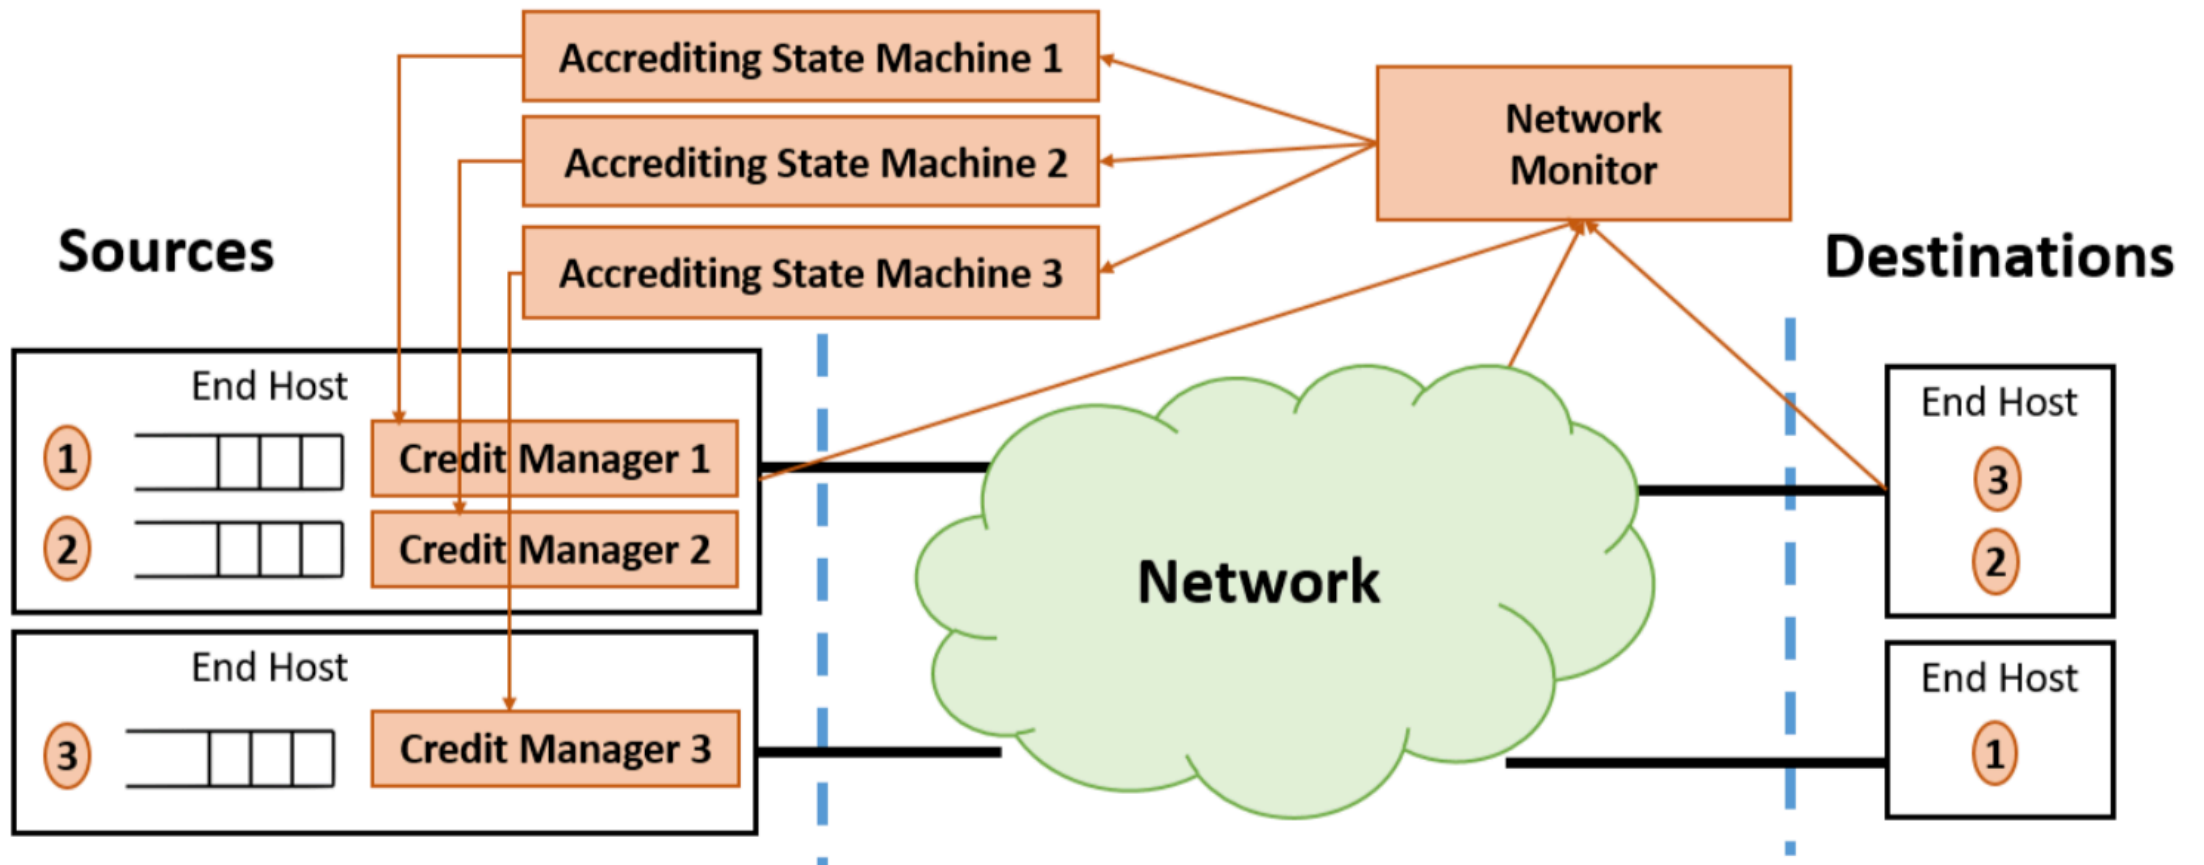
Move congestion control to NICs

- Hypervisor may retain control of the congestion state of connections in a virtualized cluster
- Example: Rate limiting
- Example: Receiver-driven flow control
- How can the hypervisor offload congestion control processing to NIC hardware
 - ... While retaining programmability to implement new schemes?

Congestion control abstractions

- Credit manager: determine when a flow may transmit packets
 - Rate-based control
 - Window-based control
 - ACK clocking
- Accrediting state machine: allocate credits based on network
 - Example: increase congestion window by one every RTT
 - State machine has “triggers” events on network conditions
- Network monitor: external entity that may trigger state transitions

HotCocoa abstractions



An example ASM

```
ASM ex_asm (Flock f, Params p) {  
  
    State Normal{};  
    State Congestion {f.cm.rate = f.rcvd_rate/f.t_rr;};  
  
    Init Normal; // Declare initial State  
  
    Transitions {  
        | Normal on f.pkt_rcvd if f.t_rr > p.max_rate  
          ==> Congestion;  
        | Congestion on f.pkt_rcvd if f.t_rr <= p.max_rate  
    };  
}
```

Conclusions

- Programmable packet processing is an active research area
 - Lots of interest from industry and academic
- The final verdict isn't out on the right offload substrate or partition of functions
 - The answer will depend on the use cases
- Achieving programmability with high performance may require deep expertise in different areas
 - Processor and interconnect architectures
 - Compilers
 - Packet-compute scheduling

Backup Slides

Outline (2/2)

- Substrates for offload: ASIC, multicore SoCs, FPGAs, no offload
- Azure NIC
 - Virtual filtering platform offload: needs?
 - Control-data plane separation
 - FPGA pipeline design
 - Future work: QoS. Better scheduling. Better monitoring
- HotCocoa
 - Congestion control at the hypervisor without the overheads
 - credit management (e.g., window vs. rate) – data plane
 - accrediting state machine (ASM) – control plane