


Congestion Control in Data Centers

Lecture 16, Computer Networks (198:552)



100Kbps–100Mbps links
~100ms latency


Transport
inside the DC


INTERNET

Fabric

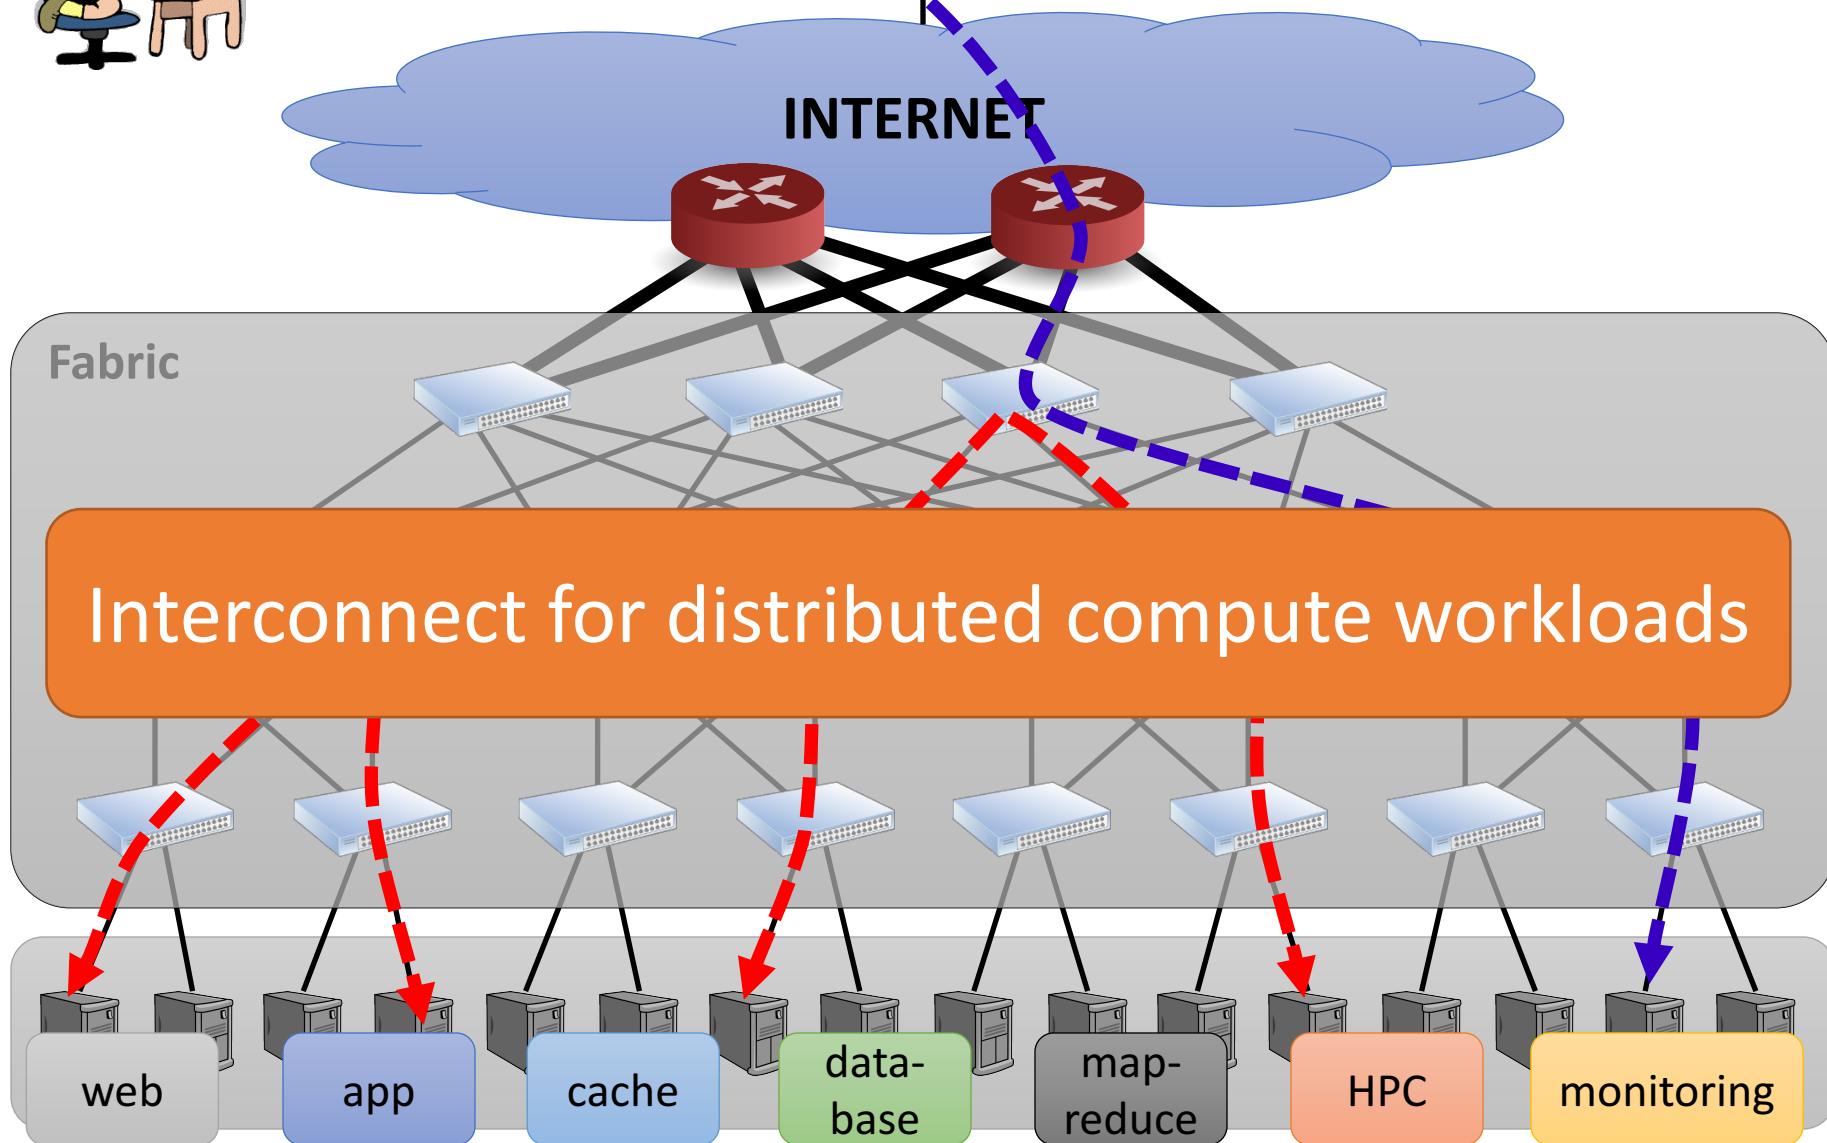
10–40Gbps links
~10–100 μ s latency

Servers





Transport
inside the DC



What's different about DC transport?

- Network characteristics
 - Very high link speeds (Gb/s); very low latency (microseconds)
- Application characteristics
 - Large-scale distributed computation
- Challenging traffic patterns
 - Diverse mix of mice & elephants
 - Incast
- Cheap switches
 - Single-chip shared-memory devices; shallow buffers

Additional degrees of flexibility

- Flow priorities and deadlines
- Preemption and termination of flows
- Coordination with switches
- Packet header changes to propagate information

Data center workloads

- Mice and Elephants!
- Short messages
(e.g., query, coordination)
- Large flows
(e.g., data update, backup)



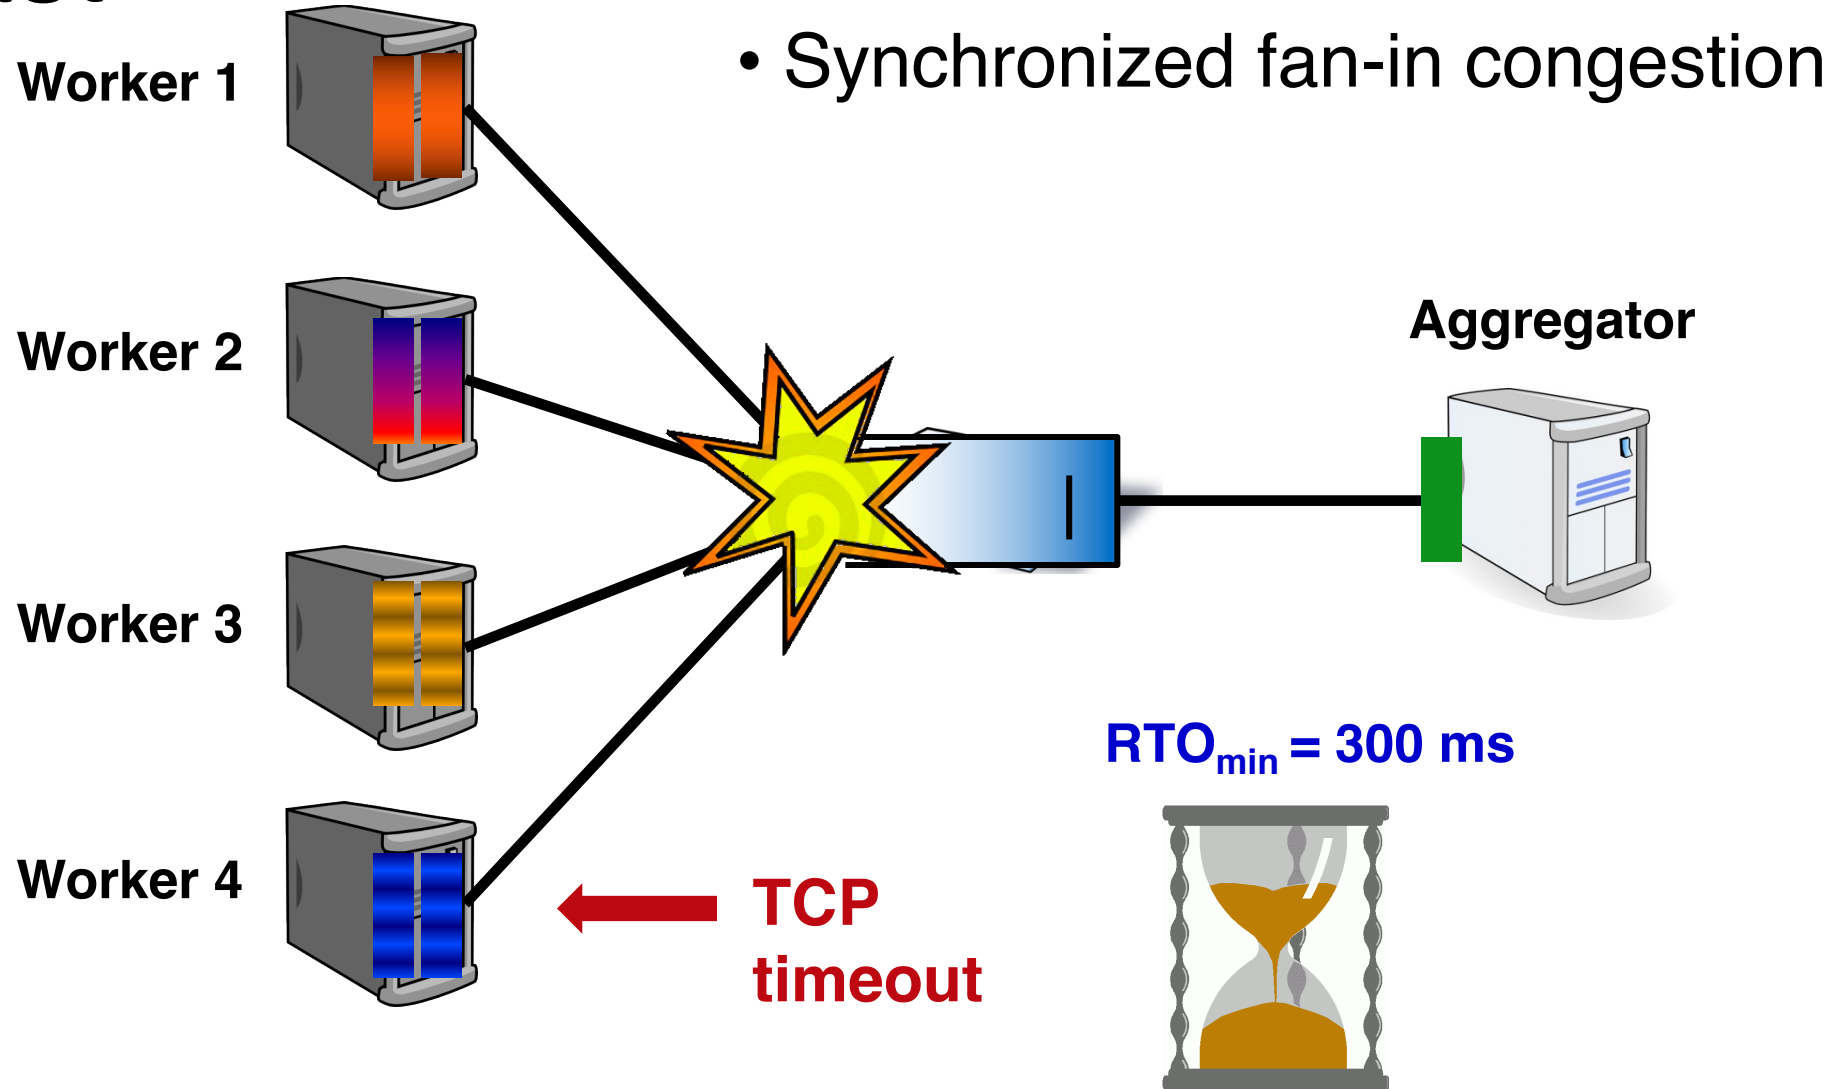
Low Latency




High Throughput



Incast



Incast in Microsoft Bing



Jittering trades of median for high percentiles

DC transport requirements

1. Low Latency

- Short messages, queries

2. High Throughput

- Continuous data updates, backups

3. High Burst Tolerance

- Incast

The challenge is to achieve these *together*

Data Center TCP

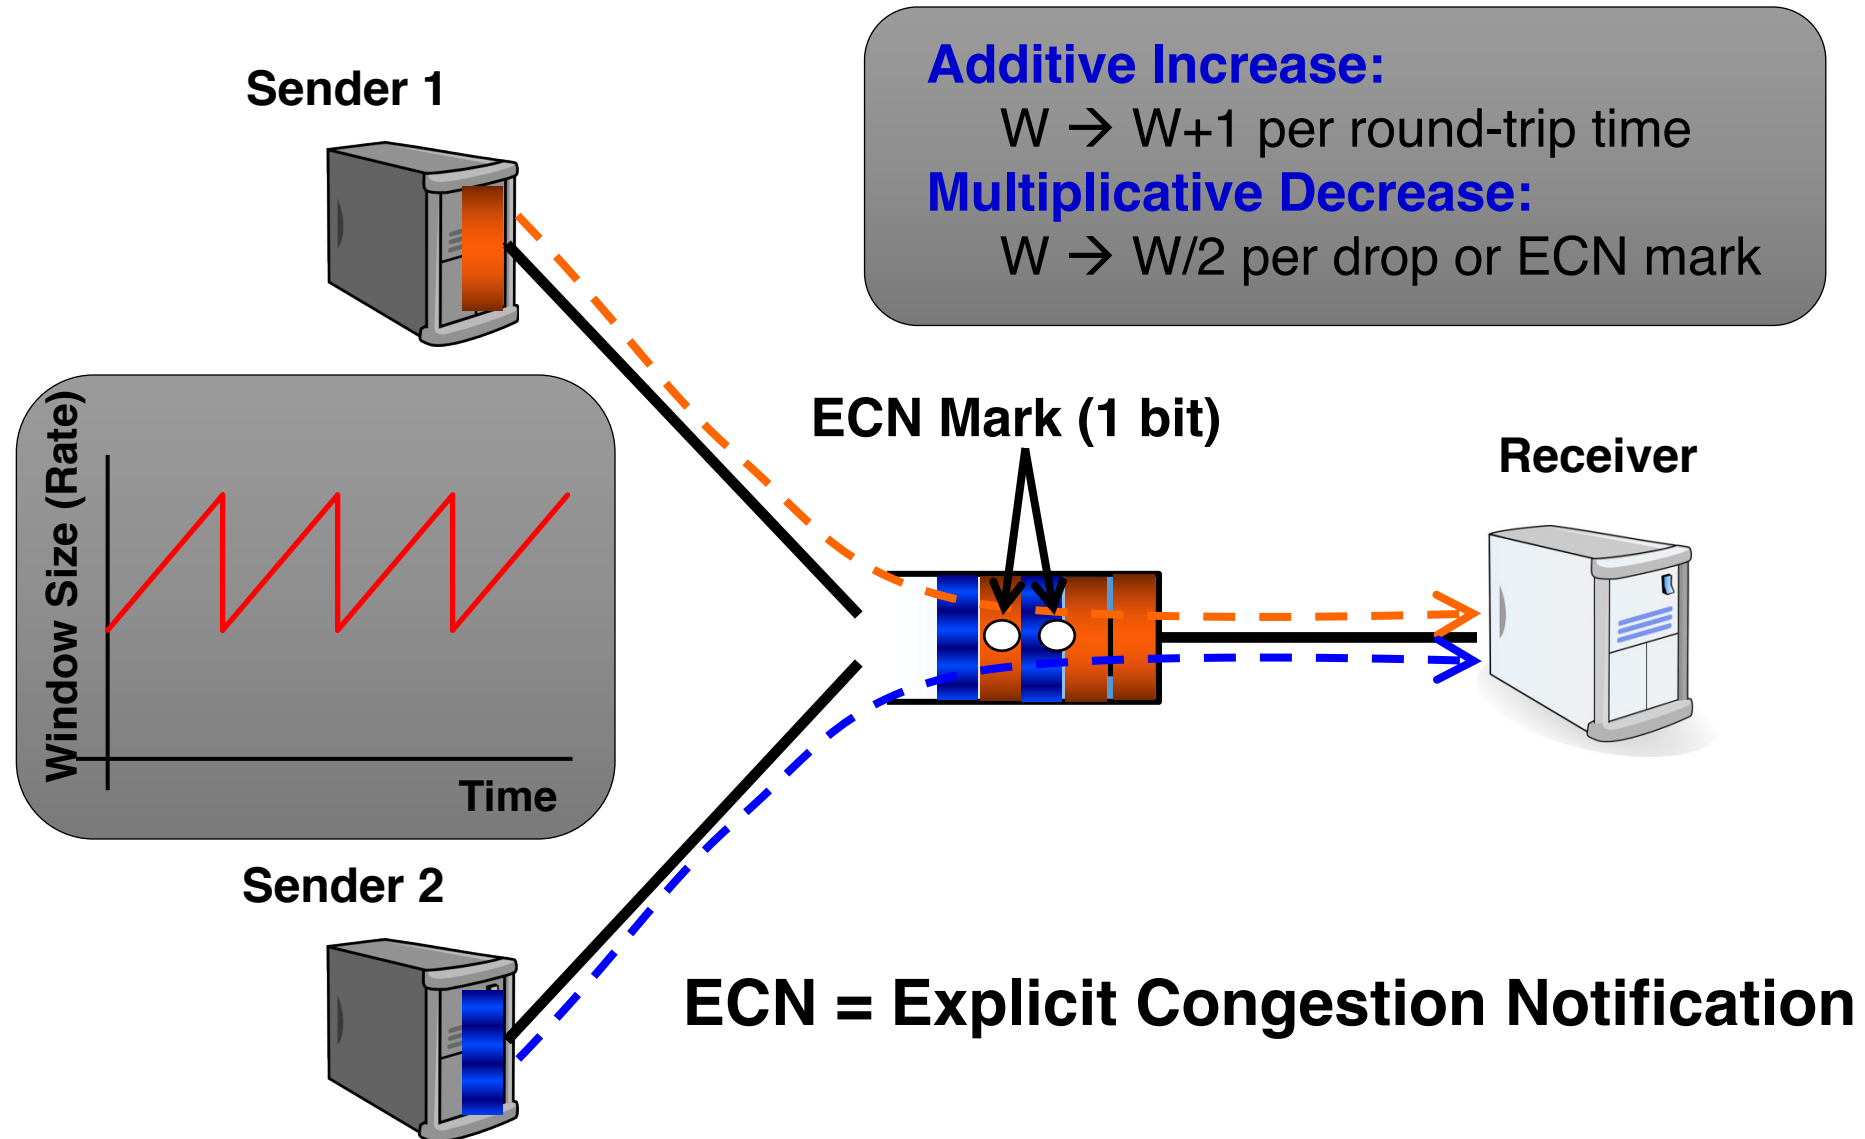
Mohammad Alizadeh et al., SIGCOMM'10

TCP widely used in the data center

- Apps use familiar interfaces
 - TCP is deeply ingrained in the apps
 - ... And developers' minds
- However, TCP not really designed for data center environments
 - Complex to work around TCP problems
 - Ad-hoc, inefficient, often expensive solutions


Practical deployment is hard
→ keep it simple!

Review: TCP algorithm




TCP buffer requirement

- Bandwidth-delay product rule of thumb:
 - A single flow needs $C \times RTT$ buffers for **100% Throughput**.



Reducing buffer requirements

- Appenzeller et al. (SIGCOMM '04):
 - Large # of flows: $C \times RTT / \sqrt{N}$ is enough.



Reducing buffer requirements

- Appenzeller et al. (SIGCOMM '04):
 - Large # of flows: $C \times RTT / \sqrt{N}$ is enough.
- Can't rely on stat-mux benefit in the DC
 - Measurements show typically **only 1-2 large flows** at each server


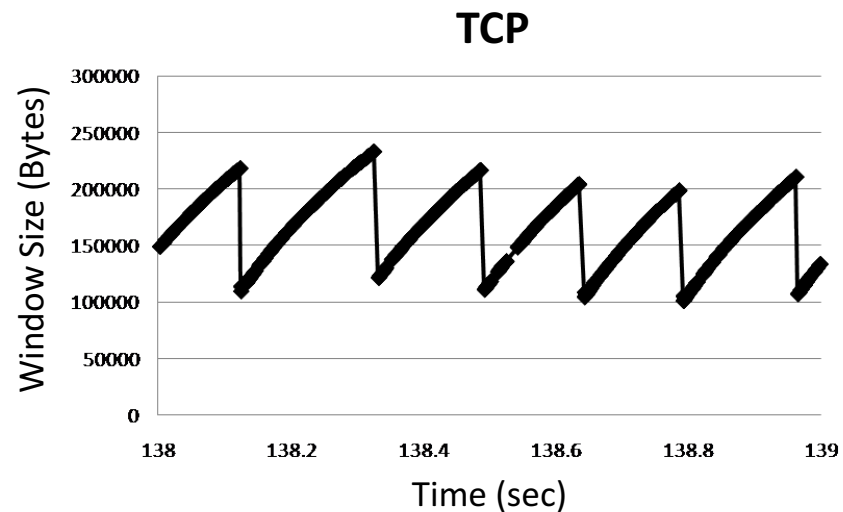
Key observation:
Low variance in sending rate → Small buffers suffice

DCTCP: Main idea

- Extract multi-bit feedback from single-bit stream of ECN marks
 - Reduce window size based on **fraction** of marked packets

DCTCP: Main idea


ECN Marks	TCP	DCTCP
1 0 1 1 1 1 0 1 1 1	Cut window by 50%	Cut window by 40%
0 0 0 0 0 0 0 0 0 1	Cut window by 50%	Cut window by 5%



DCTCP algorithm

Switch side:

- Mark packets when **Queue Length** > **K**.



Sender side:

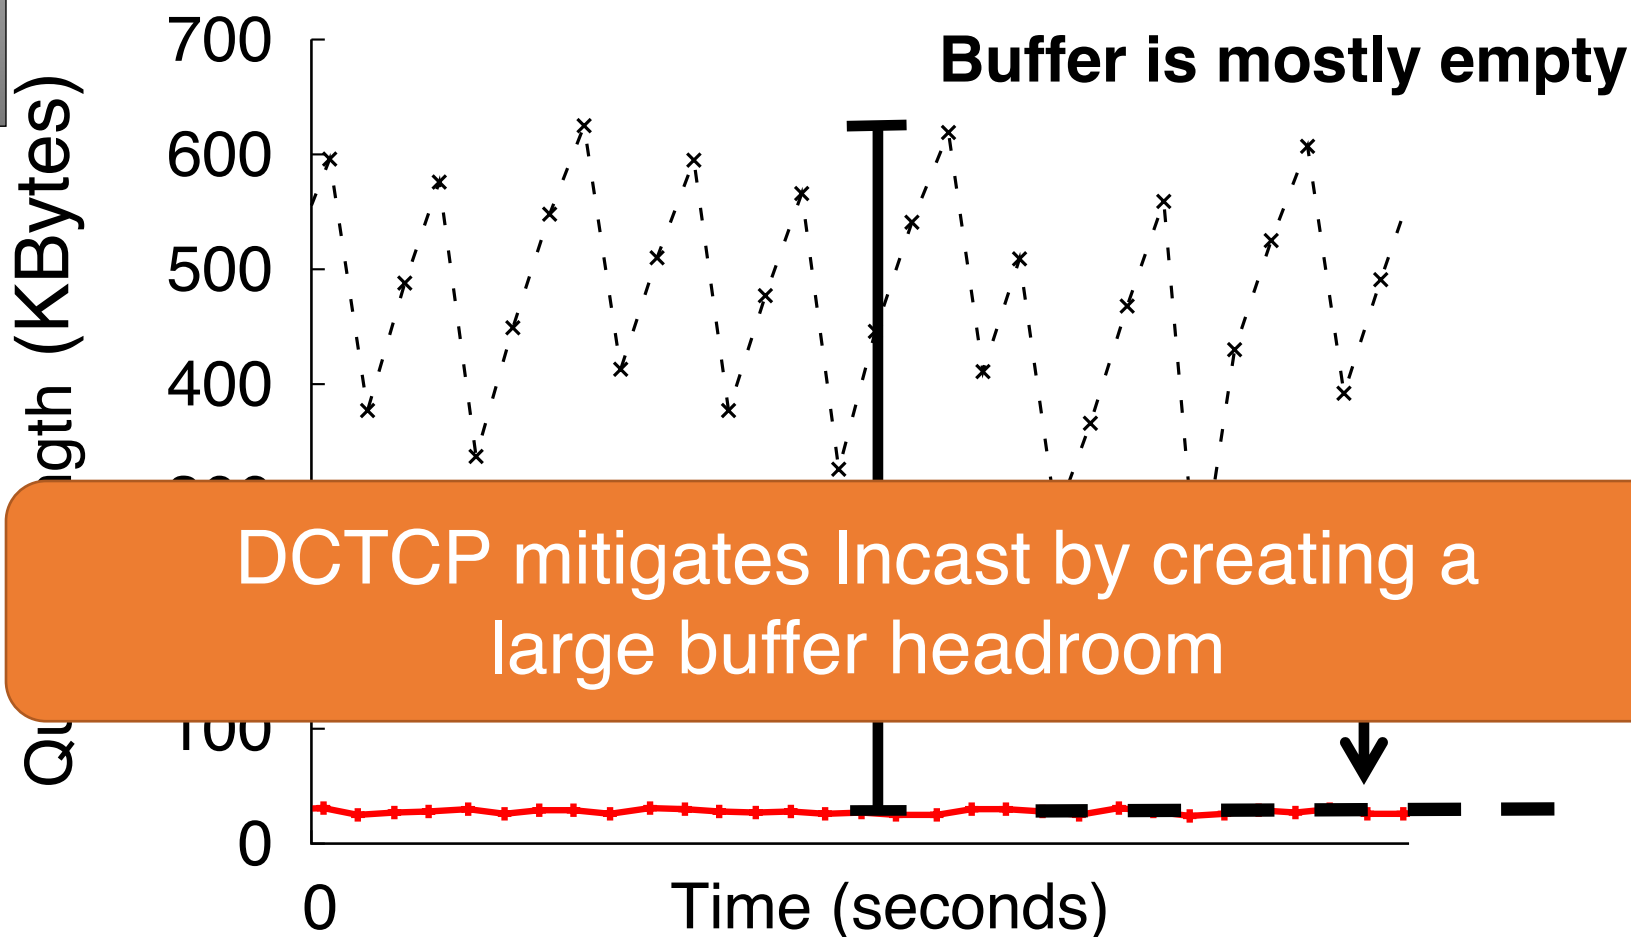

- Maintain running average of ***fraction*** of packets marked (***α***).

$$\text{each RTT: } F = \frac{\# \text{ of marked ACKs}}{\text{Total \# of ACKs}} \Rightarrow \alpha \leftarrow (1 - g)\alpha + gF$$

- Adaptive window decreases:** $W \leftarrow (1 - \frac{\alpha}{2})W$
 - Note: decrease factor between 1 and 2.

DCTCP vs TCP

Experiment: 2 flows (Win 7 stack), Broadcom 1Gbps Switch



Why it works

1. Low Latency

- ✓ **Small buffer occupancies** → low queuing delay

2. High Throughput

- ✓ **ECN averaging** → smooth rate adjustments, low variance



3. High Burst Tolerance

- ✓ **Large buffer headroom** → bursts fit
- ✓ **Aggressive marking** → sources react before packets are dropped

Setting parameters: A bit of analysis

- How much buffering does DCTCP need for 100% throughput?


➤ Need to quantify queue size oscillations (Stability).



Setting parameters: A bit of analysis

- How small can queues be without loss of throughput?

➤ Need to quantify queue size oscillations (Stability).





$$K > (1/7) C \times \text{RTT}$$




for TCP:
$$K > C \times \text{RTT}$$

Convergence time

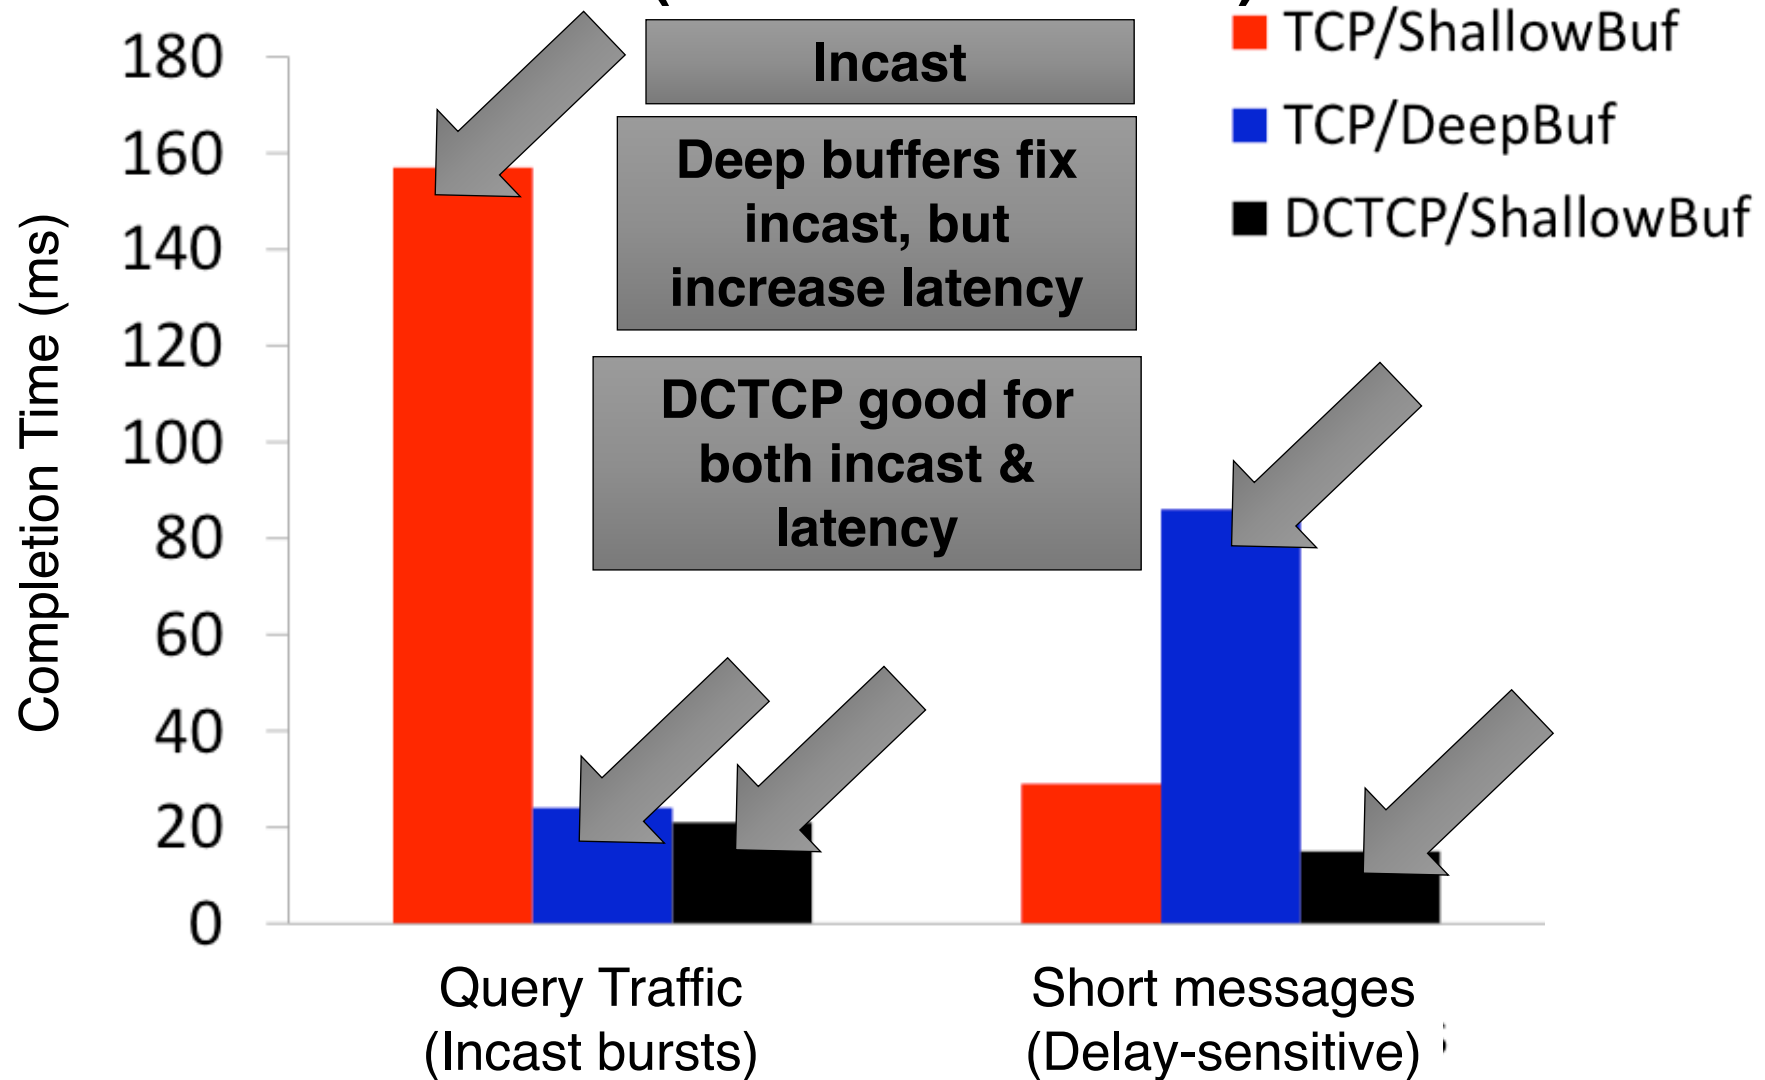
- DCTCP takes at most ~40% more **RTTs** than TCP
 - “Analysis of DCTCP”, SIGMETRICS 2011
- **Intuition:** DCTCP makes smaller adjustments than TCP, but makes them much more frequently



Bing benchmark (baseline)



Bing benchmark (scaled 10x)



Discussion

- Between throughput, delay, and convergence time, what metrics are you willing to give up? Why?
- Are there other factors that may determine choice of K and B besides loss of throughput and max queue size?
- How would you improve on DCTCP?
- How could you add on flow prioritization over DCTCP?

Acknowledgment

- Slides heavily adapted from material by Mohammad Alizadeh

