

Backbone Traffic Engineering


Lecture 20, Computer Networks (198:552)

Do IP Networks Manage Themselves?

- In some sense, yes:
 - TCP senders send less traffic during congestion
 - Routing protocols adapt to topology changes
- But, does the network run *efficiently*?
 - Congested link when idle paths exist?
 - High-delay path when a low-delay path exists?
- How should routing adapt to the traffic?
 - Avoiding congested links in the network
 - Satisfying application requirements (e.g., delay)
- ... these are the essential questions of **traffic engineering**


Routing With “Static” Link Weights

- Routers flood information to learn topology
 - Determine “next hop” to reach other routers...
 - Compute shortest paths based on link weights
- Link weights configured by network operator

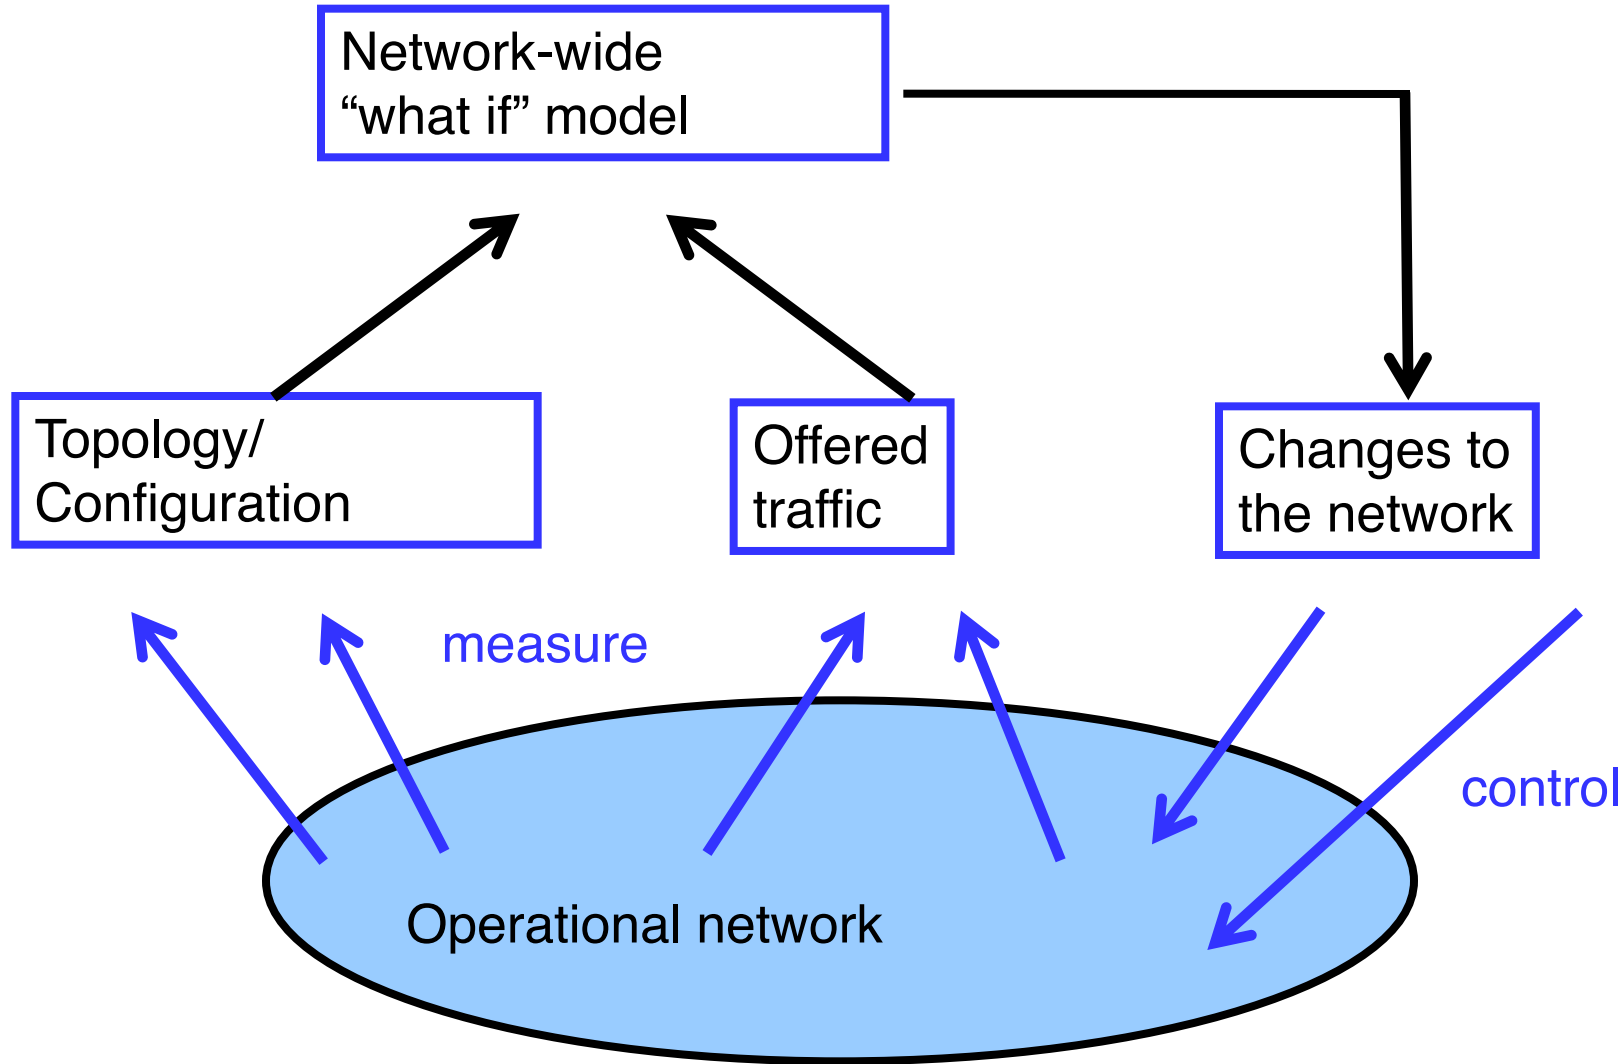


Setting the Link Weights

- How to set the weights
 - Inversely proportional to link capacity?
 - Proportional to propagation delay?
 - Network-wide optimization based on traffic?



Measure, Model, and Control




Key Ingredients

- Measurement
 - Topology: monitoring of the routing protocols
 - Traffic matrix: passive traffic measurement
- Network-wide models
 - Representations of topology and traffic
 - “What-if” models of shortest-path routing
- Network optimization
 - Efficient algorithms to find good configurations
 - Operational experience to identify constraints

Optimization Problem

- Input: graph $G(R, L)$
 - R is the set of routers
 - L is the set of unidirectional links
 - c_l is the capacity of link l
- Input: traffic matrix
 - $M_{i,j}$ is traffic load from router i to j
- Output: setting of the link weights
 - w_l is weight on unidirectional link l
 - $P_{i,j,l}$ is fraction of traffic from i to j traversing link l


Equal-Cost Multipath (ECMP)



Values of $P_{i,j,l}$

Objective Function

- Computing the link utilization
 - Link load: $u_l = \sum_{i,j} M_{i,j} P_{i,j,l}$
 - Utilization: u_l/c_l
- Objective functions
 - $\min(\max_l(u_l/c_l))$
 - $\min(\sum_l f(u_l/c_l))$



Complexity of Optimization Problem

- NP-complete optimization problem
 - No efficient algorithm to find the link weights
 - Even for simple objective functions
- What are the implications?
 - Have to resort to *searching* through weight settings
- Clearly suboptimal, but effective in practice
 - Fast computation of the link weights
 - Good performance, compared to “optimal” solution

Incorporating Operational Realities

- Minimize number of changes to the network
 - Changing just 1 or 2 link weights is often enough
- Tolerate failure of network equipment
 - Weights settings usually remain good after failure
 - ... or can be fixed by changing one or two weights
- Limit dependence on measurement accuracy
 - Good weights remain good, despite random noise
- Limit frequency of changes to the weights
 - Joint optimization for day & night traffic matrices

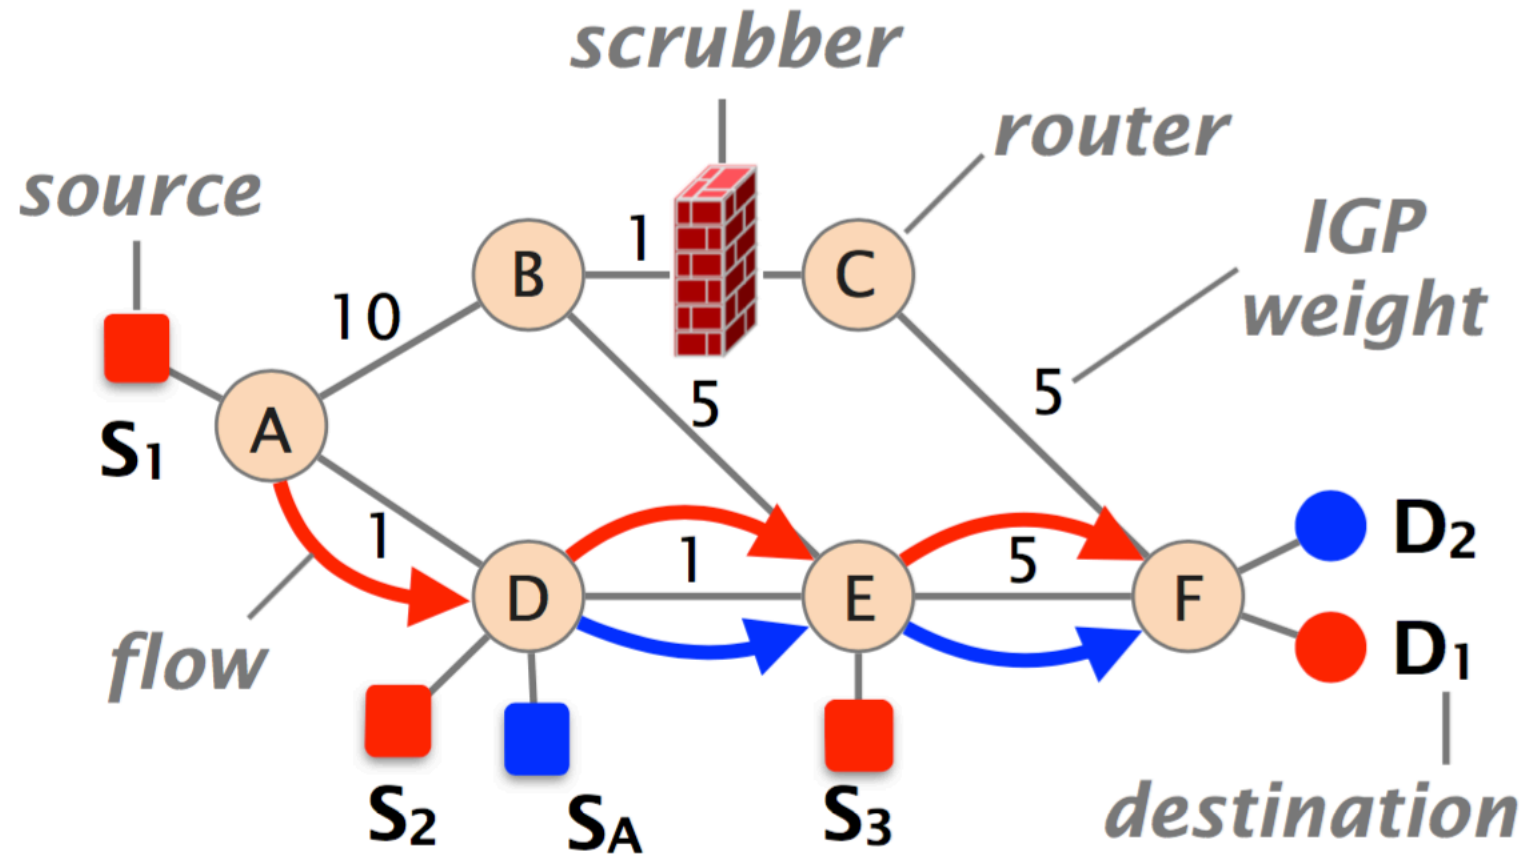
Central Control over Distributed Routing

Stefano Visicchio et al., ACM SIGCOMM'15


Context: Scalable & robust backbone TE

- Traditional IGPs perform distributed computations
 - Scalable
 - Robust to link and node failures
 - But not flexible in terms of expressed paths
- SDNs perform (logically) centralized path computations
 - Highly flexible
 - But handling large networks is challenging
 - Handling topology updates is challenging
- Could we combine the best of both worlds?

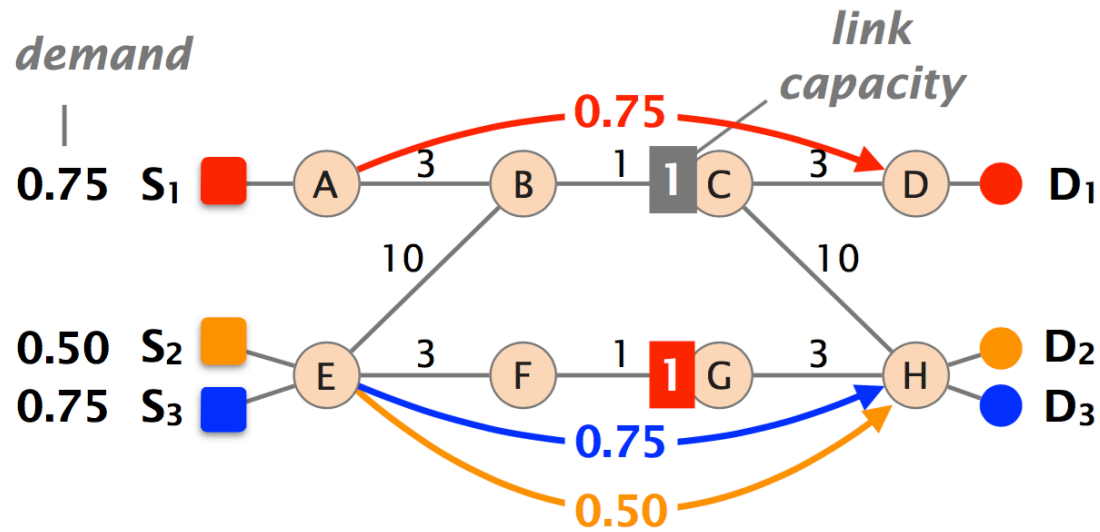
An example




Key idea: Lies to the routers



Another example

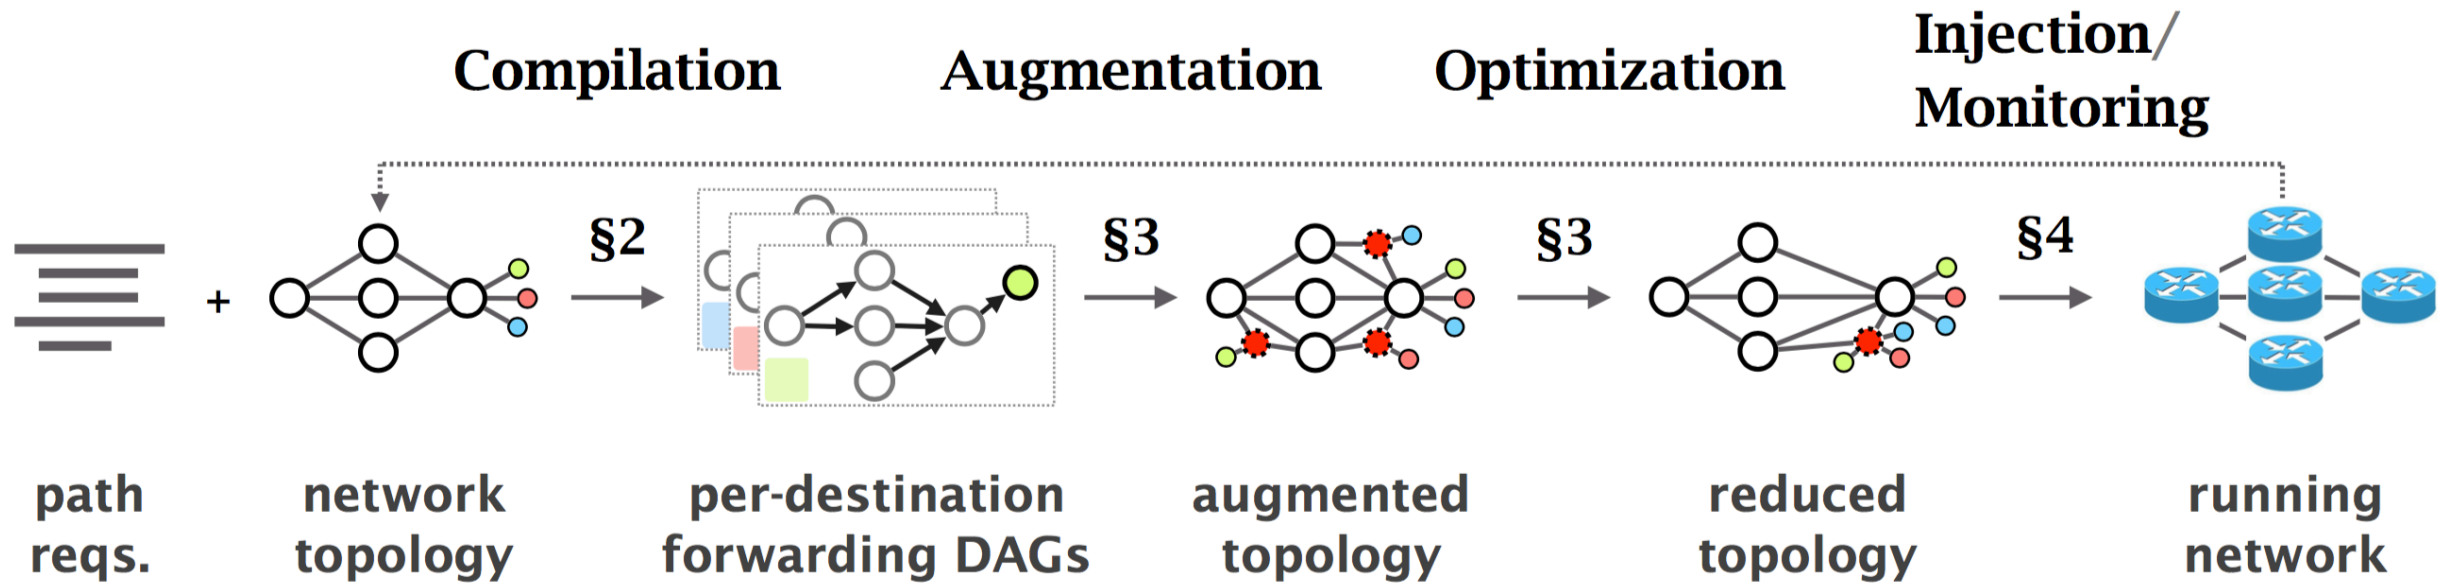


(a) Initial topology



(b) Augmented topology

Workflow




Fibbing: Expressiveness

THEOREM 1. *Any set of per-destination forwarding DAGs can always be enforced by augmenting a Fibbing-compliant topology even only with globally-scoped lies.*


Key primitives required

- Initial setting of static routing weights: “Fibbing-compliance”
- Local and global scoping of IGP announcements
- Forwarding traffic to fake nodes on any desired link


Augmentation algorithms



(a) Requirements





(b) Simple augmentation



(c) Merger augmentation

Cross-destination optimization



A comparison of approaches

	<i>centralized/SDN</i> OpenFlow [2], PCE [3], SR [4]	<i>distributed/traditional</i> IGP [5, 6], RSVP-TE [1]	<i>hybrid</i> Fibbing
<i>forwarding paths:</i>			
- <i>configuration</i>	simple (declarative & global)	complex (indirect & per-device)	simple (declarative & global)
- <i>manageability</i>	high (direct control)	low [7, 8] (need for coordination)	high (direct control)
- <i>path installation</i>	slow (by controller, per-device)	fast (by device, distributed)	fast (by device, distributed)
<i>robustness:</i>			
- <i>network failures</i>	slow (by controller)	fast (local)	fast (local)
- <i>controller failures</i>	hard (ad-hoc synch)	native (distributed)	easy (synch via IGP)
- <i>partitions</i>	hard (uncontrollable devices)	best (distributed)	best (fallback on distributed)
<i>routing policies:</i>	highest (any path)	- low for IGP (shortest paths) - highest for RSVP (any path)	high (any non-loopo paths)

Acknowledgment

- Slides heavily adapted from material by Jennifer Rexford and Stefano Visicchio