

Transport Security

Lecture 12, Computer Networks (198:552)

Fall 2019

Why security?

- Malicious people share your network
 - People who want to snoop, corrupt, destroy, pretend, steal, ...
- Problem made more severe as Internet becomes more commercialized
- Active and passive attacks

Key aspects of network security

confidentiality: only sender, intended receiver should “understand” message contents

- sender encrypts message
- receiver decrypts message


integrity: sender, receiver want to ensure message not altered (in transit, or afterwards) without detection

authentication: sender, receiver want to confirm identity of each other

non-repudiation: Once someone sends a message, or conducts a transaction, she can't later deny the contents of that message

Friends and enemies: Alice, Bob, Trudy

- well-known in network security world
- Bob and Alice want to communicate “securely”
- Trudy (intruder) may intercept, delete, add messages



Who might Bob and Alice be?

- Real humans 😊
- Web browser/server for electronic transactions (e.g., on-line purchases)
- on-line banking client/server
- DNS servers
- routers exchanging routing table updates

What can bad actors do?

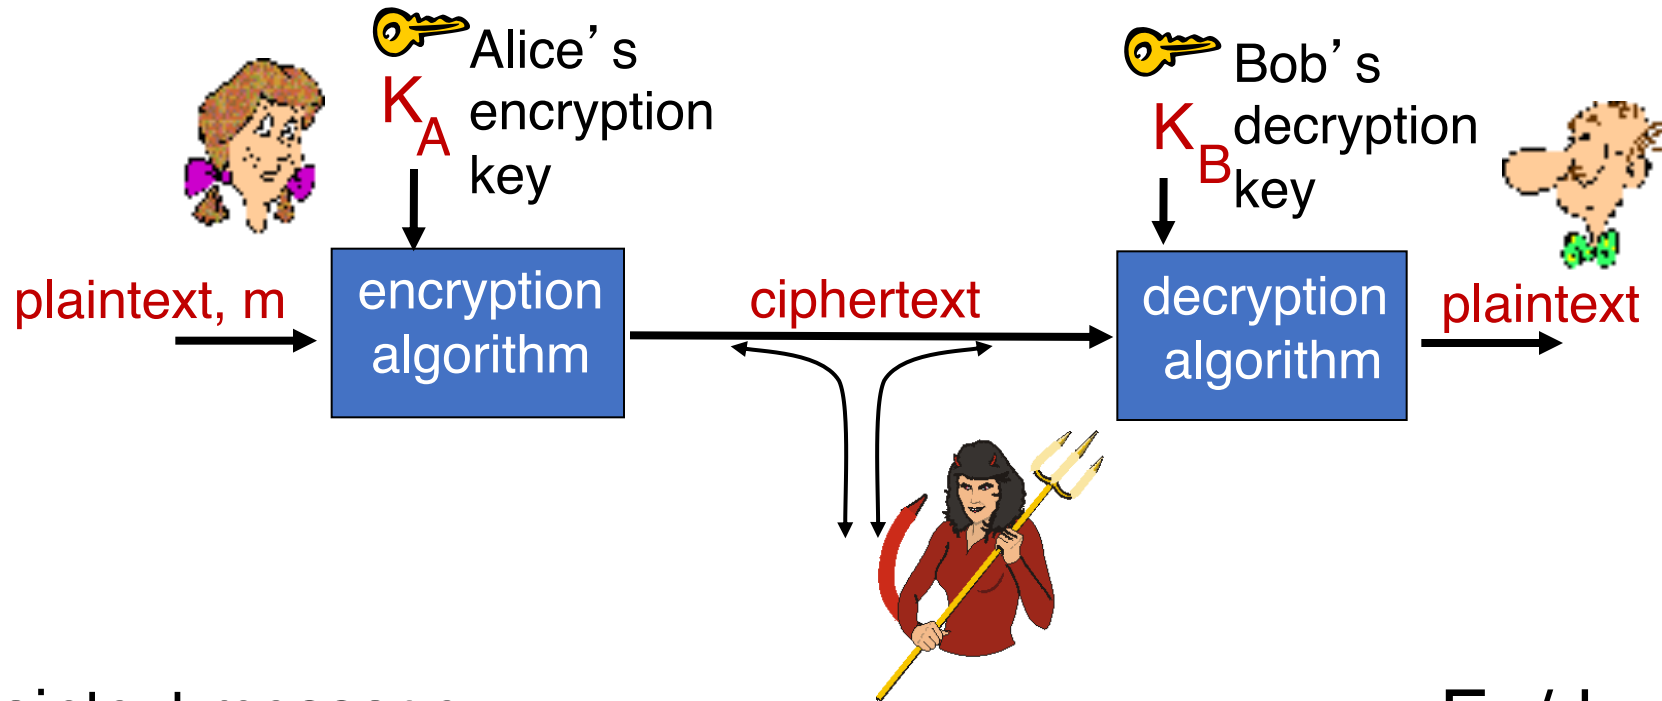
A lot!

- *eavesdrop*: intercept messages
- actively *insert* messages into connection
- *impersonation*: can fake (spoof) source address in packet (or any field in packet)
- *hijacking*: “take over” ongoing connection by removing sender or receiver, inserting itself in place
- *denial of service*: prevent service from being used by others (e.g., by overloading resources)

Confidentiality

Cryptography: preventing adversaries from reading private messages

Cryptography: Terminology



m plaintext message

$c = K_A(m)$, $K_A(m)$ ciphertext, encrypted with key K_A

$m' = K_B(c)$, $K_B(c)$ decrypted plaintext with key K_B

Want: $m = K_B(K_A(m))$

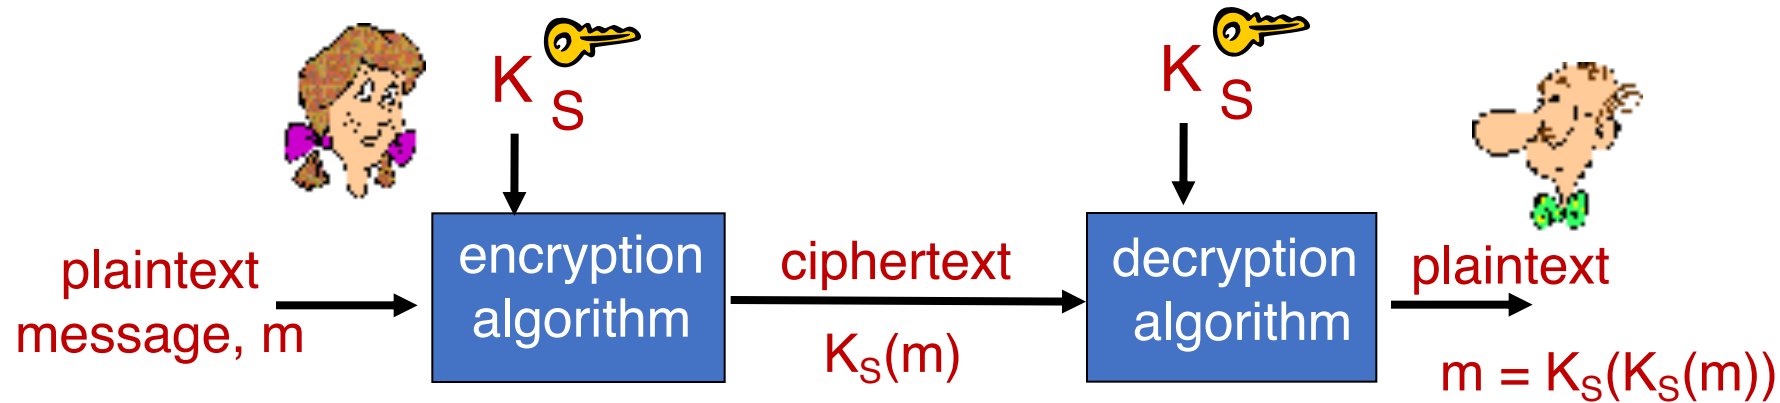
Want: $K_A(m)$ to be uncorrelated with m

En/decryption algorithms are also called **ciphers**.

Cryptography: Algorithms and Keys

- Cryptography requires both an en-/decryption **algorithm** and **keys**
 - Key is a string known only to Alice and Bob, which controls how algorithm works
- Algorithm should be public and known to all
 - Inspires trust that the algorithm works
- Keys
 - Should be long enough to prevent easy breaking of the encryption
 - Should be short enough to keep algorithm efficient
 - Typical key lengths: 56-bit, 128-bit, 256-bit, 512-bit

Symmetric key cryptography



Symmetric keys: Bob and Alice share same (symmetric) key: S

Main techniques of symmetric key cryptography:

Substitution and **Permutation**

Q: how do Bob and Alice agree on key value?

How to agree on a shared secret key?

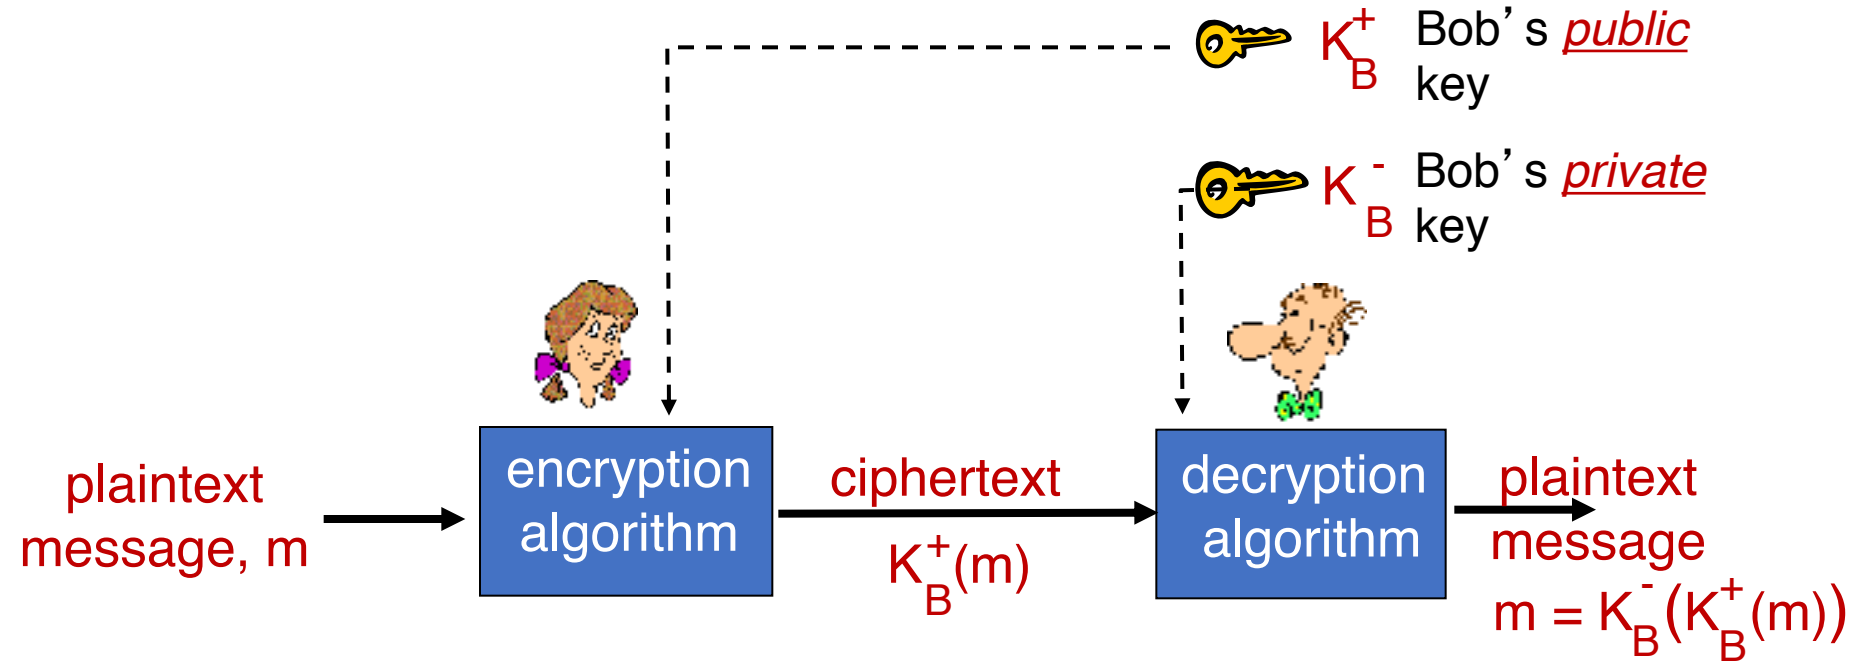
- In reality: two parties may meet in person or communicate “out of band” to exchange shared key
- But communicating parties may never meet in person
 - Example: An online retailer and customer
 - Much more common for a network 😊
- What if the shared secret is stolen?
 - All secret communications can now be decrypted and are visible
 - Including earlier ones that were encrypted using that secret
- How to communicate without necessitating key exchange?

Public key cryptography

Public Key Cryptography

- Sender and receiver do **not** share secret key
- **public** encryption key known to **all**
- **private** decryption key known only to the receiver

Public key cryptography (eg: RSA)




Diffie Hellman Merkle key exchange

- Alice and Bob agree on a modulus p and base g
- Alice chooses secret a , sends bob $A = g^a \bmod p$
- Bob chooses secret b , sends bob $B = g^b \bmod p$
- Alice computes $B^a \bmod p$
- Bob computes $A^b \bmod p$
- Is the common key computed by Alice and Bob the same?
- In what sense is D-H-M key exchange secure?

Public vs. Symmetric key crypto

- Public key crypto
- Expensive to encrypt using just modular exponentiation operations
- No need to exchange keys
- Symmetric key crypto
- Encryption and decryption are fast
- But need to solve the key exchange problem

Crypto in practice: session keys



Use public key crypto or key exchange to agree on a **symmetric session key**

Use symmetric key to protect the rest of the session **efficiently**

Integrity

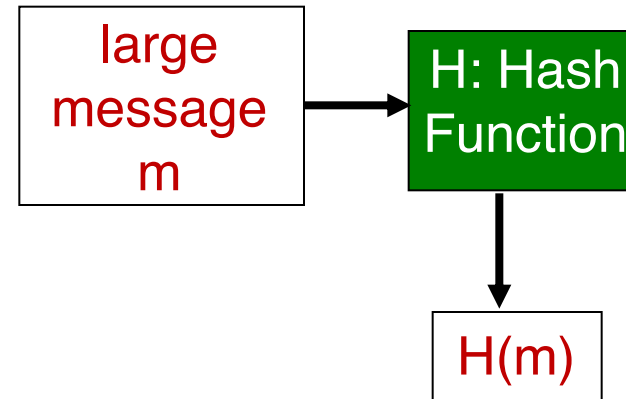
Did messages get across without tampering?

Message digests

Can we ensure that a receiver can detect message tampering?

Idea: fixed-length, easy-to-compute digital “fingerprint” of a message

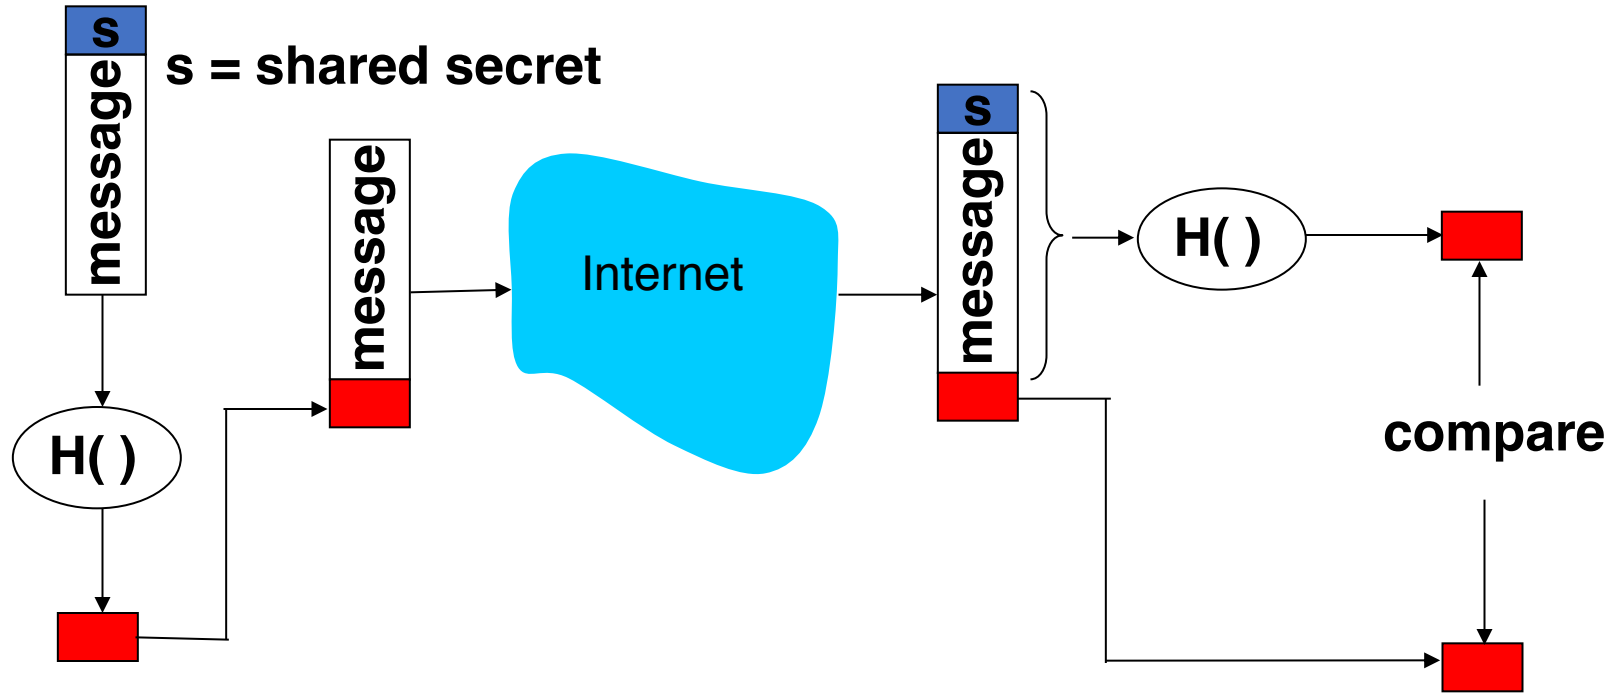
- apply hash function H to m , get fixed size message digest, $H(m)$.



Cryptographic hash function properties:

- Easy to calculate
- Produces fixed-size msg digest (fingerprint)
- Hard to reverse: given msg digest x ,
 - computationally infeasible to find m such that $x = H(m)$
 - Or another m' such that $H(m) = H(m')$

Using message digests for integrity



- Verifies message integrity
- Requires a secret shared key
- No encryption

Message digest algorithms

- You'll see the term "MAC" or Message Authentication Codes
 - I find it confusing (medium access); I will avoid using it.
- MD5 hash function widely used (RFC 1321)
 - computes 128-bit message digest in 4-step process.
 - arbitrary 128-bit string x , appears difficult to construct msg m whose MD5 hash is equal to x
- SHA-1 is also used
 - US standard [NIST, FIPS PUB 180-1]
 - 160-bit message digest

Digital signatures

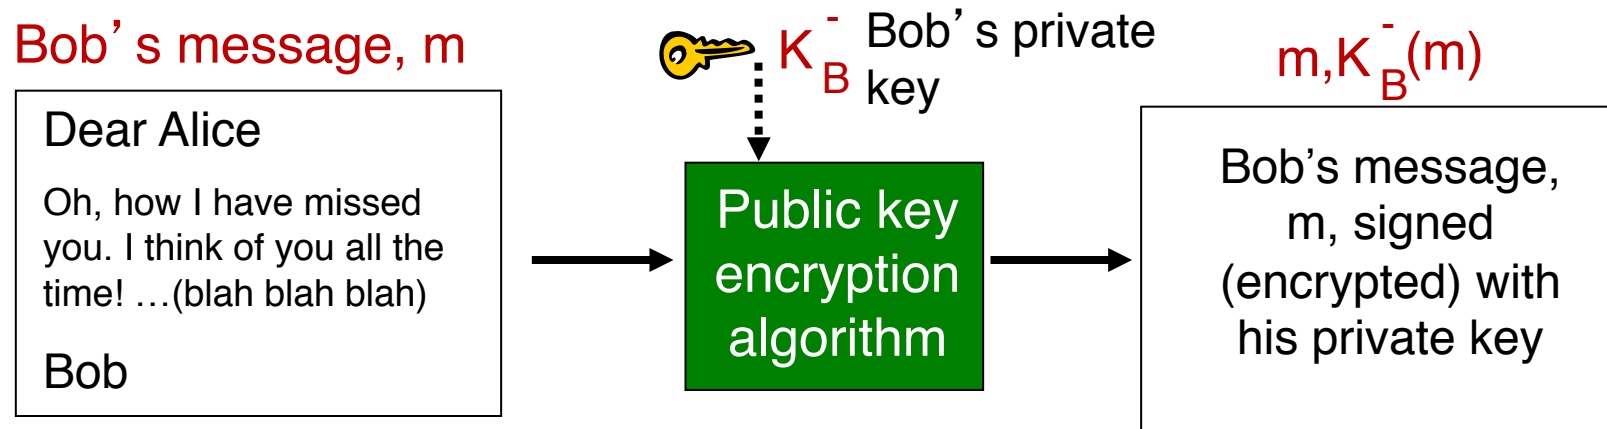
Cryptographic technique analogous to handwritten signatures:

- sender (Bob) digitally signs document, establishing he is document owner/creator.
- **verifiable, nonforgeable**: recipient (Alice) can prove to someone that Bob, and no one else (including Alice), must have signed document

Digital signatures

simple digital signature for message m :

- Bob signs m by encrypting with his private key K_B^- , creating “signed” message, $K_B^-(m)$

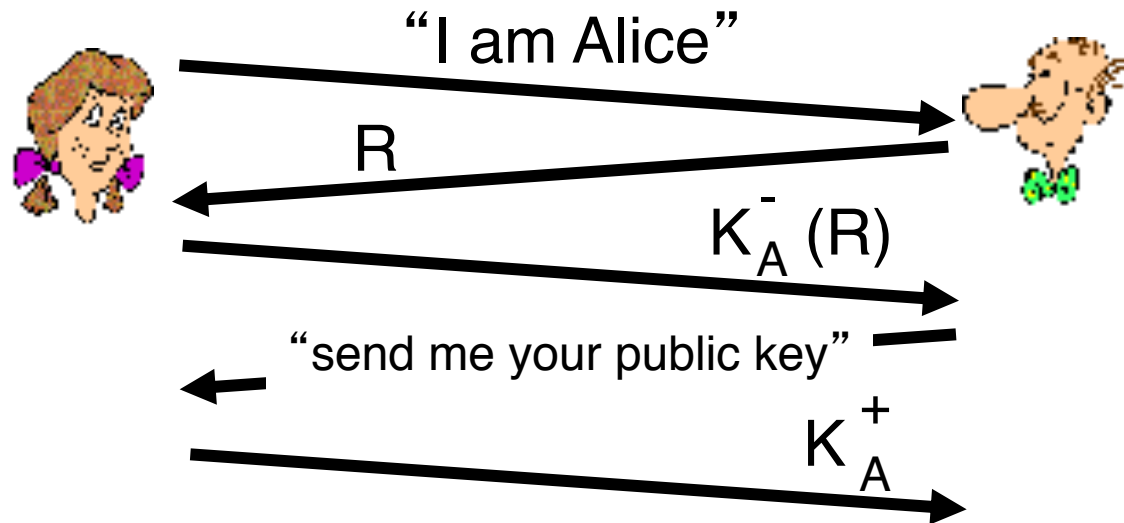


Authentication

How do I know you are who you say you are?

Authentication using public key crypto

Idea: Use **nonce** and public key cryptography



Bob computes

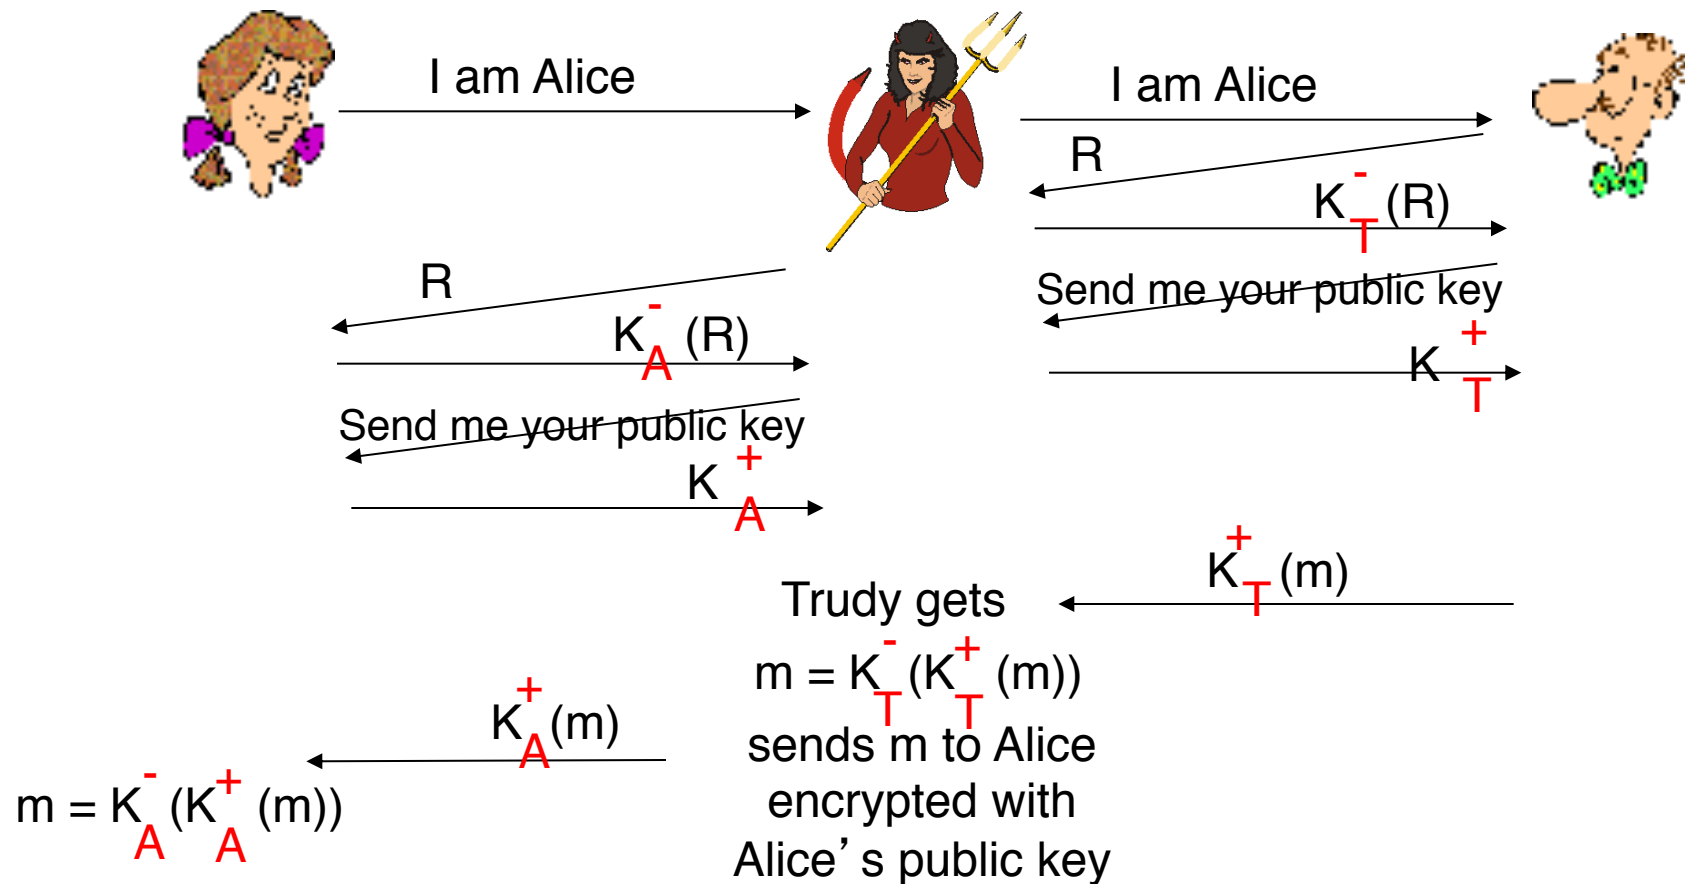
$$K_A^+(K_A^-(R)) = R$$

and knows only Alice could have the private key, that encrypted R such that

$$K_A^+(K_A^-(R)) = R$$

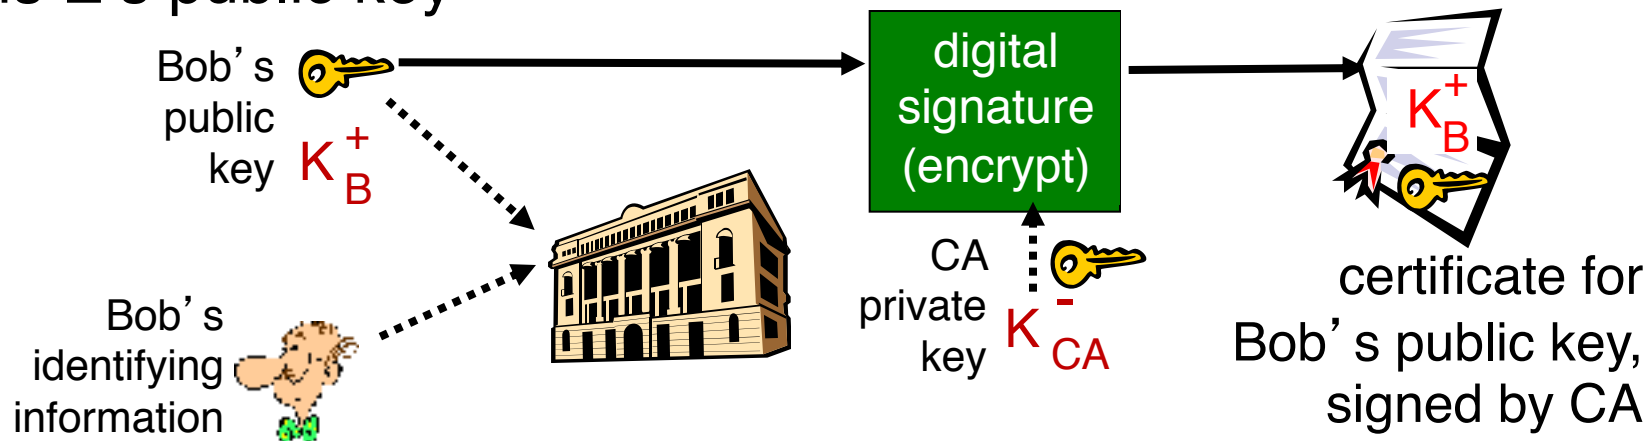
Security hole -- if you ask for public keys!

man (or woman) in the middle attack: Trudy poses as Alice (to Bob) and as Bob (to Alice)




Certification authorities

- *certification authority (CA)*: binds public key to particular entity, E.
- E (person, router) registers its public key with CA.
 - E provides “proof of identity” to CA.
 - CA creates certificate binding E to its public key.
 - certificate containing E’s public key digitally signed by CA – CA says “this is E’s public key”



Certification authorities

- When Alice wants Bob's public key:
 - gets Bob's certificate (from Bob or elsewhere).
 - apply CA's public key to Bob's certificate, get Bob's public key




Transport Layer Security (TLS)

Providing security properties in a practical protocol

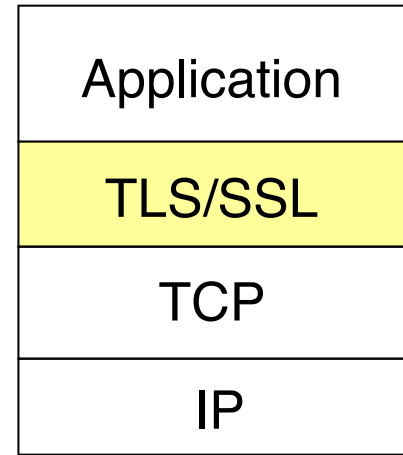
Goals of TLS

- Confidentiality
- Message integrity
- Server authentication (optionally, client authentication)
- Must work in the context of the existing network protocol stack

TLS/SSL and the rest of the protocol stack



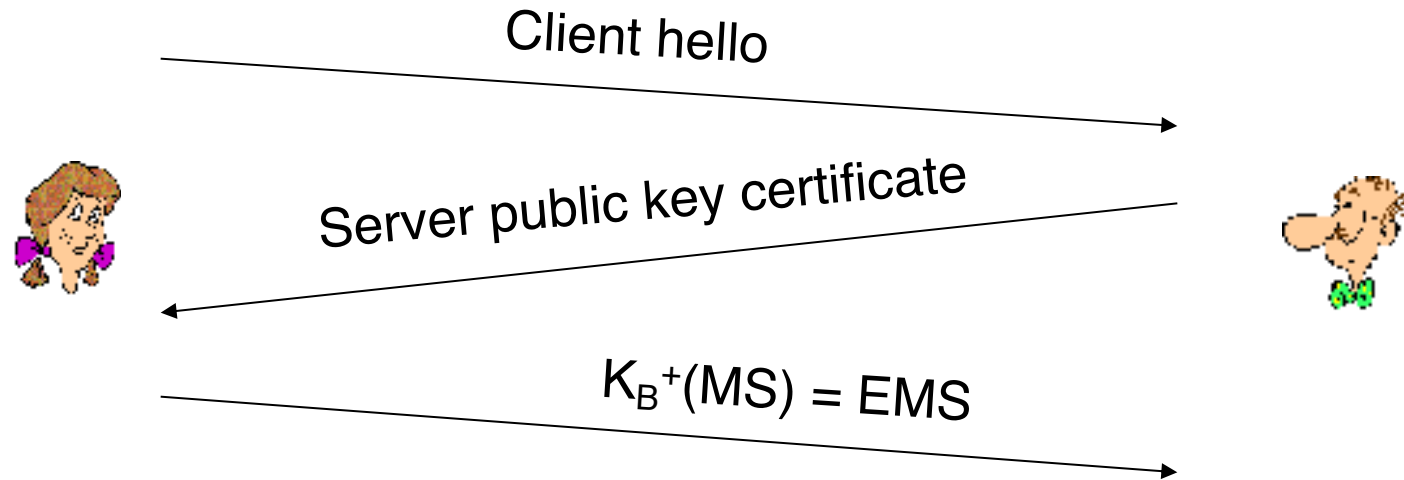
normal application



application with TLS/SSL

- TLS/SSL provides application programming interface (API) to applications
- C and Java libraries/classes readily available
 - Ex: OpenSSL

Step (1): a simple handshake



MS: master secret

EMS: encrypted master secret

Q: What all might the “master secret” be used for?

Step (2): key derivation

- considered bad to use same key for more than one cryptographic operation
 - use different keys for message integrity and encryption
- four keys:
 - K_c = encryption key for data sent from client to server
 - M_c = integrity digest key for data sent from client to server
 - K_s = encryption key for data sent from server to client
 - M_s = integrity digest key for data sent from server to client
- keys derived from key derivation function (KDF)
 - Takes master secret and (possibly) some additional random data and creates the keys

Step (3): Data records

- why not encrypt data in constant stream as we write it to TCP?
 - where would we put the message digest? If at end, no message integrity until all data processed.
 - e.g., with instant messaging, how can we do integrity check over all bytes sent before displaying?
- instead, break stream in series of records
 - each record carries a message digest
 - receiver can act on each record as it arrives
- How does receiver distinguish the digest from data within a record?
 - want to use variable-length records



TLS/SSL “Cipher Suite”

- Cipher suite
 - public-key algorithm
 - symmetric encryption algorithm
 - Integrity hashing algorithm
- TLS/SSL supports several cipher suites
- **Negotiation**: client, server agree on cipher suite
 - client offers choices
 - server picks one

Common symmetric ciphers

- AES – Advanced Encryption Standard
- DES – Data Encryption Standard: block
- 3DES – Triple strength: block
- ChaCha: stream
- RC4 – Rivest Cipher 4: stream

SSL Public key encryption

- RSA with DH

Integrity hashing algorithms

- HMAC-MD5 and others

Improved Handshake with Negotiation

1. server authentication
2. negotiation: agree on cipher suite
3. establish necessary keys
4. client authentication (optional)

All this takes a few round trip times to accomplish!

QUIC handshake

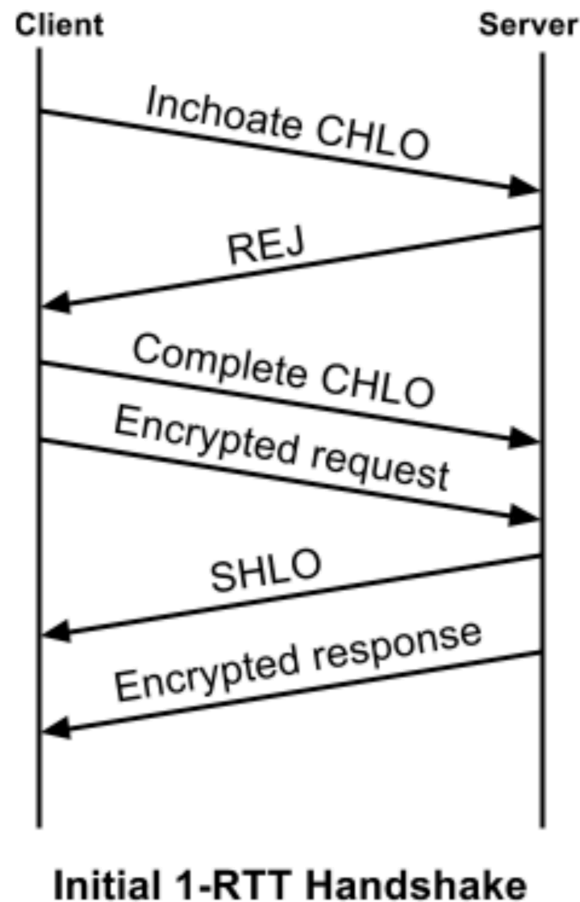
Adding security in one shot to a handshake

Goals of QUIC handshake

- Reduce round-trip times to get set up with secure connection
- Initial handshake takes 1 RTT, starting from the first byte received from the client, before server can send data
- Later, server can send data immediately upon receiving the first byte from the client (“0-RTT” handshake)

Initial 1-RTT handshake

- Client first sends an “inchoate” (incomplete) client hello message



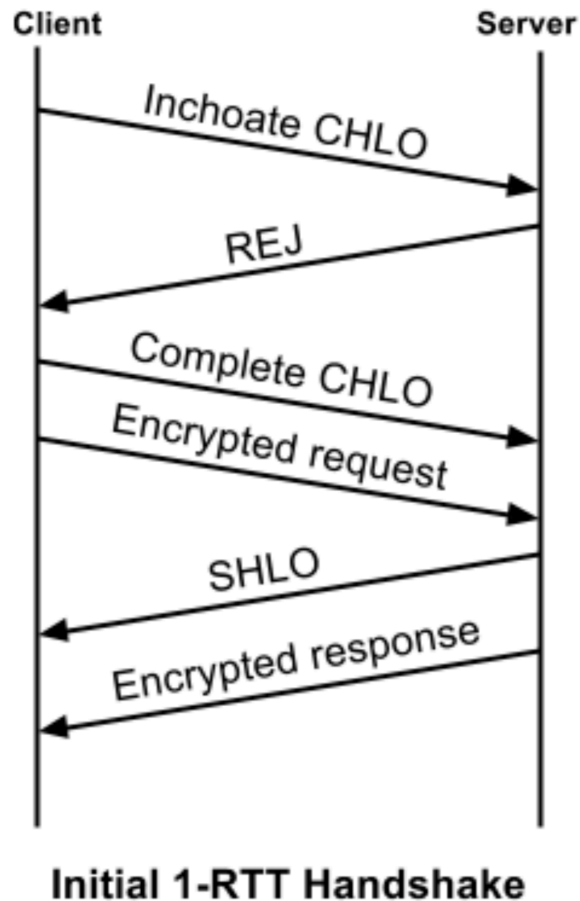
Server REJ message contains:

- (1) server's long-term DHM public value
- (2) Certificate chain authenticating the server
- (3) Signature of DHM public value using the certified (private) key
- (4) Client source address token signed by server

What properties does each of these (help) establish?

Initial 1-RTT handshake

- Finish the handshake



Client sends its ephemeral DHM public value to server in a complete CHLO.

Server then responds with its ephemeral DHM public value.

Ephemeral DH values provide **forward secrecy**

Other handshake details

- In later handshakes, client optimistically sends its source address token to the server
- Optimistic version negotiation: client proposes a cipher suite and encrypts first request with that cipher suite
 - If server can't speak that cipher suite, it forces version renegotiation, similar to TLS
- Both techniques optimize the common case when clients speak to a known server with an agreed-upon cipher suite

