

Multi-Path Congestion Control

Lecture 20, Computer Networks (198:552)

Fall 2019

Review: TCP congestion control

- Keep some **in-flight** (un-ACK'ed) packets: **congestion window**
- Adjust window based on several algorithms:
 - Startup: slow start
 - Steady state: AIMD
 - Loss: fast retransmission, fast recovery
- Classically, TCP uses a **single** path provided by the underlying network routing

If a TCP conn could use multiple paths...

- Better **resilience**
 - If one path becomes unavailable, keep traffic flowing over others
- Higher **throughput**
 - Use multiple paths to overcome single-path bottlenecks
- Seamless **mobility**
 - More paths as they become available, “handing off” as needed
- Example uses
 - Mobile phone (WiFi/cellular)
 - High-end servers (multiple NICs)
 - Data centers (many paths available)


Goal: Do all this without application-level changes

TCP throughput equation

Steady-state behavior


Model of single-path TCP

- **AIMD:** $W := W + 1$ (RTT), $W := W/2$ (drop)
- Only a single flow using the bottleneck link




Model of single-path TCP

- Idealization: Repeat **exactly same window evolution** over multiple epochs of a **single packet loss**




Model of single-path TCP

- Window goes from W_{\min} to $2 * W_{\min}$
- Loss rate p : number of packets dropped per packet sent




Model of single-path TCP

- Loss rate assumed to be **independent of sending rate**
 - In reality: more you send, more links traversed, more you drop
- Loss assumed **deterministic** (e.g., buffer full)
 - In reality: AQM, stochastic channel loss (e.g., cellular)



Model of single-path TCP

- Goal: Find TCP's throughput as a function of link rate, RTT, and packet loss rate
- Throughput = (#packets sent per epoch) / (time per epoch)



TCP throughput equation


- Throughput depends **inversely on sqrt of loss rate**, $p^{-1/2}$
 - Throughput drops steeply with increasing loss rate
 - Ideal: want it to be linear drop, ie: $C \cdot (1-p)$
- Throughput depends inversely on the RTT
 - An issue known as **RTT unfairness**: lower-RTT connection on a bottleneck gets higher throughput than higher-RTT connection
 - Ideal: independent of RTT
- Is throughput independent of the link rate and buffer size?

Multipath TCP

Design of the congestion control algorithm

Slides adapted from Damon Wischik


Goal #1: Fairness at Shared Bottlenecks



Why not just open multiple TCP connections?

To be fair, Multipath TCP should take as much capacity as TCP at a bottleneck link, no matter how many paths it is using.


Goal #2: Use Efficient Paths



Each flow has a choice of a 1-hop and a 2-hop path.

How should each flow split its traffic?


Use Efficient Paths



Equal-window TCP (EWTCP): split flow traffic 1:1 over paths


Achieve fairness using $W_s := W_s + a \text{ (RTT)}$, $a < 1$

Use Efficient Paths




Move some traffic to better paths:
What if each flow split its traffic 2:1?

Use Efficient Paths



Best: Each connection on its one-hop path
→ Least congested path

Use Efficient Paths




Goal: Get each flow to send all its traffic on the least-congested path

Achieve by **coupling** the subflow TCP window updates

$$W_s := W_s + 1/W_{\text{total}} (\text{ack}), \quad W_s := W_s - W_{\text{total}}/2 (\text{drop})$$

Use Efficient Paths




$$W_s := W_s + 1/W_{\text{total}} (\text{ack}), \quad W_s := W_s - W_{\text{total}}/2 (\text{drop})$$

Consequence: the more drops a subflow sees, the faster its window W_s reduces (note: increments same across all paths)

More lossy paths have zero window in the limit

Use Efficient Paths




$$W_s := W_s + 1/W_{\text{total}} (\text{ack}), \quad W_s := W_s - W_{\text{total}}/2 (\text{drop})$$

Consequence: Remaining paths have balanced loss rate → **equal window** if RTTs same

If loss not balanced, window would drop to zero


Coupled CC can get trapped



Keep a little traffic on each path (even if congested) to
probe for capacity always

i.e., keep your options open 😊

Coupled CC can get trapped



Semi-coupled TCP

$$W_s := W_s + \textcolor{red}{a}/W_{\text{total}} \text{ (ack)}, W_s := W_s - \textcolor{red}{W_s}/2 \text{ (drop)}$$

Compare with coupled:

$$W_s := W_s + 1/W_{\text{total}} \text{ (ack)}, W_s := W_s - W_{\text{total}}/2 \text{ (drop)}$$

Goal #3: Be Fair Compared to TCP

- Least-congested paths may not be best
 - Low loss rate, but low throughput due to differences in RTT
- Example: Two paths
 - WiFi: high loss, low RTT
 - Cellular: low loss, high RTT
- Using the least-congested path
 - Choose the cellular path, due to low loss
 - But, the RTT is high
 - So throughput is low!
- Formalize fairness requirement using actual throughput

Be Fair Compared to TCP

- To be fair, Multipath TCP should give a connection **at least as much throughput as it would get with a single-path TCP** on the best of its paths, **given the same loss rate**
 - Ensure incentive for deploying MPTCP
- A Multipath TCP should **take no more capacity on any path (or collection of paths) than if it was a single-path TCP flow** using the best of those paths, given the same loss rate
 - Do no harm

Achieving These Goals

- Regular TCP
 - Maintain a congestion window w
 - On an ACK, increase by $1/w$ (increase 1 per window)
 - On a loss, decrease by $w/2$
- MPTCP
 - Maintain a congestion window per path w_s
 - On an ACK on path s , increase w_s
 - On a loss on path s , decrease by $w_s/2$
- How much to increase w_s on an ACK?
 - If s is the only path at that bottleneck, increase by $1/w_s$

If Multiple Paths Share Bottleneck?

- Don't take any more bandwidth on a link than the best of the TCP paths would
 - But, where might the bottlenecks be?
 - Multiple paths might share the same bottleneck
 - This is hard to know across the Internet
- So, consider all possible subsets of the paths
 - Set R of paths
 - Subset S of R that includes path r
- E.g., consider path 3
 - Suppose paths 1, 3, and 4 share a bottleneck
 - ... but, path 2 does not
 - Then, we care about $S = \{1,3,4\}$

Achieving These Goals

- What is the *best* of these subflows achieving?
 - Path s is achieving throughput of w_s/RTT_s
 - So best path is getting $\max_s(w_s/\text{RTT}_s)$
- What *total* bandwidth are these subflows getting?
 - Across *all* subflows sharing that bottleneck
 - Sum over s in S of w_s/RTT_s
- Consider the *ratio* of the two
 - Increase by less if many subflows are sharing
- And pick the results for the set S with min ratio
 - To account for the *most* paths sharing a bottleneck

$$\frac{\max_{s \in S} w_s / \text{RTT}_s^2}{\left(\sum_{s \in S} w_s / \text{RTT}_s \right)^2}$$

MPTCP Implementation

Implementation Issues

- Buffer space: per-subflow or shared for entire connection?
- Reassembly across multiple paths
 - Different sequence spaces across subflows
 - But shared flow control
 - Ensure packets across subflows reach at approx. the same time
- Middleboxes
 - Avoid impact due to rewritten sequence numbers
- Initiating new subflow: new TCP flag; auth token