

# Internet Architecture

# Computer Networks

Collections of machines exchanging information  
with each other

Principles and algorithms by which communication  
software & hardware are organized




Design of foundational artifacts of the Internet  
and other modern networks

# There's a science to communication

- What language should machines speak?
- How to partition functionality? “Who” should do “what”?
- How to make communication effective?
  - Achieve better scale, performance, resource efficiency
  - Evolve to address new needs over time
- How to grow organically? include more communicating parties
  - Make it easy for humans to build and manage?
- How to make communication worthy of societal trust?

# A brief history of Internet Architecture

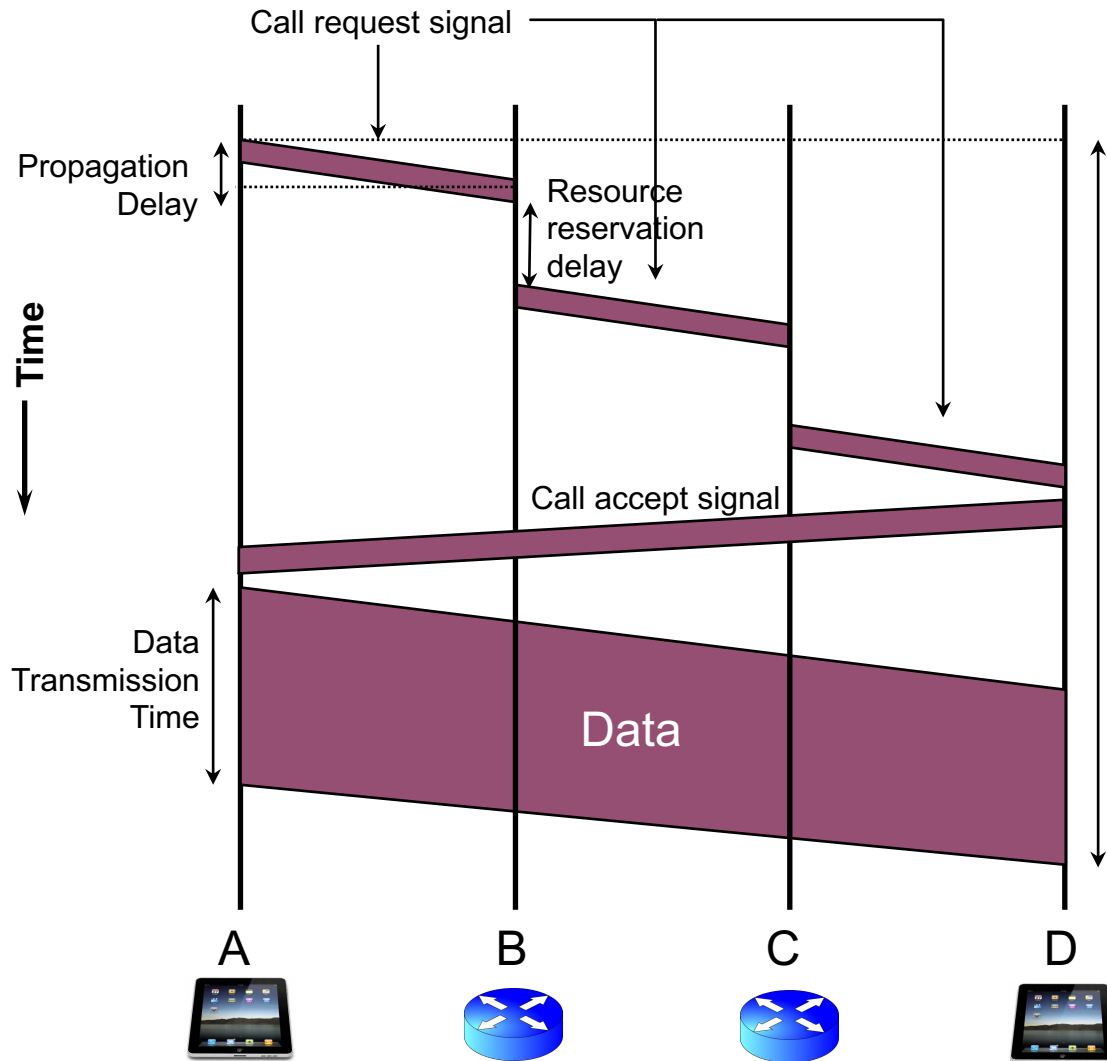
# Prior to the Internet



Different technologies:  
wireless/wired  
slow/fast,  
un/reliable,  
switching techniques  
**Different administration**

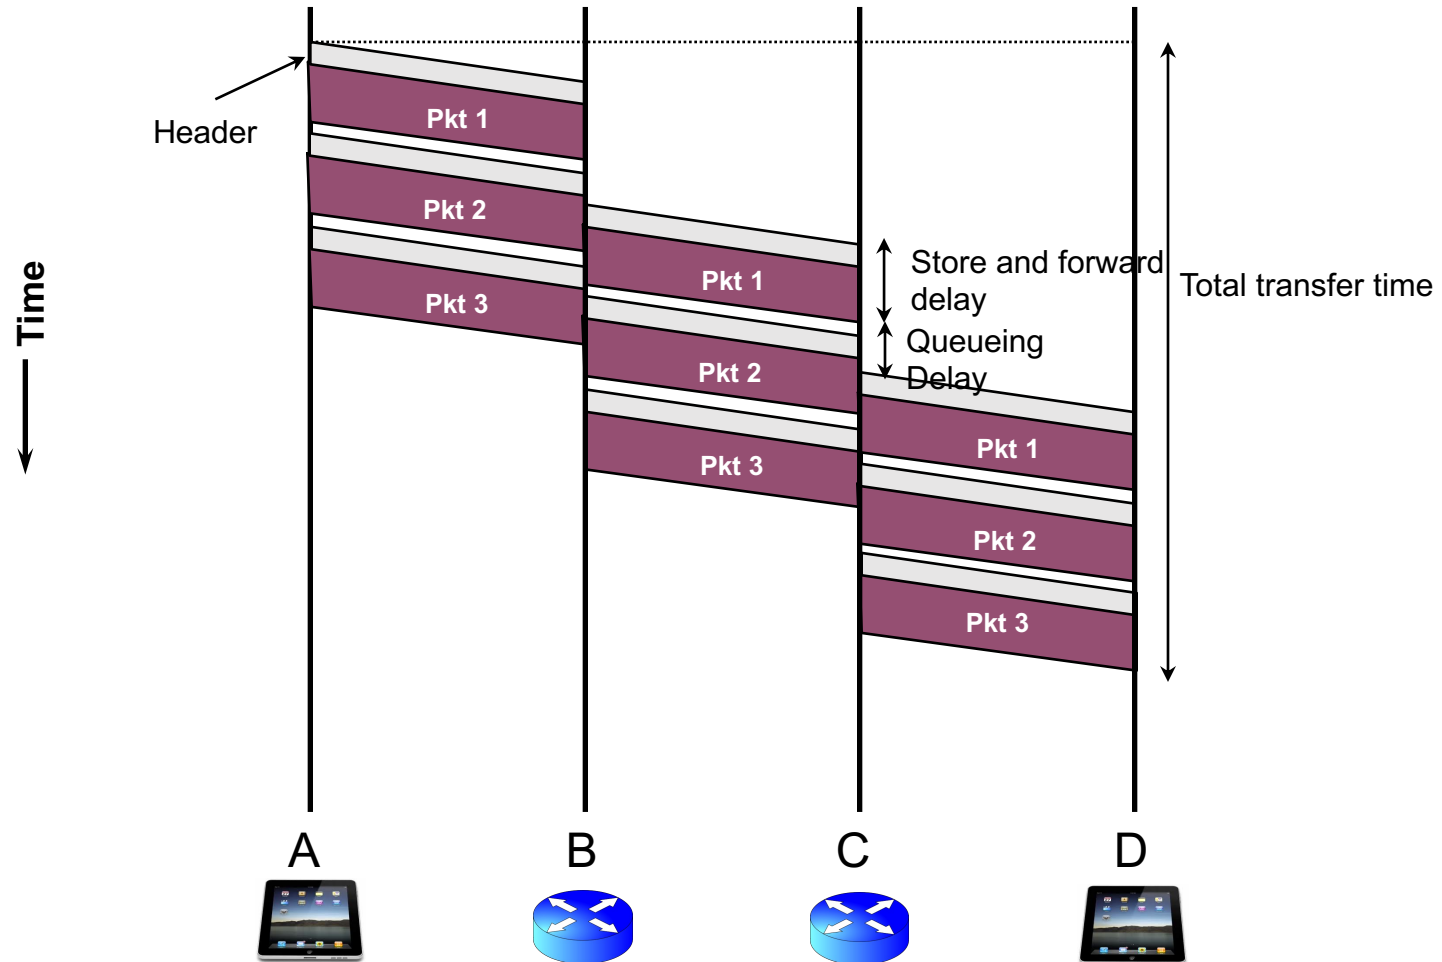
How to interconnect these existing  
networks?  
Focus more on practicality and  
usefulness rather than “clean design”

# Circuit switching



1. **Setup:** Control message sets up a path from origin to destination
2. Accept signal informs source that data transmission may proceed
3. **Data transmission** begins
4. Entire path remains allocated to the transmission (whether used or not)
5. When transmission is complete, source releases the circuit


# Packet switching



Packet switching is simpler and required less from existing networks that were interconnected.



# Survivability



Should applications care?

**No**

Who has the “context” of the communication? Endpoint vs. Network

**Tradeoffs:**

Replicate, still can't guarantee  
Instead, **fate sharing**. Simpler to engineer

The Internet puts the context of conversations in the endpoints. Call this **transport** information.

Networks provide **datagram** service.

Endpoints are **trusted** to implement the right algorithms to provide any sort of guarantees about communication.

# Distributed Management



Don't require single administrative entity.  
The Internet is federated.

Two-tier management: **within** an  
administration, **across** administrations

# Consequences of datagram service

- A simple building block for different kinds of applications:
  - Bulk file transfer
  - Conversational (real time)
  - Try very hard not to constraint what can be implemented atop network
- No explicit guarantees required from interconnecting networks
  - Accommodate a variety of existing networks
- Services explicitly not assumed from the network:
  - reliable or sequenced delivery, network level broadcast or multicast, priority ranking of transmitted packet, support for multiple types of service, and internal knowledge of failures, speeds, or delays.
  - Only see packets, no higher-level context

# Layering and Protocols

# Software/hardware organization at hosts



Communication functions broken up and “stacked”

Each layer depends on the one below it.

Each layer supports the one above it.

The interfaces between layers are well-defined and standardized.

Internet software and hardware  
are arranged in **layers**.

Layering provides **modularity**

Each layer: well-defined **function**  
& **interfaces** to layers above & below it.

Functionality is implemented in **protocols**.



# Protocols: The “rules” of networking



- Protocols consist of two things
- **Message format**
  - structure of messages exchanged with an endpoint
- **Actions**
  - operations upon receiving, or not receiving, messages
- Example of a Zoom conversation:
  - Message format: English words and sentences
  - Actions: when a word is heard, say “yes”; when nothing is heard for more than 3 seconds, say “can you hear me?”







# The protocols of the Internet

- Standardized by the Internet Engineering Task Force (IETF)
  - through documents called **RFCs** (“Request For Comments”)

- **Layering of protocols**













Packet starts as an app "payload"







Packet takes on **headers** (metadata) at each layer







Routers have network and link layers too!



# Layering


- Communication over the Internet is a complex problem.
- Layering simplifies understanding, testing, maintaining
- Easy to improve or replace protocol at one layer without affecting others

# Many open and partially solved problems

- Resource management in distributed routing
- Not cost effective: small packets have high header overheads
- Retransmitting lost packets end to end can be inefficient (network recovery could have made things simpler)
- Cost of attaching a host to the network is somewhat high: all necessary software must reside on the host (e.g., transport).
- No inherent network mechanisms to account for or control resource usage
  - e.g., Putting important packets ahead in queues

Some fundamental problems


# Some definitions





- The Internet is an example of a computer network
- **Endpoint or Host:** Machine running user application
- **Packet:** a unit of data transmission (ex: 1500 bytes)
- **Link:** physical communication channel between two or more machines
- **Router:** A machine that processes packets moving them from one link to another towards a destination
- **Network:** Collection of interconnected machines
- **Address:** a unique name given to a machine



# (0) Naming & Resolution



- Communication requires **naming** the endpoints
  - **Addresses**
- Internet addresses (IP addresses) allocated **hierarchically**
  - Machine readable, not easy for humans to remember
- Link addresses are tied to the hardware on the endpoint
- **Name resolution:** how to turn human-readable names (google.com) into routable addresses?

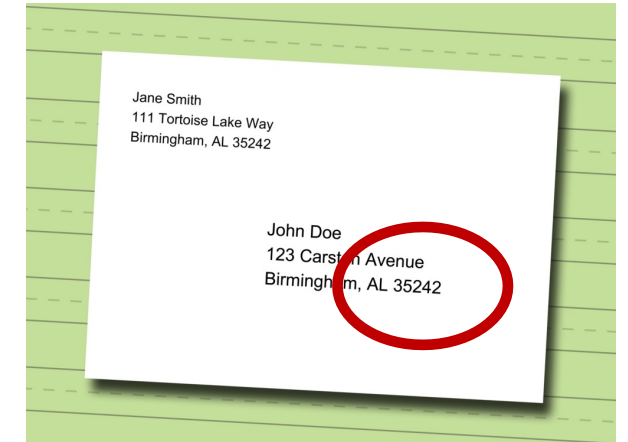


Machines communicate using **IP addresses** and **ports**

IP addresses: ~12 digits (IPv4) or more

Ports: fixed based on application (e.g., 80: web)

Need a way to turn human-readable addresses into Internet addresses.



Ask someone

**Directory service**

Ask everyone

**Query broadcast**

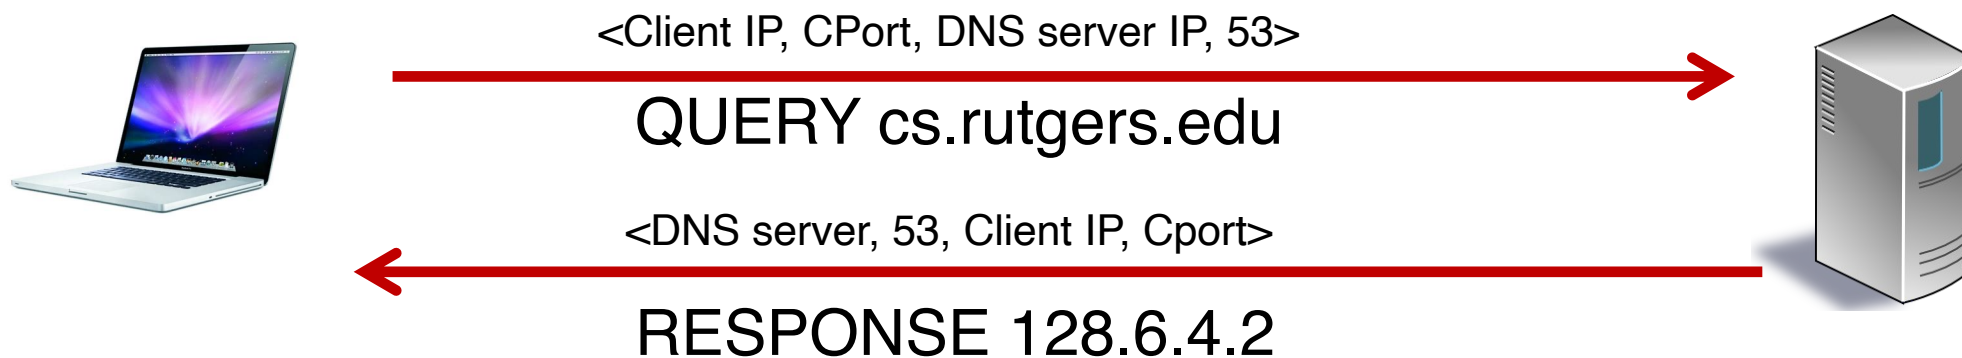
Tell everyone

**Information flooding**

Asking “someone” could involve asking many machines...


# Domain Name Service

DOMAIN NAME	IP ADDRESS
spotify.com	98.138.253.109
cs.rutgers.edu	128.6.4.2
www.google.com	74.125.225.243
www.princeton.edu	128.112.132.86




- Key idea: Implement a **server** that looks up a table.
- Will this scale?
  - Every new (changed) host needs to be (re)entered in this table
  - Performance: can the server serve billions of Internet users?
  - Failure: what if the server or the database crashes?
  - Security: What if someone “takes over” this server?

# Distributed and hierarchical database



# DNS name resolution


- Host at cs.rutgers.edu wants IP address for gaia.cs.umass.edu
- Local DNS server
- Root DNS server
- TLD DNS server
- **Authoritative** DNS server



# Example DNS interactions

- `dig <domain-name>`
- `dig +trace <domain-name>`
- `dig @<dns-server> <domain-name>`

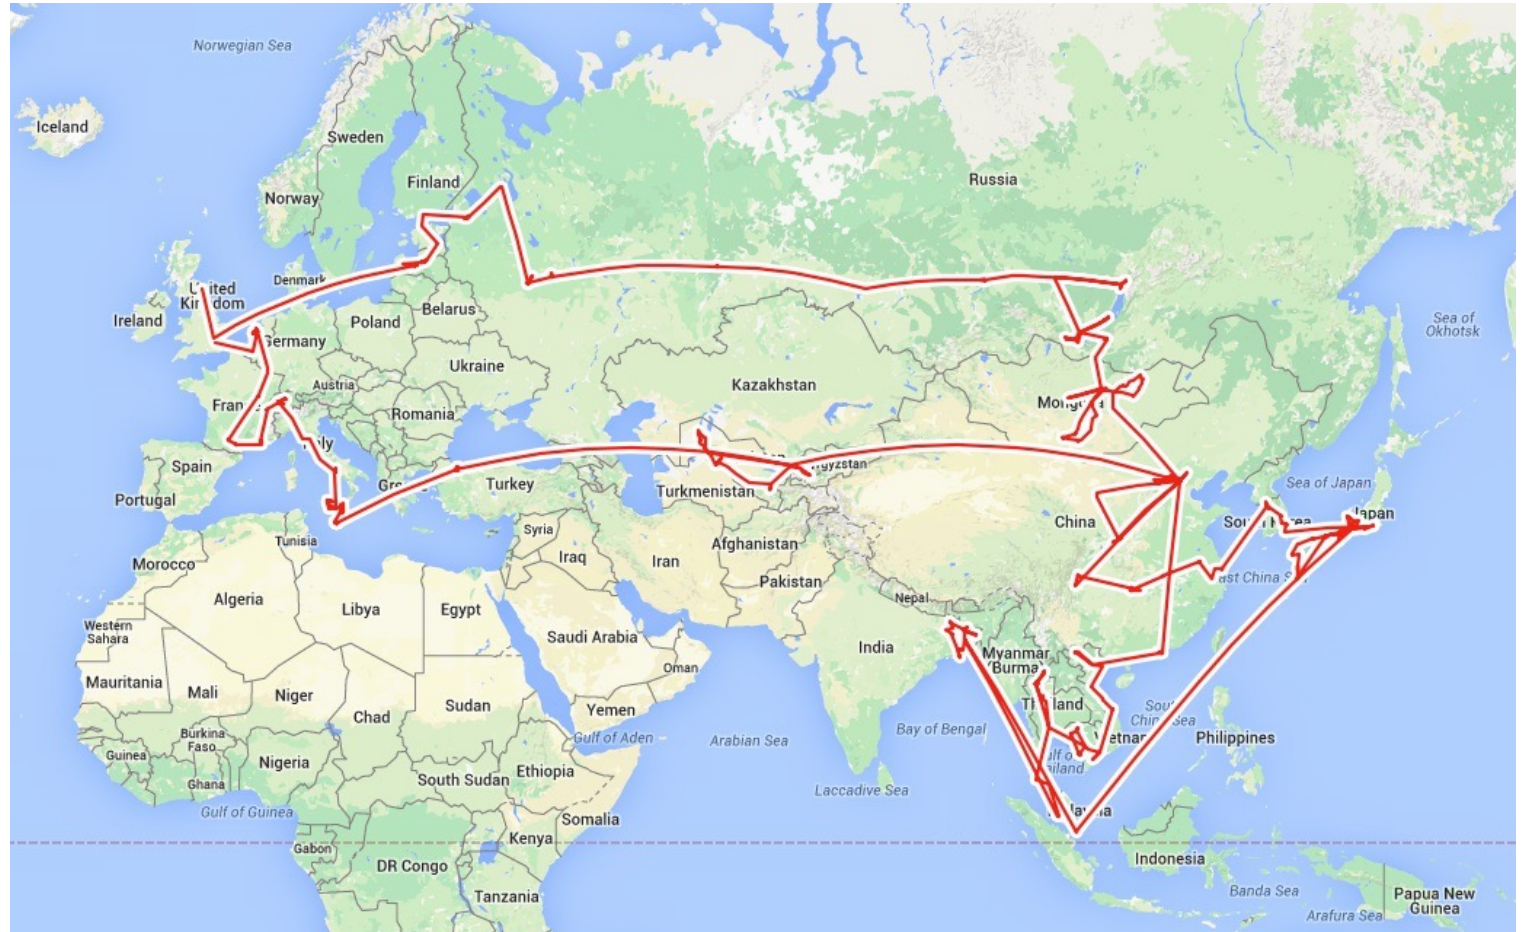
# (1) Routing



- Networks must move data between different hosts
- Need to figure out how to move packets from one host to another host, e.g., how to reach google.com from your laptop
- Known as the **routing problem**



# Routing



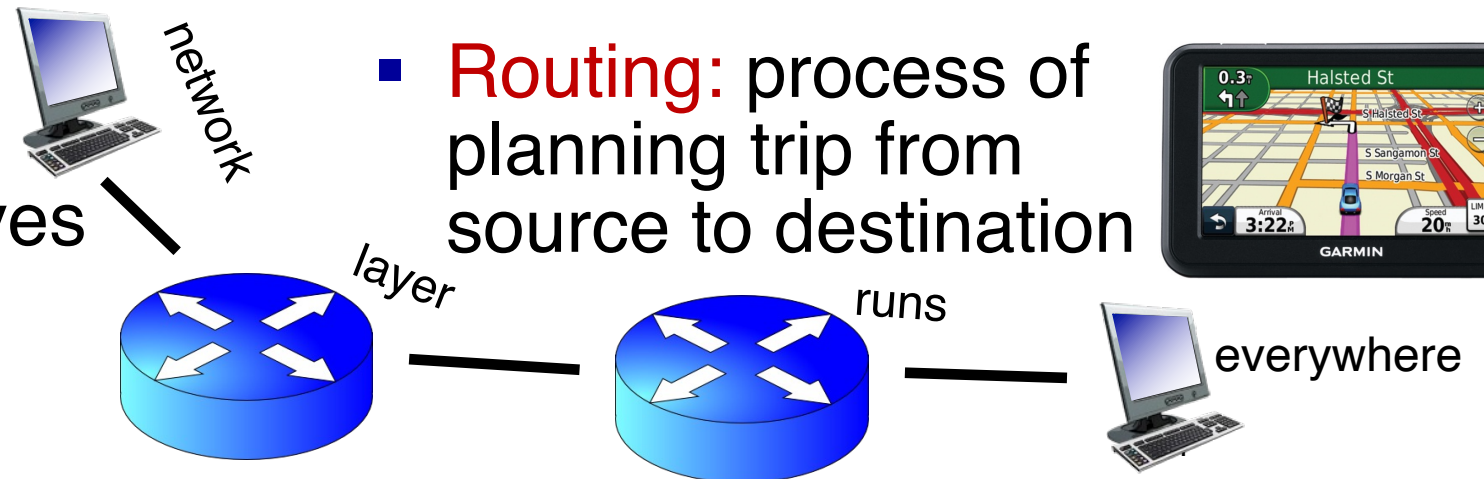


# Two key network-layer functions

- **Forwarding:** move packets from router's input to appropriate router output
- **Routing:** determine route taken by packets from source to destination
  - routing algorithms
- The network layer solves the routing problem.

Analogy: taking a road trip

- **Forwarding:** process of getting through single exit
- **Routing:** process of planning trip from source to destination




# Control/Data Planes

## Data plane = Forwarding

- local, per-router function
- determines how datagram arriving on router input port is forwarded to router output port

values in arriving  
packet header



## Control plane = Routing

- network-wide logic
- determines how datagram is routed along end-to-end path from source to destination endpoint
- two control-plane approaches:
  - **Distributed routing** algorithm running on each router
  - **Centralized routing** algorithm running on a (logically) centralized machine