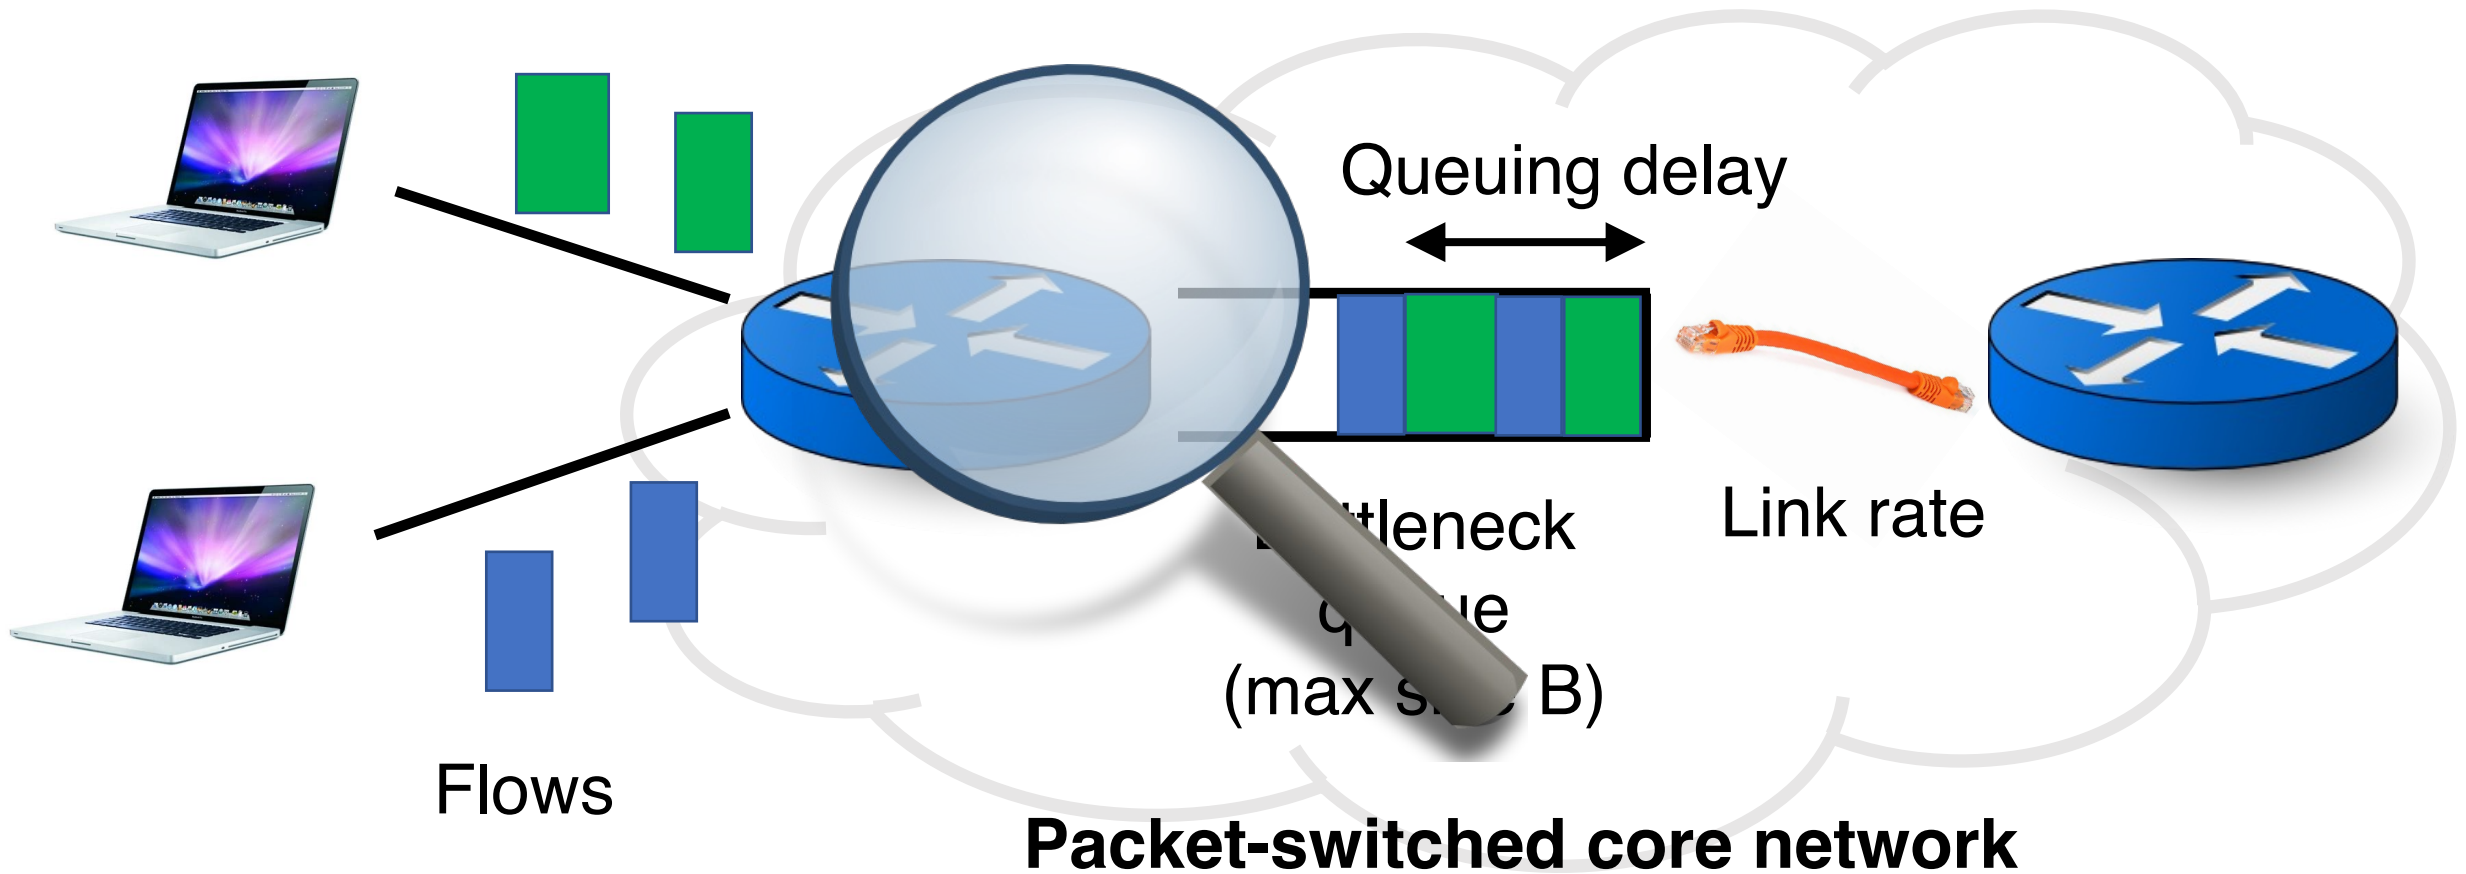


Network


What's in a router?



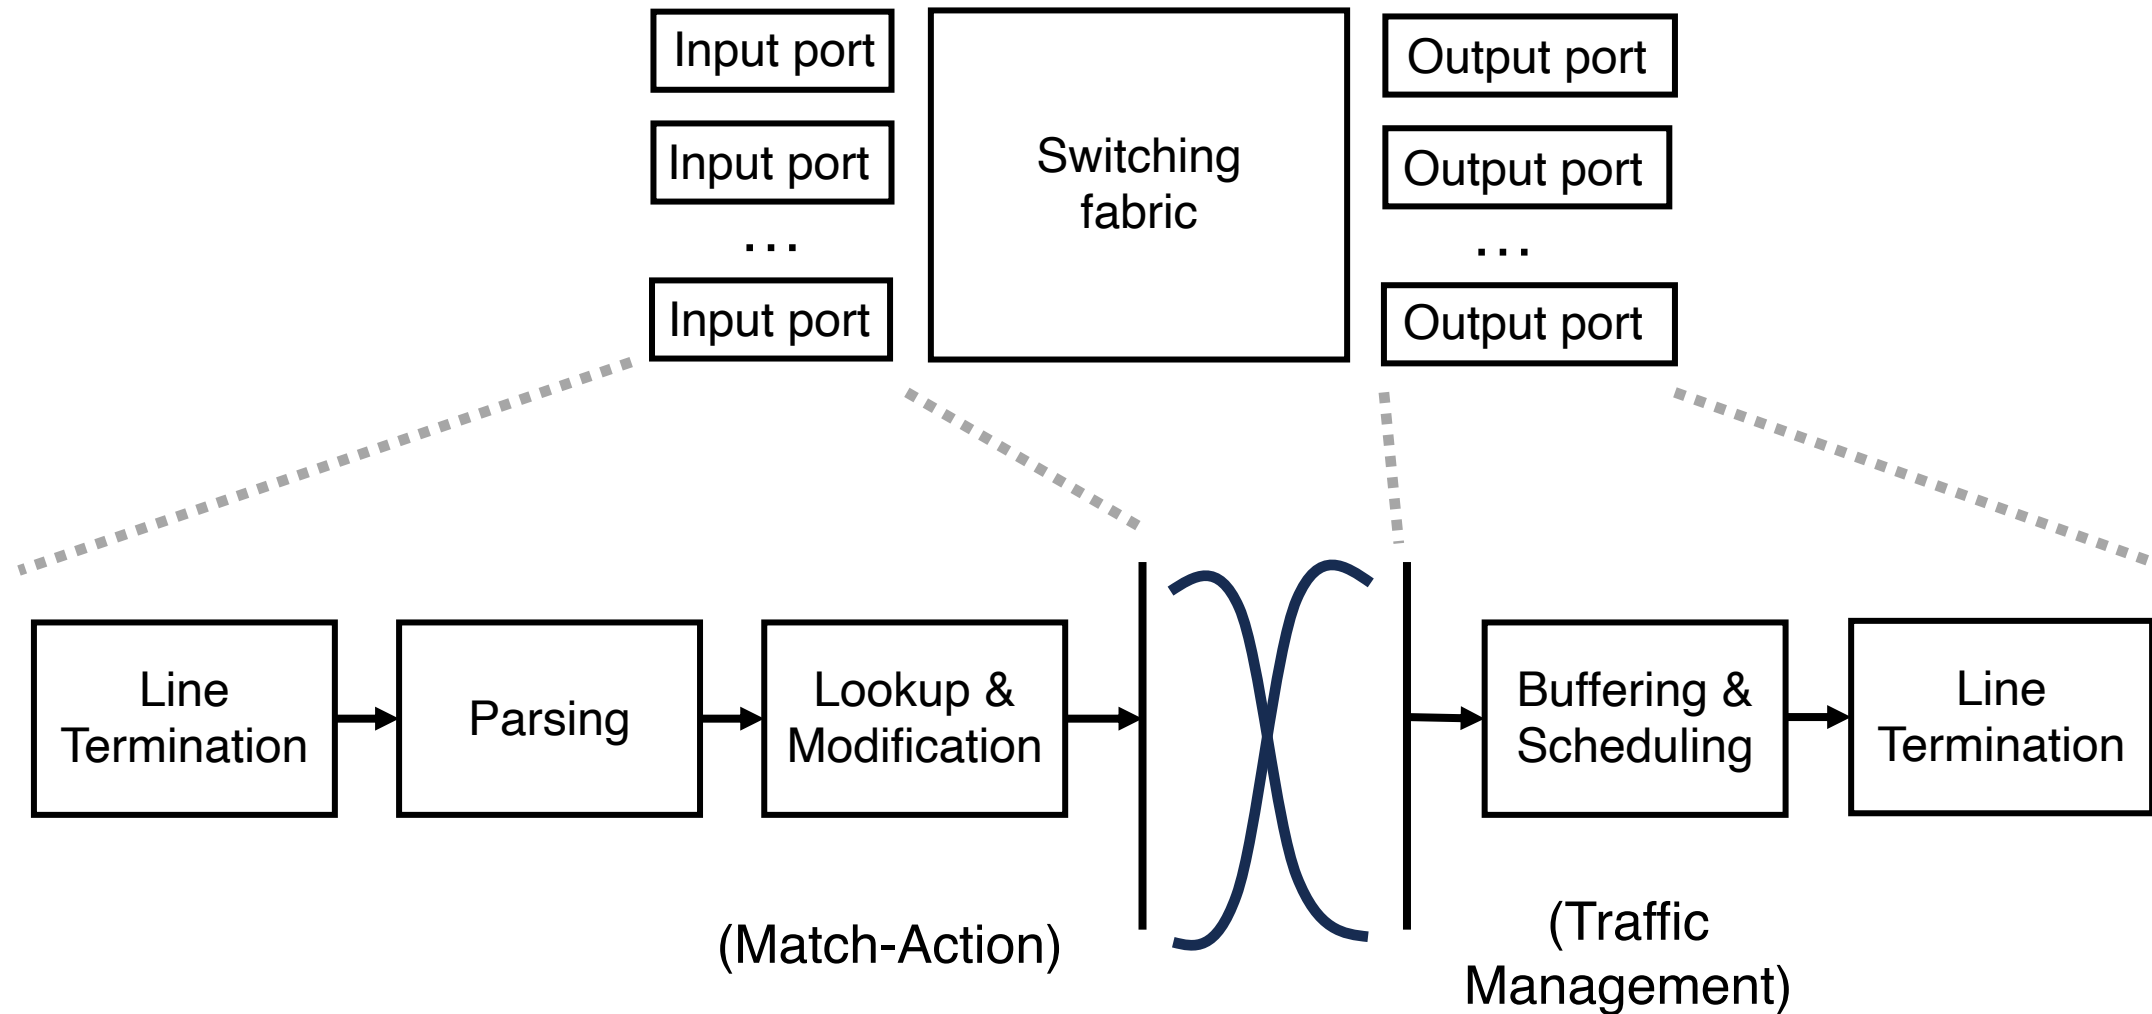
Control & Data Planes inside a router

Control plane
per route-change
processing
(~ a few seconds)

Data plane
per-packet processing
(~ tens of
nanoseconds)



Hardware Router overview



Study 2 designs

- Historically evolving, multiple concurrent router designs
- 2 exemplars:
 - **MGR**: router from the late 1990s (50Gbit/s router, Partridge et al)
 - **RMT**: router from the late 2010s

Life of a Packet

(2) Packet parsing

- Dividing a sentence into its grammatical parts:
 - “I ate an apple”
 - Sentence := Subject (I) Verb (ate) Object (an apple)
 - Object:= Article (an) Noun (apple)
- Packet parsing: dividing sequence of bits into header fields
 - 1100100101011
 - ➔ ethernet destination (1100) | source (1001) | ethertype (01) | ...
 - Unlike parsing English, parsing packets is quite **mechanical**
 - Series of **extractions** and **branches** to assign **fields to packet bits**
- Parsing could be implemented in software (MGR) or hardware (RMT)

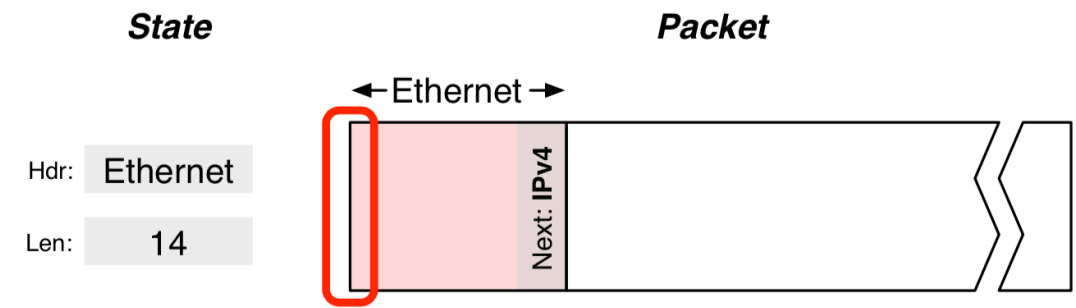
(2) Packet parsing

- A key principle: **Separate header and payload data paths**
 - Router functionality is “header-heavy” but “payload-light”
 - Don’t move the payload around too much
 - **Conserve bandwidth & resources for data moved inside the router**
- Header goes on as input to route lookup/packet modification
- Payload sits on a buffer until router knows what to do with pkt
 - Buffer could be on the ingress line card (MGR) or a buffer shared between line cards (RMT)


Parsing state machine

- Parsing is an inherently **sequential** process
- Previous header determines the next header type
- Current header length determines the start of the next header
- Parser state: tracks the current header and its length
 - Help jump to the next state


Source: Design principles for packet parsers, Gibb et al.




(a) Parsing a new packet.



(b) The Ethernet next-header field identifies the IPv4 header.



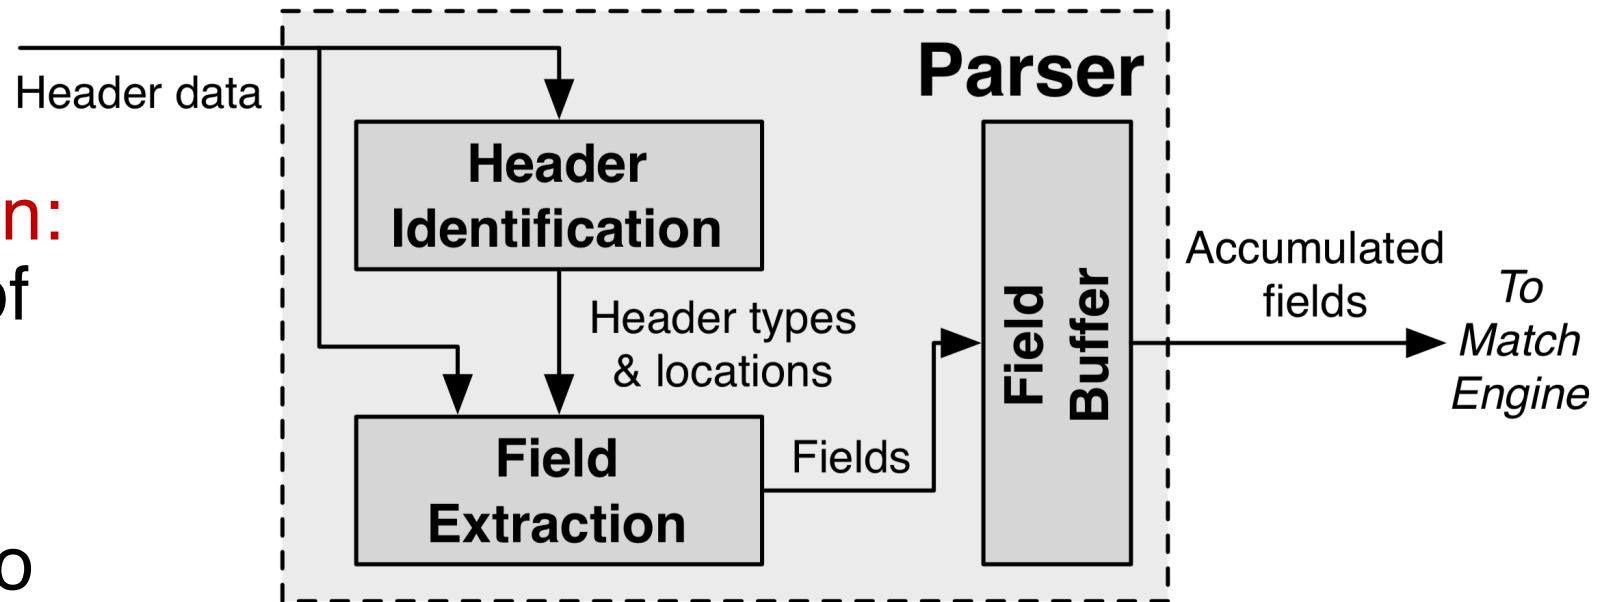
(c) The IPv4 length field identifies the IPv4 header length.



(d) The IPv4 next-header field identifies the TCP header.

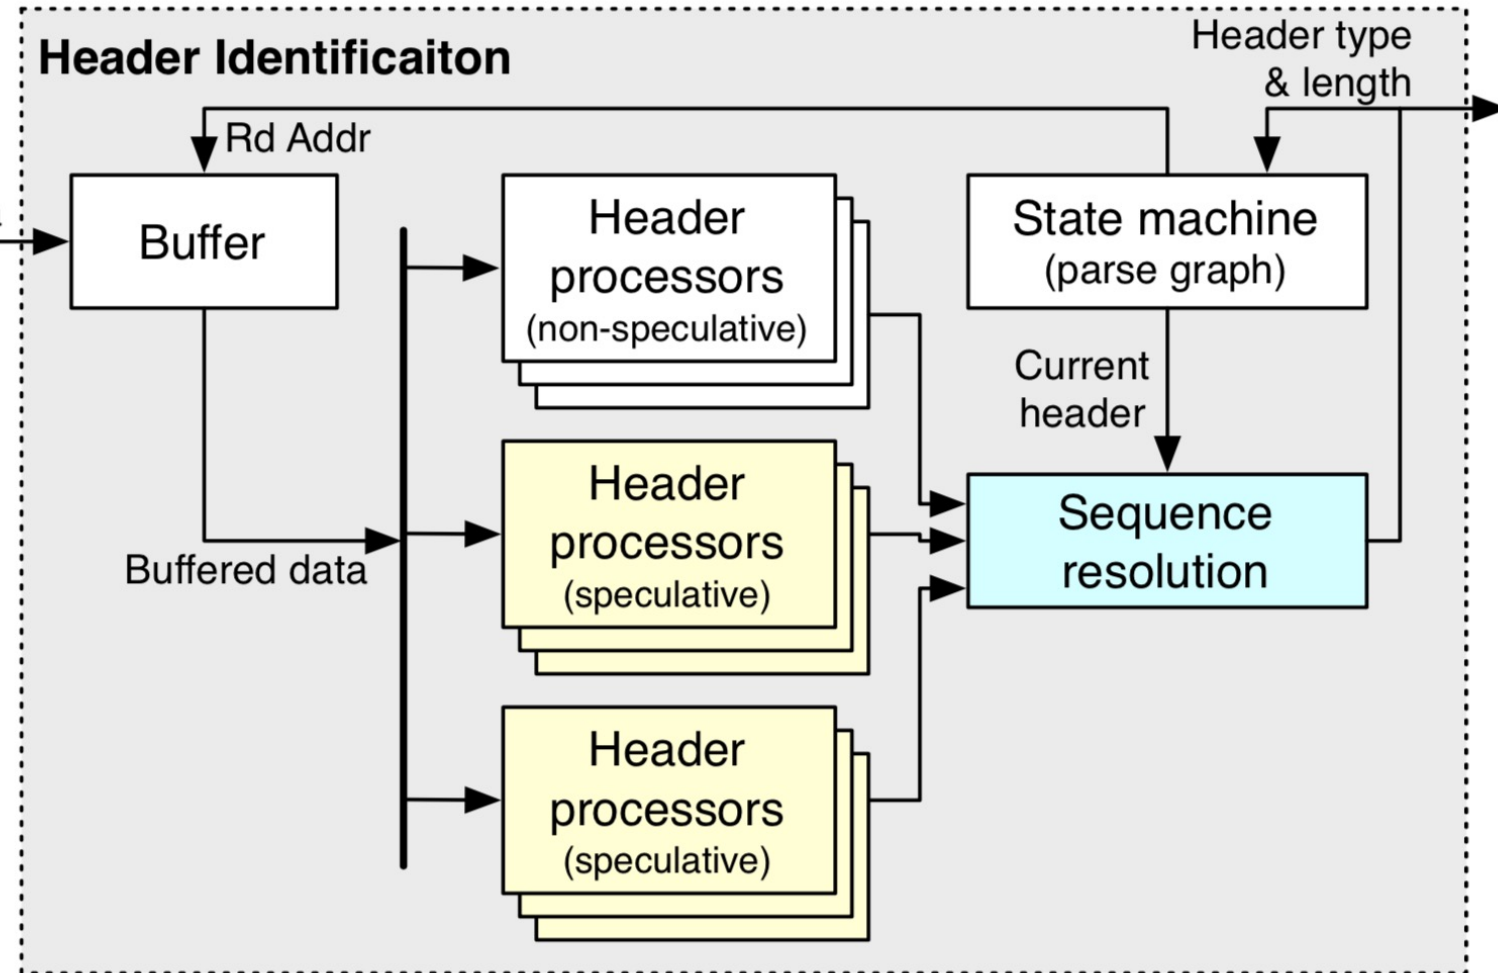
A hardware (fixed function) parser

- Two steps:
- **Header identification:**
Identify sequence of headers
- **Extract fields** from identified headers to send to matching components of tables
- Design digital circuits with a high **clock rate**?



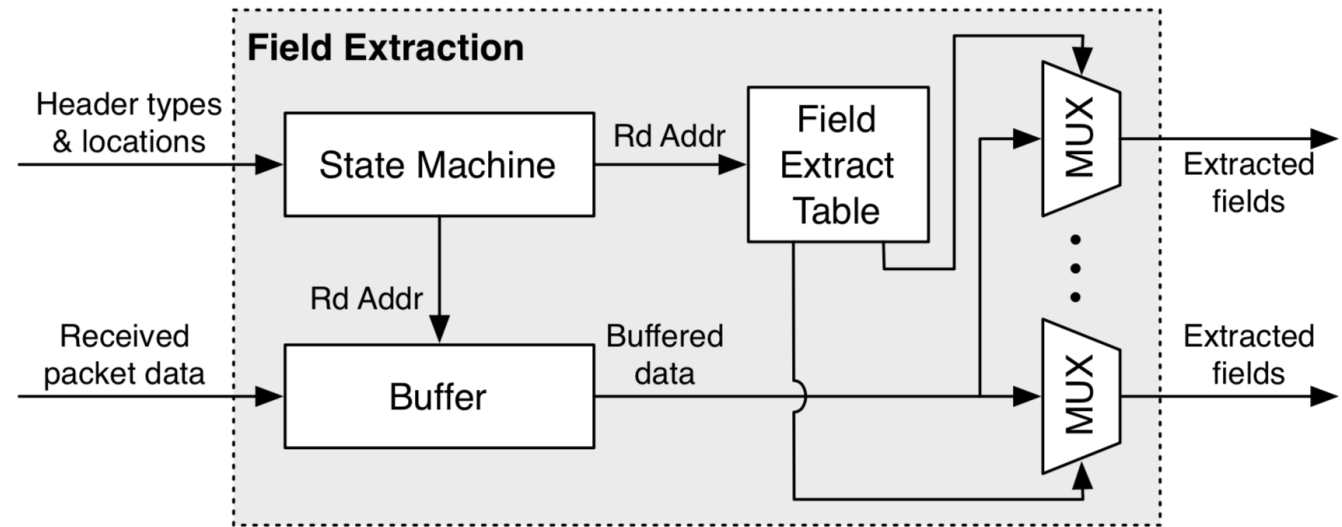
(Fixed function) header identification

- Identify headers through fixed-function **header processors**
- Simple design: extract one header per clock cycle
- **Speculate** to extract multiple headers/cycle
- Sequence resolution picks one




(Fixed function) field extraction

- Extract fields using **fixed offsets** into the packet, depending on parser state
- Field-extraction table is **hard-coded** for specific fields and protocols
 - E.g., IPv4 length: always at bytes 3 & 4



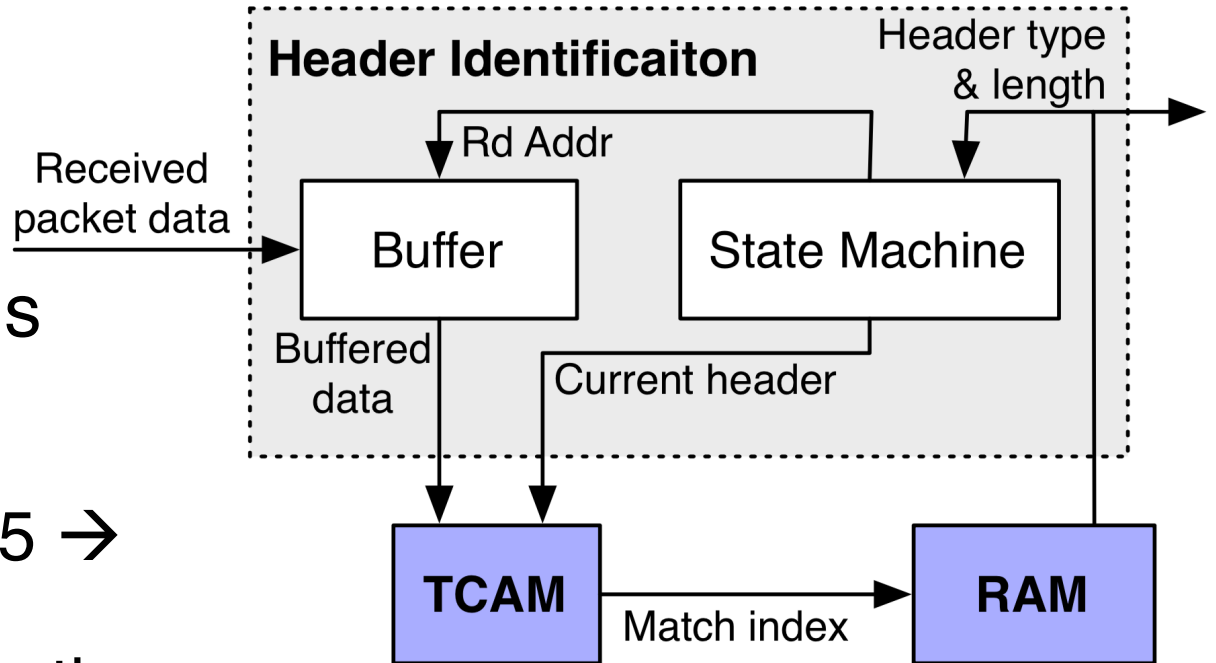
RMT makes parsing **programmable**

- **Parse graph:** Representation of **all valid** header sequences you expect



Programmable parser

- Reprogrammable **TCAM** allows us to identify headers **based on the protocol parse graph**
- Reprogrammable **SRAM** contains the protocol-dependent **field extractions**:
 - e.g., **IPv4**, byte 8 → TTL, 12–15 → IPv4 source address, etc.
 - Extracting the next header & length is a special case of generic field extractions for each protocol



Output of programmable parser

- Vector of extracted packet fields and labels. Example:
 - Ethernet -> src: 1010..; dst: 0101...; ethertype: IPv4
 - IPv4 -> version: ...; ...
 - ... *<other headers>*
- Termed the **packet header vector (PHV)**
- Fed into the packet lookup and modification engines

(3) Packet Lookup & Modification

- The main job of a router is to forward packets to the correct output port(s)
- Typically done by looking up a table of entries that were pre-computed by the control plane
- Look-up tables with a range of sizes (# entries), widths (headers)
- Packets may also need to be modified
 - RFC 1812 TTL decrement, recompute IP checksum, MAC rewrite, ...
 - Virtualized public cloud: encap/decap headers
- Outcome: a (set of) **output ports** + (possibly modified) **headers**

Typical table structures

- Sequence of tables: L2, L3, ACL
- Ethernet (L2) headers to forward packets within one IP network
- IP (L3) headers to forward packets across IP networks
- Access control lists (ACL) to implement firewalls & other policies
- Different kinds of lookups possible:
 - Exact match
 - Longest prefix match
 - Wildcard match

Packet lookup in MGR

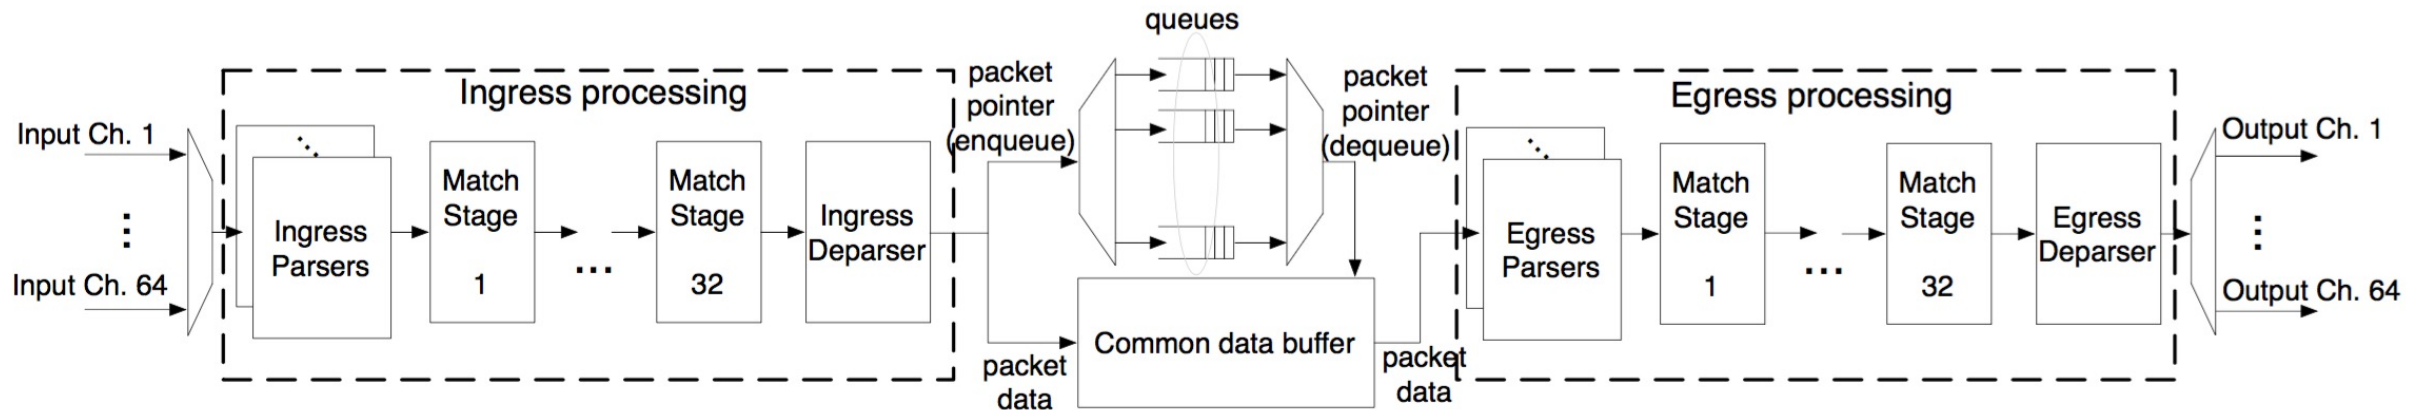
- A **forwarding engine card** separate from line cards
 - Scale forwarding and interface capacity separately
- **Software:** Use Alpha 21164 (a 415MHz generic processor)
- Programmed in assembly to do route lookup and other processing.
- Many optimizations to improve performance
 - Need for such optimizations continues today for software

MGR Software lookup performance

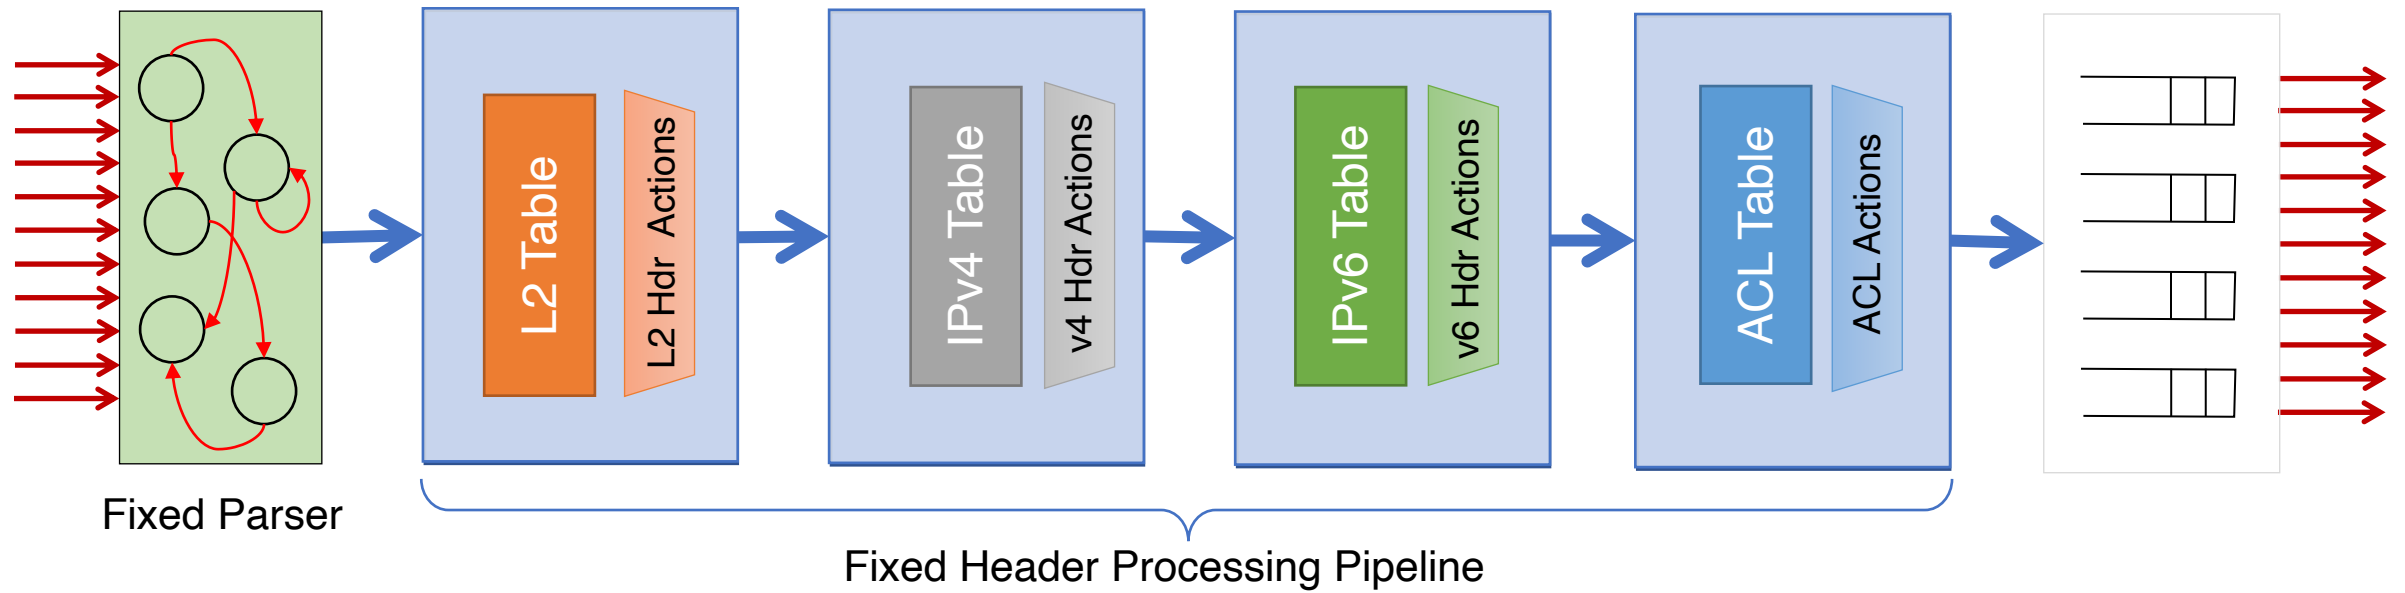
- **Separate fast path from slow path** (optimize the common case)
 - ARP lookup, fragmentation, error handling
- Try to fit all code into the processor instruction cache
- Heavily use **caching for table entries** across different memories in a hierarchy
 - **Traffic locality**: a small fast memory can service “most” traffic
- Two copies of table in external memory to support downtime-less updates to the forwarding table
- However, can't guarantee deterministic throughput for packets
 - Packet might access slower memories in the memory hierarchy

Packet lookup in RMT: *Pipelined parallelism*

- Different functionalities (ex: L2, L3) in different table stages
- Highly parallel over packets: admit 1 packet/cycle
- Pipeline circuitry *clocked* at a high rate: ex: RMT@1 GHz
- **Deterministic throughput**



Traditional pipelined hardware: fixed-function (Multiple Match-Action Tables)



MMT isn't enough!

- Operators want new protocols and services
- **Virtualization** and the cloud
 - VMs have their own address space (e.g., 10.0.x.x)
 - Physical networks route traffic using a different set of addresses (e.g., 128.6.x.x)
 - Keep them separate so you can place VMs wherever you can
 - Need: **dedicated new packet headers** to use for forwarding within the “core” fabric of the network
 - E.g., VXLAN, NVGRE
- Research experimentation and domain-specific headers
 - E.g., finance; university campuses; network feedback signals

MMT isn't enough!

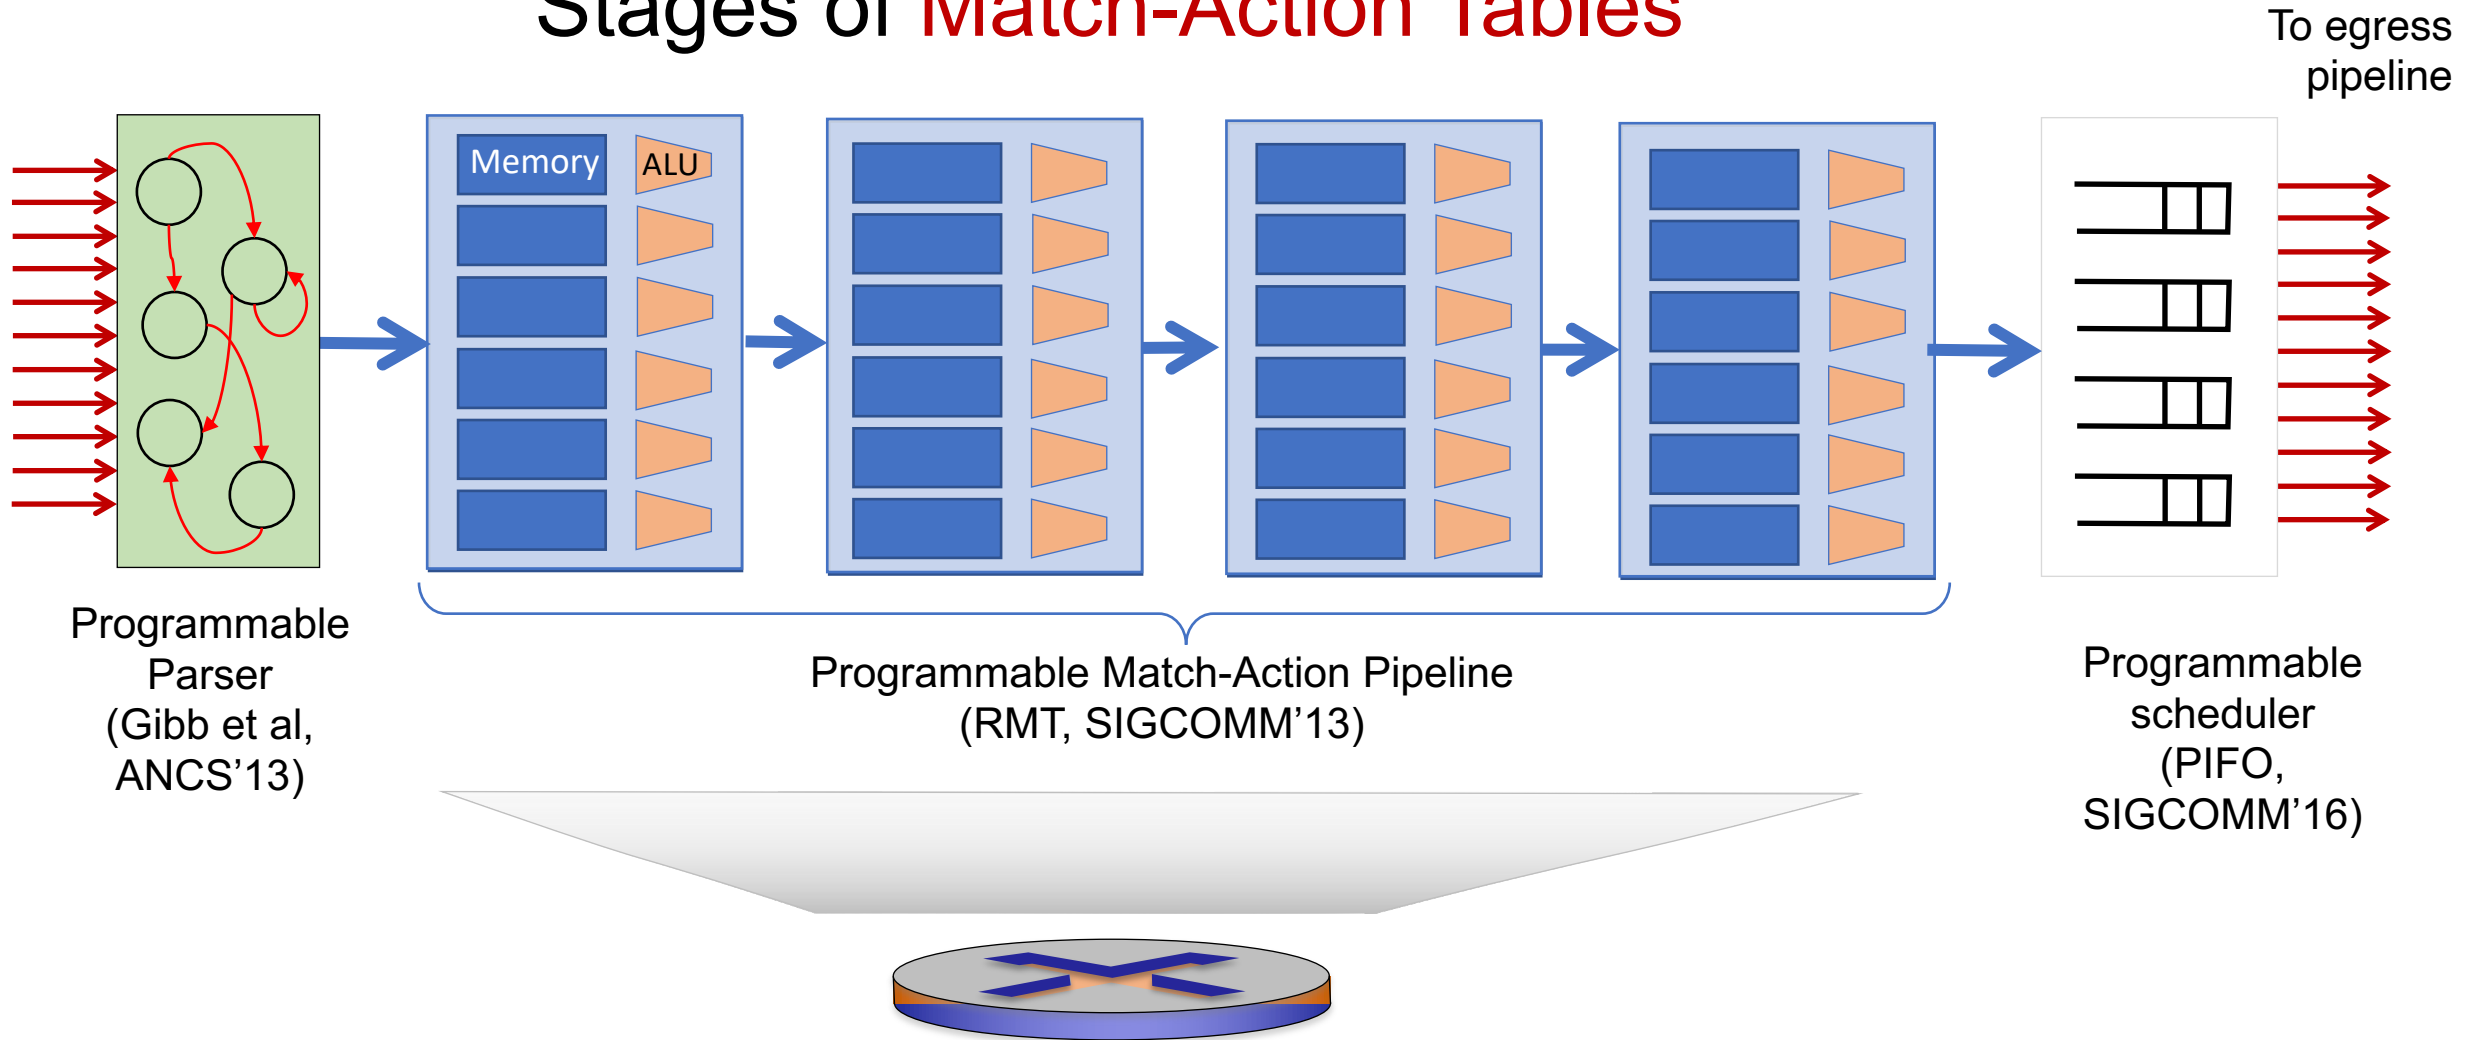
- Want to use **table memory** flexibly
 - Different environments need tables of different shapes and sizes
- Enterprises: ACL-heavy (“students can see information sent to other students from professors, but cannot see info from professors to printers”)
- Tier-1 ISP like ATT: L3-heavy (~1 million Internet IPv4 prefixes)
- **Static table sizes don't work well**
 - Can't use another table's memory, even if it is empty
 - Heterogenous devices to support different scenarios: complexity
 - Even device for a specific market may have insufficient resources

MMT isn't enough!

- All of this might be supported if switch hardware can be upgraded as often as software in general (e.g., on smartphone)
- Unfortunately, the reality is very far from it:
- Each device generation hardware upgrade: 3--5 years
 - New ASIC design, verification, fabrication, testing
- Even software upgrades take time:
 - Features requested by other customers stand in the way of releasing new feature that one customer wants

RMT: Protocol Independent Switch Arch

Stages of Match-Action Tables




A primer on high-speed memories

- Computers have different kinds of memory
 - Fastest caches (L1, L2, ...) made of SRAM
 - Fast main memory made of DRAM
 - Storage (HDD) made of magnetic disks, tape, SSDs, and so on
 - Translation Lookaside Buffer (TLB) with CAM/TCAM
- Today's high-speed routers use **SRAM** and **TCAM**
 - Static Random Access Memory for exact match
 - Ternary Content Addressable Memory for wildcard ACL match (also longest prefix match)


SRAM principles of operation

- Memory is organized into **banks**
- Each bank can be independently accessed through an **address**
- Data in the memory row at the address can be read/written each clock cycle
- Banks of larger sizes are denser (fewer wires to run), but you can only read/write one data row per bank per clock cycle (reduced parallelism)
- **Looking up** (ex: IPv4 dst) involves computing address(es) & reading mem




TCAM principles of operation

- Banks are accessed using **content**, not addresses (CAM)
- Contents of the memory are **ternary**: values can be 0, 1, or x (don't care)
- Incoming keys are matched against TCAM, with any bit accepted at the location of the wildcard bits
- Ternary logic is power- and area-hungry relative to SRAM




Match-Action Table memory design

- Match and Action units supplied with the Packet Header Vector
- Each stage accesses its own memory containing tables
- RMT uses a **crossbar** to pick PHV fields for matching against contents of SRAM/TCAM banks
 - **Flexible key generation** for lookup
- Distinction from fixed-function:
 - User-programmed fields stored in mem
 - Note: PHVs contain fields which are programmatically chosen




Match-Action Table memory design

- Entry looked up in the memory (SRAM/TCAM) contains a pointer to the **instruction** (action) for that entry
- Instructions implemented through **programmable ALUs**
 - More general ALUs than fixed-function hardware devices
- **Feasible since compute components are extremely cheap**
- **VLIW**: operate on multiple headers simultaneously
- Entry also contains pointers to:
 - Action data (e.g., port),
 - statistics (counters), if needed




Achieving reconfigurable match-action

Separately configurable and addressable memory banks is the key to using tables flexibly



Fixed-function matching




VS.

Flexible match function


Achieved by running **extra wires** (for read/write) to each memory bank + extra crossbars for multiple parallel tables

Achieving reconfigurable match-action

Separately configurable and addressable memory banks is the key to using tables flexibly



Fixed-function matching




VS.

Flexible match function


Different memory banks can contain entries for different tables
(e.g., IPv4 matching, ACL, ...)

Achieving reconfigurable match-action

Separately configurable and addressable memory banks is the key to using tables flexibly



Fixed-function matching




VS.

Flexible match function


Each table in a match-action stage may use a different crossbar configuration to extract a different set of fields from the PHV

Achieving reconfigurable match-action

Separately configurable and addressable memory banks is the key to using tables flexibly



Fixed-function matching




VS.

Flexible match function

The cost: 14% extra area (& power) for the fatter wires

Memory layout and use matters

- RMT: flexible partitioning of memory across **SRAM** and **TCAM**
- Numerous fixed size memory **blocks**
 - Memory fragmentation possible
- Circuitry for independent block-level access
- Deterministic access times
 - All of it is SRAM or TCAM
- Interesting compiler issues
 - “Packing” tables





Packet after route lookup

- What we have after lookup and modifications:
- One or more **output ports**, or a decision to **drop**
- **packet headers**, possibly modified from ingress


(4) Interconnect/Switching Fabric

- Move headers and packet from one interface to another
- Kinds of fabrics: memory, bus, crossbar



(4) Crossbars: The scheduling problem

- Want to utilize fabric capacity regardless of demand pattern
 - Crossbar is non-blocking
- MGR considers strategies to pair incoming demands & output ports
 - Greedy, wavefront, block wavefront
 - Need to address **fairness** across ports




(4) RMT switching fabric

- RMT uses **memory** as the fabric to hold packet headers and payloads between any two interfaces
- In the late 90s and early 2000s, there was considerable research on building high-speed packet buffers
- Today: **shared memory** switches & routers (shared → across ports)
 - Fast memory can be clocked at 1 GHz
- Fundamental tradeoff: **faster memories are not very dense**
 - Can't make the memory too large
 - Can't hold too many packets!
- Workaround: use memory access patterns: e.g., each queue is FIFO
- **Traffic manager** implements scheduling & buffer management


(5) Traffic Manager

- Where should the packets not currently serviced wait?
- Two designs: Input-queued vs. output-queued
- Output queueing avoids **HOL blocking** exhibited by input queueing.
 - Suppose port 1 wants to send to both 2 and 3 but port 2 busy
 - Packets from p1 towards p3 need not be delayed




(5) Traffic Manager

- RMT: Queueing represents output port contention
- A single output port can be represented by multiple **queues**
 - e.g., to implement weighted fair queueing
- Each queue is just a **linked list** in the shared memory
 - Maximum flexibility in queue sizes, but pointer overhead
 - Separate memory to maintain per-queue heads and tails



(5) Traffic Manager: Scheduling policies

- Q: how to dequeue packets from the buffer?
- Fair queueing across ports or flows
- Strict prioritization of some ports over others
- Rate limiting per port



(5) Traffic Manager: Buffer Management

- Q: how to enqueue packets into buffer?
 - If buffer is full, which packet should be dropped?
- Typical buffer management: **Tail-drop**
- Want **fairness**: if queue 1 has too many buffered pkts, don't tail-drop q2
 - Share memory by **partitioning** (carving memory out) across queues
- And **efficiency**
 - if q1 has no pkts, q2 should be able to use (nearly) all buffer memory
 - Interesting space for **dynamic buffer management algorithms**
 - e.g., Dynamic queue length thresholds for shared-memory packet switches (Choudhury and Hahne)



(6) Egress line termination

- Combine headers with payload for transmission
 - Must incorporate effect of header modifications
 - Also called **deparsing** or **serialization**
- **Multicast**: egress-specific packet processing
 - Ex: different source MAC address for each output port
- Multicast makes almost everything inside the switch (interconnect, lookups, queueing) more complex

Summary

- Packet buffer memory is a precious resource
 - Modern routers have a few tens of MegaBytes buffer
 - ... for all packets across all ports
- Traffic Manager: scheduling and buffer-management algorithms
- Buffer memory costs substantial amount: capital, area, power
 - Data centers: cheap routers with shallow buffers

Three kinds of data plane programmability

- Packet header formats, i.e., the packet parser
 - Example: Go from IPv4 -> IPv6
 - Custom packet format to carry financial info at high speed on a point-to-point link
- Table formats, actions, sizes, i.e., the match-action tables
 - Change which fields in the packet can be processed by a table
 - Control the table sizes, i.e., # entries, and hence the memory resource footprint according to use case.

Three kinds of data plane programmability

- Packet scheduling, i.e., the traffic manager
 - Flexible classification of packets
 - Flexible assignment of ordering and timing of when packets are transmitted from an outgoing link

... distinct from control plane programmability

- The control plane must compute the packet-processing rules put into the memory on the router ASIC
 - Example: packet with IPv4 destination 10.0.0.1 must go out of port 4
- Data plane programmability refers to the flexibility in the allowed set of packet headers, tables, and actions themselves, not the actual rules.
 - Example: There is a table that matches on IPv4 destination addr whose action is to determine the output port