

# The Network Layer: NAT, IPv6, Routing Algorithms

CS 352, Lecture 15, Spring 2020

<http://www.cs.rutgers.edu/~sn624/352>

Srinivas Narayana

# Course announcements

- Take-home mid-term this weekend
  - Practice problems and review mid-term released soon
- Open-book and open-notes, but not open-Internet
  - Not permitted to search the Internet for answers
  - No collaboration permitted; all answers and work must be your own
  - Calculators are allowed
- This class will follow absolute grading with generous thresholds
  - i.e., no curve, you're only competing with yourself
  - We trust you to do the right thing


# Review of concepts

- Internet Protocol:
  - Headers: src/dst, upper layer, fragmentation, time to live
- Dynamic host configuration protocol (DHCP):
  - How does an endpoint get its IP address?
  - **Broadcast-based**: endpoint asks entire network for answers
  - DHCP server returns IP address, subnet mask, local DNS server address, gateway router address
- Internet Control Message Protocol (ICMP):
  - Network troubleshooting protocol
  - Ping: **reachability** using ICMP echo request and reply
  - Traceroute: **router-level path** using ICMP time-exceeded messages and incrementing TTL

# Network Address Translation (NAT)

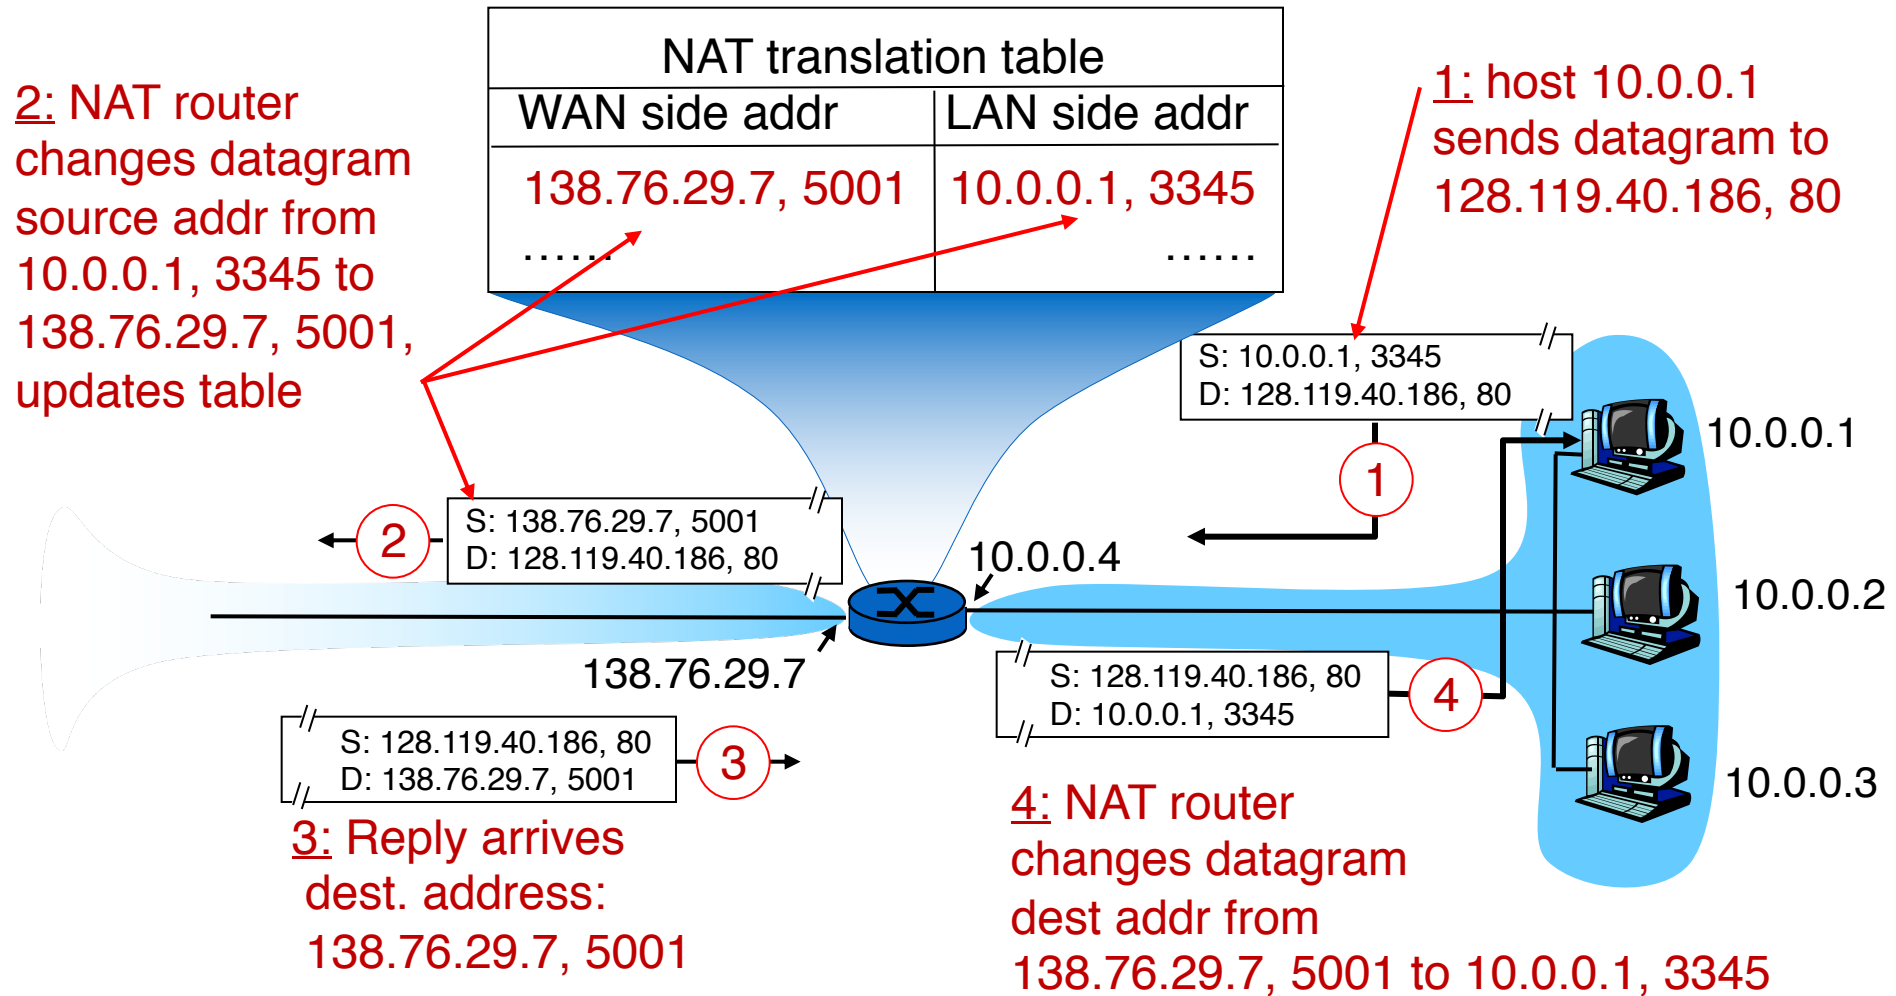
How do you survive in a world where names are scarce?

# NAT: Network Address Translation



**All** datagrams **leaving** local network have **same** source IP address: 138.76.29.7, with **different source IP port** numbers

# NAT: Network Address Translation



# NAT: Network Address Translation

- **Features:** local network uses just one IP address as far as outside world is concerned:
  - range of addresses not needed from ISP: **just one IP address for all devices**
  - can change addresses of devices in local network without notifying outside world
  - can change ISP without changing addresses of devices in local network
  - Devices inside local network not explicitly addressable
  - Devices inside local network invisible to the outside world (a security plus) unless the device inside connects first.

**Your home WiFi router implements NAT-ting.**

**If you're at home, you're almost surely behind a NAT gateway right now.**

# The impact of NATs

```
[flow:352-S20]$ ifconfig en0
en0: flags=8863<UP,BROADCAST,SMART,RUNNING,SIMPLEX,MULTICAST> mtu 1500
    ether f0:18:98:1c:fc:36
    inet6 fe80::1036:7dea:82ee:e868%en0 prefixlen 64 secured scopeid 0xa
    inet 192.168.1.151 netmask 0xffffffff broadcast 192.168.1.255
    nd6 options=201<PERFORMNUD,DAD>
    media: autoselect
    status: active
[flow:352-S20]$ █
```



what's my ip address



All Images Videos News Maps | Answer

Settings ▼

Your IP address is 74.102.79.209 in [New Brunswick, New Jersey, United States \(08901\)](#)



# NAT: Network Address Translation

- 16-bit port-number field:
  - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - Routers should only work upto the network layer, not transport ports!
  - violates “end-to-end argument”
    - NAT must be taken into account by app designers
    - e.g., P2P applications like skype
  - Purists: address shortage should instead be solved by IPv6

# Think about...

- How do the hosts inside the home network get their IP addresses?
- How does your home router get its externally visible IP address?


# Poll #1

- The length of the TCP port field is 16 bits. Assume that a NAT router has a single external IP address. In principle, how many TCP connections can the router support between the NAT and the outside world?
  - (a) 1
  - (b)  $2^{16}$
  - (c) none of the above

# Internet Protocol v6 (IPv6)

# Recent Developments: IPv6

- IPv4 has limited address space (32 bits) and is running out of addresses. 32 bits are not enough!
- More devices: phones, watches, your refrigerator(!), ...
- Real-time traffic and mobile users are also becoming more common



IP version 6

# IPv6: Main changes from IPv4

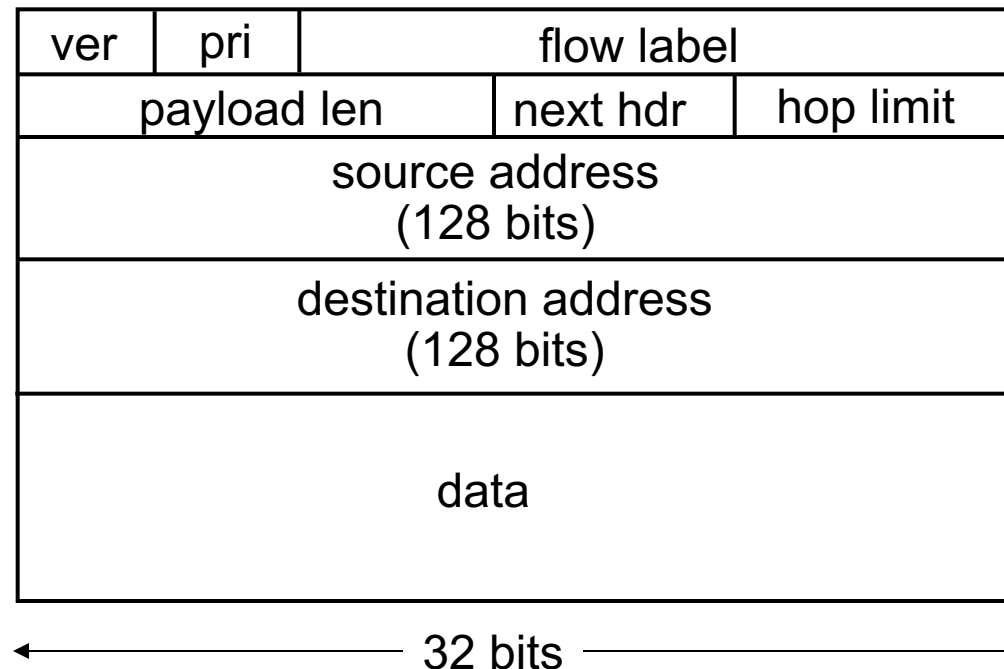
- Large address space:
  - 128-bit addresses (16 bytes)
  - Allows up to 340,282,366,920,938,463,463,374,607,431,768,211,456 unique addresses ( $3.4 \times 10^{38}$ )
- Fixed length headers (40 bytes)
  - Improves the speed of packet processing in routers
- IPv6 “options” processing happens through a separate mechanism

# IPv6 datagram format

*priority:* identify priority among datagrams in flow

*flow Label:* identify datagrams in same “flow”  
(concept of “flow” left undefined)

*next header:* identify upper layer protocol for data

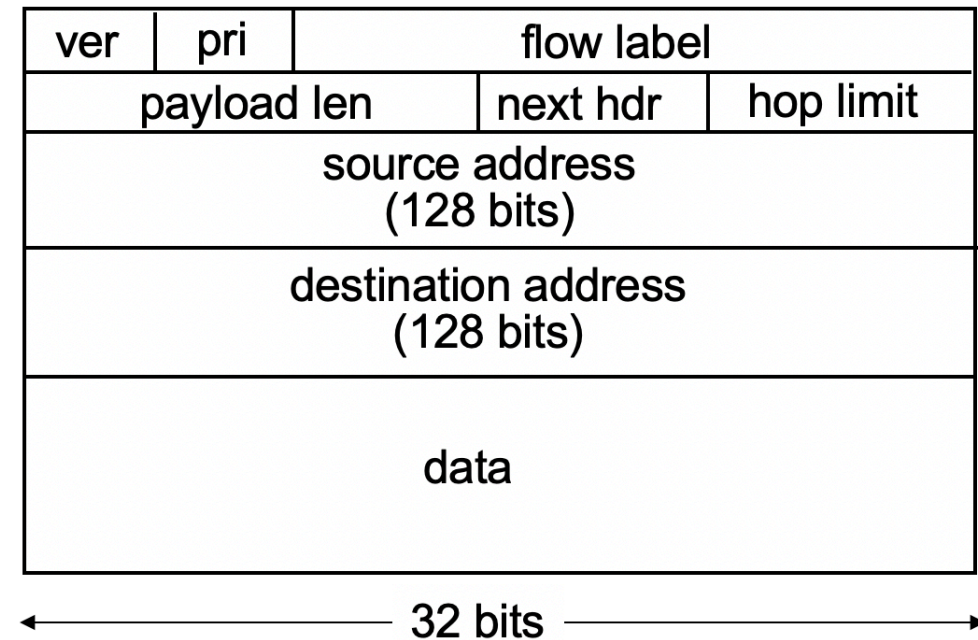
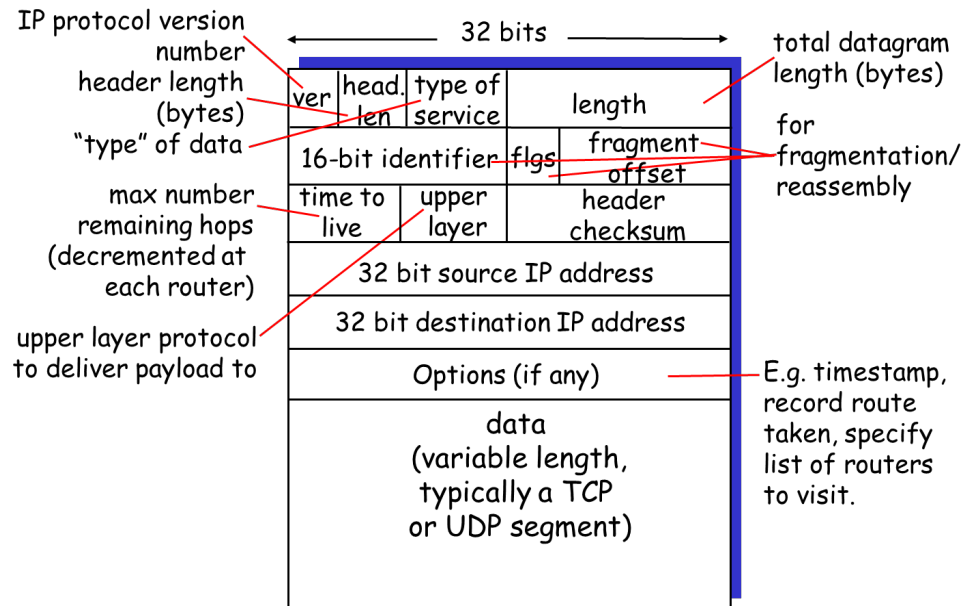


# Other changes from IPv4

- *checksum*: removed entirely to reduce processing time at each hop
- *options*: allowed, but outside of header, indicated by “Next Header” field
- *ICMPv6*: new version of ICMP
  - additional message types, e.g. “Packet Too Big”
  - multicast group management functions



# IPv4 vs IPv6: Can you tell the differences?



# IPv6 Flows

- Support for “flows”
  - Flows help support real-time service in the Internet
  - A “flow” is a number in the IPv6 header that can be used by routers to see which packets belong to the same stream
  - Guarantees can then be assigned to certain flows
  - Example:
    - Packets from flow 10 (e.g., audio call) should receive rapid delivery
    - Packets from flow 12 (e.g., web traffic) should receive reliable delivery

# IPv6 Addresses


- Classless addressing/routing (similar to CIDR)
- Notation: `xx:xx:xx:xx:xx:xx:xx:xx`
  - x = 4-bit hex number
  - contiguous 0s are compressed: `47CD::A456:0124`
  - IPv6 compatible IPv4 address: `::128.64.18.87`
    - First 96 bits are 0
  - Global unicast addresses start with `001....`
  - `2000::/3` prefix

# IPv6 adoption

## IPv6 Adoption

We are continuously measuring the availability of IPv6 connectivity among Google users. The graph shows the percentage of users that access Google over IPv6.

Native: 0.07% 6to4/Teredo: 0.16% Total IPv6: 0.23% | Jan 1, 2009



# IPv6: Adoption

- Google: ~1/3 of clients access services via IPv6 (Apr 2020)
- *Long (long!) time for deployment, use*
  - 20 years and counting!
- Think of application-level changes in last 20 years: WWW, Facebook, Skype, video streaming, AR, telesurgery,...
  - *Why?*

# Poll #2

- What's the primary difference between IPv4 and IPv6?
  - (a) Larger address space
  - (b) Fixed-length, fewer headers by default
  - (c) No fragmentation
  - (d) All of the above

# Routing Protocols

How do we design a “Google Maps” navigator for the Internet?

# Network-layer functions

*Recall: two network-layer functions:*

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to destination

*data plane*

*control plane*


*Two approaches to structuring network control plane:*

- per-router control (traditional)
- logically centralized control (software defined networking)



# Per-router control plane

Individual routing algorithm components *in each and every router* interact with each other in control plane to compute forwarding tables




# Routing protocols

*Routing protocol goal:* determine “good” paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- path: sequence of routers packets will traverse in going from given initial source host to given final destination host
- “good”: least “cost”, “fastest”, “least congested”
- routing: a “top-10” networking challenge!

# Graph abstraction



Graph:  $G = (N, E)$


$N$  = set of routers =  $\{ u, v, w, x, y, z \}$

$E$  = set of links =  $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Remark: Graph abstraction is useful in other network contexts

Example: P2P, where  $N$  is set of peers and  $E$  is set of TCP connections

# Graph abstraction: costs



- $c(x,x')$  = cost of link  $(x,x')$

- e.g.,  $c(w,z) = 5$

- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Cost of path  $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Question: What's the least-cost path between u and z ?

Routing algorithm: find “good” paths from source to destination router.

# Routing algorithm classification

*Q: global or decentralized information?*

*global:*

- all routers have complete topology, link cost info
- “link state” algorithms

*decentralized:*

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- “distance vector” algorithms

*Q: static or dynamic?*

*static:*

- routes change slowly over time

*dynamic:*

- routes change more quickly
  - periodic update
  - in response to link cost changes

# Poll #3

- What is a good candidate to use to fix edge weights in a network graph representation?
  - (a) 1/link rate
  - (b) round-trip time
  - (c) congestion
  - (d) any of the above

# Link State Algorithms

# A Link-State Routing Algorithm

## Dijkstra's algorithm

- net topology, link costs known to all nodes
  - accomplished via **link state broadcast**
  - all nodes have same info
- computes least cost paths from one node ('source') to all other nodes
  - gives **forwarding table** for that node
- iterative: after k iterations, know least cost path to k dest.'s

## Notation:

- **$c(x,y)$** : link cost from node x to y;  
=  $\infty$  if not direct neighbors
- **$D(v)$** : current value of cost of path from source to dest. v
- **$p(v)$** : predecessor node along path from source to v
- **$N'$** : set of nodes whose least cost path definitively known



# Dijkstra's Algorithm

1 **Initialization:**

2  $N' = \{u\}$

3 for all nodes  $v$

4 if  $v$  adjacent to  $u$

5 then  $D(v) = c(u,v)$

6 else  $D(v) = \infty$

7

8 **Loop**

9 find  $w$  not in  $N'$  such that  $D(w)$  is a minimum

10 add  $w$  to  $N'$

11 update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$  :

12  $D(v) = \min( D(v), D(w) + c(w,v) )$


13 /\* new cost to  $v$  is either old cost to  $v$  or known

14 shortest path cost to  $w$  plus cost from  $w$  to  $v$  \*/

15 **until all nodes in  $N'$**

# Dijkstra's algorithm: example

| Step | N'     | D(v),p(v) | D(w),p(w) | D(x),p(x) | D(y),p(y) | D(z),p(z) |
|------|--------|-----------|-----------|-----------|-----------|-----------|
| 0    | u      | 2,u       | 5,u       | 1,u       | $\infty$  | $\infty$  |
| 1    | ux     | 2,u       | 4,x       |           | 2,x       | $\infty$  |
| 2    | uxy    | 2,u       | 3,y       |           |           | 4,y       |
| 3    | uxyv   |           | 3,y       |           |           | 4,y       |
| 4    | uxyvw  |           |           |           |           | 4,y       |
| 5    | uxyvwz |           |           |           |           |           |



# Poll #4

- Link-state information of a router is sent **to all routers** before computing shortest paths in a link-state protocol.
  - (a) true
  - (b) false

# Poll #5

- Link-state protocols don't compute shortest paths when there are cycles in the network topology.
  - (a) true
  - (b) false

# Distance Vector Algorithms

# Distance Vector Algorithm

- $D_x(y)$  = estimate of least cost from  $x$  to  $y$
- Distance vector:  $\mathbf{D}_x = [D_x(y) : y \in N]$
- Node  $x$  knows cost to each neighbor  $v$ :  $c(x,v)$
- Node  $x$  maintains  $\mathbf{D}_x$
- Node  $x$  also maintains its neighbors' distance vectors
  - For each neighbor  $v$ ,  $x$  maintains  $\mathbf{D}_v = [D_v(y) : y \in N]$

# Distance vector algorithm

*Bellman-Ford equation (dynamic programming)*

let

$d_x(y) :=$  cost of least-cost path from  $x$  to  $y$

then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

cost from neighbor  $v$  to destination  $y$

cost to neighbor  $v$

$\min$  taken over all neighbors  $v$  of  $x$

# Distance vector algorithm

Basic idea:

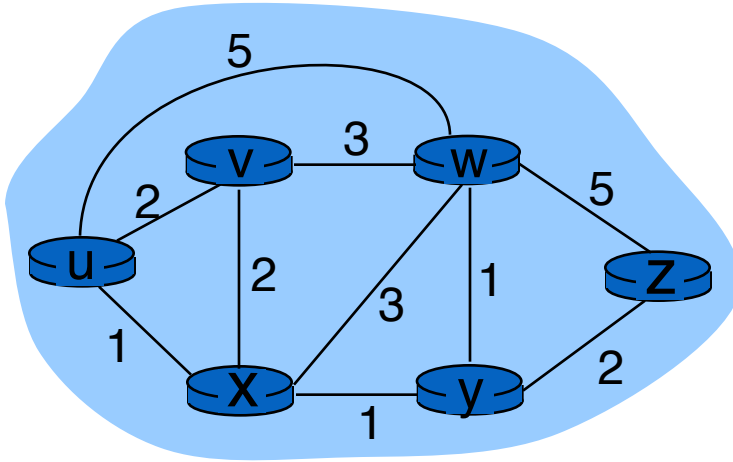
- Each node periodically sends its own distance vector estimate **to neighbors**
- When node a node  $x$  receives new DV estimate from neighbor, it updates its own DV using **Bellman-Ford** equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$

- Under some conditions, the estimate  $D_x(y)$  *converge the actual least cost*  $d_x(y)$



# Distance vector: example



Start with  $d_v(z) = 5$ ,  $d_x(z) = 3$ ,  $d_w(z) = 3$

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node that achieves minimum is next  
hop in shortest path → forwarding table

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

|      |   | cost to  |          |          |
|------|---|----------|----------|----------|
|      |   | x        | y        | z        |
| from | x | 0        | 2        | 7        |
|      | y | $\infty$ | $\infty$ | $\infty$ |
|      | z | $\infty$ | $\infty$ | $\infty$ |

node y table

|      |   | cost to  |          |          |
|------|---|----------|----------|----------|
|      |   | x        | y        | z        |
| from | x | $\infty$ | $\infty$ | $\infty$ |
|      | y | 2        | 0        | 1        |
|      | z | $\infty$ | $\infty$ | $\infty$ |

node z table

|      |   | cost to  |          |          |
|------|---|----------|----------|----------|
|      |   | x        | y        | z        |
| from | x | $\infty$ | $\infty$ | $\infty$ |
|      | y | $\infty$ | $\infty$ | $\infty$ |
|      | z | 7        | 1        | 0        |

|      |   | cost to |   |   |
|------|---|---------|---|---|
|      |   | x       | y | z |
| from | x | 0       | 2 | 3 |
|      | y | 2       | 0 | 1 |
|      | z | 7       | 1 | 0 |


|      |   | cost to |   |   |
|------|---|---------|---|---|
|      |   | x       | y | z |
| from | x | 0       | 2 | 7 |
|      | y | 2       | 0 | 1 |
|      | z | 7       | 1 | 0 |

|      |   | cost to |   |   |
|------|---|---------|---|---|
|      |   | x       | y | z |
| from | x | 0       | 2 | 7 |
|      | y | 2       | 0 | 1 |
|      | z | 3       | 1 | 0 |

|      |   | cost to |   |   |
|------|---|---------|---|---|
|      |   | x       | y | z |
| from | x | 0       | 2 | 3 |
|      | y | 2       | 0 | 1 |
|      | z | 3       | 1 | 0 |

|      |   | cost to |   |   |
|------|---|---------|---|---|
|      |   | x       | y | z |
| from | x | 0       | 2 | 3 |
|      | y | 2       | 0 | 1 |
|      | z | 3       | 1 | 0 |


|      |   | cost to |   |   |
|------|---|---------|---|---|
|      |   | x       | y | z |
| from | x | 0       | 2 | 3 |
|      | y | 2       | 0 | 1 |
|      | z | 3       | 1 | 0 |



# Distance vector: link cost changes

## *link cost changes:*

- ❖ node detects local link cost change
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors



“good  
news  
travels  
fast”

$t_0$ :  $y$  detects link-cost change, updates its DV, informs its neighbors.


$t_1$ :  $z$  receives update from  $y$ , updates its table, computes new least cost to  $x$ , sends its neighbors its DV.

$t_2$ :  $y$  receives  $z$ 's update, updates its distance table.  $y$ 's least costs do *not* change, so  $y$  does *not* send a message to  $z$ .

# Problem: Count-to-Infinity

- With distance vector routing, good news travels fast, but bad news travels slowly
- When a router goes down, it takes can take a really long time before all the other routers become aware of it

# Count-to-Infinity



1 2 3 4

Initially

3 2 3 4

After 1 exchange

3 4 3 4

After 2 exchanges

5 4 5 4

After 3 exchanges

5 6 5 6

After 4 exchanges

7 6 7 6

After 5 exchanges

etc... to infinity

# Count-to-infinity

*“Bad news travels slowly”*

*Poisoned reverse:*

- ❖ If Z routes through Y to get to X :
  - Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)
- ❖ Will this completely solve count to infinity problem?

# Comparison of LS and DV algorithms

## *message complexity*

- **LS:** with  $n$  nodes,  $E$  links,  $O(nE)$  msgs sent
- **DV:** exchange between neighbors only
  - convergence time varies

## *speed of convergence*

- **LS:**  $O(n^2)$  algorithm requires  $O(nE)$  msgs
- **DV:** convergence time varies
  - may be routing loops
  - count-to-infinity problem

**robustness:** what happens if router malfunctions?

## **LS:**

- node can advertise incorrect *link* cost
- each node computes only its *own* table

## **DV:**

- DV node can advertise incorrect *path* cost
- each node's table used by others
  - error propagate thru network

# Poll #6

- Which routing protocol(s) ensure that routers have a view of the network that is consistent with each other at all times?
  - (a) Link-state protocols
  - (b) Distance-vector protocols
  - (c) All of the above
  - (d) None of the above



# Routing protocols are widely deployed

- OSPF: a link-state protocol
  - “Open Shortest Path First”
  - IS-IS, nearly identical to OSPF
- RIP: a distance-vector protocol
- LS and DV deployed **inside** an **autonomous system**
- Additional tricks to scale the protocols with network size:
  - Areas; Hierarchy
  - “Border routers” that summarize each area
- Next lecture: Routing **across** autonomous systems