

Transport

Part II

Lecture 6, Computer Networks (198:552)
Fall 2019

Two Main Transport Layers

- User Datagram Protocol (UDP)
 - Abstraction of independent messages between endpoints
 - Just provides demultiplexing and error detection
 - Header fields: port numbers, checksum, and length
 - Low overhead, good for query/response and multimedia
- Transmission Control Protocol (TCP)
 - Provides support for a **stream of bytes** abstraction


Transmission Control Protocol (TCP)

- Multiplexing/demultiplexing
 - Determine which conversation a given packet belongs to
 - All transports need to do this
- Reliability and flow control
 - Ensure that data sent is delivered to the receiver application
 - Ensure that receiver buffer doesn't overflow
- Ordered delivery
 - Ensure bits pushed by sender arrive at receiver app **in order**
 - Q: why would packets ever be received out of order?
- Congestion control
 - Ensure that data sent doesn't overwhelm **network resources**
 - Q: which network resource?

Ordered Delivery


Reordering packets at the receiver side

- Let's suppose receiver gets packets 1, 2, and 4, but not 3 (dropped)
- Suppose you're trying to download a Word document containing a report
- What would happen if transport at the receiver directly presents packets 1, 2, and 4 to the Word application?




Reordering at the receiver side

- Reordering can also happen due to packets taking different paths through a network
- Receiver needs a general strategy to ensure that data is presented to the application **in the same order of sender side bytes pushed**



Buffering at the receiver side



Buffering at the receiver side

- The TCP receiver uses a **memory buffer** to hold packets until they can be read by the application in order
- This process is known as **TCP reassembly**

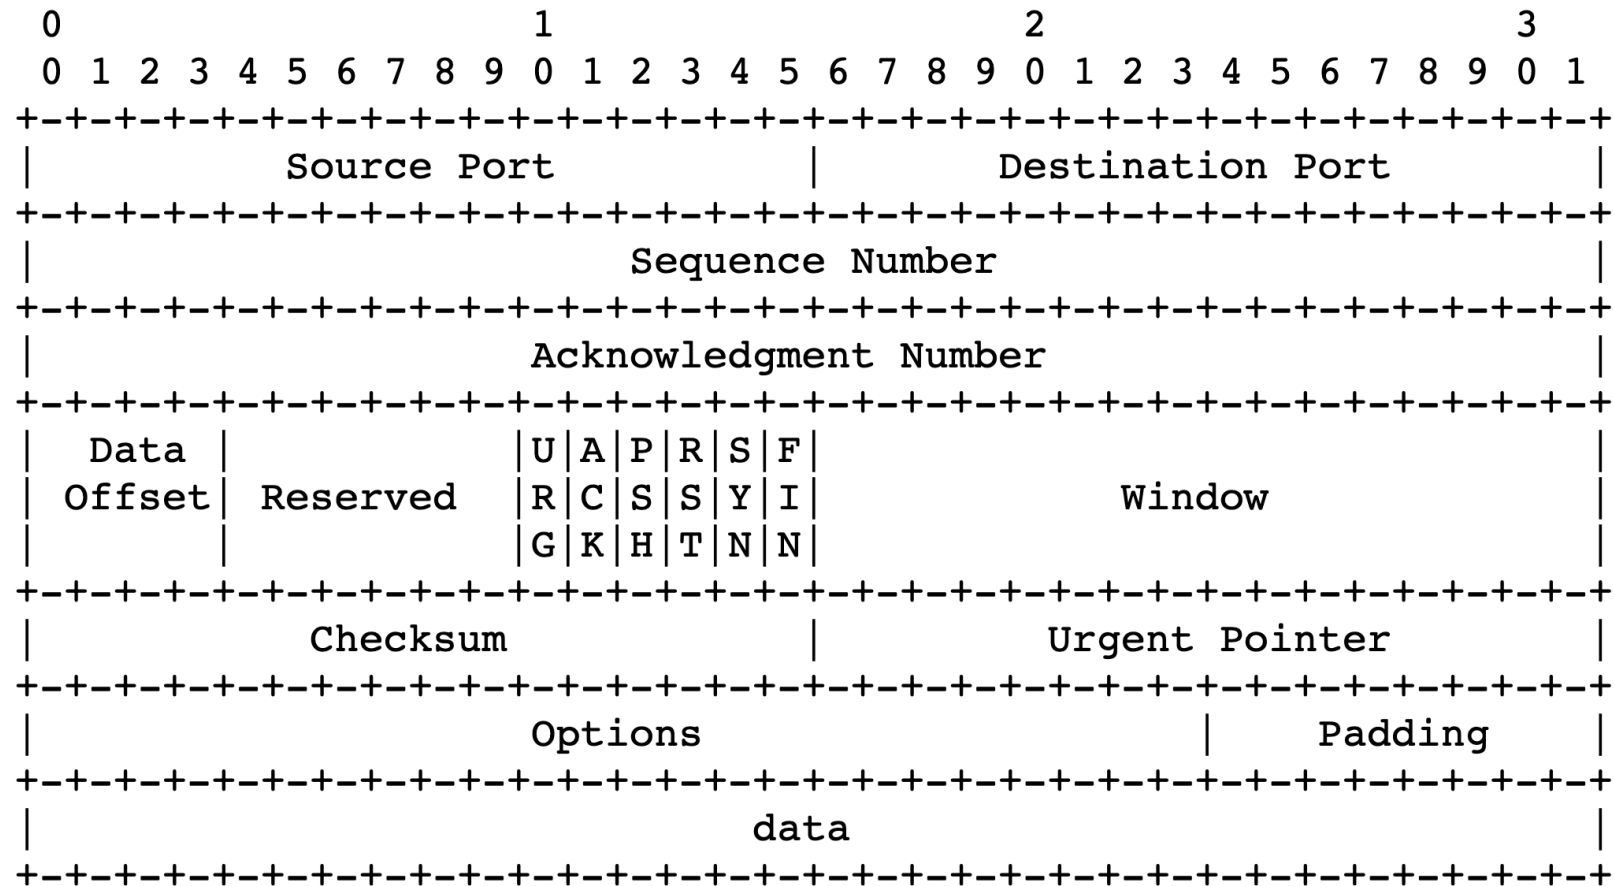
Implications of ordered delivery

- Packets cannot be delivered to the application if there is an **in-order packet missing** from the receiver's buffer
 - The receiver can only buffer so much out-of-order data
 - Subsequent out-of-order packets dropped
- It doesn't matter that the packets successfully arrive at the receiver NIC from the sender over the network
- **TCP application throughput will suffer** if there is too much packet "reordering" in the network

Implications of ordered delivery

- A TCP sender can only send as much as the **free receiver buffer space** available before packets are dropped at the receiver
 - This number is called the **receiver window size**
 - TCP is said to implement **flow control**

Flow control headers



TCP Header Format

Note that one tick mark represents one bit position.

Implications of buffering at receiver side

- Flow control has implications for TCP data delivery, even when packets arrive in order
- Typically TCP packets arrive **in a burst**
 - Why?
 - The receiver application won't read this data in one shot
- Q: What's the size of the maximum burst?

Sizing the receiver window

- **Bandwidth-delay product:** enough data to “fill the pipe”
- Implications to achieve high throughput:
 - More memory required at high bandwidth
 - More memory required at high RTT
 - Consider 100 Gbit/s connection at 100 ms RTT
- With this window size, can you guarantee that you will never “block” the connection due to a filled-up receiver buffer?
 - **You can't!** If app never reads from receiver buffer, it will fill up and not allow any more data to come in.

Implications of ordered delivery


- What if packets travel from sender to receiver over **multiple paths**?
- Imagine a situation where one path is much faster than another
- First (faster) path sends packets: 1, 3, 5, ...
- Second (slower) path sends packets: 2, 4, 6, ...
- Reassembly will require dropping the connection's throughput to match the slower one

Implications of ordered delivery

- Reordering and reassembly are **bad** for application throughput
- Most network-level load balancing mechanisms **avoid per-packet multi-path forwarding**
 - Balance load at **per-flow** or **per-flowlet** (burst) granularity
- Multi-path TCP variants exist, but all need to solve the **path scheduling** problem:
 - Schedule outgoing packet transmissions and adjust windows
 - ... so that the packets arrive at the receiver (roughly) in order

Congestion control

How should multiple endpoints share net?



- It is difficult to know where the **bottleneck** link is
- It is difficult to know how many other endpoints are using that link
- Endpoints may join and leave at any time
- Network paths may change over time, leading to different bottleneck links (with different link rates) over time

The approach that the Internet takes is to use a distributed algorithm to converge to an efficient and fair outcome.

The approach that the Internet takes is to use a **distributed algorithm** to converge to an efficient and fair outcome.

No one can centrally view or control all the endpoints and bottlenecks in the Internet.

Every endpoint must try to reach a globally good outcome by itself: i.e., in a distributed fashion.

This also puts a lot of **trust in endpoints**.

The approach that the Internet takes is to use a distributed algorithm to converge to an **efficient** and fair outcome.

If there is spare capacity in the bottleneck link, the endpoints should use it.

The approach that the Internet takes is to use a distributed algorithm to converge to an efficient and **fair** outcome.

If there are N endpoints sharing a bottleneck link, they should be able to get **equitable** shares of the link's capacity.

The approach that the Internet takes is to use a distributed algorithm to converge to an efficient and fair outcome.

So, how to achieve this?

Feedback from network offers clues...

- **Signals**


- Packets being dropped (ex, RTO fires)
- Packets being delayed
- Rate of incoming ACKs

} “Implicit” feedback signals
(more on explicit signals
later)

- **Knobs**

- What can you change to “probe” the sending rate?
- Suppose receiver buffer is unbounded:
- Let’s call the amount of in-flight data per RTT the **congestion window**
- Increase congestion window: e.g., by x or by a factor of x
- Decrease congestion window: e.g., by x or by a factor of x


Steady state: ACK clocking



Time for an activity


How to get to steady state?

- Slow start
- Congestion avoidance: Additive increase, multiplicative decrease (**AIMD**)




How to get to steady state?

- Upon a **timeout**, drop the window to a small fixed value (IW)
- Upon **idling**, drop the window to a small fixed value (RW)



Why AIMD?

- Converges to fairness
- Converges to efficiency
- Increments to rate smaller as fairness increases



TCP's steady state is not static

- As stated so far, TCP probes network capacity iteratively
 - ... Until it induces a loss
 - ... and then probes network capacity again
-
- It is important to have efficient mechanisms to detect and recover from packet loss

Loss detection & recovery in TCP

- Detecting loss before timeouts occur through **fast retransmit**
- Basic idea:
 - if the receiver did not receive a packet
 - but did receive a subsequent few packets (**duplicate ACKs**)
 - ... the unreceived packet must have been dropped.
- **Fast recovery**: Don't drop the window too much
 - If you're receiving dup ACKs, packets are being delivered
 - Do congestion avoidance instead of slow start from IW
- Many more details in RFC 2581 and follow-on work


Packet Scheduling

Are endpoint algorithms alone enough?




- What if an endpoint is malicious or buggy?
- Want the network core to do something more about **resource allocation** than best effort


Network model




First-in first-out (FIFO) queue + tail-drop



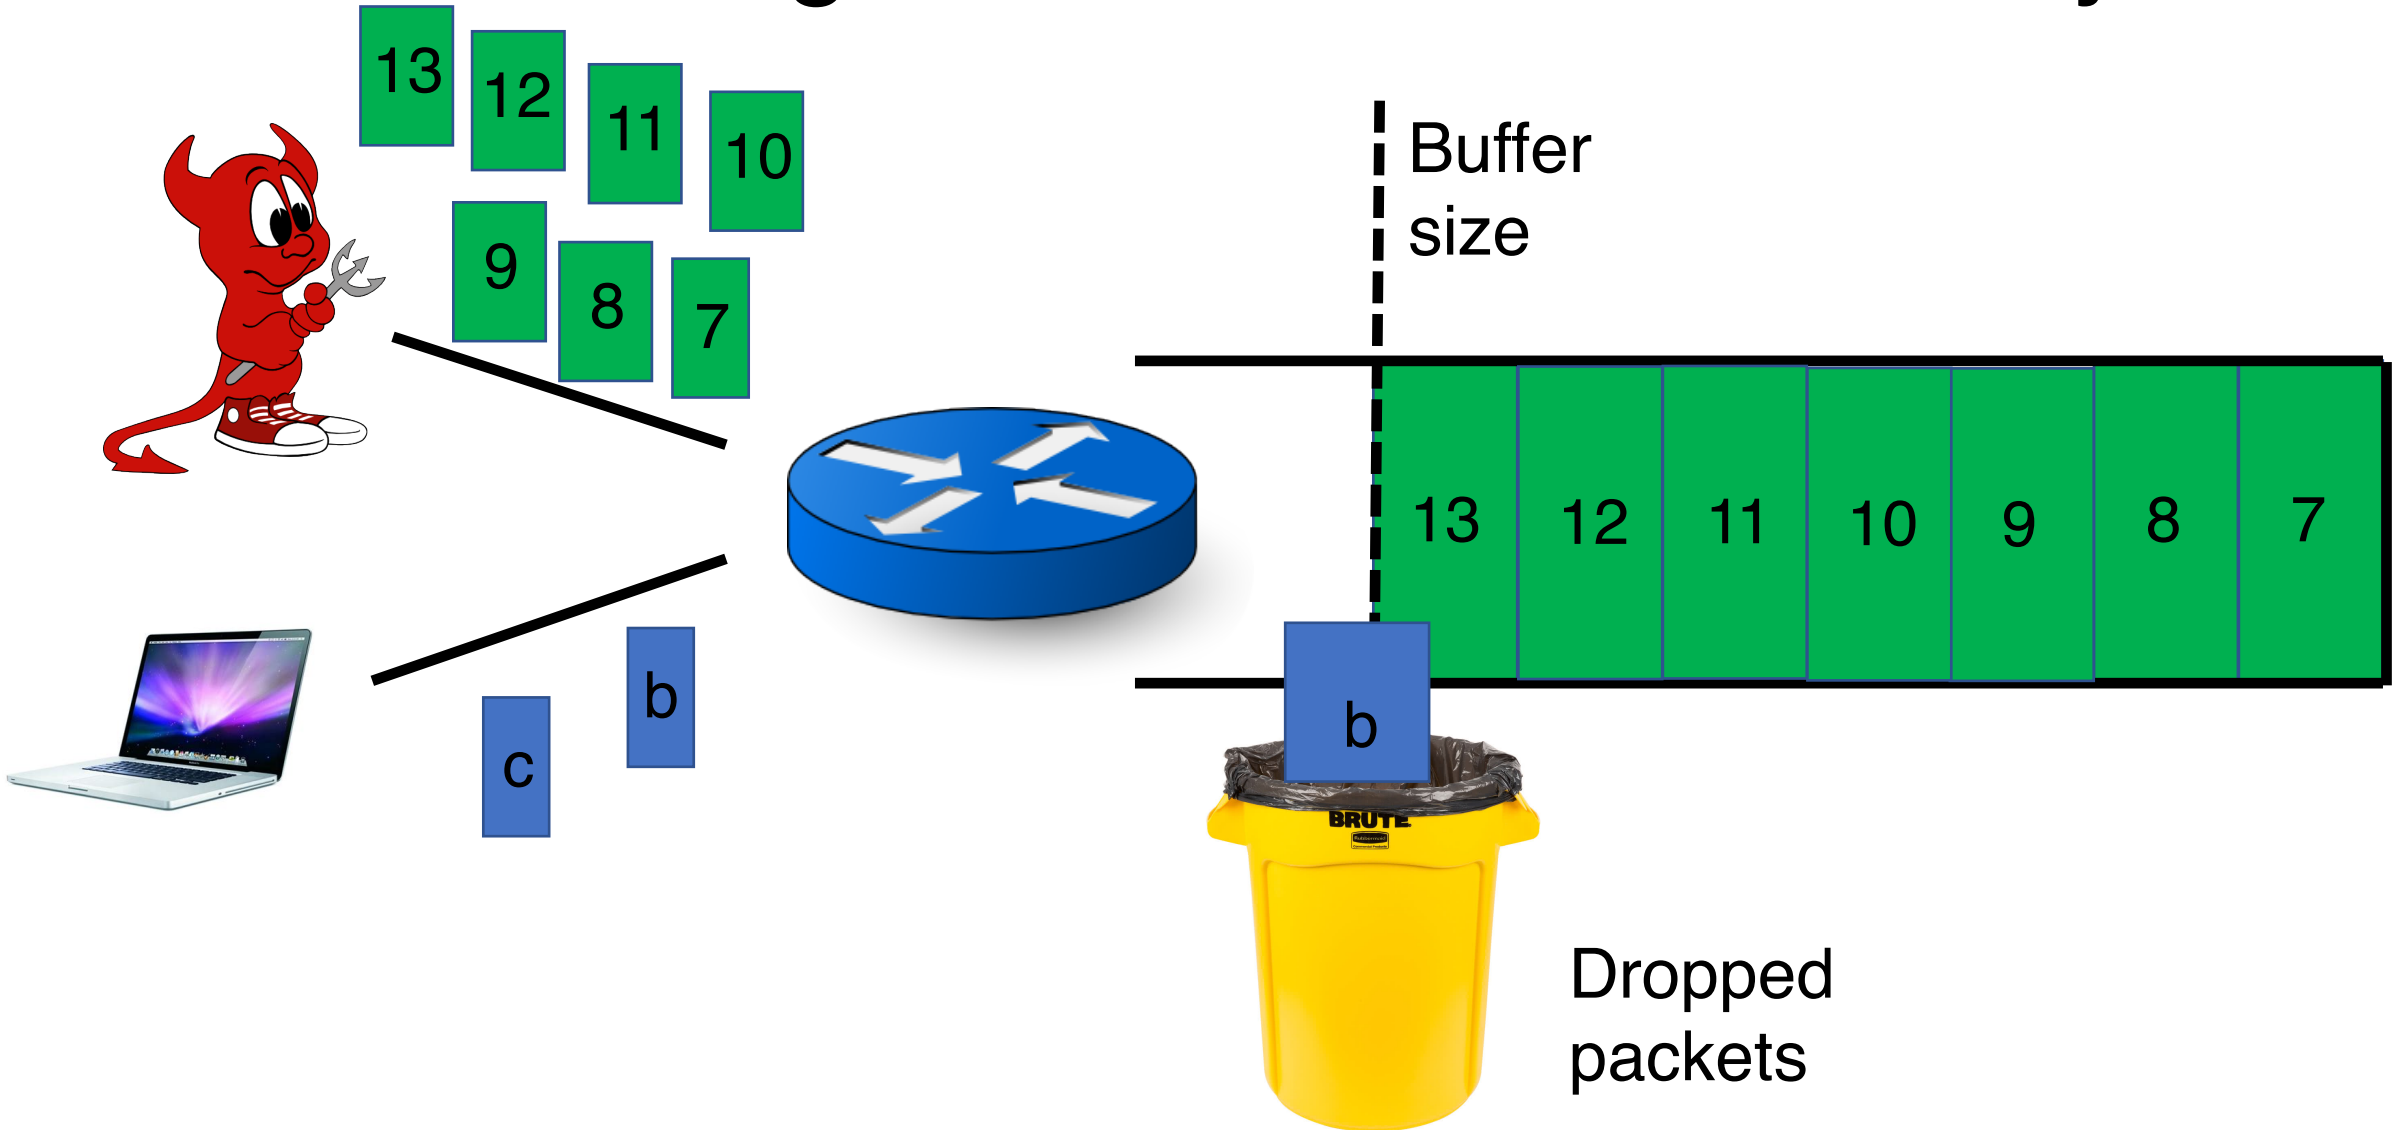
First-in first-out (FIFO) queue + tail-drop



ACK-clocking makes it worse: lucky case



ACK-clocking makes it worse: unlucky case



Network monopolized by “bad” endpoints

- An ACK signals the source of a free router buffer slot
 - Further, ACK clocking means that the source transmits again
- Contending packet arrivals may not be random enough
 - Blue flow can't capture buffer space for *a few* round-trips
- Sources which sent successfully earlier get to send again
- A FIFO tail-drop queue *incentivizes* sources to misbehave!

Packet scheduling on routers

- We will discuss packet scheduling algorithms implemented on routers in detail later in this course.
- Goal: Achieve a predetermined resource allocation **regardless of endpoint behavior**
- How to make such allocation “efficient”?
 - Implement on routers at high speeds and low cost
 - Achieve equitable sharing of network bandwidth & queues
 - Use available bandwidth effectively