

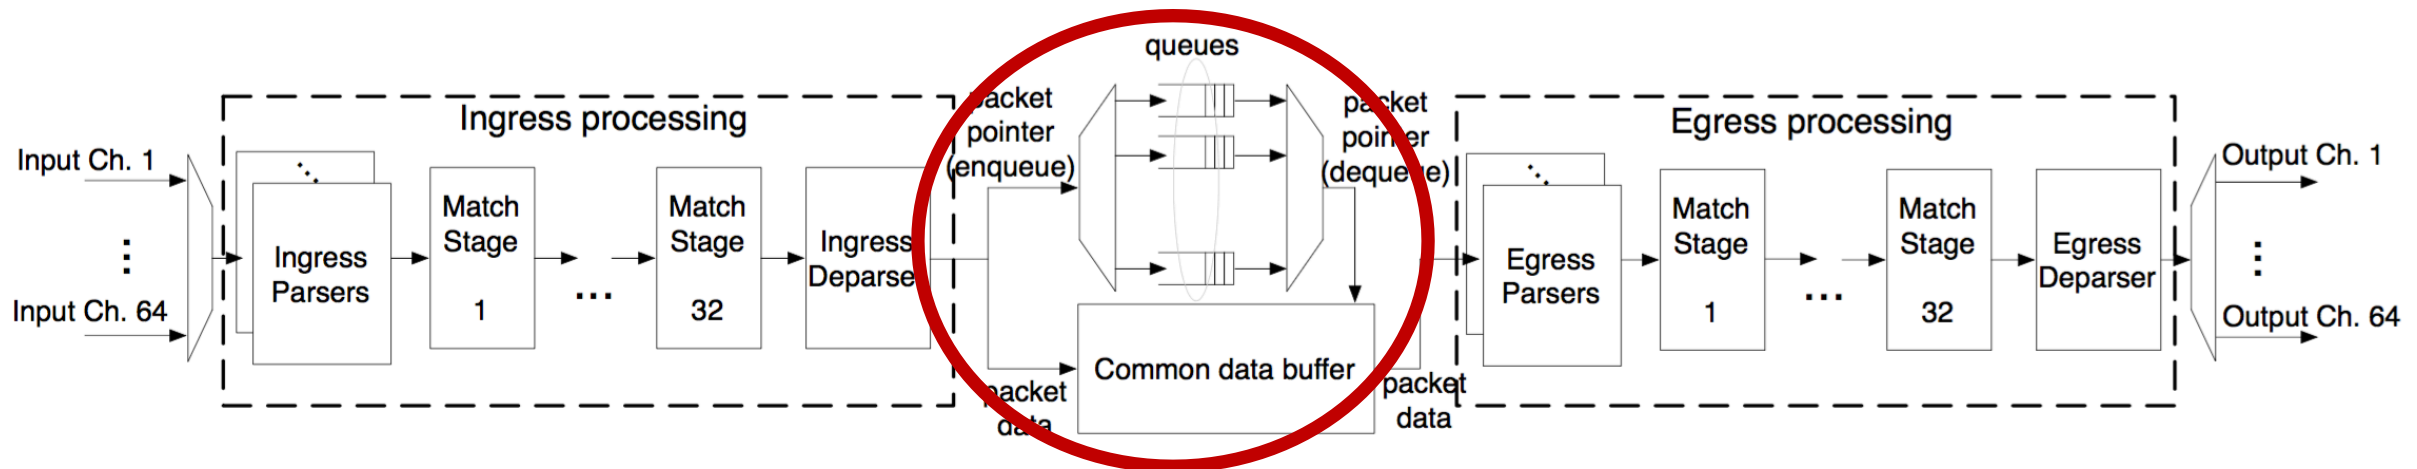
Programmable Scheduling

Lecture 18, Computer Networks (198:552)

Fall 2019

Scheduling in switch pipelines

- Packets wait in buffers/queues until serviced
- Two possibilities: **Input-queued** vs. **output-queued**
- Suppose there are pkts on port 1 to both 2 and 3
 - But suppose port 2 is clogged
 - Port 1's packets towards port 3 should not be delayed (**HOL block**)
- Better to have queues represent output port contention



Why care about packet scheduling?

- Significantly influences how packets are treated regardless of the endpoint transport
 - Implementations of **Quality of Service (QoS)** within large networks
 - Implications for **net neutrality** debates
- Intellectually interesting and influential (“top 10”) question
 - Classic Demers et al paper (WFQ) has **~ 1500** citations
 - Important connections to sched literature (e.g., job scheduling)
- Scheduling algorithms influence many daily life decisions 😊

Scheduling vs. Buffer Management

- How packets **enter** vs. how packets **leave** the switch buffer
 - Typical buffer management: Tail-drop
- How should buffer memory be partitioned across ports?
- Static partitioning?
 - **Inefficient**: even if port 1 has nothing to send, might drop port 2
- Also want **fair sharing** of buffer
 - If output port 1 is congested, why should port 2 traffic suffer?
- State of the art: **dynamic buffer sharing algorithms**



Fair Resource Allocation

Allocate *how?* among *who?*

Fair and efficient use of a resource

- Suppose n users share a **single resource**
 - Like the bandwidth on a single link
 - E.g., 3 users sharing a 30 Gbit/s link
- What is a **fair** allocation of bandwidth?
 - Suppose user demand is “elastic” (i.e., unlimited)
 - Allocate each a $1/n$ share (e.g., 10 Gbit/s each)
- But **fairness is not enough**
 - Which allocation is best: [5, 5, 5] or [18, 6, 6]?
 - [5, 5, 5] is fair but [18, 6, 6] is **more efficient**
 - What about [5, 5, 5] vs. [22, 4, 4]?

Fair use of a single resource

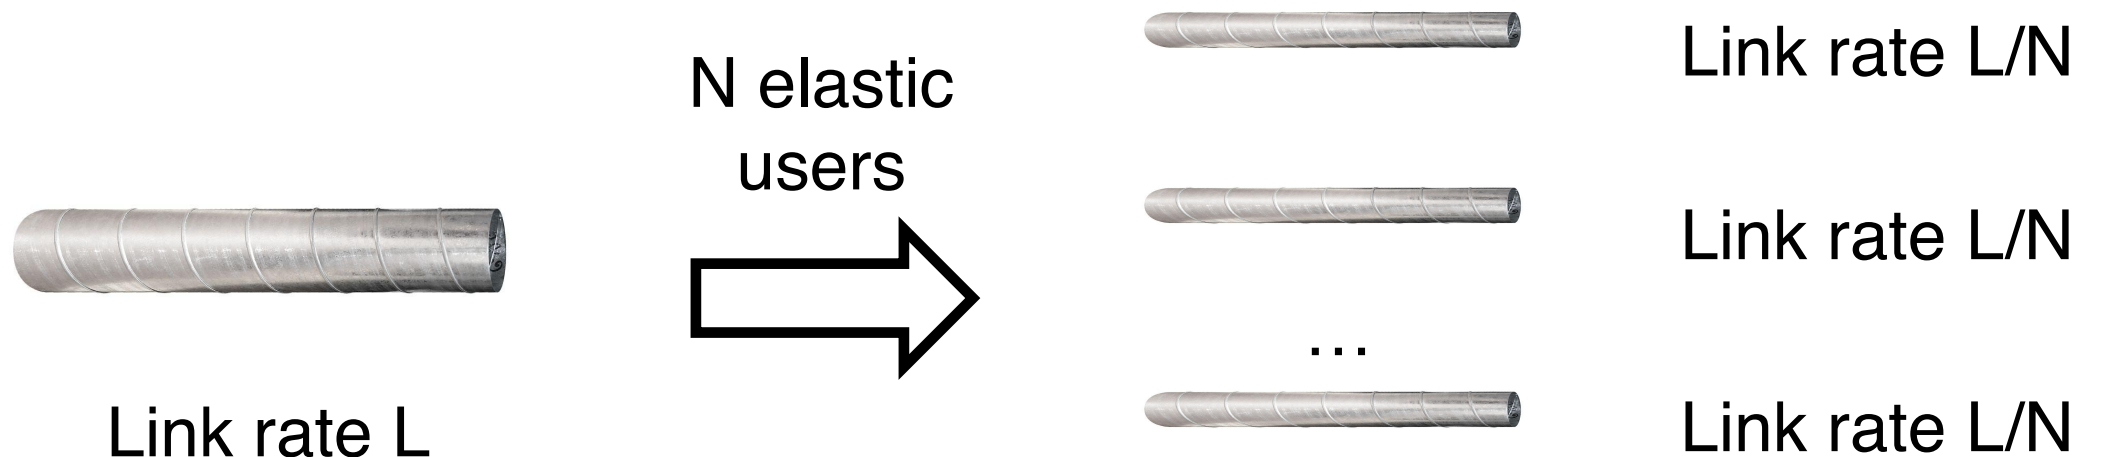
- What if some users have **inelastic** demand?
 - E.g., 3 users where 1 user only wants 6 Gbit/s
 - And the total link capacity is 30 Gbit/s
- Should we still do an “equal” allocation?
 - E.g., [6, 6, 6]
 - But that leaves 12 Gbps **unused**
- Should we allocate **in proportion to demand**?
 - E.g., 1 user wants 6 Gbps, and 2 each want 20 Gbit/s
 - Allocate [4, 13, 13]?
- Or, give the **least demanding user** all she wants?
 - E.g., allocate [6, 12, 12]?

Max-min fairness

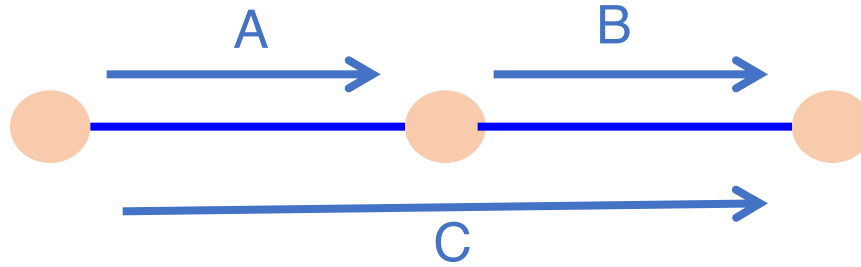
- Protect the less fortunate
 - Any attempt to *increase* the allocation of one user necessarily *decreases* the allocation of another user with equal or lower allocation
- Fully utilize a bottleneck resource
 - If demand exceeds capacity, the link is fully used

Max-min fairness for a single resource

- **Progressive filling algorithm** (also called **waterfilling**)
 - Grow all rates until some users stop having demand
 - Continue increasing all remaining rates until link is fully utilized
- If all users have elastic demands, single resource shared evenly



Allocation over multiple resources



Three users A, B, and C
Two 30 Gbit/s links

- Maximum throughput: **[30, 30, 0]**
 - Unfair: total throughput of 60, but user C starves
- Max-min fairness: **[15, 15, 15]**
 - Inefficient: everyone gets equal share, but throughput is just 45
- *Proportional fairness*: **[20, 20, 10]**
 - Allocate inversely proportional to **resource use per bit**
 - C is penalized for using two busy links, as opposed to one

Allocate fairly among *who*?

Abstract entity:
a flow

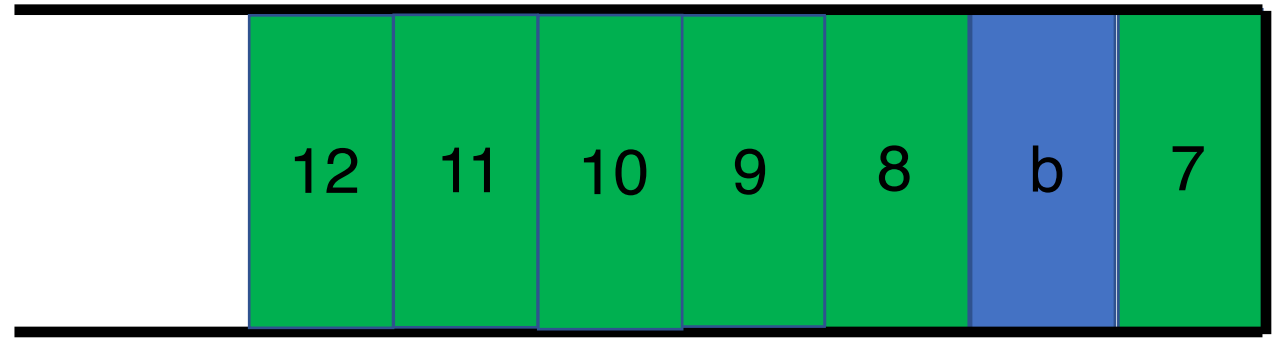
- Traffic sources?
 - Web servers, video servers, etc. need more than their fair share
- Traffic destinations?
 - Vulnerable to malicious sources denying service to receivers
- Source-destination pairs?
 - Can open up connections to many destinations
- Application flows? (i.e., src + dst + transport ports)
 - Malicious app can start up many such flows
- Administrative entities? (e.g., Rutgers NetID, ISP, ...)
 - How should a router identify packets belonging to an entity?

Packet Scheduling Algorithms

Which packet to send next? (order)

When to send the next packet? (timing)

A taxonomy



- Granularity of allocation
 - Per-packet vs. per-flow vs bit-by-bit
- Pre-emptive vs. non-pre-emptive
 - Do you interrupt the current packet/flow if another shows up?
- Size-aware vs. unaware
 - Do you consider flow or packet sizes in scheduling?
- Class-based (strict priority) vs. shared
 - Are some flows strictly higher priority than others?
- Work-conserving vs. non-work-conserving
 - Do you always use spare link capacity when there is demand?

Examples of scheduling algorithms (1/3)

- FIFO over packets
- Round-robin over packets of different flows
- **Shortest Remaining Processing Time (SRPT)**
 - Flow-size-aware allocation which strictly prioritizes short flows
 - Also called **shortest flow first** in some contexts
 - Flow-size-unaware variant may predict demand using known flow size distribution

Examples of scheduling algorithms (2/3)

- Processor sharing

- Assume each flow gets a fair share of the link every unit of time
- Ideal: each flow starts receiving service immediately upon arrival

- Rate limiting

- Non-work-conserving: flow can't send even if more demand than limit

- Class-based strict prioritization

- Pre-determined flow classes with strict priorities over each other
- Starve low priority flows if higher priority flows are always sending

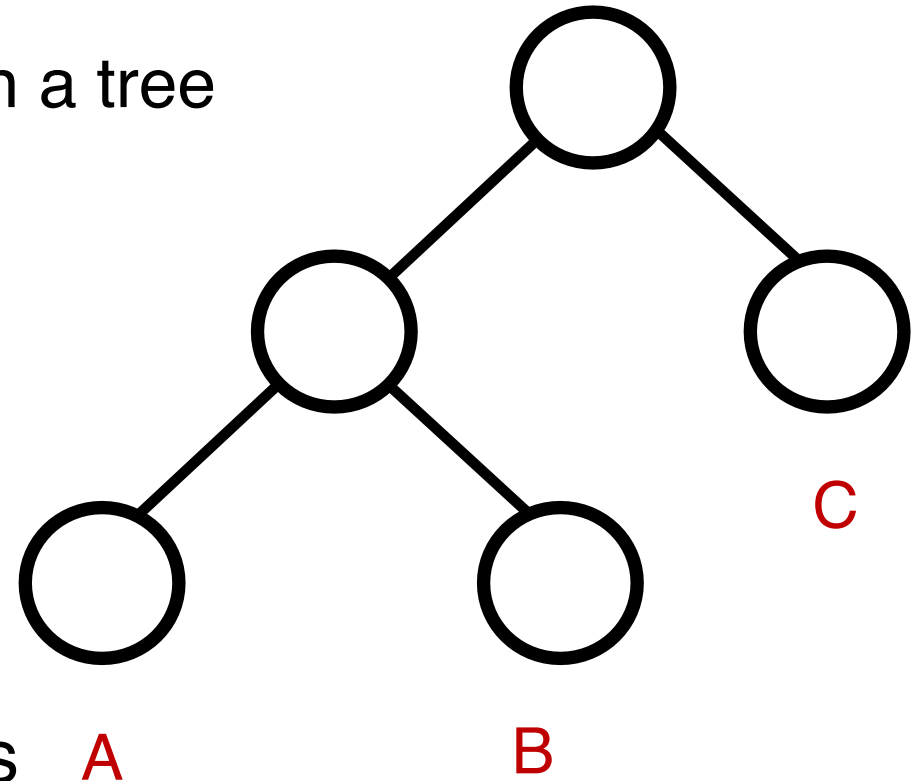
Examples of scheduling algorithms (3/3)

- Hierarchical policies

- Arrange existing scheduling policies in a tree

- Example:

- Rate-limit $A + B$
- Fair-share among A and B within limit
- Fair-share among $A+B$ and C

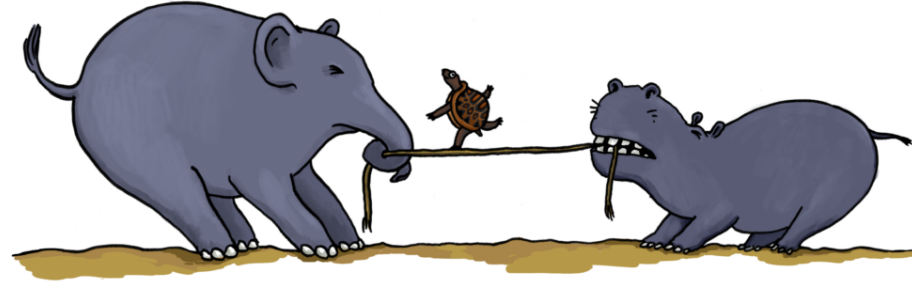


- Complex multi-tenant isolation policies

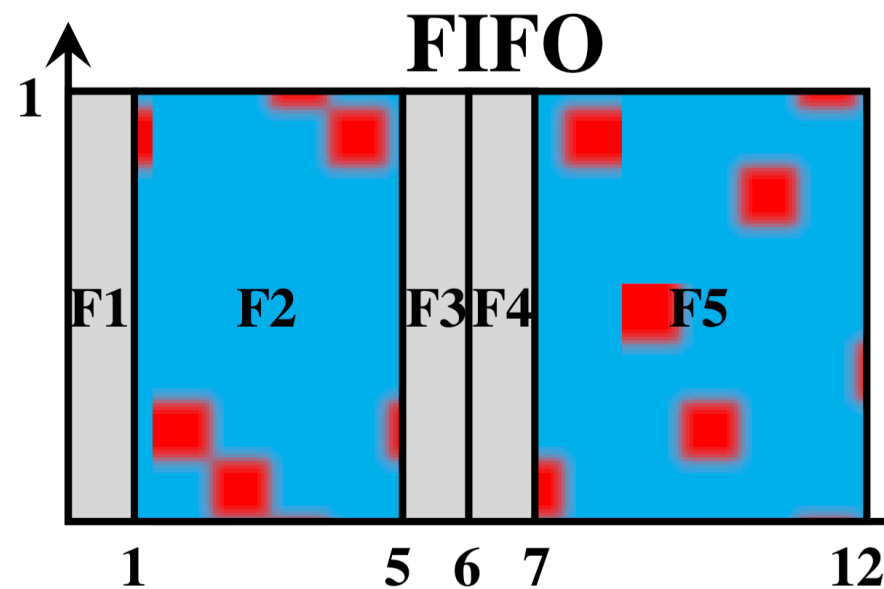
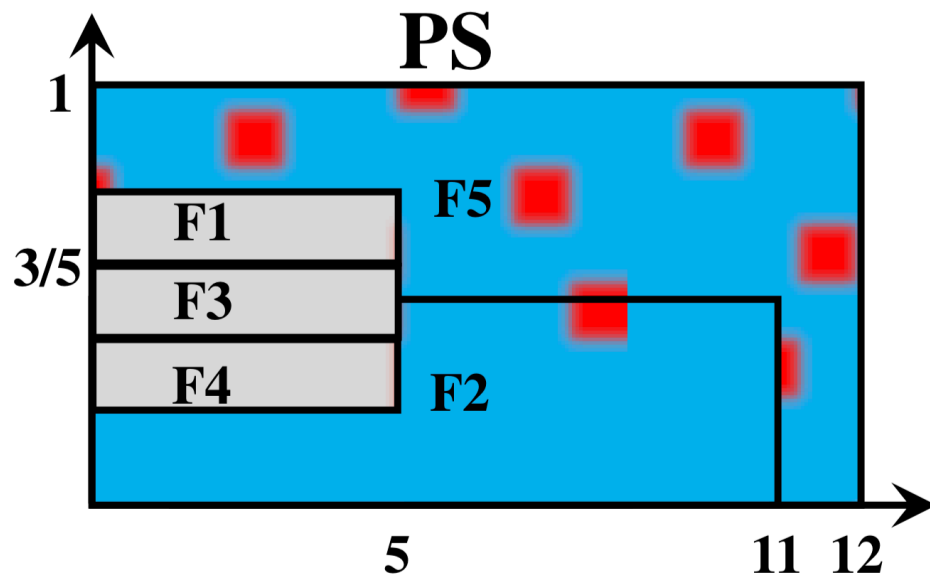
There's not one optimal scheduling

Flow ID	Size	Class
F1	1	1
F2	4	2
F3	1	1
F4	1	1
F5	5	2

Multiplexing
Avoids HOL
blocking



Serialization
Reduces flow
completion
time



Exercise: When does a flow finish?

- Consider a mix of “long” and “short” flows arriving at a Q
 - Ex: A flow may have as few as 2 packets or as many as 10^5
- Suppose a scheduling algorithm provides each flow:
 - An average **per-packet delay** d (e.g., 50 ms)
 - An average **link bandwidth share** t (e.g., 10 Mbit/s)
- Which among d & t determines
 - when a short flow finishes?
 - when a long flow finishes?

Push In First Out (PIFO)

A common primitive for many scheduling algorithms

Key ideas

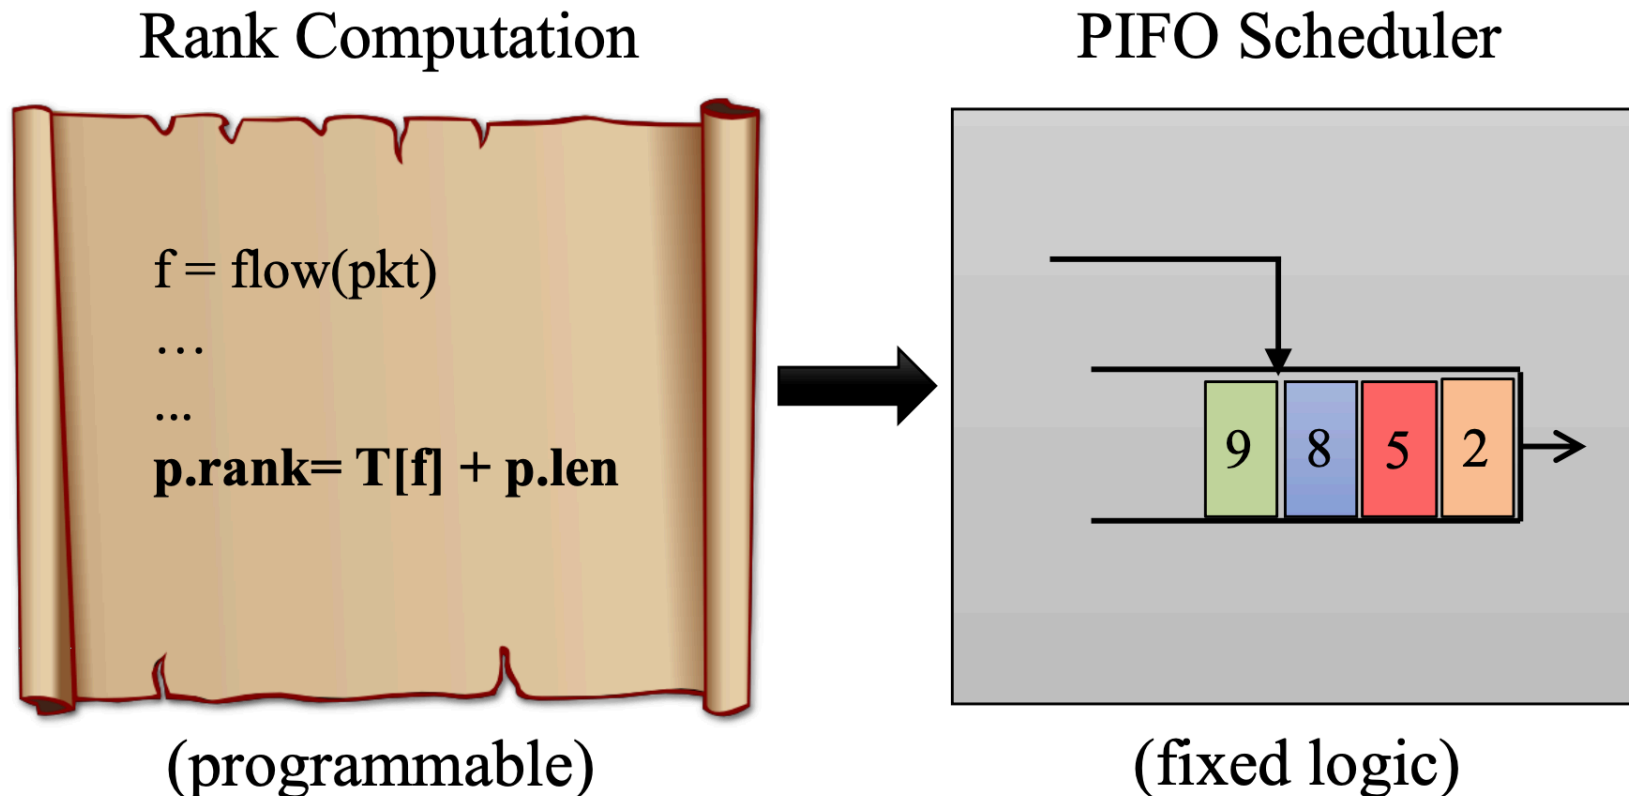
- Scheduling algorithms determine **order** and **timing** of packet departures from a queue
- Typically, **relative order of buffered packets** doesn't change upon new packet arrivals



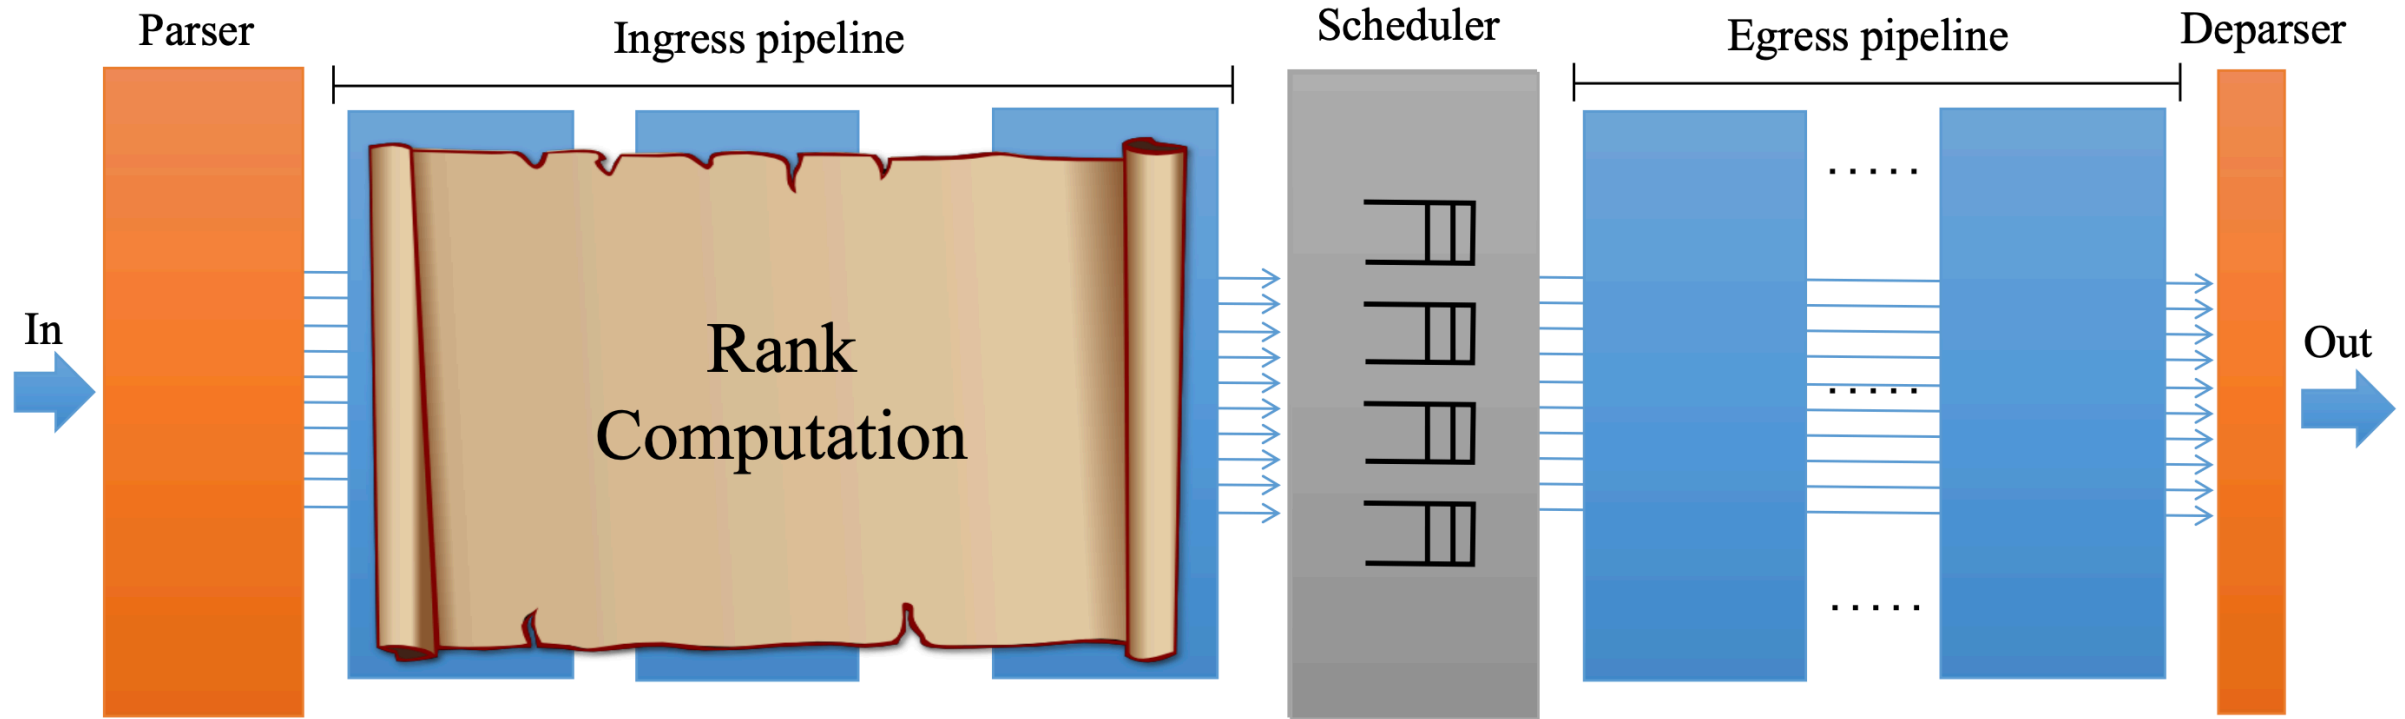
- Implement scheduling through a **priority-queue-based data structure (PIFO)**
 - Push-In: pkts have arbitrary **ranks**; push anywhere into queue
 - First-Out: always dequeue from the head of the queue

Programmable Scheduler

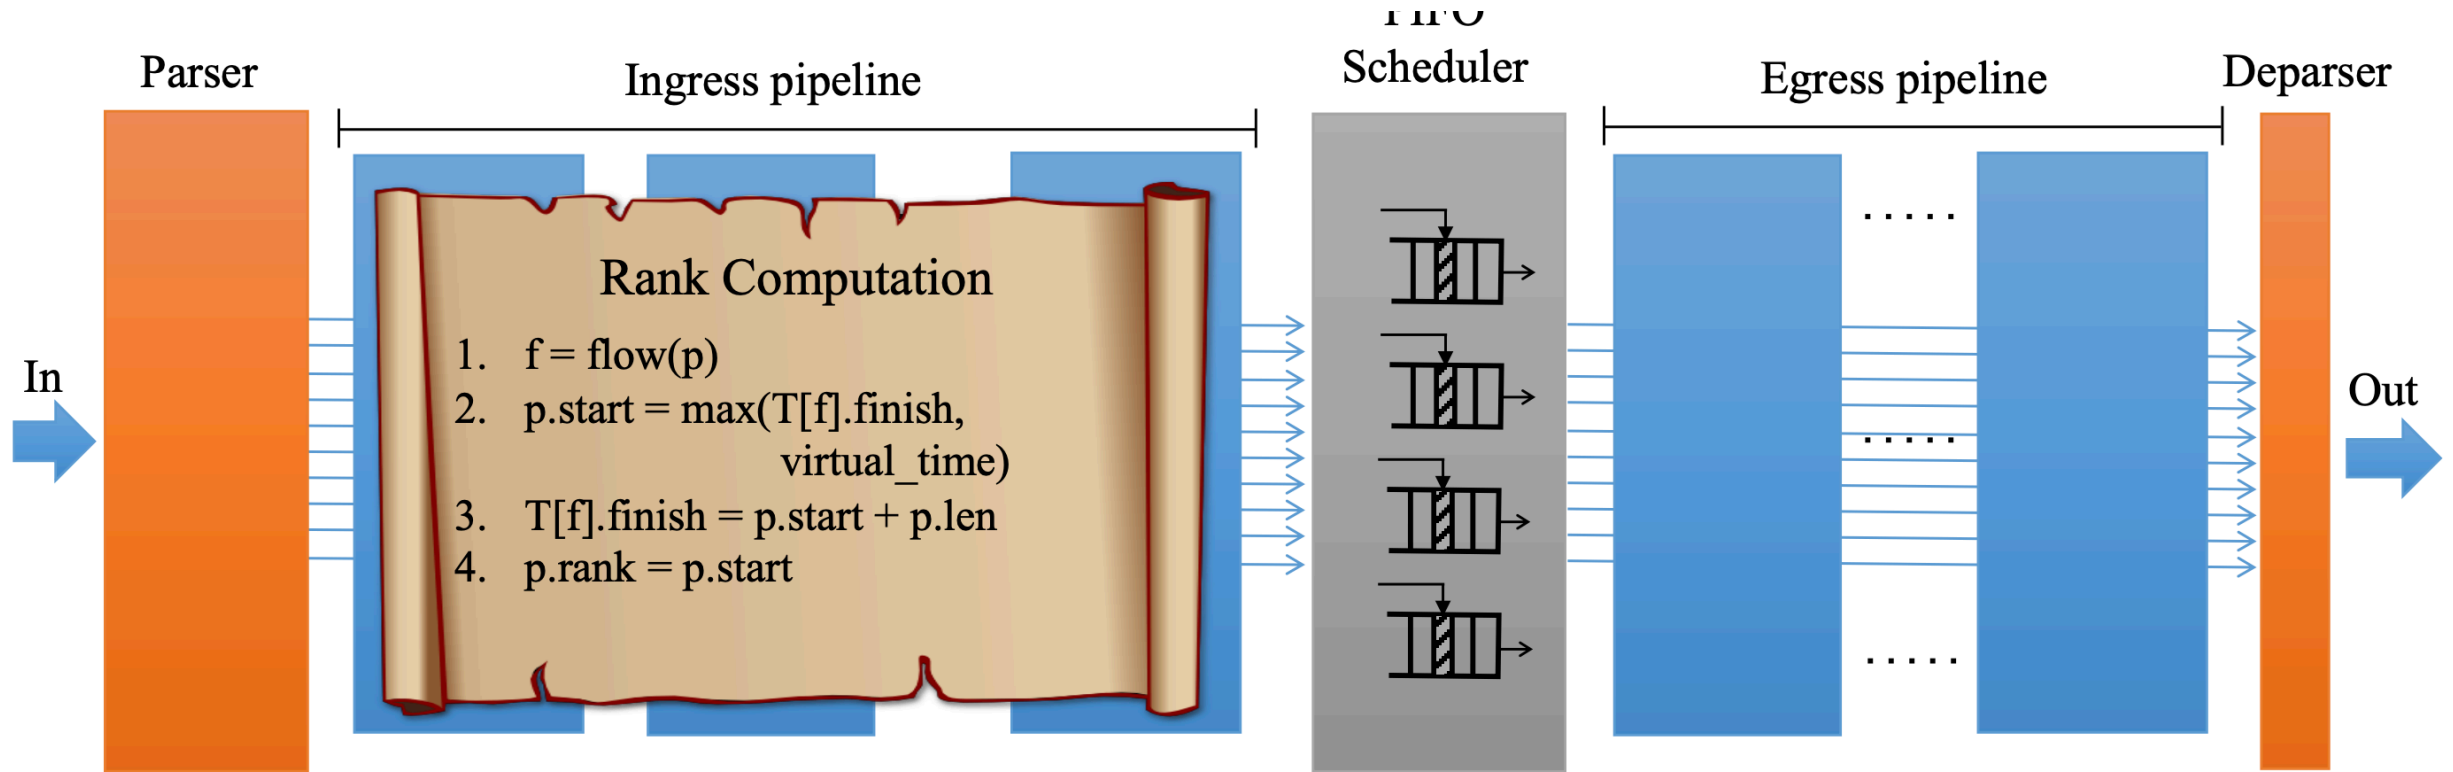
- To program the scheduler, program the **rank computation**



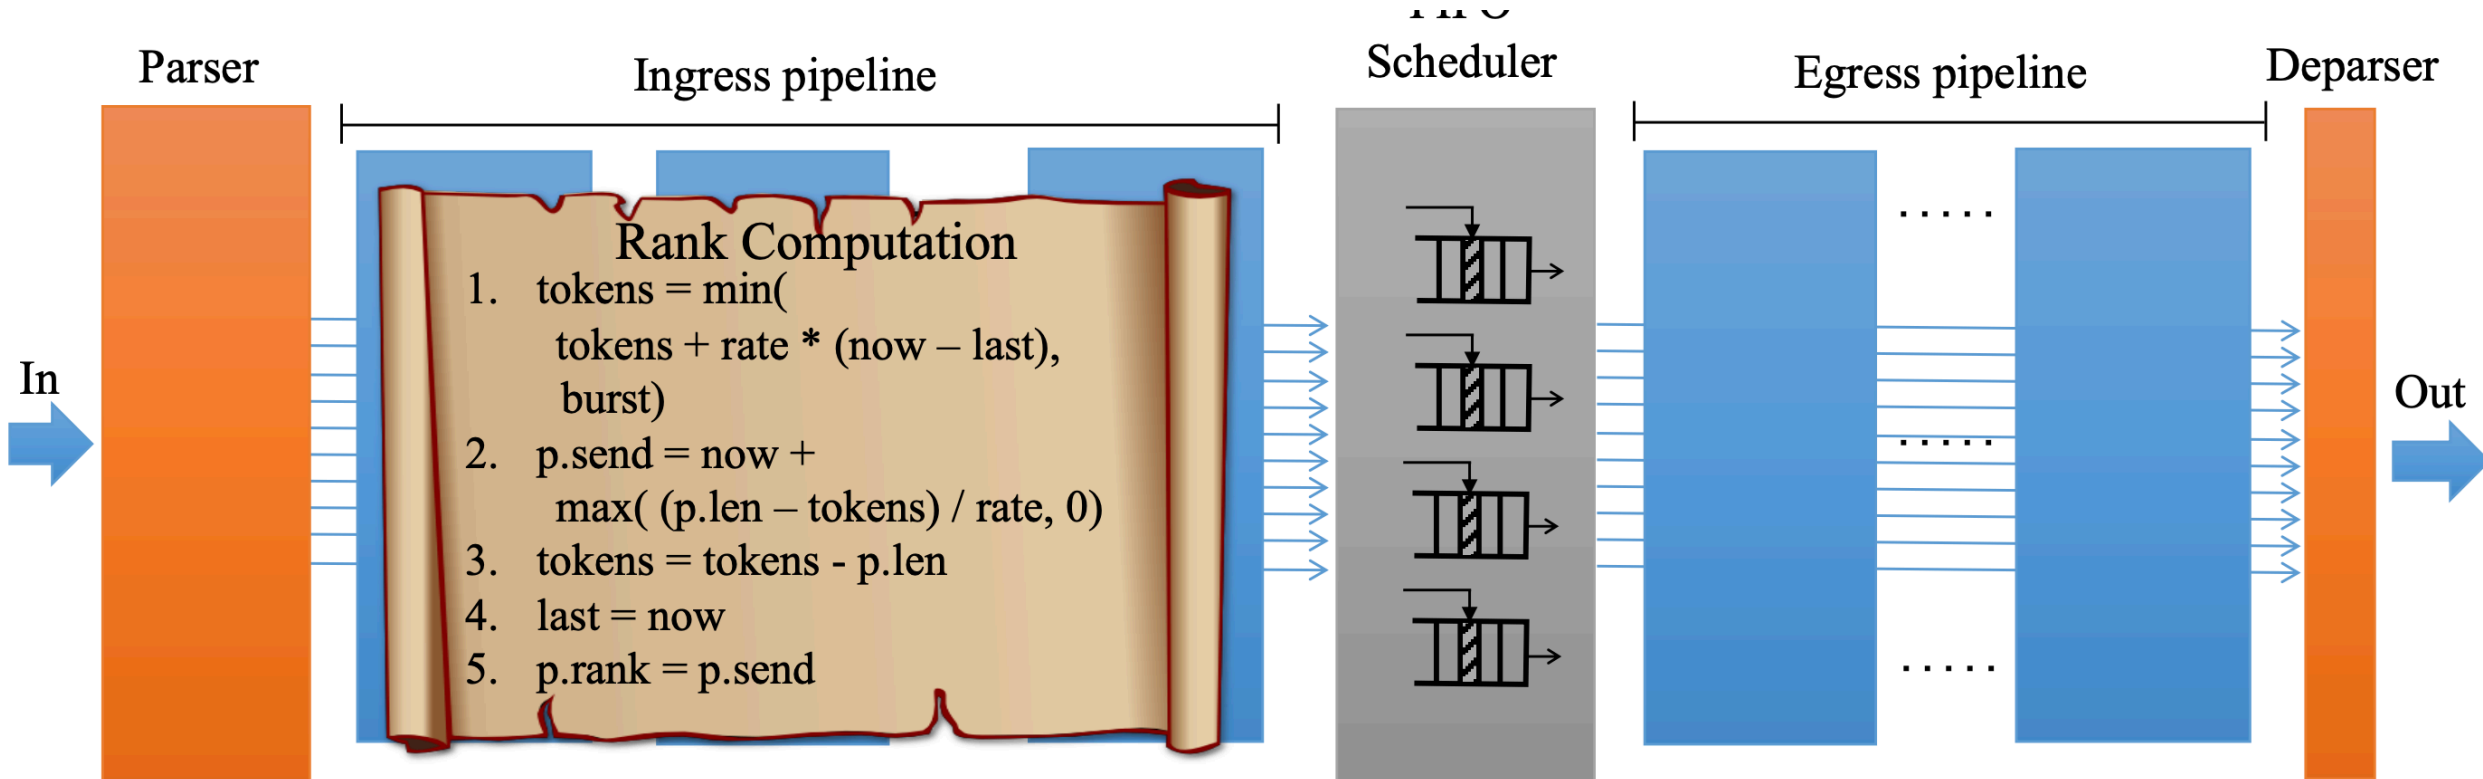
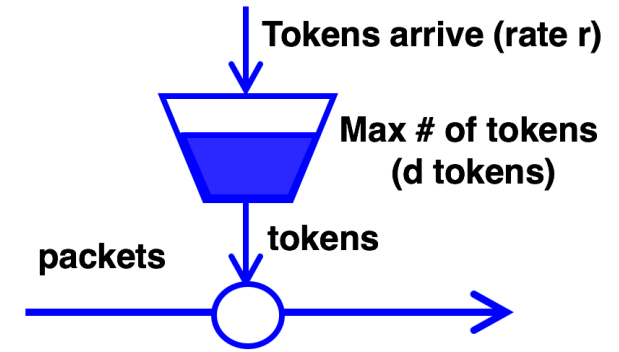
Programmable scheduling in the pipeline



Fair queueing

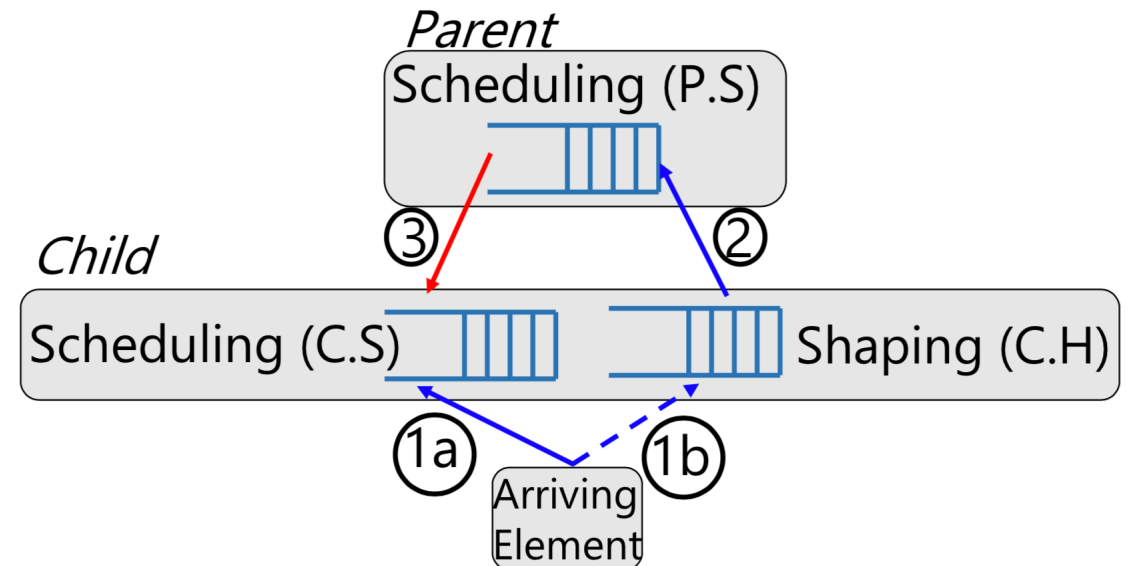
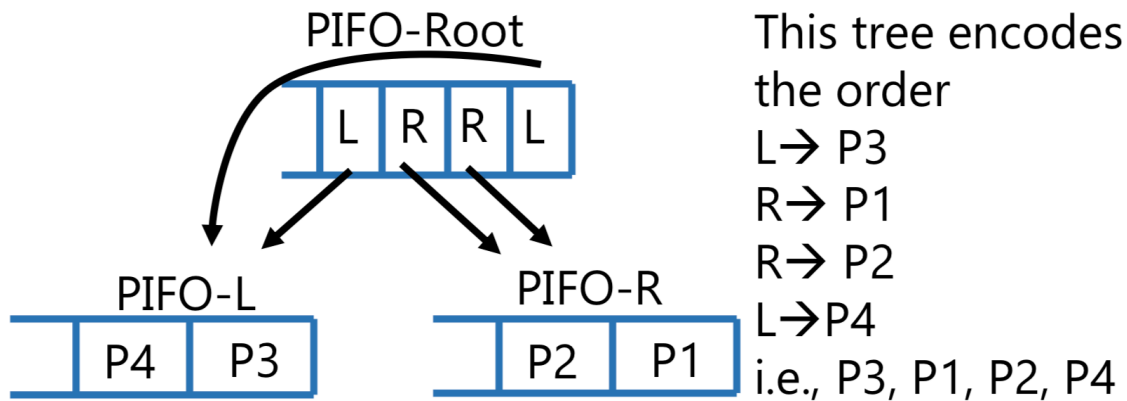


Token-bucket rate limiting



Generalizations

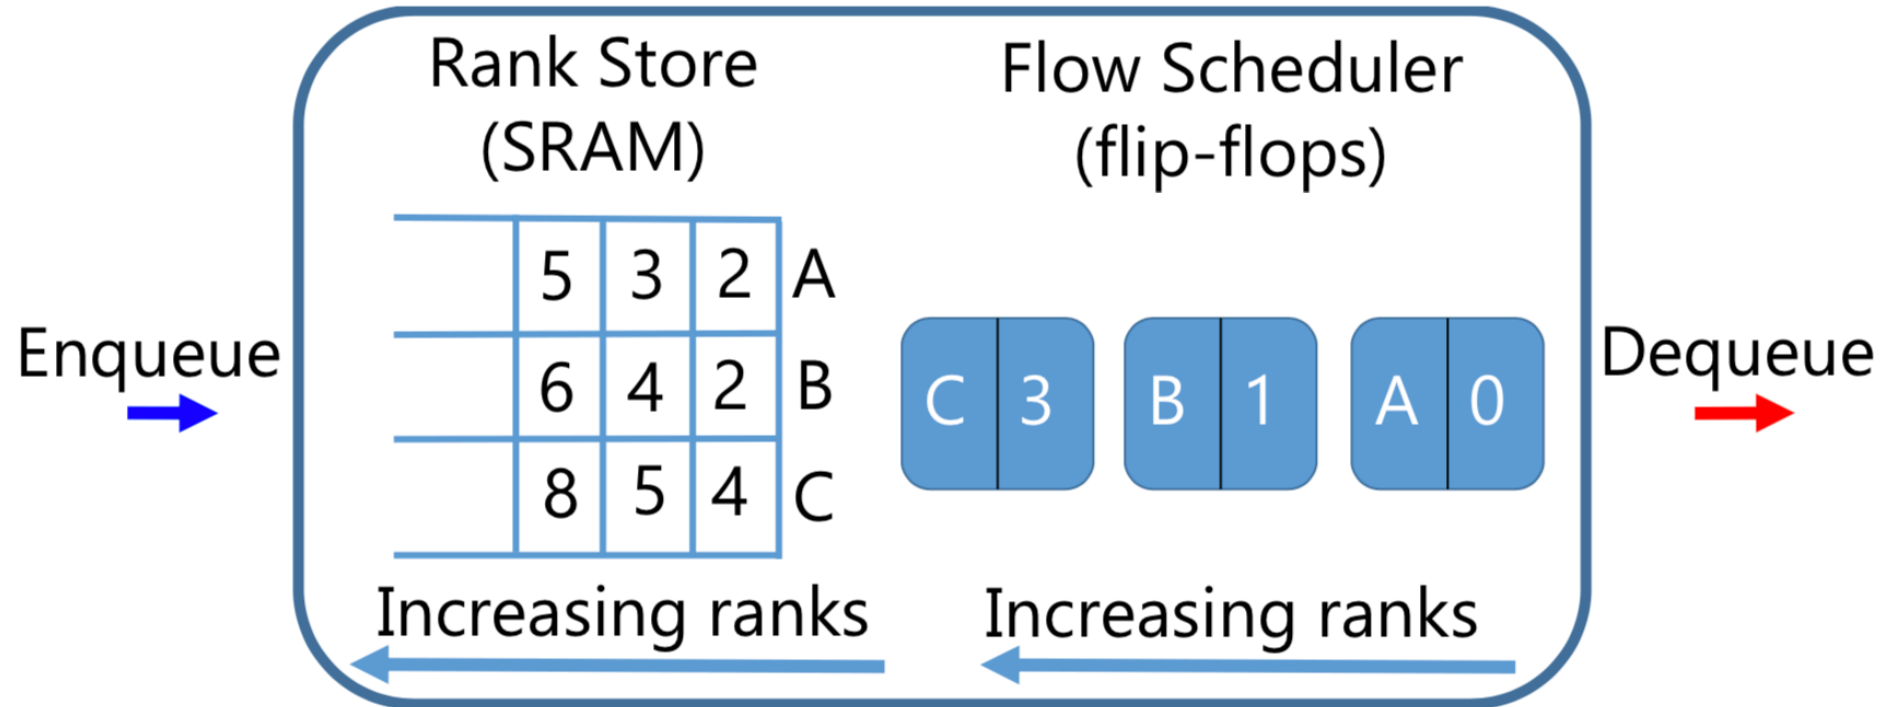
- Use a **hierarchy of PIFOs** to implement hierarchical policies
- Use a **shaping PIFO** to implement non-work-conserving scheduling policies



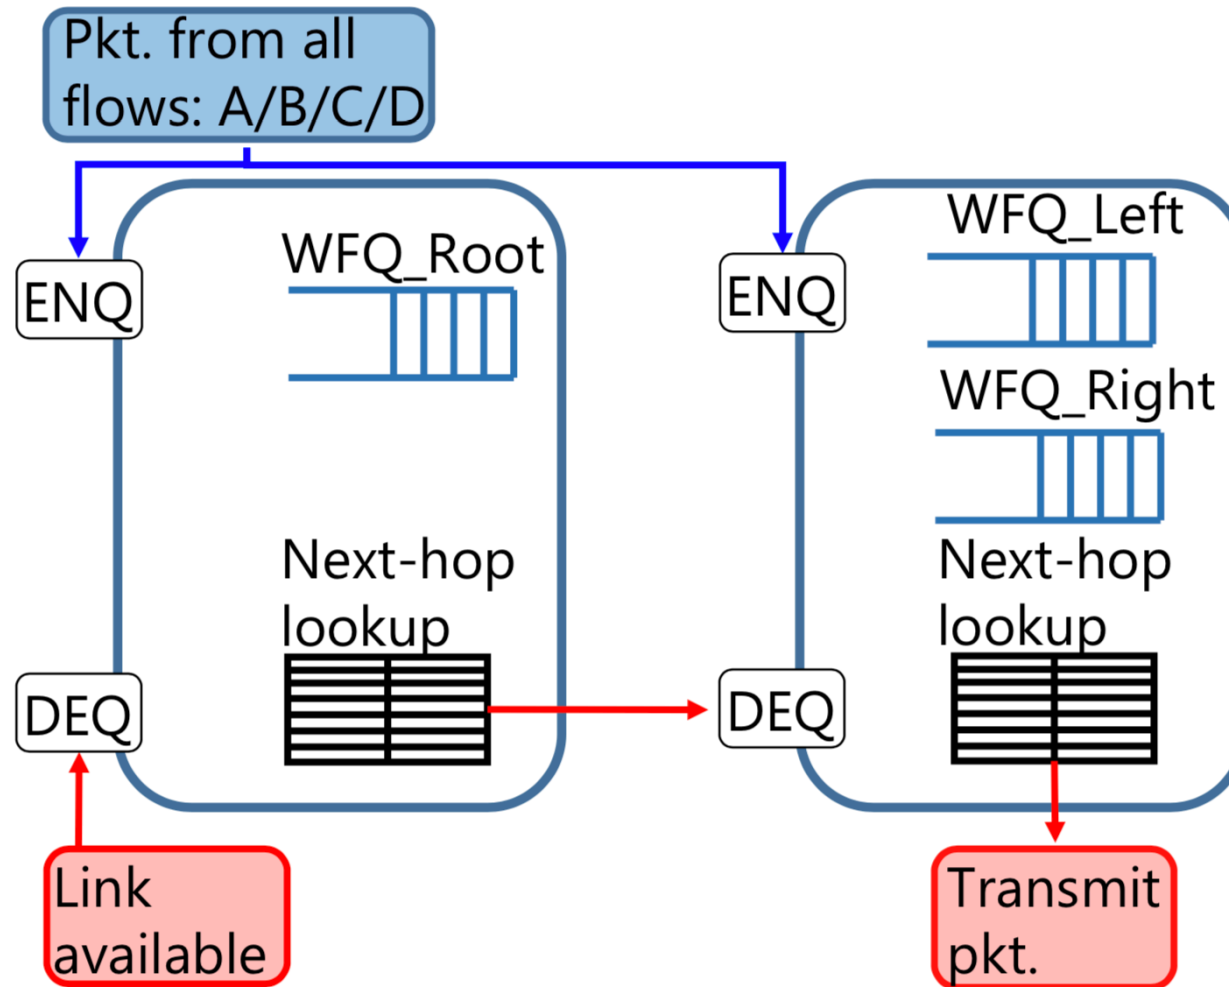
Implementation: Challenges

- Packet buffer is big! Schedule among all those packets?
 - Maintaining a sorted list of 64K packets?
 - Instead, make flow-level scheduling decisions
 - With FIFO order among packets of a given flow
- Sorting even just flows at line rate
 - Line-rate insertion and removal from hardware priority queue with 1000s of flow elements
 - Use fast flip-flops and pipelined logic

Implementation: Flow-level PIFOs



Encoding HPFQ in a PIFO mesh



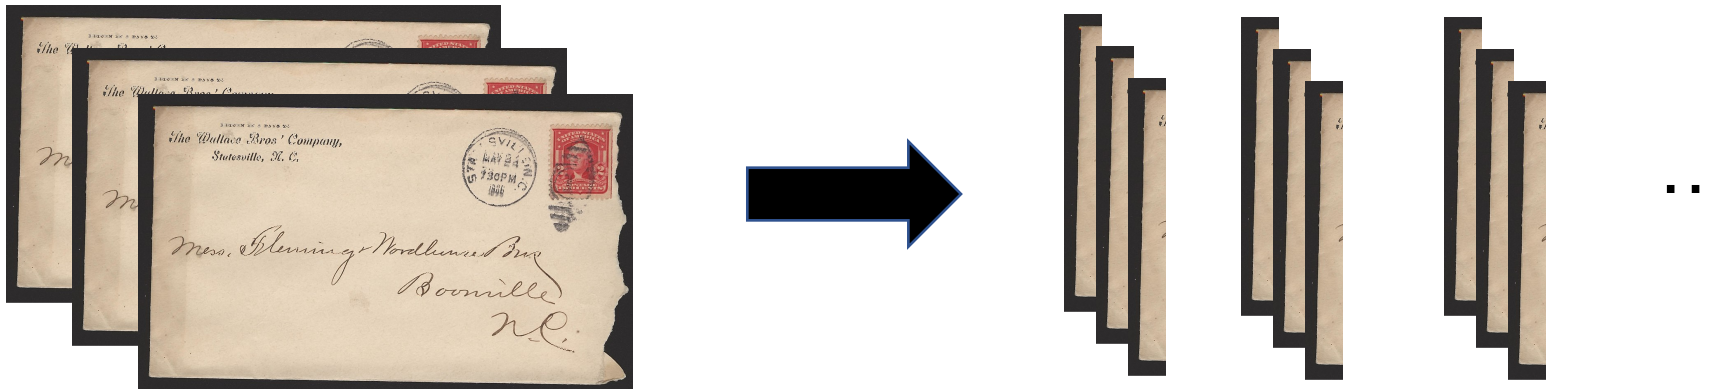
Fair Queueing

ACM SIGCOMM '89

Alan Demers, Srinivasan Keshav, and Scott Shenker

An ideal to emulate: Processor sharing

- Fair-share bandwidth in the most fine-grained fashion possible
 - If there are N active flows, each flow gets $1/N^{\text{th}}$ of the link rate
 - “Bit by bit round robin” (BR)
- Implementing BR directly on routers is unrealistic. Why?
 - One reason: consider the processing of the bit downstream
 - E.g., where to route the bit?



Emulate bit-by-bit round robin (BR)?

- How about round robin over packets?
- Unfair! A flow can use larger packets and gain larger bandwidth
- Instead, determine when a packet would finish with BR
 - Depends only on packet arrival time & # of active flows
 - Let's call this the “virtual finish time”
- FQ: Transmit packets in the order of the virtual finish times
 - Buffer management: drop pkt of the flow with the largest backlog

Deficit Round Robin

- Router-friendly implementation of a WFQ scheme

