

# Network Program Synthesis


Lecture 23, Computer Networks (198:552)

Fall 2019

# Review: Control/data plane separation

**Control plane**  
Determine how to move packets

**Data plane**  
Move packets



# Review: RMT router data plane

## Parser Program

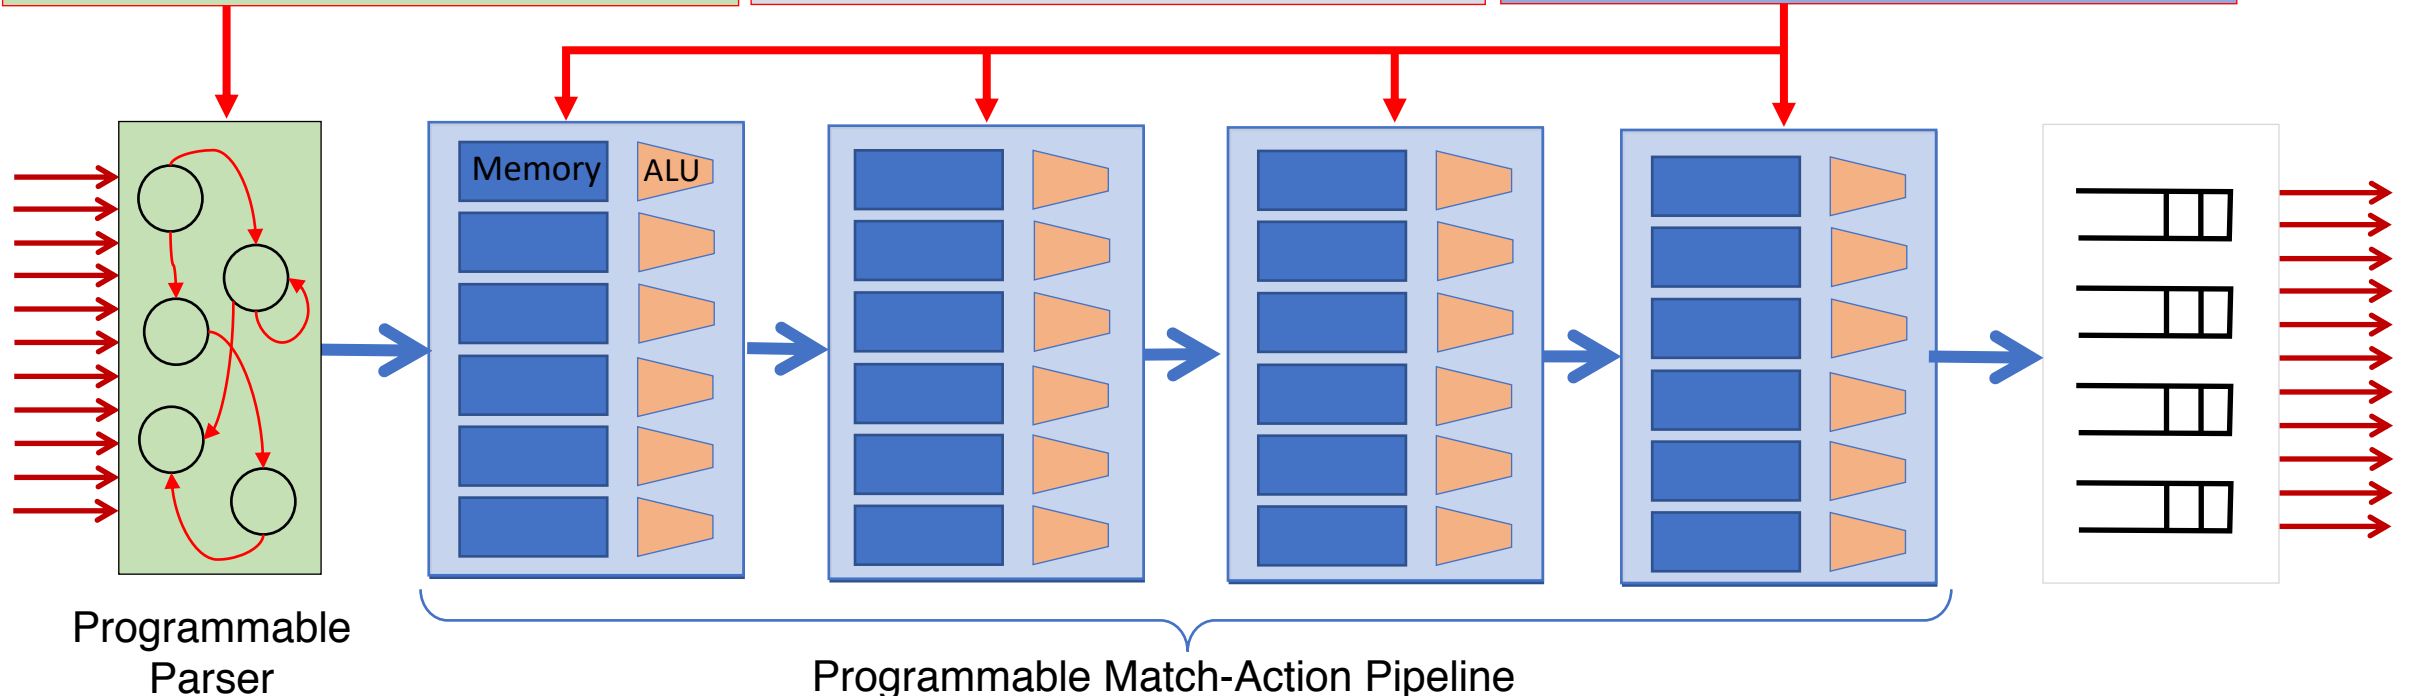
```
parser parse_ethernet {  
  extract(ethernet);  
  return switch(ethernet.ethertype) {  
    0x8100 : parse_vlan_tag;  
    0x0800 : parse_ipv4;  
    0x8847 : parse_mpls;  
    default: ingress;  
  }  
}
```

## Header and Data Declarations

```
header_type ethernet_t { ... }  
header_type l2_metadata_t { ... }  
  
header ethernet_t ethernet;  
header vlan_tag_t  
vlan_tag[2];  
metadata l2_metadata_t l2_meta;
```

## Tables and Control Flow

```
table port_table { ... }  
  
control ingress {  
  apply(port_table);  
  if (l2_meta.vlan_tags == 0) {  
    process_assign_vlan();  
  }  
}
```




# Review: Verification

for all M, does N satisfy P?




Sequence of messages:

Packets,  
Routing protocol  
Link failures



Network representation:

Data plane  
Control plane



Property of interest:

Loop freedom  
Blackholes  
Equivalence  
Many complex props...

**Decision Procedure:** An algorithm that answers yes/no

# Review: Different guises, Similar question

- **Verification:** for all  $M$ , does  $N$  satisfy  $P$ ?
- **Testing:** For the given  $M$ , does  $N$  satisfy  $P$ ?
- **Synthesis:** Given  $P$ , can you produce an  $N$  that satisfies it
  - For a given set of  $M$ ? (including for all  $M$ )
- Let  $N'$  be another network representation
- **Equivalence checking:** For all  $M$ , do  $N$  and  $N'$  behave in the same way with respect to  $P$ ?, i.e.,
  - i.e., either both satisfy  $P$  or both violate it

# Stateful data plane programs

A case study in programming networks

# Stateful processing

- Action on a packet depends on previously seen packets
- Example: send every 100<sup>th</sup> packet to a measurement server
- Example: Track DNS TTL changes
- Example: `if (pkt.field > max) { counter++; max = pkt.field }`
- Example: Flowlet-based load balancing

# Flowlet load balancing

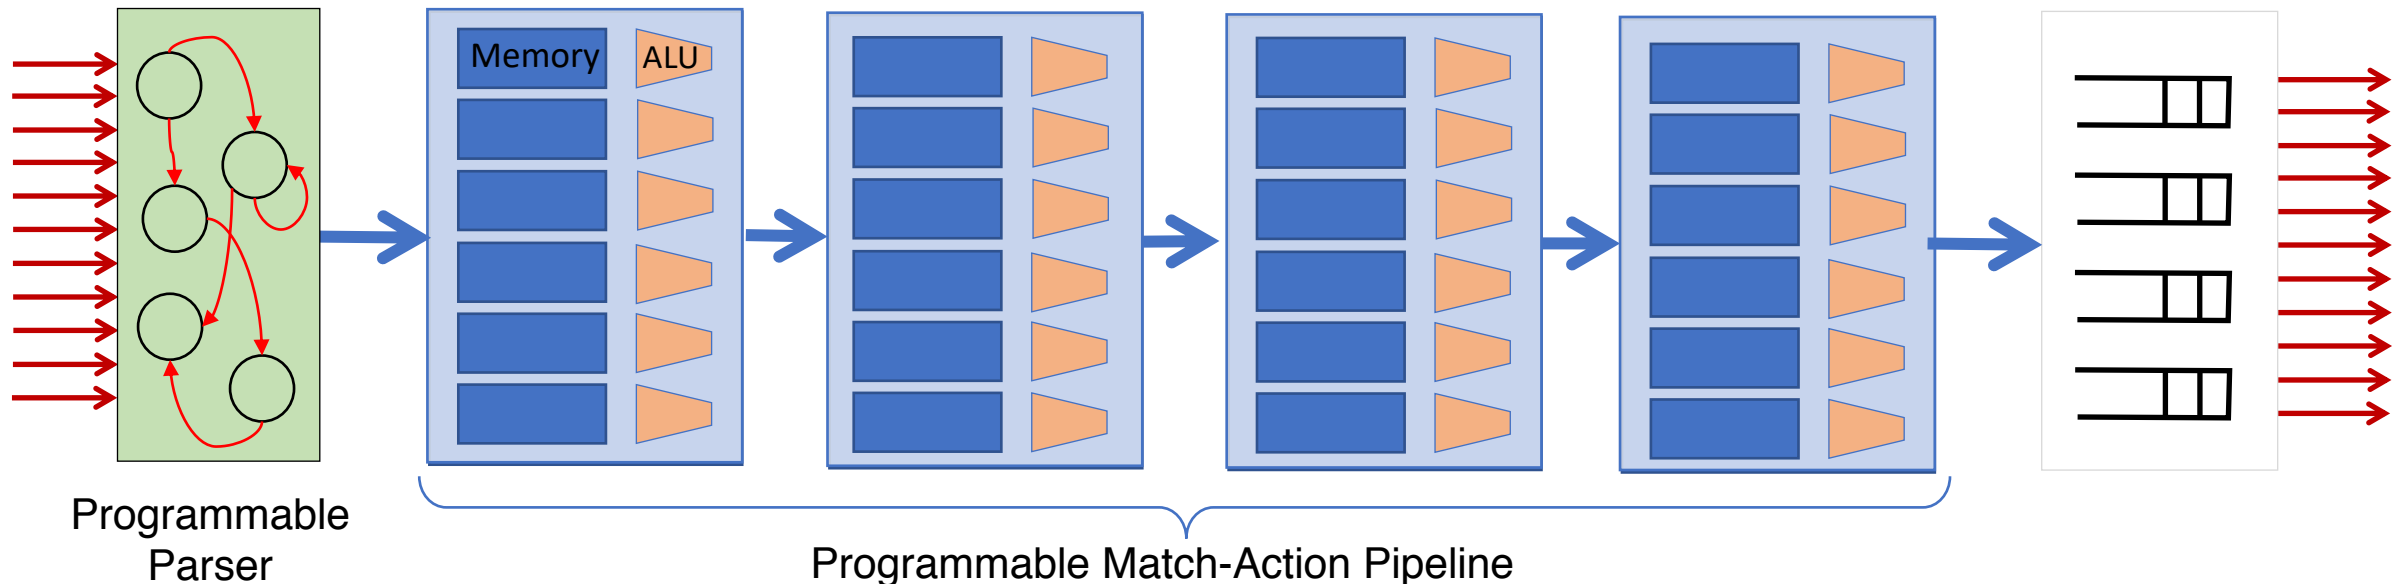
- Consider the **time of arrival** of the current packet and the last packet **of the same flow**
- If current packet arrives **threshold** later than the last packet did, consider rerouting the packet to balance load
  - Otherwise, use the same path as before
  - Q: Why?

```
if (pkt.ts - saved_ts > threshold) {  
    pkt.outport = new_random(); saved_port = pkt.outport;  
} else pkt.outport = saved_port;  
saved_ts = pkt.ts;
```



# How would you implement Flowlet LB?

- Recall that pipeline is clocked at 1 GHz
- i.e., a new packet admitted every 1 ns
- Where is the processing? Where is the state?
- How about interactions across packets?




# Abstraction: Packet transaction

- A piece of **code** along with **state** that **runs to completion** on each packet before processing the next:
  - C-like domain-specific language, Domino [sigcomm'16]
- Assume **every packet encounters the same set of actions**
  - i.e., the match on each stage is the same: superset of our needs
- Implementation must respect abstraction
  - A **compiler** must enforce atomicity for the implementation
  - Challenge: **transaction code may not fit in one pipeline stage!**
  - Challenge: **state updates must happen within 1 ns**

# Insight #1: Pipeline the stateless actions

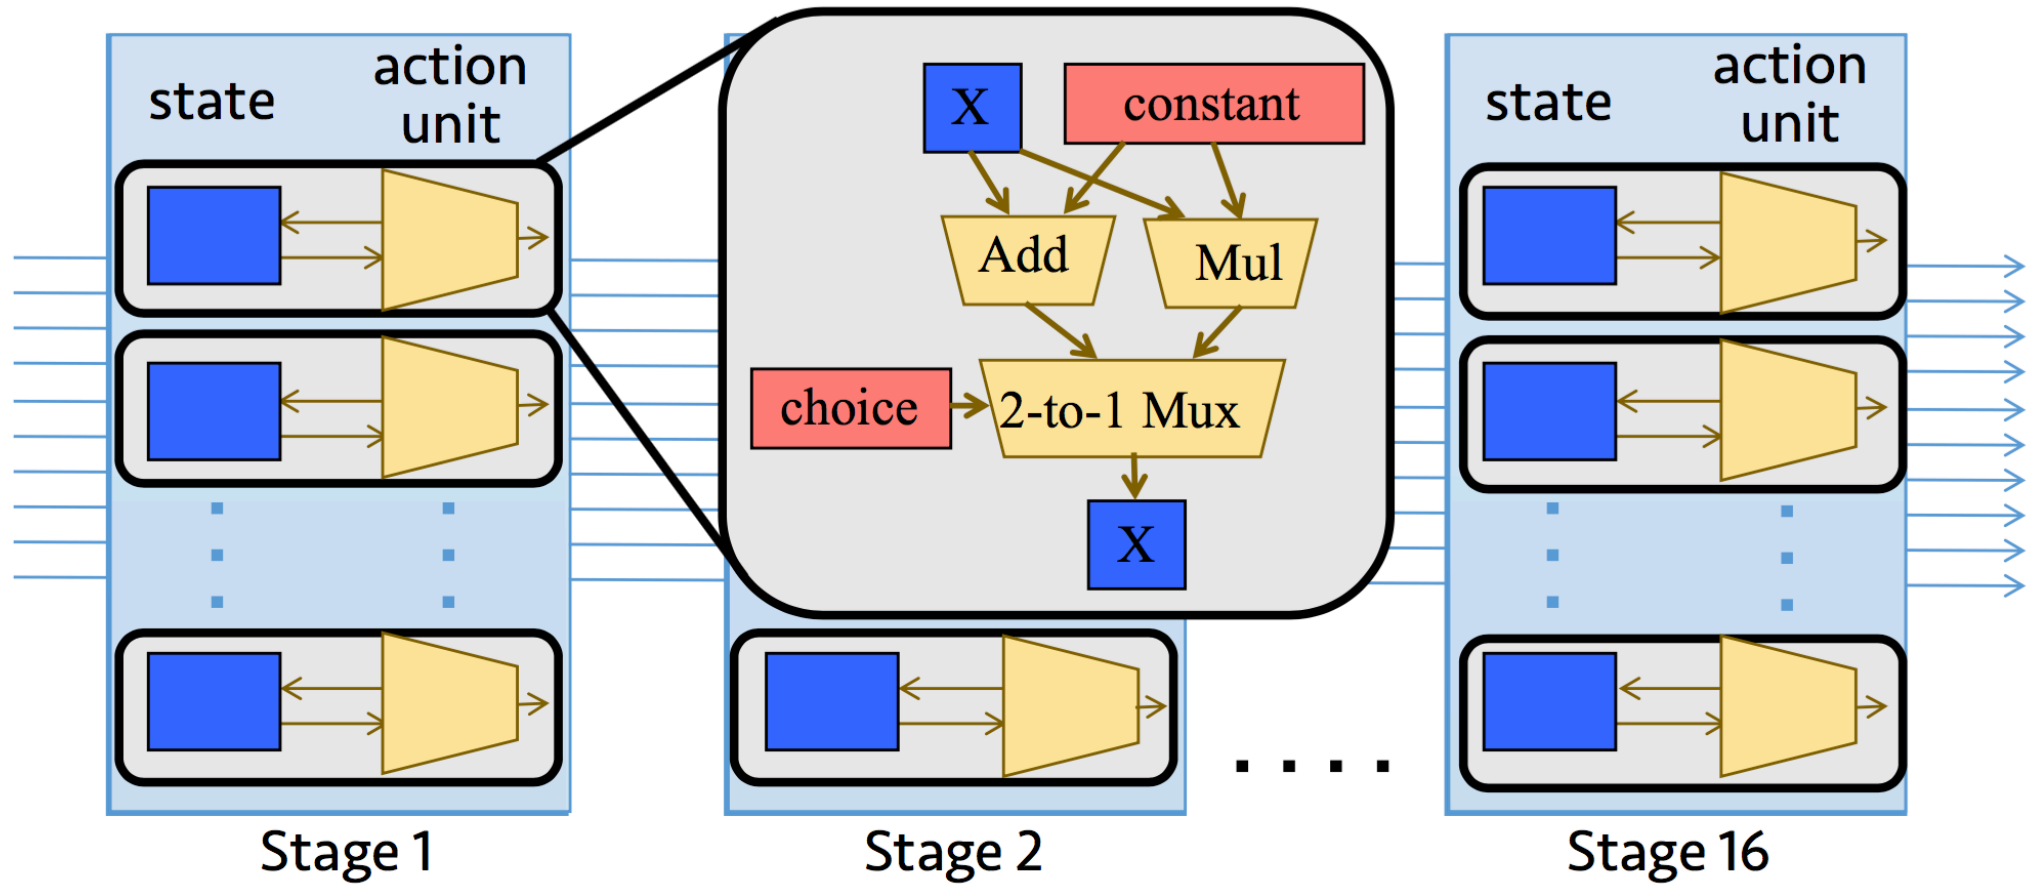
- `if (pkt.field1 + pkt.field2 == 10) { counter ++; }`



Stateless operations (whose results depend only on the current packet) can **execute over multiple stages**

**Only the stateful operation must run atomically in one pipeline stage**

# Insight #2: Fit state updates to ALUs



The **atoms** constitute the switch's action instruction set. **Run under 1 ns**

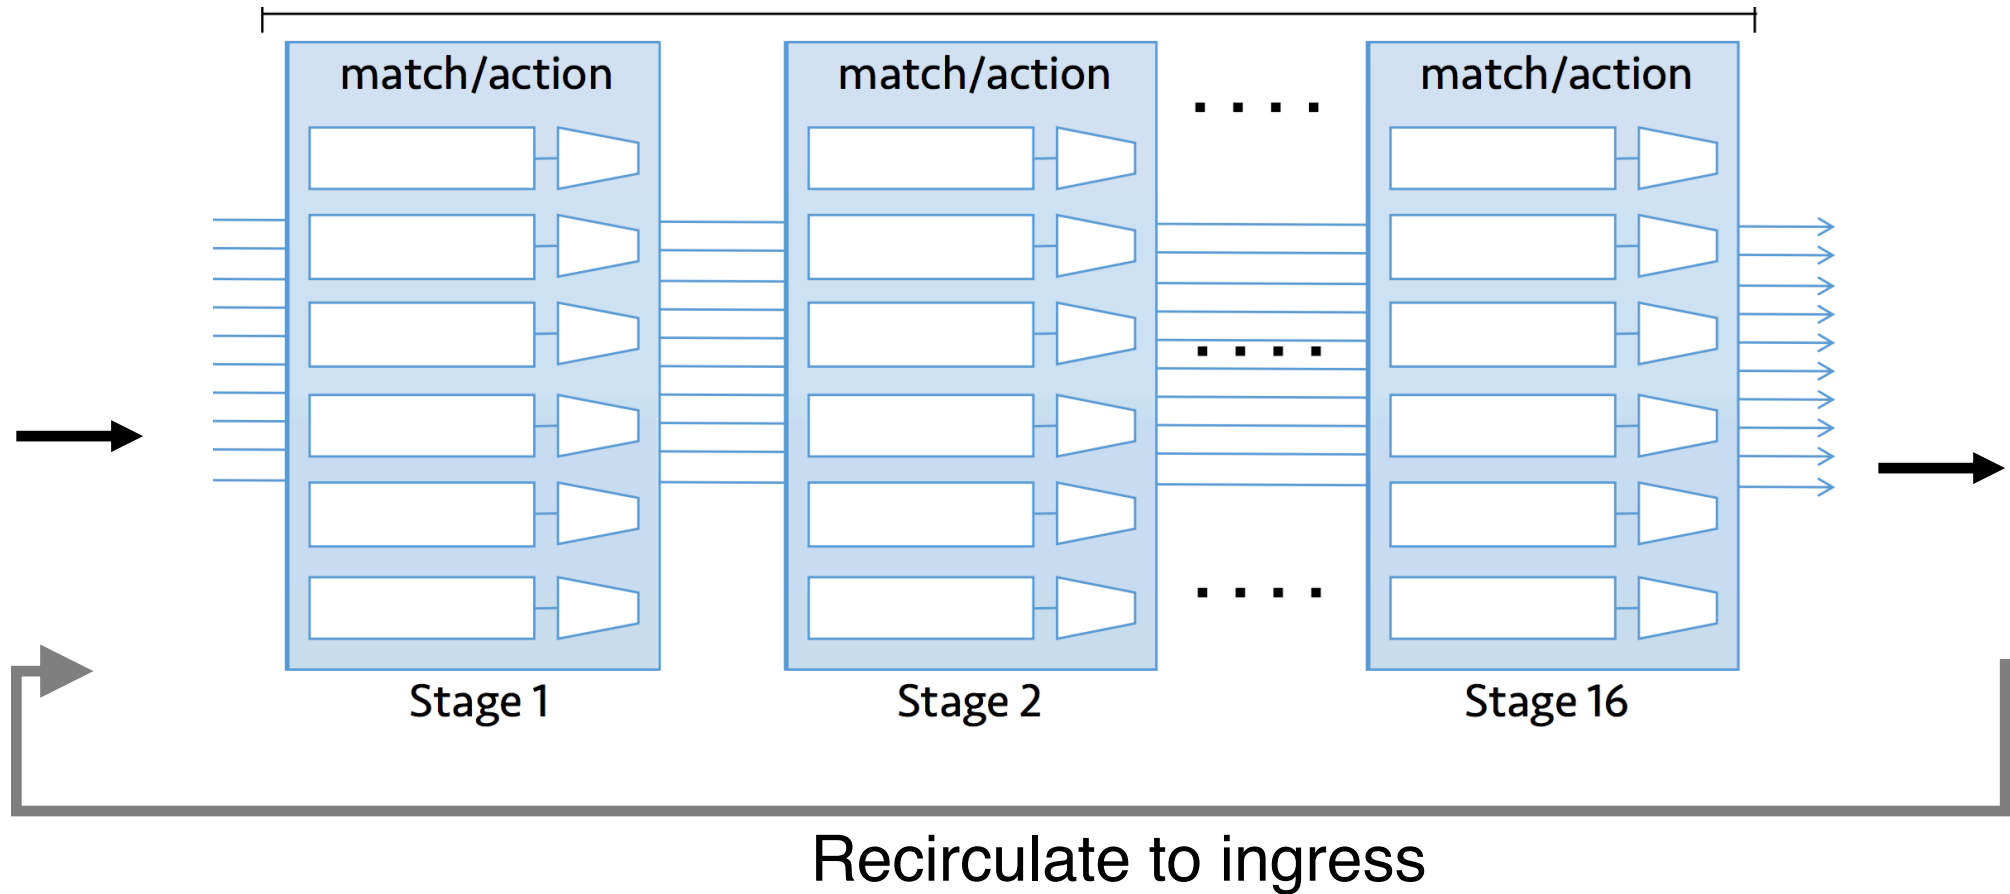
# All-or-nothing compilation

Problems with running stateful network programs

# Implementing complex policies

- What if you have a very large P4 or Domino program?
- Example: too many logical tables
  - Not enough physical stages to hold tables
- Example: logical table keys are too wide
  - Sharing memory across stages leads to paucity in physical stages
- Example: Action on each packet, e.g., domino pgm, is too large
  - Sharing compute across stages leads to paucity in physical stages
- Solution?
  - Re-circulation

# Re-circulation “extends” the pipeline



If you do this for every packet, packet throughput drops by 2X!

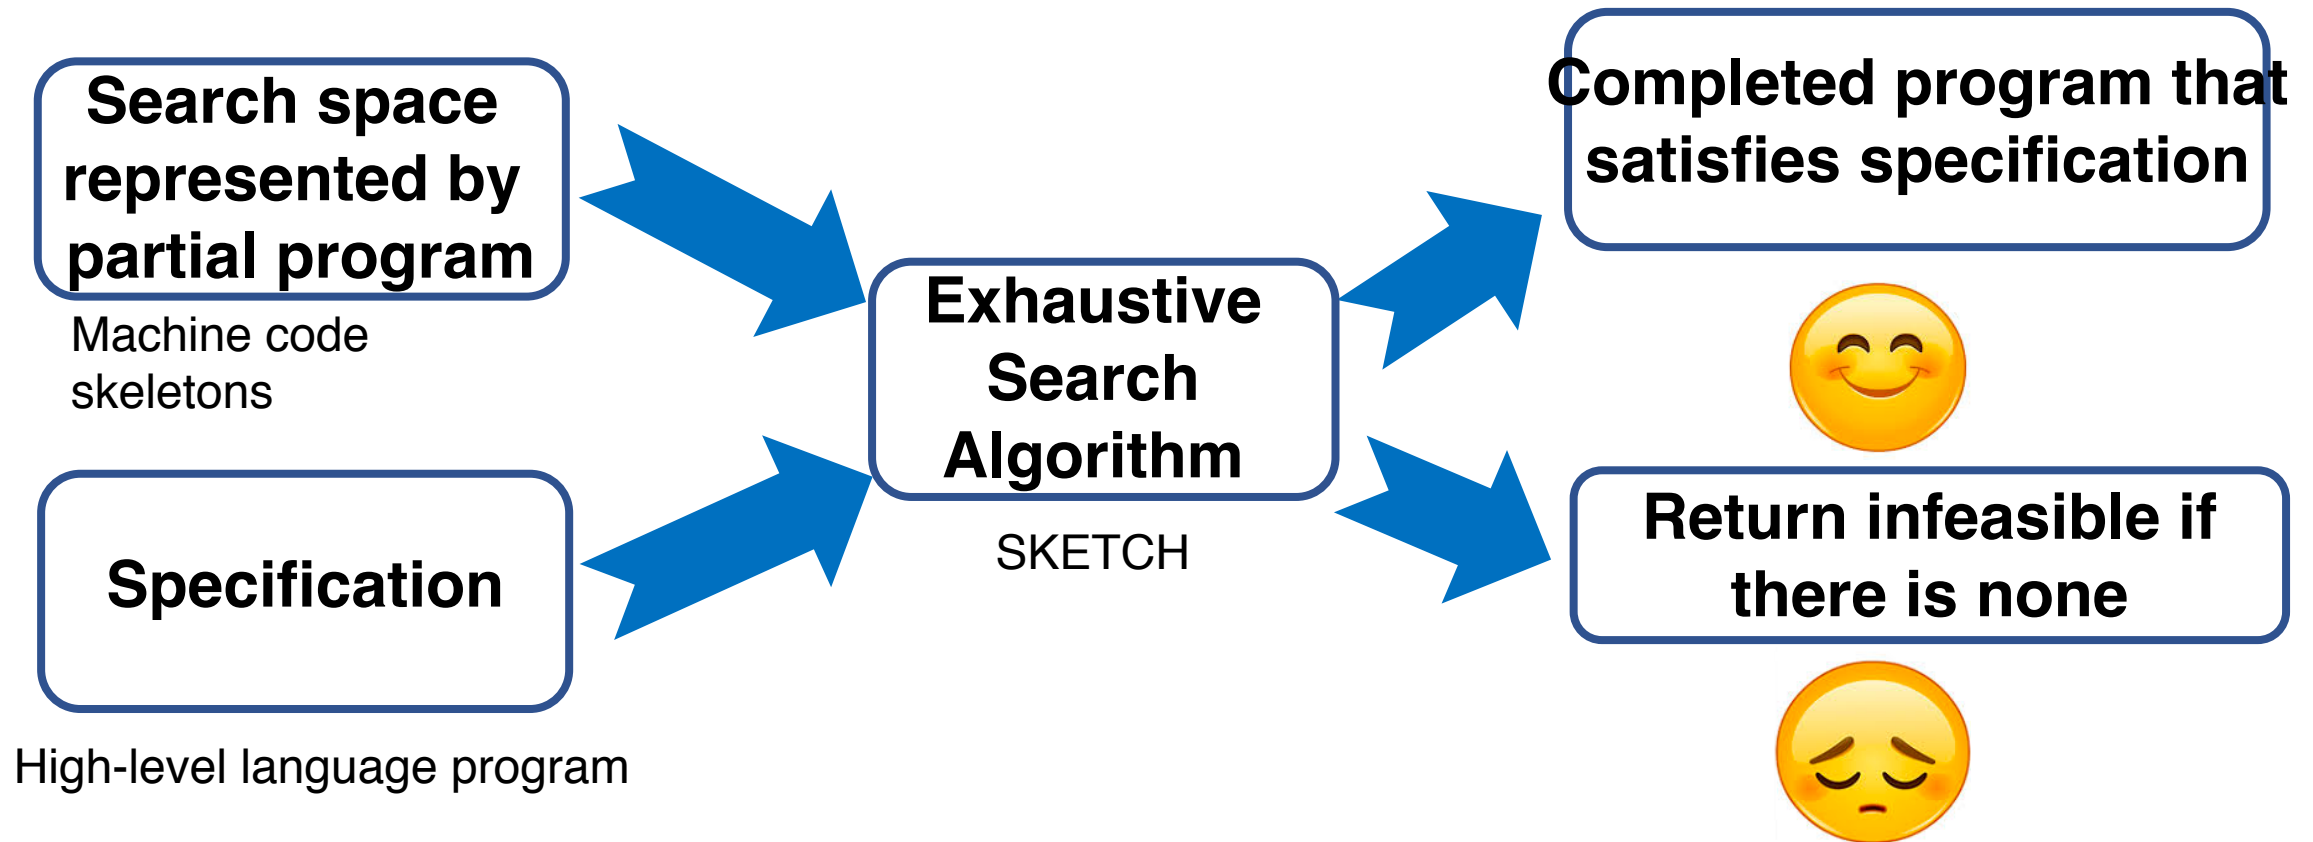
# Synthesis of stateful data plane programs

Chipmunk



# Program synthesis for code generation

Auto-generate a program from a specification **by exhaustive search**



Program synthesis as an enabling technology for pkt-processing compilers

# Example: Using SKETCH to do synthesis

```
int spec(int x) {  
    return x + x + x;  
}
```

Specification

```
int sketch1(int x)  
implements spec {  
    return x * ??;  
}
```

Partial program

```
int sketch2(int x)  
implements spec {  
    return x + ??;  
}
```

Infeasible partial program  
no hole assignment

```
int sketch1(int x)  
implements spec {  
    return x * 3;  
}
```

Completed program

Hole representing  
unknown constant





Programmer

# Chipmunk: An Overview



Compiler developer

Program as a packet transaction in Domino

```
if (count == 10):  
    count = 0  
    pkt.sample = 1  
else:  
    count++  
    pkt.sample = 0
```

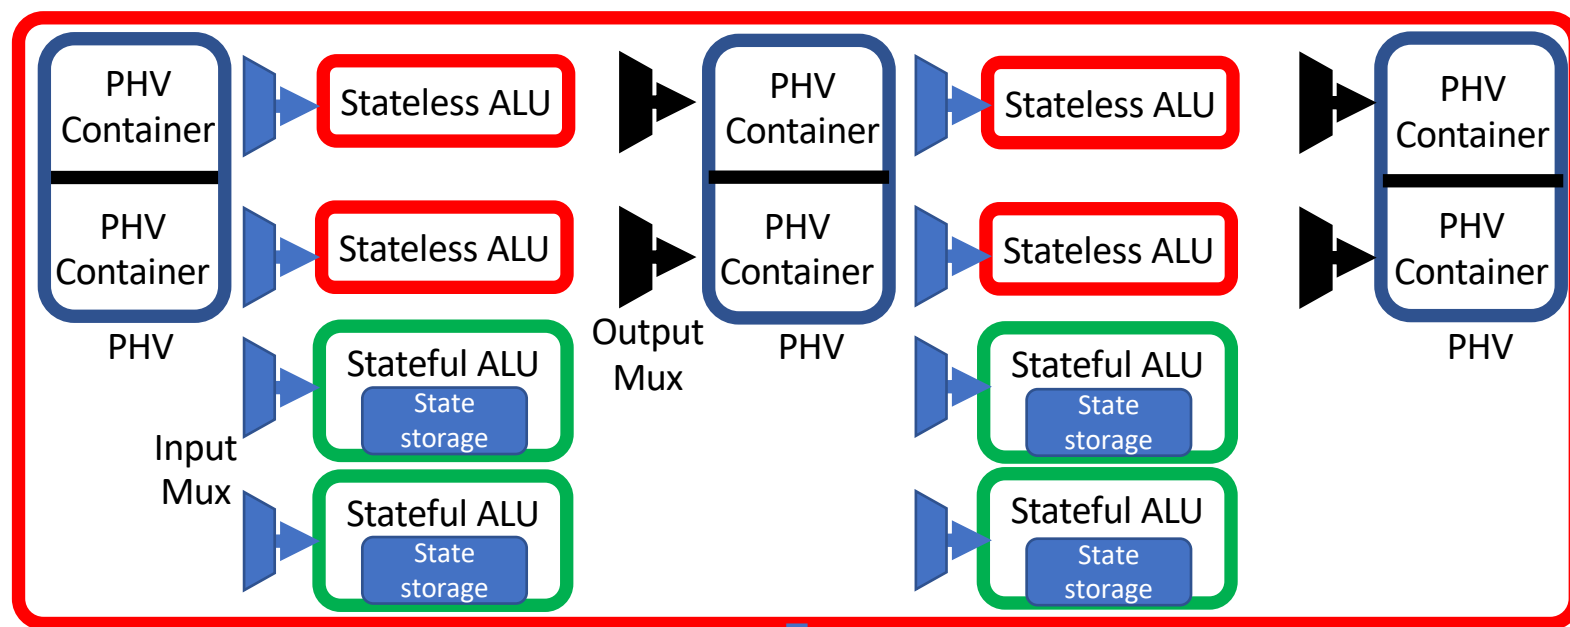
State  
Variable

Packet  
Field

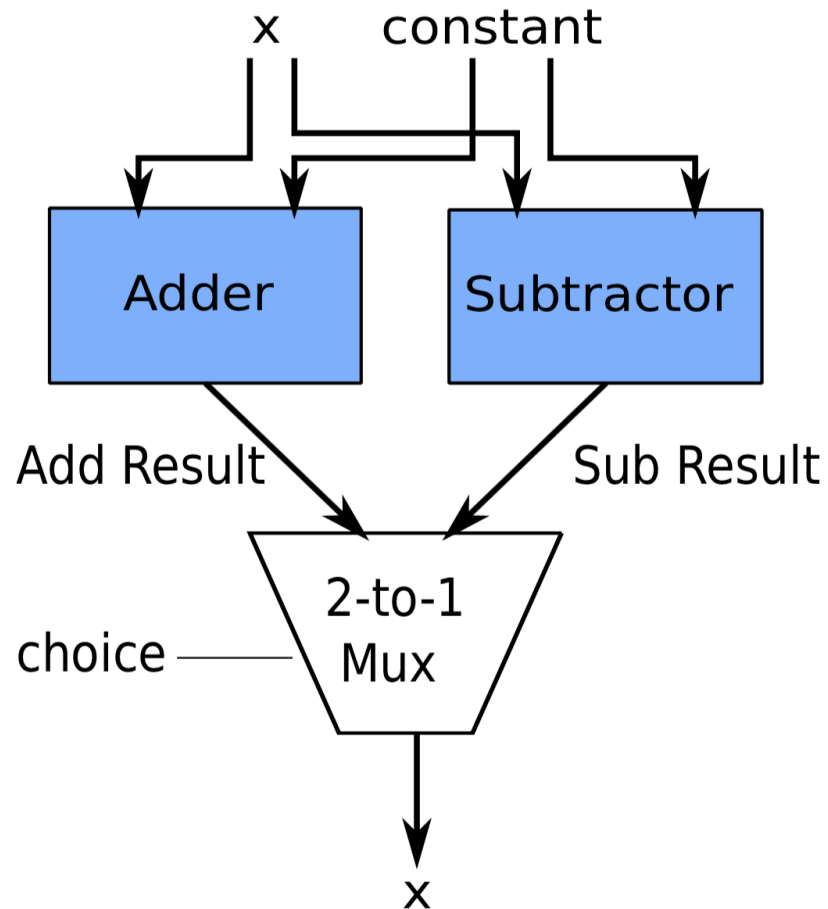
Chipmunk

Machine code for PISA simulator

Partial program for PISA simulator

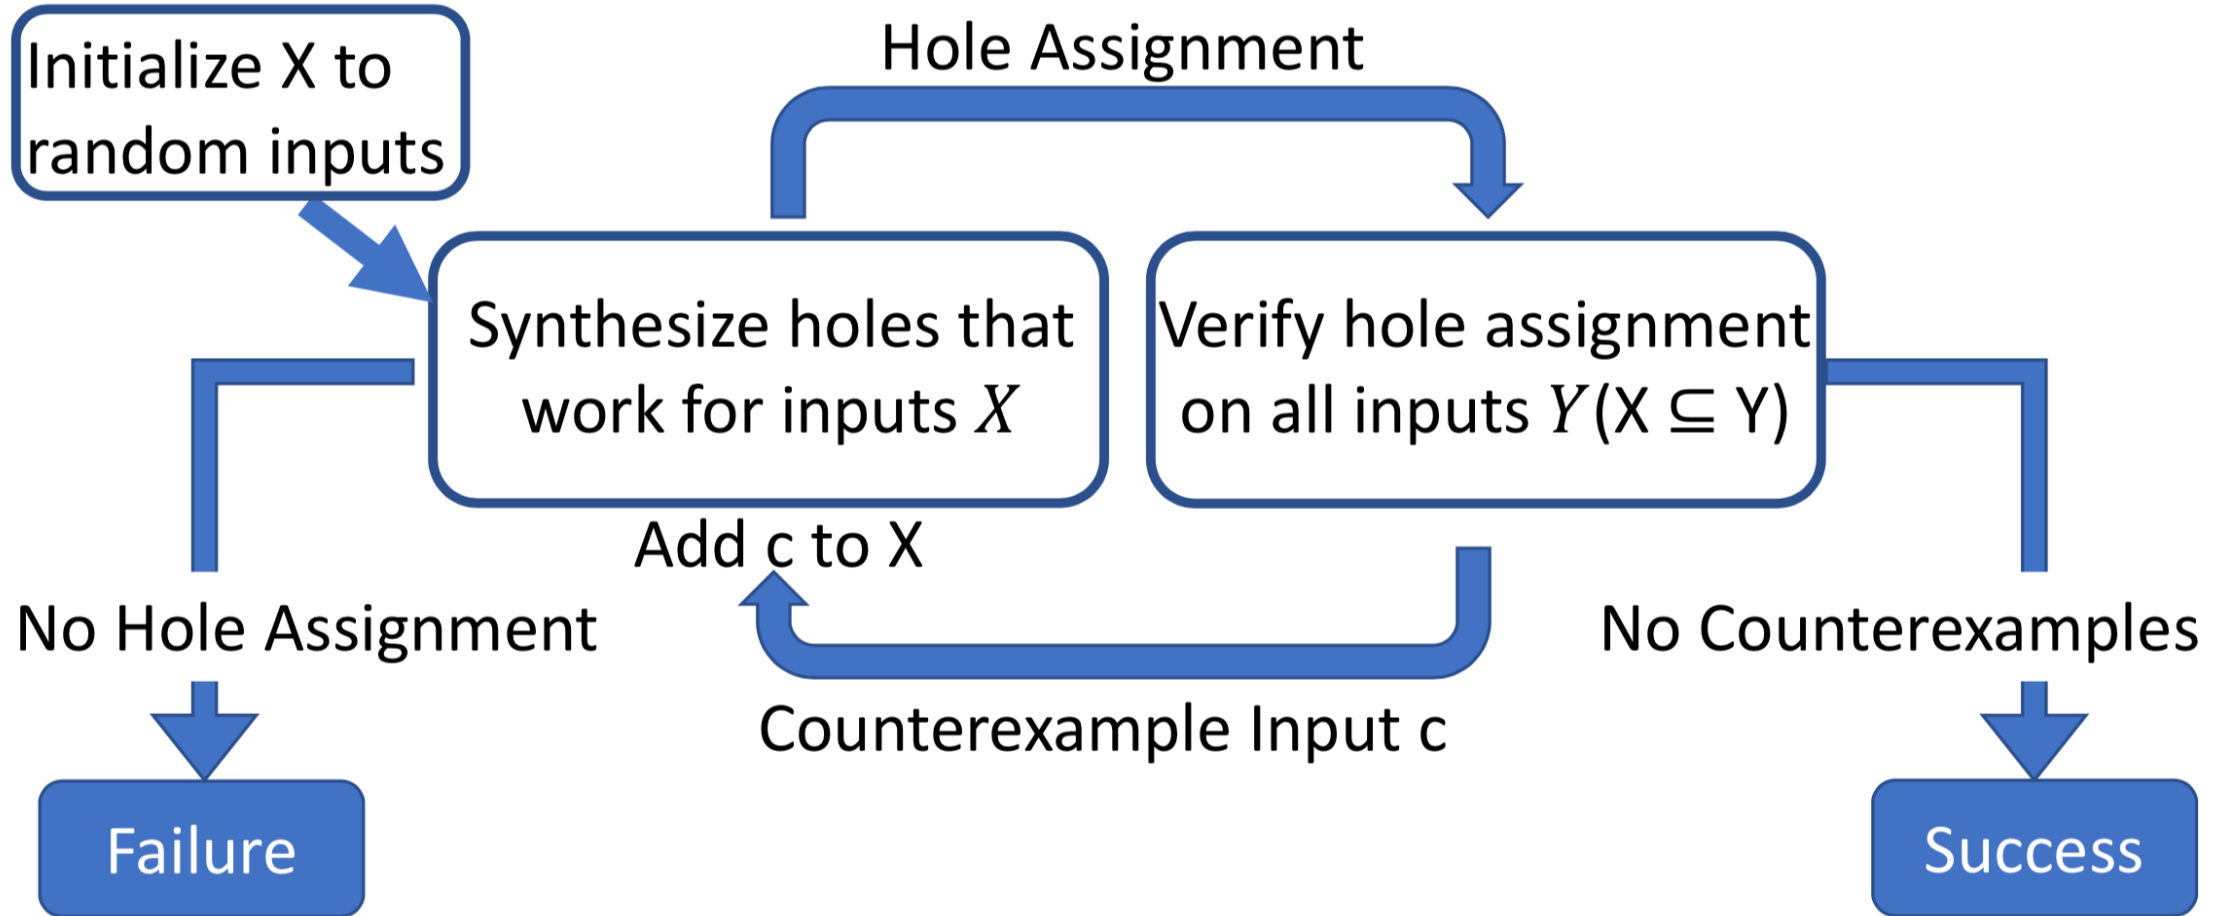


# Partial program representation: Atoms



```
bit choice = ??;  
int constant = ??;  
if (choice) {  
    x = x + constant;  
} else {  
    x = x - constant;  
}
```

# CEGIS: Finding PISA configurations (holes)



Counterexample-Guided Inductive Synthesis

# Other apps of synthesis in networking

- Generating control plane rules with desired properties
- Data plane programs for other targets
  - e.g., Network processors (CPUs)
- Program repair and troubleshooting
- Approximation